

ASPECT

Advanced Solver for Problems in Earth's ConvecTion

User Manual

Version 2.0.1

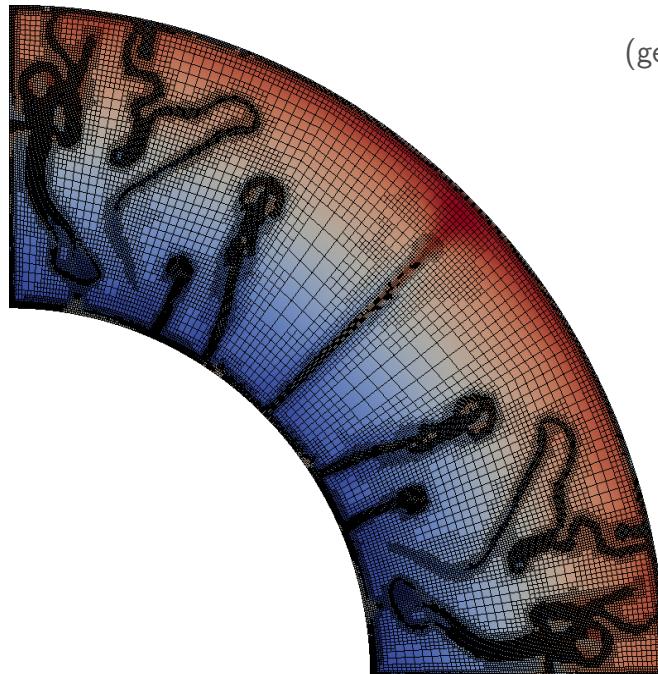
(generated June 22, 2018)

Wolfgang Bangerth

Juliane Dannberg

Rene Gassmöller

Timo Heister



with contributions by:

Jacqueline Austermann, Markus Bürg, Samuel Cox, William Durkin, Grant Euen, Menno Fraters, Thomas Geenen, Anne Glerum, Ryan Grove, Eric Heien, Louise Kellogg, Scott King, Martin Kronbichler, Marine Lasbleis, Shangxin Liu, Elvira Mulyukova, John Naliboff, Jonathan Perry-Houts, Elbridge Gerry Puckett, Tahiry Rajaonarison, Ian Rose, D. Sarah Stamps, Cedric Thieulot, Iris van Zelst, Siqi Zhang

Contents

1	Introduction	8
1.1	Referencing ASPECT	9
1.2	Acknowledgments	9
2	Equations, models, coefficients	9
2.1	Basic equations	9
2.1.1	A comment on adiabatic heating	11
2.1.2	Boundary conditions	11
2.1.3	Two-dimensional models	12
2.1.4	Comments on the final set of equations	13
2.2	Coefficients	13
2.3	Dimensional or non-dimensionalized equations?	15
2.4	Static or dynamic pressure?	16
2.5	Pressure normalization	18
2.6	Initial conditions and the adiabatic pressure/temperature	18
2.7	Compositional fields	19
2.8	Constitutive laws	20
2.9	Numerical methods	21
2.10	Approximate equations	22
2.10.1	The anelastic liquid approximation (ALA)	23
2.10.2	The truncated anelastic liquid approximation (TALA)	23
2.10.3	The Boussinesq approximation (BA)	24
2.10.4	The isothermal compression approximation (ICA)	26
2.11	Choosing a formulation in ASPECT	26
2.11.1	Mass conservation approximation	27
2.11.2	Temperature equation approximation	28
2.11.3	Approximation of the buoyancy term	28
2.11.4	Reference state: The adiabatic profile	28
2.11.5	Combined formulations	28
2.12	Free surface calculations	29
2.12.1	Arbitrary Lagrangian-Eulerian implementation	29
2.12.2	Free surface stabilization	30
2.13	Calculations with melt transport	30
2.14	Nullspace removal	32
2.15	Particles	33
3	Installation	34
3.1	Docker Container	35
3.1.1	Installing Docker and downloading the ASPECT image	35
3.1.2	Running ASPECT models	35
3.1.3	Developing ASPECT within a container	37
3.2	Virtual Machine	37
3.2.1	Installing VM software and setting up the virtual machine	37
3.2.2	Running ASPECT models	37
3.3	Local installation	38
3.3.1	System prerequisites	38
3.3.2	Using candi to compile dependencies	39
3.3.3	Obtaining ASPECT and initial configuration	39
3.3.4	Compiling ASPECT and generating documentation	40

4	Running ASPECT	40
4.1	Overview	40
4.2	Selecting between 2d and 3d runs	44
4.3	Debug or optimized mode	45
4.4	Visualizing results	46
4.4.1	Visualization the graphical output using <i>Visit</i>	47
4.4.2	Visualizing statistical data	49
4.4.3	Large data issues for parallel computations	52
4.5	Checkpoint/restart support	52
4.6	Making ASPECT run faster	53
4.6.1	Debug vs. optimized mode	53
4.6.2	Adjusting solver tolerances	53
4.6.3	Adjusting solver preconditioner tolerances	54
4.6.4	Using lower order elements for the temperature/compositional discretization	54
4.6.5	Limiting postprocessing	54
4.6.6	Switching off pressure normalization	55
4.6.7	Regularizing models with large coefficient variation	55
4.6.8	Using multithreading	55
4.7	Input parameter files	55
4.7.1	The structure of parameter files	56
4.7.2	Categories of parameters	57
4.7.3	A note on the syntax of formulas in input files	58
4.7.4	Compatibility of input files with newer ASPECT versions	58
4.8	A graphical user interface for editing ASPECT parameter files	58
4.8.1	Installing parameter-GUI	59
4.8.2	Using ASPECT-GUI	59
5	Cookbooks	60
5.1	How to set up computations	60
5.2	Simple setups	62
5.2.1	Convection in a 2d box	62
5.2.2	Convection in a 3d box	74
5.2.3	Convection in a box with prescribed, variable velocity boundary conditions	78
5.2.4	Using passive and active compositional fields	82
5.2.5	Using particles	88
5.2.6	Using a free surface	93
5.2.7	Using a free surface in a model with a crust	94
5.2.8	Averaging material properties	96
5.2.9	Prescribed internal velocity constraints	101
5.2.10	Artificial viscosity smoothing	108
5.2.11	Tracking finite strain	110
5.2.12	Reading in compositional initial composition files generated with geomIO	112
5.3	Geophysical setups	115
5.3.1	Simple convection in a quarter of a 2d annulus	117
5.3.2	Simple convection in a spherical 3d shell	123
5.3.3	Postprocessing spherical 3D convection	127
5.3.4	3D convection with an Earth-like initial condition	128
5.3.5	Using reconstructed surface velocities by GPlates	132
5.3.6	2D compressible convection with a reference profile and material properties from Burn-Man	136
5.3.7	Reproducing rheology of Morency and Doin, 2004	141
5.3.8	Crustal deformation	143

5.3.9	Continental extension	148
5.3.10	Inner core convection	152
5.3.11	Melt migration in a 2D mantle convection model	159
5.3.12	Melt migration in a 2D mid-ocean ridge model	169
5.4	Benchmarks	174
5.4.1	Running benchmarks that require code	174
5.4.2	Onset of convection benchmark	175
5.4.3	The van Keken thermochemical composition benchmark	176
5.4.4	The SolCx Stokes benchmark	184
5.4.5	The SolKz Stokes benchmark	188
5.4.6	The “inclusion” Stokes benchmark	190
5.4.7	The Burstedde variable viscosity benchmark	192
5.4.8	The hollow sphere benchmark	194
5.4.9	The 2D annulus benchmark	196
5.4.10	The “Stokes’ law” benchmark	198
5.4.11	Latent heat benchmark	203
5.4.12	The 2D cylindrical shell benchmarks by Davies et al.	208
5.4.13	The Crameri et al. benchmarks	213
5.4.14	The solitary wave benchmark	217
5.4.15	Benchmarks for operator splitting	221
6	Extending and contributing to ASPECT	223
6.1	The idea of plugins and the <code>SimulatorAccess</code> and <code>Introspection</code> classes	226
6.2	How to write a plugin	231
6.3	Materials, geometries, gravitation and other plugin types	232
6.3.1	Material models	232
6.3.2	Heating models	234
6.3.3	Geometry models	235
6.3.4	Gravity models	239
6.3.5	Initial conditions	240
6.3.6	Prescribed velocity boundary conditions	240
6.3.7	Temperature boundary conditions	241
6.3.8	Postprocessors: Evaluating the solution after each time step	242
6.3.9	Visualization postprocessors	245
6.3.10	Mesh refinement criteria	247
6.3.11	Criteria for terminating a simulation	247
6.4	Compatibility of plugins with newer ASPECT versions	248
6.5	Extending ASPECT through the signals mechanism	248
6.6	Extending the basic solver	251
6.7	Testing ASPECT	252
6.7.1	Running tests	252
6.7.2	Writing tests	254
6.8	Contributing to ASPECT’s development	254
7	Future plans for ASPECT	255
8	Finding answers to more questions	255

A Run-time input parameters	256
A.1 Global parameters	256
A.2 Parameters in section Adiabatic conditions model	260
A.3 Parameters in section Adiabatic conditions model/Ascii data model	261
A.4 Parameters in section Adiabatic conditions model/Compute profile	262
A.5 Parameters in section Adiabatic conditions model/Compute profile/Surface condition function	263
A.6 Parameters in section Adiabatic conditions model/Function	264
A.7 Parameters in section Boundary composition model	264
A.8 Parameters in section Boundary composition model/Ascii data model	267
A.9 Parameters in section Boundary composition model/Box	268
A.10 Parameters in section Boundary composition model/Box with lithosphere boundary indicators	269
A.11 Parameters in section Boundary composition model/Function	270
A.12 Parameters in section Boundary composition model/Initial composition	271
A.13 Parameters in section Boundary composition model/Spherical constant	271
A.14 Parameters in section Boundary fluid pressure model	272
A.15 Parameters in section Boundary fluid pressure model/Density	272
A.16 Parameters in section Boundary temperature model	272
A.17 Parameters in section Boundary temperature model/Ascii data model	275
A.18 Parameters in section Boundary temperature model/Box	277
A.19 Parameters in section Boundary temperature model/Box with lithosphere boundary indicators	277
A.20 Parameters in section Boundary temperature model/Constant	278
A.21 Parameters in section Boundary temperature model/Dynamic core	278
A.22 Parameters in section Boundary temperature model/Dynamic core/Geotherm parameters	281
A.23 Parameters in section Boundary temperature model/Dynamic core/Other energy source	282
A.24 Parameters in section Boundary temperature model/Dynamic core/Radioactive heat source	282
A.25 Parameters in section Boundary temperature model/Function	283
A.26 Parameters in section Boundary temperature model/Initial temperature	284
A.27 Parameters in section Boundary temperature model/Spherical constant	284
A.28 Parameters in section Boundary traction model	285
A.29 Parameters in section Boundary traction model/Ascii data model	285
A.30 Parameters in section Boundary traction model/Function	287
A.31 Parameters in section Boundary traction model/Initial lithostatic pressure	288
A.32 Parameters in section Boundary velocity model	288
A.33 Parameters in section Boundary velocity model/Ascii data model	290
A.34 Parameters in section Boundary velocity model/Function	291
A.35 Parameters in section Boundary velocity model/GPlates model	292
A.36 Parameters in section Checkpointing	294
A.37 Parameters in section Compositional fields	294
A.38 Parameters in section Discretization	295
A.39 Parameters in section Discretization/Stabilization parameters	297
A.40 Parameters in section Formulation	299
A.41 Parameters in section Free surface	300
A.42 Parameters in section Geometry model	301
A.43 Parameters in section Geometry model/Box	303
A.44 Parameters in section Geometry model/Box with lithosphere boundary indicators . . .	304
A.45 Parameters in section Geometry model/Chunk	307
A.46 Parameters in section Geometry model/Ellipsoidal chunk	308
A.47 Parameters in section Geometry model/Initial topography model	310
A.48 Parameters in section Geometry model/Initial topography model/Ascii data model . .	310
A.49 Parameters in section Geometry model/Initial topography model/Prm polygon	311

A.50 Parameters in section Geometry model/Sphere	311
A.51 Parameters in section Geometry model/Spherical shell	312
A.52 Parameters in section Gravity model	313
A.53 Parameters in section Gravity model/Ascii data model	313
A.54 Parameters in section Gravity model/Function	314
A.55 Parameters in section Gravity model/Radial constant	315
A.56 Parameters in section Gravity model/Radial linear	315
A.57 Parameters in section Gravity model/Vertical	315
A.58 Parameters in section Heating model	316
A.59 Parameters in section Heating model/Adiabatic heating	317
A.60 Parameters in section Heating model/Adiabatic heating of melt	317
A.61 Parameters in section Heating model/Compositional heating	317
A.62 Parameters in section Heating model/Constant heating	317
A.63 Parameters in section Heating model/Function	318
A.64 Parameters in section Heating model/Latent heat melt	319
A.65 Parameters in section Heating model/Radioactive decay	319
A.66 Parameters in section Initial composition model	320
A.67 Parameters in section Initial composition model/Ascii data model	321
A.68 Parameters in section Initial composition model/Function	322
A.69 Parameters in section Initial temperature model	323
A.70 Parameters in section Initial temperature model/Adiabatic	327
A.71 Parameters in section Initial temperature model/Adiabatic/Function	328
A.72 Parameters in section Initial temperature model/Adiabatic boundary	329
A.73 Parameters in section Initial temperature model/Ascii data model	330
A.74 Parameters in section Initial temperature model/Function	330
A.75 Parameters in section Initial temperature model/Harmonic perturbation	331
A.76 Parameters in section Initial temperature model/Inclusion shape perturbation	332
A.77 Parameters in section Initial temperature model/S4ORTS perturbation	333
A.78 Parameters in section Initial temperature model/SAVANI perturbation	335
A.79 Parameters in section Initial temperature model/Spherical gaussian perturbation	337
A.80 Parameters in section Initial temperature model/Spherical hexagonal perturbation	338
A.81 Parameters in section Material model	338
A.82 Parameters in section Material model/Ascii reference profile	346
A.83 Parameters in section Material model/Ascii reference profile/Ascii data model	347
A.84 Parameters in section Material model/Averaging	347
A.85 Parameters in section Material model/Compositing	348
A.86 Parameters in section Material model/Composition reaction model	350
A.87 Parameters in section Material model/Depth dependent model	352
A.88 Parameters in section Material model/Depth dependent model/Viscosity depth function	353
A.89 Parameters in section Material model/Diffusion dislocation	354
A.90 Parameters in section Material model/Drucker Prager	358
A.91 Parameters in section Material model/Drucker Prager/Viscosity	359
A.92 Parameters in section Material model/Dynamic Friction	359
A.93 Parameters in section Material model/Dynamic Friction/Viscosities	361
A.94 Parameters in section Material model/Grain size model	362
A.95 Parameters in section Material model/Latent heat	370
A.96 Parameters in section Material model/Latent heat melt	374
A.97 Parameters in section Material model/Melt global	379
A.98 Parameters in section Material model/Melt simple	382
A.99 Parameters in section Material model/Multicomponent	389
A.100 Parameters in section Material model/Nondimensional model	390

A.101	Parameters in section Material model/Simple compressible model	391
A.102	Parameters in section Material model/Simple model	392
A.103	Parameters in section Material model/Simpler model	394
A.104	Parameters in section Material model/Steinberger model	395
A.105	Parameters in section Material model/Visco Plastic	397
A.106	Parameters in section Melt settings	403
A.107	Parameters in section Mesh refinement	404
A.108	Parameters in section Mesh refinement/Artificial viscosity	410
A.109	Parameters in section Mesh refinement/Boundary	410
A.110	Parameters in section Mesh refinement/Compaction length	410
A.111	Parameters in section Mesh refinement/Composition	411
A.112	Parameters in section Mesh refinement/Composition approximate gradient	411
A.113	Parameters in section Mesh refinement/Composition gradient	411
A.114	Parameters in section Mesh refinement/Composition threshold	412
A.115	Parameters in section Mesh refinement/Maximum refinement function	412
A.116	Parameters in section Mesh refinement/Minimum refinement function	413
A.117	Parameters in section Nullspace removal	414
A.118	Parameters in section Postprocess	415
A.119	Parameters in section Postprocess/Command	418
A.120	Parameters in section Postprocess/Depth average	419
A.121	Parameters in section Postprocess/Dynamic core statistics	420
A.122	Parameters in section Postprocess/Dynamic topography	420
A.123	Parameters in section Postprocess/Geoid	421
A.124	Parameters in section Postprocess/Global statistics	422
A.125	Parameters in section Postprocess/Particles	422
A.126	Parameters in section Postprocess/Particles/Function	427
A.127	Parameters in section Postprocess/Particles/Generator	428
A.128	Parameters in section Postprocess/Particles/Generator/Ascii file	428
A.129	Parameters in section Postprocess/Particles/Generator/Probability density function	428
A.130	Parameters in section Postprocess/Particles/Generator/Reference cell	430
A.131	Parameters in section Postprocess/Particles/Generator/Uniform box	430
A.132	Parameters in section Postprocess/Particles/Generator/Uniform radial	431
A.133	Parameters in section Postprocess/Particles/Interpolator	432
A.134	Parameters in section Postprocess/Particles/Interpolator/Bilinear least squares	432
A.135	Parameters in section Postprocess/Particles/Melt particle	433
A.136	Parameters in section Postprocess/Point values	433
A.137	Parameters in section Postprocess/Topography	433
A.138	Parameters in section Postprocess/Visualization	434
A.139	Parameters in section Postprocess/Visualization/Artificial viscosity composition	441
A.140	Parameters in section Postprocess/Visualization/Compositional fields as vectors	441
A.141	Parameters in section Postprocess/Visualization/Material properties	441
A.142	Parameters in section Postprocess/Visualization/Melt fraction	442
A.143	Parameters in section Postprocess/Visualization/Melt material properties	445
A.144	Parameters in section Postprocess/Visualization/Vp anomaly	445
A.145	Parameters in section Postprocess/Visualization/Vs anomaly	445
A.146	Parameters in section Prescribed Stokes solution	446
A.147	Parameters in section Prescribed Stokes solution/Ascii data model	446
A.148	Parameters in section Prescribed Stokes solution/Compaction pressure function	447
A.149	Parameters in section Prescribed Stokes solution/Fluid pressure function	448
A.150	Parameters in section Prescribed Stokes solution/Fluid velocity function	448
A.151	Parameters in section Prescribed Stokes solution/Pressure function	449

A.152Parameters in section Prescribed Stokes solution/Velocity function	450
A.153Parameters in section Solver parameters	451
A.154Parameters in section Solver parameters/AMG parameters	451
A.155Parameters in section Solver parameters/Newton solver parameters	452
A.156Parameters in section Solver parameters/Operator splitting parameters	454
A.157Parameters in section Solver parameters/Stokes solver parameters	455
A.158Parameters in section Termination criteria	456
A.159Parameters in section Termination criteria/Steady state velocity	457
A.160Parameters in section Termination criteria/User request	457
References	459
Index of run-time parameter entries	463
Index of run-time parameters with section names	470

1 Introduction

ASPECT — short for Advanced Solver for Problems in Earth’s ConvecTion — is a code intended to solve the equations that describe thermally driven convection with a focus on doing so in the context of convection in the earth mantle. It is primarily developed by computational scientists at Texas A&M University based on the following principles:

- *Usability and extensibility:* Simulating mantle convection is a difficult problem characterized not only by complicated and nonlinear material models but, more generally, by a lack of understanding which parts of a much more complicated model are really necessary to simulate the defining features of the problem. To name just a few examples:
 - Mantle convection is often solved in a spherical shell geometry, but the earth is not a sphere – its true shape on the longest length scales is dominated by polar oblateness, but deviations from spherical shape relevant to convection patterns may go down to the length scales of mountain belts, mid-ocean ridges or subduction trenches. Furthermore, processes outside the mantle like crustal depression during glaciations can change the geometry as well.
 - Rocks in the mantle flow on long time scales, but on shorter time scales they behave more like a visco-elasto-plastic material as they break and as their crystalline structure heals again. The mathematical models discussed in Section 2 can therefore only be approximations.
 - If pressures are low and temperatures high enough, rocks melt, leading to all sorts of new and interesting behavior.

This uncertainty in what problem one actually wants to solve requires a code that is easy to extend by users to support the community in determining what the essential features of convection in the earth mantle are. Achieving this goal also opens up possibilities outside the original scope, such as the simulation of convection in exoplanets or the icy satellites of the gas giant planets in our solar system.

- *Modern numerical methods:* We build ASPECT on numerical methods that are at the forefront of research in all areas – adaptive mesh refinement, linear and nonlinear solvers, stabilization of transport-dominated processes. This implies complexity in our algorithms, but also guarantees highly accurate solutions while remaining efficient in the number of unknowns and with CPU and memory resources.
- *Parallelism:* Many convection processes of interest are characterized by small features in large domains – for example, mantle plumes of a few tens of kilometers diameter in a mantle almost 3,000 km deep. Such problems can not be solved on a single computer but require dozens or hundreds of processors to work together. ASPECT is designed from the start to support this level of parallelism.
- *Building on others’ work:* Building a code that satisfies above criteria from scratch would likely require several 100,000 lines of code. This is outside what any one group can achieve on academic time scales. Fortunately, most of the functionality we need is already available in the form of widely used, actively maintained, and well tested and documented libraries, and we leverage these to make ASPECT a much smaller and easier to understand system. Specifically, ASPECT builds immediately on top of the DEAL.II library (see <https://www.dealii.org/>) for everything that has to do with finite elements, geometries, meshes, etc.; and, through DEAL.II on Trilinos (see <http://trilinos.org/>) for parallel linear algebra and on P4EST (see <http://www.p4est.org/>) for parallel mesh handling.
- *Community:* We believe that a large project like ASPECT can only be successful as a community project. Every contribution is welcome and we want to help you so we can improve ASPECT together.

Combining all of these aspects into one code makes for an interesting challenge. We hope to have achieved our goal of providing a useful tool to the geodynamics community and beyond!

Note: ASPECT is a community project. As such, we encourage contributions from the community to improve this code over time. Natural candidates for such contributions are implementations of new plugins as discussed in Section 6.3 since they are typically self-contained and do not require much knowledge of the details of the remaining code. Obviously, however, we also encourage contributions to the core functionality in any form! If you have something that might be of general interest, please contact us.

Note: ASPECT will only solve problems relevant to the community if we get feedback from the community on things that are missing or necessary for what you want to do. Let us know by personal email to the developers, or the mantle convection or aspect-devel mailing lists hosted at <http://lists.geodynamics.org/cgi-bin/mailman/listinfo/aspect-devel>!

1.1 Referencing ASPECT

As with all scientific work, funding agencies have a reasonable expectation that if we ask for continued funding for this work, we need to demonstrate relevance. In addition, many have contributed to the development of ASPECT and deserve credit for their work. To this end, we ask that if you publish results that were obtained to some part using ASPECT, please see the information at <https://aspect.geodynamics.org/cite.html>, which includes suggestions for acknowledgments and citations.

1.2 Acknowledgments

The development of ASPECT has been funded through a variety of grants to the authors. Most immediately, it has been supported through the Computational Infrastructure in Geodynamics grant, initially by the CIG-I grant (National Science Foundation Award No. EAR-0426271, via The California Institute of Technology) and later by the CIG-II and CIG-III grants (National Science Foundation Awards No. EAR-0949446 and EAR-1550901, via The University of California – Davis). In addition, the libraries upon which ASPECT builds heavily have been supported through many other grants that are equally gratefully acknowledged.

Please acknowledge CIG as follows:

ASPECT is hosted by the Computational Infrastructure for Geodynamics (CIG) which is supported by the National Science Foundation award EAR-1550901.

2 Equations, models, coefficients

2.1 Basic equations

ASPECT solves a system of equations in a $d = 2$ - or $d = 3$ -dimensional domain Ω that describes the motion of a highly viscous fluid driven by differences in the gravitational force due to a density that depends on the temperature. In the following, we largely follow the exposition of this material in Schubert, Turcotte and Olson [STO01].

Specifically, we consider the following set of equations for velocity \mathbf{u} , pressure p and temperature T , as

well as a set of advected quantities c_i that we call *compositional fields*:

$$-\nabla \cdot \left[2\eta \left(\varepsilon(\mathbf{u}) - \frac{1}{3}(\nabla \cdot \mathbf{u})\mathbf{1} \right) \right] + \nabla p = \rho g \quad \text{in } \Omega, \quad (1)$$

$$\nabla \cdot (\rho \mathbf{u}) = 0 \quad \text{in } \Omega, \quad (2)$$

$$\begin{aligned} \rho C_p \left(\frac{\partial T}{\partial t} + \mathbf{u} \cdot \nabla T \right) - \nabla \cdot k \nabla T &= \rho H \\ &\quad + 2\eta \left(\varepsilon(\mathbf{u}) - \frac{1}{3}(\nabla \cdot \mathbf{u})\mathbf{1} \right) : \left(\varepsilon(\mathbf{u}) - \frac{1}{3}(\nabla \cdot \mathbf{u})\mathbf{1} \right) \\ &\quad + \alpha T (\mathbf{u} \cdot \nabla p) \\ &\quad + \rho T \Delta S \left(\frac{\partial X}{\partial t} + \mathbf{u} \cdot \nabla X \right) \end{aligned} \quad \text{in } \Omega, \quad (3)$$

$$\frac{\partial c_i}{\partial t} + \mathbf{u} \cdot \nabla c_i = q_i \quad \text{in } \Omega, i = 1 \dots C \quad (4)$$

where $\varepsilon(\mathbf{u}) = \frac{1}{2}(\nabla \mathbf{u} + \nabla \mathbf{u}^T)$ is the symmetric gradient of the velocity (often called the *strain rate*).¹

In this set of equations, (1) and (2) represent the compressible Stokes equations in which $\mathbf{u} = \mathbf{u}(\mathbf{x}, t)$ is the velocity field and $p = p(\mathbf{x}, t)$ the pressure field. Both fields depend on space \mathbf{x} and time t . Fluid flow is driven by the gravity force that acts on the fluid and that is proportional to both the density of the fluid and the strength of the gravitational pull.

Coupled to this Stokes system is equation (3) for the temperature field $T = T(\mathbf{x}, t)$ that contains heat conduction terms as well as advection with the flow velocity \mathbf{u} . The right hand side terms of this equation correspond to

- internal heat production for example due to radioactive decay;
- friction heating;
- adiabatic compression of material;
- phase change.

The last term of the temperature equation corresponds to the latent heat generated or consumed in the process of phase change of material. The latent heat release is proportional to changes in the fraction of material X that has already undergone the phase transition (also called phase function) and the change of entropy ΔS . This process applies both to solid-state phase transitions and to melting/solidification. Here, ΔS is positive for exothermic phase transitions. As the phase of the material, for a given composition, depends on the temperature and pressure, the latent heat term can be reformulated:

$$\frac{\partial X}{\partial t} + \mathbf{u} \cdot \nabla X = \frac{DX}{Dt} = \frac{\partial X}{\partial T} \frac{DT}{Dt} + \frac{\partial X}{\partial p} \frac{Dp}{Dt} = \frac{\partial X}{\partial T} \left(\frac{\partial T}{\partial t} + \mathbf{u} \cdot \nabla T \right) + \frac{\partial X}{\partial p} \mathbf{u} \cdot \nabla p.$$

The last transformation results from the assumption that the flow field is always in equilibrium and consequently $\partial p / \partial t = 0$ (this is the same assumption that underlies the fact that equation (1) does not have a term

¹There is no consensus in the sciences on the notation used for strain and strain rate. The symbols ε , $\dot{\varepsilon}$, $\varepsilon(\mathbf{u})$, and $\dot{\varepsilon}(\mathbf{u})$, can all be found. In this manual, and in the code, we will consistently use ε as an *operator*, i.e., the symbol is not used on its own but only as applied to a field. In other words, if \mathbf{u} is the velocity field, then $\varepsilon(\mathbf{u}) = \frac{1}{2}(\nabla \mathbf{u} + \nabla \mathbf{u}^T)$ will denote the strain rate. On the other hand, if \mathbf{d} is the displacement field, then $\varepsilon(\mathbf{d}) = \frac{1}{2}(\nabla \mathbf{d} + \nabla \mathbf{d}^T)$ will denote the strain.

$\partial\mathbf{u}/\partial t$). With this reformulation, we can rewrite (3) in the following way in which it is in fact implemented:

$$\begin{aligned} \left(\rho C_p - \rho T \Delta S \frac{\partial X}{\partial T} \right) \left(\frac{\partial T}{\partial t} + \mathbf{u} \cdot \nabla T \right) - \nabla \cdot k \nabla T &= \rho H \\ &\quad + 2\eta \left(\varepsilon(\mathbf{u}) - \frac{1}{3} (\nabla \cdot \mathbf{u}) \mathbf{1} \right) : \left(\varepsilon(\mathbf{u}) - \frac{1}{3} (\nabla \cdot \mathbf{u}) \mathbf{1} \right) \\ &\quad + \alpha T (\mathbf{u} \cdot \nabla p) \\ &\quad + \rho T \Delta S \frac{\partial X}{\partial p} \mathbf{u} \cdot \nabla p \end{aligned} \quad \text{in } \Omega. \quad (5)$$

The last of the equations above, equation (4), describes the evolution of additional fields that are transported along with the velocity field \mathbf{u} and may react with each other and react to other features of the solution, but that do not diffuse. We call these fields c_i *compositional fields*, although they can also be used for other purposes than just tracking chemical compositions. We will discuss this equation in more detail in Section 2.7.

2.1.1 A comment on adiabatic heating

Other codes and texts sometimes make a simplification to the adiabatic heating term in the previous equation. If you assume the vertical component of the gradient of the *dynamic* pressure to be small compared to the gradient of the *total* pressure (in other words, the gradient is dominated by the gradient of the hydrostatic pressure), then $-\rho\mathbf{g} \approx \nabla\mathbf{p}$, and we have the following relation (the negative sign is due to \mathbf{g} pointing downwards)

$$\alpha T (\mathbf{u} \cdot \nabla \mathbf{p}) \approx -\alpha \rho T \mathbf{u} \cdot \mathbf{g}.$$

While this simplification is possible, it is not necessary if you have access to the total pressure. ASPECT therefore by default implements the original term without this simplification, but allows to simplify this term by setting the “Use simplified adiabatic heating” parameter in section A.59.

2.1.2 Boundary conditions

Having discussed (3), let us come to the last one of the original set of equations, (4). It describes the motion of a set of advected quantities $c_i(\mathbf{x}, t), i = 1 \dots C$. We call these *compositional fields* because we think of them as spatially and temporally varying concentrations of different elements, minerals, or other constituents of the composition of the material that convects. As such, these fields participate actively in determining the values of the various coefficients of these equations. On the other hand, ASPECT also allows the definition of material models that are independent of these compositional fields, making them passively advected quantities. Several of the cookbooks in Section 5 consider compositional fields in this way, i.e., essentially as tracer quantities that only keep track of where material came from.

These equations are augmented by boundary conditions that can either be of Dirichlet, Neumann, or tangential type on subsets of the boundary $\Gamma = \partial\Omega$:

$$\mathbf{u} = 0 \quad \text{on } \Gamma_{0,\mathbf{u}}, \quad (6)$$

$$\mathbf{u} = \mathbf{u}_{\text{prescribed}} \quad \text{on } \Gamma_{\text{prescribed},\mathbf{u}}, \quad (7)$$

$$\mathbf{n} \cdot \mathbf{u} = 0 \quad \text{on } \Gamma_{\parallel,\mathbf{u}}, \quad (8)$$

$$(2\eta\varepsilon(\mathbf{u}) - pI)\mathbf{n} = \mathbf{t} \quad \text{on } \Gamma_{\text{traction},\mathbf{u}}, \quad (9)$$

$$T = T_{\text{prescribed}} \quad \text{on } \Gamma_{D,T}, \quad (10)$$

$$\mathbf{n} \cdot k \nabla T = 0 \quad \text{on } \Gamma_{N,T}. \quad (11)$$

$$c_i = c_{i,\text{prescribed}} \quad \text{on } \Gamma_{\text{in}} = \{\mathbf{x} : \mathbf{u} \cdot \mathbf{n} < 0\}. \quad (12)$$

Here, the boundary conditions for velocity and temperature are subdivided into disjoint parts:

- $\Gamma_{0,\mathbf{u}}$ corresponds to parts of the boundary on which the velocity is fixed to be zero.
- $\Gamma_{\text{prescribed},\mathbf{u}}$ corresponds to parts of the boundary on which the velocity is prescribed to some value (which could also be zero). It is possible to restrict prescribing the velocity to only certain components of the velocity vector.
- $\Gamma_{\parallel,\mathbf{u}}$ corresponds to parts of the boundary on which the velocity may be nonzero but must be parallel to the boundary, with the tangential component undetermined.
- $\Gamma_{\text{traction},\mathbf{u}}$ corresponds to parts of the boundary on which the traction is prescribed to some surface force density (a common application being $\mathbf{t} = -p\mathbf{n}$ if one just wants to prescribe a pressure component). It is possible to restrict prescribing the traction to only certain vector components.
- $\Gamma_{D,T}$ corresponds to places where the temperature is prescribed (for example at the inner and outer boundaries of the earth mantle).
- $\Gamma_{N,T}$ corresponds to places where the temperature is unknown but the heat flux across the boundary is zero (for example on symmetry surfaces if only a part of the shell that constitutes the domain the Earth mantle occupies is simulated).

We require that one of these boundary conditions hold at each point for both velocity and temperature, i.e., $\Gamma_{0,\mathbf{u}} \cup \Gamma_{\text{prescribed},\mathbf{u}} \cup \Gamma_{\parallel,\mathbf{u}} \cup \Gamma_{\text{traction},\mathbf{u}} = \Gamma$ and $\Gamma_{D,T} \cup \Gamma_{N,T} = \Gamma$.

Boundary conditions have to be imposed for the compositional fields only at those parts of the boundary where flow points inward, see equation (12), but not where it is either tangential to the boundary or points outward. The difference in treatment between temperature and compositional boundary conditions is due to the fact that the temperature equation contains a (possibly small) diffusion component, whereas the compositional equations do not.

There are other equations that ASPECT can optionally solve. For example, it can deal with free surfaces (see Section 2.12), melt generation and transport (see Section 2.13), and it can advect along particles (see Section 2.15). These optional models are discussed in more detail in the indicated sections.

2.1.3 Two-dimensional models

ASPECT allows solving both two- and three-dimensional models via a parameter in the input files, see also Section 4.2. At the same time, the world is unambiguously three-dimensional. This raises the question what exactly we mean when we say that we want to solve two-dimensional problems.

The notion we adopt here – in agreement with that chosen by many other codes – is to think of two-dimensional models in the following way: We assume that the domain we want to solve on is a two-dimensional cross section (parameterized by x and y coordinates) that extends infinitely far in both negative and positive z direction. Further, we assume that the velocity is zero in z direction and that all variables have no variation in z direction. As a consequence, we ought to really think of these two-dimensional models as three-dimensional ones in which the z component of the velocity is zero and so are all z derivatives.

If one adopts this point of view, the Stokes equations (1)–(2) naturally simplify in a way that allows us to reduce the $3 + 1$ equations to only $2 + 1$, but it makes clear that the correct description of the compressible strain rate is still $\varepsilon(\mathbf{u}) - \frac{1}{3}(\nabla \cdot \mathbf{u})\mathbf{1}$, rather than using a factor of $\frac{1}{2}$ for the second term. (A derivation of why the compressible strain rate tensor has this form can be found in [STO01, Section 6.5].)

It is interesting to realize that this compressible strain rate indeed requires a 3×3 tensor: While under the assumptions above we have

$$\varepsilon(\mathbf{u}) = \begin{pmatrix} \frac{\partial u_x}{\partial x} & \frac{1}{2} \frac{\partial u_x}{\partial y} + \frac{1}{2} \frac{\partial u_y}{\partial x} & 0 \\ \frac{1}{2} \frac{\partial u_x}{\partial y} + \frac{1}{2} \frac{\partial u_y}{\partial x} & \frac{\partial u_y}{\partial y} & 0 \\ 0 & 0 & 0 \end{pmatrix}$$

with the expected zeros in the last row and column, the full compressible strain rate tensor reads

$$\varepsilon(\mathbf{u}) - \frac{1}{3}(\nabla \cdot \mathbf{u})\mathbf{1} = \begin{pmatrix} \frac{2}{3}\frac{\partial u_x}{\partial x} - \frac{1}{3}\frac{\partial u_y}{\partial y} & \frac{1}{2}\frac{\partial u_x}{\partial y} + \frac{1}{2}\frac{\partial u_y}{\partial x} & 0 \\ \frac{1}{2}\frac{\partial u_x}{\partial y} + \frac{1}{2}\frac{\partial u_y}{\partial x} & \frac{2}{3}\frac{\partial u_y}{\partial y} - \frac{1}{3}\frac{\partial u_x}{\partial x} & 0 \\ 0 & 0 & -\frac{1}{3}\frac{\partial u_y}{\partial y} - \frac{1}{3}\frac{\partial u_x}{\partial x} \end{pmatrix}.$$

The entry in the (3, 3) position of this tensor may be surprising. It disappears, however, when taking the (three-dimensional) divergence of the stress, as is done in (1), because the divergence applies the z derivative to all elements of the last row – and the assumption above was that all z derivatives are zero; consequently whatever lives in the third row of the strain rate tensor does not matter.

2.1.4 Comments on the final set of equations

ASPECT solves these equations in essentially the form stated. In particular, the form given in (1) implies that the pressure p we compute is in fact the *total pressure*, i.e., the sum of hydrostatic pressure and dynamic pressure (however, see Section 2.4 for more information on this, as well as the extensive discussion of this issue in [KHB12]). Consequently, it allows the direct use of this pressure when looking up pressure dependent material parameters.

2.2 Coefficients

The equations above contain a significant number of coefficients that we will discuss in the following. In the most general form, many of these coefficients depend nonlinearly on the solution variables pressure p , temperature T and, in the case of the viscosity, on the strain rate $\varepsilon(\mathbf{u})$. If compositional fields $\mathbf{c} = \{c_1, \dots, c_C\}$ are present (i.e., if $C > 0$), coefficients may also depend on them. Alternatively, they may be parameterized as a function of the spatial variable \mathbf{x} . ASPECT allows both kinds of parameterizations.

Note that below we will discuss examples of the dependence of coefficients on other quantities; which dependence is actually implemented in the code is a different matter. As we will discuss in Sections A and 6, some versions of these models are already implemented and can be selected from the input parameter file; others are easy to add to ASPECT by providing self-contained descriptions of a set of coefficients that the rest of the code can then use without a need for further modifications.

Concretely, we consider the following coefficients and dependencies:

- *The viscosity* $\eta = \eta(p, T, \varepsilon(\mathbf{u}), \mathbf{c}, \mathbf{x})$: Units $\text{Pa} \cdot \text{s} = \text{kg} \frac{1}{\text{m} \cdot \text{s}}$.

The viscosity is the proportionality factor that relates total forces (external gravity minus pressure gradients) and fluid velocities \mathbf{u} . The simplest models assume that η is constant, with the constant often chosen to be on the order of 10^{21} Pa s .

More complex (and more realistic) models assume that the viscosity depends on pressure, temperature and strain rate. Since this dependence is often difficult to quantify, one modeling approach is to make η spatially dependent.

- *The density* $\rho = \rho(p, T, \mathbf{c}, \mathbf{x})$: Units $\frac{\text{kg}}{\text{m}^3}$.

In general, the density depends on pressure and temperature, both through pressure compression, thermal expansion, and phase changes the material may undergo as it moves through the pressure-temperature phase diagram.

The simplest parameterization for the density is to assume a linear dependence on temperature, yielding the form $\rho(T) = \rho_{\text{ref}}[1 - \alpha(T - T_{\text{ref}})]$ where ρ_{ref} is the reference density at temperature T_{ref} and α is the linear thermal expansion coefficient. For the earth's mantle, typical values for this parameterization would be $\rho_{\text{ref}} = 3300 \frac{\text{kg}}{\text{m}^3}$, $T_{\text{ref}} = 293 \text{ K}$, $\alpha = 2 \cdot 10^{-5} \frac{1}{\text{K}}$.

- The gravity vector $\mathbf{g} = \mathbf{g}(\mathbf{x})$: Units $\frac{\text{m}}{\text{s}^2}$.

Simple models assume a radially inward gravity vector of constant magnitude (e.g., the surface gravity of Earth, $9.81 \frac{\text{m}}{\text{s}^2}$), or one that can be computed analytically assuming a homogeneous mantle density.

A physically self-consistent model would compute the gravity vector as $\mathbf{g} = -\nabla\varphi$ with a gravity potential φ that satisfies $-\Delta\varphi = 4\pi G\rho$ with the density ρ from above and G the universal constant of gravity. This would provide a gravity vector that changes as a function of time. Such a model is not currently implemented.

- The specific heat capacity $C_p = C_p(p, T, \mathbf{c}, \mathbf{x})$: Units $\frac{\text{J}}{\text{kg}\cdot\text{K}} = \frac{\text{m}^2}{\text{s}^2\cdot\text{K}}$.

The specific heat capacity denotes the amount of energy needed to increase the temperature of one kilogram of material by one degree. Wikipedia lists a value of $790 \frac{\text{J}}{\text{kg}\cdot\text{K}}$ for granite². For the earth mantle, a value of $1250 \frac{\text{J}}{\text{kg}\cdot\text{K}}$ is within the range suggested by the literature.

- The thermal conductivity $k = k(p, T, \mathbf{c}, \mathbf{x})$: Units $\frac{\text{W}}{\text{m}\cdot\text{K}} = \frac{\text{kg}\cdot\text{m}}{\text{s}^3\cdot\text{K}}$.

The thermal conductivity denotes the amount of thermal energy flowing through a unit area for a given temperature gradient. It depends on the material and as such will from a physical perspective depend on pressure and temperature due to phase changes of the material as well as through different mechanisms for heat transport (see, for example, the partial transparency of perovskite, the most abundant material in the earth mantle, at pressures above around 120 GPa [BRV⁺⁰⁴]).

As a rule of thumb for its order of magnitude, Wikipedia quotes values of $1.83\text{--}2.90 \frac{\text{W}}{\text{m}\cdot\text{K}}$ for sandstone and $1.73\text{--}3.98 \frac{\text{W}}{\text{m}\cdot\text{K}}$ for granite.³ The values in the mantle are almost certainly higher than this though probably not by much. The exact value is not really all that important: heat transport through convection is several orders of magnitude more important than through thermal conduction.

The thermal conductivity k is often expressed in terms of the *thermal diffusivity* κ using the relation $k = \rho C_p \kappa$.

- The intrinsic specific heat production $H = H(\mathbf{x})$: Units $\frac{\text{W}}{\text{kg}} = \frac{\text{m}^2}{\text{s}^3}$.

This term denotes the intrinsic heating of the material, for example due to the decay of radioactive material. As such, it depends not on pressure or temperature, but may depend on the location due to different chemical composition of material in the earth mantle. The literature suggests a value of $\gamma = 7.4 \cdot 10^{-12} \frac{\text{W}}{\text{kg}}$.

- The thermal expansion coefficient $\alpha = \alpha(p, T, \mathbf{c}, \mathbf{x})$: Units $\frac{1}{\text{K}}$.

This term denotes by how much the material under consideration expands due to temperature increases. This coefficient is defined as $\alpha = -\frac{1}{\rho} \frac{\partial \rho}{\partial T}$, where the negative sign is due the fact that the density *decreases* as a function of temperature. Alternatively, if one considers the *volume* $V = V(T)$ a piece of material of mass M occupies, $V = \frac{M}{\rho}$, then the thermal expansion coefficient is defined as the relative increase in volume, $\alpha = \frac{1}{V} \frac{\partial V(T)}{\partial T}$, because $\frac{\partial V(T)}{\partial T} = \frac{\partial \frac{M}{\rho}}{\partial T} = -\frac{M}{\rho^2} \frac{\partial \rho}{\partial T} = -\frac{V}{\rho} \frac{\partial \rho}{\partial T}$.

The literature suggests that values of $\alpha = 1 \cdot 10^{-5} \frac{1}{\text{K}}$ at the core-mantle boundary and $\alpha = 4 \cdot 10^{-5} \frac{1}{\text{K}}$ are appropriate for Earth.

- The change of entropy ΔS at a phase transition together with the derivatives of the phase function $X = X(p, T, \mathbf{c}, \mathbf{x})$ with regard to temperature and pressure: Units $\frac{\text{J}}{\text{kg}\text{K}^2}$ ($-\Delta S \frac{\partial X}{\partial T}$) and $\frac{\text{m}^3}{\text{kg}\text{K}}$ ($\Delta S \frac{\partial X}{\partial p}$).

When material undergoes a phase transition, the entropy changes due to release or consumption of latent heat. However, phase transitions occur gradually and for a given chemical composition it depends

²See http://en.wikipedia.org/wiki/Specific_heat.

³See http://en.wikipedia.org/wiki/Thermal_conductivity and http://en.wikipedia.org/wiki/List_of_thermal_conductivities.

on temperature and pressure which phase prevails. Thus, the latent heat release can be calculated from the change of entropy ΔS and the derivatives of the phase function $\frac{\partial X}{\partial T}$ and $\frac{\partial X}{\partial p}$. These values have to be provided by the material model, separately for the coefficient $-\Delta S \frac{\partial X}{\partial T}$ on the left-hand side and $\Delta S \frac{\partial X}{\partial p}$ on the right-hand side of the temperature equation. However, they may be either approximated with the help of an analytic phase function, employing data from a thermodynamic database or in any other way that seems appropriate to the user.

2.3 Dimensional or non-dimensionalized equations?

Equations (1)–(3) are stated in the physically correct form. One would usually interpret them in a way that the various coefficients such as the viscosity, density and thermal conductivity η, ρ, κ are given in their correct physical units, typically expressed in a system such as the meter, kilogram, second (MKS) system that is part of the SI system. This is certainly how we envision ASPECT to be used: with geometries, material models, boundary conditions and initial values to be given in their correct physical units. As a consequence, when ASPECT prints information about the simulation onto the screen, it typically does so by using a postfix such as m/s to indicate a velocity or W/m^2 to indicate a heat flux.

Note: For convenience, output quantities are sometimes provided in units meters per *year* instead of meters per *second* (velocities) or in *years* instead of *seconds* (the current time, the time step size); this conversion happens at the time output is generated, and is not part of the solution process. Whether this conversion should happen is determined by the flag “`Use years in output instead of seconds`” in the input file, see Section A.1. Obviously, this conversion from seconds to years only makes sense if the model is described in physical units rather than in non-dimensionalized form, see below.

That said, in reality, ASPECT has no preferred system of units as long as every material constant, geometry, time, etc., are all expressed in the same system. In other words, it is entirely legitimate to implement geometry and material models in which the dimension of the domain is one, density and viscosity are one, and the density variation as a function of temperature is scaled by the Rayleigh number – i.e., to use the usual non-dimensionalization of the equations (1)–(3). Some of the cookbooks in Section 5 use this non-dimensional form; for example, the simplest cookbook in Section 5.2.1 as well as the SolCx, SolKz and inclusion benchmarks in Sections 5.4.4, are such cases. Whenever this is the case, output showing units m/s or W/m^2 clearly no longer have a literal meaning. Rather, the unit postfix must in this case simply be interpreted to mean that the number that precedes the first is a velocity and a heat flux in the second case.

In other words, whether a computation uses physical or non-dimensional units really depends on the geometry, material, initial and boundary condition description of the particular case under consideration – ASPECT will simply use whatever it is given. Whether one or the other is the more appropriate description is a decision we purposefully leave to the user. There are of course good reasons to use non-dimensional descriptions of realistic problems, rather than to use the original form in which all coefficients remain in their physical units. On the other hand, there are also downsides:

- Non-dimensional descriptions, such as when using the Rayleigh number to indicate the relative strength of convective to diffusive thermal transport, have the advantage that they allow to reduce a system to its essence. For example, it is clear that we get the same behavior if one increases both the viscosity and the thermal expansion coefficient by a factor of two because the resulting Rayleigh number; similarly, if we were to increase the size of the domain by a factor of 2 and thermal diffusion coefficient by a factor of 8. In both of these cases, the non-dimensional equations are exactly the same. On the other hand, the equations in their physical unit form are different and one may not see that the result of this variations in coefficients will be exactly the same as before. Using non-dimensional variables therefore reduces the space of independent parameters one may have to consider when doing parameter studies.

- From a practical perspective, equations (1)–(3) are often ill-conditioned in their original form: the two sides of each equation have physical units different from those of the other equations, and their numerical values are often vastly different.⁴ Of course, these values can not be compared: they have different physical units, and the ratios between these values depends on whether we choose to measure lengths in meters or kilometers, for example. Nevertheless, when implementing these equations in software, at one point or another, we have to work with numbers and at this point the physical units are lost. If one does not take care at this point, it is easy to get software in which all accuracy is lost due to round-off errors. On the other hand, non-dimensionalization typically avoids this since it normalizes all quantities so that values that appear in computations are typically on the order of one.
- On the downside, the numbers non-dimensionalized equations produce are not immediately comparable to ones we know from physical experiments. This is of little concern if all we have to do is convert every output number of our program back to physical units. On the other hand, it is more difficult and a source of many errors if this has to be done inside the program, for example, when looking up the viscosity as a pressure-, temperature- and strain-rate-dependent function: one first has to convert pressure, temperature and strain rate from non-dimensional to physical units, look up the corresponding viscosity in a table, and then convert the viscosity back to non-dimensional quantities. Getting this right at every one of the dozens or hundreds of places inside a program and using the correct (but distinct) conversion factors for each of these quantities is both a challenge and a possible source of errors.
- From a mathematical viewpoint, it is typically clear how an equation needs to be non-dimensionalized if all coefficients are constant. However, how is one to normalize the equations if, as is the case in the earth mantle, the viscosity varies by several orders of magnitude? In cases like these, one has to choose a reference viscosity, density, etc. While the resulting non-dimensionalization retains the universality of parameters in the equations, as discussed above, it is not entirely clear that this would also retain the numerical stability if the reference values are poorly chosen.

As a consequence of such considerations, most codes in the past have used non-dimensionalized models. This was aided by the fact that until recently and with notable exceptions, many models had constant coefficients and the difficulties associated with variable coefficients were not a concern. On the other hand, our goal with ASPECT is for it to be a code that solves realistic problems using complex models and that is easy to use. Thus, we allow users to input models in physical or non-dimensional units, at their discretion. We believe that this makes the description of realistic models simpler. On the other hand, ensuring numerical stability is not something users should have to be concerned about, and is taken care of in the implementation of ASPECT’s core (see the corresponding section in [KHB12]).

2.4 Static or dynamic pressure?

One could reformulate equation (1) somewhat. To this end, let us say that we would want to represent the pressure p as the sum of two parts that we will call static and dynamic, $p = p_s + p_d$. If we assume that p_s is already given, then we can replace (1) by

$$-\nabla \cdot 2\eta \nabla \mathbf{u} + \nabla p_d = \rho \mathbf{g} - \nabla p_s.$$

One typically chooses p_s as the pressure one would get if the whole medium were at rest – i.e., as the hydrostatic pressure. This pressure can be computed noting that (1) reduces to

$$\nabla p_s = \rho(p_s, T_s, \mathbf{x}) \mathbf{g} = \bar{\rho} \mathbf{g}$$

⁴To illustrate this, consider convection in the Earth as a back-of-the-envelope example. With the length scale of the mantle $L = 3 \cdot 10^6$ m, viscosity $\eta = 10^{24}$ kg/m/s, density $\rho = 3 \cdot 10^3$ kg/m³ and a typical velocity of $U = 0.1$ m/year = $3 \cdot 10^{-9}$ m/s, we get that the friction term in (1) has size $\eta U / L^2 \approx 3 \cdot 10^2$ kg/m²/s². On the other hand, the term $\nabla \cdot (\rho \mathbf{u})$ in the continuity equation (2) has size $\rho U / L \approx 3 \cdot 10^{-12}$ kg/s/m³. In other words, their *numerical values* are 14 orders of magnitude apart.

in the absence of any motion where T_s is some static temperature field (see also Section 2.6). This, our rewritten version of (1) would look like this:

$$-\nabla \cdot 2\eta \nabla \mathbf{u} + \nabla p_d = [\rho(p, T, \mathbf{x}) - \rho(p_s, T_s, \mathbf{x})] \mathbf{g}.$$

In this formulation, it is clear that the quantity that drives the fluid flow is in fact the *buoyancy* caused by the *variation* of densities, not the density itself.

This reformulation has a number of advantages and disadvantages:

- One can notice that in many realistic cases, the dynamic component p_d of the pressure is orders of magnitude smaller than the static component p_s . For example, in the earth, the two are separated by around 6 orders of magnitude at the bottom of the earth mantle. Consequently, if one wants to solve the linear system that arises from discretization of the original equations, one has to solve it a significant degree of accuracy (6–7 digits) to get the dynamic part of the pressure correct to even one digit. This entails a very significant numerical effort, and one that is not necessary if we can split the pressure in a way so that the pre-computed static pressure p_s (or, rather, the density using the static pressure and temperature from which p_s results) absorbs the dominant part and one only has to compute the remaining, dynamic pressure to 2 or 3 digits of accuracy, rather than the corresponding 7–8 for the total pressure.
- On the other hand, the pressure p_d one computes this way is not immediately comparable to quantities that we use to look up pressure-dependent quantities such as the density. Rather, one needs to first find the static pressure as well (see Section 2.6) and add the two together before they can be used to look up material properties or to compare them with experimental results. Consequently, if the pressure a program outputs (either for visualization, or in the internal interfaces to parts of the code where users can implement pressure- and temperature-dependent material properties) is only the dynamic component, then all of the consumers of this information need to convert it into the total pressure when comparing with physical experiments. Since any code implementing realistic material models has a great many of these places, there is a large potential for inadvertent errors and bugs.
- Finally, the definition of a reference density $\rho(p_s, T_s, \mathbf{x})$ derived from static pressures and temperatures is only simple if we have incompressible models and under the assumption that the temperature-induced density variations are small compared to the overall density. In this case, we can choose $\rho(p_s, T_s, \mathbf{x}) = \rho_0$ with a constant reference density ρ_0 . On the other hand, for more complicated models, it is not a priori clear which density to choose since we first need to compute static pressures and temperatures – quantities that satisfy equations that introduce boundary layers, may include phase changes releasing latent heat, and where the density may have discontinuities at certain depths, see Section 2.6.

Thus, if we compute adiabatic pressures and temperatures \bar{p}_s, \bar{T}_s under the assumption of a thermal boundary layer worth 900 Kelvin at the top, and we get a corresponding density profile $\bar{\rho} = \rho(\bar{p}_s, \bar{T}_s, \mathbf{x})$, but after running for a few million years the temperature turns out to be so that the top boundary layer has a jump of only 800 Kelvin with corresponding adiabatic pressures and temperatures \hat{p}_s, \hat{T}_s , then a more appropriate density profile would be $\hat{\rho} = \rho(\hat{p}_s, \hat{T}_s, \mathbf{x})$.

The problem is that it may well be that the erroneously computed density profile $\hat{\rho}$ does *not* lead to a separation where $|p_d| \ll |p_s|$ because, especially if the material undergoes phase changes, there will be entire areas of the computational domain in which $|\rho - \hat{\rho}_s| \ll |\rho|$ but $|\rho - \bar{\rho}_s| \not\ll |\rho|$. Consequently the benefits of lesser requirements on the iterative linear solver would not be realized.

We do note that most of the codes available today and that we are aware of split the pressure into static and dynamic parts nevertheless, either internally or require the user to specify the density profile as the difference between the true and the hydrostatic density. This may, in part, be due to the fact that historically most codes were written to solve problems in which the medium was considered incompressible, i.e., where the definition of a static density was simple.

On the other hand, we intend ASPECT to be a code that can solve more general models for which this definition is not as simple. As a consequence, we have chosen to solve the equations as stated originally – i.e., we solve for the *full* pressure rather than just its *dynamic* component. With most traditional methods, this would lead to a catastrophic loss of accuracy in the dynamic pressure since it is many orders of magnitude smaller than the total pressure at the bottom of the earth mantle. We avoid this problem in ASPECT by using a cleverly chosen iterative solver that ensures that the full pressure we compute is accurate enough so that the dynamic pressure can be extracted from it with the same accuracy one would get if one were to solve for only the dynamic component. The methods that ensure this are described in detail in [KHB12] and in particular in the appendix of that paper.

Note: By default, ASPECT uses the full pressure in the equations, and only prescribing density deviations from a reference state on the right-hand side of (1) would lead to negative densities in the energy equation (3). However, when using one of the approximations described in Section 2.10, the energy balance uses the reference density $\bar{\rho}$ instead of the full density, which makes it possible to formulate the Stokes system in terms of the dynamic instead of the full pressure. In order to do this, one would have to use a material model (see Section 6.3.1) in which the density is in fact a density variation, and then the pressure solution variable would only be the dynamic pressure.

2.5 Pressure normalization

The equations described above, (1)–(3), only determine the pressure p up to an additive constant. On the other hand, since the pressure appears in the definition of many of the coefficients, we need a pressure that has some sort of *absolute* definition. A physically useful definition would be to normalize the pressure in such a way that the average pressure along the “surface” has a prescribed value where the geometry description (see Section 6.3.3) has to determine which part of the boundary of the domain is the “surface” (we call a part of the boundary the “surface” if its depth is “close to zero”).

Typically, one will choose this average pressure to be zero, but there is a parameter “Surface pressure” in the input file (see Section A.1) to set it to a different value. One may want to do that, for example, if one wants to simulate the earth mantle without the overlying lithosphere. In that case, the “surface” would be the interface between mantle and lithosphere, and the average pressure at the surface to which the solution of the equations will be normalized should in this case be the hydrostatic pressure at the bottom of the lithosphere.

An alternative is to normalize the pressure in such a way that the *average* pressure throughout the domain is zero or some constant value. This is not a useful approach for most geodynamics applications but is common in benchmarks for which analytic solutions are available. Which kind of normalization is chosen is determined by the “Pressure normalization” flag in the input file, see Section A.1.

2.6 Initial conditions and the adiabatic pressure/temperature

Equations (1)–(3) require us to pose initial conditions for the temperature, and this is done by selecting one of the existing models for initial conditions in the input parameter file, see Section A.69. The equations themselves do not require that initial conditions are specified for the velocity and pressure variables (since there are no time derivatives on these variables in the model).

Nevertheless, a nonlinear solver will have difficulty converging to the correct solution if we start with a completely unphysical pressure for models in which coefficients such as density ρ and viscosity η depend on the pressure and temperature. To this end, ASPECT uses pressure and temperature fields $p_{\text{ad}}(z), T_{\text{ad}}(z)$ computed in the adiabatic conditions model (see Section A.2). By default, these fields satisfy adiabatic

conditions:

$$\rho C_p \frac{d}{dz} T_{\text{ad}}(z) = \frac{\partial \rho}{\partial T} T_{\text{ad}}(z) g_z, \quad (13)$$

$$\frac{d}{dz} p_{\text{ad}}(z) = \rho g_z, \quad (14)$$

where strictly speaking g_z is the magnitude of the vertical component of the gravity vector field, but in practice we take the magnitude of the entire gravity vector.

These equations can be integrated numerically starting at $z = 0$, using the depth dependent gravity field and values of the coefficients $\rho = \rho(p, T, z), C_p = C_p(p, T, z)$. As starting conditions at $z = 0$ we choose a pressure $p_{\text{ad}}(0)$ equal to the average surface pressure (often chosen to be zero, see Section 2.5), and an adiabatic surface temperature $T_{\text{ad}}(0)$ that is also selected in the input parameter file.

However, users can also supply their own adiabatic conditions models or define an arbitrary profile using the “function” plugin.

Note: The adiabatic surface temperature is often chosen significantly higher than the actual surface temperature. For example, on earth, the actual surface temperature is on the order of 290 K, whereas a reasonable adiabatic surface temperature is maybe 1600 K. The reason is that the bulk of the mantle is more or less in thermal equilibrium with a thermal profile that corresponds to the latter temperature, whereas the very low actual surface temperature and the very high bottom temperature at the core-mantle boundary simply induce a thermal boundary layer. Since the temperature and pressure profile we compute using the equations above are simply meant to be good starting points for nonlinear solvers, it is important to choose this profile in such a way that it covers most of the mantle well; choosing an adiabatic surface temperature of 290 K would yield a temperature and pressure profile that is wrong almost throughout the entire mantle.

2.7 Compositional fields

The last of the basic equations, (4), describes the evolution of a set of variables $c_i(\mathbf{x}, t), i = 1 \dots C$ that we typically call *compositional fields* and that we often aggregate into a vector \mathbf{c} .

Compositional fields were originally intended to track what their name suggest, namely the chemical composition of the convecting medium. In this interpretation, the composition is a non-diffusive quantity that is simply advected along passively, i.e., it would satisfy the equation

$$\frac{\partial \mathbf{c}}{\partial t} + \mathbf{u} \cdot \nabla \mathbf{c} = 0.$$

However, the compositional fields may also participate in determining the values of the various coefficients as discussed in Section 2.2, and in this sense the equation above describes a composition that is *passively advected*, but an *active participant* in the equations.

That said, over time compositional fields have shown to be a much more useful tool than originally intended. For example, they can be used to track where material comes from and goes to (see Section 5.2.4) and, if one allows for a reaction rate q on the right hand side,

$$\frac{\partial \mathbf{c}}{\partial t} + \mathbf{u} \cdot \nabla \mathbf{c} = q,$$

then one can also model interaction between species – for example to simulate phase changes where one compositional field, indicating a particular phase, transforms into another phase depending on pressure and temperature, or where several phases combine to other phases. Another example of using a right hand side – quite outside what the original term *compositional field* was supposed to indicate – is to track the accumulation of finite strain, see Section 5.2.11.

In actual practice, one finds that it is often useful to allow \mathbf{q} to be a function that has both a smooth (say, continuous) in time component, and one that is singular in time (i.e., contains Dirac delta, or “impulse” functions). Typical time integrators require the evaluation of the right hand side at specific points in time, but this would preclude the use of delta functions. Consequently, the integrators in ASPECT only require material models to provide an *integrated* value $\int_t^{t+\Delta t} \mathbf{q}(\tau) d\tau$ through the `reaction_term` output variable. Implementations often approximate this as $\Delta t \cdot \mathbf{q}(t)$, or similar formulas.

A second application for only providing integrated right hand sides comes from the fact that modeling reactions between different compositional fields often involves finding an equilibrium state between different fields because chemical reactions happen on a much faster time scale than transport. In other words, one then often assumes that there is a $\mathbf{c}^*(p, T)$ so that

$$\mathbf{q}(p, T, \varepsilon(\mathbf{u}), \mathbf{c}^*(p, T)) = 0.$$

Consequently, the material model methods that deal with source terms for the compositional fields need to compute an *increment* $\Delta \mathbf{c}$ to the previous value of the compositional fields so that the sum of the previous values and the increment equals \mathbf{c}^* . This corresponds to an *impulse change* in the compositions at every time step, as opposed to the usual approach of evaluating the right hand side term \mathbf{q} as a continuous function in time, which corresponds to a *rate*.

On the other hand, there are other uses of compositional fields that do not actually have anything to do with quantities that can be considered related to compositions. For example, one may define a field that tracks the grain size of rocks. If the strain rate is high, then the grain size decreases as the rocks break. If the temperature is high enough, then grains heal and their size increases again. Such “damage” models would then call for an equation of the form (assuming one uses only a single compositional field)

$$\frac{\partial c}{\partial t} + \mathbf{u} \cdot \nabla c = q(T, c),$$

where in the simplest case one could postulate

$$q(T, c) = -Ac + B \max\{T - T_{\text{healing}}, 0\}c.$$

One would then use this compositional field in the definition of the viscosity of the material: more damage means lower viscosity because the rocks are weaker.

In cases like this, there is only a single compositional field and it is not in permanent equilibrium. Consequently, the increment implementations of material models in ASPECT need to compute is typically the rate $q(T, c)$ times the time step. In other words, if you compute a reaction rate inside the material model you need to multiply it by the time step size before returning the value.

Compositional fields have proven to be surprisingly versatile tools to model all sorts of components of models that go beyond the simple Stokes plus temperature set of equations. Play with them!

2.8 Constitutive laws

Equation (1) describes buoyancy-driven flow in an isotropic fluid where strain rate is related to stress by a scalar (possibly spatially variable) multiplier, η . For some material models it is useful to generalize this relationship to anisotropic materials, or other exotic constitutive laws. For these cases ASPECT can optionally include a generalized, fourth-order tensor field as a material model state variable which changes equation (1) to

$$-\nabla \cdot \left[2\eta \left(C\varepsilon(\mathbf{u}) - \frac{1}{3}(tr(C\varepsilon(\mathbf{u})))\mathbf{1} \right) \right] + \nabla p = \rho\mathbf{g} \quad \text{in } \Omega \quad (15)$$

and the shear heating term in equation (3) to

$$\dots + 2\eta \left(C\varepsilon(\mathbf{u}) - \frac{1}{3}(tr(C\varepsilon(\mathbf{u})))\mathbf{1} \right) : \left(\varepsilon(\mathbf{u}) - \frac{1}{3}(\nabla \cdot \mathbf{u})\mathbf{1} \right) \dots \quad (16)$$

where $C = C_{ijkl}$ is defined by the material model. For physical reasons, C needs to be a symmetric rank-4 tensor: i.e., when multiplied by a symmetric (strain rate) tensor of rank 2 it needs to return another symmetric tensor of rank 2. In mathematical terms, this means that $C_{ijkl} = C_{jikl} = C_{ijlk} = C_{jilk}$. Energy considerations also require that C is positive definite: i.e., for any $\varepsilon \neq 0$, the scalar $\varepsilon : (C\varepsilon)$ must be positive.

This functionality can be optionally invoked by any material model that chooses to define a C field, and falls back to the default case ($C = \mathbb{I}$) if no such field is defined. It should be noted that η still appears in equations (15) and (16). C is therefore intended to be thought of as a “director” tensor rather than a replacement for the viscosity field, although in practice either interpretation is okay.

2.9 Numerical methods

There is no shortage in the literature for methods to solve the equations outlined above. The methods used by ASPECT use the following, interconnected set of strategies in the implementation of numerical algorithms:

- *Mesh adaptation:* Mantle convection problems are characterized by widely disparate length scales (from plate boundaries on the order of kilometers or even smaller, to the scale of the entire earth). Uniform meshes can not resolve the smallest length scale without an intractable number of unknowns. Fully adaptive meshes allow resolving local features of the flow field without the need to refine the mesh globally. Since the location of plumes that require high resolution change and move with time, meshes also need to be adapted every few time steps.
- *Accurate discretizations:* The equations upon which most models for the earth mantle are based have a number of intricacies that make the choice of discretization non-trivial. In particular, the finite elements chosen for velocity and pressure need to satisfy the usual compatibility condition for saddle point problems. This can be worked around using pressure stabilization schemes for low-order discretizations, but high-order methods can yield better accuracy with fewer unknowns and offer more reliability. Equally important is the choice of a stabilization method for the highly advection-dominated temperature equation. ASPECT uses a nonlinear artificial diffusion method for the latter.
- *Efficient linear solvers:* The major obstacle in solving the system of linear equations that results from discretization is the saddle-point nature of the Stokes equations. Simple linear solvers and preconditioners can not efficiently solve this system in the presence of strong heterogeneities or when the size of the system becomes very large. ASPECT uses an efficient solution strategy based on a block triangular preconditioner utilizing an algebraic multigrid that provides optimal complexity even up to problems with hundreds of millions of unknowns.
- *Parallelization of all of the steps above:* Global mantle convection problems frequently require extremely large numbers of unknowns for adequate resolution in three dimensional simulations. The only realistic way to solve such problems lies in parallelizing computations over hundreds or thousands of processors. This is made more complicated by the use of dynamically changing meshes, and it needs to take into account that we want to retain the optimal complexity of linear solvers and all other operations in the program.
- *Modularity of the code:* A code that implements all of these methods from scratch will be unwieldy, unreadable and unusable as a community resource. To avoid this, we build our implementation on

widely used and well tested libraries that can provide researchers interested in extending it with the support of a large user community. Specifically, we use the DEAL.II library [BHK07, BHK12] for meshes, finite elements and everything discretization related; the TRILINOS library [HBH⁺05, H⁺11] for scalable and parallel linear algebra; and P4EST [BWG11] for distributed, adaptive meshes. As a consequence, our code is freed of the mundane tasks of defining finite element shape functions or dealing with the data structures of linear algebra, can focus on the high-level description of what is supposed to happen, and remains relatively compact. The code will also automatically benefit from improvements to the underlying libraries with their much larger development communities. ASPECT is extensively documented to enable other researchers to understand, test, use, and extend it.

Rather than detailing the various techniques upon which ASPECT is built, we refer to the paper by Kronbichler, Heister and Bangerth [KHB12] that gives a detailed description and rationale for the various building blocks.

2.10 Approximate equations

There are a number of common variations to equations (1)–(3) that are frequently used in the geosciences. For example, one frequently finds references to the anelastic liquid approximation (ALA), truncated anelastic liquid approximation (TALA), and the Boussinesq approximation (BA). These can all be derived from the basic equations (1)–(3) via various approximations, and we will discuss them in the following. Since they are typically only provided considering velocity, pressure and temperature, we will in the following omit the dependence on the compositional fields used in previous sections, though this dependence can easily be added back into the equations stated below. A detailed discussion of the approximations introduced below can also be found in [KLVK⁺10].

The three approximations mentioned all start by writing the pressure and temperature as the sum of a (possibly depth dependent) reference state plus a perturbation, i.e., we will write

$$\begin{aligned} p(\mathbf{x}, t) &= \bar{p}(z) + p'(\mathbf{x}, t), \\ T(\mathbf{x}, t) &= \bar{T}(z) + T'(\mathbf{x}, t). \end{aligned}$$

Here, barred quantities are reference states and may depend on the depth z (not necessarily the third component of \mathbf{x}) whereas primed quantities are the spatially and temporally variable deviations of the temperature and pressure fields from this reference state. In particular, the reference pressure is given by solving the hydrostatic equation,

$$\nabla \bar{p} = \bar{\rho} \mathbf{g}, \quad (17)$$

where $\bar{\rho} = \rho(\bar{p}, \bar{T})$ is a *reference density* that depends on depth and represents a typical change of material parameters and solution variables with depth. $\bar{T}(z)$ is chosen as an adiabatic profile accounting for the fact that the temperature increases as the pressure increases. With these definitions, equations (1)–(2) can equivalently be written as follows:

$$-\nabla \cdot \left[2\eta \left(\varepsilon(\mathbf{u}) - \frac{1}{3}(\nabla \cdot \mathbf{u})\mathbf{1} \right) \right] + \nabla p' = (\rho - \bar{\rho})\mathbf{g} \quad \text{in } \Omega, \quad (18)$$

$$\nabla \cdot (\rho \mathbf{u}) = 0 \quad \text{in } \Omega. \quad (19)$$

The temperature equation, when omitting entropic effects, still reads as

$$\begin{aligned} \rho C_p \left(\frac{\partial T}{\partial t} + \mathbf{u} \cdot \nabla T \right) - \nabla \cdot k \nabla T \\ = \rho H + 2\eta \left(\varepsilon(\mathbf{u}) - \frac{1}{3}(\nabla \cdot \mathbf{u})\mathbf{1} \right) : \left(\varepsilon(\mathbf{u}) - \frac{1}{3}(\nabla \cdot \mathbf{u})\mathbf{1} \right) + \alpha T (\mathbf{u} \cdot \nabla p) \quad \text{in } \Omega, \end{aligned} \quad (20)$$

where the right-hand side includes radiogenic heat production, shear heating and adiabatic heating (in that order).

Starting from these equations, the approximations discussed in the next few subsections make use of the fact that for the flows for which these approximations are valid, the perturbations p' , T' are much smaller than typical values of the reference quantities \bar{p} , \bar{T} . The terms influenced by these approximations are $\nabla \cdot (\rho u) = 0$ in the continuity equation, and all occurrences of $\rho(p, T)$ in the temperature equation, and we will discuss them separately below. The equations for these approximations are almost always given in terms of non-dimensionalized quantities. We will for now stick with the dimensional form because it expresses in a clearer way the approximations that are made. The non-dimensionalization can then be done on each of the forms below separately.

2.10.1 The anelastic liquid approximation (ALA)

The *anelastic liquid approximation (ALA)* is based on two assumptions. First, that the density variations $\rho(p, T) - \bar{\rho}$ are small and in particular can be accurately described by a Taylor expansion:

$$\rho(p, T) \approx \bar{\rho} + \frac{\partial \rho(\bar{p}, \bar{T})}{\partial p} p' + \frac{\partial \rho(\bar{p}, \bar{T})}{\partial T} T'.$$

Here, $\frac{\partial \rho(\bar{p}, \bar{T})}{\partial T}$ is related to the thermal expansion coefficient $\alpha = -\frac{1}{\rho} \frac{\partial \rho}{\partial T}$, and $\frac{\partial \rho(\bar{p}, \bar{T})}{\partial p}$ to the compressibility $\kappa = \frac{1}{\rho} \frac{\partial \rho}{\partial p}$.

The second assumption is that the variation of the density from the reference density can be neglected in the mass balance and temperature equations. This yields the following system of equations for the velocity and pressure equations:

$$-\nabla \cdot \left[2\eta \left(\varepsilon(\mathbf{u}) - \frac{1}{3}(\nabla \cdot \mathbf{u})\mathbf{1} \right) \right] + \nabla p' = \left(\frac{\partial \rho(\bar{p}, \bar{T})}{\partial p} p' + \frac{\partial \rho(\bar{p}, \bar{T})}{\partial T} T' \right) \mathbf{g} \quad \text{in } \Omega, \quad (21)$$

$$\nabla \cdot (\bar{\rho} \mathbf{u}) = 0 \quad \text{in } \Omega. \quad (22)$$

For the temperature equation, using the definition of the hydrostatic pressure gradient (17), we arrive at the following:

$$\begin{aligned} \bar{\rho} C_p \left(\frac{\partial T}{\partial t} + \mathbf{u} \cdot \nabla T \right) - \nabla \cdot k \nabla T \\ = \bar{\rho} H + 2\eta \left(\varepsilon(\mathbf{u}) - \frac{1}{3}(\nabla \cdot \mathbf{u})\mathbf{1} \right) : \left(\varepsilon(\mathbf{u}) - \frac{1}{3}(\nabla \cdot \mathbf{u})\mathbf{1} \right) + \alpha \bar{\rho} T (\mathbf{u} \cdot \mathbf{g}) \quad \text{in } \Omega. \end{aligned} \quad (23)$$

Note: Our energy equation is formulated in terms of T , while in the literature, the equation has sometimes been formulated in terms of T' , which yields additional terms containing \bar{T} on the right-hand side. Both ways of writing the equation are equivalent.

2.10.2 The truncated anelastic liquid approximation (TALA)

The *truncated anelastic liquid approximation (TALA)* further simplifies the ALA by assuming that the variation of the density due to pressure variations is small, i.e., that

$$\rho(p, T) \approx \bar{\rho} + \frac{\partial \rho(\bar{p}, \bar{T})}{\partial T} T'.$$

This does not mean that the density is not pressure dependent – it will, for example, continue to be depth dependent because the hydrostatic pressure grows with depth. It simply means that the deviations from the

reference pressure are assumed to be so small that they do not matter in describing the density. Because the pressure variation p' is induced by the flow field (the static component pressure is already taken care of by the hydrostatic pressure), this assumption in essence means that we assume the flow to be very slow, even beyond the earlier assumption that we can neglect inertial terms when deriving (1)–(2).

This further assumption then transforms (21)–(22) into the following equations:

$$-\nabla \cdot \left[2\eta \left(\varepsilon(\mathbf{u}) - \frac{1}{3}(\nabla \cdot \mathbf{u})\mathbf{1} \right) \right] + \nabla p' = \frac{\partial \rho(\bar{p}, \bar{T})}{\partial T} T' \mathbf{g} \quad \text{in } \Omega, \quad (24)$$

$$\nabla \cdot (\bar{\rho} \mathbf{u}) = 0 \quad \text{in } \Omega. \quad (25)$$

The energy equation is the same as in the ALA case.

2.10.3 The Boussinesq approximation (BA)

If we further assume that the reference temperature and the reference density are constant, $\bar{T}(z) = T_0$, $\bar{\rho}(\bar{p}, \bar{T}) = \rho_0$, – in other words, density variations are so small that they are negligible everywhere except for in the right-hand side of the velocity equation (the buoyancy term), which describes the driving force of the flow, then we can further simplify the mass conservation equations of the TALA to $\nabla \cdot \mathbf{u} = 0$. This means that the density in all other parts of the equations is not only independent of the pressure variations p' as assumed in the TALA, but also does not depend on the much larger hydrostatic pressure \bar{p} nor on the reference temperature \bar{T} . We then obtain the following set of equations that also uses the incompressibility in the definition of the strain rate:

$$-\nabla \cdot [2\eta\varepsilon(\mathbf{u})] + \nabla p' = \frac{\partial \rho(\bar{p}, \bar{T})}{\partial T} T' \mathbf{g} \quad \text{in } \Omega, \quad (26)$$

$$\nabla \cdot \mathbf{u} = 0 \quad \text{in } \Omega. \quad (27)$$

In addition, as the reference temperature is constant, one needs to neglect the adiabatic and shear heating in the energy equation

$$\bar{\rho}C_p \left(\frac{\partial T}{\partial t} + \mathbf{u} \cdot \nabla T \right) - \nabla \cdot k \nabla T = \bar{\rho}H \quad \text{in } \Omega. \quad (28)$$

On incompressibility. The Boussinesq approximation assumes that the density can be considered constant in all occurrences in the equations with the exception of the buoyancy term on the right hand side of (1). The primary result of this assumption is that the continuity equation (2) will now read

$$\nabla \cdot \mathbf{u} = 0.$$

This makes the equations *much* simpler to solve: First, because the divergence operation in this equation is the transpose of the gradient of the pressure in the momentum equation (1), making the system of these two equations symmetric. And secondly, because the two equations are now linear in pressure and velocity (assuming that the viscosity η and the density ρ are considered fixed). In addition, one can drop all terms involving $\nabla \cdot \mathbf{u}$ from the left hand side of the momentum equation (1); while dropping these terms does not affect the solution of the equations, it makes assembly of linear systems faster.

From a physical perspective, the assumption that the density is constant in the continuity equation but variable in the momentum equation is of course inconsistent. However, it is justified if the variation is small since the momentum equation can be rewritten to read

$$-\nabla \cdot 2\eta\varepsilon(\mathbf{u}) + \nabla p' = (\rho - \rho_0)\mathbf{g},$$

where p' is the *dynamic* pressure and ρ_0 is the constant reference density. This makes it clear that the true driver of motion is in fact the *deviation* of the density from its background value, however small this value is: the resulting velocities are simply proportional to the density variation, not to the absolute magnitude of the density.

As such, the Boussinesq approximation can be justified. On the other hand, given the real pressures and temperatures at the bottom of the Earth's mantle, it is arguable whether the density can be considered to be almost constant. Most realistic models predict that the density of mantle rocks increases from somewhere around 3300 at the surface to over 5000 kilogram per cubic meters at the core mantle boundary, due to the increasing lithostatic pressure. While this appears to be a large variability, if the density changes slowly with depth, this is not in itself an indication that the Boussinesq approximation will be wrong. To this end, consider that the continuity equation can be rewritten as $\frac{1}{\rho} \nabla \cdot (\rho \mathbf{u}) = 0$, which we can multiply out to obtain

$$\nabla \cdot \mathbf{u} + \frac{1}{\rho} \mathbf{u} \cdot \nabla \rho = 0.$$

The question whether the Boussinesq approximation is valid is then whether the second term (the one omitted in the Boussinesq model) is small compared to the first. To this end, consider that the velocity can change completely over length scales of maybe 10 km, so that $\nabla \cdot \mathbf{u} \approx \|u\|/10\text{km}$. On the other hand, given a smooth dependence of density on pressure, the length scale for variation of the density is the entire earth mantle, i.e., $\frac{1}{\rho} \mathbf{u} \cdot \nabla \rho \approx \|u\|0.5/3000\text{km}$ (given a variation between minimal and maximal density of 0.5 times the density itself). In other words, for a smooth variation, the contribution of the compressibility to the continuity equation is very small. This may be different, however, for models in which the density changes rather abruptly, for example due to phase changes at mantle discontinuities.

On almost linear models. A further simplification can be obtained if one assumes that all coefficients with the exception of the density do not depend on the solution variables but are, in fact, constant. In such models, one typically assumes that the density satisfies a relationship of the form $\rho = \rho(T) = \rho_0(1 - \alpha(T - T_0))$ with a small thermal expansion coefficient α and a reference density ρ_0 that is attained at temperature T_0 . Since the thermal expansion is considered small, this naturally leads to the following variant of the Boussinesq model discussed above:

$$-\nabla \cdot [2\eta\varepsilon(\mathbf{u})] + \nabla p' = -\alpha\rho_0 T g \quad \text{in } \Omega, \quad (29)$$

$$\nabla \cdot \mathbf{u} = 0 \quad \text{in } \Omega, \quad (30)$$

$$\rho_0 C_p \left(\frac{\partial T}{\partial t} + \mathbf{u} \cdot \nabla T \right) - \nabla \cdot k \nabla T = \rho H \quad \text{in } \Omega. \quad (31)$$

Note that the right hand side forcing term in (29) is now only the deviation of the gravitational force from the force that would act if the material were at temperature T_0 .

Under the assumption that all other coefficients are constant, one then arrives at equations in which the only nonlinear term is the advection term, $\mathbf{u} \cdot \nabla T$ in the temperature equation (31). This facilitates the use of a particular class of time stepping schemes in which one does not solve the whole set of equations at once, iterating out nonlinearities as necessary, but instead in each time step solves first the Stokes system with the previous time step's temperature, and then uses the so-computed velocity to solve the temperature equation. These kind of time stepping schemes are often referred to as *operator splitting* methods.

Note: ASPECT does not solve the equations in the way described in this paragraph, however, a particular operator splitting method was used in earlier ASPECT versions. It first solves the Stokes equations and then uses a semi-explicit time stepping method for the temperature equation where diffusion is handled implicitly and advection explicitly. This algorithm is often called *IMPES* (it originated in the porous media flow community, where the acronym stands for *Implicit Pressure, Explicit Saturation*) and is explained in more detail in [KHB12]. Since then the algorithm in ASPECT has been rewritten to use an implicit time stepping algorithm also for the temperature equation because this allows to use larger time steps.

2.10.4 The isothermal compression approximation (ICA)

In the compressible case and without the assumption of a reference state, the conservation of mass equation in equation (2) is $\nabla \cdot (\rho \mathbf{u}) = 0$, which is nonlinear and not symmetric to the ∇p term in the force balance equation (1), making solving and preconditioning the resulting linear and nonlinear systems difficult. To make this work in ASPECT, we consequently reformulate this equation. Dividing by ρ and applying the product rule of differentiation gives

$$\frac{1}{\rho} \nabla \cdot (\rho \mathbf{u}) = \nabla \cdot \mathbf{u} + \frac{1}{\rho} \nabla \rho \cdot \mathbf{u}.$$

We will now make two basic assumptions: First, the variation of the density $\rho(p, T, \mathbf{x}, \mathbf{c})$ is dominated by the dependence on the (total) pressure; in other words, $\nabla \rho \approx \frac{\partial \rho}{\partial p} \nabla p$. This assumption is primarily justified by the fact that, in the Earth's mantle, the density increases by at least 50% between Earth's crust and the core-mantle boundary due to larger pressure there. Secondly, we assume that the pressure is dominated by the static pressure, which implies that $\nabla p \approx \nabla p_s \approx \rho \mathbf{g}$. This is justified, because the viscosity in the Earth is large and velocities are small, hence $\nabla p' \ll \nabla p_s$. This finally allows us to write

$$\frac{1}{\rho} \nabla \rho \cdot \mathbf{u} \approx \frac{1}{\rho} \frac{\partial \rho}{\partial p} \nabla p \cdot \mathbf{u} \approx \frac{1}{\rho} \frac{\partial \rho}{\partial p} \nabla p_s \cdot \mathbf{u} \approx \frac{1}{\rho} \frac{\partial \rho}{\partial p} \rho \mathbf{g} \cdot \mathbf{u}$$

so we get

$$\nabla \cdot \mathbf{u} = -\frac{1}{\rho} \frac{\partial \rho}{\partial p} \rho \mathbf{g} \cdot \mathbf{u} \quad (32)$$

where $\frac{1}{\rho} \frac{\partial \rho}{\partial p}$ is often referred to as the compressibility.

For this approximation, Equation (32) replaces Equation (2). It has the advantage that it retains the symmetry of the Stokes equations if we can treat the right hand side of (32) as known. We do so by evaluating ρ and \mathbf{u} using the solution from the last time step (or values extrapolated from previous time steps), or using a nonlinear solver scheme.

Note: This is the default approximation ASPECT uses to model compressible convection, see Section 2.11.5.

2.11 Choosing a formulation in ASPECT

After discussing different reasonable approximations for modeling compressible or incompressible mantle convection, we will now describe the different steps one has to take to use one of these approximations in a computation. This includes

1. Choosing an approximation for the mass conservation equation;
2. Choosing an approximation for the density in the energy balance, and deciding which heating terms should be included;
3. Formulating the buoyancy term in the material model to be used on the right-hand side of the momentum equation;
4. Prescribing a suitable reference state for the temperature, pressure, and density; i.e. the adiabatic profile, if necessary for the approximations chosen in the first three steps.

All of these choices can be made in the input file by selecting the corresponding parameters (see Sections A.40 and A.2). A description of how to run ASPECT and the basic structure of the input file can be found in Section 4.

2.11.1 Mass conservation approximation

First, we have to choose how to approximate the conservation of mass: $\nabla \cdot (\rho \mathbf{u}) = 0$ (see Equation (2)). We provide the following options, which can be selected in the parameter file in the subsection **Formulation/Mass conservation** (see also A.40):

- “incompressible”:

$$\nabla \cdot \mathbf{u} = 0,$$

- “isothermal compression”:

$$\nabla \cdot \mathbf{u} = -\frac{1}{\rho} \frac{\partial \rho}{\partial p} \rho \mathbf{g} \cdot \mathbf{u},$$

where $\frac{1}{\rho} \frac{\partial \rho}{\partial p} = \kappa$ is the compressibility, and is defined in the material model. This is the explicit compressible mass equation where the velocity \mathbf{u} on the right-hand side is an extrapolated velocity from the last timesteps.

- “hydrostatic compression”:

$$\nabla \cdot \mathbf{u} = -\left(\frac{1}{\rho} \frac{\partial \rho}{\partial p} \rho \mathbf{g} + \frac{1}{\rho} \frac{\partial \rho}{\partial T} \nabla T\right) \cdot \mathbf{u} = -(\kappa \rho \mathbf{g} - \alpha \nabla T) \cdot \mathbf{u}$$

where $\frac{1}{\rho} \frac{\partial \rho}{\partial p} = \kappa$ is the compressibility, $-\frac{1}{\rho} \frac{\partial \rho}{\partial T} = \alpha$ is the thermal expansion coefficient, and both are defined in the material model.

- “reference density profile”:

$$\nabla \cdot \mathbf{u} = -\frac{1}{\bar{\rho}} \frac{\partial \bar{\rho}}{\partial z} \frac{\mathbf{g}}{\|\mathbf{g}\|} \cdot \mathbf{u},$$

where the reference profiles for the density $\bar{\rho}$ and the density gradient $\frac{\partial \bar{\rho}}{\partial z}$ provided by the adiabatic conditions model (2.6) are used. Note that the gravity is assumed to point downwards in depth direction. This is the explicit mass equation where the velocity \mathbf{u} on the right-hand side is an extrapolated velocity from the last timesteps.

- “implicit reference density profile”:

$$\nabla \cdot \mathbf{u} + \frac{1}{\bar{\rho}} \frac{\partial \bar{\rho}}{\partial z} \frac{\mathbf{g}}{\|\mathbf{g}\|} \cdot \mathbf{u} = 0,$$

which uses the same approximation for the density as “reference density profile”, but implements this term on the left-hand side instead of the right-hand side of the mass conservation equation. This effectively uses the current velocity \mathbf{u} instead of an explicitly extrapolated velocity from the last timesteps.

- “ask material model”, which uses “isothermal compression” if the material model reports that it is compressible and “incompressible” otherwise.

Note: The stress tensor approximation.

Incompressibility in the mass conservation equation automatically simplifies the shear strain rate in the momentum and temperature equation from

$$\nabla \cdot \tau = \nabla \cdot \left[2\eta \left(\varepsilon(\mathbf{u}) - \frac{1}{3}(\nabla \cdot \mathbf{u})\mathbf{1} \right) \right]$$

to

$$\nabla \cdot \tau = \nabla \cdot [2\eta \varepsilon(\mathbf{u})],$$

as $\nabla \cdot \mathbf{u} = 0$.

2.11.2 Temperature equation approximation

The density occurs multiple times in the temperature equation. Depending on the selected approximation it is computed in one of two different ways. Which of these options is used can be chosen in the parameter file in the subsection `Formulation/Temperature equation` (see also [A.40](#)):

- “real density”: Use the full density $\rho(p, T)$ that equals the one also used in the buoyancy term of the force balance equation; this is also the value that is computed by the material models when asked for the density,
- “reference density profile”: Use the density as computed for the reference profile (which can be constant, an adiabatic profile, or an entirely different function, and is determined by the adiabatic conditions model).

2.11.3 Approximation of the buoyancy term

The buoyancy term (right-hand side of the momentum equation) always uses the density that is provided by the material model (see [Section 6.3.1](#)). Depending on the material model, this density could for example depend on temperature and pressure (such as in ALA), or on temperature and depth (as in TALA); and the model can also be set up in a way that it uses density deviations from a reference state instead of a full density (see [Section 2.4](#)).

Note: In the current version of ASPECT, it is the responsibility of the user to select a material model that is consistent with the formulation they want to use in their model. In the future, we plan to make it more obvious which approximations are supported by a particular material model.

2.11.4 Reference state: The adiabatic profile

The reference temperature profile \bar{T} , reference density profile $\bar{\rho}$ and the reference pressure \bar{p} are computed in the adiabatic conditions model (provided by the class `AdiabaticConditions`, see [Section 2.6](#)). By default, these fields satisfy adiabatic conditions (if adiabatic heating is included in the model, see [Section A.59](#)):

$$\frac{d\bar{T}(z)}{dz} = \frac{\alpha\bar{T}(z)g_z}{C_p}, \quad (33)$$

$$\frac{d\bar{p}(z)}{dz} = \bar{\rho}g_z, \quad (34)$$

$$\bar{\rho} = \bar{\rho}(\bar{p}, \bar{T}, z) \quad (\text{as defined by the material model}), \quad (35)$$

where strictly speaking g_z is the magnitude of the vertical component of the gravity vector field, but in practice we take the magnitude of the entire gravity vector. If there is no adiabatic heating in the model, \bar{T} is constant by default and set to the adiabatic surface temperature. The density gradient is always computed by a simple finite difference approximation of the depth derivative of $\bar{\rho}$.

However, users can also supply their own adiabatic conditions models or define an arbitrary profile using the “function” plugin, which allows the user to define arbitrary functions for $\bar{T}(z)$, $\bar{p}(z)$ and $\bar{\rho}(z)$, see [Section A.2](#).

2.11.5 Combined formulations

Not all combinations of the different approximations discussed above are physically reasonable, and to help users choose between these options, we provide a number of combined “Formulations” that are equivalent to the approximate equations discussed above ([Section 2.10](#)). They can be selected in the subsection `Formulation/Formulation` (see also [A.40](#)):

- “anelastic liquid approximation”: This formulation sets the mass conservation approximation to “reference density profile”, the temperature equation approximation to “reference density profile” and checks that both adiabatic and shear heating are included in the list of heating plugins used in the model, using the simplified version of the adiabatic heating term (see Section A.59). The default setting for the adiabatic conditions is an adiabatic temperature profile, and hydrostatic pressure and density profiles. This option should be chosen together with a material model that defines a density that depends on temperature and pressure (and potentially depth), which would be equivalent to the anelastic liquid approximation (Section 2.10.1), or with a material model that defines a density that depends on temperature and depth (and not on the pressure), which would be equivalent to the truncated anelastic liquid approximation (Section 2.10.2).
- “Boussinesq approximation”: This formulation sets the mass conservation approximation to “incompressible”, the temperature equation approximation to “reference density profile” and checks that neither adiabatic nor shear heating are included in the list of heating plugins used in the model. The default setting for the adiabatic conditions is a constant temperature, and hydrostatic pressure and density profiles. This option should be chosen together with a material model that defines a density that only depends on temperature and depth (and not on the pressure). This is equivalent to the Boussinesq approximation (Section 2.10.3).
- “isothermal compression”: This formulation sets the mass conservation approximation to “isothermal compression”, the temperature equation approximation to “real density” and checks that both adiabatic and shear heating are included in the list of heating plugins used in the model. The default setting for the adiabatic conditions is an adiabatic temperature profile, and hydrostatic pressure and density profiles. The density can depend on any of the solution variables. This is equivalent to the isothermal compression approximation (Section 2.10.4).
- “custom”: By default, this formulation sets the mass conservation approximation to “ask material model” and the temperature equation approximation to “real density”. The adiabatic conditions model uses an adiabatic temperature profile if adiabatic heating is included in the model, and a constant temperature if adiabatic heating is not included. Pressure and density profiles are hydrostatic. The density can depend on any of the solution variables. However, this option can also be used to arbitrarily combine the different approximations described in this section. Users should be careful when using this option, as some combinations may lead to unphysical model behaviour.

An example cookbook that shows a comparison between different approximations is discussed in Section 5.3.6.

2.12 Free surface calculations

In reality the boundary conditions of a convecting Earth are not no-slip or free slip (i.e., no normal velocity). Instead, we expect that a free surface is a more realistic approximation, since air and water should not prevent the flow of rock upward or downward. This means that we require zero stress on the boundary, or $\sigma \cdot \mathbf{n} = 0$, where $\sigma = 2\eta\varepsilon(\mathbf{u})$. In general there will be flow across the boundary with this boundary condition. To conserve mass we must then advect the boundary of the domain in the direction of fluid flow. Thus, using a free surface necessitates that the mesh be dynamically deformable.

2.12.1 Arbitrary Lagrangian-Eulerian implementation

The question of how to handle the motion of the mesh with a free surface is challenging. Eulerian meshes are well behaved, but they do not move with the fluid motions, which makes them difficult for use with free surfaces. Lagrangian meshes do move with the fluid, but they quickly become so distorted that remeshing is required. ASPECT implements an Arbitrary Lagrangian-Eulerian (ALE) framework for handling motion of the mesh. The ALE approach tries to retain the benefits of both the Lagrangian and the Eulerian approaches

by allowing the mesh motion \mathbf{u}_m to be largely independent of the fluid. The mass conservation condition requires that $\mathbf{u}_m \cdot \mathbf{n} = \mathbf{u} \cdot \mathbf{n}$ on the free surface, but otherwise the mesh motion is unconstrained, and should be chosen to keep the mesh as well behaved as possible.

ASPECT uses a Laplacian scheme for calculating the mesh velocity. The mesh velocity is calculated by solving

$$-\Delta \mathbf{u}_m = 0 \quad \text{in } \Omega, \quad (36)$$

$$\mathbf{u}_m = (\mathbf{u} \cdot \mathbf{n}) \mathbf{n} \quad \text{on } \partial\Omega_{\text{free surface}}, \quad (37)$$

$$\mathbf{u}_m \cdot \mathbf{n} = 0 \quad \text{on } \partial\Omega_{\text{free slip}}, \quad (38)$$

$$\mathbf{u}_m = 0 \quad \text{on } \partial\Omega_{\text{Dirichlet}}. \quad (39)$$

After this mesh velocity is calculated, the mesh vertices are time-stepped explicitly. This scheme has the effect of choosing a minimally distorting perturbation to the mesh. Because the mesh velocity is no longer zero in the ALE approach, we must then correct the Eulerian advection terms in the advection system with the mesh velocity (see, e.g. [DHPRF04]). For instance, the temperature equation (31) becomes

$$\rho C_p \left(\frac{\partial T}{\partial t} + (\mathbf{u} - \mathbf{u}_m) \cdot \nabla T \right) - \nabla \cdot k \nabla T = \rho H \quad \text{in } \Omega.$$

2.12.2 Free surface stabilization

Small disequilibria in the location of a free surface can cause instabilities in the surface position and result in a “sloshing” instability. This may be countered with a quasi-implicit free surface integration scheme described in [KMM10]. This scheme enters the governing equations as a small stabilizing surface traction that prevents the free surface advection from overshooting its true position at the next time step. ASPECT implements this stabilization, the details of which may be found in [KMM10].

An example of a simple model which uses a free surface may be found in Section 5.2.6.

2.13 Calculations with melt transport

The original formulation of the equations in Section 2.1 describes the movement of solid mantle material. These computations also allow for taking into account how partially molten material changes the material properties and the energy balance through the release of latent heat. However, this will not consider melt extraction or any relative movement between melt and solid and there might be problems where the transport of melt is of interest. Thus, ASPECT allows for solving additional equations describing the behavior of silicate melt percolating through and interacting with a viscously deforming host rock. This requires the advection of a compositional field representing the volume fraction of melt present at any given time (the porosity ϕ), and also a change of the mechanical part of the system. The latter is implemented using the approach of [KMK13] and changes the Stokes system to

$$-\nabla \cdot \left[2\eta \left(\varepsilon(\mathbf{u}_s) - \frac{1}{3}(\nabla \cdot \mathbf{u}_s)\mathbf{1} \right) \right] + \nabla p_f + \nabla p_c = \rho \mathbf{g} \quad \text{in } \Omega, \quad (40)$$

$$\begin{aligned} \nabla \cdot \mathbf{u}_s - \nabla \cdot K_D \nabla p_f - K_D \nabla p_f \cdot \frac{\nabla \rho_f}{\rho_f} &= -\nabla \cdot K_D \rho_f \mathbf{g} \\ &\quad + \Gamma \left(\frac{1}{\rho_f} - \frac{1}{\rho_s} \right) \\ &\quad - \frac{\phi}{\rho_f} \mathbf{u}_s \cdot \nabla \rho_f - \frac{1-\phi}{\rho_s} \mathbf{u}_s \cdot \nabla \rho_s \\ &\quad - K_D \mathbf{g} \cdot \nabla \rho_f \end{aligned} \quad \text{in } \Omega, \quad (41)$$

$$\nabla \cdot \mathbf{u}_s + \frac{p_c}{\xi} = 0. \quad (42)$$

We use the indices s to indicate properties of the solid and f for the properties of the fluid. The equations are solved for the solid velocity \mathbf{u}_s , the fluid pressure p_f , and an additional variable, the compaction pressure p_c , which is related to the fluid and solid pressure through the relation $p_c = (1-\phi)(p_s - p_f)$. K_D is the Darcy coefficient, which is defined as the quotient of the permeability and the fluid viscosity and Γ is the melting rate. η and ξ are the shear and compaction viscosities and can depend on the porosity, temperature, pressure, strain rate and composition. However, there are various laws for these quantities and so they are implemented in the material model. Common formulations for the dependence on porosity are $\eta = (1-\phi)\eta_0 e^{-\alpha_\phi \phi}$ with $\alpha_\phi \approx 25\dots30$ and $\xi = \eta_0 \phi^{-n}$ with $n \approx 1$.

To avoid the density gradients in Equation (41), which would have to be specified individually for each material model by the user, we can use the same method as for the mass conservation (described in Section 2.10.3) and assume the change in solid density is dominated by the change in static pressure, which can be written as $\nabla p_s \approx \nabla p_{\text{static}} \approx \rho_s \mathbf{g}$. This finally allows us to write

$$\frac{1}{\rho_s} \nabla \rho_s \approx \frac{1}{\rho_s} \frac{\partial \rho_s}{\partial p_s} \nabla p_s \approx \frac{1}{\rho_s} \frac{\partial \rho_s}{\partial p_s} \nabla p_s \approx \frac{1}{\rho_s} \frac{\partial \rho_s}{\partial p_s} \rho_s \mathbf{g} \approx \kappa_s \rho_s \mathbf{g}.$$

For the fluid pressure, choosing a good approximation depends on the model parameters and setup (see [DH16]). Hence, we make $\nabla \rho_f$ a model input parameter, which can be adapted based on the forces that are expected to be dominant in the model. We can then replace the second equation by

$$\begin{aligned} \nabla \cdot \mathbf{u}_s - \nabla \cdot K_D \nabla p_f - K_D \nabla p_f \cdot \frac{\nabla \rho_f}{\rho_f} &= -\nabla \cdot (K_D \rho_f \mathbf{g}) \\ &\quad + \Gamma \left(\frac{1}{\rho_f} - \frac{1}{\rho_s} \right) \\ &\quad - \frac{\phi}{\rho_f} \mathbf{u}_s \cdot \nabla \rho_f - (\mathbf{u}_s \cdot \mathbf{g})(1-\phi)\kappa_s \rho_s \\ &\quad - K_D \mathbf{g} \cdot \nabla \rho_f. \end{aligned}$$

The melt velocity is computed as

$$\mathbf{u}_f = \mathbf{u}_s - \frac{K_D}{\phi} (\nabla p_f - \rho_f g),$$

but is only used for postprocessing purposes and for computing the time step length.

Note: Here, we do not use the visco-elasto-plastic rheology of the [KMK13] formulation. Hence, we do not consider the elastic deformation terms that would appear on the right hand side of Equation (40) and Equation (42) and that include the elastic and compaction stress evolution parameters ξ_τ and ξ_p . Moreover, our viscosity parameters η and ξ only cover viscous deformation instead of combining visco-elasticity and plastic failure. This would require a modification of the rheologic law using effective shear and compaction viscosities η_{eff} and ξ_{eff} combining a failure criterion and shear and compaction visco-elasticities.

Moreover, melt transport requires an advection equation for the porosity field ϕ :

$$\rho_s \frac{\partial(1-\phi)}{\partial t} + \nabla \cdot [\rho_s(1-\phi)\mathbf{u}_s] = -\Gamma \quad \text{in } \Omega, i = 1 \dots C \quad (43)$$

In order to solve this equation in the same way as the other advection equations, we replace the second term of the equation by:

$$\nabla \cdot [\rho_s(1-\phi)\mathbf{u}_s] = (1-\phi)(\rho_s \nabla \cdot \mathbf{u}_s + \nabla \rho_s \cdot \mathbf{u}_s) - \nabla \phi \cdot \rho_s \mathbf{u}_s$$

Then we use the same method as described above and assume again that the change in density is dominated by the change in static pressure

$$\frac{1}{\rho_s} \nabla \rho_s \cdot \mathbf{u}_s \approx \kappa_s \rho_s \mathbf{g} \cdot \mathbf{u}_s$$

so we get

$$\frac{\partial \phi}{\partial t} + \mathbf{u}_s \cdot \nabla \phi = \frac{\Gamma}{\rho_s} + (1-\phi)(\nabla \cdot \mathbf{u}_s + \kappa_s \rho_s \mathbf{g} \cdot \mathbf{u}_s).$$

More details on the implementation can be found in [DH16]. A benchmark case demonstrating the propagation of solitary waves can be found in Section 5.4.14.

2.14 Nullspace removal

The Stokes equation (1) only involves symmetric gradients of the velocity, and as such the velocity is determined only up to rigid-body motions (that is to say, translations and rotations). For many simulations the boundary conditions will fully specify the velocity solution, but for some combinations of geometries and boundary conditions the solution will still be underdetermined. In the language of linear algebra, the Stokes system may have a nullspace.

Usually the user will be able to determine beforehand whether their problem has a nullspace. For instance, a model in a spherical shell geometry with free-slip boundary conditions at the top and bottom will have a rigid-body rotation in its nullspace (but not translations, as the boundary conditions do not allow flow through them). That is to say, the solver may be able to come up with a solution to the Stokes operator, but that solution plus an arbitrary rotation is also an equally valid solution.

Another example is a model in a Cartesian box with periodic boundary conditions in the x -direction, and free slip boundaries on the top and bottom. This setup has arbitrary translations along the x -axis in its nullspace, so any solution plus an arbitrary x -translation is also a solution.

A solution with some small power in these nullspace modes should not affect the physics of the simulation. However, the timestepping of the model is based on evaluating the maximum velocities in the solution, and having unnecessary motions can severely shorten the time steps that ASPECT takes. Furthermore, rigid body motions can make postprocessing calculations and visualization more difficult to interpret.

ASPECT allows the user to specify if their model has a nullspace. If so, any power in the nullspace is calculated and removed from the solution after every timestep. There are two varieties of nullspace removal implemented: removing net linear/angular momentum, and removing net translations/rotations.

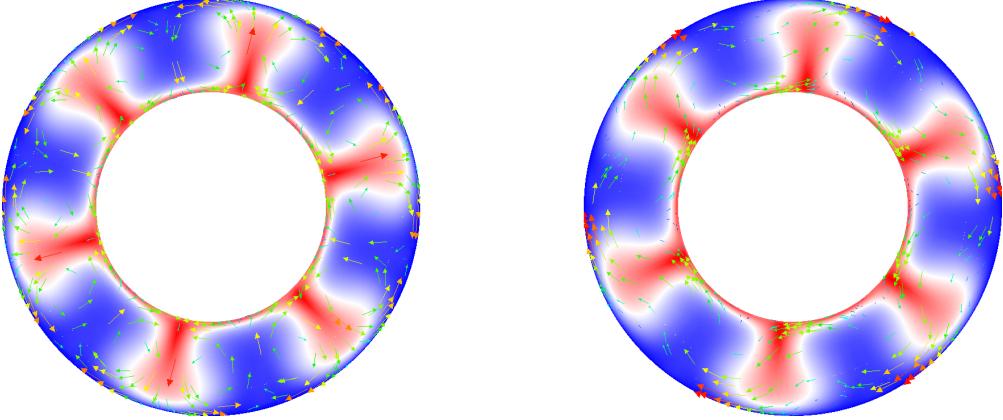


Figure 1: Example of nullspace removal. On the left the nullspace (a rigid rotation) is removed, and the velocity vectors accurately show the mantle flow. On the right there is a significant clockwise rotation to the velocity solution which is making the more interesting flow features difficult to see.

For removing linear momentum we search for a constant velocity vector \mathbf{c} such that

$$\int_{\Omega} \rho(\mathbf{u} - \mathbf{c}) = 0$$

This may be solved by realizing that $\int_{\Omega} \rho \mathbf{u} = \mathbf{p}$, the linear momentum, and $\int_{\Omega} \rho = M$, the total mass of the model. Then we find

$$\mathbf{c} = \mathbf{p}/M$$

which is subtracted off of the velocity solution.

Removing the angular momentum is similar, though a bit more complicated. We search for a rotation vector ω such that

$$\int_{\Omega} \rho(\mathbf{x} \times (\mathbf{u} - \omega \times \mathbf{x})) = 0$$

Recognizing that $\int_{\Omega} \rho \mathbf{x} \times \mathbf{u} = \mathbf{H}$, the angular momentum, and $\int_{\Omega} \rho \mathbf{x} \times \omega \times \mathbf{x} = \mathbf{I} \cdot \omega$, the moment of inertia dotted into the sought-after vector, we can solve for ω :

$$\omega = \mathbf{I}^{-1} \cdot \mathbf{H}$$

A rotation about the rotation vector ω is then subtracted from the velocity solution.

Removing the net translations/rotations are identical to their momentum counterparts, but for those the density is dropped from the formulae. For most applications the density should not vary so wildly that there will be an appreciable difference between the two varieties, though removing linear/angular momentum is more physically motivated.

The user can flag the nullspace for removal by setting the `Remove nullspace` option, as described in Section A.117. Figure 1 shows the result of removing angular momentum from a convection model in a 2D annulus with free-slip velocity boundary conditions.

2.15 Particles

ASPECT can, optionally, also deal with particles (sometimes called “tracers”). Particles can be thought of as point-like objects that are simply advected along with the flow. In other words, if $\mathbf{u}(\mathbf{x}, t)$ is the flow field that results from solving equations (1)–(2), then the k th particle’s position satisfies the equations

$$\frac{\partial}{\partial t} \mathbf{x}_k(t) = \mathbf{u}(\mathbf{x}_k(t), t). \quad (44)$$

The initial positions of all particles also need to be given and are usually either chosen randomly, based on a fixed pattern, or are read from a file.

Particles are typically used to track visually where material that starts somewhere ends up after some time of a simulation. It can also be used to track the *history* of the volume of the fluid that surrounds a particle, for example by tracking how much strain has accumulated, or what the minimal or maximal temperature may have been in the medium along the trajectory of a particle. To this end, particles can carry *properties*. These are scalar- or vector-valued quantities that are attached to each particle, that are initialized at the beginning of a simulation, and that are then updated at each time step. In other words, if we denote by $\mathbf{p}_{k,m}(t)$ the value of the m th property attached to the k th particle, then $\mathbf{p}_{k,m}(t)$ will satisfy a differential equation of the form

$$\frac{\partial}{\partial t} \mathbf{p}_{k,m}(t) = \mathbf{g}_m(\mathbf{p}_{k,m}, p(\mathbf{x}_k(t), t), T(\mathbf{x}_k(t), t), \varepsilon(\mathbf{u}(\mathbf{x}_k(t), t)), \mathbf{c}(\mathbf{x}_k(t), t)).$$

The exact form of \mathbf{g}_m of course depends on what exactly a particular property represents. Like with compositional fields (see Section 2.7), it is possible to describe the right hand side \mathbf{g}_m in ways that also allows for impulse (delta) functions in time.

How particles are used in practice is probably best explained using examples. To this end, see in particular Section 5.2.5. All particle-related input parameters are listed in Section A.125. The implementation of particles is discussed in great detail in [GHPB17].

3 Installation

There are three distinct ways to install ASPECT – compilation from source, installing a virtual machine, and using a Docker container – each providing distinct advantages and disadvantages. In this section we describe all three options and start with a summary of their properties to guide users to an informed decision about the best option for their purpose.

Feature	Compile & Install	Virtual Machine	Docker Container
Speed overhead	0%	30%	0–5%
Disk overhead	0 GB	1 GB	200 MB
Knowledge required	Much	Very Little	Little
Root privileges required	No	No (installed VM software)	Partially
Embedded in native environment	Yes	No	Partially
MacOS support	Yes	Yes	Yes
Windows support	No	Yes	Yes
Local parallelization	Yes	Yes	Yes
Massively parallel computations	Yes	No	No
Modifying ASPECT	Possible	Possible	Possible
Configuring dependencies	Possible	No	No

Table 1: Features of the different installation options of ASPECT.

The available options can be best presented in form of typical use cases:

1. Virtual Machine (ASPECT beginner and tutorial participant): The virtual machine image provides a fully prepared user environment that contains installations of ASPECT, all required libraries, and visualization software on top of a full Linux environment. This way beginning users and tutorial participants can work in a unified environment, thus minimizing installation time and technical problems. Due to the overhead of virtualizing a full operating system this installation typically needs more space, and is approximately 30 % slower than a native installation. Additionally working in a virtual machine ‘feels’ differently from working in your usual desktop environment. The virtual machine can be run on

all host operating systems that can run a virtualization software like VirtualBox (e.g. Linux, Apple MacOS, Microsoft Windows).

2. Docker Container (advanced user with no need to configure/change the underlying libraries, possibly changing parts of ASPECT): Docker containers are lightweight packages that only encapsulate the minimal dependencies to run an application like ASPECT on top of the host operating system. They allow easy installation and usage of ASPECT in a unified environment, while relying on the user's operating system to provide visualization software and model input data. When compared to the virtual machine it is simple to exchange files between the host operating system and the docker container, and it provides the benefit to work in the desktop environment you are used to. They have very little overhead in terms of memory and speed compared to virtual machines, and allow for reproducible computations. The container is set up with a standard ASPECT installation, but this can be modified by advanced users (source code development within the container is possible).
3. Compile & Install (advanced users and developers with the need to reconfigure underlying libraries or running massively parallel models): The most advanced option is to compile and install ASPECT from source. This allows maximal control over the underlying libraries like TRILINOS and DEAL.II, as well as easy modifications to ASPECT by recompiling a modified source directory. Our installation instructions cover most Linux and MacOS operating systems, but incompatibilities on individual systems can always occur and make the installation more cumbersome. If you are planning to run massively parallel computations on a compute cluster this is likely your only option. Since clusters usually have a very individual setup, it is always a good idea to ask IT support staff for help when installing ASPECT, to avoid hard to reproduce setup problems, and performance penalties.

3.1 Docker Container

3.1.1 Installing Docker and downloading the ASPECT image

Docker is a lightweight virtualization software that allows to ship applications with all their dependencies in a simple way. It is outside of the scope of this manual to explain all possible applications of Docker, and we refer to the introduction (<https://www.docker.com/what-docker>) and installation and quickstart guides (<https://www.docker.com/products/docker>) on the Docker website for more detailed descriptions of how to set up and use the docker engine. More importantly Docker provides a marketplace for exchanging prepared docker images (called Docker Hub). After setting up the docker engine downloading a precompiled ASPECT image from Docker Hub is as simple as typing in a terminal:

```
docker pull gassmoeller/aspect
```

Note that the transfer size of the compressed image containing ASPECT and all its dependencies is about 900 MB. When extracted the image requires about 3.2 GB of disk space.

3.1.2 Running ASPECT models

Although it is possible to use the downloaded ASPECT docker image in a number of different ways, we recommend the following workflow:

1. Create your ASPECT input file in a folder of your choice (possibly also containing any input data that is required by your model) and navigate in a terminal into that directory.
2. Run the docker image and mount the current directory as a read-only volume into the docker container⁵. This is accomplished by specifying the -v flag followed by the absolute path on the host machine, colon,

⁵Note that it is possible to mount a directory as writeable into the container. However, this is often associated with file permission conflicts between the host system and the container. Therefore, we recommend this slightly more cumbersome, but also more reliable workflow.

absolute path within the docker container, colon, and specifying read-only permissions as in the example below.

Make sure your parameter file specifies a model output directory *other* than the input directory, e.g. `/home/dealii/aspect/model_output`. When you have started the container run the aspect model inside the container. Note that there are two ASPECT executables in the work directory of the container: `aspect` and `aspect-release`. For a discussion of the different versions see Section 4.3, in essence: You should run `aspect` first to check your model for errors, then run `aspect-release` for a faster model run.

To sum up, the steps you will want to execute are:

```
docker run -it -v "$(pwd):/home/dealii/aspect/model_input:ro" \
gassmoeller/aspect:latest bash
```

Within the container, simply run your model by executing:

```
./aspect model_input/your_input_file.prm
```

3. After the model has finished (or during the model run if you want to check intermediate results) copy the model output out of the container into your current directory. For this you need to find the name or ID of the docker container by running `docker ps -a` in a separate terminal first. Look for the most recently started container to identify your current ASPECT container.

Commands that copy the model output to the current directory could be:

```
docker ps -a # Find the name of the running / recently closed container in the output
docker cp CONTAINER_NAME:/home/dealii/aspect/model_output .
```

4. The output data is saved inside your container even after the computation finishes and even when you stop the container. After you have copied the data out of the container you should therefore delete the container to avoid duplication of output data. Even after deleting you will always be able to start a new container from the downloaded image following step 2. Deleting the finished container can be achieved by the `docker container prune` command that removes any container that is not longer running.

Note: If you own other finished containers that you want to keep use `docker container rm CONTAINER_NAME` to only remove the container named `CONTAINER_NAME`.

To remove all finished containers use the following command:

```
docker container prune
```

Alternatively only remove a particular container:

```
docker container rm CONTAINER_NAME
```

You are all set. Repeat steps 1-4 of this process as necessary when updating your model parameters.

3.1.3 Developing ASPECT within a container

The above given workflow does not include advice on how to modify ASPECT inside the container. We recommend a slightly different workflow for advanced users that want to modify parts of ASPECT. The ASPECT docker container itself is build on top of a DEAL.II container that contains all dependencies for compiling ASPECT. Therefore it is possible to run the deal.II container, mount an ASPECT source directory from your host system and compile it inside of the container. An example workflow could look as following (assuming you navigated in a terminal into the modified ASPECT source folder):

```
docker pull dealii/dealii:v8.5.pre.4-gcc-mpi-fulldepsmanual-debugrelease
docker run -it -v "$(pwd):/home/dealii/aspect:ro" \
    dealii/dealii:v8.5.pre.4-gcc-mpi-fulldepsmanual-debugrelease bash
```

Inside of the container you now find a read-only ASPECT directory that contains your modified source code. You can compile and run a model inside the container, e.g. in the following way:

```
mkdir aspect-build
cd aspect-build
cmake -DCMAKE_BUILD_TYPE=Debug -DDEAL_II_DIR=$HOME/deal.II-install $HOME/aspect
./aspect $HOME/aspect/cookbooks/shell_simple_2d.prm
```

To avoid repeated recompilations of the ASPECT source folder we recommend to reuse the so prepared container instead of starting new containers based on the DEAL.II image. This can be achieved by the following commands outside of the container:

```
docker ps -a # Find the name of the running / recently closed container in the output
docker restart CONTAINER_NAME
docker attach CONTAINER_NAME
```

For more information on the differences between using images and containers, and how to attach additional terminals to a running container, we refer to the docker documentation (e.g. https://docs.docker.com/engine/getstarted/step_two/).

3.2 Virtual Machine

3.2.1 Installing VM software and setting up the virtual machine

The ASPECT project provides an experimental virtual machine containing a fully configured version of ASPECT. To use this machine, you will need to install VirtualBox (<http://www.virtualbox.org/>) on your machine, and then import a virtual machine image that can be downloaded from <http://www.math.clemson.edu/~heister/dealvm/>. Note, however, that the machine image is several gigabytes in size and downloading will take a while. After downloading and installing the virtual image it is convenient to set up a shared folder between your host system and the virtual machine to exchange model files and outputs.

3.2.2 Running ASPECT models

The internal setup of the virtual machine is similar to the Docker container discussed above, except that it contains a full-featured desktop environment. Also note that the user name is `ubuntu`, not `dealii` as in the Docker container. Again there are multiple ways to use the virtual machine, but we recommend the following workflow:

1. Create your ASPECT input file in the shared folder and start the virtual machine.
2. Navigate in a terminal to your model directory.
3. Run your model using the provided ASPECT executable:

```
~/aspect/aspect your_input_file.prm
```

4. The model output should automatically appear on your host machine in the shared directory.
5. After you have verified that your model setup is correct, you might want to consider recompiling ASPECT in release mode to increase the speed of the computation. See Section 4.3 for a discussion of debug and release mode.
6. Visualize your model output either inside of the virtual machine (ParaView and VisIt are pre-installed), or outside on your host system.

You are all set. Repeat steps 1-6 of this process as necessary when updating your model parameters.

3.3 Local installation

This is a brief explanation of how to compile and install the required dependencies and ASPECT itself. This installation procedure guarantees fastest runtimes, and largest flexibility, but usually requires more work than the options mentioned in the previous sections. While it is possible to install ASPECT's dependencies in particular P4EST, TRILINOS, and DEAL.II manually, we recommend to use the `candi` software (see <https://github.com/dealii/candi>). `candi` was written as an installation program for deal.II, and includes a number of system specific instructions that will be listed when starting the program. It can be flexibly configured to allow for non-default compilers or libraries (e.g. Intel's MKL instead of LAPACK) by changing entries in the configuration file `candi.cfg`, or by providing platform specific installation files.

In case you encounter problems during the installation, please consult our wiki (<https://github.com/geodynamics/aspect/wiki>) for frequently asked questions and special instructions for MacOS users, before posting your questions on the mailing list.

3.3.1 System prerequisites

`candi` will show system specific instructions on startup, but its prerequisites are relatively widely used and packaged for most operating systems. You will need compilers for C, C++ and Fortran, the GNU make system, the CMake build system, and the libraries and header files of BLAS, LAPACK and zlib, which is used for compressing the output data. To use more than one process for your computations you will need to install a MPI library, its headers and the necessary executables to run MPI programs. There are some optional packages for additional features, like the HDF5 libraries for additional output formats, PETSc for alternative solvers, and Numdiff for checking ASPECT's test results with reasonable accuracy, but these are not strictly required, and in some operating systems they are not available as packages but need to be compiled from scratch. Finally, for obtaining a recent development version of ASPECT you will need the git version control system.

An exemplary command to obtain all required packages on Ubuntu 14.04 would be:

```
sudo apt-get install build-essential \
    cmake \
    gcc \
    g++ \
    gfortran \
    git \
    libblas-dev \
    liblapack-dev \
    libopenmpi-dev \
    numdiff \
    openmpi-bin \
    zlib1g-dev
```

3.3.2 Using candi to compile dependencies

In its default configuration `candi` downloads and compiles a DEAL.II configuration that is able to run ASPECT, but it also contains a number of packages that are not required (and that can be safely disabled if problems occur during the installation). We require at least the packages P4EST, TRILINOS (or as an experimental alternative PETSC), and finally DEAL.II.

At the time of this writing `candi` will install P4EST 2.0, TRILINOS 12.10.1, PETSc 3.6.4, and DEAL.II 8.5.0. We strive to keep the development version of ASPECT compatible with the latest release of DEAL.II and the current DEAL.II development version at any time, and we usually support several older versions of P4EST, TRILINOS, and PETSc.

1. *Obtaining candi:* Download `candi` by running

```
git clone https://github.com/dealii/candi
```

in a directory of your choice.

2. *Installing DEAL.II and its dependencies:* Execute `candi` by running

```
cd candi  
./candi.sh -p INSTALL_PATH
```

(here we assume you replace `INSTALL_PATH` by the path where you want to install all dependencies and DEAL.II, typically a directory inside `$HOME/bin` or a similar place). This step might take a long time, but can be parallelized by adding `-jN`, where `N` is the number of CPU cores available on your computer. Further configuration options and parameters are listed at <https://github.com/dealii/candi>.

3. You may now want to configure your environment to make it aware of the newly installed packages. This can be achieved by adding the line `source INSTALL_PATH/configuration/enable.sh` to the file responsible for setting up your shell environment⁶ (again we assume you replace `INSTALL_PATH` by the patch chosen in the previous step). Then close the terminal and open it again to activate the change.
4. *Testing your installation:* Test that your installation works by compiling the `step-32` example that you can find in `$DEAL_II_DIR/examples/step-32`. Prepare and compile by running `cmake . && make` and run with `mpirun -n 2 ./step-32`.

Congratulations, you are now set up for compiling ASPECT itself.

3.3.3 Obtaining ASPECT and initial configuration

The development version of ASPECT can be downloaded by executing the command

```
git clone https://github.com/geodynamics/aspect.git
```

If `$DEAL_II_DIR` points to your DEAL.II installation, you can configure ASPECT by running

```
cmake .
```

in the ASPECT directory created by the `git clone` command above. If you did not set `$DEAL_II_DIR` you have to supply `cmake` with the location:

```
cmake -DDEAL_II_DIR=/u/username/deal-installed/ .
```

An alternative would be to configure ASPECT as an out-of-source build. You would need to create a separate build directory and specify ASPECT's source directory using `cmake PATH_TO_ASPECT_SOURCE` from within the build directory. The instructions in the following sections assume an in-source build.

⁶For bash this would be the file `~/.bashrc`.

3.3.4 Compiling ASPECT and generating documentation

After downloading ASPECT and having built the libraries it builds on, you can compile it by typing

```
make
```

on the command line (or `make -jN` if you have multiple processors in your machine, where N is the number of processors). This builds the ASPECT executable which will reside in the main directory and will be named `./aspect`. If you intend to modify ASPECT for your own experiments, you may want to also generate documentation about the source code. This can be done using the command

```
cd doc; make
```

which assumes that you have the `doxygen` documentation generation tool installed. Most Linux distributions have packages for `doxygen`. The result will be the file `doc/doxygen/index.html` that is the starting point for exploring the documentation.

4 Running ASPECT

4.1 Overview

After compiling ASPECT as described above, you should have an executable file in the main directory. It can be called as follows:

```
./aspect parameter-file.prm
```

or, if you want to run the program in parallel, using something like

```
mpirun -np 32 ./aspect parameter-file.prm
```

to run with 32 processors. In either case, the argument denotes the (path and) name of a file that contains input parameters.⁷ When you download ASPECT, there are a number of sample input files in the `cookbooks` directory, corresponding to the examples discussed in Section 5, and input files for some of the benchmarks discussed in Section 5.4 are located in the `benchmarks` directory. A full description of all parameters one can specify in these files is given in Section A.

Running ASPECT with an input file will produce output that will look something like this (numbers will all be different, of course):

```
-----  
-- This is ASPECT, the Advanced Solver for Problems in Earth's ConvecTion.  
-- . version 2.0.0-pre (include_dealii_version, c20eba0)  
-- . using deal.II 9.0.0-pre (master, 952baa0)  
-- . using Trilinos 12.10.1  
-- . using p4est 2.0.0  
-- . running in DEBUG mode  
-- . running with 1 MPI process  
-----
```

⁷As a special case, if you call ASPECT with an argument that consists of two dashes, “-”, then the arguments will be read from the standard input stream of the program. In other words, you could type the input parameters into your shell window in this case (though that would be cumbersome, ASPECT would seem to hang until you finish typing all of your input into the window and then terminating the input stream by typing Ctrl-D). A more common case would be to use Unix pipes so that the default input of ASPECT is the output of another program, as in a command like `cat parameter-file.prm.in | mypreprocessor | ./aspect -`, where `mypreprocessor` would be a program of your choice that somehow transforms the file `parameter-file.prm.in` into a valid input file, for example to systematically vary one of the input parameters.

If you want to run ASPECT in parallel, you can do something like `cat parameter-file.prm.in | mypreprocessor | mpirun -np 4 ./aspect -`. In cases like this, `mpirun` only forwards the output of `mypreprocessor` to the first of the four MPI processes, which then sends the text to all other processors.

```

Number of active cells: 1,536 (on 5 levels)
Number of degrees of freedom: 20,756 (12,738+1,649+6,369)

*** Timestep 0: t=0 years

Rebuilding Stokes preconditioner...
Solving Stokes system... 30+3 iterations.
Solving temperature system... 8 iterations.

Number of active cells: 2,379 (on 6 levels)
Number of degrees of freedom: 33,859 (20,786+2,680+10,393)

*** Timestep 0: t=0 years

Rebuilding Stokes preconditioner...
Solving Stokes system... 30+4 iterations.
Solving temperature system... 8 iterations.

Postprocessing:
Writing graphical output: output/solution/solution-00000
RMS, max velocity: 0.0946 cm/year, 0.183 cm/year
Temperature min/avg/max: 300 K, 3007 K, 6300 K
Inner/outer heat fluxes: 1.076e+05 W, 1.967e+05 W

*** Timestep 1: t=1.99135e+07 years

Solving Stokes system... 30+3 iterations.
Solving temperature system... 8 iterations.

Postprocessing:
Writing graphical output: output/solution/solution-00001
RMS, max velocity: 0.104 cm/year, 0.217 cm/year
Temperature min/avg/max: 300 K, 3008 K, 6300 K
Inner/outer heat fluxes: 1.079e+05 W, 1.988e+05 W

*** Timestep 2: t=3.98271e+07 years

Solving Stokes system... 30+3 iterations.
Solving temperature system... 8 iterations.

Postprocessing:
RMS, max velocity: 0.111 cm/year, 0.231 cm/year
Temperature min/avg/max: 300 K, 3008 K, 6300 K
Inner/outer heat fluxes: 1.083e+05 W, 2.01e+05 W

*** Timestep 3: t=5.97406e+07 years

...

```

The output starts with a header that lists the used ASPECT, DEAL.II, TRILINOS and P4EST versions as well as the mode you compiled ASPECT in (see 4.3), and the number of parallel processes used⁸. With

⁸If you used the git version control system to download ASPECT and/or DEAL.II, as in this example, you will also get the current branch, and unique revision identifier for the current version. This is very important if you modify either software between releases, or you use a development version that is not an official release. Note that this revision can not track changes you made to the software that are not part of a git commit.

this information we strive to make ASPECT models as reproducible as possible.

The following output depends on the model, and in this case was produced by a parameter file that, among other settings, contained the following values (we will discuss many such input files in Section 5):

```

set Dimension          = 2
set End time           = 1.5e9
set Output directory   = output

subsection Geometry model
    set Model name        = spherical shell
end

subsection Mesh refinement
    set Initial global refinement = 4
    set Initial adaptive refinement = 1
end

subsection Postprocess
    set List of postprocessors = visualization, velocity statistics, temperature statistics,
                                ↪ heat flux statistics, depth average
end

```

In other words, these run-time parameters specify that we should start with a geometry that represents a spherical shell (see Sections A.42 and A.51 for details). The coarsest mesh is refined 4 times globally, i.e., every cell is refined into four children (or eight, in 3d) 4 times. This yields the initial number of 1,536 cells on a mesh hierarchy that is 5 levels deep. We then solve the problem there once and, based on the number of adaptive refinement steps at the initial time set in the parameter file, use the solution so computed to refine the mesh once adaptively (yielding 2,379 cells on 6 levels) on which we start the computation over at time $t = 0$.

Within each time step, the output indicates the number of iterations performed by the linear solvers, and we generate a number of lines of output by the postprocessors that were selected (see Section A.118). Here, we have selected to run all postprocessors that are currently implemented in ASPECT which includes the ones that evaluate properties of the velocity, temperature, and heat flux as well as a postprocessor that generates graphical output for visualization.

While the screen output is useful to monitor the progress of a simulation, its lack of a structured output makes it not useful for later plotting things like the evolution of heat flux through the core-mantle boundary. To this end, ASPECT creates additional files in the output directory selected in the input parameter file (here, the `output/` directory relative to the directory in which ASPECT runs). In a simple case, this will look as follows:

```

aspect> ls -l output/
total 932
-rw-rw-r-- 1 bangerth bangerth 11134 Dec 11 10:08 depth_average.gnuplot
-rw-rw-r-- 1 bangerth bangerth 11294 Dec 11 10:08 log.txt
-rw-rw-r-- 1 bangerth bangerth 326074 Dec 11 10:07 parameters.prm
-rw-rw-r-- 1 bangerth bangerth 577138 Dec 11 10:07 parameters.tex
drwxr-xr-x 2 bangerth bangerth 4096 Dec 11 10:08 solution
-rw-rw-r-- 1 bangerth bangerth 484 Dec 11 10:08 solution.pvd
-rw-rw-r-- 1 bangerth bangerth 451 Dec 11 10:08 solution.visit
-rw-rw-r-- 1 bangerth bangerth 8267 Dec 11 10:08 statistics

```

The purpose of these files is as follows:

- *A listing of all run-time parameters:* The `output/parameters.prm` file contains a complete listing of all run-time parameters. In particular, this includes the ones that have been specified in the input parameter file passed on the command line, but it also includes those parameters for which defaults

have been used. It is often useful to save this file together with simulation data to allow for the easy reproduction of computations later on.

There is a second file, `output/parameters.prm`, that lists these parameters in L^AT_EX format.

- *Graphical output files:* One of the postprocessors chosen in the parameter file used for this computation is the one that generates output files that represent the solution at certain time steps. The screen output indicates that it has run at time step 0, producing output files that start with `output/solution/solution-00000`. Depending on the settings in the parameter file, output will be generated every so many seconds or years of simulation time, and subsequent output files will then start with `output/solution/solution-00001`, all placed in the `output/solution` subdirectory. This is because there are often *a lot* of output files: For many time steps, times the number of processors, so they are placed in a subdirectory so as not to make it more difficult than necessary to find the other files.

At the current time, the default is that ASPECT generates this output in VTK format⁹ as that is widely used by a number of excellent visualization packages and also supports parallel visualization.¹⁰ If the program has been run with multiple MPI processes, then the list of output files will be `output/solution/solution-XXXXX.YYYY` denoting that this the XXXXXth time we create output files and that the file was generated by the YYYYth processor.

VTK files can be visualized by many of the large visualization packages. In particular, the [Visit](#) and [ParaView](#) programs, both widely used, can read the files so created. However, while VTK has become a de-facto standard for data visualization in scientific computing, there doesn't appear to be an agreed upon way to describe which files jointly make up for the simulation data of a single time step (i.e., all files with the same XXXXX but different YYYY in the example above). Visit and Paraview both have their method of doing things, through `.pvtu` and `.visit` files. To make it easy for you to view data, ASPECT simply creates both kinds of files in each time step in which graphical data is produced, and these are then also placed into the subdirectories as `output/solution/solution-XXXXX.pvtu` and `output/solution/solution-XXXXX.visit`.

The final two files of this kind, `output/solution.pvd` and `output/solution.visit`, are files that describes to Paraview and Visit, respectively, which `output/solution/solution-XXXXX.pvtu` and `output/solution/solution-XXXXX.YYYY.vtu` jointly form a complete simulation. In the former case, the file lists the `.pvtu` files of all timesteps together with the simulation time to which they correspond. In the latter case, it actually lists all `.vtu` that belong to one simulation, grouped by the timestep they correspond to. To visualize an entire simulation, not just a single time step, it is therefore simplest to just load one of these files, depending on whether you use Paraview or Visit.¹¹ Because loading an *entire* simulation is the most common use case, these are the two files you will most often load, and so they are placed in the `output` directory, not the subdirectory where the actual `.vtu` data files are located.

For more on visualization, see also Section 4.4.

- *A statistics file:* The `output/statistics` file contains statistics collected during each time step, both from within the simulator (e.g., the current time for a time step, the time step length, etc.) as well as from the postprocessors that run at the end of each time step. The file is essentially a table that

⁹The output is in fact in the VTU version of the VTK file format. This is the XML-based version of this file format in which contents are compressed. Given that typical file sizes for 3d simulation are substantial, the compression saves a significant amount of disk space.

¹⁰The underlying DEAL.II package actually supports output in around a dozen different formats, but most of them are not very useful for large-scale, 3d, parallel simulations. If you need a different format than VTK, you can select this using the run-time parameters discussed in Section A.138.

¹¹At the time of writing this, current versions of Visit (starting with version 2.5.1) actually have a bug that prevents them from successfully reading the `output/solution.visit` or `output/solution/solution-XXXXX.visit` files – Visit believes that each of these files corresponds to an individual time step, rather than that a whole group of files together form one time step. This bug is not fixed in Visit 2.6.3, but may be fixed in later versions.

allows for the simple production of time trends. In the example above, and at the time when we are writing this section, it looks like this:

```
# 1: Time step number
# 2: Time (years)
# 3: Iterations for Stokes solver
# 4: Time step size (year)
# 5: Iterations for temperature solver
# 6: Visualization file name
# 7: RMS velocity (m/year)
# 8: Max. velocity (m/year)
# 9: Minimal temperature (K)
# 10: Average temperature (K)
# 11: Maximal temperature (K)
# 12: Average nondimensional temperature (K)
# 13: Core-mantle heat flux (W)
# 14: Surface heat flux (W)
0 0.000e+00 33 2.9543e+07 8           "" 0.0000 0.0000 0.0000 0.0000 ...
0 0.000e+00 34 1.9914e+07 8 output/solution/solution-00000 0.0946 0.1829 300.00 3007.2519 ...
1 1.991e+07 33 1.9914e+07 8 output/solution/solution-00001 0.1040 0.2172 300.00 3007.8406 ...
2 3.982e+07 33 1.9914e+07 8           "" 0.1114 0.2306 300.00 3008.3939 ...
```

The actual columns you have in your statistics file may differ from the ones above, but the format of this file should be obvious. Since the hash mark is a comment marker in many programs (for example, `gnuplot` ignores lines in text files that start with a hash mark), it is simple to plot these columns as time series. Alternatively, the data can be imported into a spreadsheet and plotted there.

Note: As noted in Section 2.3, ASPECT can be thought of as using the meter-kilogram-second (MKS, or SI) system. Unless otherwise noted, the quantities in the output file are therefore also in MKS units.

A simple way to plot the contents of this file is shown in Section 4.4.2.

- *Output files generated by other postprocessors:* Similar to the `output/statistics` file, several of the existing postprocessors one can select from the parameter file generate their data in their own files in the output directory. For example, ASPECT’s “depth average” postprocessor will write depth-average statistics into the file `output/depth_average.gnuplot`. Input parameters chosen in the input file control how often this file is updated by the postprocessor, as well as what graphical file format to use (if anything other than `gnuplot` is desired).

By default, the data is written in text format that can be easily visualized, see for example Figure 2. The plot shows how an initially linear temperature profile forms upper and lower boundary layers.

There are other parts of ASPECT that may also create files in the output directory. For example, if your simulation includes advecting along particles (see Section 2.15), then visualization information for these particles will also appear in this file. See Section 5.2.5 for an example of how this looks like.

4.2 Selecting between 2d and 3d runs

ASPECT can solve both two- and three-dimensional problems.¹² You select which one you want by putting a line like the following into the parameter file (see Section A):

¹²For a description of what exactly we mean when we consider two-dimensional models, see Section 2.1.3.

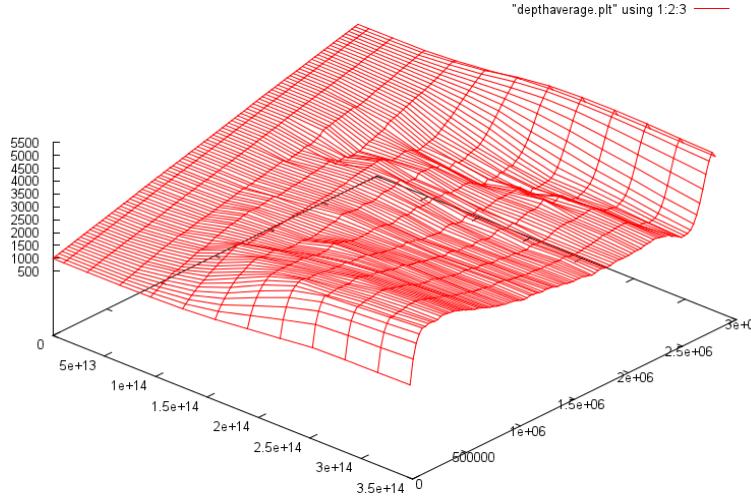


Figure 2: Example output for depth average statistics. On the left axis are 13 time steps, on the right is the depth (from the top at 0 to the bottom of the mantle on the far right), and the upwards pointing axis is the average temperature. This plot is generated by gnuplot, but the depth averages can be written in many other output formats as well, if preferred (see Section A.120).

```
set Dimension
```

```
= 2
```

Internally, dealing with the dimension builds on a feature in DEAL.II, upon which ASPECT is based, that is called *dimension-independent programming*. In essence, what this does is that you write your code only once in a way so that the space dimension is a variable (or, in fact, a template parameter) and you can compile the code for either 2d or 3d. The advantage is that codes can be tested and debugged in 2d where simulations are relatively cheap, and the same code can then be re-compiled and executed in 3d where simulations would otherwise be prohibitively expensive for finding bugs; it is also a useful feature when scoping out whether certain parameter settings will have the desired effect by testing them in 2d first, before running them in 3d. This feature is discussed in detail in the [DEAL.II tutorial program step-4](#). Like there, all the functions and classes in ASPECT are compiled for both 2d and 3d. Which dimension is actually called internally depends on what you have set in the input file, but in either case, the machine code generated for 2d and 3d results from the same source code and should, thus, contain the same set of features and bugs. Running in 2d and 3d should therefore yield comparable results. Be prepared to wait much longer for computations to finish in the latter case, however.

4.3 Debug or optimized mode

ASPECT utilizes a DEAL.II feature called *debug mode*. By default, ASPECT uses debug mode, i.e., it calls a version of the DEAL.II library that contain lots of checks for the correctness of function arguments, the consistency of the internal state of data structure, etc. If you program with DEAL.II, for example to extend ASPECT, it has been our experience over the years that, by number, most programming errors are of the kind where one forgets to initialize a vector, one accesses data that has not been updated, one tries to write into a vector that has ghost elements, etc. If not caught, the result of these bugs is that parts of the program use invalid data (data written into ghost elements is not communicated to other processors), that operations simply make no sense (adding vectors of different length), that memory is corrupted (writing past the end of an array) or, in rare and fortunate cases, that the program simply crashes.

Debug mode is designed to catch most of these errors: It enables some 7,300 assertions (as of late 2011) in DEAL.II where we check for errors like the above and, if the condition is violated, abort the program

with a detailed message that shows the failed check, the location in the source code, and a stacktrace how the program got there. The downside of debug mode is, of course, that it makes the program much slower – depending on application by a factor of 4–10. An example of the speedup one can get is shown in Section 5.2.1.

ASPECT by default uses debug mode because most users will want to play with the source code, and because it is also a way to verify that the compilation process worked correctly. If you have verified that the program runs correctly with your input parameters, for example by letting it run for the first 10 time steps, then you can switch to optimized mode by compiling ASPECT with the command¹³

```
make release
```

and then compile using

```
make
```

To switch back to debug mode type:

```
make debug
```

Note: It goes without saying that if you make significant modifications to the program, you should do the first runs in debug mode to verify that your program still works as expected.

4.4 Visualizing results

Among the postprocessors that can be selected in the input parameter file (see Sections 4.1 and A.138) are some that can produce files in a format that can later be used to generate a graphical visualization of the solution variables \mathbf{u} , p and T at select time steps, or of quantities derived from these variables (for the latter, see Section 6.3.9).

By default, the files that are generated are in VTU format, i.e., the XML-based, compressed format defined by the VTK library, see <http://public.kitware.com/VTK/>. This file format has become a broadly accepted pseudo-standard that many visualization program support, including two of the visualization programs used most widely in computational science: Visit (see <https://visit.llnl.gov/>) and ParaView (see <http://www.paraview.org/>). The VTU format has a number of advantages beyond being widely distributed:

- It allows for compression, keeping files relatively small even for sizeable computations.
- It is a structured XML format, allowing other programs to read it without too much trouble.
- It has a degree of support for parallel computations where every processor would only write that part of the data to a file that this processor in fact owns, avoiding the need to communicate all data to a single processor that then generates a single file. This requires a master file for each time step that then contains a reference to the individual files that together make up the output of a single time step. Unfortunately, there doesn't appear to be a standard for these master records; however, both ParaView and Visit have defined a format that each of these programs understand and that requires placing a file with ending .pvtu or .visit into the same directory as the output files from each processor. Section 4.1 gives an example of what can be found in the output directory.

Note: You can select other formats for output than VTU, see the run-time parameters in Section A.138. However, none of the numerous formats currently implemented in DEAL.II other than the VTK/VTU formats allows for splitting up data over multiple files in case of parallel computations, thus making subsequent visualization of the entire volume impossible. Furthermore, given the amount of data ASPECT can produce, the compression that is part of the VTU format is an important part of keeping data manageable.

¹³Note that this procedure also changed with the switch to cmake.

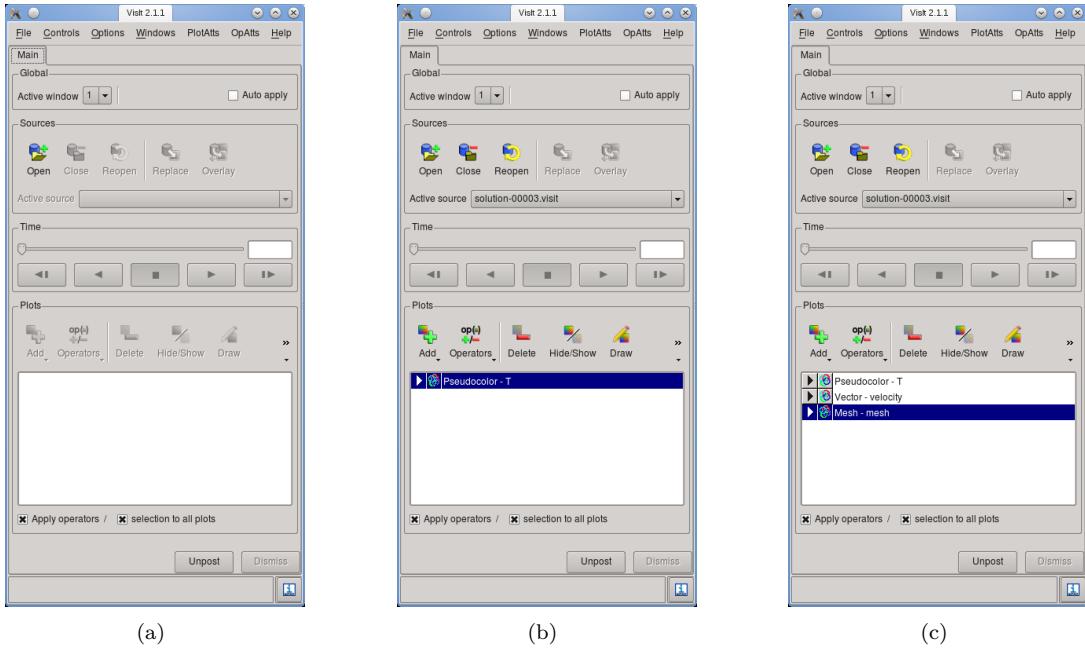


Figure 3: Main window of *Visit*, illustrating the different steps of adding content to a visualization.

4.4.1 Visualization the graphical output using *Visit*

In the following, let us discuss the process of visualizing a 2d computation using *Visit*. The steps necessary for other visualization programs will obviously differ but are, in principle, similar.

To this end, let us consider a simulation of convection in a box-shaped, 2d region (see the “cookbooks” section, Section 5, and in particular Section 5.2.1 for the input file for this particular model). We can run the program with 4 processors using

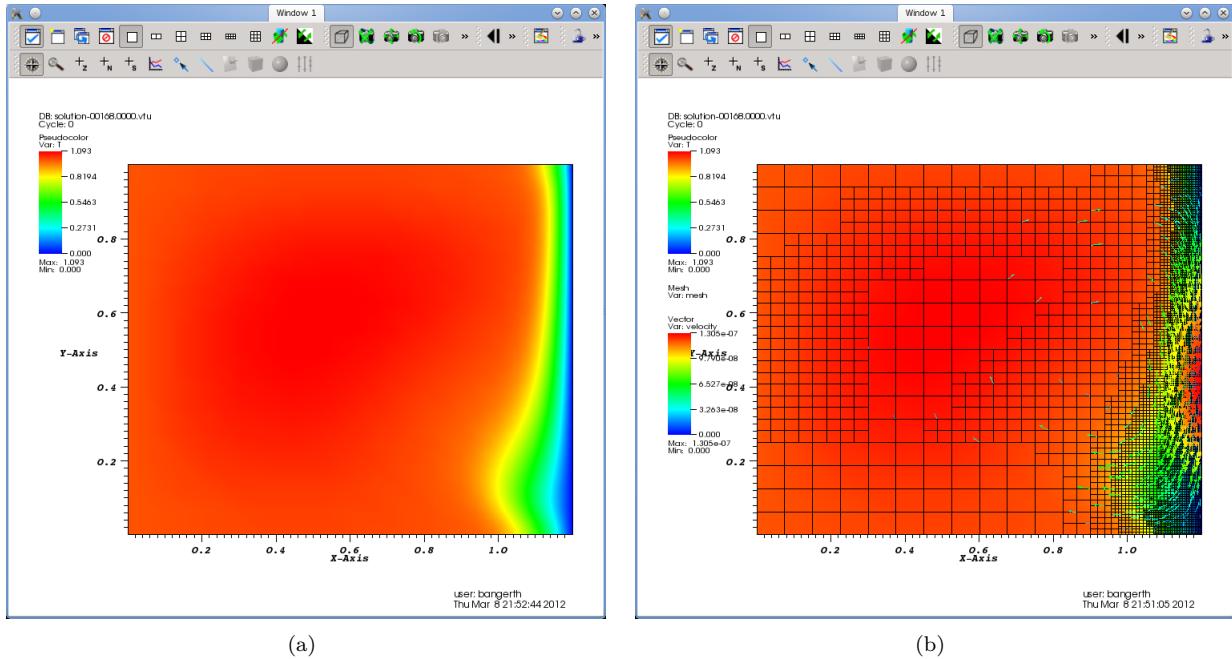
```
mpirun -np 4 ./aspect cookbooks/convection-box.prm
```

Letting the program run for a while will result in several output files as discussed in Section 4.1 above.

In order to visualize one time step, follow these steps:¹⁴

- *Selecting input files:* As mentioned above, in parallel computations we usually generate one output file per processor in each time step for which visualization data is produced (see, however, Section 4.4.3). To tell *Visit* which files together make up one time step, ASPECT creates a `output/solution/solution-XXXX.visit` file in the output directory. To open it, start *Visit*, click on the “Open” button in the “Sources” area of its main window (see Fig. 3(a)) and select the file you want. Alternatively, you can also select files using the “File > Open” menu item, or hit the corresponding keyboard short-cut. After adding an input source, the “Sources” area of the main window should list the selected file name. More easily, you can also just open `output/solution.visit` which references *all* output files for all time steps. If you open this, *Visit* will display a slider that allows you to select which time step you want to visualize, along with forward, backward, and play buttons that allow you to move between time steps.
- *Selecting what to plot:* ASPECT outputs all sorts of quantities that characterize the solution, such as temperature, pressure, velocity, and many others on demand (see Section A.138). Once an input file has been opened, you will want to add graphical representations of some of this data to the still empty

¹⁴The instructions and screenshots were generated with *Visit* 2.1. Later versions of *Visit* differ slightly in the arrangement of components of the graphical user interface, but the workflow and general idea remains unchanged.



(a)

(b)

Figure 4: *Display window of Visit, showing a single plot and one where different data is overlaid.*

canvas. To this end, click on the “Add” button of the “Plots” area. The resulting menu provides a number of different kinds of plots. The most important for our purpose are: (i) “Pseudocolor” allows the visualization of a scalar field (e.g., temperature, pressure, density) by using a color field. (ii) “Vector” displays a vector-valued field (e.g., velocity) using arrows. (iii) “Mesh” displays the mesh. The “Contour”, “Streamline” and “Volume” options are also frequently useful, in particular in 3d.

Let us choose the “Pseudocolor” item and select the temperature field as the quantity to plot. Your main window should now look as shown in Fig. 3(b). Then hit the “Draw” button to make Visit generate data for the selected plots. This will yield a picture such as shown in Fig. 4(a) in the display window of Visit.

- *Overlaying data:* Visit can overlay multiple plots in the same view. To this end, add another plot to the view using again the “Add” button to obtain the menu of possible plots, then the “Draw” button to actually draw things. For example, if we add velocity vectors and the mesh, the main window looks as in Fig. 3(c) and the main view as in Fig. 4(b).
- *Adjusting how data is displayed:* Without going into too much detail, if you double click onto the name of a plot in the “Plots” window, you get a dialog in which many of the properties of this plot can be adjusted. Further details can be changed by using “Operators” on a plot.
- *Making the output prettier:* As can be seen in Fig. 4, Visit by default puts a lot of clutter around the figure – the name of the user, the name of the input file, color bars, axes labels and ticks, etc. This may be useful to explore data in the beginning but does not yield good pictures for presentations or publications. To reduce the amount of information displayed, go to the “Controls > Annotations” menu item to get a dialog in which all of these displays can be selectively switched on and off.
- *Saving figures:* To save a visualization into a file that can then be included into presentations and publications, go to the menu item “File > Save window”. This will create successively numbered files

in the directory from which Visit was started each time a view is saved. Things like the format used for these files can be chosen using the “File > Set save options” menu item. We have found that one can often get better looking pictures by selecting the “Screenshot” method in this dialog.

More information on all of these topics can be found in the Visit documentation, see <https://visit.llnl.gov/>. We have also recorded video lectures demonstrating this process interactively at <http://www.youtube.com/watch?v=3ChnUxqtt08> for Visit, and at <http://www.youtube.com/watch?v=w-65jufR-bc> for Paraview.

4.4.2 Visualizing statistical data

In addition to the graphical output discussed above, ASPECT produces a statistics file that collects information produced during each time step. For the remainder of this section, let us assume that we have run ASPECT with the input file discussed in Section 5.2.1, simulating convection in a box. After running ASPECT, you will find a file called `statistics` in the output directory that, at the time of writing this, looked like this: This file has a structure that looks (at the time of writing this section) like this:

```
# 1: Time step number
# 2: Time (seconds)
# 3: Number of mesh cells
# 4: Number of Stokes degrees of freedom
# 5: Number of temperature degrees of freedom
# 6: Iterations for temperature solver
# 7: Iterations for Stokes solver
# 8: Velocity iterations in Stokes preconditioner
# 9: Schur complement iterations in Stokes preconditioner
# 10: Time step size (seconds)
# 11: RMS velocity (m/s)
# 12: Max. velocity (m/s)
# 13: Minimal temperature (K)
# 14: Average temperature (K)
# 15: Maximal temperature (K)
# 16: Average nondimensional temperature (K)
# 17: Outward heat flux through boundary with indicator 0 ("left") (W)
# 18: Outward heat flux through boundary with indicator 1 ("right") (W)
# 19: Outward heat flux through boundary with indicator 2 ("bottom") (W)
# 20: Outward heat flux through boundary with indicator 3 ("top") (W)
# 21: Visualization file name
0 0.0000e+00 256 2467 1089 0 29 30 29 1.2268e-02 1.79026783e+00 2.54322608e+00
1 1.2268e-02 256 2467 1089 32 29 30 30 3.7388e-03 5.89844152e+00 8.35160076e+00
2 1.6007e-02 256 2467 1089 20 28 29 29 2.0239e-03 1.09071922e+01 1.54298908e+01
3 1.8031e-02 256 2467 1089 15 27 28 28 1.3644e-03 1.61759153e+01 2.28931189e+01
4 1.9395e-02 256 2467 1089 13 26 27 27 1.0284e-03 2.14465789e+01 3.03731397e+01
5 2.0424e-02 256 2467 1089 11 25 26 26 8.2812e-04 2.66110761e+01 3.77180480e+01
```

In other words, it first lists what the individual columns mean with a hash mark at the beginning of the line and then has one line for each time step in which the individual columns list what has been explained above.¹⁵

This file is easy to visualize. For example, one can import it as a whitespace separated file into a spreadsheet such as Microsoft Excel or OpenOffice/LibreOffice Calc and then generate graphs of one column against another. Or, maybe simpler, there is a multitude of simple graphing programs that do not need the

¹⁵With input files that ask for initial adaptive refinement, the first time step may appear twice because we solve on a mesh that is globally refined and we then start the entire computation over again on a once adaptively refined mesh (see the parameters in Section A.107 for how to do that).

overhead of a full fledged spreadsheet engine and simply plot graphs. One that is particularly simple to use and available on every major platform is **Gnuplot**. It is extensively documented at <http://www.gnuplot.info/>.

Gnuplot is a command line program in which you enter commands that plot data or modify the way data is plotted. When you call it, you will first get a screen that looks like this:

```
/home/user/aspect/output gnuplot

G N U P L O T
Version 4.6 patchlevel 0      last modified 2012-03-04
Build System: Linux x86_64

Copyright (C) 1986-1993, 1998, 2004, 2007-2012
Thomas Williams, Colin Kelley and many others

gnuplot home:      http://www.gnuplot.info
faq, bugs, etc:    type "help FAQ"
immediate help:   type "help"  (plot window: hit 'h')

Terminal type set to 'qt'
gnuplot>
```

At the prompt on the last line, you can then enter commands. Given the description of the individual columns given above, let us first try to plot the heat flux through boundary 2 (the bottom boundary of the box), i.e., column 19, as a function of time (column 2). This can be achieved using the following command:

```
plot "statistics" using 2:19
```

The left panel of Fig. 5 shows what **Gnuplot** will display in its output window. There are many things one can configure in these plots (see the **Gnuplot** manual referenced above). For example, let us assume that we want to add labels to the x - and y -axes, use not just points but lines and points for the curves, restrict the time axis to the range $[0, 0.2]$ and the heat flux axis to $[-10 : 10]$, plot not only the flux through the bottom but also through the top boundary (column 20) and finally add a key to the figure, then the following commands achieve this:

```
set xlabel "Time"
set ylabel "Heat flux"
set style data linespoints
plot [0:0.2] [-10:10] "statistics" using 2:19 title "Bottom boundary", \
                    "statistics" using 2:20 title "Top boundary"
```

If a line gets too long, you can continue it by ending it in a backslash as above. This is rarely used on the command line but useful when writing the commands above into a script file, see below. We have done it here to get the entire command into the width of the page.

For those who are lazy, **Gnuplot** allows to abbreviate things in many different ways. For example, one can abbreviate most commands. Furthermore, one does not need to repeat the name of an input file if it is the same as the previous one in a plot command. Thus, instead of the commands above, the following abbreviated form would have achieved the same effect:

```
se xl "Time"
se yl "Heat flux"
se sty da lp
pl [:0.2] [-10:10] "statistics" us 2:19 t "Bottom boundary", "" us 2:20 t "Top boundary"
```

This is of course unreadable at first but becomes useful once you become more familiar with the commands offered by this program.

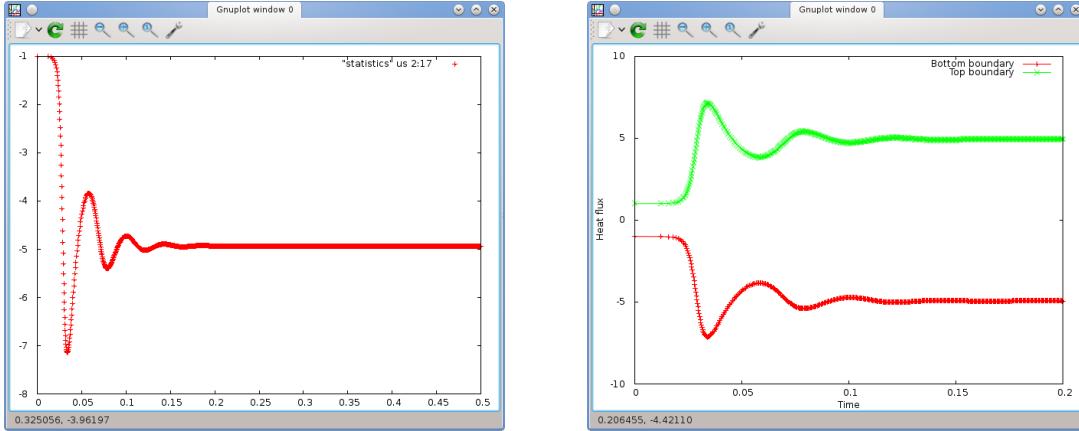


Figure 5: Visualizing the statistics file obtained from the example in Section 5.2.1 using **Gnuplot**: Output using simple commands.

Once you have gotten the commands that create the plot you want right, you probably want to save it into a file. **Gnuplot** can write output in many different formats. For inclusion in publications, either **eps** or **png** are the most common. In the latter case, the commands to achieve this are

```
set terminal png
set output "heatflux.png"
replot
```

The last command will simply generate the same plot again but this time into the given file. The result is a graphics file similar to the one shown in Fig. 8 on page 69.

Note: After setting output to a file, *all* following plot commands will want to write to this file. Thus, if you want to create more plots after the one just created, you need to reset output back to the screen. On Linux, this is done using the command **set terminal X11**. You can then continue experimenting with plots and when you have the next plot ready, switch back to output to a file.

What makes **Gnuplot** so useful is that it doesn't just allow entering all these commands at the prompt. Rather, one can write them all into a file, say **plot-heatflux.gnuplot**, and then, on the command line, call

```
gnuplot plot-heatflux.gnuplot
```

to generate the **heatflux.png** file. This comes in handy if one wants to create the same plot for multiple simulations while playing with parameters of the physical setup. It is also a very useful tool if one wants to generate the same kind of plot again later with a different data set, for example when a reviewer requested additional computations to be made for a paper or if one realizes that one has forgotten or misspelled an axis label in a plot.¹⁶

Gnuplot has many many more features we have not even touched upon. For example, it is equally happy to produce three-dimensional graphics, and it also has statistics modules that can do things like curve fits, statistical regression, and many more operations on the data you provide in the columns of an input file. We will not try to cover them here but instead refer to the manual at <http://www.gnuplot.info/>. You can

¹⁶In my own work, I usually save the ASPECT input file, the **statistics** output file and the **Gnuplot** script along with the actual figure I want to include in a paper. This way, it is easy to either re-run an entire simulation, or just tweak the graphic at a later time. Speaking from experience, you will not believe how often one wants to tweak a figure long after it was first created. In such situations it is outstandingly helpful if one still has both the actual data as well as the script that generated the graphic.

also get a good amount of information by typing `help` at the prompt, or a command like `help plot` to get help on the `plot` command.

4.4.3 Large data issues for parallel computations

Among the challenges in visualizing the results of parallel computations is dealing with the large amount of data. The first bottleneck this presents is during run-time when ASPECT wants to write the visualization data of a time step to disk. Using the compressed VTU format, ASPECT generates on the order of 10 bytes of output for each degree of freedom in 2d and more in 3d; thus, output of a single time step can run into the range of gigabytes that somehow have to get from compute nodes to disk. This stresses both the cluster interconnect as well as the data storage array.

There are essentially two strategies supported by ASPECT for this scenario:

- If your cluster has a fast interconnect, for example Infiniband, and if your cluster has a fast, distributed file system, then ASPECT can produce output files that are already located in the correct output directory (see the options in Section A.1) on the global file system. ASPECT uses MPI I/O calls to this end, ensuring that the local machines do not have to access these files using slow NFS-mounted global file systems.
- If your cluster has a slow interconnect, e.g., if it is simply a collection of machines connected via Ethernet, then writing data to a central file server may block the rest of the program for a while. On the other hand, if your machines have fast local storage for temporary file systems, then ASPECT can write data first into such a file and then move it in the background to its final destination while already continuing computations. To select this mode, set the appropriate variables discussed in Section A.138. Note, however, that this scheme only makes sense if every machine on which MPI processes run has fast local disk space for temporary storage.

Note: An alternative would be if every processor directly writes its own files into the global output directory (possibly in the background), without the intermediate step of the temporary file. In our experience, file servers are quickly overwhelmed when encountering a few hundred machines wanting to open, fill, flush and close their own file via NFS mounted file system calls, sometimes completely blocking the entire cluster environment for extended periods of time.

4.5 Checkpoint/restart support

If you do long runs, especially when using parallel computations, there are a number of reasons to periodically save the state of the program:

- If the program crashes for whatever reason, the entire computation may be lost. A typical reason is that a program has exceeded the requested wallclock time allocated by a batch scheduler on a cluster.
- Most of the time, no realistic initial conditions for strongly convecting flow are available. Consequently, one typically starts with a somewhat artificial state and simply waits for a long while till the convective state enters the phase where it shows its long-term behavior. However, getting there may take a good amount of CPU time and it would be silly to always start from scratch for each different parameter setting. Rather, one would like to start such parameter studies with a saved state that has already passed this initial, unphysical, transient stage.

To this end, ASPECT creates a set of files in the output directory (selected in the parameter file) every N time steps (controlled by the number of steps or wall time as specified in subsection Checkpointing, see Section A.36) in which the entire state of the program is saved so that a simulation can later be continued at this point. The previous checkpoint files will then be deleted. To resume operations from the last saved state, you need to set the `Resume computation` flag in the input parameter file to `true`, see Section A.1.

Note: It is not imperative that the parameters selected in the input file are exactly the same when resuming a program from a saved state than what they were at the time when this state was saved. For example, one may want to choose a different parametrization of the material law, or add or remove postprocessors that should be run at the end of each time step. Likewise, the end time, the times at which some additional mesh refinement steps should happen, etc., can be different.

Yet, it is clear that some other things can't be changed: For example, the geometry model that was used to generate the coarse mesh and describe the boundary must be the same before and after resuming a computation. Likewise, you can not currently restart a computation with a different number of processors than initially used to checkpoint the simulation. Not all invalid combinations are easy to detect, and ASPECT may not always realize immediate what is going on if you change a setting that can't be changed. However, you will almost invariably get nonsensical results after some time.

4.6 Making ASPECT run faster

When developing ASPECT, we are guided by the principle that the default for all settings should be *safe*. In particular, this means that you should get errors when something goes wrong, the program should not let you choose an input file parameter so that it doesn't make any sense, and we should solve the equations to best ability without cutting corners. The goal is that when you start working with ASPECT that we give you the best answer we can. The downside is that this also makes ASPECT run slower than may be possible. This section describes ways of making ASPECT run faster – assuming that you know what you are doing and are making conscious decisions.

4.6.1 Debug vs. optimized mode

Both DEAL.II and ASPECT by default have a great deal of internal checking to make sure that the code's state is valid. For example, if you write a new postprocessing plugin (see Section 6.1)) in which you need to access the solution vector, then DEAL.II's `Vector` class will make sure that you are only accessing elements of the vector that actually exist and are available on the current machine if this is a parallel computation. We do so because it turns out that by far the most bugs one introduces in programs are of the kind where one tries to do something that obviously doesn't make sense (such as accessing vector element 101 when it only has 100 elements). These kinds of bugs are more frequent than implementing a wrong algorithm, but they are fortunately easy to find if you have a sufficient number of assertions in your code. The downside is that assertions cost run time.

As mentioned above, the default is to have all of these assertions in the code to catch those places where we may otherwise silently access invalid memory locations. However, once you have a plugin running and verified that your input file runs without problems, you can switch off all of these checks by switching from debug to optimized mode. This means re-compiling ASPECT and linking against a version of the DEAL.II library without all of these internal checks. Because this is the first thing you will likely want to do, we have already discussed how to do all of this in Section 4.3.

4.6.2 Adjusting solver tolerances

At the heart of every time step lies the solution of linear systems for the Stokes equations, the temperature field, and possibly for compositional fields. In essence, each of these steps requires us to solve a linear system of the form $Ax = b$ which we do through iterative solvers, i.e., we try to find a sequence of approximations $x^{(k)}$ where $x^{(k)} \rightarrow x = A^{-1}b$. This iteration is terminated at iteration k if the approximation is “close enough” to the exact solution. The solvers we use determine this by testing after every iteration whether the residual, $r^{(k)} = A(x - x^{(k)}) = b - Ax^{(k)}$, satisfies $\|r^{(k)}\| \leq \varepsilon \|r^{(0)}\|$ where ε is called the (relative) *tolerance*.

Obviously, the smaller we choose ε , the more accurate the approximation $x^{(k)}$ will be. On the other hand, it will also take more iterations and, consequently, more CPU time to reach the stopping criterion with a

smaller tolerance. The default value of these tolerances are chosen so that the approximation is typically sufficient. You can make ASPECT run faster if you choose these tolerances larger. The parameters you can adjust are all listed in Section A.1 and are located at the top level of the input file. In particular, the parameters you want to look at are `Linear solver tolerance`, `Temperature solver tolerance` and `Composition solver tolerance`.

All this said, it is important to understand the consequences of choosing tolerances larger. In particular, if you choose tolerances too large, then the difference between the exact solution of a linear system x and the approximation $x^{(k)}$ may become so large that you do not get an accurate output of your model any more. A rule of thumb in choosing tolerances is to start with a small value and then increase the tolerance until you come to a point where the output quantities start to change significantly. This is the point where you will want to stop.

4.6.3 Adjusting solver preconditioner tolerances

To solve the Stokes equations it is necessary to lower the condition number of the Stokes matrix by preconditioning it. In ASPECT a right preconditioner $Y^{-1} = \begin{pmatrix} \widetilde{A}^{-1} & -\widetilde{A}^{-1}B^T\widetilde{S}^{-1} \\ 0 & \widetilde{S}^{-1} \end{pmatrix}$ is used to precondition the system, where \widetilde{A}^{-1} is the approximate inverse of the A block and \widetilde{S}^{-1} is the approximate inverse of the Schur complement matrix. Matrix \widetilde{A}^{-1} and \widetilde{S}^{-1} are calculated through a CG solve, which requires a tolerance to be set. In comparison with the solver tolerances of the previous section, these parameters are relatively safe to use, since they only change the preconditioner, but can speed up or slow down solving the Stokes system considerably.

In practice \widetilde{A}^{-1} takes by far the most time to compute, but is also very important in conditioning the system. The accuracy of the computation of \widetilde{A}^{-1} is controlled by the parameter `Linear solver A block tolerance` which has a default value of $1e-2$. Setting this tolerance to a less strict value will result in more outer iterations, since the preconditioner is not as good, but the amount of time to compute \widetilde{A}^{-1} can drop significantly resulting in a reduced total solve time. The cookbook crustal deformation (Section 5.3.8) for example can be computed much faster by setting the `Linear solver A block tolerance` to $5e-1$. The calculation of \widetilde{S}^{-1} is usually much faster and the conditioning of the system is less sensitive to the parameter `Linear solver S block tolerance`, but for some problems it might be worth it to investigate.

4.6.4 Using lower order elements for the temperature/compositional discretization

The default settings of ASPECT use quadratic finite elements for the velocity. Given that the temperature and compositional fields essentially (up to material parameters) satisfy advection equations of the kind $\partial_t T + \mathbf{u} \cdot \nabla T = \dots$, it seems appropriate to also use quadratic finite element shape functions for the temperature and compositional fields.

However, this is not mandatory. If you do not care about high accuracy in these fields and are mostly interested in the velocity or pressure field, you can select lower-order finite elements in the input file. The polynomial degrees are controlled with the parameters in the *discretization* section of the input file, see Section A.38, in particular by `Temperature polynomial degree` and `Composition polynomial degree`.

As with the other parameters discussed above and below, it is worthwhile comparing the results you get with different values of these parameters when making a decision whether you want to save on accuracy in order to reduce compute time. An example of how this choice affects the accuracy you get is discussed in Section 5.2.1.

4.6.5 Limiting postprocessing

ASPECT has a lot of postprocessing capabilities, from generating graphical output to computing average temperatures or temperature fluxes. To see what all is possible, take a look at the `List of postprocessors` parameter that can be set in the input file, see Section A.118.

Many of these postprocessors take a non-negligible amount of time. How much they collectively use can be inferred from the timing report ASPECT prints periodically among its output, see for example the output shown in Section 5.2.1. So, if your computations take too long, consider limiting which postprocessors you run to those you really need. Some postprocessors – for example those that generate graphical output, see Section A.138 – also allow you to run them only once every once in a while, rather than at every time step.

4.6.6 Switching off pressure normalization

In most practically relevant cases, the Stokes equations (1)–(2) only determine the pressure up to a constant because only the pressure gradient appears in the equations, not the actual value of it. However, unlike this “mathematical” pressure, we have a very specific notion of the “physical” pressure: namely a well-defined quantity that at the surface of Earth equals the air pressure, which compared to the hydrostatic pressure inside Earth is essentially zero.

As a consequence, the default in ASPECT is to normalize the computed “mathematical” pressure in such a way that either the mean pressure at the surface is zero (where the geometry model describes where the “surface” is, see Section 6.3.3), or that the mean pressure in the domain is zero. This normalization is important if your model describes densities, viscosities and other quantities in dependence of the pressure – because you almost certainly had the “physical” pressure in mind, not some unspecified “mathematical” one. On the other hand, if you have a material model in which the pressure does not enter, then you don’t need to normalize the pressure at all – simply go with whatever the solver provides. In that case, you can switch off pressure normalization by looking at the `Pressure normalization` parameter at the top level of the input file, see Section A.1.

4.6.7 Regularizing models with large coefficient variation

Models with large jumps in viscosity and other coefficients present significant challenges to both discretizations and solvers. In particular, they can lead to very long solver times. Section 5.2.8 presents parameters that can help regularize models and these typically also include significant improvements in run-time.

4.6.8 Using multithreading

In most cases using as many MPI processes as possible is the optimal parallelization strategy for ASPECT models, but if you are limited by the amount of MPI communication it can be beneficial to use multiple threads per MPI process. While not utilized by our linear solvers, this parallelization can speed up the assembly of the system matrices, e.g. by around 10-15% if you utilize unused logical cores, or nearly linearly if you use otherwise unused physical cores. This can also reduce the performance cost if you are memory limited and need to run your model on less than the available number of cores per node on a cluster to increase the available memory per core. Running with for example two threads per process will offset some of the performance loss you will see in these situations.

Multithreading is controlled by setting the command line parameter `-j` or `--threads`. If the parameter is not set, ASPECT will create exactly one thread per MPI process, i.e. multithreading is disabled. Appending the parameter allows ASPECT to spawn several threads per MPI process. Note that the internally used TBB library will determine the number of threads based on the number of available cores, i.e., if you start 2 MPI processes on a quadcore machine with hyperthreading (8 logical cores), ASPECT will spawn 4 threads on each MPI process. Also note that there is no guarantee that the final number of threads will exactly match the number of available logical cores if you start with a number of processes that is not a divisor of your logical cores (e.g. 3 MPI processes for 8 logical cores).

4.7 Input parameter files

What ASPECT computes is driven by two things:

- The models implemented in ASPECT. This includes the geometries, the material laws, or the initial conditions currently supported. Which of these models are currently implemented is discussed below; Section 6 discusses in great detail the process of implementing additional models.
- Which of the implemented models is selected, and what their run-time parameters are. For example, you could select a model that prescribes constant coefficients throughout the domain from all the material models currently implemented; you could then select appropriate values for all of these constants. Both of these selections happen from a parameter file that is read at run time and whose name is specified on the command line. (See also Section 4.1.)

In this section, let us give an overview of what can be selected in the parameter file. Specific parameters, their default values, and allowed values for these parameters are documented in Section A. An index with page numbers for all run-time parameters can be found on page 463.

4.7.1 The structure of parameter files

Most of the run-time behavior of ASPECT is driven by a parameter file that looks in essence like this:

```
set Dimension          = 2
set Resume computation = false
set End time           = 1e10
set CFL number         = 1.0
set Output directory   = output

subsection Mesh refinement
  set Initial adaptive refinement = 1
  set Initial global refinement  = 4
end

subsection Material model
  set Model name           = simple

  subsection Simple model
    set Reference density     = 3300
    set Reference temperature = 293
    set Viscosity             = 5e24
  end
end

...
```

Some parameters live at the top level, but most parameters are grouped into subsections. An input parameter file is therefore much like a file system: a few files live in the root directory; others are in a nested hierarchy of sub-directories. And just as with files, parameters have both a name (the thing to the left of the equals sign) and a content (what's to the right).

All parameters you can list in this input file have been *declared* in ASPECT. What this means is that you can't just list anything in the input file, and expect that entries that are unknown are simply ignored. Rather, if your input file contains a line setting a parameter that is unknown, you will get an error message. Likewise, all declared parameters have a description of possible values associated with them – for example, some parameters must be non-negative integers (the number of initial refinement steps), can either be true or false (whether the computation should be resumed from a saved state), or can only be a single element from a selection (the name of the material model). If an entry in your input file doesn't satisfy these constraints, it will be rejected at the time of reading the file (and not when a part of the program actually accesses the value and the programmer has taken the time to also implement some error checking at this location). Finally, because parameters have been declared, you do not *need* to specify a parameter in the input file: if a parameter isn't listed, then the program will simply use the default provided when declaring the parameter.

Note: In cases where a parameter requires a significant amount of text, you can end a line in the input file with a backslash. This indicates that the following line will simply continue to be part of the text of the current line, in the same way as the C/C++ preprocessor expands lines that end in backslashes. The underlying implementation always eats whitespace at the beginning of each continuing line, but not before the backslash. This means that the parameter file

```
set Some parameter = abc\
    def
```

is equivalent to

```
set Some parameter = abcdef
```

that is, with no space between `abc` and `def` despite the leading whitespace at the beginning of the second line. If you do want space between these two parts, you need to add it before the backslash in the first of the two lines.

4.7.2 Categories of parameters

The parameters that can be provided in the input file can roughly be categorized into the following groups:

- Global parameters (see Section A.1): These parameters determine the overall behavior of the program. Primarily they describe things like the output directory, the end time of the simulation, or whether the computation should be resumed from a previously saved state.
- Parameters for certain aspects of the numerical algorithm: These describe, for example, the specifics of the spatial discretization. In particular, this is the case for parameters concerning the polynomial degree of the finite element approximation (Section A.38), some details about the stabilization (Section A.39), and how adaptive mesh refinement is supposed to work (Section A.107).
- Parameters that describe certain global aspects of the equations to be solved: This includes, for example, a description if certain terms in the model should be omitted or not. See Section A.40 for the list of parameters in this category.
- Parameters that characterize plugins: Certain behaviors of ASPECT are described by what we call *plugins* – self-contained parts of the code that describe one particular aspect of the simulation. An example would be which of the implemented material models to use, and the specifics of this material model. The sample parameter file above gives an indication of how this works: within a subsection of the file that pertains to the material models, one can select one out of several plugins (or, in the case of the postprocessors, any number, including none, of the available plugins), and one can then specify the specifics of this model in a sub-subsection dedicated to this particular model.

A number of components of ASPECT are implemented via plugins. Some of these, together with the sections in which their parameters are declared, are the following:

- The material model: Sections A.81 and following.
- The geometry: Sections A.42 and following.
- The gravity description: Sections A.52 and following.
- Initial conditions for the temperature: Sections A.69 and following.
- Temperature boundary conditions: Sections A.16 and following.
- Postprocessors: Sections A.118 and following for most postprocessors, section A.138 and following for postprocessors related to visualization.

The details of parameters in each of these categories can be found in the sections linked to above. Some of them will also be used in the cookbooks in Section 5.

4.7.3 A note on the syntax of formulas in input files

Input files have different ways of describing certain things to ASPECT. For example, you could select a plugin for the temperature initial values that prescribes a constant temperature, or a plugin that implements a particular formula for these initial conditions in C++ in the code of the plugin, or a plugin that allows you to describe this formula in a symbolic way in the input file (see Section A.69). An example of this latter case is this snippet of code discussed in Section 5.2.2:

```
subsection Initial temperature model
  set Model name = function

  subsection Function
    set Variable names      = x,y,z
    set Function constants = p=0.01, L=1, pi=3.1415926536, k=1
    set Function expression = (1.0-z) - p*cos(k*pi*x/L)*sin(pi*z)*y^3
  end
end
```

The formulas you can enter here need to use a syntax that is understood by the functions and classes that interpret what you write. Internally, this is done using the muparser library, see <http://muparser.beltoforion.de/>. The syntax is mostly self-explanatory in that it allows to use the usual symbols `x`, `y` and `z` to reference coordinates (unless a particular plugin uses different variables, such as the depth), the symbol `t` for time in many situations, and allows you to use all of the typical mathematical functions such as sine and cosine. Another common case is an if-statement that has the general form `if(condition, true-expression, false-expression)`. For more examples of the syntax understood, reference the documentation of the muparser library linked to above.

4.7.4 Compatibility of input files with newer ASPECT versions

We strive to maintain compatibility for options in input files as long as possible. However, occasionally we have to reorder, rename, or remove options from parameter files to improve ASPECT further. This is in particular true for new major versions. In order to allow running old parameter files with newer ASPECT versions we provide scripts that can automatically update existing parameter files to the new syntax. Executing `doc/update_prm_files.sh` with one or more parameter files as arguments will create a backup of the old parameter file (named `old_filename.bak`), and replace the existing file with a version that should work with the current ASPECT version. Using this script would look like this:

```
bash doc/update_prm_files.sh cookbooks/convection_box.prm
```

Note: Not all text replacements are unique, and the structure of input files allows for constructions the script can not properly parse. Also we can not guarantee to preserve the structure and position of comments, as it is not always clear to which part of the input file they refer. Thus, it is important that you check your updated input file for errors. That being said, all input files in the main ASPECT repository are updated successfully using this script.

4.8 A graphical user interface for editing ASPECT parameter files

Preparing a parameter file in a text editor can be a tedious task, not only because the number of input parameters has grown considerably during the development of ASPECT, but also because remembering the names of commonly used options is a waste of (human) memory. Therefore, we provide a graphical user interface that builds upon an available program for DEAL.II. This GUI allows to investigate existing input parameters, including their default values, modify these values in a spreadsheet like environment and then save a formatted parameter file that can be used to start an ASPECT model. In the following subsections we describe installing and using this user interface.

4.8.1 Installing parameter-GUI

The DEAL.II parameter-GUI program can be downloaded at https://github.com/dealii/parameter_gui, and is compiled using the `cmake` program just like ASPECT itself. The program has no dependencies except for the Qt development libraries that should be available as packages for most Linux distributions and can also be obtained for all major operating systems at <https://www.qt.io/download-open-source/>.

Example steps for installing the parameter-gui could look as follows:

1. Download the program from https://github.com/dealii/parameter_gui.
2. Prepare a Makefile by running `cmake .` in the source folder.
3. Compile the program by running `make`.
4. Make sure to set the environment variable `PARAMETER_GUI_DIR` to the directory that contains the parameter-GUI executable (optional). This will allow ASPECT to automatically enable the GUI during configuration.

Installing on macOS On a mac machine with recent macOS Sierra 10.12.4, Qt development libraries of version 4.x.x at the libraries' official website <https://www.qt.io/download-open-source/> may fail to install. Alternatively, you can install `qt4` through Homebrew (also see instruction here <https://github.com/cartr/homebrew-qt4>)

```
brew tap cartr/qt4
brew tap-pin cartr/qt4
brew install qt@4
```

or install it through Mac Ports (<https://www.macports.org/>)

```
sudo port install qt4-mac
```

Then you can follow the Linux user instructions provided previously to download and install dealii parameter-GUI. Before running `cmake ..`, you may need to either pass the path of `qt4` and specify the value of variable `QT_LIBRARIES` to the directory that contains the libraries of `qt4` or add those information into your `.bash_profile`. For example, for installation through Mac Ports, you can set the following into your `.bash_profile`

```
export PATH="$PATH:/opt/local/libexec/qt4"
export QT_LIBRARIES="/opt/local/libexec/qt4"
```

4.8.2 Using ASPECT-GUI

When configuring ASPECT after executing the above steps, it should automatically pick up the location of the parameter-GUI, and will create a new script named `aspect-gui` within the build folder. If this does not happen, it is possible to hand over the location of the parameter-GUI as a `cmake` variable during configuration (e.g. `cmake -D PARAMETER_GUI_EXECUTABLE=path_to_your_executable`).

The `aspect-gui` script can be executed from any folder either with no argument or with one argument that contains the path to an existing parameter file. The script will run ASPECT with the given (or an empty) parameter file, generate a database of existing input parameters, and open the parameter-GUI program with this database. The resulting window looks similar to Fig. 6. If an existing parameter file was given, all parameter fields are pre-filled with the values set in the file instead of the default values. In the program's main window you can change parameters as necessary, and then save the file as a new parameter file. This parameter file can then be used to start an ASPECT model as usual. Note that it is possible to prepare and execute the parameter file with different versions of ASPECT, e.g. if you prepare parameter files on a local machine, and execute the model on a remote compute cluster. Note however that if the two ASPECT versions contain different default values or parameter names have changed, this can lead to unexpected model behavior or even unusable parameter files.

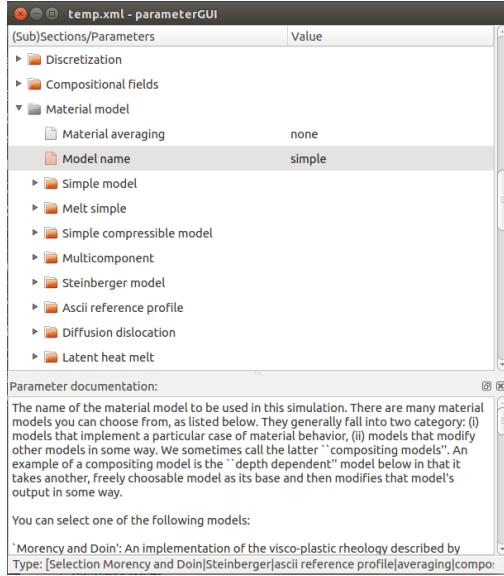


Figure 6: *The parameter GUI lists all available parameter options, and allows to change and save them into a new parameter file. Input fields know about the type of the variable and will display useful options to change them (e.g. drop-down menus, file dialogs, text fields).*

5 Cookbooks

In this section, let us present a number of “cookbooks” – examples of how to use ASPECT in typical or less typical ways. As discussed in Sections 4 and A, ASPECT is driven by run-time parameter files, and so setting up a particular situation primarily comes down to creating a parameter file that has the right entries. Thus, the subsections below will discuss in detail what parameters to set and to what values. Note that parameter files need not specify *all* parameters – of which there is a bewildering number – but only those that are relevant to the particular situation we would like to model. All parameters not listed explicitly in the input file are simply left at their default value (the default values are also documented in Section A).

Of course, there are situations where what you want to do is not covered by the models already implemented. Specifically, you may want to try a different geometry, a different material or gravity model, or different boundary conditions. In such cases, you will need to implement these extensions in the actual source code. Section 6 provides information on how to do that.

The remainder of this section shows a number of applications of ASPECT. They are grouped into three categories: Simple setups of examples that show thermal convection (Section 5.2), setups that try to model geophysical situations (Section 5.3) and setups that are used to benchmark ASPECT to ensure correctness or to test accuracy of our solvers (Section 5.4). Before we get there, however, we will review how one usually approaches setting up computations in Section 5.1.

Note: The input files discussed in the following sections can generally be found in the `cookbooks/` directory of your ASPECT installation.

5.1 How to set up computations

ASPECT’s computations are controlled by input parameter files such as those we will discuss in the following sections.¹⁷ Basically, these are just regular text files you can edit with programs like `gedit`, `kwrite` or `kate`

¹⁷You can also extend ASPECT using plugins – i.e., pieces of code you compile separately and either link into the ASPECT executable itself, or reference from the input file. This is discussed in Section 6.

when working on Linux, or something as simple as NotePad on Windows. When setting up these input files, you basically have to describe everything that characterizes the computation you want to do. In particular, this includes the following:

- What internal forces act on the medium (the equation)?
- What external forces do we have (the right hand side)
- What is the domain (geometry)?
- What happens at the boundary for each variable involved (boundary conditions)?
- How did it look at the beginning (initial conditions)?

For each of these questions, there are one or more input parameters (sometimes grouped into sections) that allow you to specify what you want. For example, to choose a geometry, you will typically have a block like this in your input file:

```
set Dimension = 2
subsection Geometry model
    set Model name = box

    subsection Box
        set X extent = 1
        set Y extent = 1
    end
end
```

This indicates that you want to do a computation in 2d, using a rectangular geometry (a “box”) with edge length equal to one in both the x - and y -directions. Of course, there are other geometries you can choose from for the `Model name` parameter, and consequently other subsections that specify the details of these geometries.

Similarly, you describe boundary conditions using parameters such as this:

```
subsection Boundary temperature model
    set Fixed temperature boundary indicators = bottom, top
end

subsection Boundary velocity model
    set Tangential velocity boundary indicators = left, right, bottom, top
end
```

This snippet describes which of the four boundaries of the two-dimensional box we have selected above should have a prescribed temperature or an insulating boundary, and at which parts of the boundary we want zero, tangential or prescribed velocities.¹⁸

If you go down the list of questions about the setup above, you have already done the majority of the work describing your computation. The remaining parameters you will typically want to specify have to do with the computation itself. For example, what variables do you want to output and how often? What statistics do you want to compute. The following sections will give ample examples for all of this, but using the questions above as a guideline is already a good first step.

¹⁸Internally, the geometry models ASPECT uses label every part of the boundary with what is called a *boundary indicator* – a number that identifies pieces of the boundary. If you know which number each piece has, you can list these numbers on the right hand sides of the assignments of boundary types above. For example, the left boundary of the box has boundary indicator zero (see Section A.42), and using this number instead of the `left` would have been equally valid. However, numbers are far more difficult to remember than names, and consequently every geometry model provides string aliases such as “`left`” for each boundary indicator describing parts of the boundary. These symbolic aliases are specific to the geometry – for the box, they are “`left`”, “`right`”, “`bottom`”, etc., whereas for a spherical shell they are “`inner`” and “`outer`” – but are described in the documentation of every geometry model, see Section A.42.

Note: It is of course possible to set up input files for computations completely from scratch. However, in practice, it is often simpler to go through the list of cookbooks already provided and find one that comes close to what you want to do. You would then modify this cookbook until it does what you want to do. The advantage is that you can start with something you already know works, and you can inspect how each change you make – changing the details of the geometry, changing the material model, or changing what is being computed at the end of each time step – affects what you get.

5.2 Simple setups

5.2.1 Convection in a 2d box

In this first example, let us consider a simple situation: a 2d box of dimensions $[0, 1] \times [0, 1]$ that is heated from below, insulated at the left and right, and cooled from the top. We will also consider the simplest model, the incompressible Boussinesq approximation with constant coefficients $\eta, \rho_0, \mathbf{g}, C_p, k$, for this testcase. Furthermore, we assume that the medium expands linearly with temperature. This leads to the following set of equations:

$$-\nabla \cdot [2\eta\varepsilon(\mathbf{u})] + \nabla p = \rho_0(1 - \alpha(T - T_0))\mathbf{g} \quad \text{in } \Omega, \quad (45)$$

$$\nabla \cdot \mathbf{u} = 0 \quad \text{in } \Omega, \quad (46)$$

$$\rho_0 C_p \left(\frac{\partial T}{\partial t} + \mathbf{u} \cdot \nabla T \right) - \nabla \cdot k \nabla T = 0 \quad \text{in } \Omega. \quad (47)$$

It is well known that we can non-dimensionalize this set of equations by introducing the Rayleigh number $Ra = \frac{\rho_0 g \alpha \Delta T h^3}{\eta \kappa}$, where h is the height of the box, $\kappa = \frac{k}{\rho_0 C_p}$ is the thermal diffusivity and ΔT is the temperature difference between top and bottom of the box. Formally, we can obtain the non-dimensionalized equations by using the above form and setting coefficients in the following way:

$$\rho_0 = C_p = \kappa = \alpha = \eta = h = \Delta T = 1, \quad T_0 = 0, \quad g = Ra,$$

where $\mathbf{g} = -g\mathbf{e}_z$ is the gravity vector in negative z -direction. We will see all of these values again in the input file discussed below. One point to note is that for the Boussinesq approximation, as described above, the density in the temperature equation is chosen as the reference density ρ_0 rather than the full density $\rho(1 - \alpha(T - T_0))$ as we see it in the buoyancy term on the right hand side of the momentum equation. As ASPECT is able to handle different approximations of the equations (see Section 2.10), we also have to specify in the input file that we want to use the Boussinesq approximation. The problem is completed by stating the velocity boundary conditions: tangential flow along all four of the boundaries of the box.

This situation describes a well-known benchmark problem for which a lot is known and against which we can compare our results. For example, the following is well understood:

- For values of the Rayleigh number less than a critical number $Ra_c \approx 780$, thermal diffusion dominates convective heat transport and any movement in the fluid is damped exponentially. If the Rayleigh number is moderately larger than this threshold then a stable convection pattern forms that transports heat from the bottom to the top boundaries. The simulations we will set up operates in this regime. Specifically, we will choose $Ra = 10^4$.

On the other hand, if the Rayleigh number becomes even larger, a series of period doublings starts that makes the system become more and more unstable. We will investigate some of this behavior at the end of this section.

- For certain values of the Rayleigh number, very accurate values for the heat flux through the bottom and top boundaries are available in the literature. For example, Blankenbach *et al.* report a non-dimensional heat flux of 4.884409 ± 0.00001 , see [BBC⁺89]. We will compare our results against this value below.

With this said, let us consider how to represent this situation in practice.

The input file. The verbal description of this problem can be translated into an ASPECT input file in the following way (see Section A for a description of all of the parameters that appear in the following input file, and the indices at the end of this manual if you want to find a particular parameter; you can find the input file to run this cookbook example in [cookbooks/convection-box.prm](#)):

```
# At the top, we define the number of space dimensions we would like to
# work in:
set Dimension = 2

# There are several global variables that have to do with what
# time system we want to work in and what the end time is. We
# also designate an output directory.
set Use years in output instead of seconds = false
set End time = 0.5
set Output directory = output-convection-box

# Then there are variables that describe how the pressure should
# be normalized. Here, we choose a zero average pressure
# at the surface of the domain (for the current geometry, the
# surface is defined as the top boundary).
set Pressure normalization = surface
set Surface pressure = 0

# Then come a number of sections that deal with the setup
# of the problem to solve. The first one deals with the
# geometry of the domain within which we want to solve.
# The sections that follow all have the same basic setup
# where we select the name of a particular model (here,
# the box geometry) and then, in a further subsection,
# set the parameters that are specific to this particular
# model.
subsection Geometry model
    set Model name = box

    subsection Box
        set X extent = 1
        set Y extent = 1
    end
end

# The next section deals with the initial conditions for the
# temperature (there are no initial conditions for the
# velocity variable since the velocity is assumed to always
# be in a static equilibrium with the temperature field).
# There are a number of models with the 'function' model
# a generic one that allows us to enter the actual initial
# conditions in the form of a formula that can contain
# constants. We choose a linear temperature profile that
# matches the boundary conditions defined below plus
# a small perturbation:
subsection Initial temperature model
    set Model name = function
```

```

subsection Function
  set Variable names      = x,z
  set Function constants  = p=0.01, L=1, pi=3.1415926536, k=1
  set Function expression = (1.0-z) - p*cos(k*pi*x/L)*sin(pi*z)
end
end

# Then follows a section that describes the boundary conditions
# for the temperature. The model we choose is called 'box' and
# allows to set a constant temperature on each of the four sides
# of the box geometry. In our case, we choose something that is
# heated from below and cooled from above.
# All other parts of the boundary are insulated (i.e., no heat flux through
# these boundaries; this is also often used to specify symmetry boundaries).
subsection Boundary temperature model
  set Fixed temperature boundary indicators = bottom, top
  set List of model names = box

subsection Box
  set Bottom temperature = 1
  set Left temperature   = 0
  set Right temperature  = 0
  set Top temperature    = 0
end
end

# We then also have to prescribe several other parts of the model such as
# which boundaries actually carry a prescribed boundary temperature, whereas
# The next parameters then describe on which parts of the
# boundary we prescribe a zero or nonzero velocity and
# on which parts the flow is allowed to be tangential.
# Here, all four sides of the box allow tangential
# unrestricted flow but with a zero normal component:
subsection Boundary velocity model
  set Tangential velocity boundary indicators = left, right, bottom, top
end

# The following two sections describe first the
# direction (vertical) and magnitude of gravity and the
# material model (i.e., density, viscosity, etc). We have
# discussed the settings used here in the introduction to
# this cookbook in the manual already.
subsection Gravity model
  set Model name = vertical

subsection Vertical
  set Magnitude = 1e4  # = Ra
end
end

subsection Material model

```

```

set Model name = simple

subsection Simple model
  set Reference density          = 1
  set Reference specific heat    = 1
  set Reference temperature      = 0
  set Thermal conductivity       = 1
  set Thermal expansion coefficient = 1
  set Viscosity                  = 1
end
end

# We also have to specify that we want to use the Boussinesq
# approximation (assuming the density in the temperature
# equation to be constant, and incompressibility).
subsection Formulation
  set Formulation = Boussinesq approximation
end

# The settings above all pertain to the description of the
# continuous partial differential equations we want to solve.
# The following section deals with the discretization of
# this problem, namely the kind of mesh we want to compute
# on. We here use a globally refined mesh without
# adaptive mesh refinement.
subsection Mesh refinement
  set Initial global refinement      = 4
  set Initial adaptive refinement    = 0
  set Time steps between mesh refinement = 0
end

# The final part is to specify what ASPECT should do with the
# solution once computed at the end of every time step. The
# process of evaluating the solution is called 'postprocessing'
# and we choose to compute velocity and temperature statistics,
# statistics about the heat flux through the boundaries of the
# domain, and to generate graphical output files for later
# visualization. These output files are created every time
# a time step crosses time points separated by 0.01. Given
# our start time (zero) and final time (0.5) this means that
# we will obtain 50 output files.
subsection Postprocess
  set List of postprocessors = velocity statistics, temperature statistics, ...
                                ... heat flux statistics, visualization

subsection Visualization
  set Time between graphical output = 0.01
end
end

subsection Solver parameters
  set Temperature solver tolerance = 1e-10
end

```

Running the program. When you run this program for the first time, you are probably still running ASPECT in debug mode (see Section 4.3) and you will get output like the following:

```
Number of active cells: 256 (on 5 levels)
Number of degrees of freedom: 3,556 (2,178+289+1,089)

*** Timestep 0: t=0 seconds
Solving temperature system... 0 iterations.
Rebuilding Stokes preconditioner...
Solving Stokes system... 31+0 iterations.

[. . .]

*** Timestep 1085: t=0.5 seconds
Solving temperature system... 0 iterations.
Solving Stokes system... 5 iterations.

Postprocessing:
RMS, max velocity:          43.5 m/s, 70.3 m/s
Temperature min/avg/max:     0 K, 0.5 K, 1 K
Heat fluxes through boundary parts: 0.01977 W, -0.01977 W, -4.787 W, 4.787 W

Termination requested by criterion: end time

+-----+
| Total wallclock time elapsed since start | 66.5s | | |
|                                           |         |
|                                           |         |
| Section           | no. calls | wall time | % of total |
+-----+-----+-----+
| Assemble Stokes system | 1086 | 8.63s | 13% |
| Assemble temperature system | 1086 | 32s | 48% |
| Build Stokes preconditioner | 1 | 0.0225s | 0% |
| Build temperature preconditioner | 1086 | 1.52s | 2.3% |
| Solve Stokes system | 1086 | 7.7s | 12% |
| Solve temperature system | 1086 | 0.729s | 1.1% |
| Initialization | 1 | 0.0316s | 0% |
| Postprocessing | 1086 | 7.76s | 12% |
| Setup dof systems | 1 | 0.0104s | 0% |
| Setup initial conditions | 1 | 0.00621s | 0% |
+-----+
```

If you've read up on the difference between debug and optimized mode (and you should before you switch!) then consider disabling debug mode. If you run the program again, every number should look exactly the same (and it does, in fact, as I am writing this) except for the timing information printed every hundred time steps and at the end of the program:

```
+-----+
| Total wallclock time elapsed since start | 25.8s | | |
|                                           |         |
|                                           |         |
| Section           | no. calls | wall time | % of total |
+-----+-----+-----+
| Assemble Stokes system | 1086 | 2.51s | 9.7% |
| Assemble temperature system | 1086 | 9.88s | 38% |
| Build Stokes preconditioner | 1 | 0.0271s | 0.11% |
| Build temperature preconditioner | 1086 | 1.58s | 6.1% |
| Solve Stokes system | 1086 | 6.38s | 25% |
+-----+
```

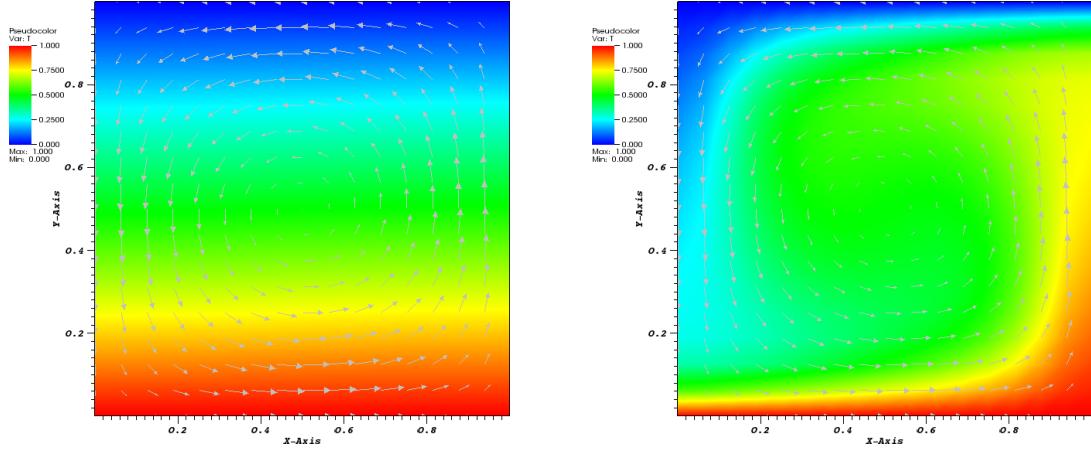


Figure 7: Convection in a box: Initial temperature and velocity field (left) and final state (right).

Solve temperature system	1086	0.542s	2.1%	
Initialization	1	0.219s	0.85%	
Postprocessing	1086	2.79s	11%	
Setup dof systems	1	0.23s	0.89%	
Setup initial conditions	1	0.107s	0.41%	
+-----+-----+-----+-----+-----+				

In other words, the program ran more than 2 times faster than before. Not all operations became faster to the same degree: assembly, for example, is an area that traverses a lot of code both in ASPECT and in DEAL.II and so encounters a lot of verification code in debug mode. On the other hand, solving linear systems primarily requires lots of matrix vector operations. Overall, the fact that in this example, assembling linear systems and preconditioners takes so much time compared to actually solving them is primarily a reflection of how simple the problem is that we solve in this example. This can also be seen in the fact that the number of iterations necessary to solve the Stokes and temperature equations is so low. For more complex problems with non-constant coefficients such as the viscosity, as well as in 3d, we have to spend much more work solving linear systems whereas the effort to assemble linear systems remains the same.

Visualizing results. Having run the program, we now want to visualize the numerical results we got. ASPECT can generate graphical output in formats understood by pretty much any visualization program (see the parameters described in Section A.138) but we will here follow the discussion in Section 4.4 and use the default VTU output format to visualize using the Visit program.

In the parameter file we have specified that graphical output should be generated every 0.01 time units. Looking through these output files (which can be found in the folder `output-convection-box`, as specified in the input file), we find that the flow and temperature fields quickly converge to a stationary state. Fig. 7 shows the initial and final states of this simulation.

There are many other things we can learn from the output files generated by ASPECT, specifically from the statistics file that contains information collected at every time step and that has been discussed in Section 4.4.2. In particular, in our input file, we have selected that we would like to compute velocity, temperature, and heat flux statistics. These statistics, among others, are listed in the statistics file whose head looks like this for the current input file:

```
# 1: Time step number
# 2: Time (seconds)
# 3: Time step size (seconds)
# 4: Number of mesh cells
```

```

# 5: Number of Stokes degrees of freedom
# 6: Number of temperature degrees of freedom
# 7: Iterations for temperature solver
# 8: Iterations for Stokes solver
# 9: Velocity iterations in Stokes preconditioner
# 10: Schur complement iterations in Stokes preconditioner
# 11: RMS velocity (m/s)
# 12: Max. velocity (m/s)
# 13: Minimal temperature (K)
# 14: Average temperature (K)
# 15: Maximal temperature (K)
# 16: Average nondimensional temperature (K)
# 17: Outward heat flux through boundary with indicator 0 ("left") (W)
# 18: Outward heat flux through boundary with indicator 1 ("right") (W)
# 19: Outward heat flux through boundary with indicator 2 ("bottom") (W)
# 20: Outward heat flux through boundary with indicator 3 ("top") (W)
# 21: Visualization file name
... lots of numbers arranged in columns ...

```

Fig. 8 shows the results of visualizing the data that can be found in columns 2 (the time) plotted against columns 11 and 12 (root mean square and maximal velocities). Plots of this kind can be generated with Gnuplot by typing (see Section 4.4.2 for a more thorough discussion):

```
plot "output-convection-box/statistics" using 2:11 with lines
```

Fig. 8 shows clearly that the simulation enters a steady state after about $t \approx 0.1$ and then changes very little. This can also be observed using the graphical output files from which we have generated Fig. 7. One can look further into this data to find that the flux through the top and bottom boundaries is not exactly the same (up to the obvious difference in sign, given that at the bottom boundary heat flows into the domain and at the top boundary out of it) at the beginning of the simulation until the fluid has attained its equilibrium. However, after $t \approx 0.2$, the fluxes differ by only $5 \cdot 10^{-5}$, i.e., by less than 0.001% of their magnitude.¹⁹ The flux we get at the last time step, 4.787, is less than 2% away from the value reported in [BBC⁺89] (≈ 4.88) although we compute on a 16×16 mesh and the values reported by Blankenbach are extrapolated from meshes of size up to 72×72 . This shows the accuracy that can be obtained using a higher order finite element. Secondly, the fluxes through the left and right boundary are not exactly zero but small. Of course, we have prescribed boundary conditions of the form $\frac{\partial T}{\partial n} = 0$ along these boundaries, but this is subject to discretization errors. It is easy to verify that the heat flux through these two boundaries disappears as we refine the mesh further.

Furthermore, ASPECT automatically also collects statistics about many of its internal workings. Fig. 9 shows the number of iterations required to solve the Stokes and temperature linear systems in each time step. It is easy to see that these are more difficult to solve in the beginning when the solution still changes significant from time step to time step. However, after some time, the solution remains mostly the same and solvers then only need 9 or 10 iterations for the temperature equation and 4 or 5 iterations for the Stokes equations because the starting guess for the linear solver – the previous time step’s solution – is already pretty good. If you look at any of the more complex cookbooks, you will find that one needs many more iterations to solve these equations.

Play time 1: Different Rayleigh numbers. After showing you results for the input file as it can be found in [cookbooks/convection-box.prm](#), let us end this section with a few ideas on how to play with it and what to explore. The first direction one could take this example is certainly to consider different Rayleigh numbers. As mentioned above, for the value $Ra = 10^4$ for which the results above have been produced, one gets a stable convection pattern. On the other hand, for values $Ra < Ra_c \approx 780$, any movement of the fluid

¹⁹This difference is far smaller than the numerical error in the heat flux on the mesh this data is computed on.

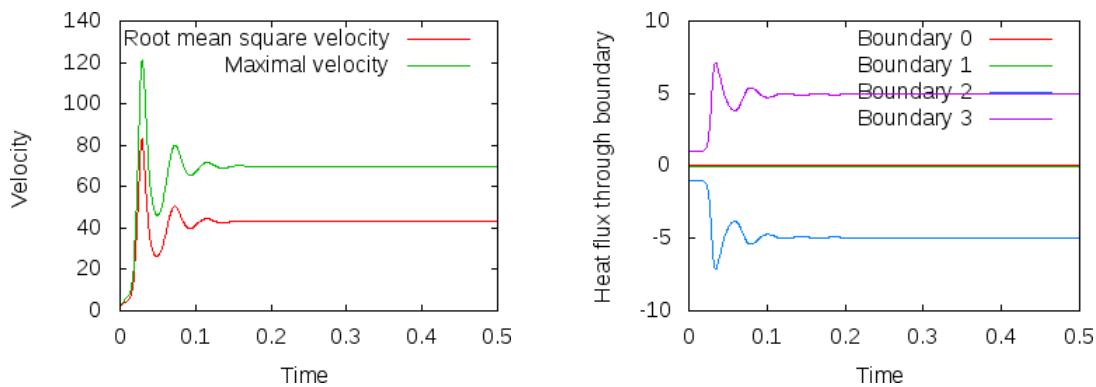


Figure 8: *Convection in a box: Root mean square and maximal velocity as a function of simulation time (left). Heat flux through the four boundaries of the box (right).*

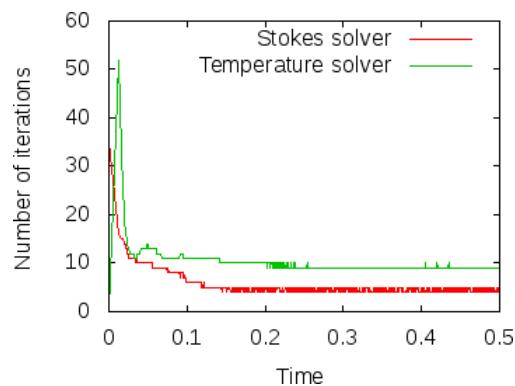


Figure 9: *Convection in a box: Number of linear iterations required to solve the Stokes and temperature equations in each time step.*

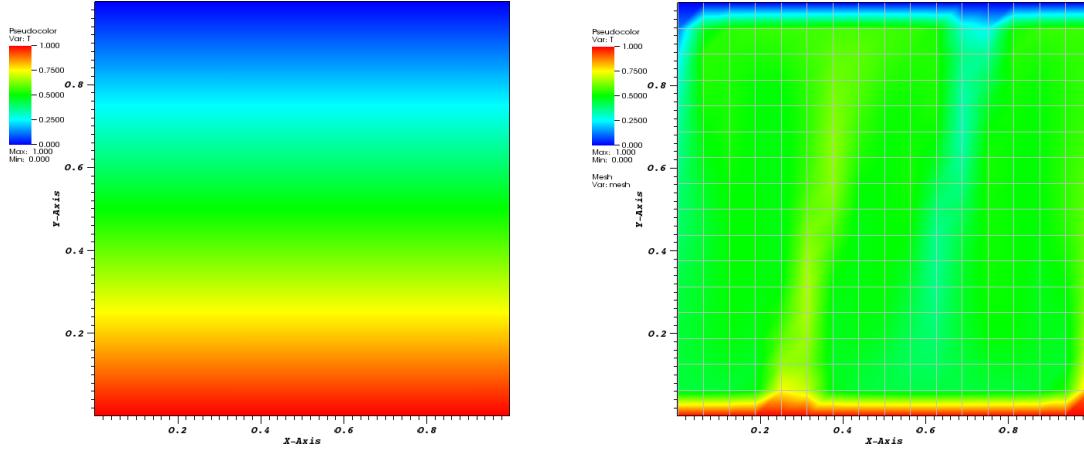


Figure 10: Convection in a box: Temperature fields at the end of a simulation for $Ra = 10^2$ where thermal diffusion dominates (left) and $Ra = 10^6$ where convective heat transport dominates (right). The mesh on the right is clearly too coarse to resolve the structure of the solution.

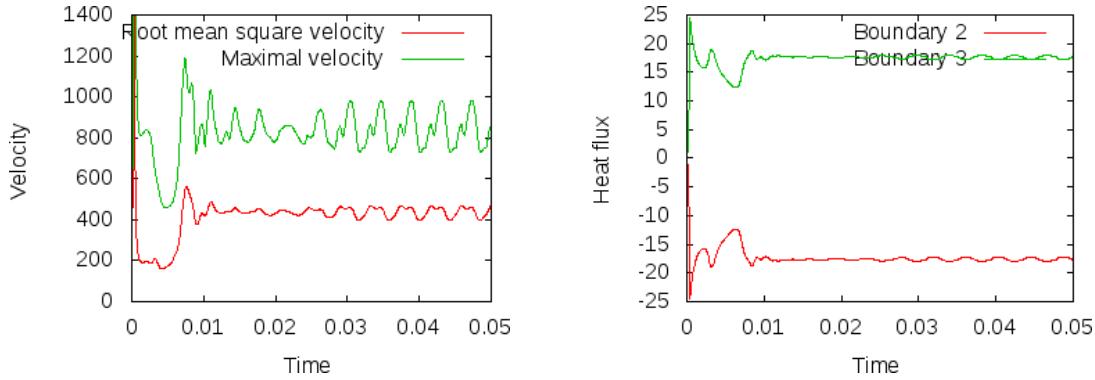


Figure 11: Convection in a box: Velocities (left) and heat flux across the top and bottom boundaries (right) as a function of time at $Ra = 10^6$.

dies down exponentially and we end up with a situation where the fluid doesn't move and heat is transported from the bottom to the top only through heat conduction. This can be explained by considering that the Rayleigh number in a box of unit extent is defined as $Ra = \frac{\rho_0 g \alpha \Delta T}{\eta k}$. A small Rayleigh number means that the viscosity is too large (i.e., the buoyancy given by the product of the magnitude of gravity times the density anomalies caused by temperature – $\rho_0 \alpha \Delta T$ – is not strong enough to overcome friction forces within the fluid).

On the other hand, if the Rayleigh number is large (i.e., the viscosity is small or the buoyancy large) then the fluid develops an unsteady convection period. As we consider fluids with larger and larger Ra , this pattern goes through a sequence of period-doubling events until flow finally becomes chaotic. The structures of the flow pattern also become smaller and smaller.

We illustrate these situations in Fig.s 10 and 11. The first shows the temperature field at the end of a simulation for $Ra = 10^2$ (below Ra_c) and at $Ra = 10^6$. Obviously, for the right picture, the mesh is not fine enough to accurately resolve the features of the flow field and we would have to refine it more. The second of the figures shows the velocity and heatflux statistics for the computation with $Ra = 10^6$; it is obvious here that the flow no longer settles into a steady state but has a periodic behavior. This can also be seen by looking at movies of the solution.

To generate these results, remember that we have chosen $g = Ra$ in our input file. In other words, changing the input file to contain the parameter setting

```
subsection Gravity model
  subsection Vertical
    set Magnitude = 1e6  # = Ra
  end
end
```

will achieve the desired effect of computing with $Ra = 10^6$.

Play time 2: Thinking about finer meshes. In our computations for $Ra = 10^4$ we used a 16×16 mesh and obtained a value for the heat flux that differed from the generally accepted value from Blankenbach *et al.* [BBC⁺89] by less than 2%. However, it may be interesting to think about computing even more accurately. This is easily done by using a finer mesh, for example. In the parameter file above, we have chosen the mesh setting as follows:

```
subsection Mesh refinement
  set Initial global refinement      = 4
  set Initial adaptive refinement   = 0
  set Time steps between mesh refinement = 0
end
```

We start out with a box geometry consisting of a single cell that is refined four times. Each time we split each cell into its 4 children, obtaining the 16×16 mesh already mentioned. The other settings indicate that we do not want to refine the mesh adaptively at all in the first time step, and a setting of zero for the last parameter means that we also never want to adapt the mesh again at a later time. Let us stick with the never-changing, globally refined mesh for now (we will come back to adaptive mesh refinement again at a later time) and only vary the initial global refinement. In particular, we could choose the parameter `Initial global refinement` to be 5, 6, or even larger. This will get us closer to the exact solution albeit at the expense of a significantly increased computational time.

A better strategy is to realize that for $Ra = 10^4$, the flow enters a steady state after settling in during the first part of the simulation (see, for example, the graphs in Fig. 8). Since we are not particularly interested in this initial transient process, there is really no reason to spend CPU time using a fine mesh and correspondingly small time steps during this part of the simulation (remember that each refinement results in four times as many cells in 2d and a time step half as long, making reaching a particular time at least 8 times as expensive, assuming that all solvers in ASPECT scale perfectly with the number of cells). Rather, we can use a parameter in the ASPECT input file that lets us increase the mesh resolution at later times. To this end, let us use the following snippet for the input file:

```
subsection Mesh refinement
  set Initial global refinement      = 3
  set Initial adaptive refinement   = 0
  set Time steps between mesh refinement = 0
  set Additional refinement times   = 0.2, 0.3, 0.4
  set Refinement fraction           = 1.0
  set Coarsening fraction          = 0.0
end
```

What this does is the following: We start with an 8×8 mesh (3 times globally refined) but then at times $t = 0.2, 0.3$ and 0.4 we refine the mesh using the default refinement indicator (which one this is is not important because of the next statement). Each time, we refine, we refine a fraction 1.0 of the cells, i.e., *all* cells and we coarsen a fraction of 0.0 of the cells, i.e. no cells at all. In effect, at these additional refinement times, we do another global refinement, bringing us to refinement levels 4, 5 and finally 6.

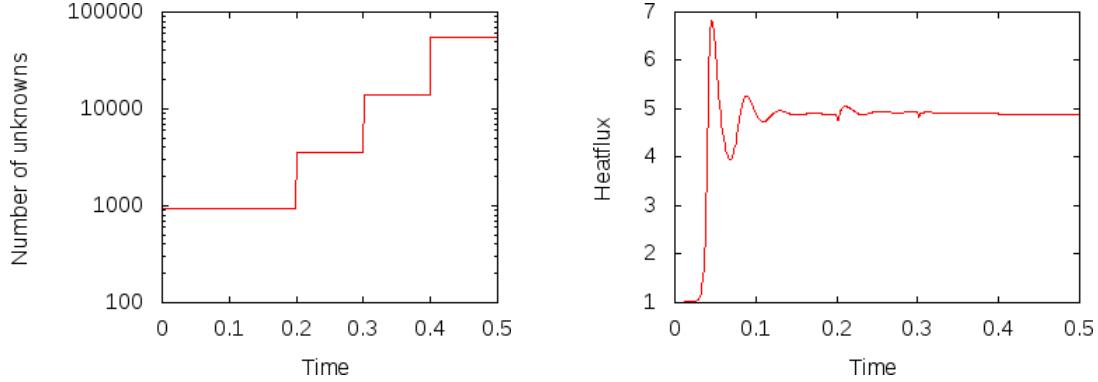


Figure 12: *Convection in a box: Refinement in stages. Total number of unknowns in each time step, including all velocity, pressure and temperature unknowns (left) and heat flux across the top boundary (right).*

Fig. 12 shows the results. In the left panel, we see how the number of unknowns grows over time (note the logscale for the y -axis). The right panel displays the heat flux. The jumps in the number of cells is clearly visible in this picture as well. This may be surprising at first but remember that the mesh is clearly too coarse in the beginning to really resolve the flow and so we should expect that the solution changes significantly if the mesh is refined. This effect becomes smaller with every additional refinement and is barely visible at the last time this happens, indicating that the mesh before this refinement step may already have been fine enough to resolve the majority of the dynamics.

In any case, we can compare the heat fluxes we obtain at the end of these computations: With a globally four times refined mesh, we get a value of 4.787 (an error of approximately 2% against the accepted value from Blankenbach, 4.884409 ± 0.00001). With a globally five times refined mesh we get 4.879, and with a globally six times refined mesh we get 4.89 (an error of almost 0.1%). With the mesh generated using the procedure above we also get 4.89 with the digits printed on the screen²⁰ (also corresponding to an error of almost 0.1%). In other words, our simple procedure of refining the mesh during the simulation run yields the same accuracy as using the mesh that is globally refined in the beginning of the simulation, while needing a much lower compute time.

Play time 3: Changing the finite element in use. Another way to increase the accuracy of a finite element computation is to use a higher polynomial degree for the finite element shape functions. By default, ASPECT uses quadratic shape functions for the velocity and the temperature and linear ones for the pressure. However, this can be changed with a single number in the input file.

Before doing so, let us consider some aspects of such a change. First, looking at the pictures of the solution in Fig. 7, one could surmise that the quadratic elements should be able to resolve the velocity field reasonably well given that it is rather smooth. On the other hand, the temperature field has a boundary layer at the top and bottom. One could conjecture that the temperature polynomial degree is therefore the limiting factor and not the polynomial degree for the flow variables. We will test this conjecture below. Secondly, given the nature of the equations, increasing the polynomial degree of the flow variables increases the cost to solve these equations by a factor of $\frac{22}{9}$ in 2d (you can get this factor by counting the number of degrees of freedom uniquely associated with each cell) but leaves the time step size and the cost of solving the temperature system unchanged. On the other hand, increasing the polynomial degree of the temperature variable from 2 to 3 requires $\frac{9}{4}$ times as many degrees of freedom for the temperature and also requires us to reduce the size of the time step by a factor of $\frac{2}{3}$. Because solving the temperature system is not a dominant factor in each time step (see the timing results shown at the end of the screen output above), the reduction

²⁰The statistics file gives this value to more digits: 4.89008498. However, these are clearly more digits than the result is accurate.

in time step is the only important factor. Overall, increasing the polynomial degree of the temperature variable turns out to be the cheaper of the two options.

Following these considerations, let us add the following section to the parameter file:

```
subsection Discretization
  set Stokes velocity polynomial degree      = 2
  set Temperature polynomial degree        = 3
end
```

This leaves the velocity and pressure shape functions at quadratic and linear polynomial degree but increases the polynomial degree of the temperature from quadratic to cubic. Using the original, four times globally refined mesh, we then get the following output:

```
Number of active cells: 256 (on 5 levels)
Number of degrees of freedom: 4,868 (2,178+289+2,401)

*** Timestep 0: t=0 seconds
Solving temperature system... 0 iterations.
Rebuilding Stokes preconditioner...
Solving Stokes system... 30+0 iterations.

[... ...]

*** Timestep 1621: t=0.5 seconds
Solving temperature system... 0 iterations.
Solving Stokes system... 1+0 iterations.

Postprocessing:
RMS, max velocity:           42.9 m/s, 69.5 m/s
Temperature min/avg/max:     0 K, 0.5 K, 1 K
Heat fluxes through boundary parts: -0.004602 W, 0.004602 W, -4.849 W, 4.849 W

Termination requested by criterion: end time
```

Total wallclock time elapsed since start	53.6s		
Section	no. calls	wall time	% of total
Assemble Stokes system	1622	4.04s	7.5%
Assemble temperature system	1622	24.4s	46%
Build Stokes preconditioner	1	0.0121s	0%
Build temperature preconditioner	1622	8.05s	15%
Solve Stokes system	1622	8.92s	17%
Solve temperature system	1622	1.67s	3.1%
Initialization	1	0.0327s	0%
Postprocessing	1622	4.27s	8%
Setup dof systems	1	0.00418s	0%
Setup initial conditions	1	0.00236s	0%

The heat flux through the top and bottom boundaries is now computed as 4.878. Using the five times globally refined mesh, it is 4.8837 (an error of 0.015%). This is 6 times more accurate than the once more globally refined mesh with the original quadratic elements, at a cost significantly smaller. Furthermore, we can of course combine this with the mesh that is gradually refined as simulation time progresses, and we then get a heat flux that is equal to 4.884446, also only 0.01% away from the accepted value!

As a final remark, to test our hypothesis that it was indeed the temperature polynomial degree that was the limiting factor, we can increase the Stokes polynomial degree to 3 while leaving the temperature polynomial degree at 2. A quick computation shows that in that case we get a heat flux of 4.747 – almost the same value as we got initially with the lower order Stokes element. In other words, at least for this testcase, it really was the temperature variable that limits the accuracy.

5.2.2 Convection in a 3d box

The world is not two-dimensional. While the previous section introduced a number of the knobs one can play with with ASPECT, things only really become interesting once one goes to 3d. The setup from the previous section is easily adjusted to this and in the following, let us walk through some of the changes we have to consider when going from 2d to 3d. The full input file that contains these modifications and that was used for the simulations we will show subsequently can be found at [cookbooks/convection-box-3d.prm](#).

The first set of changes has to do with the geometry: it is three-dimensional, and we will have to address the fact that a box in 3d has 6 sides, not the 4 we had previously. The documentation of the “box” geometry (see Section A.42) states that these sides are numbered as follows: “*in 3d, boundary indicators 0 through 5 indicate left, right, front, back, bottom and top boundaries.*” Recalling that we want tangential flow all around and want to fix the temperature to known values at the bottom and top, the following will make sense:

```
set Dimension          = 3

subsection Geometry model
  set Model name = box

  subsection Box
    set X extent = 1
    set Y extent = 1
    set Z extent = 1
  end
end

subsection Boundary temperature model
  set Fixed temperature boundary indicators = bottom, top
  set List of model names = box

  subsection Box
    set Bottom temperature = 1
    set Top temperature   = 0
  end
end

subsection Boundary velocity model
  set Tangential velocity boundary indicators = left, right, front, back, bottom, top
end
```

The next step is to describe the initial conditions. As before, we will use an unstably layered medium but the perturbation now needs to be both in x - and y -direction

```
subsection Initial temperature model
  set Model name = function

  subsection Function
    set Variable names      = x,y,z
    set Function constants  = p=0.01, L=1, pi=3.1415926536, k=1
    set Function expression = (1.0-z) - p*cos(k*pi*x/L)*sin(pi*z)*y^3
```

```

end
end

```

The third issue we need to address is that we can likely not afford a mesh as fine as in 2d. We choose a mesh that is refined 3 times globally at the beginning, then 3 times adaptively, and is then adapted every 15 time steps. We also allow one additional mesh refinement in the first time step following $t = 0.003$ once the initial instability has given way to a more stable pattern:

```

subsection Mesh refinement
  set Initial global refinement      = 3
  set Initial adaptive refinement   = 3
  set Time steps between mesh refinement = 15

  set Additional refinement times    = 0.003
end

```

Finally, as we have seen in the previous section, a computation with $Ra = 10^4$ does not lead to a simulation that is exactly exciting. Let us choose $Ra = 10^6$ instead (the mesh chosen above with up to 7 refinement levels after some time is fine enough to resolve this). We can achieve this in the same way as in the previous section by choosing $\alpha = 10^{-10}$ and setting

```

subsection Gravity model
  set Model name = vertical

  subsection Vertical
    set Magnitude = 1e16  # = Ra / Thermal expansion coefficient
  end
end

```

This has some interesting implications. First, a higher Rayleigh number makes time scales correspondingly smaller; where we generated graphical output only once every 0.01 time units before, we now need to choose the corresponding increment smaller by a factor of 100:

```

subsection Postprocess
  set List of postprocessors = velocity statistics, temperature statistics, ...
                            ...heat flux statistics, visualization

  subsection Visualization
    set Time between graphical output = 0.0001
  end
end

```

Secondly, a simulation like this – in 3d, with a significant number of cells, and for a significant number of time steps – will likely take a good amount of time. The computations for which we show results below was run using 64 processors by running it using the command `mpirun -n 64 ./aspect convection-box-3d.prm`. If the machine should crash during such a run, a significant amount of compute time would be lost if we had to run everything from the start. However, we can avoid this by periodically checkpointing the state of the computation:

```

subsection Checkpointing
  set Steps between checkpoint = 50
end

```

If the computation does crash (or if a computation runs out of the time limit imposed by a scheduling system), then it can be restarted from such checkpointing files (see the parameter `Resume computation` in Section A.1).

Running with this input file requires a bit of patience²¹ since the number of degrees of freedom is just so large: it starts with a bit over 330,000...

```
Running with 64 MPI tasks.
Number of active cells: 512 (on 4 levels)
Number of degrees of freedom: 20,381 (14,739+729+4,913)

*** Timestep 0: t=0 seconds
  Solving temperature system... 0 iterations.
  Rebuilding Stokes preconditioner...
  Solving Stokes system... 18 iterations.

Number of active cells: 1,576 (on 5 levels)
Number of degrees of freedom: 63,391 (45,909+2,179+15,303)

*** Timestep 0: t=0 seconds
  Solving temperature system... 0 iterations.
  Rebuilding Stokes preconditioner...
  Solving Stokes system... 19 iterations.

Number of active cells: 3,249 (on 5 levels)
Number of degrees of freedom: 122,066 (88,500+4,066+29,500)

*** Timestep 0: t=0 seconds
  Solving temperature system... 0 iterations.
  Rebuilding Stokes preconditioner...
  Solving Stokes system... 20 iterations.

Number of active cells: 8,968 (on 5 levels)
Number of degrees of freedom: 331,696 (240,624+10,864+80,208)

*** Timestep 0: t=0 seconds
  Solving temperature system... 0 iterations.
  Rebuilding Stokes preconditioner...
  Solving Stokes system... 21 iterations.
[...]
```

...but then increases quickly to around 2 million as the solution develops some structure and, after time $t = 0.003$ where we allow for an additional refinement, increases to over 10 million where it then hovers between 8 and 14 million with a maximum of 15,147,534. Clearly, even on a reasonably quick machine, this will take some time: running this on a machine bought in 2011, doing the 10,000 time steps to get to $t = 0.0219$ takes approximately 484,000 seconds (about five and a half days).

The structure or the solution is easiest to grasp by looking at isosurfaces of the temperature. This is shown in Fig. 13 and you can find a movie of the motion that ensues from the heating at the bottom at http://www.youtube.com/watch?v=_bKqU_P4j48. The simulation uses adaptively changing meshes that are fine in rising plumes and sinking blobs and are coarse where nothing much happens. This is most easily seen in the movie at <http://www.youtube.com/watch?v=CzCKYyR-cmg>. Fig. 14 shows some of these meshes, though still pictures do not do the evolving nature of the mesh much justice. The effect of increasing the Rayleigh number is apparent when comparing the size of features with, for example, the picture at the right of Fig. 7. In contrast to that picture, the simulation is also clearly non-stationary.

As before, we could analyze all sorts of data from the statistics file but we will leave this to those interested in specific data. Rather, Fig. 15 only shows the upward heat flux through the bottom and top boundaries

²¹For computations of this size, one should test a few time steps in debug mode but then, of course, switch to running the actual computation in optimized mode – see Section 4.3.

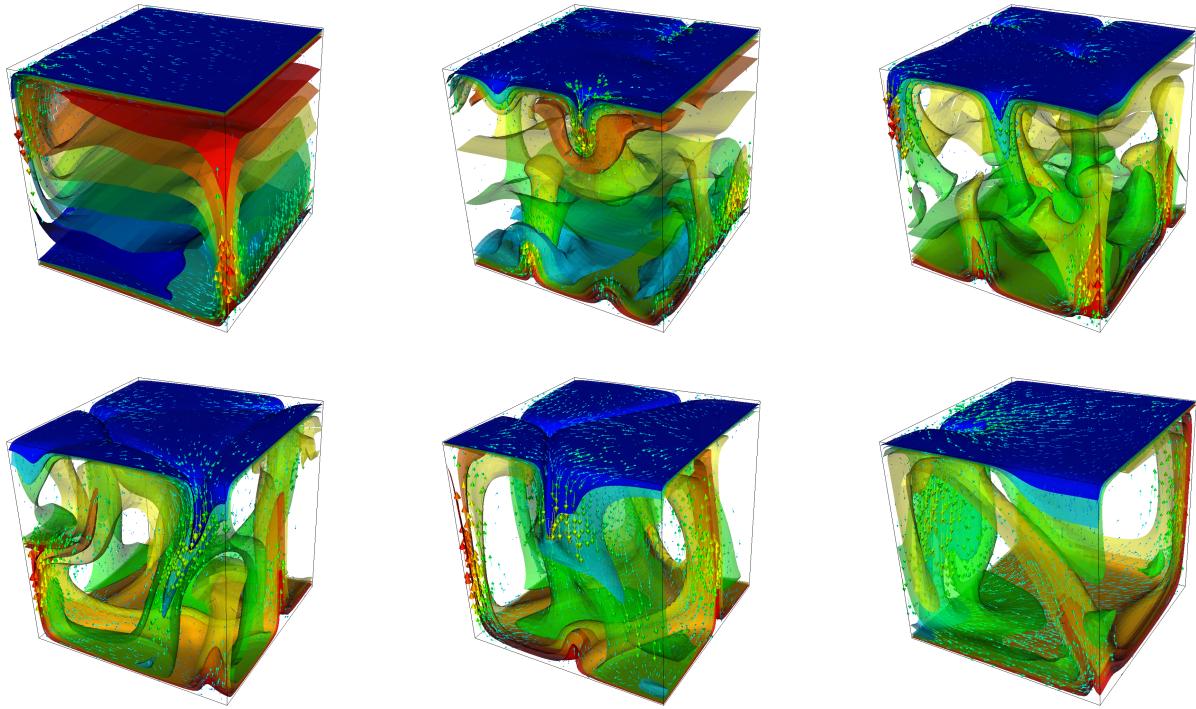


Figure 13: Convection in a 3d box: Temperature isocontours and some velocity vectors at the first time step after times $t = 0.001, 0.004, 0.006$ (top row, left to right) and $t = 0.01, 0.013, 0.018$ (bottom row).

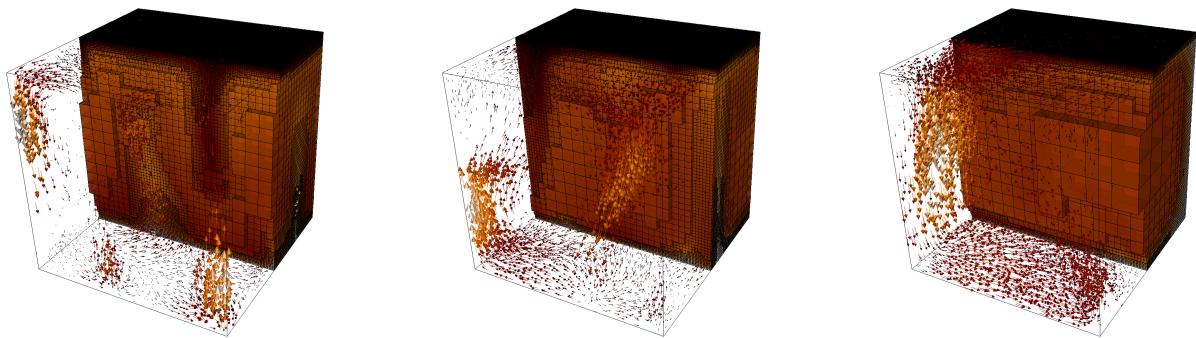


Figure 14: Convection in a 3d box: Meshes and large-scale velocity field for the third, fourth and sixth of the snapshots shown in Fig. 13.

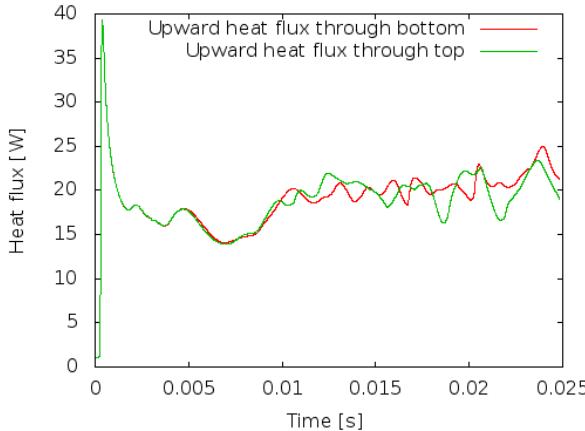


Figure 15: *Convection in a 3d box: Upward heat flux through the bottom and top boundaries as a function of time.*

of the domain as a function of time.²² The figure reinforces a pattern that can also be seen by watching the movie of the temperature field referenced above, namely that the simulation can be subdivided into three distinct phases. The first phase corresponds to the initial overturning of the unstable layering of the temperature field and is associated with a large spike in heat flux as well as large velocities (not shown here). The second phase, until approximately $t = 0.01$ corresponds to a relative hull: some plumes rise up, but not very fast because the medium is now stably layered but not fully mixed. This can be seen in the relatively low heat fluxes, but also in the fact that there are almost horizontal temperature isosurfaces in the second of the pictures in Fig. 13. After that, the general structure of the temperature field is that the interior of the domain is well mixed with a mostly constant average temperature and thin thermal boundary layers at the top and bottom from which plumes rise and sink. In this regime, the average heat flux is larger but also more variable depending on the number of plumes currently active. Many other analyses would be possible by using what is in the statistics file or by enabling additional postprocessors.

5.2.3 Convection in a box with prescribed, variable velocity boundary conditions

A similarly simple setup to the ones considered in the previous subsections is to equip the model we had with a different set of boundary conditions. There, we used slip boundary conditions, i.e., the fluid can flow tangentially along the four sides of our box but this tangential velocity is unspecified. On the other hand, in many situations, one would like to actually prescribe the tangential flow velocity as well. A typical application would be to use boundary conditions at the top that describe experimentally determined velocities of plates. This cookbook shows a simple version of something like this. To make it slightly more interesting, we choose a 2×1 domain in 2d.

Like for many other things, ASPECT has a set of plugins for prescribed velocity boundary values (see Sections A.3.2 and 6.3.6). These plugins allow one to write sophisticated models for the boundary velocity on parts or all of the boundary, but there is also one simple implementation that just takes a formula for the components of the velocity.

²²Note that the statistics file actually contains the *outward* heat flux for each of the six boundaries, which corresponds to the *negative* of upward flux for the bottom boundary. The figure therefore shows the negative of the values available in the statistics file.

To illustrate this, let us consider the `cookbooks/plate-like-boundary.prm` input file. It essentially extends the input file considered in the previous example. The part of this file that we are particularly interested in in the current context is the selection of the kind of velocity boundary conditions on the four sides of the box geometry, which we do using a section like this:

```

subsection Boundary velocity model
  set Tangential velocity boundary indicators = left, right, bottom
  set Prescribed velocity boundary indicators = top: function

  subsection Function
    set Variable names      = x,z,t
    set Function constants  = pi=3.1415926
    set Function expression = if(x>1+sin(0.5*pi*t), 1, -1); 0
  end
end

```

We use tangential flow at boundaries named left, right and bottom. Additionally, we specify a comma separated list (here with only a single element) of pairs consisting of the name of a boundary and the name of a prescribed velocity boundary model. Here, we use the `function` model on the `top` boundary, which allows us to provide a function-like notation for the components of the velocity vector at the boundary.

The second part we need is that we actually describe the function that sets the velocity. We do this in the subsection `Function`. The first of these parameters gives names to the components of the position vector (here, we are in 2d and we use `x` and `z` as spatial variable names) and the time. We could have left this entry at its default, `x,y,t`, but since we often think in terms of “depth” as the vertical direction, let us use `z` for the second coordinate. In the second parameter we define symbolic constants that can be used in the formula for the velocity that is specified in the last parameter. This formula needs to have as many components as there are space dimensions, separated by semicolons. As stated, this means that we prescribe the (horizontal) x -velocity and set the vertical velocity to zero. The horizontal component is here either 1 or -1 , depending on whether we are to the right or the left of the point $1 + \sin(\pi t/2)$ that is moving back and forth with time once every four time units. The `if` statement understood by the parser we use for these formulas has the syntax `if(condition, value-if-true, value-if-false)`.

Note: While you can enter most any expression into the parser for these velocity boundary conditions, not all make sense. In particular, if you use an incompressible medium like we do here, then you need to make sure that either the flow you prescribe is indeed tangential, or that at least the flow into and out of the boundary this function applies to is balanced so that in sum the amount of material in the domain stays constant.

It is in general not possible for ASPECT to verify that a given input is sensible. However, you will quickly find out if it isn't: The linear solver for the Stokes equations will simply not converge. For example, if your function expression in the input file above read

```
if(x>1+sin(0.5*pi*t), 1, -1); 1
```

then at the time of writing this you would get the following error message:

```
*** Timestep 0: t=0 seconds
Solving temperature system... 0 iterations.
Rebuilding Stokes preconditioner...
Solving Stokes system...

...some timing output ...
```

Exception on processing:

Iterative method reported convergence failure in step 9539 with residual 6.0552
Aborting!

The reason is, of course, that there is no incompressible (divergence free) flow field that allows for a constant vertical outflow component along the top boundary without corresponding inflow anywhere else.

The remainder of the setup is described in the following, complete input file:

```
#####
# Global parameters

set Dimension = 2
set Start time = 0
set End time = 20
set Use years in output instead of seconds = false
set Output directory = output-plate-like-boundary

#####
# Parameters describing the model
# Let us here choose again a box domain of size 2x1
# where we fix the temperature at the bottom and top,
# allow free slip along the bottom, left and right,
# and prescribe the velocity along the top using the
# 'function' description.

subsection Geometry model
  set Model name = box

  subsection Box
    set X extent = 2
    set Y extent = 1
  end
end
```

```

# We then set the temperature to one at the bottom and zero
# at the top:
subsection Boundary temperature model
    set Fixed temperature boundary indicators = bottom, top
    set List of model names = box

    subsection Box
        set Bottom temperature = 1
        set Top temperature     = 0
    end
end

# The velocity along the top boundary models a spreading
# center that is moving left and right:
subsection Boundary velocity model
    set Tangential velocity boundary indicators = left, right, bottom
    set Prescribed velocity boundary indicators = top: function

    subsection Function
        set Variable names      = x,z,t
        set Function constants  = pi=3.1415926
        set Function expression = if(x>1+sin(0.5*pi*t), 1, -1); 0
    end
end

# We then choose a vertical gravity model and describe the
# initial temperature with a vertical gradient. The default
# strength for gravity is one. The material model is the
# same as before.
subsection Gravity model
    set Model name = vertical
end

subsection Initial temperature model
    set Model name = function

    subsection Function
        set Variable names      = x,z
        set Function expression = (1-z)
    end
end

subsection Material model
    set Model name = simple

    subsection Simple model
        set Thermal conductivity      = 1e-6
        set Thermal expansion coefficient = 1e-4
        set Viscosity                 = 1
    end
end

```

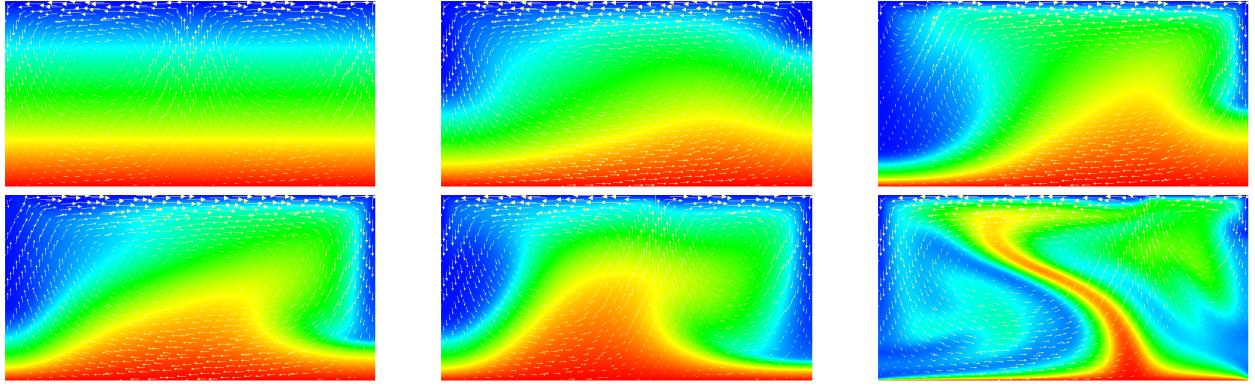


Figure 16: *Variable velocity boundary conditions: Temperature and velocity fields at the initial time (top left) and at various other points in time during the simulation.*

```
# The final part of this input file describes how many times the
# mesh is refined and what to do with the solution once computed
subsection Mesh refinement
    set Initial adaptive refinement      = 0
    set Initial global refinement       = 5
    set Time steps between mesh refinement = 0
end

subsection Postprocess
    set List of postprocessors = visualization, temperature statistics, heat flux statistics

    subsection Visualization
        set Time between graphical output = 0.1
    end
end
```

This model description yields a setup with a Rayleigh number of 200 (taking into account that the domain has size 2). It would, thus, be dominated by heat conduction rather than convection if the prescribed velocity boundary conditions did not provide a stirring action. Visualizing the results of this simulation²³ yields images like the ones shown in Fig. 16.

5.2.4 Using passive and active compositional fields

One frequently wants to track where material goes, either because one simply wants to see where stuff ends up (e.g., to determine if a particular model yields mixing between the lower and upper mantle) or because the material model in fact depends not only pressure, temperature and location but also on the mass fractions of certain chemical or other species. We will refer to the first case as *passive* and the latter as *active* to indicate the role of the additional quantities whose distribution we want to track. We refer to the whole process as *compositional* since we consider quantities that have the flavor of something that denotes the composition of the material at any given point.

²³In fact, the pictures are generated using a twice more refined mesh to provide adequate resolution. We keep the default setting of five global refinements in the parameter file as documented above to keep compute time reasonable when using the default settings.

There are basically two ways to achieve this: one can advect a set of particles (“tracers”) along with the velocity field, or one can advect along a field. In the first case, where the closest particle came from indicates the value of the concentration at any given position. In the latter case, the concentration(s) at any given position is simply given by the value of the field(s) at this location.

ASPECT implements both strategies, at least to a certain degree. In this cookbook, we will follow the route of advected fields.

The passive case. We will consider the exact same situation as in the previous section but we will ask where the material that started in the bottom 20% of the domain ends up, as well as the material that started in the top 20%. For the moment, let us assume that there is no material between the materials at the bottom, the top, and the middle. The way to describe this situation is to simply add the following block of definitions to the parameter file (you can find the full parameter file in [cookbooks/composition-passive.prm](#):

```
# This is the new part: We declare that there will
# be two compositional fields that will be
# advected along. Their initial conditions are given by
# a function that is one for the lowermost 0.2 height
# units of the domain and zero otherwise in the first case,
# and one in the top most 0.2 height units in the latter.
subsection Compositional fields
  set Number of fields = 2
end

subsection Initial composition model
  set Model name = function

  subsection Function
    set Variable names      = x,y
    set Function expression = if(y<0.2, 1, 0) ; if(y>0.8, 1, 0)
  end
end
```

Running this simulation yields results such as the ones shown in Fig. 17 where we show the values of the functions $c_1(\mathbf{x}, t)$ and $c_2(\mathbf{x}, t)$ at various times in the simulation. Because these fields were one only inside the lowermost and uppermost parts of the domain, zero everywhere else, and because they have simply been advected along with the flow field, the places where they are larger than one half indicate where material has been transported to so far.²⁴

Fig. 17 shows one aspect of compositional fields that occasionally makes them difficult to use for very long time computations. The simulation shown here runs for 20 time units, where every cycle of the spreading center at the top moving left and right takes 4 time units, for a total of 5 such cycles. While this is certainly no short-term simulation, it is obviously visible in the figure that the interface between the materials has diffused over time. Fig. 18 shows a zoom into the center of the domain at the final time of the simulation. The figure only shows values that are larger than 0.5, and it looks like the transition from red or blue to the edge of the shown region is no wider than 3 cells. This means that the computation is not overly diffusive but it is nevertheless true that this method has difficulty following long and thin filaments.²⁵ This is an area in which ASPECT may see improvements in the future.

²⁴Of course, this interpretation suggests that we could have achieved the same goal by encoding everything into a single function – that would, for example, have had initial values one, zero and minus one in the three parts of the domain we are interested in.

²⁵We note that this is no different for particles where the position of particles has to be integrated over time and is subject to numerical error. In simulations, their location is therefore not the exact one but also subject to a diffusive process resulting from numerical inaccuracies. Furthermore, in long thin filaments, the number of particles per cell often becomes too small and new particles have to be inserted; their properties are then interpolated from the surrounding particles, a process that also incurs a smoothing penalty.

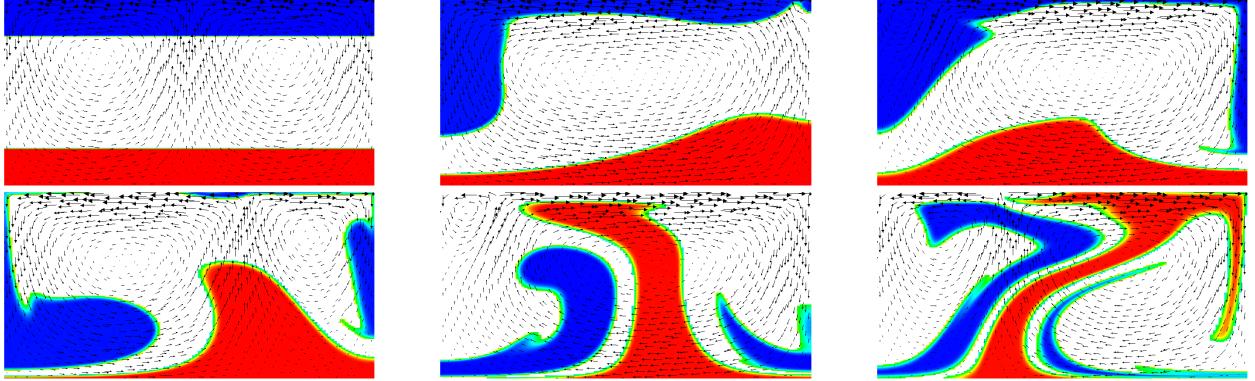


Figure 17: *Passive compositional fields:* The figures show, at different times in the simulation, the velocity field along with those locations where the first compositional field is larger than 0.5 (in red, indicating the locations where material from the bottom of the domain has gone) as well as where the second compositional field is larger than 0.5 (in blue, indicating material from the top of the domain). The results were obtained with two more global refinement steps compared to the [cookbooks/composition-passive.prm](#) input file.

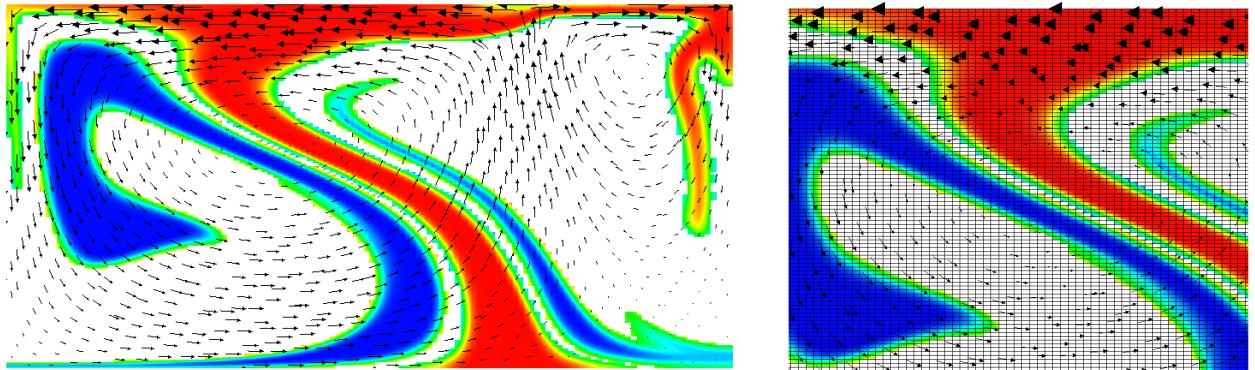


Figure 18: *Passive compositional fields:* A later image of the simulation corresponding to the sequence shown in Fig. 17 (left) and zoom-in on the center, also showing the mesh (right).

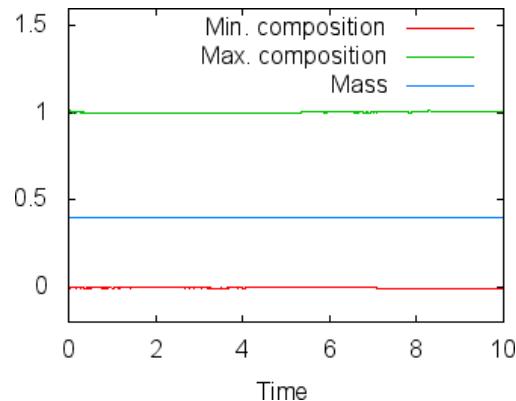


Figure 19: *Passive compositional fields:* Minimum and maximum of the first compositional variable over time, as well as the mass $m_1(t) = \int_{\Omega} c_1(\mathbf{x}, t)$ stored in this variable.

A different way of looking at the quality of compositional fields as opposed to particles is to ask whether they conserve mass. In the current context, the mass contained in the i th compositional field is $m_i(t) = \int_{\Omega} c_i(\mathbf{x}, t)$. This can easily be achieved in the following way, by adding the `composition statistics` postprocessor:

```
subsection Postprocess
  set List of postprocessors = visualization, temperature statistics, composition statistics
end
```

While the scheme we use to advect the compositional fields is not strictly conservative, it is almost perfectly so in practice. For example, in the computations shown in this section (using two additional global mesh refinements over the settings in the parameter file `cookbooks/composition-passive.prm`), Fig. 19 shows the maximal and minimal values of the first compositional fields over time, along with the mass $m_1(t)$ (these are all tabulated in columns of the statistics file, see Sections 4.1 and 4.4.2). While the maximum and minimum fluctuate slightly due to the instability of the finite element method in resolving discontinuous functions, the mass appears stable at a value of 0.403646 (the exact value, namely the area that was initially filled by each material, is 0.4; the difference is a result of the fact that we can't exactly represent the step function on our mesh with the finite element space). In fact, the maximal difference in this value between time steps 1 and 500 is only $1.1 \cdot 10^{-6}$. In other words, these numbers show that the compositional field approach is almost exactly mass conservative.

The active case. The next step, of course, is to make the flow actually depend on the composition. After all, compositional fields are not only intended to indicate where material come from, but also to indicate the properties of this material. In general, the way to achieve this is to write material models where the density, viscosity, and other parameters depend on the composition, taking into account what the compositional fields actually denote (e.g., if they simply indicate the origin of material, or the concentration of things like olivine, perovskite, ...). The construction of material models is discussed in much greater detail in Section 6.3.1; we do not want to revisit this issue here and instead choose – once again – the simplest material model that is implemented in ASPECT: the `simple` model.

The place where we are going to hook in a compositional dependence is the density. In the `simple` model, the density is fundamentally described by a material that expands linearly with the temperature; for small density variations, this corresponds to a density model of the form $\rho(T) = \rho_0(1 - \alpha(T - T_0))$. This is, by virtue of its simplicity, the most often considered density model. But the `simple` model also has a hook to make the density depend on the first compositional field $c_1(\mathbf{x}, t)$, yielding a dependence of the form $\rho(T) = \rho_0(1 - \alpha(T - T_0)) + \gamma c_1$. Here, let us choose $\rho_0 = 1, \alpha = 0.01, T_0 = 0, \gamma = 100$. The rest of our model setup will be as in the passive case above. Because the temperature will be between zero and one, the temperature induced density variations will be restricted to 0.01, whereas the density variation by origin of the material is 100. This should make sure that dense material remains at the bottom despite the fact that it is hotter than the surrounding material.²⁶

This setup of the problem can be described using an input file that is almost completely unchanged from the passive case. The only difference is the use of the following section (the complete input file can be found in `cookbooks/composition-active.prm`):

```
subsection Material model
  set Model name = simple

  subsection Simple model
    set Thermal conductivity = 1e-6
    set Thermal expansion coefficient = 0.01
```

²⁶The actual values do not matter as much here. They are chosen in such a way that the system – previously driven primarily by the velocity boundary conditions at the top – now also feels the impact of the density variations. To have an effect, the buoyancy induced by the density difference between materials must be strong enough to balance or at least approach the forces exerted by whatever is driving the velocity at the top.

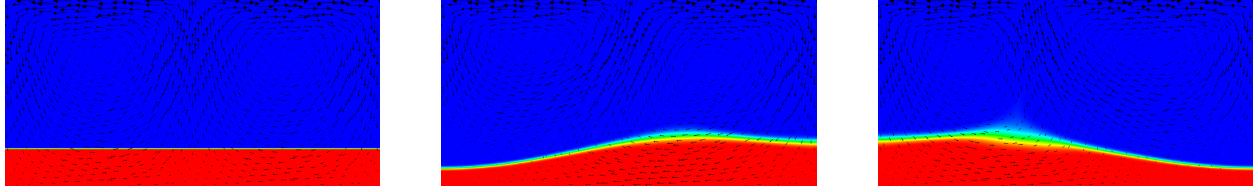


Figure 20: *Active compositional fields: Compositional field 1 at the time $t = 0, 10, 20$.* Compared to the results shown in Fig. 17 it is clear that the heavy material stays at the bottom of the domain now. The effect of the density on the velocity field is also clearly visible by noting that at all three times the spreading center at the top boundary is in exactly the same position; this would result in exactly the same velocity field if the density and temperature were constant.

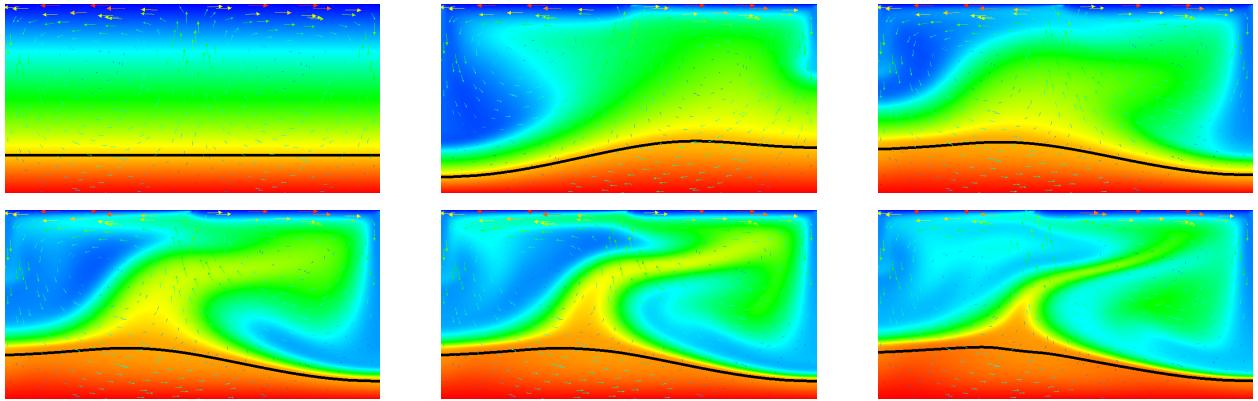


Figure 21: *Active compositional fields: Temperature fields at $t = 0, 2, 4, 8, 12, 20$. The black line is the isocontour line $c_1(\mathbf{x}, t) = 0.5$ delineating the position of the dense material at the bottom.*

```

set Viscosity          = 1
set Reference density = 1
set Reference temperature = 0
set Density differential for compositional field 1 = 0.1
end
end

```

To debug the model, we will also want to visualize the density in our graphical output files. This is done using the following addition to the postprocessing section, using the `density` visualization plugin:

```

subsection Postprocess
  set List of postprocessors = visualization, temperature statistics, composition statistics

  subsection Visualization
    set List of output variables = density
    set Time between graphical output = 0.1
  end
end

```

Results of this model are visualized in Fig.s 20 and 21. What is visible is that over the course of the simulation, the material that starts at the bottom of the domain remains there. This can only happen if the circulation is significantly affected by the high density material once the interface starts to become non-horizontal, and this is indeed visible in the velocity vectors. As a second consequence, if the material at the bottom does not move away, then there needs to be a different way for the heat provided at the bottom

to get through the bottom layer: either there must be a secondary convection system in the bottom layer, or heat is simply conducted. The pictures in the figure seem to suggest that the latter is the case.

It is easy, using the outline above, to play with the various factors that drive this system, namely:

- The magnitude of the velocity prescribed at the top.
- The magnitude of the velocities induced by thermal buoyancy, as resulting from the magnitude of gravity and the thermal expansion coefficient.
- The magnitude of the velocities induced by compositional variability, as described by the coefficient γ and the magnitude of gravity.

Using the coefficients involved in these considerations, it is trivially possible to map out the parameter space to find which of these effects is dominant. As mentioned in discussing the values in the input file, what is important is the *relative* size of these parameters, not the fact that currently the density in the material at the bottom is 100 times larger than in the rest of the domain, an effect that from a physical perspective clearly makes no sense at all.

The active case with reactions. *This section was contributed by Juliane Dannberg and René Gäßmöller.*

In addition, there are setups where one wants the compositional fields to interact with each other. One example would be material upwelling at a mid-ocean ridge and changing the composition to that of oceanic crust when it reaches a certain depth. In this cookbook, we will describe how this kind of behaviour can be achieved by using the `composition reaction` function of the material model.

We will consider the exact same setup as in the previous paragraphs, except for the initial conditions and properties of the two compositional fields. There is one material that initially fills the bottom half of the domain and is less dense than the material above. In addition, there is another material that only gets created when the first material reaches the uppermost 20% of the domain, and that has a higher density. This should cause the first material to move upwards, get partially converted to the second material, which then sinks down again. This means we want to change the initial conditions for the compositional fields:

```
subsection Initial composition model
  set Model name = function

  subsection Function
    set Variable names      = x,z
    set Function expression = if(z<0.5, 1, 0); 0
  end
end
```

Moreover, instead of the `simple` material model, we will use the `composition reaction` material model, which basically behaves in the same way, but can handle two active compositional field and a reaction between those two fields. In the input file, the user defines a depth and above this `reaction depth` the first compositional fields is converted to the second field. This can be done by changing the following section (the complete input file can be found in [cookbooks/composition-reaction.prm](#)).

```
subsection Material model
  set Model name = composition reaction

  subsection Composition reaction model
    set Thermal conductivity          = 1e-6
    set Thermal expansion coefficient = 0.01
    set Viscosity                   = 1
    set Density differential for compositional field 1 = -5
    set Density differential for compositional field 2 = 5
    set Reaction depth              = 0.2
  end
```

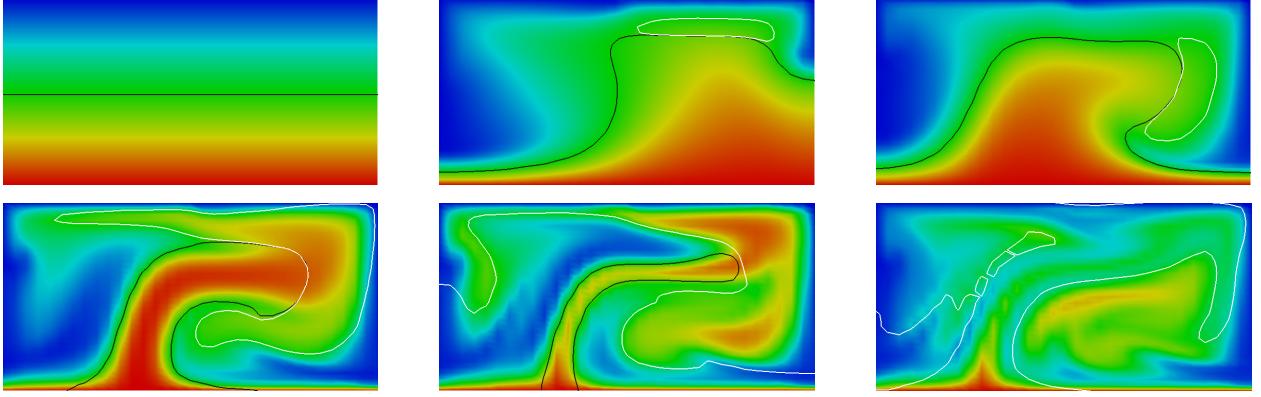


Figure 22: Reaction between compositional fields: Temperature fields at $t = 0, 2, 4, 8, 12, 20$. The black line is the isocontour line $c_1(\mathbf{x}, t) = 0.5$ delineating the position of the material starting at the bottom and the white line is the isocontour line $c_2(\mathbf{x}, t) = 0.5$ delineating the position of the material that is created by the reaction.

end

Results of this model are visualized in Fig 22. What is visible is that over the course of the simulation, the material that starts at the bottom of the domain ascends, reaches the reaction depth and gets converted to the second material, which starts to sink down.

5.2.5 Using particles

Using compositional fields to trace where material has come from or is going to has many advantages from a computational point of view. For example, the numerical methods to advect along fields are well developed and we can do so at a cost that is equivalent to one temperature solve for each of the compositional fields. Unless you have many compositional fields, this cost is therefore relatively small compared to the overall cost of a time step. Another advantage is that the value of a compositional field is well defined at every point within the domain. On the other hand, compositional fields over time diffuse initially sharp interfaces, as we have seen in the images of the previous section.

At the same time, the geodynamics community has a history of using particles for this purpose. Historically, this may have been because it is conceptually simpler to advect along individual particles rather than whole fields, since it only requires an ODE integrator rather than the stabilization techniques necessary to advect fields. They also provide the appearance of no diffusion, though this is arguable. Leaving aside the debate whether fields or particles are the way to go, ASPECT supports both: using fields and using particles.

In order to advect particles along with the flow field, one just needs to add the `particles` postprocessor to the list of postprocessors and specify a few parameters. We do so in the [cookbooks/composition-passive-particles.prm](#) input file, which is otherwise just a minor variation of the [cookbooks/composition-passive.prm](#) case discussed in the previous Section 5.2.4. In particular, the postprocess section now looks like this:

```

subsection Postprocess
  set List of postprocessors = visualization, particles

  subsection Visualization
    set Time between graphical output = 0.1
  end

  subsection Particles

```

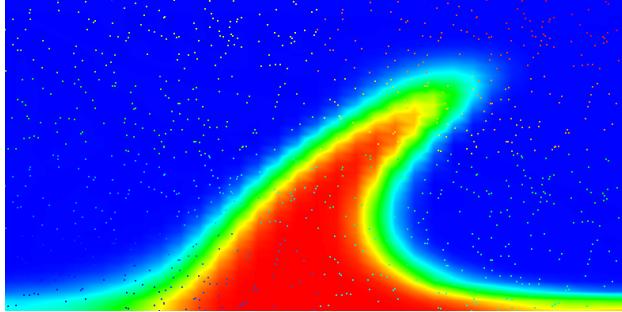


Figure 23: *Passively advected quantities visualized through both a compositional field and a set of 1,000 particles, at $t = 7.2$.*

```

set Number of particles      = 1000
set Time between data output = 0.1
set Data output format      = vtu
end
end

```

The 1000 particles we are asking here are initially uniformly distributed throughout the domain and are, at the end of each time step, advected along with the velocity field just computed. (There are a number of options to decide which method to use for advecting particles, see Section A.125.)

If you run this cookbook, information about all particles will be written into the output directory selected in the input file (as discussed in Section 4.1). In the current case, in addition to the files already discussed there, a directory listing at the end of a run will show several particle related files:

```

aspect> ls -l output/
total 932
-rw-rw-r-- 1 bangerth bangerth 11134 Dec 11 10:08 depth_average.gnuplot
-rw-rw-r-- 1 bangerth bangerth 11294 Dec 11 10:08 log.txt
-rw-rw-r-- 1 bangerth bangerth 326074 Dec 11 10:07 parameters.prm
-rw-rw-r-- 1 bangerth bangerth 577138 Dec 11 10:07 parameters.tex
drwxrwxr-x 2 bangerth bangerth 4096 Dec 11 18:40 particles
-rw-rw-r-- 1 bangerth bangerth 335 Dec 11 18:40 particles.pvd
-rw-rw-r-- 1 bangerth bangerth 168 Dec 11 18:40 particles.visit
drwxr-xr-x 2 bangerth bangerth 4096 Dec 11 10:08 solution
-rw-rw-r-- 1 bangerth bangerth 484 Dec 11 10:08 solution.pvd
-rw-rw-r-- 1 bangerth bangerth 451 Dec 11 10:08 solution.visit
-rw-rw-r-- 1 bangerth bangerth 8267 Dec 11 10:08 statistics

```

Here, the `particles.pvd` and `particles.visit` files contain a list of all visualization files from all processors and time steps. These files can be loaded in much the same way as the `solution.pvd` and `solution.visit` files that were discussed in Section 4.4. The actual data files – possibly a large number, but not of much immediate interest to users – are located in the `particles` subdirectory.

Coming back to the example at hand, we can visualize the particles that were advected along by opening both the field-based output files and the ones that correspond to particles (for example, `output/solution.visit` and `output/particles.visit`) and using a pseudo-color plot for the particles, selecting the “id” of particles to color each particle. By going to, for example, the output from the 72nd visualization output, this then results in a plot like the one shown in Fig. 23.

The particles shown here are not too impressive in still pictures since they are colorized by their particle number, which does not carry any particular meaning other than the fact that it enumerates the particles.²⁷

²⁷Particles are enumerated in a way so that first the first processor in a parallel computations numbers all of the particles

The particle “id” can, however, be useful when viewing an animation of time steps. There, the different colors of particles allows the eye to follow the motion of a single particle. This is especially true if, after some time, particles have become well mixed by the flow field and adjacent particles no longer have similar colors. In any case, viewing such animations makes it rather intuitive to understand a flow field, but it can of course not be reproduced in a static medium such as this manual.

In any case, we will see in the next section how to attach more interesting information to particles, and how to visualize these.

Using particle properties. The particles in the above example only fulfil the purpose of visualizing the convection pattern. A more meaningful use for particles is to attach “properties” to them. A property consists of one or more numbers (or vectors or tensors) that may either be set at the beginning of the model run and stay constant, or are updated during the model runtime. These properties can then be used for many applications, e.g., storing an initial property (like the position, or initial composition), evaluating a property at a defined particle path (like the pressure-temperature evolution of a certain piece of rock), or by integrating a quantity along a particle path (like the integrated strain a certain domain has experienced). We illustrate these properties in the cookbook [cookbooks/composition-passive-particles-properties.prm](#), in which we add the following lines to the `Particles` subsection (we also increase the number of particles compared to the previous section to make the visualizations below more obvious):

```
subsection Postprocess
  subsection Particles
    set Number of particles      = 50000

    set List of particle properties = function, initial composition, initial position, pT path

    subsection Function
      set Variable names      = x,y
      set Function expression = if(y<0.2, 1, 0)
    end
  end
end
```

These commands make sure that every particle will carry four different properties (`function`, `pT path`, `initial position` and `initial composition`), some of which may be scalars and others can have multiple components. (A full list of particle properties that can currently be selected can be found in Section [A.125](#), and new particle properties can be added as plugins as described in Section [6.2](#).) The properties selected above do the following:

- **initial position:** This particle property simply stores the initial coordinates of the particle and then never changes them. This can be useful to compare the final position of a particle with its initial position and therefore determine how far certain domains traveled during the model runtime. Alternatively, one may want to simply visualize the norm of this vector property (i.e., the norm of the initial position, which is of course equal to the distance from the origin at which a particle started): in mantle simulations in spherical coordinates, the radius is indicative of which part of the mantle a particle comes from, and can therefore be used to visualize where material gets transported over the course of a simulation.
- **initial composition:** This property uses the same method to initialize particle properties as is used to initialize the corresponding compositional fields. Using this, it stores the compositional field initialization values at the location where the particle started, and again never changes them. This is

on its first cell, then its second cell, and so on; then the second processor does the same with particles in the order of the cells it owns; etc. Thus, the “id” shown in the picture is not just a random number, but rather shows the order of cells and how they belonged to the processors that participated in the computation at the time when particles were created. After some time, particles may of course have become well mixed. In any case, this ordering is of no real practical use.

useful in the same context as shown for the field-based example in Section 5.2.4 where we would like to figure where materials ends up. In this case, one would set the initial composition to an indicator function for certain parts of the domain, and then set the initial composition property for the particles to match this composition. Letting the particles advect and at a later time visualizing this particle property will then show where particles came from. In cases where compositional variables undergo changes, e.g., by describing phase changes or chemical reactions, the “initial composition” property can also be useful to compare the final composition of a particle with its initial composition and therefore determine which regions underwent reactions such as those described in Section 5.2.4, and where the material that underwent this reaction got transported to.

- **function:** This particle property can be used to assign to each particle values that are described based on a function of space. It provides an alternative way to set initial values if you don’t want to first set a compositional field’s initial values based on a function, and then copy these values via the “initial composition” property to particles. In the example above, we use the same function as for the compositional initial composition of field number one in Section 5.2.4. Therefore, this property should behave identical to the compositional field (except that the compositional field may have a reaction term that this particle property does not), although the two are of course advected using very different methods. This allows to compare the error in particle position to the numerical diffusion of the compositional field.
- **pT path:** This property is interesting in that the particle property’s values always exactly mirror the pressure and temperature at the particle’s current location. This does not seem to be very useful since the information is already available. However, because each particle has a unique id, one can select a particular particle and output its properties (including pressure and temperature based on the pT path property) at all time steps. This allows for the creation of a pressure-temperature curve of a certain piece of rock. This property is interesting in many lithosphere and crustal scale models, because it is determining the metamorphic reactions that happen during deformation processes (e.g., in a subduction zone).

The results of all of these properties can of course be visualized. Fig. 24 shows some of the pictures one can create with particles. The top row shows both the composition field C_1 (along with the mesh on which it is defined) and the corresponding “initial C_1 ” particle property, at $t = 7.2$. Because the compositional field does not undergo any reactions, it should of course simply be the initial composition advected along with the flow field, and therefore equal the values of the corresponding particle property. However, field-based compositions suffer from diffusion. On the other hand, the amount of diffusion can easily be decreased by mesh refinement.

The bottom of the figure shows the norm of the “initial position” property at the initial time and at time $t = 20$. These images therefore show how far from the origin each of the particles shown was at the initial time.

Using active particles. In the examples above, particle properties passively track distinct model properties. These particle properties, however, may also be used to actively influence the model as it runs. For instance, a composition-dependent material model may use particles’ initial composition rather than an advected compositional field. To make this work – i.e., to get information from particles that are located at unpredictable locations, to the quadrature points at which material models and other parts of the code need to evaluate these properties – we need to somehow get the values from particles back to fields that can then be evaluated at any point where this is necessary. A slightly modified version of the active-composition cookbook ([cookbooks/composition-active.prm](#)) illustrates how to use ‘active particles’ in this manner.

This cookbook, [cookbooks/composition-active-particles.prm](#), modifies two sections of the input file. First, particles are added under the Postprocess section:

```
subsection Postprocess
  subsection Particles
```

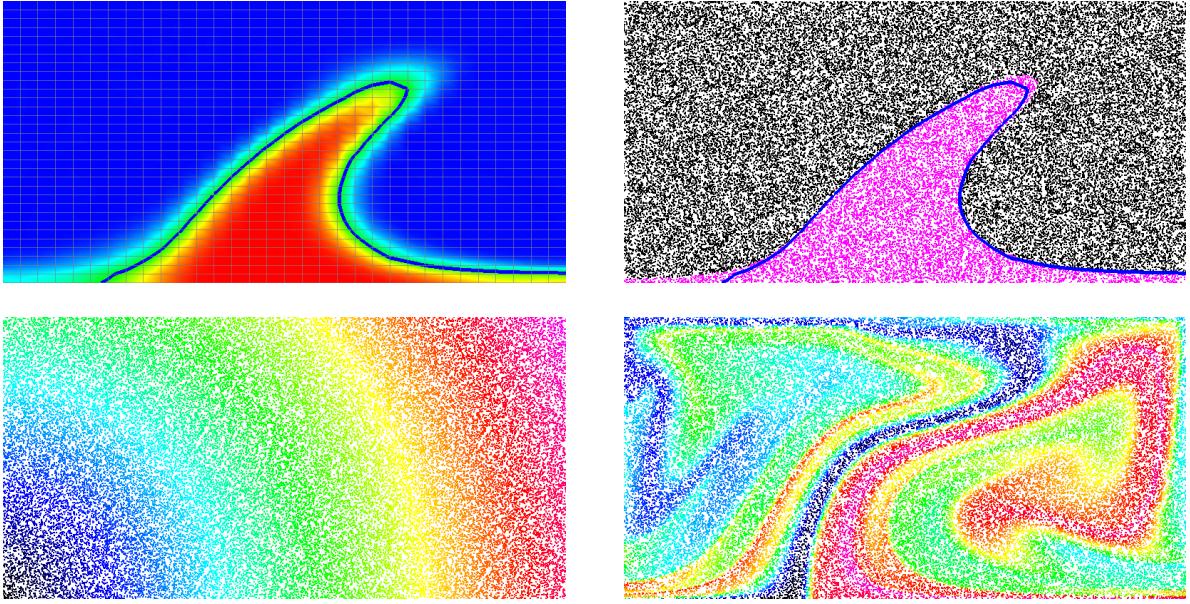


Figure 24: *Passively advected particle properties visualized.* Top row: Composition C_1 and particle property “initial C_1 ”. The blue line in both figures is the 0.5-isocontour for the C_1 field. Bottom row: Norm of the “initial position” of particles at $t = 0$ and $t = 20$.

```

set Number of particles      = 100000
set Time between data output = 0
set Data output format      = vtu
set List of particle properties = velocity, initial composition
set Interpolation scheme    = cell average
set Particle generator name = random uniform
end
end

```

Here, each particle will carry the `velocity` and `initial composition` properties. In order to use the particle `initial composition` value to modify the flow through the material model, we now modify the `Composition` section:

```

subsection Compositional fields
  set Number of fields      = 2
  set Names of fields       = lower, upper
  set Compositional field methods = particles, particles
  set Mapped particle properties = lower:initial lower, upper:initial upper
end

```

What this does is the following: It says that there will be two compositional fields, called `lower` and `upper` (because we will use them to indicate material that comes from either the lower or upper part of the domain). Next, the `Compositional field methods` states that each of these fields will be computed by interpolation from the particles (if we had left this parameter at its default value, `field`, for each field, then it would have solved an advection PDE in each time step, as we have done in all previous examples).

In this case, we specify that both of the compositional fields are in fact interpolated from particle properties in each time step. How this is done is described in the fourth line. To understand it, it is important to realize that particles and fields have matching names: We have named the fields `lower` and `upper`, whereas the properties that result from the `initial composition` entry in the particles section are

called `initial lower` and `initial upper`, since they inherit the names of the fields.

The syntax for interpolation from particles to fields then states that the `lower` field will be set to the interpolated value of the `initial lower` particle property at the end of each time step, and similarly for the `upper` field. In turn, the `initial composition` particle property was using the same method that one would have used for the compositional field initialization if these fields were actually advected along in each time step.

In this model the given global refinement level (5), associated number of cells (1024) and 100,000 total particles produces an average particle-per-cell count slightly below 100. While on the high end compared to most geodynamic studies using active particles, increasing the number of particles per cell further may alter the solution. As with the numerical resolution, any study using active particles should systematically vary the number of particles per cell in order to determine this parameter's influence on the simulation.

Note: ASPECT's particle implementation is in a preliminary state. While the accuracy and scalability of the implementation is benchmarked, other limitations remain. This in particular means that it is not optimized for performance, and more than a few thousand particles per process can slow down a model significantly. Moreover, models with a highly adaptive mesh and many particles do encounter a significant slowdown, because ASPECT only considers the number of degrees of freedom for load balancing across processes and not the number of particles. Therefore processes that compute the solution for coarse-grid regions have to process many more particles than other processes. Additionally, the checkpoint/restart functionality for particles is only implemented in models with a constant number of processes before and after the checkpoint and when the selected particle properties do not change. These limitations might be removed over time, but for current models the user should be aware of them.

5.2.6 Using a free surface

This section was contributed by Ian Rose.

Free surfaces are numerically challenging but can be useful for self consistently tracking dynamic topography and may be quite important as a boundary condition for tectonic processes like subduction. The parameter file [cookbooks/free-surface.prm](#) provides a simple example of how to set up a model with a free surface, as well as demonstrates some of the challenges associated with doing so.

ASPECT supports models with a free surface using an Arbitrary Lagrangian-Eulerian framework (see Section 2.12). Most of this is done internally, so you do not need to worry about the details to run this cookbook. Here we demonstrate the evolution of surface topography that results when a blob of hot material rises in the mantle, pushing up the free surface as it does. Usually the amplitude of free surface topography will be small enough that it is difficult to see with the naked eye in visualizations, but the `topography` postprocessor can help by outputting the maximum and minimum topography on the free surface at every time step.

The bulk of the parameter file for this cookbook is similar to previous ones in this manual. We use initial temperature conditions that set up a hot blob of rock in the center of the domain.

The main addition is the `Free surface` subsection. In this subsection you need to give ASPECT a comma separated list of the free surface boundary indicators. In this case, we are dealing with the 'top' boundary of a box in 2D. There is another significant parameter that needs to be set here: the value for the stabilization parameter "theta". If this parameter is zero, then there is no stabilization, and you are likely to see instabilities develop in the free surface. If this parameter is one then it will do a good job of stabilizing the free surface, but it may overly damp its motions. The default value is 0.5.

Also worth mentioning is the change to the CFL number. Stability concerns typically mean that when making a model with a free surface you will want to take smaller time steps. In general just how much smaller will depend on the problem at hand as well as the desired accuracy.

Following are the sections in the input file specific to this testcase. The full parameter file may be found at [cookbooks/free-surface.prm](#).

```

set CFL number          = 0.1

subsection Initial temperature model
  set Model name = function
  subsection Function
    set Variable names      = x,y
    set Function expression = if( sqrt( (x-250.e3)^2 + (y-100.e3)^2 ) < 25.e3, 200.0, 0.0)
  end
end

subsection Free surface
  set Free surface boundary indicators      = top
  set Free surface stabilization theta = 0.5
end

subsection Boundary temperature model
  set Fixed temperature boundary indicators = left, right, bottom, top
end

subsection Boundary velocity model
  set Tangential velocity boundary indicators = left, right, bottom
end

subsection Postprocess
  set List of postprocessors = visualization,topography,velocity statistics,
  subsection Visualization
    set Time between graphical output = 1.e6
  end
end

```

Running this input file will produce results like those in Figure 25. The model starts with a single hot blob of rock which rises in the domain. As it rises, it pushes up the free surface in the middle, creating a topographic high there. This is similar to the kind of dynamic topography that you might see above a mantle plume on Earth. As the blob rises and diffuses, it loses the buoyancy to push up the boundary, and the surface begins to relax.

After running the cookbook, you may modify it in a number of ways:

- Add a more complicated initial temperature field to see how that affects topography.
- Add a high-viscosity lithosphere to the top using a compositional field to tamp down on topography.
- Explore different values for the stabilization theta and the CFL number to understand the nature of when and why stabilization is necessary.
- Try a model in a different geometry, such as spherical shells.

5.2.7 Using a free surface in a model with a crust

This section was contributed by William Durkin.

This cookbook is a modification of the previous example that explores changes in the way topography develops when a highly viscous crust is added. In this cookbook, we use a material model in which the material changes from low viscosity mantle to high viscosity crust at $z = z_j = \text{jump height}$, i.e., the piecewise viscosity function is defined as

$$\eta(z) = \begin{cases} \eta_U & \text{for } z > z_j, \\ \eta_L & \text{for } z \leq z_j. \end{cases}$$



Figure 25: *Evolution of surface topography due to a rising blob.* On the left is a snapshot of the model setup. The right shows the value of the highest topography in the domain over 18 Myr of model time. The topography peaks at 165 meters after 5.2 Myr. This cookbook may be run with the [cookbooks/free-surface.prm](#) input file.

where η_U and η_L are the viscosities of the upper and lower layers, respectively. This viscosity model can be implemented by creating a plugin that is a small modification of the `simpler` material model (from which it is otherwise simply copied). We call this material model “SimplerWithCrust”. In particular, what is necessary is an evaluation function that looks like this:

```
template <int dim>
void
SimplerWithCrust<dim>::
evaluate(const typename Interface<dim>::MaterialModelInputs &in,
         typename Interface<dim>::MaterialModelOutputs &out) const
{
    for (unsigned int i=0; i<in.position.size(); ++i)
    {
        const double z = in.position[i][1];

        if (z>jump_height)
            out.viscosities[i] = eta_U;
        else
            out.viscosities[i] = eta_L;

        out.densities[i] = reference_rho*(1.0-thermal_alpha*(in.temperature[i]-reference_T));
        out.thermal_expansion_coefficients[i] = thermal_alpha;
        out.specific_heat[i] = reference_specific_heat;
        out.thermal_conductivities[i] = k_value;
        out.compressibilities[i] = 0.0;
    }
}
```

Additional changes make the new parameters `Jump height`, `Lower viscosity`, and `Upper viscosity` available to the input parameter file, and corresponding variables available in the class and used in the code snippet above. The entire code can be found in [cookbooks/free-surface-with-crust/plugin/simpler-with-crust.cc](#).

[cc.](#) Refer to Section 6.1 for more information about writing and running plugins.

The following changes are necessary compared to the input file from the cookbook shown in Section 5.2.6 to include a crust:

- Load the plugin implementing the new material model:

<code>set Additional shared libraries</code>	<code>= ./plugin/libsimpler-with-crust.so</code>
----------------------------------------------	--------------------------------------------------

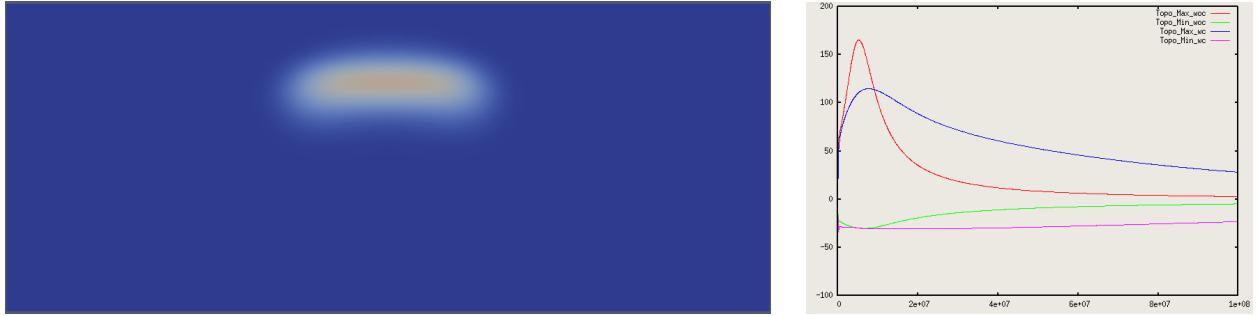


Figure 26: *Adding a viscous crust to a model with surface topography. The thermal anomaly spreads horizontally as it collides with the highly viscous crust (left). The addition of a crustal layer both dampens and delays the appearance of the topographic maximum and minimum (right).*

- Declare values for the new parameters:

```

subsection Material model
  set Model name = simpler with crust
  subsection Simpler with crust model
    set Reference density          = 3300
    set Reference specific heat   = 1250
    set Reference temperature     = 0.0
    set Thermal conductivity      = 1.0 # low thermal conductivity for a sharp blob
    set Thermal expansion coefficient = 4e-5

    # Parameters added for this cookbook:
    # The box is 200km high and has its origin set at the bottom left corner.
    # Setting the jump height to 170km creates a 30km thick crust
    set Lower viscosity           = 1e20
    set Upper viscosity           = 1e23
    set Jump height               = 170e3
  end
end

```

Note that the height of the interface at 170km is interpreted in the coordinate system in which the box geometry of this cookbook lives. The box has dimensions 500km × 200km, so an interface height of 170km implies a depth of 30km.

The entire script is located in [cookbooks/free-surface-with-crust/free-surface-with-crust.prm](#). Running this input file yields a crust that is 30km thick and 1000 times as viscous as the lower layer. Figure 26 shows that adding a crust to the model causes the maximum topography to both decrease and occur at a later time. Heat flows through the system primarily by advection until the temperature anomaly reaches the base of the crustal layer (approximately at the time for which Fig 26 shows the temperature profile). The crust's high viscosity reduces the temperature anomaly's velocity substantially, causing it to affect the surface topography at a later time. Just as the cookbook shown in Section 5.2.6, the topography returns to zero after some time.

5.2.8 Averaging material properties

The original motivation for the functionality discussed here, as well as the setup of the input file, were provided by Cedric Thieulot.

Geophysical models are often characterized by abrupt and large jumps in material properties, in particular in the viscosity. An example is a subducting, cold slab surrounded by the hot mantle: Here, the strong

temperature-dependence of the viscosity will lead to a sudden jump in the viscosity between mantle and slab. The length scale over which this jump happens will be a few or a few tens of kilometers. Such length scales cannot be adequately resolved in three-dimensional computations with typical meshes for global computations.

Having large viscosity variations in models poses a variety of problems to numerical computations. First, you will find that they lead to very long compute times because our solvers and preconditioners break down. This may be acceptable if it would at least lead to accurate solutions, but large viscosity gradients lead also to large pressure gradients, and this in turn leads to over- and undershoots in the numerical approximation of the gradient. We will demonstrate both of these issues experimentally below.

One of the solution to such problems is the realization that one can mitigate some of the effects by averaging material properties on each cell somehow (see, for example, [SBE⁺08, DK08, DMGT11, Thi15, TMK14]). Before going into detail, it is important to realize that if we choose material properties not per quadrature point when doing the integrals for forming the finite element matrix, but per cell, then we will lose accuracy in the solution in those cases where the solution is smooth. More specifically, we will likely lose one or more orders of convergence. In other words, it would be a bad idea to do this averaging unconditionally. On the other hand, if the solution has essentially discontinuous gradients and kinks in the velocity field, then at least at these locations we cannot expect a particularly high convergence order anyway, and the averaging will not hurt very much either. In cases where features of the solution that are due to strongly varying viscosities or other parameters, dominate, we may then as well do the averaging per cell.

To support such cases, ASPECT supports an operation where we evaluate the material model at every quadrature point, given the temperature, pressure, strain rate, and compositions at this point, and then either (i) use these values, (ii) replace the values by their arithmetic average $\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$, (iii) replace the values by their harmonic average $\bar{x} = \left(\frac{1}{N} \sum_{i=1}^N \frac{1}{x_i} \right)^{-1}$, (iv) replace the values by their geometric average $\bar{x} = \left(\prod_{i=1}^N \frac{1}{x_i} \right)^{-1/N}$, or (v) replace the values by the largest value over all quadrature points on this cell. Option (vi) is to project the values from the quadrature points to a bi- (in 2d) or trilinear (in 3d) Q_1 finite element space on every cell, and then evaluate this finite element representation again at the quadrature points. Unlike the other five operations, the values we get at the quadrature points are not all the same here.

We do this operation for all quantities that the material model computes, i.e., in particular, the viscosity, the density, the compressibility, and the various thermal and thermodynamic properties. In the first 4 cases, the operation guarantees that the resulting material properties are bounded below and above by the minimum and maximum of the original data set. In the last case, the situation is a bit more complicated: The nodal values of the Q_1 projection are not necessarily bounded by the minimal or maximal original values at the quadrature points, and then neither are the output values after re-interpolation to the quadrature points. Consequently, after projection, we limit the nodal values of the projection to the minimal and maximal original values, and only then interpolate back to the quadrature points.

We demonstrate the effect of all of this with the “sinker” benchmark. This benchmark is defined by a high-viscosity, heavy sphere at the center of a two-dimensional box. This is achieved by defining a compositional field that is one inside and zero outside the sphere, and assigning a compositional dependence to the viscosity and density. We run only a single time step for this benchmark. This is all modeled in the following input file that can also be found in [cookbooks/sinker-with-averaging/sinker-with-averaging.prm](#):

```

set Dimension          = 2
set Start time        = 0
set End time          = 0
set Output directory  = output_sinker-with-averaging

set Pressure normalization = volume

subsection Geometry model

```

```

set Model name = box
subsection Box
  set X extent   = 1.0000
  set Y extent   = 1.0000
end
end

subsection Boundary velocity model
  set Zero velocity boundary indicators      = left, right, bottom, top
end

subsection Material model
  set Model name = simple

  subsection Simple model
    set Reference density          = 1
    set Viscosity                  = 1
    set Thermal expansion coefficient = 0
    set Composition viscosity prefactor = 1e6
    set Density differential for compositional field 1 = 10
  end

  set Material averaging = none
end

subsection Gravity model
  set Model name = vertical
  subsection Vertical
    set Magnitude = 1
  end
end

#####
# Parameters describing the temperature field
# Note: The temperature plays no role in this model

subsection Boundary temperature model
  set List of model names = box
end

subsection Initial temperature model
  set Model name = function
  subsection Function
    set Function expression = 0
  end
end

#####
# Parameters describing the compositional field
# Note: The compositional field is what drives the flow
# in this example

subsection Compositional fields
  set Number of fields = 1
end

```

```

subsection Initial composition model
  set Model name = function
  subsection Function
    set Variable names      = x,y
    set Function expression = if( (sqrt((x-0.5)^2+(y-0.5)^2)>0.22) , 0 , 1 )
  end
end

#####
# Parameters describing the discretization

subsection Mesh refinement
  set Initial global refinement      = 6
  set Initial adaptive refinement   = 0
end

#####
# Parameters describing what to do with the solution

subsection Postprocess
  set List of postprocessors = visualization, velocity statistics, composition statistics
  subsection Visualization
    set Output format           = vtu
    set Time between graphical output = 0
    set List of output variables = density, viscosity
  end
end

```

The type of averaging on each cell is chosen using this part of the input file:

```

subsection Material model
  set Material averaging = harmonic average
end

```

For the various different averaging options, and for different levels of mesh refinement, Fig. 27 shows pressure plots that illustrate the problem with oscillations of the discrete pressure. The important part of these plots is not that the solution looks discontinuous – in fact, the exact solution is discontinuous at the edge of the circle²⁸ – but the spikes that go far above and below the “cliff” in the pressure along the edge of the circle. Without averaging, these spikes are obviously orders of magnitude larger than the actual jump height. The spikes do not disappear under mesh refinement nor averaging, but they become far less pronounced with averaging. The results shown in the figure do not really allow to draw conclusions as to which averaging approach is the best; a discussion of this question can also be found in [SBE⁺08, DK08, DMGT11, TMK14]).

A very pleasant side effect of averaging is that not only does the solution become better, but it also becomes cheaper to compute. Table 2 shows the number of outer GMRES iterations when solving the Stokes equations (1)–(2).²⁹ The implication of these results is that the averaging gives us a solution that not only

²⁸This is also easy to try experimentally – use the input file from above and select 5 global and 10 adaptive refinement steps, with the refinement criteria set to `density`, then visualize the solution.

²⁹The outer iterations are only part of the problem. As discussed in [KHB12], each GMRES iteration requires solving a linear system with the elliptic operator $-\nabla \cdot 2\eta\varepsilon(\cdot)$. For highly heterogeneous models, such as the one discussed in the current section, this may require a lot of Conjugate Gradient iterations. For example, for 8 global refinement steps, the 30+188 outer iterations without averaging shown in Table 2 require a total of 22,096 inner CG iterations for the elliptic block (and a total of 837 for the approximate Schur complement). Using harmonic averaging, the 30+26 outer iterations require only 1258 iterations on the elliptic block (and 84 on the Schur complement). In other words, the number of inner iterations per outer iteration (taking into account the split into “cheap” and “expensive” outer iterations, see [KHB12]) is reduced from 117 to 47 for the elliptic block and from 3.8 to 1.5 for the Schur complement.

reduces the degree of pressure over- and undershoots, but is also significantly faster to compute: for example, the total run time for 8 global refinement steps is reduced from 5,250s for no averaging to 358s for harmonic averaging.

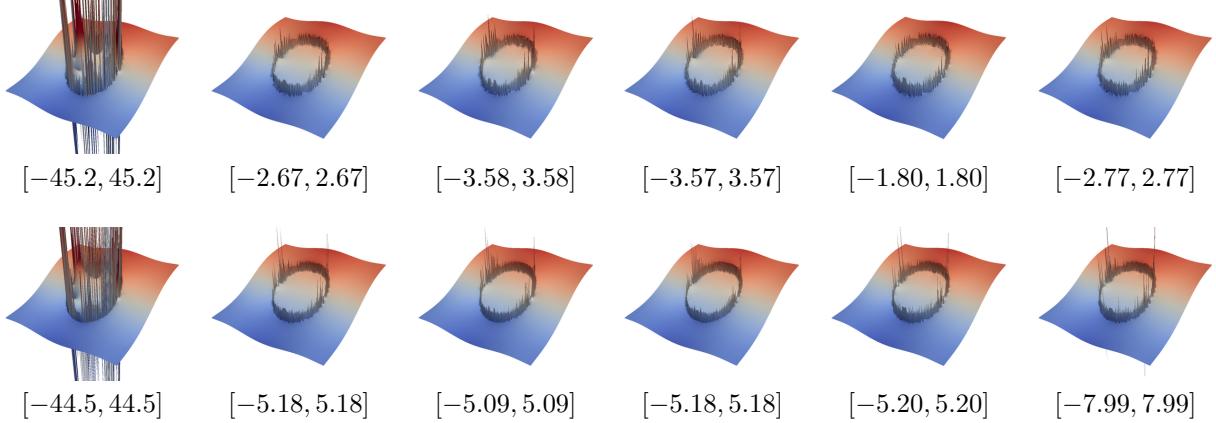


Figure 27: *Visualization of the pressure field for the “sinker” problem. Left to right: No averaging, arithmetic averaging, harmonic averaging, geometric averaging, pick largest, project to Q_1 . Top: 7 global refinement steps. Bottom: 8 global refinement steps. The minimal and maximal pressure values are indicated below every picture. This range is symmetric because we enforce that the average of the pressure equals zero. The color scale is adjusted to show only values between $p = -3$ and $p = 3$.*

# of global refinement steps	no averaging	arithmetic averaging	harmonic averaging	geometric averaging	pick largest	project to Q_1
4	30+64	30+13	30+10	30+12	30+13	30+15
5	30+87	30+14	30+13	30+14	30+14	30+16
6	30+171	30+14	30+15	30+14	30+15	30+17
7	30+143	30+27	30+28	30+26	30+26	30+28
8	30+188	30+27	30+26	30+27	30+28	30+28

Table 2: *Number of outer GMRES iterations to solve the Stokes equations for various numbers of global mesh refinement steps and for different material averaging operations. The GMRES solver first tries to run 30 iterations with a cheaper preconditioner before switching to a more expensive preconditioner (see Section A.1).*

Such improvements carry over to more complex and realistic models. For example, in a simulation of flow under the East African Rift by Sarah Stamps, using approximately 17 million unknowns and run on 64 processors, the number of outer and inner iterations is reduced from 169 and 114,482 without averaging to 77 and 23,180 with harmonic averaging, respectively. This translates into a reduction of run-time from 145 hours to 17 hours. Assessing the accuracy of the answers is of course more complicated in such cases because we do not know the exact solution. However, the results without and with averaging do not differ in any significant way.

A final comment is in order. First, one may think that the results should be better in cases of discontinuous pressures if the numerical approximation actually allowed for discontinuous pressures. This is in fact possible: We can use a finite element in which the pressure space contains piecewise constants (see Section A.38). To do so, one simply needs to add the following piece to the input file:

```
subsection Discretization
  set Use locally conservative discretization = true
end
```

Disappointingly, however, this makes no real difference: the pressure oscillations are no better (maybe even worse) than for the standard Stokes element we use, as shown in Fig. 28 and Table 3. Furthermore, as shown in Table 4, the iteration numbers are also largely unaffected if any kind of averaging is used – though they are far worse using the locally conservative discretization if no averaging has been selected. On the positive side, the visualization of the discontinuous pressure finite element solution makes it much easier to see that the true pressure is in fact discontinuous along the edge of the circle.

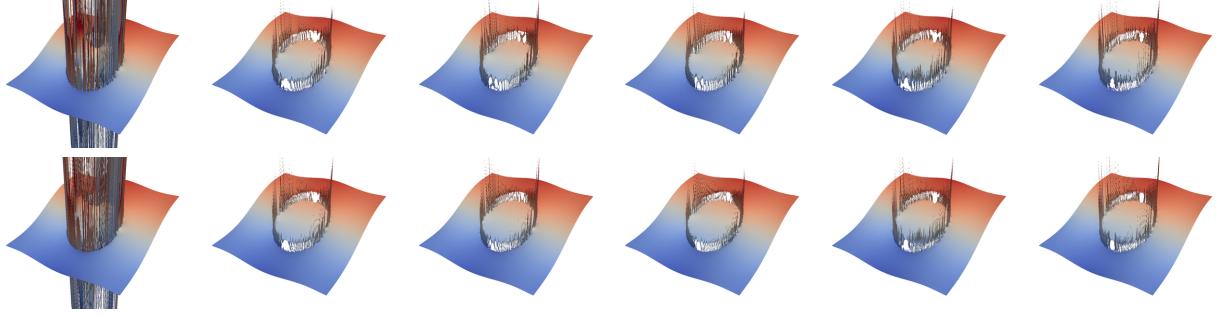


Figure 28: *Visualization of the pressure field for the “sinker” problem. Like Fig. 27 but using the locally conservative, enriched Stokes element. Pressure values are shown in Table 3.*

# of global refinement steps	no averaging	arithmetic averaging	harmonic averaging	geometric averaging	pick largest	project to Q_1
4	66.32	2.66	2.893	1.869	3.412	3.073
5	81.06	3.537	4.131	3.997	3.885	3.991
6	75.98	4.596	4.184	4.618	4.568	5.093
7	84.36	4.677	5.286	4.362	4.635	5.145
8	83.96	5.701	5.664	4.686	5.524	6.42

Table 3: *Maximal pressure values for the “sinker” benchmark, using the locally conservative, enriched Stokes element. The corresponding pressure solutions are shown in Fig. 28.*

# of global refinement steps	no averaging	arithmetic averaging	harmonic averaging	geometric averaging	pick largest	project to Q_1
4	30+376	30+16	30+12	30+14	30+14	30+17
5	30+484	30+16	30+14	30+14	30+14	30+16
6	30+583	30+16	30+17	30+14	30+17	30+17
7	30+1319	30+27	30+28	30+26	30+28	30+28
8	30+1507	30+28	30+27	30+28	30+28	30+29

Table 4: *Like Table 2, but using the locally conservative, enriched Stokes element.*

5.2.9 Prescribed internal velocity constraints

This section was contributed by Jonathan Perry-Houts

In cases where it is desirable to investigate the behavior of one part of the model domain but the controlling physics of another part is difficult to capture, such as corner flow in subduction zones, it may be useful to force the desired behavior in some parts of the model domain and solve for the resulting flow everywhere else. This is possible through the use of ASPECT’s “signal” mechanism, as documented in Section 6.5.

Internally, ASPECT adds “constraints” to the finite element system for boundary conditions and hanging nodes. These are places in the finite element system where certain solution variables are required to match some prescribed value. Although it is somewhat mathematically inadmissible to prescribe constraints on nodes inside the model domain, Ω , it is nevertheless possible so long as the prescribed velocity field fits in to the finite element’s solution space, and satisfies the other constraints (i.e., is divergence free).

Using ASPECT’s signals mechanism, we write a shared library which provides a “slot” that listens for the signal which is triggered after the regular model constraints are set, but before they are “distributed.”

As an example of this functionality, below is a plugin which allows the user to prescribe internal velocities with functions in a parameter file:

```
#include <deal.II/base/parameter_handler.h>
#include <deal.II/base/parsed_function.h>
#include <deal.II/fe/fe_values.h>
#include <aspect/global.h>
#include <aspect/simulator_signals.h>

namespace aspect
{
    using namespace dealii;

    // Global variables (to be set by parameters)
    bool prescribe_internal_velocities;

    // Because we do not initially know what dimension we're in, we need
    // function parser objects for both 2d and 3d.
    Functions::ParsedFunction<2> prescribed_velocity_indicator_function_2d (2);
    Functions::ParsedFunction<3> prescribed_velocity_indicator_function_3d (3);
    Functions::ParsedFunction<2> prescribed_velocity_function_2d (2);
    Functions::ParsedFunction<3> prescribed_velocity_function_3d (3);

    /**
     * Declare additional parameters.
     */
    void declare_parameters(const unsigned int dim,
                           ParameterHandler &prm)
    {
        prm.declare_entry ("Prescribe internal velocities", "false",
                           Patterns::Bool (),
                           "Whether or not to use any prescribed internal velocities. "
                           "Locations in which to prescribe velocities are defined "
                           "in section ‘‘Prescribed velocities/Indicator function’’ "
                           "and the velocities are defined in section ‘‘Prescribed "
                           "velocities/Velocity function’’. Indicators are evaluated "
                           "at the center of each cell, and all DOFs associated with "
                           "the specified velocity component at the indicated cells "
                           "are constrained.");
    }

    prm.enter_subsection ("Prescribed velocities");
    {
        prm.enter_subsection ("Indicator function");
        {
            if (dim == 2)
                Functions::ParsedFunction<2>::declare_parameters (prm, 2);
            else
                Functions::ParsedFunction<3>::declare_parameters (prm, 3);
        }
    }
}
```

```

        }

    prm.leave_subsection();

    prm.enter_subsection ("Velocity function");
    {
        if (dim == 2)
            Functions::ParsedFunction<2>::declare_parameters (prm, 2);
        else
            Functions::ParsedFunction<3>::declare_parameters (prm, 3);
    }
    prm.leave_subsection();
}
prm.leave_subsection();
}

template <int dim>
void parse_parameters(const Parameters<dim> &,
                      ParameterHandler &prm)
{
    prescribe_internal_velocities = prm.get_bool ("Prescribe internal velocities");
    prm.enter_subsection ("Prescribed velocities");
    {
        prm.enter_subsection("Indicator function");
        {
            try
            {
                if (dim == 2)
                    prescribed_velocity_indicator_function_2d.parse_parameters (prm);
                else
                    prescribed_velocity_indicator_function_3d.parse_parameters (prm);
            }
            catch (...)
            {
                std::cerr << "ERROR: FunctionParser failed to parse\n"
                << "\t'Prescribed velocities.Indicator function'\n"
                << "with expression\n"
                << "\t'" << prm.get("Function expression") << "'";
                throw;
            }
        }
        prm.leave_subsection();
    }
    prm.enter_subsection("Velocity function");
    {
        try
        {
            if (dim == 2)
                prescribed_velocity_function_2d.parse_parameters (prm);
            else
                prescribed_velocity_function_3d.parse_parameters (prm);
        }
        catch (...)
        {
            std::cerr << "ERROR: FunctionParser failed to parse\n"
            << "\t'Prescribed velocities.Velocity function'\n"
            << "with expression\n"
        }
    }
}

```

```

        << "\t" << prm.get("Function expression") << "";
    throw;
}
}
prm.leave_subsection();
}
prm.leave_subsection ();
}

/***
* This function retrieves the unit support points (in the unit cell) for the current element.
* The DGP element used when 'set Use locally conservative discretization = true' does not
* have support points. If these elements are in use, a fictitious support point at the cell
* center is returned for each shape function that corresponds to the pressure variable,
* whereas the support points for the velocity are correct; the fictitious points don't matter
* because we only use this function when interpolating the velocity variable, and ignore the
* evaluation at the pressure support points.
*/
template <int dim>
std::vector< Point<dim> >
get_unit_support_points_for_velocity(const SimulatorAccess<dim> &simulator_access)
{
    std::vector< Point<dim> > unit_support_points;
    if ( !simulator_access.get_parameters().use_locally_conservative_discretization )
    {
        return simulator_access.get_fe().get_unit_support_points();
    }
    else
    {
        //special case for discontinuous pressure elements, which lack unit support points
        std::vector< Point<dim> > unit_support_points;
        const unsigned int dofs_per_cell = simulator_access.get_fe().dofs_per_cell;
        for (unsigned int dof=0; dof < dofs_per_cell; ++dof)
        {
            // base will hold element, base_index holds node/shape function within that element
            const unsigned int
                base     = simulator_access.get_fe().system_to_base_index(dof).first.first,
                base_index = simulator_access.get_fe().system_to_base_index(dof).second;
            // get the unit support points for the relevant element
            std::vector< Point<dim> > my_support_points = simulator_access.get_fe().base_element(
                → base).get_unit_support_points();
            if ( my_support_points.size() == 0 )
            {
                //manufacture a support point, arbitrarily at cell center
                if ( dim==2 )
                    unit_support_points.push_back( Point<dim> (0.5,0.5) );
                if ( dim==3 )
                    unit_support_points.push_back( Point<dim> (0.5,0.5,0.5) );
            }
            else
            {
                unit_support_points.push_back(my_support_points[base_index]);
            }
        }
        return unit_support_points;
    }
}

```

```

}

namespace
{
    template <int dim>
    const Point<2> &as_2d(const Point<dim> &p);

    template <>
    const Point<2> &as_2d(const Point<2> &p)
    {
        return p;
    }

    template <int dim>
    const Point<3> &as_3d(const Point<dim> &p);

    template <>
    const Point<3> &as_3d(const Point<3> &p)
    {
        return p;
    }
}

/***
 * This function is called by a signal which is triggered after the other constraints
 * have been calculated. This enables us to define additional constraints in the mass
 * matrix on any arbitrary degree of freedom in the model space.
 */
template <int dim>
void constrain_internal_velocities (const SimulatorAccess<dim> &simulator_access,
                                    ConstraintMatrix &current_constraints)
{
    if (prescribe_internal_velocities)
    {
        const std::vector< Point<dim> > points = get_unit_support_points_for_velocity(
            simulator_access);
        const Quadrature<dim> quadrature (points);
        FEValues<dim> fe_values (simulator_access.get_fe(), quadrature, update_q_points);
        typename DoFHandler<dim>::active_cell_iterator cell;

        // Loop over all cells
        for (cell = simulator_access.get_dof_handler().begin_active();
             cell != simulator_access.get_dof_handler().end();
             ++cell)
        if (! cell->is_artificial())
        {
            fe_values.reinit (cell);
            std::vector<unsigned int> local_dof_indices(simulator_access.get_fe().dofs_per_cell);
            cell->get_dof_indices (local_dof_indices);

            for (unsigned int q=0; q<quadrature.size(); q++)
                // If it's okay to constrain this DOF
                if (current_constraints.can_store_line(local_dof_indices[q]) &&
                    !current_constraints.is_constrained(local_dof_indices[q]))

```

```

{
    // Get the velocity component index
    const unsigned int c_idx =
        simulator_access.get_fe().system_to_component_index(q).first;

    // If we're on one of the velocity DOFs
    if ((c_idx >=
        simulator_access.introspection().component_indices.velocities[0])
    &&
    (c_idx <=
        simulator_access.introspection().component_indices.velocities[dim-1]))
    {
        // Which velocity component is this DOF associated with?
        const unsigned int component_direction
            = (c_idx
                - simulator_access.introspection().component_indices.velocities[0]);

        // we get time passed as seconds (always) but may want
        // to reinterpret it in years
        double time = simulator_access.get_time();
        if (simulator_access.convert_output_to_years())
            time /= year_in_seconds;

        prescribed_velocity_indicator_function_2d.set_time (time);
        prescribed_velocity_indicator_function_3d.set_time (time);
        prescribed_velocity_function_2d.set_time (time);
        prescribed_velocity_function_3d.set_time (time);

        const Point<dim> p = fe_values.quadrature_point(q);

        // Because we defined and parsed our parameter
        // file differently for 2d and 3d we need to
        // be sure to query the correct object for
        // function values. The function parser
        // objects expect points of a certain
        // dimension, but Point p will be compiled for
        // both 2d and 3d, so we need to do some trickery
        // to make this compile.
        double indicator, u_i;
        if (dim == 2)
        {
            indicator = prescribed_velocity_indicator_function_2d.value
                (as_2d(p),
                 component_direction);
            u_i      = prescribed_velocity_function_2d.value
                (as_2d(p),
                 component_direction);
        }
        else
        {
            indicator = prescribed_velocity_indicator_function_3d.value
                (as_3d(p),
                 component_direction);
            u_i      = prescribed_velocity_function_3d.value
                (as_3d(p),
                 component_direction);
        }
    }
}

```

```

        }

        if (indicator > 0.5)
        {
            // Add a constraint of the form dof[q] = u_i
            // to the list of constraints.
            current_constraints.add_line (local_dof_indices[q]);
            current_constraints.set_inhomogeneity (local_dof_indices[q], u_i);
        }
    }
}

// Connect declare_parameters and parse_parameters to appropriate signals.
void parameter_connector ()
{
    SimulatorSignals<2>::declare_additional_parameters.connect (&declare_parameters);
    SimulatorSignals<3>::declare_additional_parameters.connect (&declare_parameters);

    SimulatorSignals<2>::parse_additional_parameters.connect (&parse_parameters<2>);
    SimulatorSignals<3>::parse_additional_parameters.connect (&parse_parameters<3>);
}

// Connect constraints function to correct signal.
template <int dim>
void signal_connector (SimulatorSignals<dim> &signals)
{
    signals.post_constraints_creation.connect (&constrain_internal_velocities<dim>);
}

// Tell Aspect to send signals to the connector functions
ASPECT_REGISTER_SIGNALS_PARAMETER_CONNECTOR(parameter_connector)
ASPECT_REGISTER_SIGNALS_CONNECTOR(signal_connector<2>, signal_connector<3>)
}

```

The above plugin can be compiled with `cmake . && make` in the [cookbooks/prescribed_velocity](#) directory. It can be loaded in a parameter file as an “Additional shared library.” By setting parameters like those shown below, it is possible to produce many interesting flow fields such as the ones visualized in (Figure 29).

```

# Load the signal library.
set Additional shared libraries = ./libprescribed_velocity.so

## Turn prescribed velocities on
set Prescribe internal velocities = true

subsection Prescribed velocities
    subsection Indicator function
        set Variable names = x,y,t
        # Return where to prescribe u_x; u_y; u_z
        # (last one only used if dimension = 3)
        # 1 if velocity should be prescribed, 0 otherwise
        set Function expression = if((x-.5)^2+(y-.5)^2<.125,1,0); \
                                if((x-.5)^2+(y-.5)^2<.125,1,0)
    end

```

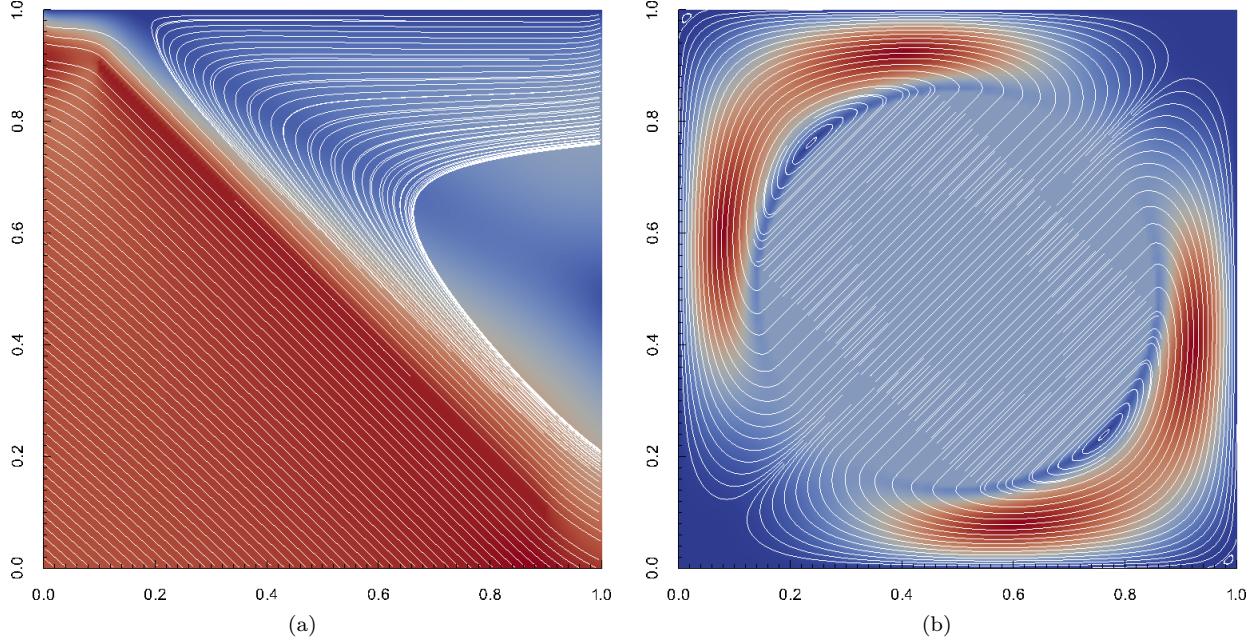


Figure 29: Examples of flows with prescribed internal velocities, as described in Section 5.2.9.

```

subsection Velocity function
  set Variable names = x,y,t
  # Return u_x; u_y; u_z (u_z only used if in 3d)
  set Function expression = 1;-1
end
end

```

5.2.10 Artificial viscosity smoothing

This section was contributed by Ryan Grove

Standard finite element discretizations of advection-diffusion equations introduce unphysical oscillations around steep gradients. Therefore, stabilization must be added to the discrete formulation to obtain correct solutions. In ASPECT, we use the Entropy Viscosity scheme developed by Guermond et al. in the paper [Jea11]. In this scheme, an artificial viscosity is calculated on every cell and used to try to combat these oscillations that cause unwanted overshoot and undershoot. More information about how ASPECT does this is located at https://dealii.org/developer/doxygen/deal.II/step_31.html.

Instead of just looking at an individual cell's artificial viscosity, improvements in the minimizing of the oscillations can be made by smoothing. Smoothing is the act of finding the maximum artificial viscosity taken over a cell T and the neighboring cells across the faces of T , i.e.,

$$\bar{v}_h(T) = \max_{K \in N(T)} v_h(K)$$

where $N(T)$ is the set containing T and the neighbors across the faces of T .

This feature can be turned on by setting the [Use artificial viscosity smoothing](#) flag inside the [Stabilization](#) subsection inside the [Discretization](#) subsection in your parameter file.

To show how this can be used in practice, let us consider the simple convection in a quarter of a 2d annulus cookbook in Section 5.3.1, a radial compositional field was added to help show the advantages of using the artificial viscosity smoothing feature.

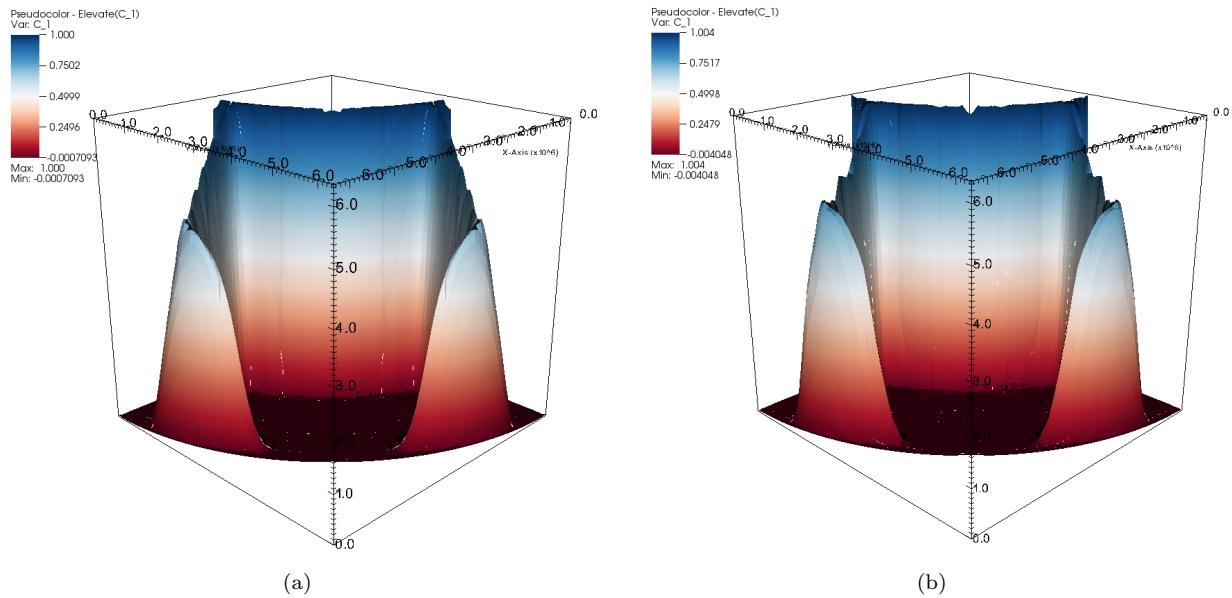


Figure 30: Artificial viscosity smoothing: Example of the output of two similar runs. The run on the left has the artificial viscosity smoothing turned on and the run on the right does not, as described in Section 5.2.10.

By applying the following changes shown below to the parameters of the already existing file cookbooks/shell_simple_2d.prm,

```

subsection Discretization
  set Temperature polynomial degree = 2

  subsection Stabilization parameters
    set Use artificial viscosity smoothing = true
  end
end

subsection Compositional fields
  set Number of fields = 1
end

subsection Initial composition model
  set Model name = function

  subsection Function
    set Variable names      = x,y
    set Function expression = if(sqrt(x*x+y*y)<4000000,1,0)
  end
end

```

it is possible to produce pictures of the simple convection in a quarter of a 2d annulus such as the ones visualized in Figure 30.

5.2.11 Tracking finite strain

This section was contributed by Juliane Dannberg and Rene Gassmöller

Note: In this section, following [BKEO03, DT98], we denote the velocity gradient tensor as \mathbf{G} , where $\mathbf{G} = \nabla \mathbf{u}^T$, and \mathbf{u} is the velocity. Note that this is different from the definition of the strain rate $\epsilon(\mathbf{u})$, which only contains the symmetric part of \mathbf{G} . We then denote the deformation gradient (or deformation) tensor by \mathbf{F} , where \mathbf{F} is the tensor that deforms an initial state \mathbf{x} into an deformed state $\mathbf{r} = \mathbf{F}\mathbf{x}$.

In many geophysical settings, material properties, and in particular the rheology, do not only depend on the current temperature, pressure and strain rate, but also on the history of the system. This can be incorporated in ASPECT models by tracking history variables through compositional fields. In this cookbook, we will show how to do this by tracking the strain that idealized little grains of finite size accumulate over time at every (Lagrangian) point in the model.

Here, we use a material model plugin that defines the compositional fields as the components of the deformation gradient tensor \mathbf{F}_{ij} , and modifies the right-hand side of the corresponding advection equations to accumulate strain over time. This is done by adjusting the `out.reaction_terms` variable:

```
for (unsigned int q=0; q < in.position.size(); ++q)
{
    // Convert the compositional fields into the tensor quantity they represent.
    Tensor<2,dim> strain;
    for (unsigned int i = 0; i < Tensor<2,dim>::n_independent_components ; ++i)
        strain[Tensor<2,dim>::unrolled_to_component_indices(i)] = in.composition[q][i];

    // Compute the strain accumulated in this timestep.
    const Tensor<2,dim> strain_increment = this->get_timestep() * (velocity_gradients[q] * strain);

    // Output the strain increment component-wise to its respective compositional field's reaction
    // → terms.
    for (unsigned int i = 0; i < Tensor<2,dim>::n_independent_components ; ++i)
        out.reaction_terms[q][i] = strain_increment[Tensor<2,dim>::unrolled_to_component_indices(i)];
}
```

Let us denote the accumulated deformation at time step n as \mathbf{F}^n . We can calculate its time derivative as the product of two tensors, namely the current velocity gradient $\mathbf{G}_{ij} = \frac{\partial u_i}{\partial x_j}$ and the deformation gradient \mathbf{F}^{n-1} accumulated up to the previous time step, in other words $\frac{\partial \mathbf{F}}{\partial t} = \mathbf{G}\mathbf{F}$, and $\mathbf{F}^0 = \mathbf{I}$, with \mathbf{I} being the identity tensor. While we refer to other studies [MJ83, DT98, BKEO03] for a derivation of this relationship, we can give an intuitive example for the necessity to apply the velocity gradient to the already accumulated deformation, instead of simply integrating the velocity gradient over time. Consider a simple one-dimensional “grain” of length 1.0, in which case the deformation tensor only has one component, the compression in x -direction. If one embeds this grain into a convergent flow field for a compressible medium where the dimensionless velocity gradient is -0.5 (e.g. a velocity of zero at its left end at $x = 0.0$, and a velocity of -0.5 at its right end at $x = 1.0$), simply integrating the velocity gradient would suggest that the grain reaches a length of zero after two units of time, and would then “flip” its orientation, which is clearly non-physical. What happens instead can be seen by solving the equation of motion for the right end of the grain $\frac{dx}{dt} = v = -0.5x$. Solving this equation for x leads to $x(t) = e^{-0.5t}$. This is therefore also the solution for \mathbf{F} since $\mathbf{F}\mathbf{x}$ transforms the initial position of $x(t=0) = 1.0$ into the deformed position of $x(t=1) = e^{-0.5}$, which is the definition of \mathbf{F} .

In more general cases a visualization of \mathbf{F} is not intuitive, because it contains rotational components that represent a rigid body rotation without deformation. Following [BKEO03] we can polar-decompose the tensor into a positive-definite and symmetric left stretching tensor \mathbf{L} , and an orthogonal rotation tensor \mathbf{Q} ,

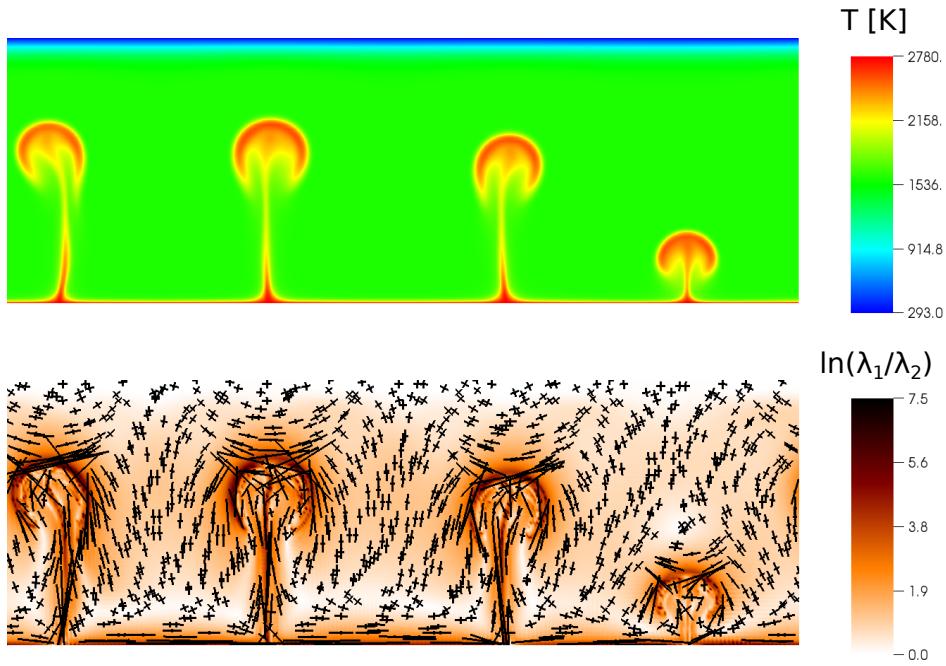


Figure 31: Accumulated finite strain in an example convection model, as described in Section 5.2.11 at a time of 67.6 Ma. Top panel: Temperature distribution. Bottom panel: Natural strain distribution. Additional black crosses are the scaled eigenvectors of the stretching tensor \mathbf{L} , showing the direction of stretching and compression.

as $\mathbf{F} = \mathbf{L}\mathbf{Q}$, therefore $\mathbf{L}^2 = \mathbf{L}\mathbf{L}^T = \mathbf{F}\mathbf{F}^T$. The left stretching tensor \mathbf{L} (or finite strain tensor) then describes the deformation we are interested in, and its eigenvalues λ_i and eigenvectors \mathbf{e}_i describe the length and orientation of the half-axes of the finite strain ellipsoid. Moreover, we will represent the amount of relative stretching at every point by the ratio $\ln(\lambda_1/\lambda_2)$, called the *natural strain* [Rib92].

The full plugin implementing the integration of \mathbf{F} can be found in [cookbooks/finite_strain/finite_strain.cc](#) and can be compiled with `cmake . && make` in the [cookbooks/finite_strain](#) directory. It can be loaded in a parameter file as an “Additional shared library”, and selected as material model. As it is derived from the “simple” material model, all input parameters for the material properties are read in from the subsection `Simple model`.

```

set Additional shared libraries          = ./libfinite_strain.so

subsection Material model
  set Model name = finite strain

  subsection Simple model
    set Thermal conductivity = 4.7
    set Reference density = 3400
    set Thermal expansion coefficient = 2e-5
    set Viscosity = 5e21
    set Thermal viscosity exponent = 7
    set Reference temperature = 1600
  end
end

```

The plugin was tested against analytical solutions for the deformation gradient tensor in simple and

pure shear as described in `benchmarks/finite_strain/pure_shear.prm` and `benchmarks/finite_strain/simple_shear.prm`.

We will demonstrate its use at the example of a 2D Cartesian convection model (Figure 31): Heating from the bottom leads to the ascent of plumes from the boundary layer (top panel), and the amount of stretching is visible in the distribution of natural strain (color in lower panel). Additionally, the black crosses show the direction of stretching and compression (the eigenvectors of \mathbf{L}). Material moves to the sides at the top of the plume head, so that it is shortened in vertical direction (short vertical lines) and stretched in horizontal direction (long horizontal lines). The sides of the plume head show the opposite effect. Shear occurs mostly at the edges of the plume head, in the plume tail, and in the bottom boundary layer (black areas in the natural strain distribution).

The example used here shows how history variables can be integrated up over the model evolution. While we do not use these variables actively in the computation (in our example, there is no influence of the accumulated strain on the rheology or any other material property), it would be trivial to extend this material model in a way that material properties depend on the integrated strain: Because the values of the compositional fields are part of what the material model gets as inputs, they can easily be used for computing material model outputs such as the viscosity.

Note: In this model we present the use of multiple compositional fields for other purposes than chemical composition. It would have been feasible to run the same model with particles that track the deformation gradient, as additionally implemented and tested in the simple shear and pure shear benchmarks mentioned in this section. Both approaches have specific advantages, and for scientific computations one needs to evaluate the more suitable strategy. Compositional fields cover the whole domain, but are affected by numerical diffusion, effectively reducing the maximum accumulated strain. Particles only provide finite strain values at discrete positions, but can, if this is desired, be used in fewer numbers and only a part of the model domain (and are much faster in this case). If however there needs to be a large number of particles (possibly because they are used for other purposes as well), then they can be much more expensive. Both approaches can be used to actively influence the rheology in the material model.

5.2.12 Reading in compositional initial composition files generated with geomIO

This section was contributed by Juliane Dannberg

Many geophysical setups require initial conditions with several different materials and complex geometries. Hence, sometimes it would be easier to generate the initial geometries of the materials as a drawing instead of by writing code. The MATLAB-based library geomIO (<http://geomio.bitbucket.org/>) provides a convenient tool to convert a drawing generated with the vector graphics editor Inkscape (<https://inkscape.org/en/>) to a data file that can be read into ASPECT. Here, we will demonstrate how this can be done for a 2D setup for a model with one compositional field, but geomIO also has the capability to create 3D volumes based on a series of 2D vector drawings using any number of different materials. Similarly, initial conditions defined in this way can also be used with particles instead of compositional fields.

To obtain the developer version of geomIO, you can clone the bitbucket repository by executing the command

```
git clone https://bitbucket.org/geomio/geomio.git
```

or you can download geomIO [here](#). You will then need to add the geomIO source folders to your MATLAB path by running the file located in `/path/to/geomio/installation/InstallGeomIO.m`. An extensive documentation for how to use geomIO can be found [here](#). Among other things, it explains [how to generate drawings in Inkscape](#) that can be read in by geomIO, which involves assigning new attributes to paths in Inkscape's XML editor. In particular, a new property 'phase' has to be added to each path, and set to a value corresponding to the index of the material that should be present in this region in the initial condition of the geodynamic model.

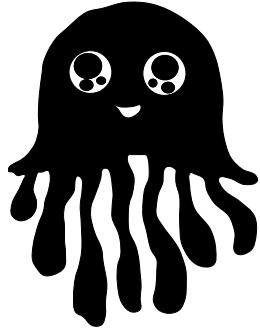


Figure 32: Vector drawing of a jellyfish.

Note: geomIO currently only supports the latest stable version of Inkscape (0.91), and other versions might not work with geomIO or cause errors. Moreover, geomIO currently does not support grouping paths (paths can still be combined using Path→Union, Path→Difference or similar commands), and only the outermost closed contour of a path will be considered. This means that, for example, for modeling a spherical annulus, you would have to draw two circles, and assign the inner one the same phase as the background of your drawing.

We will here use a drawing of a jellyfish located in [cookbooks/geomio/jellyfish.svg](#), where different phases have already been assigned to each path (Figure 32).

Note: The page of your drawing in Inkscape should already have the extents (in px) that you later want to use in your model (in m).

After geomIO is initialized in MATLAB, we [run geomIO as described in the documentation](#), loading the default options and then specifying all the option we want to change, such as the path to the input file, or the resolution:

```
% set options for geomIO
opt          = geomIO_Options();
opt.inputFileName = ['/path/to/aspect/doc/manual/cookbooks/geomio/jellyfish.svg'];
opt.DrawCoordRes = 21; % optionally change resolution with opt.DrawCoordRes = your value;

% run geomIO
[PathCoord]      = run_geomIO(opt,'2D');
```

You can view all of the options available by typing `opt` in MATLAB.

In the next step we create the grid that is used for the coordinates in the `ascii` data initial conditions file and assign a phase to each grid point:

```
% define the bounding box for the output mesh
% (this should be the X extent and Y extent in your Aspect model)
xmin = 0; xmax = opt.svg.width;
ymin = 0; ymax = opt.svg.height;

% set the resolution in the output file:
% [Xp,Yp] = ndgrid(xmin:your_stepLength_x:xmax,ymin:your_stepLength_y:ymax);
[Xp,Yp]      = ndgrid(xmin:15:xmax,ymin:15:ymax);
Phase        = zeros(size(Xp));
```

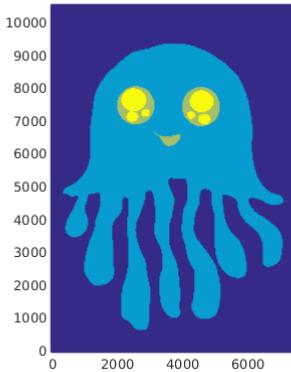


Figure 33: Plot of the Phase variable in MATLAB.

```
% assign a phase to each grid point according to your drawing
Phase = assignPhase2Markers(PathCoord, opt, Xp, Yp, Phase);

% plot your output
figure(2)
scatter(Xp(:, ),Yp(:, ),10,Phase(:, ),'filled');
axis equal
axis([xmin xmax ymin ymax])
```

You can plot the Phase variable in MATLAB to see if the drawing was read in and all phases are assigned correctly (Figure 33). Finally, we want to write output in a format that can be read in by ASPECT's ascii data compositional initial conditions plugin. We write the data into the file `jelly.txt`:

```
% the headers ASPECT needs for the ascii data plugin
header1 = 'x';
header2 = 'y';
header3 = 'phase';

% create an array in the correct format for the ascii data plugin
Vx = Xp(:, );
Vy = Yp(:, );
VPhase = Phase(:, );
[m,n] = size(Phase);

% write the data into the output file
fid=fopen('jelly.txt','w');
fprintf(fid, '# POINTS: %d %d \n',[m n]);
fprintf(fid, ['# Columns: ' header1 ' ' header2 ' ' header3 '\n']);
fprintf(fid, '%f %f %f \n', [Vx Vy VPhase']);
fclose(fid);
```

To read in the file we just created (a copy is located in ASPECT's data directory), we set up a model with a box geometry with the same extents we specified for the drawing in px and one compositional field. We choose the ascii data compositional initial conditions and specify that we want to read in our jellyfish. The relevant parts of the input file are listed below:

```
subsection Geometry model
  set Model name = box
```

```

# The extents of the box is the same as the width
# and height of the drawing in px
# (an A4 page = 7350x10500 px).
subsection Box
    set X extent = 7350
    set Y extent = 10500
end
end

# We need one compositional field that will be assigned
# the values read in from the ascii data plugin.
subsection Compositional fields
    set Number of fields = 1
end

# We use the ascii data plugin to read in the file created with geomIO.
subsection Initial composition model
    set Model name = ascii data

    subsection Ascii data model
        set Data directory      = $ASPECT_SOURCE_DIR/data/initial-composition/ascii-data/test/
        set Data file name     = jelly.txt
    end
end

# We refine the mesh where compositional gradients are
# high, i.e. at the boundary between the different phases
# assigned to the compositional field through the initial
# condition.
subsection Mesh refinement
    set Refinement fraction      = 0.99
    set Coarsening fraction      = 0
    set Initial global refinement = 5
    set Initial adaptive refinement = 4
    set Time steps between mesh refinement = 0
    set Strategy                 = composition
end

```

If we look at the output in `paraview`, we can see our jellyfish, with the mesh refined at the boundaries between the different phases (Figure 34).

For a geophysical setup, the MATLAB code could be extended to write out the phases into several different columns of the ASCII data file (corresponding to different compositional fields). This initial conditions file could then be used in ASPECT with a material model such as the `multicomponent` model, assigning each phase different material properties.

An animation of a model using the jellyfish as initial condition and assigning it a higher viscosity can be found here: <https://www.youtube.com/watch?v=YzNTubNG83Q>.

5.3 Geophysical setups

Having gone through the ways in which one can set up problems in rectangular geometries, let us now move on to situations that are directed more towards the kinds of things we want to use ASPECT for: the simulation of convection in the rocky mantles of planets or other celestial bodies.

To this end, we need to go through the list of issues that have to be described and that were outlined in Section 5.1, and address them one by one:

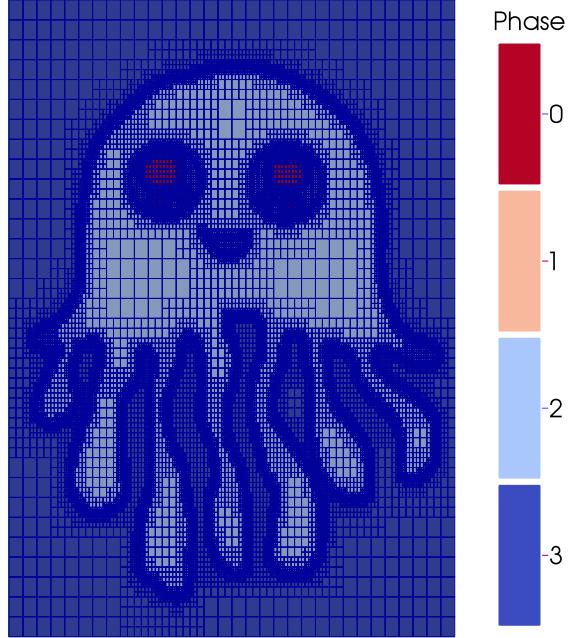


Figure 34: ASPECT model output of the jellyfish and corresponding mesh in paraview.

- *What internal forces act on the medium (the equation)?* This may in fact be the most difficult to answer part of it all. The real material in Earth’s mantle is certainly no Newtonian fluid where the stress is a linear function of the strain with a proportionality constant (the viscosity) η that only depends on the temperature. Rather, the real viscosity almost surely also depends on the pressure and the strain rate. Because the issue is complicated and the exact material model not entirely clear, for the next few subsections we will therefore ignore the issue and start with just using the “simple” material model where the viscosity is constant and most other coefficients depend at most on the temperature.
- *What external forces do we have (the right hand side)* There are of course other issues: for example, should the model include terms that describe shear heating? Should it be compressible? Adiabatic heating due to compression? Most of the terms that pertain to these questions appear on the right hand sides of the equations, though some (such as the compressibility) also affect the differential operators on the left. Either way, for the moment, let us just go with the simplest models and come back to the more advanced questions in later examples.

One right hand side that will certainly be there is that due to gravitational acceleration. To first order, within the mantle gravity points radially inward and has a roughly constant magnitude. In reality, of course, the strength and direction of gravity depends on the distribution and density of materials in Earth – and, consequently, on the solution of the model at every time step. We will discuss some of the associated issues in the examples below.

- *What is the domain (geometry)?* This question is easier to answer. To first order, the domains we want to simulate are spherical shells, and to second order ellipsoid shells that can be obtained by considering the isopotential surface of the gravity field of a homogeneous, rotating fluid. A more accurate description is of course the geoid for which several parameterizations are available. A complication arises if we ask whether we want to include the mostly rigid crust in the domain and simply assume that it is part of the convecting mantle, albeit a rather viscous part due to its low temperature and the low pressure there, or whether we want to truncate the computation at the asthenosphere.

- *What happens at the boundary for each variable involved (boundary conditions)?* The mantle has two boundaries: at the bottom where it contacts the outer core and at the top where it either touches the air or, depending on the outcome of the discussion of the previous question, where it contacts the lithospheric crust. At the bottom, a very good approximation of what is happening is certainly to assume that the velocity field is tangential (i.e., horizontal) and without friction forces due to the very low viscosity of the liquid metal in the outer core. Similarly, we can assume that the outer core is well mixed and at a constant temperature. At the top boundary, the situation is slightly more complex because in reality the boundary is not fixed but also allows vertical movement. If we ignore this, we can assume free tangential flow at the surface or, if we want, prescribe the tangential velocity as inferred from plate motion models. ASPECT has a plugin that allows to query this kind of information from the `GPlates` program.
- *How did it look at the beginning (initial conditions)?* This is of course a trick question. Convection in the mantle of earth-like planets did not start with a concrete initial temperature distribution when the mantle was already fully formed. Rather, convection already happened when primordial material was still separating into mantle and core. As a consequence, for models that only simulate convection using mantle-like geometries and materials, no physically reasonable initial conditions are possible that date back to the beginning of Earth. On the other hand, recall that we only need initial conditions for the temperature (and, if necessary, compositional fields). Thus, if we have a temperature profile at a given time, for example one inferred from seismic data at the current time, then we can use these as the starting point of a simulation.

This discussion shows that there are in fact many pieces with which one can play and for which the answers are in fact not always clear. We will address some of them in the cookbooks below. Recall in the descriptions we use in the input files that ASPECT uses physical units, rather than non-dimensionalizing everything. The advantage, of course, is that we can immediately compare outputs with actual measurements. The disadvantage is that we need to work a bit when asked for, say, the Rayleigh number of a simulation.

5.3.1 Simple convection in a quarter of a 2d annulus

Let us start this sequence of cookbooks using a simpler situation: convection in a quarter of a 2d shell. We choose this setup because 2d domains allow for much faster computations (in turn allowing for more experimentation) and because using a quarter of a shell avoids a pitfall with boundary conditions we will discuss in the next section. Because it's simpler to explain what we want to describe in pictures than in words, Fig. 35 shows the domain and the temperature field at a few time steps. In addition, you can find a movie of how the temperature evolves over this time period at <http://www.youtube.com/watch?v=d4AS1FmdarU>.³⁰

Let us just start by showing the input file (which you can find in `cookbooks/shell_simple_2d.prm`):

```
set Dimension = 2
set Use years in output instead of seconds = true
set End time = 1.5e9
set Output directory = output-shell_simple_2d

subsection Material model
  set Model name = simple

  subsection Simple model
    set Thermal expansion coefficient = 4e-5
    set Viscosity = 1e22
  end
end
```

³⁰In YouTube, click on the gear symbol at the bottom right of the player window to select the highest resolution to see all the details of this video.

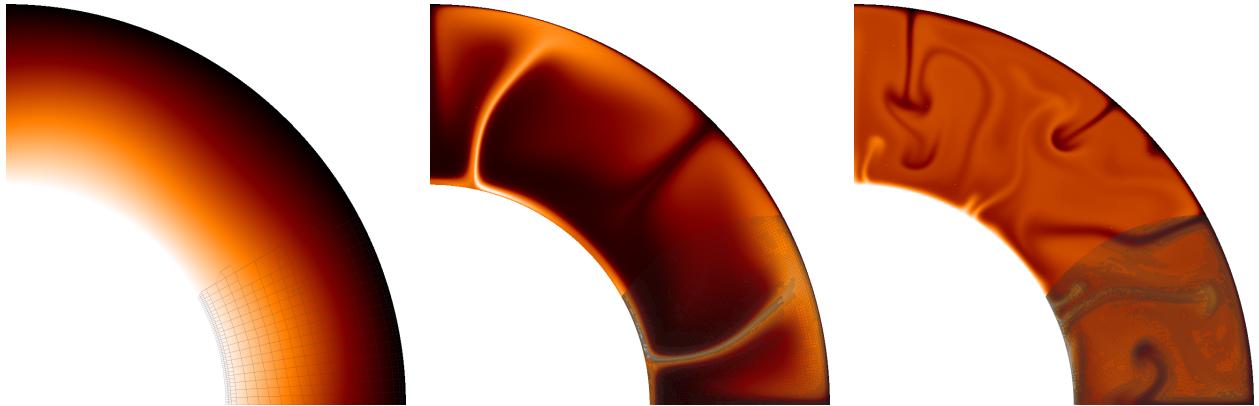


Figure 35: Simple convection in a quarter of an annulus: Snapshots of the temperature field at times $t = 0$, $t = 1.2 \cdot 10^7$ years (time step 2135), and $t = 10^9$ years (time step 25,662). The bottom right part of each figure shows an overlay of the mesh used during that time step.

```

subsection Geometry model
  set Model name = spherical shell

  subsection Spherical shell
    set Inner radius   = 3481000
    set Outer radius   = 6336000
    set Opening angle = 90
  end
end

subsection Boundary velocity model
  set Zero velocity boundary indicators      = inner
  set Tangential velocity boundary indicators = outer, left, right
end

subsection Boundary temperature model
  set Fixed temperature boundary indicators = inner, outer
  set List of model names = spherical constant

  subsection Spherical constant
    set Inner temperature = 4273
    set Outer temperature = 973
  end
end

subsection Heating model
  set List of model names = shear heating
end

subsection Initial temperature model

```

```

set Model name = spherical hexagonal perturbation
end

subsection Gravity model
  set Model name = ascii data
end

subsection Mesh refinement
  set Initial global refinement      = 5
  set Initial adaptive refinement   = 4
  set Strategy                      = temperature
  set Time steps between mesh refinement = 15
end

subsection Postprocess
  set List of postprocessors = visualization, velocity statistics, temperature statistics, ...
                             ... heat flux statistics, depth average

  subsection Visualization
    set Output format           = vtu
    set Time between graphical output = 1e6
    set Number of grouped files = 0
  end

  subsection Depth average
    set Time between graphical output = 1e6
  end
end

```

In the following, let us pick apart this input file:

1. Lines 1–4 are just global parameters. Since we are interested in geophysically realistic simulations, we will use material parameters that lead to flows so slow that we need to measure time in years, and we will set the end time to 1.5 billion years – enough to see a significant amount of motion.
2. The next block (lines 7–14) describes the material that is convecting (for historical reasons, the remainder of the parameters that describe the equations is in a different section, see the fourth point below). We choose the simplest material model ASPECT has to offer where the viscosity is constant (here, we set it to $\eta = 10^{22}$ Pa s) and so are all other parameters except for the density which we choose to be $\rho(T) = \rho_0(1 - \alpha(T - T_{\text{ref}}))$ with $\rho_0 = 3300\text{kg m}^{-3}$, $\alpha = 4 \cdot 10^{-5}\text{K}^{-1}$ and $T_{\text{ref}} = 293\text{K}$. The remaining material parameters remain at their default values and you can find their values described in the documentation of the `simple` material model in Sections A.81 and A.102.
3. Lines 17–25 then describe the geometry. In this simple case, we will take a quarter of a 2d shell (recall that the dimension had previously been set as a global parameter) with inner and outer radii matching those of a spherical approximation of Earth.
4. The second part of the model description and boundary values follows in lines 28–42. The boundary conditions require us to look up how the geometry model we chose (the `spherical shell` model) assigns boundary indicators to the four sides of the domain. This is described in Section A.42 where the model announces that boundary indicator zero is the inner boundary of the domain, boundary indicator one is the outer boundary, and the left and right boundaries for a 2d model with opening angle of 90 degrees as chosen here get boundary indicators 2 and 3, respectively. In other words, the

settings in the input file correspond to a zero velocity at the inner boundary and tangential flow at all other boundaries. We know that this is not realistic at the bottom, but for now there are of course many other parts of the model that are not realistic either and that we will have to address in subsequent cookbooks. Furthermore, the temperature is fixed at the inner and outer boundaries (with the left and right boundaries then chosen so that no heat flows across them, emulating symmetry boundary conditions) and, further down, set to values of 700 and 4000 degrees Celsius – roughly realistic for the bottom of the crust and the core-mantle boundary.

5. Lines 45–47 describe that we want a model where equation (3) contains the shear heating term $2\eta\varepsilon(\mathbf{u}) : \varepsilon(\mathbf{u})$ (noting that the default is to use an incompressible model for which the term $\frac{1}{3}(\nabla \cdot \mathbf{u})\mathbf{1}$ in the shear heating contribution is zero). Considering a reasonable choice of heating terms is not the focus of this simple cookbook, therefore we will leave a discussion of possible and reasonable heating terms to another cookbook.
6. The description of what we want to model is complete by specifying that the initial temperature is a perturbation with hexagonal symmetry from a linear interpolation between inner and outer temperatures (see Section A.66), and what kind of gravity model we want to choose (one reminiscent of the one inside the Earth mantle, see Section A.52).
7. The remainder of the input file consists of a description of how to choose the initial mesh and how to adapt it (lines 60–65) and what to do at the end of each time step with the solution that ASPECT computes for us (lines 68–81). Here, we ask for a variety of statistical quantities and for graphical output in VTU format every million years.

Note: Having described everything to ASPECT, you may want to view the video linked to above again and compare what you see with what you expect. In fact, this is what one should always do having just run a model: compare it with expectations to make sure that we have not overlooked anything when setting up the model or that the code has produced something that doesn't match what we thought we should get. Any such mismatch between expectation and observed result is typically a learning opportunity: it either points to a bug in our input file, or it provides us with insight about an aspect of reality that we had not foreseen. Either way, accepting results uncritically is, more often than not, a way to scientifically invalid results.

The model we have chosen has a number of inadequacies that make it not very realistic (some of those happened more as an accident while playing with the input file and weren't a purposeful experiment, but we left them in because they make for good examples to discuss below). Let us discuss these issues in the following.

Dimension. This is a cheap shot but it is nevertheless true that the world is three-dimensional whereas the simulation here is 2d. We will address this in the next section.

Incompressibility, adiabaticity and the initial conditions. This one requires a bit more discussion. In the model selected above, we have chosen a model that is incompressible in the sense that the density does not depend on the pressure and only very slightly depends on the temperature. In such models, material that rises up does not cool down due to expansion resulting from the pressure dropping, and material that is transported down does not adiabatically heat up. Consequently, the adiabatic temperature profile would be constant with depth, and a well-mixed model with hot inner and cold outer boundary would have a constant temperature with thin boundary layers at the bottom and top of the mantle. In contrast to this, our initial temperature field was a perturbation of a linear temperature profile.

There are multiple implications of this. First, the temperature difference between outer and inner boundary of 3300 K we have chosen in the input file is much too large. The temperature difference that drives the convection, is the difference *in addition* to the temperature increase a volume of material would experience

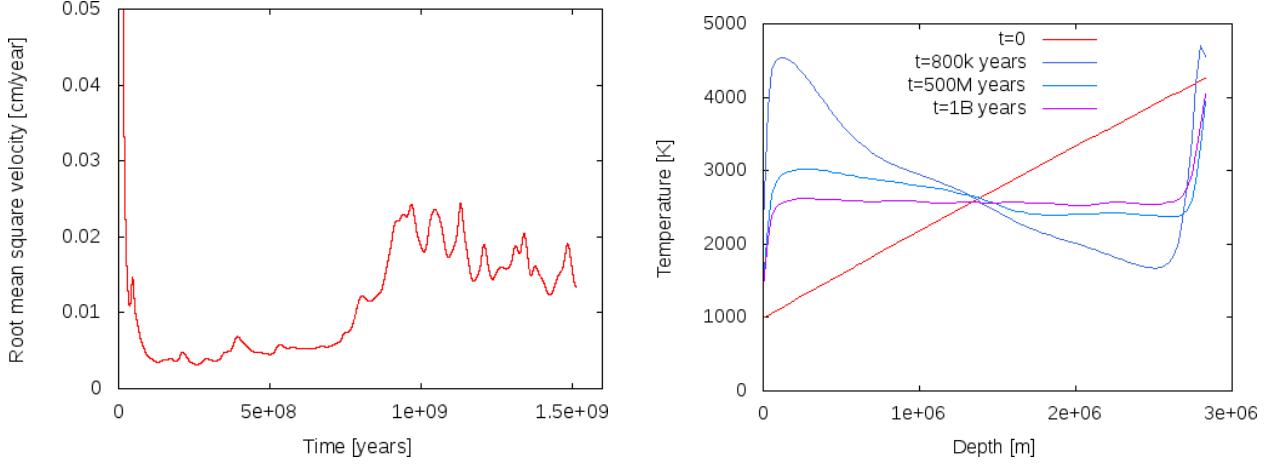


Figure 36: *Simple convection in a quarter of an annulus.* Left: Root mean square values of the velocity field. The initial spike (off the scale) is due to the overturning of the unstable layering of the temperature. Convection is suppressed for the first 800 million years due to the stable layering that results from it. The maximal velocity encountered follows generally the same trend and is in the range of 2–3 cm/year between 100 and 800 million years, and 4–8 cm/year following that. Right: Average temperature at various depths for $t = 0$, $t = 800,000$ years, $t = 5 \cdot 10^8$ years, and $t = 10^9$ years.

if it were to be transported adiabatically from the surface to the core-mantle boundary. This difference is much smaller than 3300 K in reality, and we can expect convection to be significantly less vigorous than in the simulation here. Indeed, using the values in the input file shown above, we can compute the Rayleigh number for the current case to be³¹

$$\text{Ra} = \frac{g \alpha \Delta T \rho L^3}{\kappa \eta} = \frac{10 \text{ m s}^{-2} \times 4 \cdot 10^{-5} \text{ K}^{-1} \times 3300 \text{ K} \times 3300 \text{ kg m}^{-3} \times (2.86 \cdot 10^6 \text{ m})^3}{10^{-6} \text{ m}^2 \text{s}^{-1} \times 10^{22} \text{ kg m}^{-1} \text{s}^{-1}}.$$

Second, the initial temperature profile we chose is not realistic – in fact, it is a completely unstable one: there is hot material underlying cold one, and this is not just the result of boundary layers. Consequently, what happens in the simulation is that we first overturn the entire temperature field with the hot material in the lower half of the domain swapping places with the colder material in the top, to achieve a stable layering except for the boundary layers. After this, hot blobs rise from the bottom boundary layer into the cold layer at the bottom of the mantle, and cold blobs sink from the top, but their motion is impeded about half-way through the mantle once they reach material that has roughly the same temperature as the plume material. This impedes convection until we reach a state where these plumes have sufficiently mixed the mantle to achieve a roughly constant temperature profile.

This effect is visible in the movie linked to above where convection does not penetrate the entire depth of the mantle for the first 20 seconds (corresponding to roughly the first 800 million years). We can also see this effect by plotting the root mean square velocity, see the left panel of Fig. 36. There, we can see how the average velocity picks up once the stable layering of material that resulted from the initial overturning has been mixed sufficiently to allow plumes to rise or sink through the entire depth of the mantle.

The right panel of Fig. 36 shows a different way of visualizing this, using the average temperature at various depths of the model (this is what the `depth average` postprocessor computes). The figure shows how the initially linear unstable layering almost immediately reverts completely, and then slowly equilibrates towards a temperature profile that is constant throughout the mantle (which in the incompressible model chosen here equates to an adiabatic layering) except for the boundary layers at the inner and outer boundaries.

³¹Note that the density in 2d has units kg m^{-2}

(The end points of these temperature profiles do not exactly match the boundary values specified in the input file because we average temperatures over shells of finite width.)

A conclusion of this discussion is that if we want to evaluate the statistical properties of the flow field, e.g., the number of plumes, average velocities or maximal velocities, then we need to restrict our efforts to times after approximately 800 million years in this simulation to avoid the effects of our inappropriately chosen initial conditions. Likewise, we may actually want to choose initial conditions more like what we see in the model for later times, i.e., constant in depth with the exception of thin boundary layers, if we want to stick to incompressible models.

Material model. The model we use here involves viscosity, density, and thermal property functions that do not depend on the pressure, and only the density varies (slightly) with the temperature. We know that this is not the case in nature.

Shear heating. When we set up the input file, we started with a model that includes the shear heating term $2\eta\varepsilon(\mathbf{u}) : \varepsilon(\mathbf{u})$ in eq. (3). In hindsight, this may have been the wrong decision, but it provides an opportunity to investigate whether we think that the results of our computations can possibly be correct.

We first realized the issue when looking at the heat flux that the `heat flux statistics` postprocessor computes. This is shown in the left panel of Fig. 37.³² There are two issues one should notice here. The more obvious one is that the flux from the mantle to the air is consistently higher than the heat flux from core to mantle. Since we have no radiogenic heating model selected (see the `List of model names` parameter in the `Heating model` section of the input file; see also Section A.58), in the long run the heat output of the mantle must equal the input, unless it cools. Our misconception was that after the 800 million year transition, we believed that we had reached a steady state where the average temperature remains constant and convection simply moves heat from the core-mantle boundary to the surface. One could also be tempted to believe this from the right panel in Fig. 36 where it looks like the average temperature does at least not change dramatically. But, it is easy to convince oneself that that is not the case: the `temperature statistics` postprocessor we had previously selected also outputs data about the mean temperature in the model, and it looks like shown in the left panel of Fig. 38. Indeed, the average temperature drops over the course of the 1.2 billion years shown here. We could now convince ourselves that indeed the loss of thermal energy in the mantle due to the drop in average temperature is exactly what fuels the persistently imbalanced energy outflow. In essence, what this would show is that if we kept the temperature at the boundaries constant, we would have chosen a mantle that was initially too hot on average to be sustained by the boundary values and that will cool until it will be in energetic balance and on longer time scales, in-and outflow of thermal energy would balance each other.

However, there is a bigger problem. Fig. 37 shows that at the very beginning, there is a spike in energy flux through the outer boundary. We can explain this away with the imbalanced initial temperature field that leads to an overturning and, thus, a lot of hot material rising close to the surface that will then lead to a high energy flux towards the cold upper boundary. But, worse, there is initially a *negative* heat flux into the mantle from the core – in other words, the mantle is *losing* energy to the core. How is this possible? After all, the hottest part of the mantle in our initial temperature field is at the core-mantle boundary, no thermal energy should be flowing from the colder overlying material towards the hotter material at the boundary! A glimpse of the solution can be found in looking at the average temperature in Fig. 38: At the beginning, the average temperature *rises*, and apparently there are parts of the mantle that become hotter than the 4273 K we have given the core, leading to a downward heat flux. This heating can of course only come from the shear heating term we have accidentally left in the model: at the beginning, the unstable layering leads to very large velocities, and large velocities lead to large velocity gradients that in turn lead to a lot of shear heating! Once the initial overturning has subsided, after say 100 million years (see the mean velocity in Fig. 36), the shear heating becomes largely irrelevant and the cooling of the mantle indeed begins.

³²The `heat flux statistics` postprocessor computes heat fluxes through parts of the boundary in *outward* direction, i.e., from the mantle to the air and to the core. However, we are typically interested in the flux from the core into the mantle, so the figure plots the negative of the computed quantity.

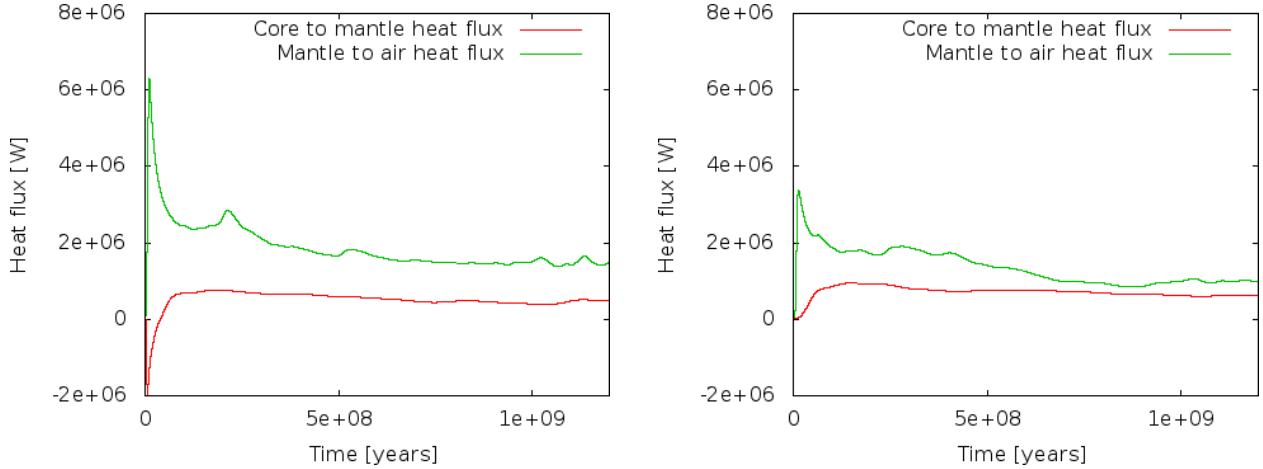


Figure 37: *Simple convection in a quarter of an annulus.* Left: Heat flux through the core-mantle and mantle-air boundaries of the domain for the model with shear heating. Right: Same for a model without shear heating.

Whether this is really the case is of course easily verified: The right panels of Fig.s 37 and 38 show heat fluxes and average temperatures for a model where we have switched off the shear heating by setting

```
subsection Heating model
  set List of model names =
end
```

Indeed, doing so leads to a model where the heat flux from core to mantle is always positive, and where the average temperature strictly drops!

Summary. As mentioned, we will address some of the issues we have identified as unrealistic in the following sections. However, despite all of this, some things are at least at the right order of magnitude, confirming that what ASPECT is computing is reasonable. For example, the maximal velocities encountered in our model (after the 800 million year boundary) are in the range of 6–7cm per year, with occasional excursions up to 11cm. Clearly, something is going in the right direction.

5.3.2 Simple convection in a spherical 3d shell

The setup from the previous section can of course be extended to 3d shell geometries as well – though at significant computational cost. In fact, the number of modifications necessary is relatively small, as we will discuss below. To show an example up front, a picture of the temperature field one gets from such a simulation is shown in Fig. 39. The corresponding movie can be found at <http://youtu.be/j63MkEc0RRw>.

The input file. Compared to the input file discussed in the previous section, the number of changes is relatively small. However, when taking into account the various discussions about which parts of the model were or were not realistic, they go throughout the input file, so we reproduce it here in its entirety, interspersed with comments (the full input file can also be found in `cookbooks/shell_simple_3d.prm`). Let us start from the top where everything looks the same except that we set the dimension to 3:

```
set Dimension          = 3
set Use years in output instead of seconds = true
set End time           = 1.5e9
set Output directory    = output-shell_simple_3d
```

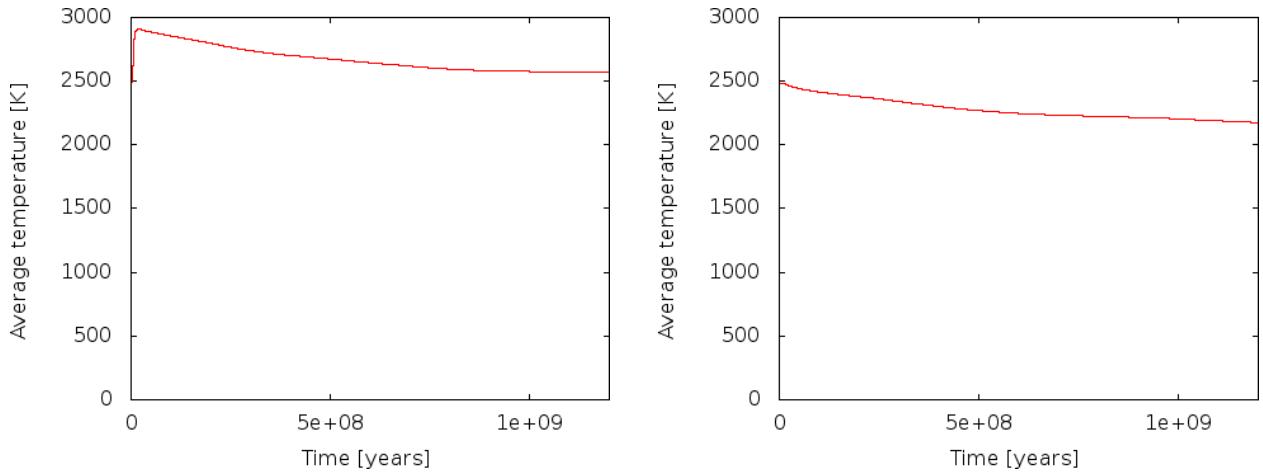


Figure 38: Simple convection in a quarter of an annulus. Left: Average temperature throughout the model for the model with shear heating. Right: Same for a model without shear heating.

```

subsection Material model
  set Model name = simple

subsection Simple model
  set Thermal expansion coefficient = 4e-5
  set Viscosity = 1e22
end
end

```

The next section concerns the geometry. The geometry model remains unchanged at “spherical shell” but we omit the opening angle of 90 degrees as we would like to get a complete spherical shell. Such a shell of course also only has two boundaries (the inner one has indicator zero, the outer one indicator one) and consequently these are the only ones we need to list in the “Boundary velocity model” section:

```

subsection Geometry model
  set Model name = spherical shell

subsection Spherical shell
  set Inner radius = 3481000
  set Outer radius = 6336000
end
end

subsection Boundary velocity model
  set Zero velocity boundary indicators = inner
  set Tangential velocity boundary indicators = outer
end

```

Next, since we convinced ourselves that the temperature range from 973 to 4273 was too large given that we do not take into account adiabatic effects in this model, we reduce the temperature at the inner edge of the mantle to 1973. One can think of this as an approximation to the real temperature there minus the amount of adiabatic heating material would experience as it is transported from the surface to the

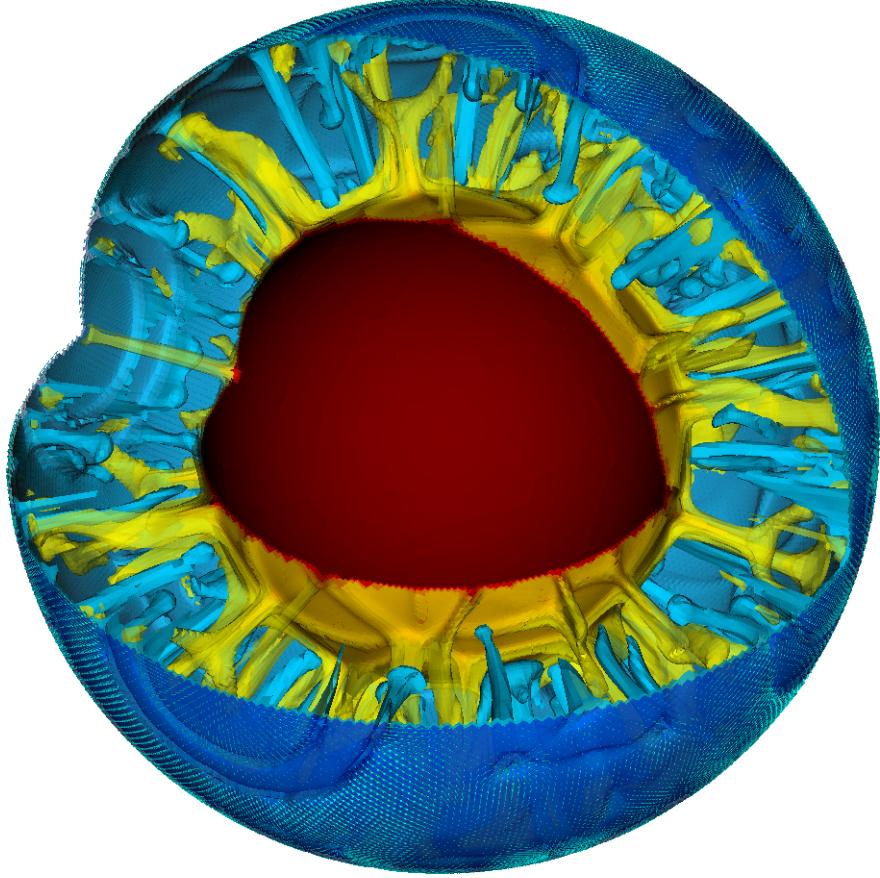


Figure 39: *Convection in a spherical shell: Snapshot of isosurfaces of the temperature field at time $t \approx 1.06 \cdot 10^9$ years with a quarter of the geometry cut away. The surface shows vectors indicating the flow velocity and direction.*

core-mantle boundary. This is, in effect, the temperature difference that drives the convection (because a completely adiabatic temperature profile is stable despite the fact that it is much hotter at the core mantle boundary than at the surface). What the real value for this temperature difference is, is unclear from current research, but it is thought to be around 1000 Kelvin, so let us choose these values.

```

subsection Boundary temperature model
  set Fixed temperature boundary indicators = inner, outer
  set List of model names = spherical constant

  subsection Spherical constant
    set Inner temperature = 1973
    set Outer temperature = 973
  end
end

```

The second component to this is that we found that without adiabatic effects, an initial temperature profile that decreases the temperature from the inner to the outer boundary makes no sense. Rather, we expected a more or less constant temperature with boundary layers at both ends. We could describe such an initial temperature field, but since any initial temperature is mostly arbitrary anyway, we opt to just assume a constant temperature in the middle between the inner and outer temperature boundary values and let the

simulation find the exact shape of the boundary layers itself:

```

subsection Initial temperature model
    set Model name = function
    subsection Function
        set Function expression = 1473
    end
end

subsection Gravity model
    set Model name = ascii data
end

```

As before, we need to determine how many mesh refinement steps we want. In 3d, it is simply not possible to have as much mesh refinement as in 2d, so we choose the following values that lead to meshes that have, after an initial transitory phase, between 1.5 and 2.2 million cells and 50–75 million unknowns:

```

subsection Mesh refinement
    set Initial global refinement      = 2
    set Initial adaptive refinement   = 3
    set Strategy                      = temperature
    set Time steps between mesh refinement = 15
end

```

Second to last, we specify what we want ASPECT to do with the solutions it computes. Here, we compute the same statistics as before, and we again generate graphical output every million years. Computations of this size typically run with 1000 MPI processes, and it is not efficient to let every one of them write their own file to disk every time we generate graphical output; rather, we group all of these into a single file to keep file systems reasonably happy. Likewise, to accommodate the large amount of data, we output depth averaged fields in VTU format since it is easier to visualize:

```

subsection Postprocess
    set List of postprocessors = visualization, velocity statistics, \
                                temperature statistics, heat flux statistics, \
                                depth average

    subsection Visualization
        set Output format          = vtu
        set Time between graphical output = 1e6
        set Number of grouped files = 1
    end

    subsection Depth average
        set Time between graphical output = 1.5e6
        set Output format          = vtu
    end
end

```

Finally, we realize that when we run very large parallel computations, nodes go down or the scheduler aborts programs because they ran out of time. With computations this big, we cannot afford to just lose the results, so we checkpoint the computations every 50 time steps and can then resume it at the last saved state if necessary (see Section 4.5):

```

subsection Checkpointing
    set Steps between checkpoint = 50
end

```

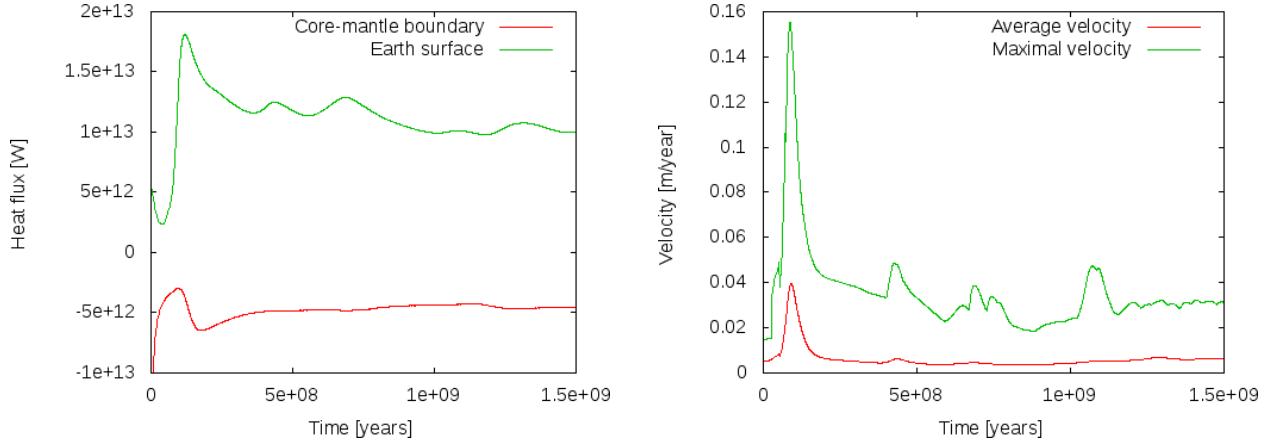


Figure 40: *Evaluating the 3d spherical shell model. Left: Outward heat fluxes through the inner and outer boundaries of the shell. Right: Average and maximal velocities in the mantle.*

Evaluation. Just as in the 2d case above, there are still many things that are wrong from a physical perspective in this setup, notably the no-slip boundary conditions at the bottom and of course the simplistic material model with its fixed viscosity and its neglect for adiabatic heating and compressibility. But there are also a number of things that are already order of magnitude correct here.

For example, if we look at the heat flux this model produces, we find that the convection here produces approximately the correct number. Wikipedia's article on [Earth's internal heat budget](#)³³ states that the overall heat flux through the Earth surface is about $47 \cdot 10^{12}$ W (i.e., 47 terawatts) of which an estimated 12–30 TW are primordial heat released from cooling the Earth and 15–41 TW from radiogenic heating.³⁴ Our model does not include radiogenic heating (though ASPECT has a number of [Heating models](#) to switch this on, see Section A.58) but we can compare what the model gives us in terms of heat flux through the inner and outer boundaries of our shell geometry. This is shown in the left panel of Fig. 40 where we plot the heat flux through boundaries zero and one, corresponding to the core-mantle boundary and Earth's surface. ASPECT always computes heat fluxes in outward direction, so the flux through boundary zero will be negative, indicating the we have a net flux *into* the mantle as expected. The figure indicates that after some initial jitters, heat flux from the core to the mantle stabilizes at around 4.5 TW and that through the surface at around 10 TW, the difference of 5.5 TW resulting from the overall cooling of the mantle. While we cannot expect our model to be quantitatively correct, this can be compared with estimates heat fluxes of 5–15 TW for the core-mantle boundary, and an estimated heat loss due to cooling of the mantle of 7–15 TW (values again taken from Wikipedia).

A second measure of whether these results make sense is to compare velocities in the mantle with what is known from observations. As shown in the right panel of Fig. 40, the maximal velocities settle to values on the order of 3 cm/year (each of the peaks in the line for the maximal velocity corresponds to a particularly large plume rising or falling). This is, again, at least not very far from what we know to be correct and we should expect that with a more elaborate material model we should be able to get even closer to reality.

5.3.3 Postprocessing spherical 3D convection

This section was contributed by Jacqueline Austermann, Ian Rose, and Shangxin Liu

³³Not necessarily the most scientific source, but easily accessible and typically about right in terms of numbers. The numbers stated here are those listed on Wikipedia at the time this section was written in March 2014.

³⁴As a point of reference, for the mantle an often used number for the release of heat due to radioactive decay is $7.4 \cdot 10^{-12}$ W/kg. Taking a density of 3300 kg/m^3 and a volume of 10^{12} m^3 would yield roughly $2.4 \cdot 10^{13}$ W of heat produced. This back of the envelope calculation lies within the uncertain range stated above.

There are several postprocessors that can be used to turn the velocity and pressure solution into quantities that can be compared to surface observations. In this cookbook ([cookbooks/shell_3d_postprocess.prm](#)) we introduce two postprocessors: dynamic topography and the geoid. We initialize the model with a harmonic perturbation of degree 4 and order 2 and calculate the instantaneous solution. Analogous to the previous setup we use a spherical shell geometry model and a simple material model.

The relevant section in the input file that determines the postprocessed output is as follows:

```
subsection Postprocess
  set List of postprocessors = velocity statistics, dynamic topography, visualization, basic
    ↳ statistics, geoid

  subsection Visualization
    set Output format          = vtu
    set List of output variables = geoid, dynamic topography, density, viscosity, gravity
    set Number of grouped files = 1
  end
end
```

This initial condition results in distinct flow cells that cause local up- and downwellings (Figure 41). This flow deflects the top and bottom boundaries of the mantle away from their reference height, a process known as dynamic topography. The deflection of the surfaces and density perturbations within the mantle also cause a perturbation in the gravitational field of the planet relative to the hydrostatic equilibrium ellipsoid.

Dynamic topography at the surface and core mantle boundary. Dynamic topography is calculated at the surface and bottom of the domain through a stress balancing approach where we assume that the radial stress at the surface is balanced by excess (or deficit) topography. We use the consistent boundary flux (CBF) method to calculate the radial stress at the surface [ZGH93]. For the bottom surface we define positive values as up (out) and negative values are down (in), analogous to the deformation of the upper surface. Dynamic topography can be outputted in text format (which writes the Euclidean coordinates followed by the corresponding topography value) or as part of the visualization. The upwelling and downwelling flow along the equator causes alternating topography high and lows at the top and bottom surface (Figure 41). In Figure 41 c, d we have subtracted the mean dynamic topography from the output field as a postprocessing step outside of Aspect. Since mass is conserved within the Earth, the mean dynamic topography should always be zero, however, the outputted values might not fulfill this constraint if the resolution of the model is not high enough to provide an accurate solution. This cookbook only uses a refinement of 2, which is relatively low resolution.

Geoid anomalies. Geoid anomalies are perturbations of the gravitational equipotential surface that are due to density variations within the mantle as well as deflections of the surface and core mantle boundary. The geoid anomalies are calculated using a spherical harmonic expansion of the respective fields. The user has the option to specify the minimum and maximum degree of this expansion. By default, the minimum degree is 2, which conserves the mass of the Earth (by removing degree 0) and chooses the Earth's center of mass as reference frame (by removing degree 1). In this model, downwellings coincide with lows in the geoid anomaly. That means the mass deficit caused by the depression at the surface is not fully compensated by the high density material below the depression that drags the surface down. The geoid postprocessor uses a spherical harmonic expansion and can therefore only be used with the 3D spherical shell geometry model.

5.3.4 3D convection with an Earth-like initial condition

This section was contributed by Jacqueline Austermann

For any model run with ASPECT we have to choose an initial condition for the temperature field. If we want to model convection in the Earth's mantle we want to choose an initial temperature distribution that

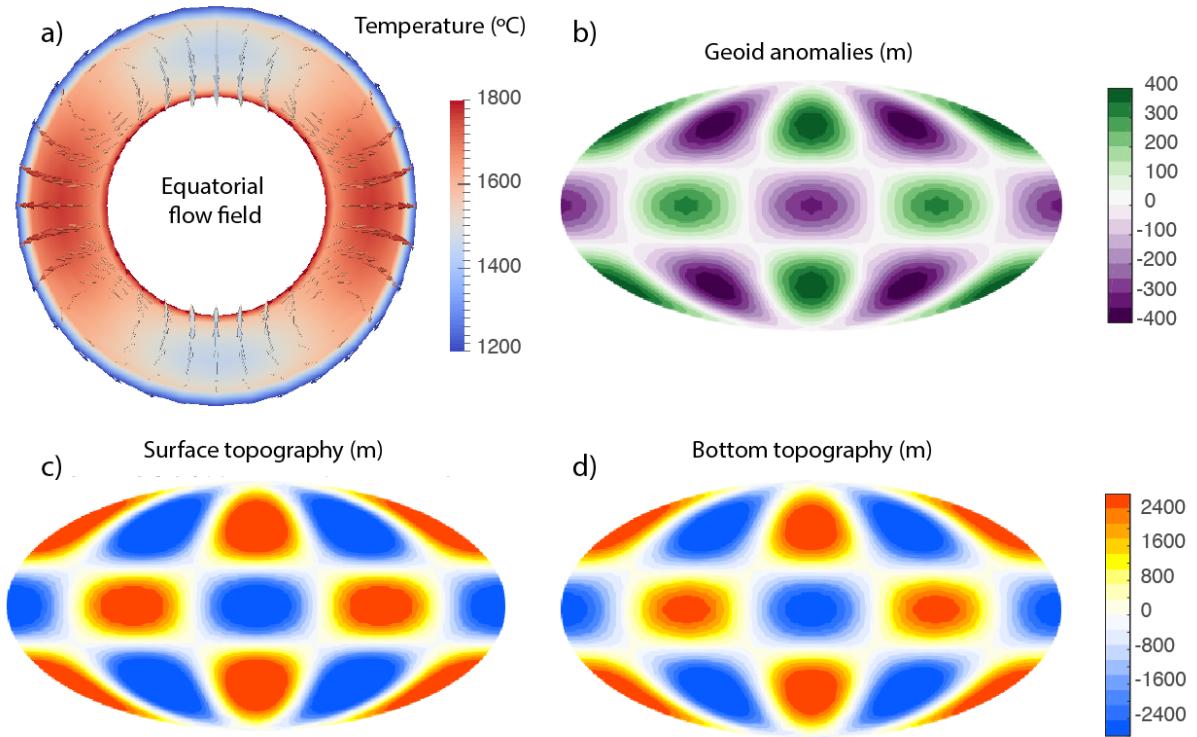


Figure 41: Panel (a) shows an equatorial cross section of the temperature distribution and resulting flow from a harmonic perturbation. Panel (b) shows the resulting geoid, and panels (c) and (d) show the resulting surface and bottom topography. Note that we have subtracted the mean surface and bottom topography in the respective panels (c and d) as a postprocessing step outside of Aspect.

captures the Earth's buoyancy structure. In this cookbook we present how to use temperature perturbations based on the shear wave velocity model S20RTS [RvH00] to initialize a mantle convection calculation.

The input shear wave model. The current version of ASPECT can read in the shear wave velocity models S20RTS [RvH00] and S40RTS [RDvHW11], which are located in `data/initial-conditions/S40RTS/`. Those models provide spherical harmonic coefficients up to degree 20 and 40, respectively, for 21 depth layers. The interpolation with depth is done through a cubic spline interpolation. The input files `S20RTS.sph` and `S40RTS.sph` were downloaded from <http://www.earth.lsa.umich.edu/~jritsema/Research.html> and have the following format (this example is S20RTS):

The first number in the first line denotes the maximum degree. This is followed in the next line by the spherical harmonic coefficients from the surface down to the CMB. The coefficients are arranged in the following way:

a_{00}				
a_{10}	a_{11}	b_{11}		
a_{20}	a_{21}	b_{21}	a_{22}	b_{22}
...				

a_{yz} is the cosine coefficient of degree y and order z and b_{yz} is the sine coefficient of degree y and order z . The depth layers are specified in the file `Spline_knots.txt` by a normalized depth value ranging from the CMB (3480km, normalized to -1) to the Moho (6346km, normalized to 1). This is the original format provided on the homepage.

Any other perturbation model in this same format can also be used, one only has to specify the different filename in the parameter file (see next section). For models with different depth layers one has to adjust the `Spline_knots.txt` file as well as the number of depth layers, which is hard coded in the current code. A further note of caution when switching to a different input model concerns the normalization of the spherical harmonics, which might differ. After reading in the shear wave velocity perturbation one has several options to scale this into temperature differences, which are then used to initialize the temperature field.

Setting up the ASPECT model. For this cookbook we will use the parameter file provided in `cookbooks/S2ORTS.prm`, which uses a 3d spherical shell geometry similar to section 5.3.2. This plugin is only sensible for a 3D spherical shell with Earth-like dimensions.

The relevant section in the input file is as follows:

```

subsection Initial temperature model
  set Model name = S4ORTS perturbation
subsection S4ORTS perturbation
  set Data directory          = $ASPECT_SOURCE_DIR/data/initial-temperature/S4ORTS/
  set Initial condition file name = S2ORTS.sph
  set Spline knots depth file name = Spline_knots.txt
  set Remove degree 0 from perturbation = false
  set Vs to density scaling      = 0.15
  set Thermal expansion coefficient in initial temperature scaling = 3e-5
  set Reference temperature       = 1600
end
end

```

For this initial condition model we need to first specify the data directory in which the input files are located as well as the initial condition file (S20RTS.sph or S40RTS.sph) and the file that contains the normalized depth layers (Spline knots depth file name). We next have the option to remove the degree 0 perturbation from the shear wave model. This might be the case if we want to make sure that the depth average temperature follows the background (adiabatic or constant) temperature.

The next input parameters describe the scaling from the shear wave velocity perturbation to the final temperature field. The shear wave velocity perturbation $\delta v_s/v_s$ (that is provided by S20RTS) is scaled into a density perturbation $\delta\rho/\rho$ with a constant that is specified in the initial condition section of the input parameter file as ‘Vs to density scaling’. Here we choose a constant scaling of 0.15. This perturbation is further translated into a temperature difference ΔT by multiplying it by the negative inverse of thermal expansion, which is also specified in this section of the parameter file as ‘Thermal expansion coefficient in initial temperature scaling’. This temperature difference is then added to the background temperature, which is the adiabatic temperature for a compressible model or the reference temperature (as specified in this section of the parameter file) for an incompressible model. Features in the upper mantle such as cratons might be chemically buoyant and therefore isostatically compensated, in which case their shear wave perturbation would not contribute buoyancy variations. We therefore included an additional option to zero out temperature perturbations within a certain depth, however, in this example we don’t make use of this functionality. The chemical variation within the mantle might require a more sophisticated ‘Vs to density’ scaling that varies for example with depth or as a function of the perturbation itself, which is not captured in this model. The described procedure provides an absolute temperature for every point, which will only be adjusted at the boundaries if indicated in the Boundary temperature model. In this example we chose a surface and core mantle boundary temperature that differ from the reference mantle temperature in order to approximate thermal boundary layers.

Visualizing 3D models. In this cookbook we calculate the instantaneous solution to examine the flow field. Figures 42 and 43 show some of the output for a resolution of 2 global refinement steps (42c and 43a, c, e) as used in the cookbook, as well as 4 global refinement steps (other panels in these figures). Computations with 4 global refinements are expensive, and consequently this is not the default for this cookbook. For example, as of 2017, it takes 64 cores approximately 2 hours of walltime to finish this cookbook with 4 global refinements. Figure 42a and b shows the density variation that has been obtained from scaling S20RTS in the way described above. One can see the two large low shear wave velocity provinces underneath Africa and the Pacific that lead to upwelling if they are assumed to be buoyant (as is done in this case). One can also see the subducting slabs underneath South America and the Philippine region that lead to local downwelling. Figure 42c and d shows the heat flux density at the surface for 2 refinement steps (c, colorbar ranges from 13 to 19 mW/m²) and for 4 refinement steps (d, colorbar ranges from 35 to 95 mW/m²). A first order correlation with upper mantle features such as high heat flow at mid ocean ridges and low heat flow at cratons is correctly initialized by the tomography model. The mantle flow and buoyancy variations produce dynamic topography on the top and bottom surface, which is shown for 2 refinement steps (43a and c, respectively) and 4 refinement steps (43b and d, respectively). One can see that subduction zones are visible as depressed surface topography due to the downward flow, while regions such as Iceland, Hawaii, or mid ocean ridges are elevated due to (deep and) shallow upward flow. The core mantle boundary topography shows that the upwelling large low shear wave velocity provinces deflect the core mantle boundary up. Lastly, Figure 43e and f show geoid perturbations for 2 and 4 global refinement steps, respectively. The geoid anomalies show a strong correlation with the surface dynamic topography. This is in part expected given that the geoid anomalies are driven by the deflection of the upper and lower surface as well as internal density variations. The relative importance of these different contributors is dictated by the Earth’s viscosity profile. Due to the isoviscous assumption in this cookbook, we don’t properly recover patterns of the observed geoid.

As discussed in the previous cookbook, dynamic topography does not necessarily average to zero if the resolution is not high enough. While one can simply subtract the mean as a postprocessing step this should be done with caution since a non-zero mean indicates that the refinement is not sufficiently high to resolve the convective flow. In Figure 43a-d we refrained from subtracting the mean but indicated it at the bottom

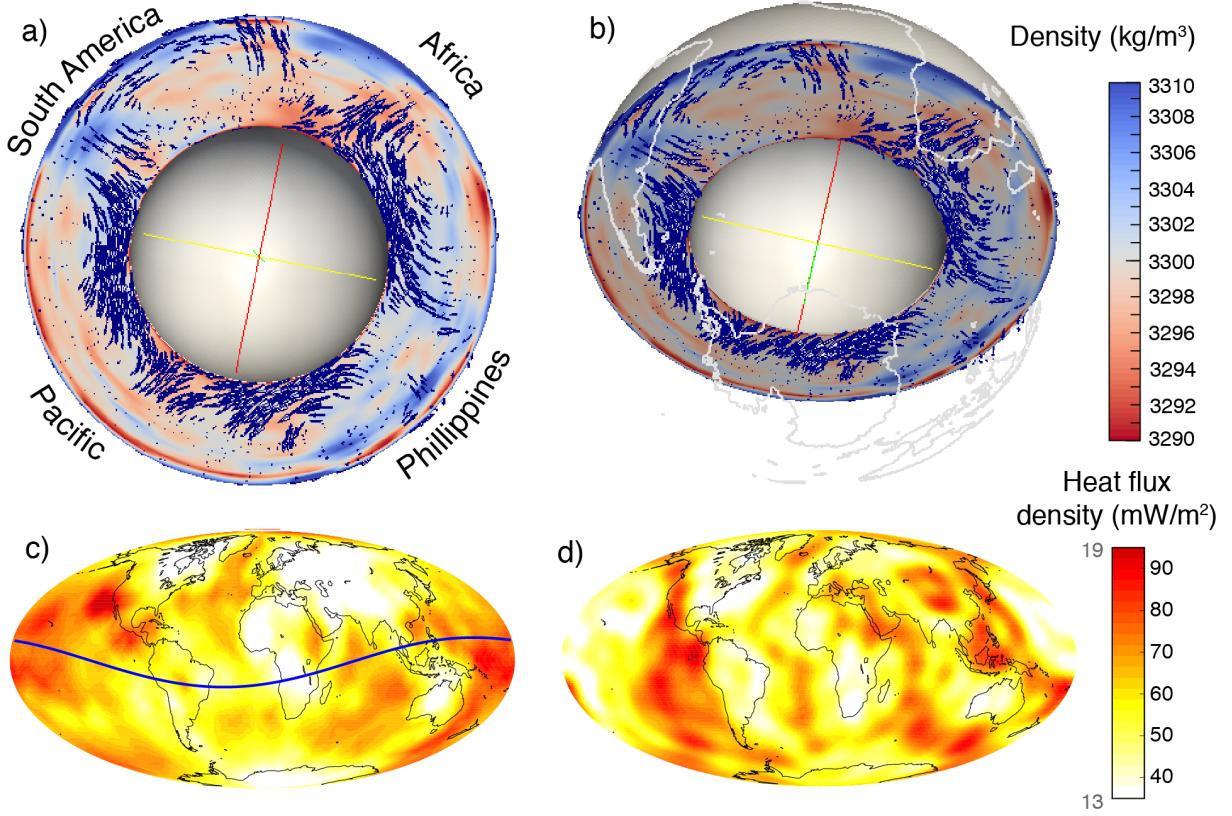


Figure 42: Panels (a) and (b) show the density distribution as prescribed from the shear wave velocity model S20RTS and the resulting flow for a global refinement of 4. This model assumes a constant scaling between shear wave and density perturbations. Panel (c) shows the great circle (dashed blue line) along which the top slices are evaluated. Panels (c) and (d) show the resulting heat flux density for a global refinement of 2 (c, cookbook) and 4 (d). The colorbar ranges from 13 to 19 mW/m² for panel (c) and from 35 to 95 mW/m² for panel (d).

left of each panel. The mean dynamic topography approaches zero for increasing refinement. Furthermore, the mean bottom dynamic topography is closer to zero than the mean top dynamic topography. This is likely due to the larger magnitude of dynamic topography at the surface and the difference in resolution between the top and bottom domain (for a given refinement, the resolution at the core mantle boundary is higher than the resolution at the surface). The average geoid height is zero since the minimum degree in the geoid anomaly expansion is set to 2.

This model uses a highly simplified material model that is incompressible and isoviscous and does therefore not represent real mantle flow. More realistic material properties, density scaling as well as boundary conditions will affect the magnitudes and patterns shown here.

5.3.5 Using reconstructed surface velocities by GPlates

This section was contributed by René Gäßmöller

In a number of model setups one may want to include a surface velocity boundary condition that prescribes the velocity according to a specific geologic reconstruction. The purpose of this kind of models is often to test

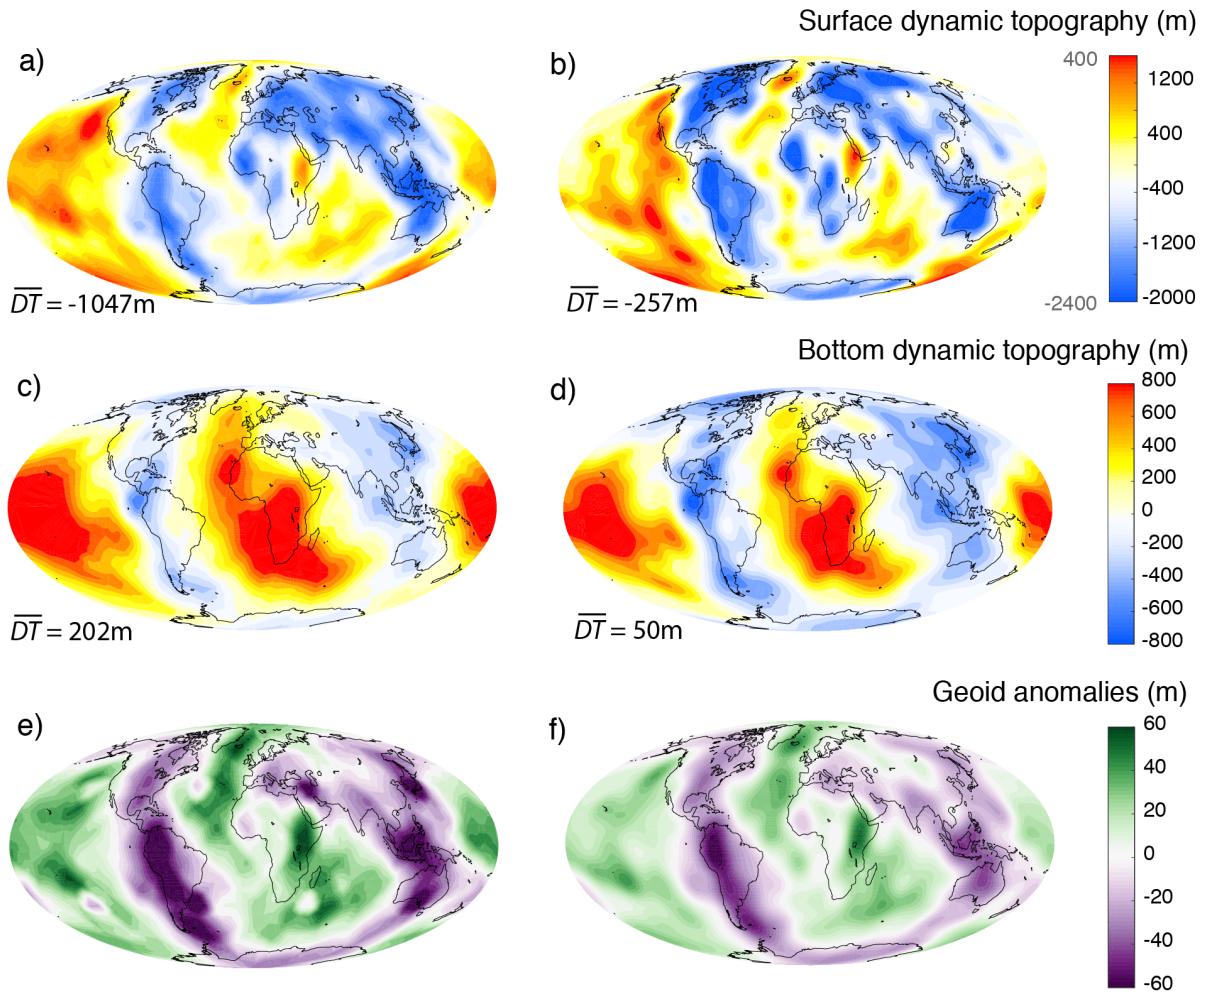


Figure 43: The first row of this figure shows the surface dynamic topography resulting from the flow shown in Figure 42 for a global refinement of 2 (a, cookbook) and 4 (b). The colorbar ranges from -2400m to 400m for panel (a) and from -2000m to 1600m for panel (b). The second row shows the dynamic topography at the core mantle boundary for the same model and a refinement of 2 (c, cookbook) and 4 (d). Averages of the dynamic topography fields are indicated at the bottom left of each panel. The bottom row shows the geoid anomalies from this model at the surface for refinement of 2 (e, cookbook) and 4 (f).

a proposed geologic model and compare characteristic convection results to present-day observables in order to gain information about the initially assumed geologic input. In this cookbook we present ASPECT’s interface to the widely used plate reconstruction software GPlates, and the steps to go from a geologic plate reconstruction to a geodynamic model incorporating these velocities as boundary condition.

Acquiring a plate reconstruction. The plate reconstruction that is used in this cookbook is included in the `data/velocity-boundary-conditions/gplates/` directory of your ASPECT installation. For a new model setup however, a user eventually needs to create her own data files, and so we will briefly discuss the process of acquiring a usable plate reconstruction and transferring it into a format usable by ASPECT. Both the necessary software and data are provided by the GPlates project. GPlates is an open-source software for interactive visualization of plate tectonics. It is developed by the EarthByte Project in the School of Geosciences at the University of Sydney, the Division of Geological and Planetary Sciences (GPS) at CalTech and the Center for Geodynamics at the Norwegian Geological Survey (NGU). For extensive documentation and support we refer to the GPlates website (<http://www.gplates.org>). Apart from the software one needs the actual plate reconstruction that consists of closed polygons covering the complete model domain. For our case we will use the data provided by [GTZ⁺¹²] that is available from the GPlates website under “Download → Download GPlates-compatible data → Global reconstructions with continuously closing plates from 140 Ma to the present”. The data is provided under a Creative Commons Attribution 3.0 Unported License (<http://creativecommons.org/licenses/by/3.0/>).

Converting GPlates data to ASPECT input. After loading the data files into GPlates (*.gpml for plate polygons, *.rot for plate rotations over time) the user needs to convert the GPlates data to velocity information usable in ASPECT. The purpose of this step is to convert from the description GPlates uses internally (namely a representation of plates as polygons that rotate with a particular velocity around a pole) to one that can be used by ASPECT (which needs velocity vectors defined at individual points at the surface).

With loaded plate polygon and rotation information the conversion from GPlates data to ASPECT-readable velocity files is rather straightforward. First the user needs to generate (or import) so-called “velocity domain points”, which are discrete sets of points at which GPlates will evaluate velocity information. This is done using the “Features → Generate Velocity Domain Points → Latitude Longitude” menu option. Because ASPECT is using an adaptive mesh it is not possible for GPlates to generate velocity information at the precise surface node positions like for CitcomS or Terra (the other currently available interfaces). Instead GPlates will output the information on a general Latitude/Longitude grid with nodes on all crossing points. ASPECT then internally interpolates this information to the current node locations during the model run. This requires the user to choose a sensible resolution of the GPlates output, which can be adjusted in the “Generate Latitude/Longitude Velocity Domain Points” dialog of GPlates. In general a resolution that resolves the important features is necessary, while a resolution that is higher than the maximal mesh size for the ASPECT model is unnecessary and only increases the computational cost and memory consumption of the model.

Important note: The Mesh creation routine in GPlates has significantly changed from version 1.3 to 1.4. In GPlates 1.4 and later the user has to make sure that the number of longitude intervals is set as twice the number of latitude intervals, the “Place node points at centre of latitude/longitude cells” box is unchecked and the “Latitude/Longitude extents” are set to “Use Global Extents”. ASPECT does check for most possible combinations that can not be read and will cancel the calculation in these cases, however some mistakes can not be checked against from the information provided in the GPlates file.

After creating the Velocity Domain Points the user should see the created points and their velocities indicated as points and arrows in GPlates. To export the calculated velocities one would use the “Reconstruction → Export” menu. In this dialog the user may specify the time instant or range at which the velocities shall be exported. The only necessary option is to include the “Velocities” data type in the “Add Export” sub-dialog. The velocities need to be exported in the native GPlates *.gpml format, which is based on XML and can be read by ASPECT. In case of a time-range the user needs to add a pattern specifier to

the name to create a series of files. The `%u` flag is especially suited for the interaction with ASPECT, since it can easily be replaced by a calculated file index (see also [5.3.5](#)).

Setting up the ASPECT model. For this cookbook we will use the parameter file provided in [cookbooks/gplates-2d.prm](#) which uses the 2d shell geometry previously discussed in Section [5.3.1](#). ASPECT’s GPlates plugin allows for the use of two- and three-dimensional models incorporating the GPlates velocities. Since the output by GPlates is three-dimensional in any case, ASPECT internally handles the 2D model by rotating the model plane to the orientation specified by the user and projecting the plate velocities into this plane. The user specifies the orientation of the model plane by prescribing two points that define a plane together with the coordinate origin (i.e. in the current formulation only great-circle slices are allowed). The coordinates need to be in spherical coordinates θ and ϕ with θ being the colatitude (0 at north pole) and ϕ being the longitude (0 at Greenwich meridian, positive eastwards) both given in radians. The approach of identifying two points on the surface of the Earth along with its center allows to run computations on arbitrary two-dimensional slices through the Earth with realistic boundary conditions.

The relevant section of the input file is then as follows:

```
subsection Boundary temperature model
  set Fixed temperature boundary indicators = inner, outer
end

subsection Boundary velocity model
  set Prescribed velocity boundary indicators = outer:gplates
  set Tangential velocity boundary indicators = inner
  subsection GPlates model
    set Data directory = $ASPECT_SOURCE_DIR/data/boundary-velocity/gplates/
    set Velocity file name = current_day.gpml
    set Time step = 1e6
    set Point one = 1.5708,4.87
    set Point two = 1.5708,5.24
    set Interpolation width = 2000000
  end
end
```

In the “Boundary velocity model” subsection the user prescribes the boundary that is supposed to use the GPlates plugin. Although currently nothing forbids the user to use GPlates plugin for other boundaries than the surface, its current usage and the provided sample data only make sense for the surface of a spherical shell (boundary number 1 in the above provided parameter file). In case you are familiar with this kind of modeling and the plugin you could however also use it to prescribe mantle movements *below* a lithosphere model. All plugin specific options may be set in section [A.32](#). Possible options include the data directory and file name of the velocity file/files, the time step (in model units, mostly seconds or years depending on the “Use years in output instead of seconds” flag) and the points that define the 2D plane. The parameter “Interpolation width” is used to smooth the provided velocity files by a moving average filter. All velocity data points within this distance are averaged to determine the actual boundary velocity at a certain mesh point. This parameter is usually set to 0 (no interpolation, use nearest velocity point data) and is only needed in case the original setting is unstable or slowly converging.

Comparing and visualizing 2D and 3D models. The implementation of plate velocities in both two- and three-dimensional model setups allows for an easy comparison and test for common sources of error in the interpretation of model results. The left top figure in Fig. [44](#) shows a modification of the above presented parameter file by setting “Dimension = 3” and “Initial global refinement = 3”. The top right plot of Fig. [44](#) shows an example of three independent two-dimensional computations of the same reduced resolution. The models were prescribed to be orthogonal slices by setting:

```
set Point one = 3.1416,0.0
```

```
set Point two = 1.5708,0.0
```

and

```
set Point one = 3.1416,1.5708  
set Point two = 1.5708,1.5708
```

The results of these models are plotted simultaneously in a single three-dimensional figure in their respective model planes. The necessary information to rotate the 2D models to their respective planes (rotation axis and angle) is provided by the GPlates plugin in the beginning of the model output. The bottom plot of Fig. 44 finally shows the results of the original [cookbooks/gplates-2d.prm](#) also in the three mentioned planes.

Now that we have model output for otherwise identical 2D and 3D models with equal resolution and additional 2D output for a higher resolution an interesting question to ask would be: What additional information can be created by either using three-dimensional geometry or higher resolution in mantle convection models with prescribed boundary velocities. As one can see in the comparison between the top right and bottom plot in Fig. 44 additional resolution clearly improves the geometry of small scale features like the shape of up- and downwellings as well as the maximal temperature deviation from the background mantle. However, the limitation to two dimensions leads to inconsistencies, that are especially apparent at the cutting lines of the individual 2D models. Note for example that the Nazca slab of the South American subduction zone is only present in the equatorial model plane and is not captured in the polar model plane west of the South American coastline. The (coarse) three-dimensional model on the other hand shows the same location of up- and downwellings but additionally provides a consistent solution that is different from the two dimensional setups. Note that the Nazca slab is subducting eastward, while all 2D models (even in high resolution) predict a westward subduction.

Finally we would like to emphasize that these models (especially the used material model) are way too simplified to draw any scientific conclusion out of it. Rather it is thought as a proof-of-concept what is possible with the dimension independent approach of ASPECT and its plugins.

Time-dependent boundary conditions. The example presented above uses a constant velocity boundary field that equals the present day plate movements. For a number of purposes one may want to use a prescribed velocity boundary condition that changes over time, for example to investigate the effect of trench migration on subduction. Therefore ASPECT's GPlates plugin is able to read in multiple velocity files and linearly interpolate between pairs of files to the current model time. To achieve this, one needs to use the `%d` wildcard in the velocity file name, which represents the current velocity file index (e.g. `time_dependent.%d.gpml`). This index is calculated by dividing the current model time by the user-defined time step between velocity files (see parameter file above). As the model time progresses the plugin will update the interpolation accordingly and if necessary read in new velocity files. In case it can not read the next velocity file, it assumes the last velocity file to be the constant boundary condition until the end of the model run. One can test this behavior with the provided data files `data/velocity_boundary_conditions/gplates/time_dependent.%d.gpml` with the index `d` ranging from 0 to 3 and representing the plate movements of the last 3 million years corresponding to the same plate reconstruction as used above. Additionally, the parameter `Velocity file start time` allows for a period of no-slip boundary conditions before starting the use of the GPlates plugin. This is a comfort implementation, which could also be achieved by using the checkpointing possibility described in section 4.5.

5.3.6 2D compressible convection with a reference profile and material properties from Burn-Man

This section was contributed by Juliane Dannberg and René Gassmöller

In this cookbook we will set up a compressible mantle convection model that uses the (truncated) anelastic liquid approximation (see Sections 2.10.1 and 2.10.2), together with a reference profile read in from an ASCII data file. The data we use here is generated with the open source mineral physics toolkit BurnMan

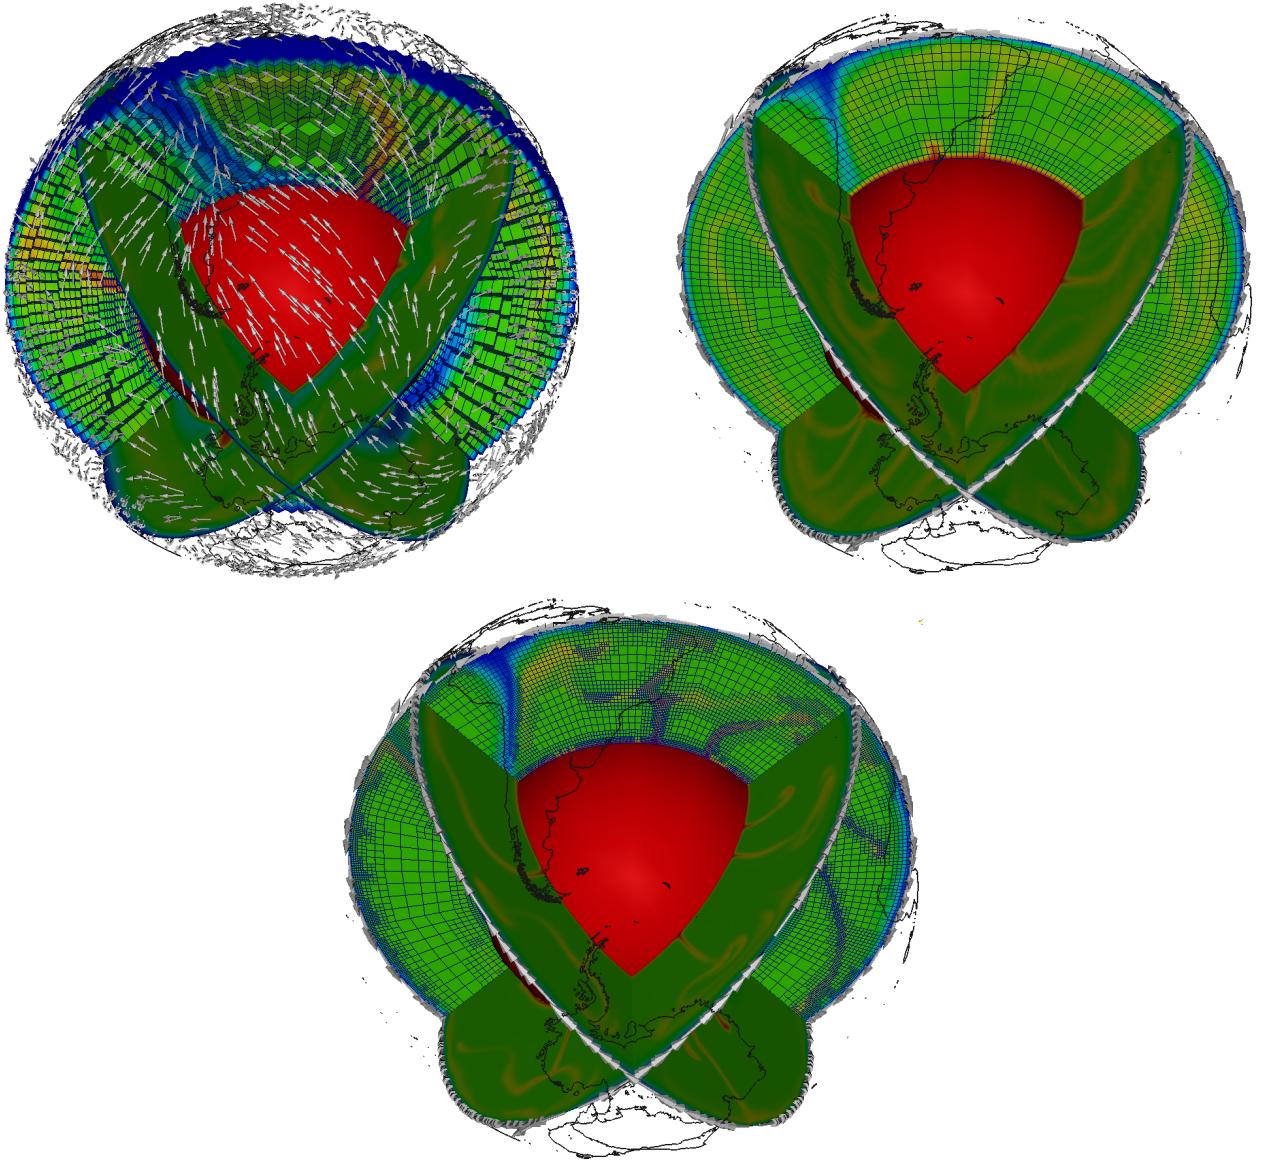


Figure 44: *Using GPlates for velocity boundary conditions:* The top left figure shows the results of a three-dimensional model using the present day plate velocities provided by GPlates as surface boundary condition. The top right figure shows three independent computations on two-dimensional slices through Earth. The boundary conditions for each of these slices (white arrows) are tangential to the slices and are projections of the three-dimensional velocity vectors into the two-dimensional plane occupied by the slice. While the two top models are created with the same mesh resolution the bottom figure shows three independent two-dimensional models using a higher resolution. The view is centered on South America with Antarctica being near the bottom of the figure (coastlines provided by NGU and the GPlates project).

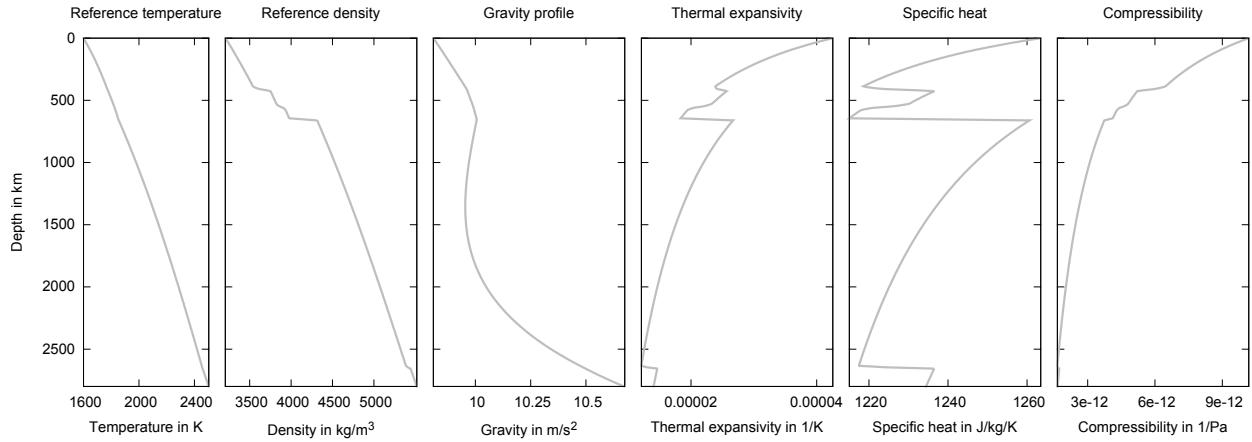


Figure 45: Reference profile generated using BurnMan.

(<http://www.burnman.org>) using the python example program `simple_adiabat.py`. This file is available as a part of BurnMan, and provides a tutorial for how to generate ASCII data files that can be used together with ASPECT. The computation is based on the Birch-Murnaghan equation of state, and uses a harzburgitic composition. However, in principle, other compositions or equations of state can be used, as long as the reference profile contains data for the reference temperature, pressure, density, gravity, thermal expansivity, specific heat capacity and compressibility. Using BurnMan to generate the reference profile has the advantage that all the material property data are consistent, for example, the gravity profile is computed using the reference density. The reference profile is shown in Figure 45, and the corresponding data file is located at [data/adiabatic-conditions/ascii-data/isentrope_properties.txt](#).

Setting up the ASPECT model. In order to use this profile, we have to import and use the data in the adiabatic conditions model, in the gravity model and in the material model, which is done using the corresponding ASCII data plugins. The input file is provided in [cookbooks/burnman.prm](#), and it uses the 2d shell geometry previously discussed in Section 5.3.1 and surface velocities imported from GPlates as explained in Section 5.3.5.

To use the BurnMan data in the material model, we have to specify that we want to use the `ascii reference profile` model. This material model makes use of the functionality provided by the `AsciiData` classes in ASPECT, which allow plugins such as material models, boundary or initial conditions models to read in ASCII data files (see for example Section 5.2.12). Hence, we have to provide the directory and file name of the data to be used in the separate subsection `Ascii data model`, and the same functionality and syntax will also be used for the adiabatic conditions and gravity model.

The viscosity in this model is computed as the product of a profile $\eta_r(z)$, where z corresponds to the depth direction of the chosen geometry model, and a term that describes the dependence on temperature:

$$\eta(z, T) = \eta_r(z) \eta_0 \exp\left(-A \frac{T - T_{\text{adi}}}{T_{\text{adi}}}\right),$$

where A and η_0 are constants determined in the input file via the parameters `Viscosity` and `Thermal viscosity exponent`, and $\eta_r(z)$ is a stepwise constant function that determines the viscosity profile. This function can be specified by providing a list of `Viscosity prefactors` and a list of depths that describe in which depth range each prefactor should be applied, in other words, at which depth the viscosity changes. By default, it is set to viscosity jumps at 150 km depth, between upper mantle and transition zone, and between transition zone and lower mantle). The prefactors used here lead to a low-viscosity asthenosphere,

and high viscosities in the lower mantle. To make sure that these viscosity jumps do not lead to numerical problems in our computation (see Section 5.2.8), we also use harmonic averaging of the material properties.

```

subsection Material model
  set Model name = ascii reference profile

  subsection Ascii data model
    set Data file name = isentrope_properties.txt
    set Data directory = $ASPECT_SOURCE_DIR/data/adiabatic-conditions/ascii-data/
  end

  subsection Ascii reference profile
    set Thermal viscosity exponent = 10.0
    set Viscosity prefactors = 1.0, 0.1, 1.0, 10.0
  end

  set Material averaging = harmonic average
end

```

As the reference profile has a depth dependent density and also contains data for the compressibility, this material model supports compressible convection models.

For the adiabatic conditions and the gravity model, we also specify that we want to use the respective `ascii` data plugin, and provide the data directory in the same way as for the material model. The gravity model automatically uses the same file as the adiabatic conditions model.

```

subsection Adiabatic conditions model
  set Model name = ascii data

  subsection Ascii data model
    set Data directory = $ASPECT_SOURCE_DIR/data/adiabatic-conditions/ascii-data/
    set Data file name = isentrope_properties.txt
  end
end

```

```

subsection Gravity model
  set Model name = ascii data
end

```

To make use of the reference state we just imported from BurnMan, we choose a formulation of the equations that employs a reference state and compressible convection, in this case the anelastic liquid approximation (see Section 2.10.1).

```

subsection Formulation
  set Formulation = anelastic liquid approximation
end

```

This means that the reference profiles are used for all material properties in the model, except for the density in the buoyancy term (on the right-hand side of the force balance equation (1), which in the limit of the anelastic liquid approximation becomes Equation (21)). In addition, the density derivative in the mass conservation equation (see Section 2.11.1) is taken from the adiabatic conditions, where it is computed as the depth derivative of the provided reference density profile (see also Section 2.11.5).

Visualizing the model output. If we look at the output of our model (for example in `ParaView`), we can see how cold, highly viscous slabs are subducted and hot plumes rise from the core-mantle boundary. The final time step of the model is shown in Figure 46, and the full model evolution can be found at <https://youtu.be/nRB0pw5kp-4>. Visualizing material properties such as density, thermal expansivity or

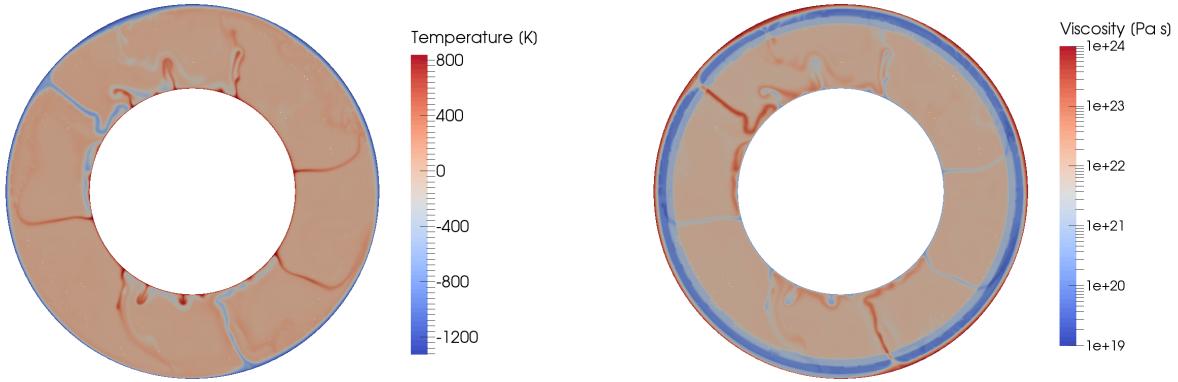


Figure 46: Compressible convection in a 2d spherical shell, using a reference profile exported form BurnMan, which is based on the Birch-Murnaghan equation of state. The figure shows the state at the end of the model evolution over 260 Ma.

specific heat shows how they change with depth, and reveals abrupt jumps at the phase transitions, where properties change from one mineral phase to the next. We can also visualize the gravity and the adiabatic profile, to ensure that the data we provided in the [data/adiabatic-conditions/ascii-data/isentrope_properties.txt](#) file is used in our model.

Comparing different model approximations. For the model described above, we have used the anelastic liquid approximation. However, one might want to use different approximations that employ a reference state, such as the truncated anelastic liquid approximation (TALA, see Section 2.10.2), which is also supported by the `ascii reference profile` material model. In this case, the only change compared to ALA is in the density used in the buoyancy term, the only place where the temperature-dependent density instead of the reference density is used. For the TALA, this density only depends on the temperature (and not on the dynamic pressure, as in the ALA). Hence, we have to make this change in the appropriate place in the material model (while keeping the formulation of the equations set to `anelastic liquid approximation`):

```
subsection Material model
  subsection Ascii reference profile
    set Use TALA = true
  end
end
```

We now want to compare these commonly used approximations to the “isothermal compression approximation” (see Section 2.10.4) that is unique to ASPECT. It does not require a reference state and uses the full density everywhere in the equations except for the right-hand side mass conservation, where the compressibility is used to compute the density derivative with regard to pressure. Nevertheless, this formulation can make use of the reference profile computed by BurnMan and compute the dependence of material properties on temperature and pressure in addition to that by taking into account deviations from the reference profile in both temperature and pressure. As this requires a modification of the equations outside of the material model, we have to specify this change in the `Formulation` (and remove the lines for the use of TALA discussed above).

```
subsection Formulation
  set Formulation = isothermal compression
end
```

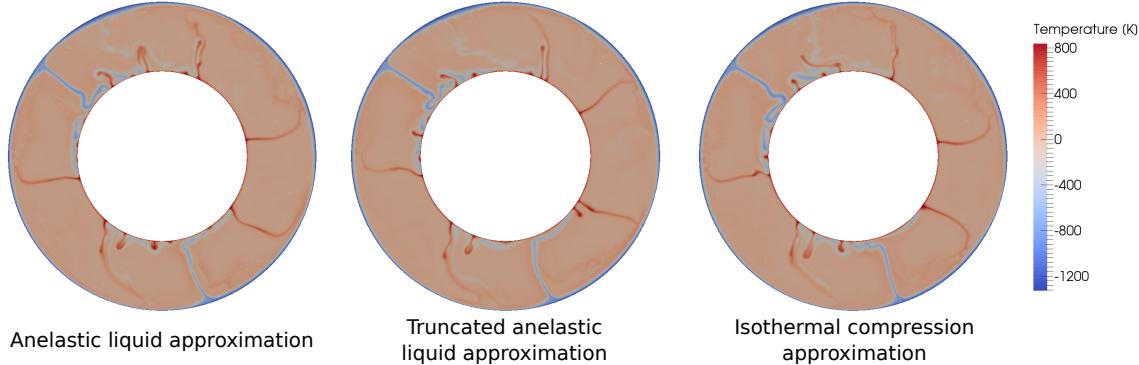


Figure 47: Comparison between the anelastic liquid approximation, the truncated anelastic liquid approximation and the isothermal compression approximation, showing the temperature distribution for the different models at the end of the model evolution at 260 Ma.

As the “isothermal compression approximation” is also ASPECT’s default for compressible models, the same model setup can also be achieved by just removing the lines that specify which `Formulation` should be used.

The Figures 47 and 48 show a comparison between the different models. They demonstrate that upwellings and downwellings may occur in slightly different places and at slightly different times when using a different approximation, but averaged model properties describing the state of the model – such as the root mean square velocity – are similar between the models.

5.3.7 Reproducing rheology of Morency and Doin, 2004

This section was contributed by Jonathan Perry-Houts

Modeling interactions between the upper mantle and the lithosphere can be difficult because of the dynamic range of temperatures and pressures involved. Many simple material models will assign very high viscosities at low temperature thermal boundary layers. The pseudo-brittle rheology described in [MD04] was developed to limit the strength of lithosphere at low temperature. The effective viscosity can be described as the harmonic mean of two non-Newtonian rheologies:

$$v_{\text{eff}} = \left(\frac{1}{v_{\text{eff}}^v} + \frac{1}{v_{\text{eff}}^p} \right)^{-1}$$

where

$$\begin{aligned} v_{\text{eff}}^v &= B \left(\frac{\dot{\epsilon}}{\dot{\epsilon}_{ref}} \right)^{-1+1/n_v} \exp \left(\frac{E_a + V_a \rho_m g z}{n_v R T} \right), \\ v_{\text{eff}}^p &= (\tau_0 + \gamma \rho_m g z) \left(\frac{\dot{\epsilon}^{-1+1/n_p}}{\dot{\epsilon}_{ref}^{1/n_p}} \right), \end{aligned}$$

where B is a scaling constant; $\dot{\epsilon}$ is defined as the quadratic sum of the second invariant of the strain rate tensor and a minimum strain rate, $\dot{\epsilon}_0$; $\dot{\epsilon}_{ref}$ is a reference strain rate; n_v , and n_p are stress exponents; E_a is the activation energy; V_a is the activation volume; ρ_m is the mantle density; R is the gas constant; T is temperature; τ_0 is the cohesive strength of rocks at the surface; γ is a coefficient of yield stress increase with depth; and z is depth.

By limiting the strength of the lithosphere at low temperature, this rheology allows one to more realistically model processes like lithospheric delamination and foundering in the presence of weak crustal layers.

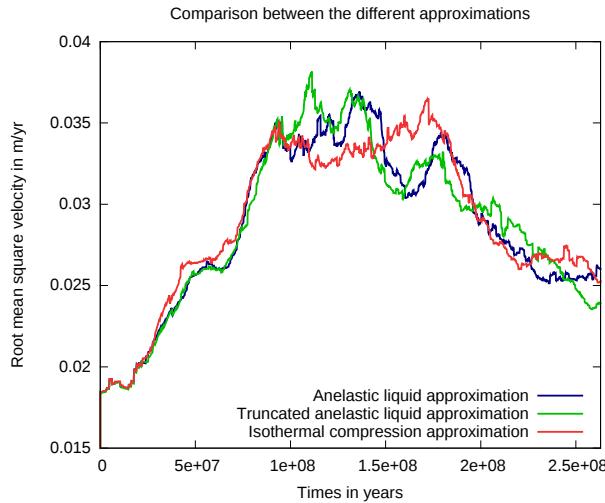


Figure 48: *Comparison between the anelastic liquid approximation, the truncated anelastic liquid approximation and the isothermal compression approximation, showing the evolution of the root mean square velocity.*

A similar model setup to the one described in [MD04] can be reproduced with the files in the directory `cookbooks/morency_doin_2004`. In particular, the following sections of the input file are important to reproduce the setup:

Note: [MD04] defines the second invariant of the strain rate in a nonstandard way. The formulation in the paper is given as $\epsilon_{II} = \sqrt{\frac{1}{2}(\epsilon_{11}^2 + \epsilon_{12}^2)}$, where ϵ is the strain rate tensor. For consistency, that is also the formulation implemented in ASPECT. Because of this irregularity it is inadvisable to use this material model for purposes beyond reproducing published results.

Note: The viscosity profile in Figure 1 of [MD04] appears to be wrong. The published parameters do not reproduce those viscosities; it is unclear why. The values used here get very close. See Figure 49 for an approximate reproduction of the original figure.

```

subsection Geometry model
  set Model name = box

  subsection Box
    set X extent      = 3000e3
    set Y extent      = 750e3
    set X repetitions = 4
  end
end

subsection Compositional fields
  set Number of fields = 2
  set Names of fields  = upper_crust, lower_crust
end

subsection Initial composition model
  set Model name = function

```

```

subsection Function
  set Variable names = x,y
  set Function expression = if(y>=725e3,1,0);if((y<725e3&y>700e3),1,0)
end
end

subsection Initial temperature model
  set Model name = function

  subsection Function
    set Variable names = x,y
    set Function constants = h=750e3, w=3000e3, mantleT=1350 # deg C
    set Function expression = \
      if( y < 100e3, \
        (100e3-y)/100e3*(1600-mantleT)+mantleT+293, \
        if(y>650e3, \
          (h-y)/(100e3)*mantleT+293, \
          mantleT+293))
    end
  end

subsection Material model
  set Model name = Morency and Doin

  subsection Morency and Doin
    set Densities = 3300,2920,2920
    set Activation energies = 500,320,320
    set Coefficient of yield stress increase with depth = 0.25
    set Thermal expansivities = 3.5e-5
    set Stress exponents for viscous rheology = 3
    set Stress exponents for plastic rheology = 30
    set Thermal diffusivity = 0.8e-6
    set Heat capacity = 1.25e3
    set Activation volume = 6.4e-6
    set Reference strain rate = 6.4e-16
    set Preexponential constant for viscous rheology law = 7e11 ## Value used in paper is 1.24e14
    set Cohesive strength of rocks at the surface = 117
    set Reference temperature = 293
    set Minimum strain rate = 5e-19 ## Value used in paper is 1.4e-20
  end
end

```

5.3.8 Crustal deformation

This section was contributed by Cedric Thieulot, and makes use of the Drucker-Prager material model written by Anne Glerum and the free surface plugin by Ian Rose.

This is a simple example of an upper-crust undergoing compression or extension. It is characterized by a single layer of visco-plastic material subjected to basal kinematical boundary conditions. In compression, this setup is somewhat analogous to [Wil99], and in extension to [AHT11].

Brittle failure is approximated by adapting the viscosity to limit the stress that is generated during deformation. This “cap” on the stress level is parameterized in this experiment by the pressure-dependent Drucker Prager yield criterion and we therefore make use of the Drucker-Prager material model (see Section A.81) in the cookbooks/crustal_model_2D.prm.

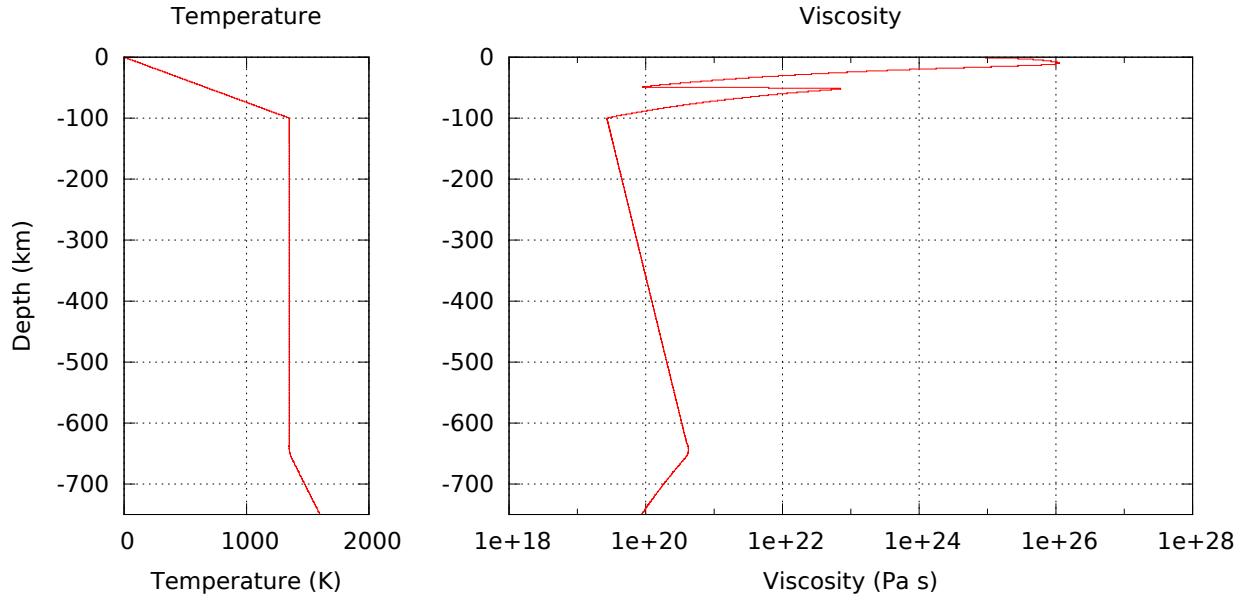


Figure 49: Approximate reproduction of figure 1 from [MD04] using the ‘morency doin’ material model with almost all default parameters. Note the low-viscosity Moho, enabled by the low activation energy of the crustal component.

The layer is assumed to have dimensions of 80km × 16km and to have a density $\rho = 2800 \text{ kg/m}^3$. The plasticity parameters are specified as follows:

```

subsection Material model
  set Model name = drucker prager
  subsection Drucker Prager
    set Reference density = 2800
    subsection Viscosity
      set Minimum viscosity = 1e19
      set Maximum viscosity = 1e25
      set Reference strain rate = 1e-20
      set Angle internal friction = 30
      set Cohesion = 20e6
    end
  end
end

```

The yield strength σ_y is a function of pressure, cohesion and angle of friction (see Drucker-Prager material model in Section A.81), and the effective viscosity is then computed as follows:

$$\mu_{\text{eff}} = \left(\frac{1}{\frac{\sigma_y}{2\dot{\epsilon}}} + \frac{1}{\mu_{\text{max}}} \right)^{-1}$$

where $\dot{\epsilon}$ is the square root of the second invariant of the deviatoric strain rate. The viscosity cutoffs insure that the viscosity remains within computationally acceptable values.

During the first iteration of the first timestep, the strain rate is zero, so we avoid dividing by zero by setting the strain rate to **Reference strain rate**.

The top boundary is a free surface while the left, right and bottom boundaries are subjected to the following boundary conditions:

```

subsection Boundary velocity model
    subsection Function
        set Variable names      = x,y
        set Function constants = cm=0.01, year=1
        set Function expression = if (x<40e3 , 1*cm/year, -1*cm/year) ; 0
    end
end

```

Note that compressive boundary conditions are simply achieved by reversing the sign of the imposed velocity.

The free surface will be advected up and down according to the solution of the Stokes solve. We have a choice whether to advect the free surface in the direction of the surface normal or in the direction of the local vertical (i.e., in the direction of gravity). For small deformations, these directions are almost the same, but in this example the deformations are quite large. We have found that when the deformation is large, advecting the surface vertically results in a better behaved mesh, so we set the following in the free surface subsection:

```

subsection Free surface
    set Surface velocity projection = vertical
end

```

We also make use of the strain rate-based mesh refinement plugin:

```

subsection Mesh refinement
    set Initial adaptive refinement      = 1
    set Initial global refinement       = 3
    set Refinement fraction           = 0.95
    set Strategy                      = strain rate
    set Coarsening fraction          = 0.05
    set Time steps between mesh refinement = 1
    set Run postprocessors on initial refinement = true
end

```

Setting `set Initial adaptive refinement = 4` yields a series of meshes as shown in Fig. (50), all produced during the first timestep. As expected, we see that the location of the highest mesh refinement corresponds to the location of a set of conjugated shear bands.

If we now set this parameter to 1 and allow the simulation to evolve for 500kyr, a central graben or plateau (depending on the nature of the boundary conditions) develops and deepens/thickens over time, nicely showcasing the unique capabilities of the code to handle free surface large deformation, localised strain rates through visco-plasticity and adaptive mesh refinement as shown in Fig. (51).

Deformation localizes at the basal velocity discontinuity and plastic shear bands form at an angle of approximately 53° to the bottom in extension and 35° in compression, both of which correspond to the reported Arthur angle [Kau10, Bui12].

Extension to 3D. We can easily modify the previous input file to produce `crustal_model_3D.prm` which implements a similar setup, with the additional constraint that the position of the velocity discontinuity varies with the y -coordinate, as shown in Fig. (52). The domain is now $128 \times 96 \times 16\text{km}$ and the boundary conditions are implemented as follows:

```

subsection Boundary velocity model
    subsection Function
        set Variable names      = x,y,z
        set Function constants = cm=0.01, year=1
        set Function expression = if (x<56e3 && y<=48e3 | x<72e3 && y>48e3,-1*cm/year,1*cm/year);0;0
    end

```

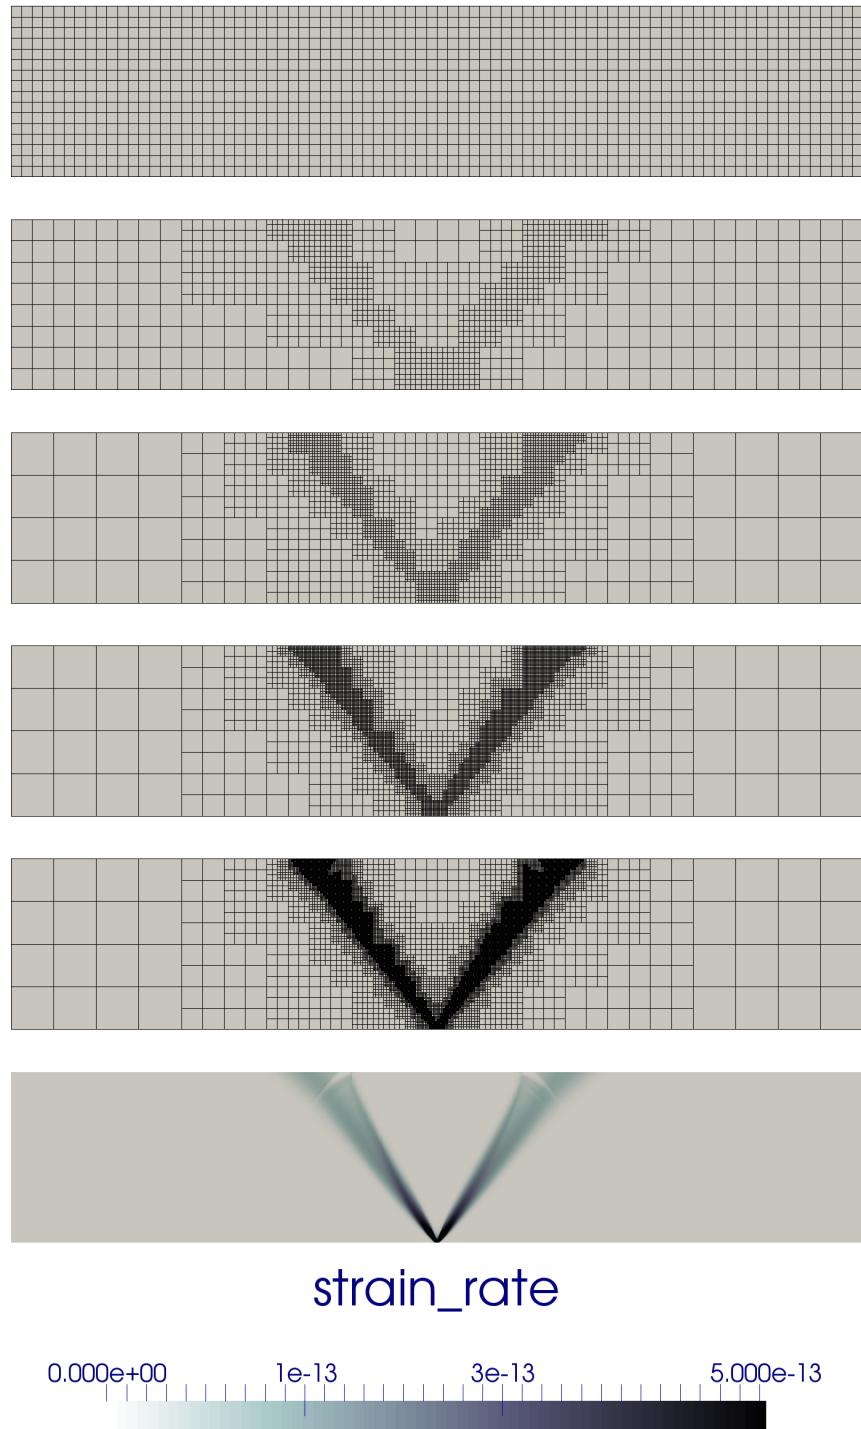


Figure 50: *Mesh evolution during the first timestep (refinement is based on strain rate).*

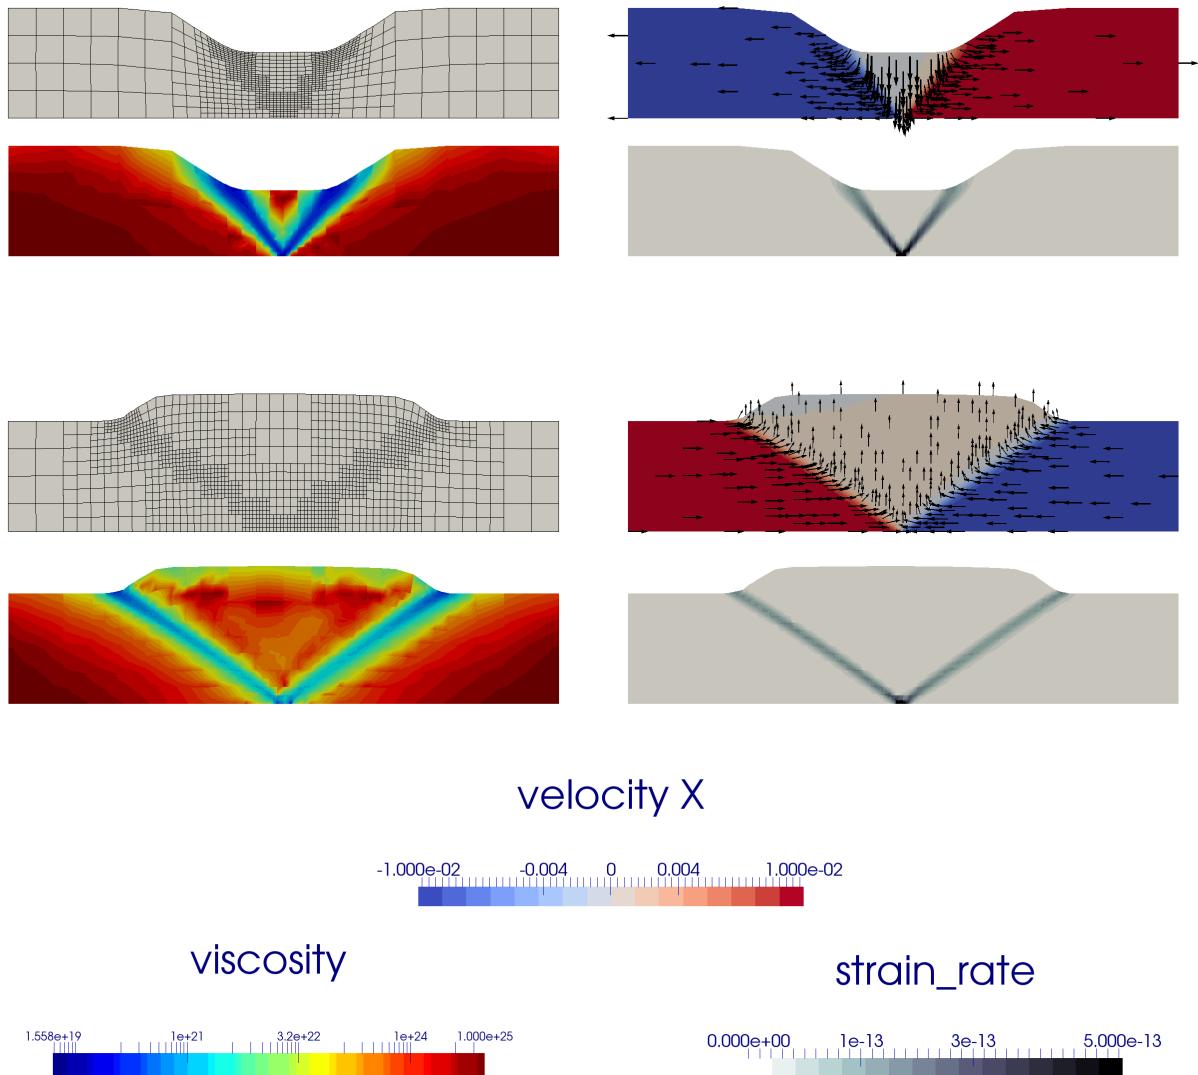


Figure 51: Finite element mesh, velocity, viscosity and strain rate fields in the case of extensional boundary conditions (top) and compressive boundary conditions (bottom) at $t=500\text{kyr}$.

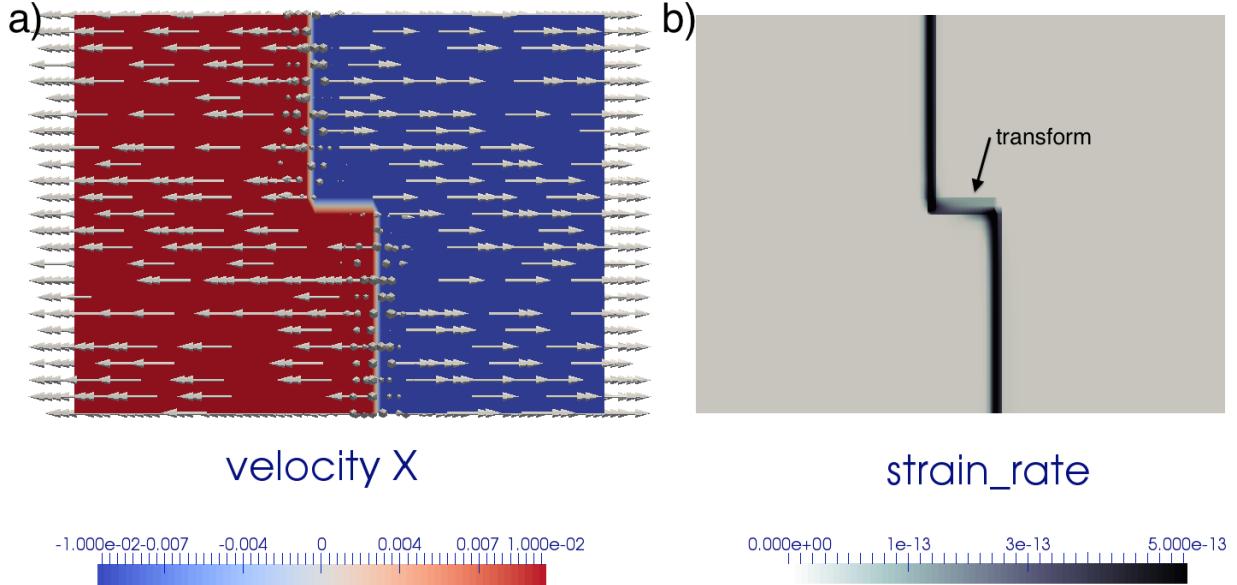


Figure 52: Basal velocity boundary conditions and corresponding strain rate field for the 3D model.

end

The presence of an offset between the two velocity discontinuity zones leads to a transform fault which connects them.

The Finite Element mesh, the velocity, viscosity and strain rate fields are shown in Fig. (53) at the end of the first time steps. The reader is encouraged to run this setup in time to look at how the two grabens interact as a function of their initial offset [AHT11, AHT12, AHFT13].

5.3.9 Continental extension

This section was contributed by John Naliboff

In the crustal deformation examples above, the viscosity depends solely on the Drucker Prager yield criterion defined by the cohesion and internal friction angle. While this approximation works reasonably well for the uppermost crust, deeper portions of the lithosphere may undergo either brittle or viscous deformation, with the latter depending on a combination of composition, temperature, pressure and strain-rate. In effect, a combination of the Drucker-Prager and Diffusion dislocation material models is required. The visco-plastic material model is designed to take into account both brittle (plastic) and non-linear viscous deformation, thus providing a template for modeling complex lithospheric processes. Such a material model can be used in ASPECT using the following set of input parameters:

```
subsection Material model
  set Model name = visco plastic
  subsection Visco Plastic
```

This cookbook provides one such example where the continental lithosphere undergoes extension. Notably, the model design follows that of numerous previously published continental extension studies [HB11, BHPeGeS14, NB15, and references therein].

Continental Extension The 2D Cartesian model spans 400 (x) by 100 (y) km and has a finite element grid with uniform 2 km spacing. Unlike the crustal deformation cookbook (see Section 5.3.8, the mesh is not refined with time).

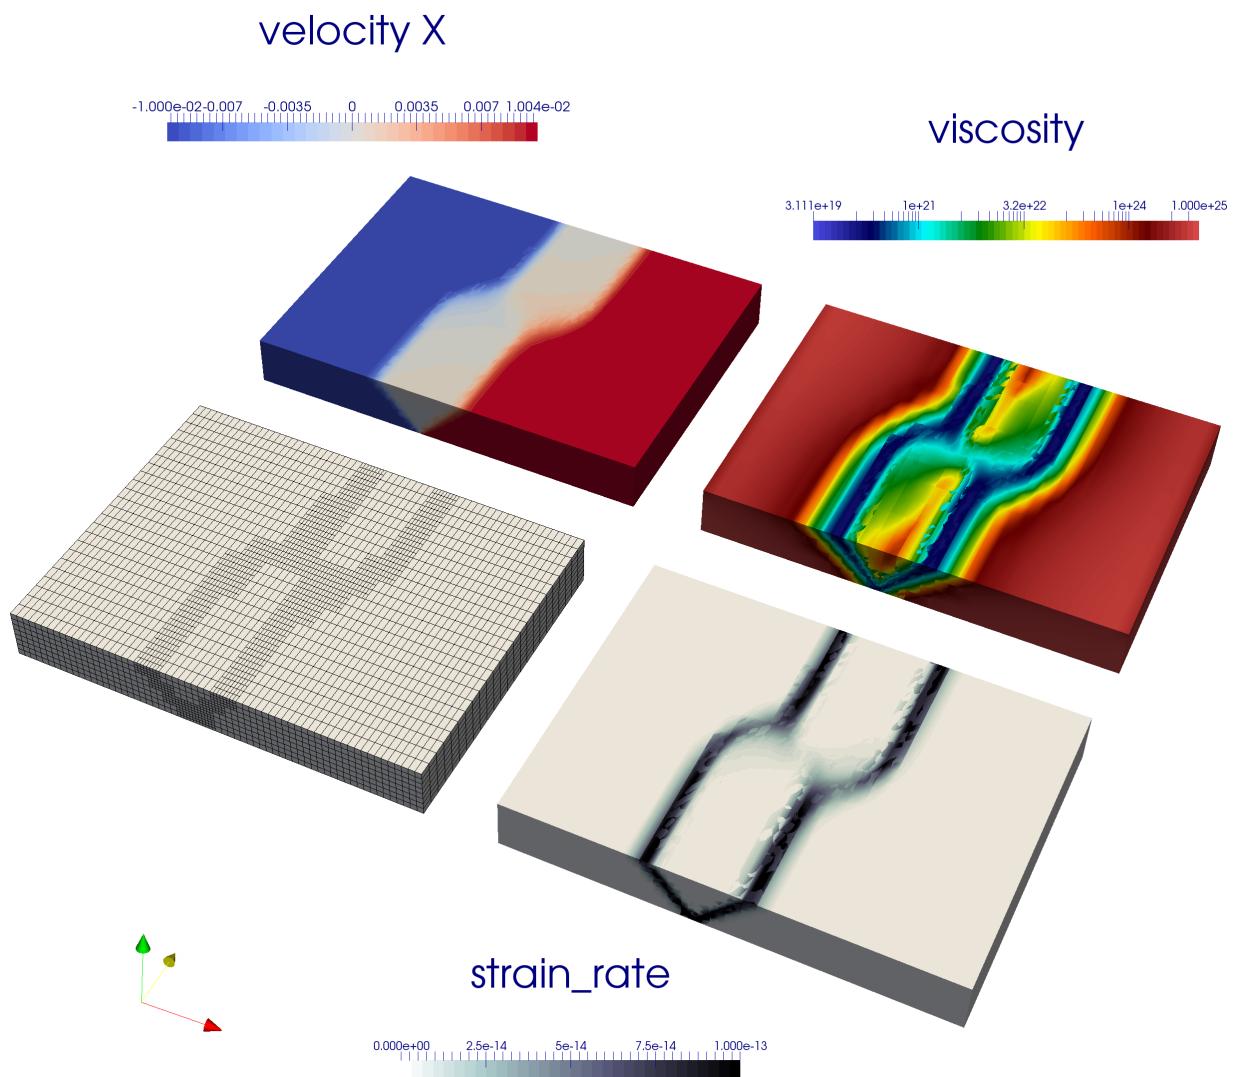


Figure 53: *Finite element mesh, velocity, viscosity and strain rate fields at the end of the first time step after one level of strain rate-based adaptive mesh refinement.*

```

subsection Geometry model
  set Model name = box
  subsection Box
    set X repetitions = 200
    set Y repetitions = 50
    set X extent      = 400e3
    set Y extent      = 100e3
  end
end

subsection Mesh refinement
  set Initial adaptive refinement      = 0
  set Initial global refinement       = 0
  set Time steps between mesh refinement = 0
end

```

Similar to the crustal deformation examples above, this model contains a free surface. Deformation is driven by constant horizontal (x -component) velocities (0.25 cm/yr) on the side boundaries (y -velocity component unconstrained), while the bottom boundary has vertical inflow to balance the lateral outflow. The top, and bottom boundaries have fixed temperatures, while the sides are insulating. The bottom boundary is also assigned a fixed composition, while the top and sides are unconstrained.

```

subsection Boundary composition model
  set Fixed composition boundary indicators = bottom
end

subsection Free surface
  set Free surface boundary indicators      = top
end

subsection Boundary velocity model
  set Prescribed velocity boundary indicators = left x: function, right x:function, bottom y:
  ↪ function
  subsection Function
    set Variable names      = x,y
    set Function constants  = cm=0.01, year=1
    set Function expression = if (x<200e3 , -0.25*cm/year, 0.25*cm/year) ; 0.125*cm/year;
  end
end

subsection Boundary temperature model
  set Fixed temperature boundary indicators = bottom, top
  set List of model names = box
  subsection Box
    set Bottom temperature = 1573
    set Top temperature   = 273
  end
end

```

Sections of the lithosphere with distinct properties are represented by compositional fields for the upper crust (20 km thick), lower crust (10 km thick) and mantle lithosphere (70 km thick). A mechanically weak seed within the mantle lithosphere helps localize deformation. Material (viscous flow law parameters, cohesion, internal friction angle) and thermodynamic properties for each compositional field are based largely on previous numerical studies. Dislocation creep viscous flow parameters are taken from published deformation experiments for wet quartzite [RB04], wet anorthite [RGWD06] and dry olivine [HK04].

```

subsection Compositional fields
    set Number of fields = 4
    set Names of fields = upper, lower, mantle, seed
end

subsection Initial composition model
    set Model name = function
    subsection Function
        set Variable names      = x,y
        set Function expression = if(y>=80.e3, 1, 0); \
                                if(y<80.e3 && y>=70.e3, 1, 0); \
                                if(y<70.e3 && y>-100.e3,1, 0); \
                                if(y<68.e3 && y>60.e3 && x>=198.e3 && x<=202.e3 , 1, 0);
    end
end

```

The initial thermal structure, radiogenic heating model and associated thermal properties are consistent with the prescribed thermal boundary conditions and produce a geotherm characteristic of the continental lithosphere. The equations defining the initial geotherm [Cha86] follow the form

$$T(z) = T_T + \frac{q_T}{k}z - \frac{Az^2}{2k} \quad (48)$$

where T is temperature, z is depth, T_T is the temperature at the layer surface (top), q_T is surface heat flux, k is thermal conductivity, and A is radiogenic heat production.

For a layer thickness Δz , the basal temperature (T_B) and heat flux (q_B) are

$$T_B = T_T + \frac{q_T}{k}\Delta z - \frac{A\Delta z^2}{2k}, \quad (49)$$

$$q_B = q_T - A\Delta z. \quad (50)$$

In this example, specifying the top (273 K) and bottom temperature (1573 K), thermal conductivity of each layer and radiogenic heat production in each layer provides enough constraints to successively solve for the temperature and heat flux at the top of the lower crust and mantle.

As noted above, the mechanically weak seed placed within the mantle localizes the majority of deformation onto two conjugate shear bands that propagate from the surface of the seed to the free surface. After 5 million years of extension background ‘stretching’ is clearly visible in the strain-rate field, but deformation is still largely focused within the set of conjugate shear bands originating at the weak seed (Fig. 54). As expected, crustal thickness and surface topography patterns reveal a relatively symmetric horst and graben structure, which arises from displacements along the shear bands (Fig. 55). While deformation along the two major shear bands dominates at this early stage of extension, additional shear bands often develop within the horst-graben system leading to small inter-graben topographic variations. This pattern is illustrated in a model with double the numerical resolution (initial 1 km grid spacing) after 10 million years of extension (Fig. 56).

With further extension for millions of years, significant crustal thinning and surface topography development should occur in response to displacement along the conjugate shear bands. However, given that the model only extends to 100 km depth, the simulation will not produce a realistic representation of continental breakup due to the lack of an upwelling asthenosphere layer. Indeed, numerical studies that examine continental breakup, rather than just the initial stages of continental extension, include an asthenospheric layer or modified basal boundary conditions (e.g. Winkler boundary condition [BHPeGeS14, for example]) as temperature variations associated with lithospheric thinning exert a first-order influence on the deformation patterns. As noted below, numerous additional parameters may also affect the temporal evolution of deformation patterns.

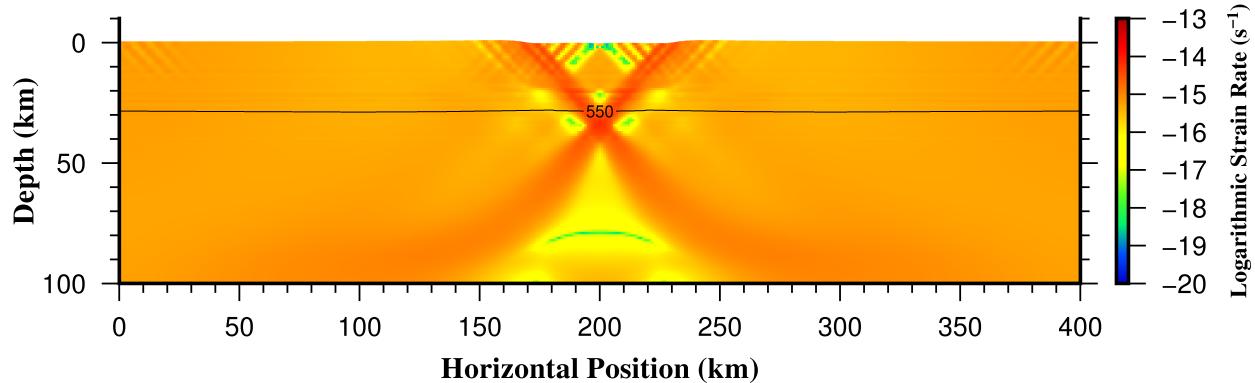


Figure 54: *Strain rate (s^{-1}) after $5e6$ years of extension. The black line marks the $550\text{ }^{\circ}\text{C}$ isotherm.*

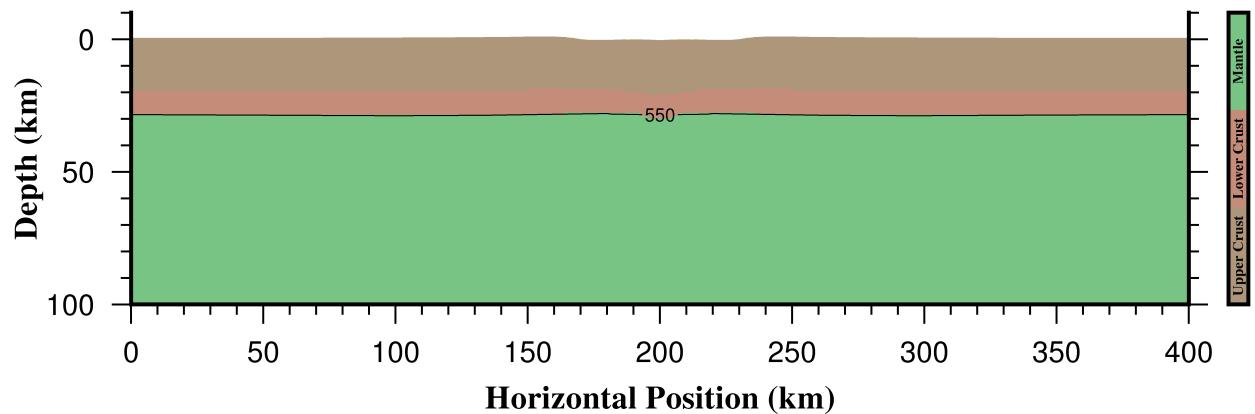


Figure 55: *Compositional field after $5e6$ years of extension. The black line marks the $550\text{ }^{\circ}\text{C}$ isotherm.*

Note: It is important to consider that the non-linearity of visco-plastic rheologies and mesh-dependence of brittle shear bands make lithospheric deformation models highly sensitive to a large number of parameters. In order to ensure the conclusions drawn from a series of numerical experiments are robust, one should complete a sensitivity test for a large range of parameters including grid resolution, model geometry, boundary conditions, initial composition and temperature conditions, material properties, composition discretization, CFL number and solver settings. If you are new to modeling lithospheric processes, a reasonable starting point is to try and reproduce results from a relevant previous study and then perform a sensitivity test for the parameters listed above. While highly time consuming, completing this procedure will prove invaluable when you design and assess the results of your own numerical study.

5.3.10 Inner core convection

This section was contributed by Juliane Dannberg, and the model setup was inspired by discussions with John Rudge. Additional materials and comments by Mathilde Kervazo and Marine Lasbleis.

This is an example of convection in the inner core of the Earth. The model is based on a spherical geometry, with a single material. Three main particularities are constitutive of this inner core dynamics modeling: it consists of a self-gravitating sphere where the gravity decreases linearly from the boundary to zero at the center of the inner core; the boundary conditions combine normal stress and normal velocity, and

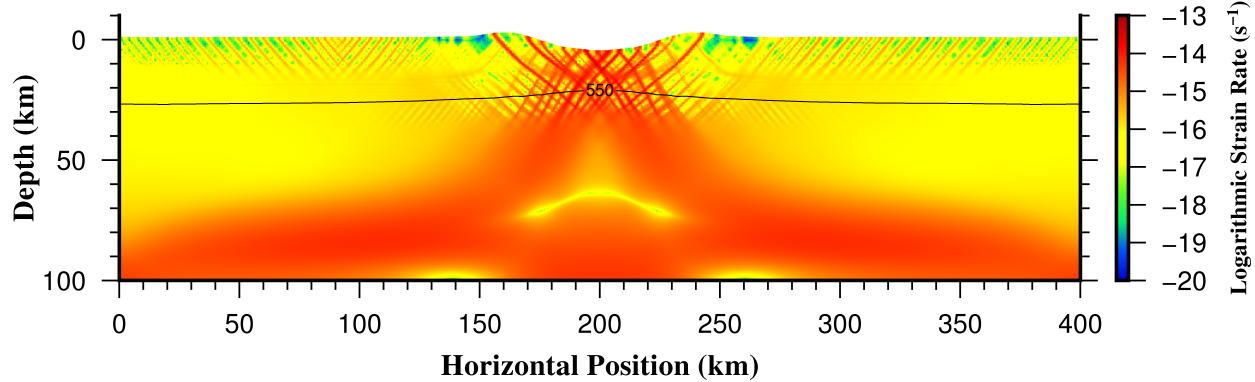


Figure 56: Strain rate(s^{-1}) after $10e6$ years of extension. The black line marks the $550\text{ }^{\circ}\text{C}$ isotherm. The numerical resolution (1 km) is double that of the previous model.

take into account the rate of phase change (melting/freezing) at the inner-outer core boundary; the material has a temperature dependent density that makes the density profile unstably stratified as temperature increases towards the center of the core. Note that we do not actually compute self-gravitation, but instead define a linear gravity profile. Since the density variations are very small, this is a good approximation.

The setup is analogous to the models described in [DAC13], and all material properties are chosen in a way so that the equations are non-dimensional.

The required heating model and changes to the material model are implemented in a shared library ([cookbooks/inner_core_convection/inner_core_convection.cc](#)).

In the non-dimensional form of the equations derived by [DAC13], we solve for the potential temperature $T = \tilde{T} - T_{is}$ (\tilde{T} is the temperature field, T_{is} the isentropic – also called adiabatic – temperature). This allows to solve the temperature field with simple boundary conditions ($T = 0$), even if the temperature of the inner core boundary evolves with time, defined as the intersection between the isentrope and the liquidus of the material in the outer core. The equations for inner core convection in the approximation of no growth (equation 59 for the potential temperature) are

$$\nabla \cdot \sigma = -RaTg, \quad (51)$$

$$\nabla \cdot \mathbf{u} = 0, \quad (52)$$

$$\left(\frac{\partial T}{\partial t} + \mathbf{u} \cdot \nabla T \right) - \nabla^2 T = H, \quad (53)$$

where Ra is the Rayleigh number and H is the 'source term', constructed when removing the adiabatic temperature from the temperature field to obtain the potential temperature T . H describes the time-evolution of the adiabatic temperature over time, due to secular cooling of the outer core. In spherical geometry, $H = 6$.

Mechanical boundary. The mechanical boundary conditions for the inner core are tangential stress-free and continuity of the normal stress at the inner-outer core boundary. For the non-dimensional equations, that means that we define a "phase change number" \mathcal{P} (see [DAC13]) so that the normal stress at the boundary is $-\mathcal{P}u_r$ with the radial velocity u_r . This number characterizes the resistance to phase change at the boundary, with $\mathcal{P} \rightarrow \infty$ corresponding to infinitely slow melting/freezing (or a free slip boundary), and $\mathcal{P} \rightarrow 0$ corresponding to instantaneous melting/freezing (or a zero normal stress, corresponding to an open boundary).

In the weak form, this results in boundary conditions of the form of a surface integral:

$$\int_S \mathcal{P}(\mathbf{u} \cdot \mathbf{n})(\mathbf{v} \cdot \mathbf{n})dS,$$

with the normal vector \mathbf{n} .

This phase change term is added to the matrix in the `cookbooks/inner_core_convection/inner_core_assembly.cc` plugin by using a signal (as described in Section 6.5). The signal connects the function `set_assemblers_phase_boundary`, which is only called once at the beginning of the model run. It creates the new assembler `PhaseBoundaryAssembler` for the boundary faces of the Stokes system and adds it to the list of assemblers executed in every time step. The assembler contains the function `phase_change_boundary_conditions` that loops over all faces at the model boundary, queries the value of \mathcal{P} from the material model, and adds the surface integral given above to the matrix:

```
#include <aspect/simulator_access.h>
#include <aspect/global.h>
#include <aspect/simulator.h>
#include <aspect/simulator/assemblers/interface.h>

#include <deal.II/base/quadrature_lib.h>
#include <deal.II/fe/fe_values.h>

#include "inner_core_convection.cc"

namespace aspect
{
    /**
     * A new assembler class that implements boundary conditions for the
     * normal stress and the normal velocity that take into account the
     * rate of phase change (melting/freezing) at the inner-outer core
     * boundary. The model is based on Deguen, Alboussiere, and Cardin
     * (2013), Thermal convection in Earth's inner core with phase change
     * at its boundary. GJI, 194, 1310-133.
     *
     * The mechanical boundary conditions for the inner core are
     * tangential stress-free and continuity of the normal stress at the
     * inner-outer core boundary. For the non-dimensional equations, that
     * means that we can define a 'phase change number'  $\mathcal{P}$  so
     * that the normal stress at the boundary is  $-\mathcal{P} u_r$  with
     * the radial velocity  $u_r$ . This number characterizes the resistance
     * to phase change at the boundary, with  $\mathcal{P} \rightarrow \infty$ 
     * corresponding to infinitely slow melting/freezing (free slip
     * boundary), and  $\mathcal{P} \rightarrow 0$  corresponding to
     * instantaneous melting/freezing (zero normal stress, open boundary).
     *
     * In the weak form, this results in boundary conditions of the form
     * of a surface integral:
     * 
$$\int_S \mathcal{P} (\mathbf{u} \cdot \mathbf{n}) (\mathbf{v} \cdot \mathbf{n}) dS,$$

     * with the normal vector  $\mathbf{n}$ .
     *
     * The function value of  $\mathcal{P}$  is taken from the inner core
     * material model.
     */
    template <int dim>
    class PhaseBoundaryAssembler :
        public aspect::Assemblers::Interface<dim>,
        public SimulatorAccess<dim>
    {
        public:
            virtual
            void
            execute (internal::Assembly::Scratch::ScratchBase<dim> &scratch_base,
                     internal::Assembly::CopyData::CopyDataBase<dim> &data_base) const
            {
                internal::Assembly::Scratch::StokesSystem<dim> &scratch =
                    dynamic_cast<internal::Assembly::Scratch::

```

```

    ↳ StokesSystem<dim>& > (scratch_base);
internal::Assembly::CopyData::StokesSystem<dim> &data = dynamic_cast<internal::Assembly::CopyData::
    ↳ StokesSystem<dim>& > (data_base);

const Introspection<dim> &introspection = this->introspection();
const FiniteElement<dim> &fe = this->get_fe();
const unsigned int stokes_dofs_per_cell = data.local_dof_indices.size();
const unsigned int n_q_points = scratch.face_finite_element_values.n_quadrature_points;

//assemble force terms for the matrix for all boundary faces
if (scratch.cell->face(scratch.face_number)->at_boundary())
{
    scratch.face_finite_element_values.reinit (scratch.cell, scratch.face_number);

    for (unsigned int q=0; q<n_q_points; ++q)
    {
        const double P = dynamic_cast<const MaterialModel::InnerCore<dim>&>
            (this->get_material_model()).resistance_to_phase_change
            .value(scratch.material_model_inputs.position[q]);

        for (unsigned int i = 0, i_stokes = 0; i_stokes < stokes_dofs_per_cell; /*increment at end of
    ↳ loop*/)
        {
            if (introspection.is_stokes_component(fe.system_to_component_index(i).first))
            {
                scratch.phi_u[i_stokes] = scratch.face_finite_element_values[introspection
                    .extractors.velocities].
    ↳ value(i, q);
                ++i_stokes;
            }
            ++i;
        }

        const Tensor<1,dim> normal_vector = scratch.face_finite_element_values.normal_vector(q);
        const double JxW = scratch.face_finite_element_values.JxW(q);

        // boundary term: P*u*n*v*n*JxW(q)
        for (unsigned int i=0; i<stokes_dofs_per_cell; ++i)
            for (unsigned int j=0; j<stokes_dofs_per_cell; ++j)
                data.local_matrix(i,j) += P *
                    scratch.phi_u[i] *
                    normal_vector *
                    scratch.phi_u[j] *
                    normal_vector *
                    JxW;
    }
}
};

template <int dim>
void set_assemblers_phase_boundary(const SimulatorAccess<dim> &simulator_access,
                                    Assemblers::Manager<dim> &assemblers)
{
    AssertThrow (dynamic_cast<const MaterialModel::InnerCore<dim>>*
        (&simulator_access.get_material_model()) != 0,
        ExcMessage ("The phase boundary assembler can only be used with the "
                    "material model 'inner core material'!"));

    PhaseBoundaryAssembler<dim> *phase_boundaryAssembler = new PhaseBoundaryAssembler<dim>();
    assemblers.stokes_system_on_boundary_face.push_back (std_cxx11::unique_ptr<PhaseBoundaryAssembler<dim> > (
    ↳ phase_boundaryAssembler));
}
}

```

```

template <int dim>
void signal_connector (aspect::SimulatorSignals<dim> &signals)
{
  signals.set_assemblers.connect (&aspect::set_assemblers_phase_boundary<dim>);
}

ASPECT_REGISTER_SIGNALS_CONNECTOR(signal_connector<2>,
                                  signal_connector<3>)

```

Instructions for how to compile and run models with a shared library are given in Section 5.4.1.

Governing parameters. Analyzing Equations (51)–(53), two parameters determine the dynamics of convection in the inner core: the Rayleigh number Ra and the phase change number \mathcal{P} . Three main areas can be distinguished: the stable area, the plume convection area and the translation mode of convection area (Figure 57). For low Rayleigh numbers (below the critical value Ra_c), there is no convection and thermal diffusion dominates the heat transport. However, if the inner core is convectively unstable ($Ra > Ra_c$), the convection regime depends mostly on \mathcal{P} . For low \mathcal{P} (< 29), the convective translation mode dominates, where material freezes at one side of the inner core and melts at the other side, so that the velocity field is uniform, pointing from the freezing to the melting side. Otherwise, at high \mathcal{P} (> 29), convection takes the usual form of thermal convection with shear free boundary and no phase change, that is the one-cell axisymmetric mode at the onset, and chaotic plume convection for larger Rayleigh number. In this case, melting and solidification at the ICB have only a small dynamic effect. At intermediate values of \mathcal{P} , the first unstable mode is a linear combination of the high- \mathcal{P} convection mode and of the small- \mathcal{P} translation mode.

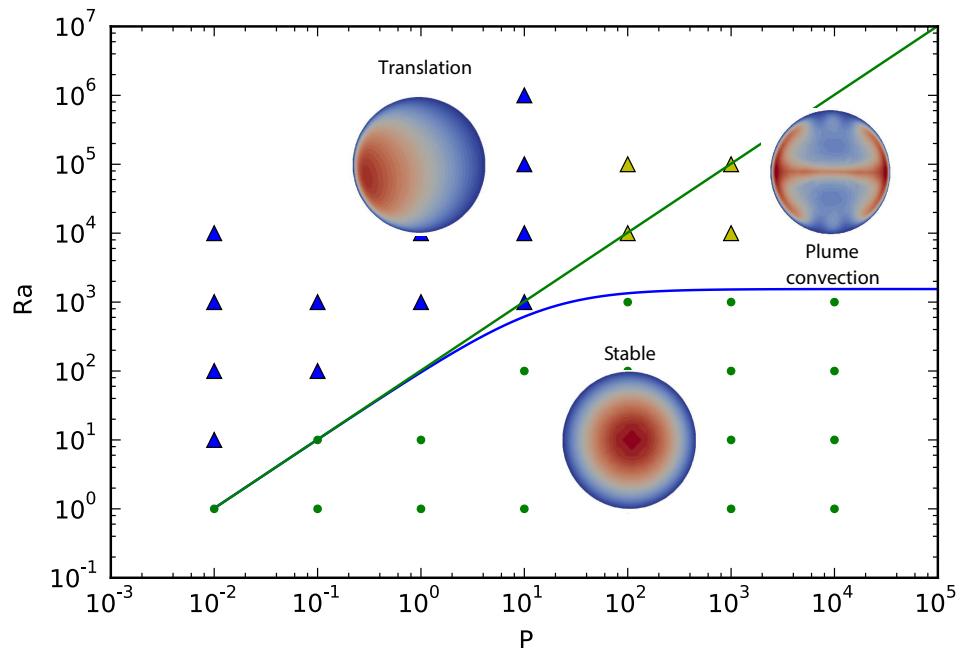


Figure 57: Stability diagram for convection in a sphere with phase change at its outer boundary. The stability curves for the first unstable mode ($l=1$) and the translation are obtained from [DAC13]. Each dot (no convection) and triangle (blue: translation, yellow: plume convection) is one model run done with ASPECT. The highest the Ra and P are, the more refinement is required (see text).

Changing the values of Ra and \mathcal{P} in the input file allows switching between the different regimes. The

Rayleigh number can be changed by adjusting the magnitude of the gravity:

```
# The gravity has its maximum value at the boundary of inner and
# outer core, and decreases approximately linearly to zero towards
# the center of the core.
# The Rayleigh number used in the model is given by the magnitude
# of the gravity at the inner core/outer core boundary.
subsection Gravity model
    set Model name = radial linear

    subsection Radial linear
        set Magnitude at bottom = 0.0
        set Magnitude at bottom = 0.0
        set Magnitude at surface = 2      # <-- Ra
    end
end
```

The phase change number is implemented as part of the material model, and as a function that can depend on the spatial coordinates and/or on time:

```
subsection Material model
    set Model name = inner core material

    # The 'inner core material' model also contains a function that
    # represents the resistance to melting/freezing at the inner core
    # boundary.
    # For P-->infinity, the boundary is a free slip boundary, and for
    # P-->0, the boundary is an open boundary (with zero normal stress).
    subsection Inner core
        subsection Phase change resistance function
            set Variable names      = x,y,z
            set Function expression = 1e-2      # <-- P
        end
    end
end
```

Figure 58 shows examples of the three regimes with $Ra = 3000, \mathcal{P} = 1000$ (plume convection), $Ra = 10^5, \mathcal{P} = 0.01$ (translation), $Ra = 10, \mathcal{P} = 30$ (no convection).

Mesh refinement. The temperature is set to 0 at the outer boundary and a large temperature gradient can develop at the boundary layer, especially for the translation regime. The adaptive mesh refinement allows it to resolve this layer at the inner core boundary. Another solution is to apply a specific initial refinement, based on the boundary layer thickness scaling law $\delta \propto Ra^{-0.236}$, and to refine specifically the uppermost part of the inner core.

In order to have a mesh that is much finer at the outer boundary than in the center of the domain, this expression for the mesh refinement subsection can be used in the input file:

```
subsection Mesh refinement
    set Initial global refinement      = 4  #this may be more expensive, and should be run on a
        ↪ cluster.
    set Initial adaptive refinement   = 1
    set Strategy                      = minimum refinement function
    set Time steps between mesh refinement = 0

    subsection Minimum refinement function
        set Variable names = depth, phi, theta
        set Function expression = if(depth>0.1,if(depth>0.2,2,5),6)
```

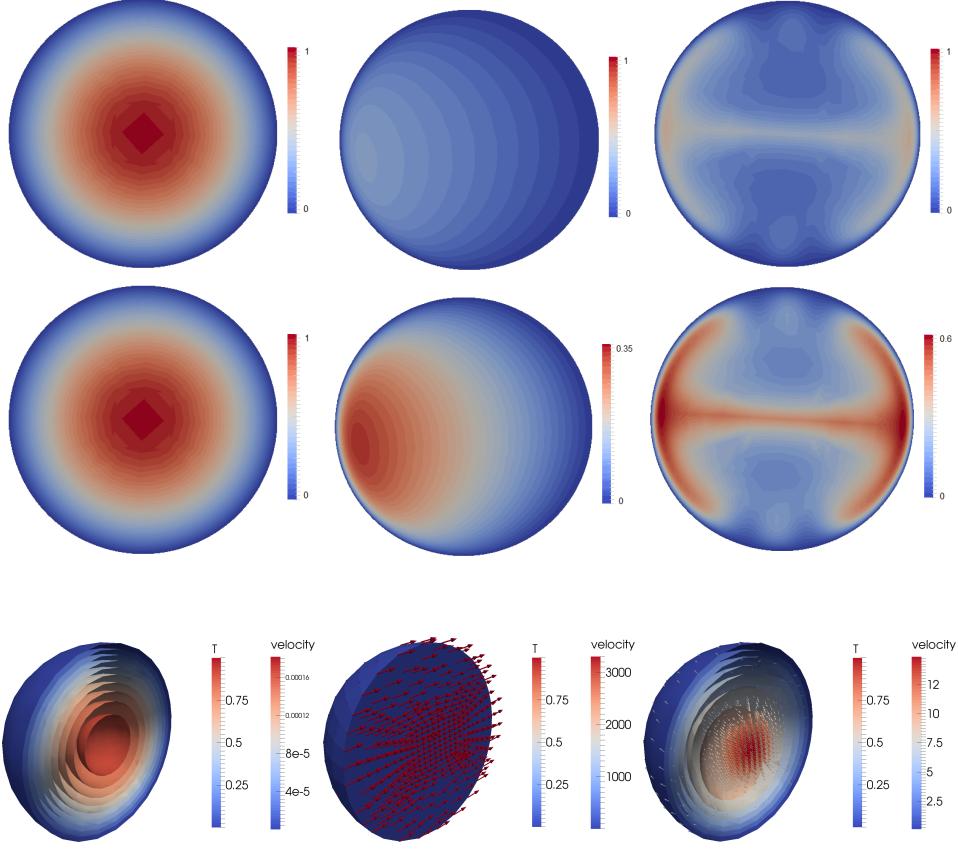


Figure 58: Convection regimes in the inner core for different values of Ra and \mathcal{P} . From left to right: no convection, translation, plume convection; the 2D slices at the top are with the default temperature scale for all panels, while in the second row an adaptive scale is used. The bottom row features slightly different model parameters (that are still in the same regime as the models in the respective panels above) and also shows the velocity as arrows.

```
    end
end
```

Scaling laws. In addition, [DAC13] give scaling laws for the velocities in each regime derived from linear stability analysis of perfect translation, and show how numerical results compare to them. In the regimes of low \mathcal{P} , translation will start at a critical ratio of Rayleigh number and phase change number $\frac{Ra}{\mathcal{P}} = \frac{175}{2}$ with steady-state translation velocities being zero below this threshold and tending to $v_0 = \frac{175}{2} \sqrt{\frac{6}{5} \frac{Ra}{\mathcal{P}}}$ going towards higher values of $\frac{Ra}{\mathcal{P}}$. In the same way, translation velocities will decrease from v_0 with increasing \mathcal{P} , with translation transitioning to plume convection at $\mathcal{P} \sim 29$. Both trends are shown in Figure 59 and can be compared to Figure 8 and 9 in [DAC13].

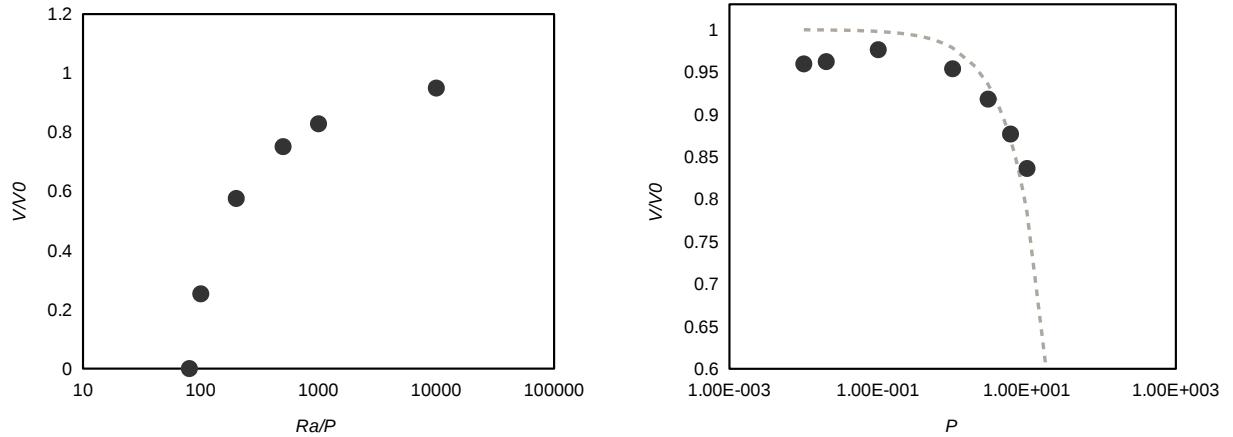


Figure 59: Translation rate (approximated by the average of the velocity component in the direction of translation), normalized to the low \mathcal{P} limit estimate given in [Dac13], as a function of $\frac{Ra}{\mathcal{P}}$ for $\mathcal{P} = 10^{-2}$ (left) and as a function of \mathcal{P} for $Ra = 10^5$ (right). The dashed gray line gives the translation velocity predicted in the limit of low \mathcal{P} . Disagreement for larger values of \mathcal{P} indicates that higher order terms (not included in the low \mathcal{P} approximation) become important. Additionally, differences between the analytical and numerical model might be the result of limited resolution (only 12 elements in radial direction).

5.3.11 Melt migration in a 2D mantle convection model

This section was contributed by Juliane Dannberg and is based on a section in [DH16] by Juliane Dannberg and Timo Heister.

The following cookbook will explain how to use ASPECT’s implementation of coupled magma/mantle dynamics (see Section 2.13) to set up a model of mantle convection that also includes melting and freezing of mantle rock, and the transport of melt according to the two-phase flow equations. The model setup is described in detail in [DH16], which can be found [here](#), and in the following we will go over a slightly simplified version in lower resolution. We will start by looking at a global mantle convection without melt migration, and will then discuss how the input file has to be modified in order to add melt transport. A movie that compares the evolution of the temperature field and the amount of melt present in both models in higher resolution can be found [online](#).

The model setup is a 2D box with dimensions of 2900×8700 km, and it is heated from the bottom and cooled from the top. A full description can be found in Section 4.7 “Influence of melt migration on a global-scale convection model” in [DH16]. In the first model we will look at, melting and freezing is only included passively: We use the `melt fraction` visualization postprocessor to compute how much melt is present for a given temperature and pressure at every given point in time and space in our model, but the presence of melt does not influence material properties like density or viscosity, and melt does not move relative to the solid. This also means that because melt is not extracted, the bulk composition of the material always stays the same, and melt only freezes again once advection or conduction causes the temperature of the solid rock to be below the solidus. The following input file (which can be found in `cookbooks/global_no_melt.prm`) contains a detailed description of the different options required to set up such a model:

```
# Model setup for mantle convection in a 2D box without melting.
# This file is used as a starting point for a cookbook that
# explains how to add melting and melt transport to a mantle
# convection simulation.

set Dimension          = 2
set Adiabatic surface temperature = 1600
```

```

set Maximum time step          = 1e6
set Output directory           = no_melt
set Use years in output instead of seconds = true

# The end time of the simulation. Because we want to see how upwellings
# and downwellings evolve over time, and if differences develop between
# the model with and without melt migration, we set the end time to 140 Ma.
set End time                  = 1.4e8

# We choose a stricter than default linear Stokes solver tolerance,
# to be consistent with the global_melt cookbook.
subsection Solver parameters
    subsection Stokes solver parameters
        set Linear solver tolerance = 1e-8
        set Number of cheap Stokes solver steps = 100
    end
end

# We prescribe free-slip boundary conditions on all
# sides.
subsection Boundary velocity model
    set Tangential velocity boundary indicators = left, right, top, bottom
end

# We also choose relatively large values for the stabilization parameters:
# The model resolution is very coarse (in order for this model to run in a
# short time), so we want to make sure that no temperature over- and
# undershoots will develop. In a model with melting this would be
# particularly problematic, as large amounts of melt could be generated
# by temperature spikes, and we want to be consistent between the model
# with and without melt transport.
subsection Discretization
    subsection Stabilization parameters
        set beta   = 0.5
        set cR     = 1
    end
end

##### Initial conditions #####
# We choose an adiabatic temperature profile as initial condition,
# with conductive temperature profiles in the top and bottom boundary
# layers, which were computed using a half space cooling model.
# The cold top boundary layer corresponds to an age of 300 Ma,
# and the hot top boundary layer corresponds to an age of 500 Ma.
# A small temperature perturbation is added at the bottom of the
# domain. To make the model asymmetric, we place it in a distance of
# x = 2900 km from the left boundary.
# Temperatures from both initial temperature models we specify are
# added (by default).
subsection Initial temperature model
    set List of model names = adiabatic, function
    subsection Adiabatic
        set Age bottom boundary layer = 5e8
        set Age top boundary layer   = 3e8

```

```

subsection Function
  set Function expression      = 0;0
end
end

subsection Function
  set Function constants      = r=350000, amplitude=50
  set Function expression      = if((x-2900000)*(x-2900000)+y*y<r,amplitude,0)
end
end

##### Boundary conditions #####
# As boundary conditions for the temperature, we just use the
# initial conditions again. This temperature is applied as a prescribed
# temperature at the top and bottom boundary, as specified above.
subsection Boundary temperature model
  set Fixed temperature boundary indicators = top, bottom
  set List of model names = initial temperature

subsection Initial temperature
  set Minimal temperature = 293
  set Maximal temperature = 3700
end
end

##### Model geometry #####
# The model geometry is a box with an aspect ratio of 3,
# extending to the base of the mantle in vertical direction.
subsection Geometry model
  set Model name = box

subsection Box
  set X extent = 8700000
  set Y extent = 2900000
  set X repetitions = 3
end

end

##### Gravity and material properties #####
# The model has a constant gravity.
subsection Gravity model
  set Model name = vertical

subsection Vertical
  set Magnitude = 9.81
end
end

# We use the melt global material model, which is one of the
# material models that works with melt transport, as it also
# specifies the material properties needed for melt migration,
# such as the permeability, the melt density and the melt

```

```

# viscosity.
# It also works without melt transport, and in this case these
# properties are not used, so we do not have to specify them
# here.
subsection Material model

    set Model name = melt global
    subsection Melt global
        set Thermal conductivity          = 4.7
        set Reference solid density       = 3400
        set Thermal expansion coefficient = 2e-5
        set Reference shear viscosity    = 5e21
        set Thermal viscosity exponent   = 7
        set Reference temperature        = 1600
        set Solid compressibility        = 4.2e-12
    end
end

##### Mesh refinement #####
# For the model without melt migration, we do not have to use
# mesh adaptivity, because time- and length scales of material
# motion does do not vary a lot across the model, and a global
# resolution of 4 is sufficient to capture the behaviour of
# upwellings and downwellings.
subsection Mesh refinement
    set Initial adaptive refinement      = 0
    set Initial global refinement        = 4
    set Time steps between mesh refinement = 0
end

# As the model is compressible and has an adiabatic temperature profile, we include
# adiabatic heating in the list of heating models.
subsection Heating model
    set List of model names = adiabatic heating
end

##### Postprocessing #####
# In addition to the visualization output, we select a number
# of postprocessors that allow us to compute some statistics
# about the output (to see how much the model without and the
# model with melt migration differ), and in particular we use
# the "depth average" postprocessor that will allow us to plot
# depth-averaged model quantities over time.
subsection Postprocess
    set List of postprocessors = visualization, composition statistics, velocity statistics,
                                ↪ temperature statistics, melt statistics, depth average

# For the model without melt migration, we only compute the
# equilibrium melt fraction in dependence of temperature and
# pressure. This is done as a postprocessing step, by adding
# "melt fraction" to the list of visualization postprocessors.

subsection Visualization
    set List of output variables      = material properties, nonadiabatic temperature

```

```

subsection Material properties
    set List of material properties = density, viscosity, melt fraction
end

set Number of grouped files      = 0
set Output format              = vtu
set Time between graphical output = 6e5
set Interpolate output         = true
end

subsection Depth average
    set Number of zones = 12
    set Time between graphical output = 6e5
end

end

# We write a checkpoint approximately every half an hour,
# so that we are able to restart the computation from that
# point.
subsection Checkpointing
    set Time between checkpoint = 1700
end

```

When we look at visualization output of this model, we can see that over time, first upwellings, and then downwellings start to form. Both are more or less stable over time, and only change their positions slowly. As melt does not move relative to the solid, broad stable zones of melting with melt fraction of 10% or more form in areas where material is upwelling.

In the second model, melt is an active component of the model. Temperature, pressure and composition control how much of the rock melts, and as soon as that happens, melt can migrate relative to the solid rock. As material only melts partially, that means that the composition of the rock changes when it melts and melt is extracted, and we track this change in composition using a compositional field with the name **peridotite**. Positive values mark depletion (the composition of the residual host rock as more and more melt is extracted), and negative values mark enrichment (the composition of generated melt, or regions where melt freezes again). Both the fraction of melt (tracked by the compositional field with the name **porosity**) and the changes in composition influence the material properties such as density and viscosity. Moreover, there are additional material properties that describe how easily melt can move through the host rock, such as the **permeability**, or properties of the melt itself, such as the **fluid viscosity**. The following input file (a complete version of which can be found in [cookbooks/global_melt.prm](#)) details the changes we have to make from the first model to set up a model with melt migration:

```

# Cookbook for a global-scale 2D box mantle convection model
# with melt migration.
# In this file we will go through all of the steps that are
# required for adding melting and melt transport to a mantle
# convection simulation.

# For models with melt migration, there is a nonlinear coupling between
# the Stokes system, the temperature, and the advection equation for the
# porosity (several material properties, such as the viscosities and the
# permeability depend nonlinearly on the porosity; and changes in temperature
# determine how much material is melting or freezing).
# Because of that, we use a nonlinear solver scheme ('iterated Advection and Stokes')
# that iterates between all of these equations, and we have to set its

```

```

# solver tolerance and the maximum number of iterations separately from
# the linear solver parameters.
set Nonlinear solver scheme = iterated Advection and Stokes
set Max nonlinear iterations = 20
set Nonlinear solver tolerance = 1e-5

# In addition, melting and freezing normally happens on a much faster
# time scale than the flow of melt, so we want to decouple the advection
# of melt (and temperature) and the melting process itself. To do that,
# we use the operator splitting scheme, and define that for every
# advection time step, we want to do at least 10 reaction time steps.
# If these time steps would be larger than 10,000 years, we will do
# more reaction time steps (so that reaction time step size never exceeds
# 10,000 years). Here, we also specify the Stokes linear solver tolerance
# and maximum number of cheap Stokes solver steps to improve the nonlinear
# convergence behavior.
set Use operator splitting = true
subsection Solver parameters
    subsection Operator splitting parameters
        set Reaction time step = 1e4
        set Reaction time steps per advection step = 10
    end
    subsection Stokes solver parameters
        set Linear solver tolerance = 1e-8
        set Number of cheap Stokes solver steps = 100
    end
end

subsection Melt settings
    # In addition, we now also specify in the model settings that we want to
    # run a model with melt transport.
    set Include melt transport = true
end

##### Settings for melt transport #####
# In models with melt transport, we always need a compositional field with
# the name 'porosity'. Only the field with that name will be advected with
# the melt velocity, all other compositional fields will continue to work
# as before. Material model will typically query for the field with the
# name porosity to compute all melt material properties.
# In addition, the 'melt global' model also requires a field with the name
# 'peridotite'. This field is used to track how much material has been
# molten at each point of the model, so it tracks the information how the
# composition of the rock changes due to partial melting events (sometimes
# also called depletion). This is important, because usually less melt is
# generated for a given temperature and pressure if the rock has undergone
# melting before. Typically, material properties like the density are also
# different for more or less depleted material.
subsection Compositional fields
    set Number of fields = 2
    set Names of fields = porosity, peridotite
end

##### Initial conditions #####

```

```

# Now that our model uses compositional fields, we also need initial
# conditions for the composition.
# We use the function plugin to set both fields to zero at the beginning
# of the model run.
subsection Initial composition model
    set Model name = function
    subsection Function
        set Function expression = 0; 0
        set Variable names      = x,y
    end
end

##### Boundary conditions #####
# We again choose the initial composition as boundary condition
# for all compositional fields.
subsection Boundary composition model
    set List of model names = initial composition
end

# Models with melt transport also need an additional boundary condition:
# the gradient of the fluid pressure at the model boundaries. This boundary
# condition indirectly also prescribes boundary conditions for the melt velocity,
# as the melt velocity is related to the fluid pressure gradient via Darcy's law.
# If we choose the fluid pressure gradient = solid density * gravity, melt will
# flow in and out of the model (even if the solid can not flow out) according to
# the dynamic fluid pressure in the model. Conversely, if we choose the
# fluid pressure gradient = fluid density * gravity, melt will flow in or out
# with the same velocity as the solid (so for a closer boundary, no melt will
# flow in or out). This is what we choose as our boundary condition here.
subsection Boundary fluid pressure model
    set Plugin name = density
    subsection Density
        set Density formulation = fluid density
    end
end

##### Material properties #####
# In addition to the material properties for the solid rock,
# we also have to specify properties for the melt.
subsection Material model

    set Model name = melt global
    subsection Melt global

        # First we describe the parameters for the solid, in the same way
        # we did in the model without melt transport
        set Thermal conductivity          = 4.7
        set Reference solid density       = 3400
        set Thermal expansion coefficient = 2e-5
        set Reference shear viscosity     = 5e21
        set Thermal viscosity exponent   = 7
        set Reference temperature        = 1600
        set Solid compressibility         = 4.2e-12

```

```

# The melt usually has a different (lower) density than the solid.
set Reference melt density = 3000

# The permeability describes how well the pores of a porous material
# are connected (and hence how fast melt can flow through the rock).
# It is computed as the product of the reference value given here
# and the porosity cubed. This means that the lower the porosity is
# the more difficult it is for the melt to flow.
set Reference permeability = 1e-8

# The bulk viscosity describes the resistance of the rock to dilation
# and compaction. Melt can only flow into a region that had no melt
# before if the matrix of the solid rock expands, so this parameter
# also limits how fast melt can flow upwards.
# The bulk viscosity is computed as the reference value given here times
# a term that scales with one over the porosity. This means that for zero
# porosity, the rock can not dilate/compact any more, which is the same
# behaviour that we have for solid mantle convection.
set Reference bulk viscosity = 1e19

# In dependence of how much melt is present, we also weaken the shear
# viscosity: The more melt is present, the weaker the rock gets.
# This scaling is exponential, following the relation
# viscosity ~ exp(-alpha * DeltaT),
# where alpha is the parameter given here, and DeltaT is the deviation from the
# reference temperature.
set Exponential melt weakening factor = 10

# In the same way the shear viscosity is reduced with increasing temperature,
# we also prescribe the temperature-dependence of the bulk viscosity.
set Thermal bulk viscosity exponent = 7

# Analogous to the compressibility of the solid rock, we also define a
# compressibility for the melt (which is generally higher than for the solid).
# As we do not want our compressibility to depend on depth, we set the
# pressure derivative to zero.
set Melt compressibility = 1.25e-11
set Melt bulk modulus derivative = 0.0

# Finally, we prescribe the viscosity of the melt, which is used in Darcy's
# law. The lower this viscosity, the faster melt can flow.
set Reference melt viscosity = 1

# change the density contrast of depleted material (in kg/m^3)
set Depletion density change = -200.0

# How much melt has been generated and subsequently extracted from a particular
# volume of rock (how 'depleted' that volume of rock is) usually changes the
# solidus. The more the material has been molten already, the less melt will be
# generated afterwards for the same pressure and temperature conditions. We
# model this using a simplified, linear relationship, saying that to melt 100%
# of the rock the temperature has to be 200 K higher than to melt it initially.
set Depletion solidus change = 200

# We also have to determine how fast melting and freezing should happen.
# Here, we choose a time scale of 10,000 years, which is a relatively long time

```

```

# (or in other words, slow melting rate), but because this is a global model
# and the time steps are big, it should be sufficient.
set Melting time scale for operator splitting = 1e4
end
end

##### Mesh refinement #####
# For the model with melt migration, we use adaptive refinement.
# We make use of two different refinement criteria: we set a minimum of 4 global
# refinements everywhere in the model (which is the same resolution as for the
# model without melt), and we refine in regions where melt is present, to be
# precise, everywhere where the porosity is bigger than 1e-5.
# We adapt the mesh every 5 time steps.
subsection Mesh refinement
    set Coarsening fraction          = 0.05
    set Refinement fraction          = 0.8

    set Initial adaptive refinement = 2
    set Initial global refinement   = 4
    set Strategy                   = composition threshold, minimum refinement function
    set Time steps between mesh refinement = 4

# minimum of 4 global refinements
subsection Minimum refinement function
    set Coordinate system = depth
    set Function expression = 4
    set Variable names     = depth,phi
end

# refine where the porosity is bigger than 1e-5
subsection Composition threshold
    set Compositional field thresholds = 1e-5,1.0
end
end

##### Postprocessing #####
# In addition to the visualization output, we select a number
# of postprocessors that allow us to compute some statistics
# about the output (to see how much the model without and the
# model with melt migration differ), and in particular we use
# the "depth average" postprocessor that will allow us to plot
# depth-averaged model quantities over time.
subsection Postprocess

    set List of postprocessors = visualization, composition statistics, velocity statistics,
                                ↪ temperature statistics, depth average

    # For the model with melt migration, also add a visualization
    # postprocessor that computes the material properties relevant
    # to migration (permeability, viscosity of the melt, etc.).

subsection Visualization
    set List of output variables      = material properties, nonadiabatic temperature, strain rate,
                                ↪ melt material properties

```

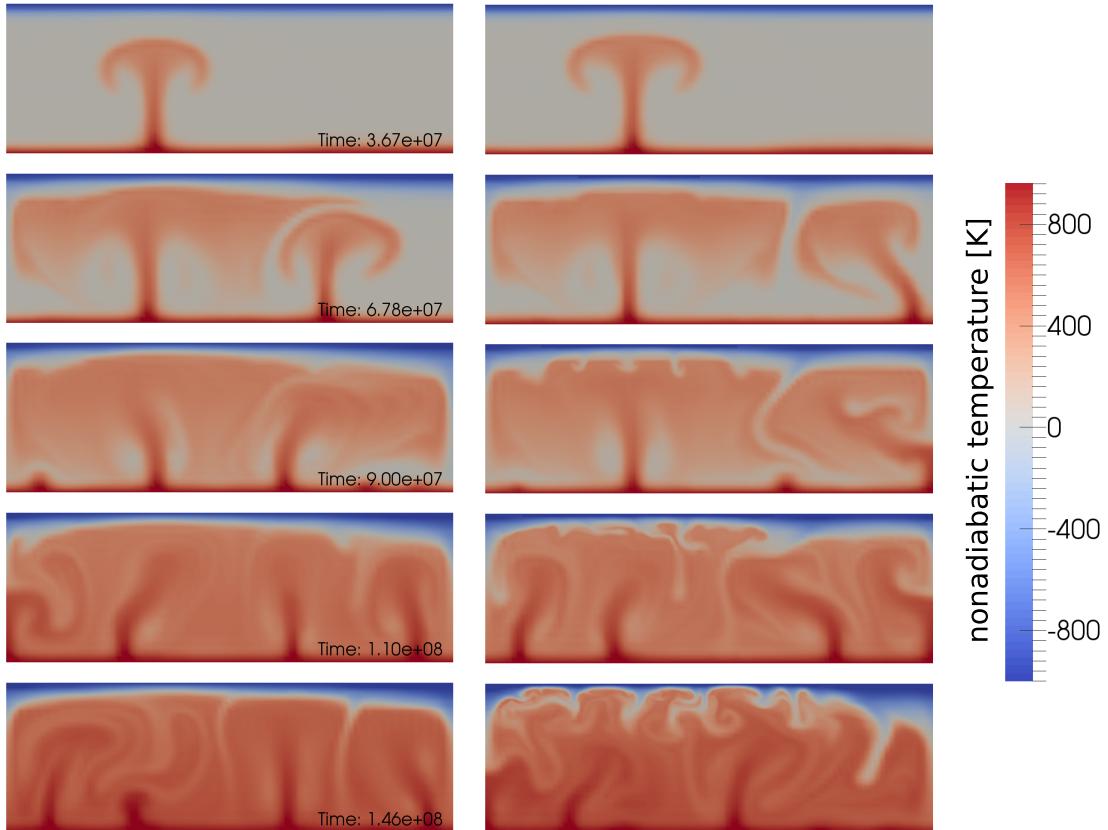


Figure 60: Evolution of the model without (left) and with (right) melt migration.

```

subsection Material properties
    set List of material properties = density, viscosity, thermal expansivity
end

subsection Melt material properties
    set List of properties = fluid density, permeability, fluid viscosity, compaction viscosity,
    ↪ p_c
end
end

```

In the first few tens of millions of years, this models evolves similarly to the model without melt migration. Upwellings rise in the same locations, and regions where material starts to melt are similar. However, once melt is formed, the model evolutions start to deviate. In the model with melt migration, melt moves upwards from the region where it is generated much faster than the flow of solid material, so that it reaches cold regions – where it freezes again – in a shorter amount of time. Because of that, the overall amount of melt is smaller in this model at any given point in time. In addition, enriched material, present in places where melt has crystallized, has a higher density than average or depleted mantle material. This means that in regions above stable upwellings, instabilities of dense, enriched material start to form, which leads to small-scale downwellings. Hence, both areas where material is partially molten and the location of the upwellings themselves have a much shorter wavelength and change much faster over time in comparison to the model without melt migration.

Figure 60 shows the time evolution of both models. A more complete comparison of the two models can be found in Section 4.7 “Influence of melt migration on a global-scale convection model” in [DH16].

5.3.12 Melt migration in a 2D mid-ocean ridge model

This section was contributed by Juliane Dannberg.

The following cookbook will explain how to set up a model of a mid-ocean ridge that uses ASPECT’s implementation of coupled magma/mantle dynamics (see Section 2.13) and melting and freezing of mantle rock. In particular, it will outline

1. how to use operator splitting to accurately compute melting and freezing of melt,
2. how to use traction boundary conditions to set up the flow field of a mid-ocean ridge,
3. useful strategies for how to refine the mesh in models with melt migration.

How to set up a model with melt migration in general is explained in the previous cookbook 5.3.11.

As the flow at mid-ocean ridges can be assumed to be roughly symmetric with respect to the ridge axis in the center, we only model one half of the ridge in a 2d Cartesian box with dimensions of 105×70 km. Solid material is flowing in from the bottom with a prescribed temperature and melting due to decompression as it rises. The model is cooled from the top so that melt freezes again as it approaches this boundary. In addition, a fixed plate velocity away from the ridge axis is prescribed at the top boundary, inducing corner flow. Material can flow out freely at the right model boundary. The model shows both how melt is focused towards the ridge axis, and how melting and freezing induces chemical heterogeneity in the mantle, generating the crust and lithosphere. A movie of the full model evolution can be found [online](#).

The input file. One important problem in models with melting and freezing (and other reactions) is that these reactions can be much faster than the time step of the model. For mid-ocean ridges, melt is generally assumed to be in equilibrium with the solid, which means that the reaction is basically instantaneous. To model these type of processes, ASPECT uses operator splitting (see also Section 5.4.15): Reactions are solved on a different time scale than advection. For this model, this means that at the beginning of each time step, all melting reactions, including their latent heat effects, are solved using several shorter sub-time steps. In the input file, we have to choose both the size of these sub-time steps and the rate (or characteristic time scale) of melting, and they have to be consistent in the sense that the operator splitting time step can not be larger than the reaction time scale. The melting model we use here is the anhydrous mantle melting model of [KSL03] for a peridotitic rock composition, as implemented in the “melt simple” material model.

```
#####
# Melting and freezing #####
#####

# Because the model includes reactions that might be on a faster time scale
# than the time step of the model (melting and the freezing of melt), we use
# the operator splitting scheme.
set Use operator splitting = true

subsection Solver parameters
  subsection Operator splitting parameters
    # We choose the size of the reaction time step as 200 years, small enough
    # so that it can accurately model melting and freezing.
    set Reaction time step = 2e2

    # Additionally, we always want to do at least 10 operator splitting time
    # steps in each model time step, to accurately compute the reactions.
    set Reaction time steps per advection step = 10
  end
end
```

```

# We use the melt simple material model that includes melting and freezing of
# melt for an average mantle composition that is characteristic for a mid-ocean
# ridge setting, and mainly use its default parameters.
# In particular, we have to define how fast melting and freezing should be.
# We assume that both reactions happen on a time scale of 200 years (or a rate
# of 5e-3/year), which should be substantially shorter than the time step size,
# so that the melt fraction will always be close to equilibrium.
# As the model includes melting and freezing, we do not have to extract any melt.

subsection Material model
  set Model name = melt simple
  subsection Melt simple
    set Reference permeability = 1e-7
    set Melt extraction depth = 0.0
    set Freezing rate      = 0.005
    set Melting time scale for operator splitting = 2e2
  end
end

```

To make sure we reproduce the characteristic triangular melting region of a mid-ocean ridge, we have to set up the boundary conditions in a way so that they will lead to corner flow. At the top boundary, we can simply prescribe the half-spreading rate, and at the left boundary we can use a free-slip boundary, as material should not cross this centerline. However, we do not know the inflow and outflow velocities at the bottom and right side of the model. Instead, what we can do here is prescribing the lithostatic pressure as a boundary condition for the stress. We accomplish this by using the “initial lithostatic pressure” model. This plugin will automatically compute a 1d lithostatic pressure profile at a given point at the time of the model start and apply it as a boundary traction.

```

#####
Velocity #####
#####


# To model the divergent velocity field of a mid-ocean ridge, we prescribe
# the plate velocity (pointing away from the ridge) at the top boundary.
# We use a closed boundary with free slip conditions as the left boundary, which
# marks the ridge axis and also acts as a center line for our model, so that
# material can not cross this boundary.
# We prescribe the velocity at the top boundary using a function:
# At the ridge axis, the velocity is zero, at a distance of 10 km from the ridge
# axis or more, the rigid plate uniformly moves away from the ridge with a constant
# speed, and close to the ridge we interpolate between these two conditions.

subsection Boundary velocity model
  set Prescribed velocity boundary indicators = top:function
  set Tangential velocity boundary indicators = left
  subsection Function
    # We choose a half-spreading rate of u0=3cm/yr.
    set Function constants = u0=0.03, x0=10000
    set Variable names     = x,z
    set Function expression = if(x<x0,(1-(x/x0-1)*(x/x0-1))*u0,u0); 0
  end
end

# We prescribe the lithostatic pressure as a boundary traction on
# the bottom and right side of the model, so that material can flow in and out
# according to the flow induced by the moving plate.
subsection Boundary traction model
  set Prescribed traction boundary indicators = right:initial lithostatic pressure, bottom:initial

```

```

    ↳ lithostatic pressure

subsection Initial lithostatic pressure
  # We calculate the pressure profile at the right model boundary.
  set Representative point      = 105000, 70000
end
end

```

Finally, we have to make sure that the resolution is high enough to model melt migration. This is particularly important in regions where the porosity is low, but still high enough that the two-phase flow equations are solved (instead of the Stokes system, which is solved if there is no melt present in a cell). At the boundary between these regions, material properties like the compaction viscosity may jump, and there may be strong gradients or jumps in some solution variables such as the melt velocity and the compaction pressure. In addition, the characteristic length scale for melt transport, the compaction length δ , depends on the porosity:

$$\delta = \sqrt{\frac{(\xi + 4\eta/3)k}{\eta_f}}. \quad (54)$$

While the melt viscosity η_f is usually assumed to be constant, and the shear and compaction viscosities η and ξ increase with decreasing porosity ϕ , the permeability $k \propto \phi^2$ or $k \propto \phi^3$ dominates this relation, so that the compaction length becomes smaller for lower porosities. As the length scale of melt migration is usually smaller than for mantle convection, we want to make sure that regions where melt is present have a high resolution, and that this high resolution extends to all cells where the two-phase flow equations are solved.

```

#####
# Mesh refinement #####
#####

# We use adaptive mesh refinement to increase the resolution in regions where
# melt is present, and otherwise use a uniform grid.

subsection Mesh refinement
  set Coarsening fraction      = 0.5
  set Refinement fraction       = 0.5

  # A refinement level of 5 (4 global + 1 adaptive refinements) corresponds to
  # a cell size of approximately 1 km.
  set Initial adaptive refinement = 1
  set Initial global refinement = 4
  set Strategy                  = minimum refinement function, composition threshold
  set Time steps between mesh refinement = 5

  subsection Minimum refinement function
    set Coordinate system = cartesian
    set Function expression = 4
    set Variable names     = x,y
  end

  # We use a very small refinement threshold for the porosity to make sure that
  # all cells where the two-phase flow equations are solved (melt cells) have
  # the higher resolution.

  subsection Composition threshold
    set Compositional field thresholds = 1e-6, 1.0
  end
end

```

ASPECT also supports an alternative method to make sure that regions with melt are sufficiently well resolved, relying directly on the compaction length, and we will discuss this method as a possible modification

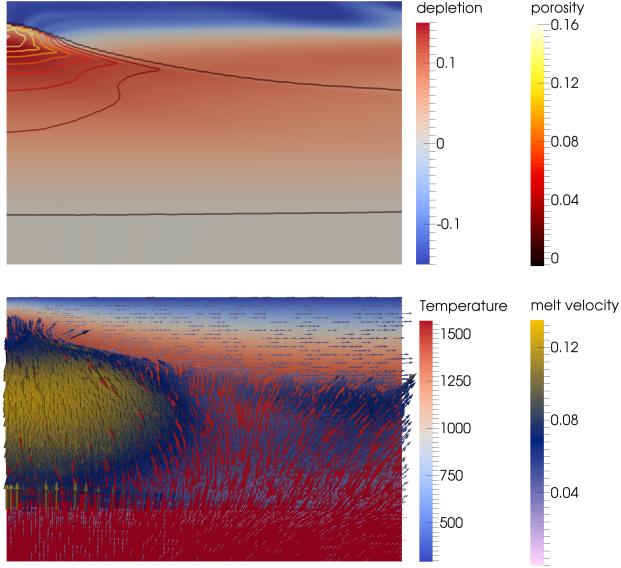


Figure 61: Mid-ocean ridge model after 8 million years. The top panel shows the depletion and porosity fields (with the characteristic triangular melting region), the bottom panel shows the temperature distribution and the melt velocity, indicated by the arrows.

to this cookbook at the end of this section.

The complete input file is located at [cookbooks/mid_ocean_ridge.prm](#).

Model evolution. When we look at the visualization output of this model (see also Figure 61), we can see how the hot material flowing in from the bottom starts to melt as it reaches lower and lower pressures and crosses the solidus. Simultaneously, melting makes the residual solid rock more depleted (as indicated by the positive values of the compositional field called ‘peridotite’). Once material approaches the surface, it is cooled from the cold boundary layer above, and melt starts to crystallize again, generating ‘enriched’ basaltic crust where it freezes (as indicated by the negative values of the compositional field called ‘peridotite’). As the temperature gradients are much sharper close to the surface, this transition from melt to solid rock is much sharper than in the melting region. Once material crystallizes, it is transported away from the ridge axis due to the flow field induced by the prescribed plate velocity at the top boundary. This way, over time, the classical triangular melting region develops at the ridge axis, and the material transported away from the ridge shows two distinct layers: The top ≈ 7 km are enriched material, and form the basaltic crust (negative peridotite field), and the ≈ 50 km below are depleted material, and form the lithosphere (positive peridotite field). A vertical profile at a distance of 80 km from the ridge axis showing this composition can be found in Figure 62.

Mesh refinement. Another option for making sure that melt migration is resolved properly in the model is using a refinement criterion that directly relates to the compaction length. This can be done in the mesh refinement section of the input file:

```

subsection Mesh refinement
  set Coarsening fraction          = 0.5
  set Refinement fraction          = 0.5

  # Note that we allow for more adaptive refinements than before, as only cells
  # with a small compaction length will be marked for refinement (as opposed to

```

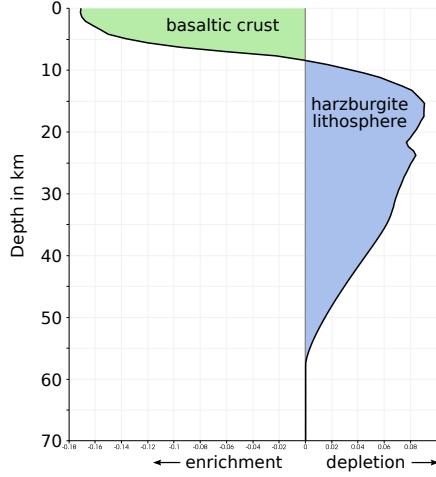


Figure 62: Vertical profile through the model domain at a distance of 80 km from the ridge axis at the end of the model run, showing the distribution of depletion and enrichment as indicated by the peridotite field.

```

# all melt cells), and we want to properly resolve the compaction length.
set Initial adaptive refinement          = 3
set Initial global refinement           = 4
set Strategy                           = minimum refinement function, compaction length
set Time steps between mesh refinement = 5

subsection Minimum refinement function
  set Coordinate system    = cartesian
  set Function expression = 4
  set Variable names      = x,y
end

# We want the cells to be 8 times smaller than the compaction length.
subsection Compaction length
  set Mesh cells per compaction length = 8.0
end
end

```

This will lead to a higher resolution particularly in regions with low (but not zero) porosity, and can be useful to resolve the strong gradients in the melt velocity and compaction pressure that are to be expected in these places (see Figure 63). Of course it is also possible to combine both methods for refining the mesh.

Extending the model. There are a number of parameters that influence the amount of melting, how fast the melt moves, and ultimately the distribution of crustal and lithospheric material. Some ideas for adapting the model setup:

- Changing the spreading rate: This can be done by choosing a different magnitude of the prescribed velocity at the top boundary, and influences the size and shape of the triangular melting region. Faster spreading allows hot material to move further away from the ridge axis, and hence facilitates a melting region that extends further in horizontal direction.
- Changing the temperature profile: This can be done by choosing a different bottom boundary temperature and influences the amount of melting, and hence the thickness of the crust. Higher temperatures lead to more melt being generated.

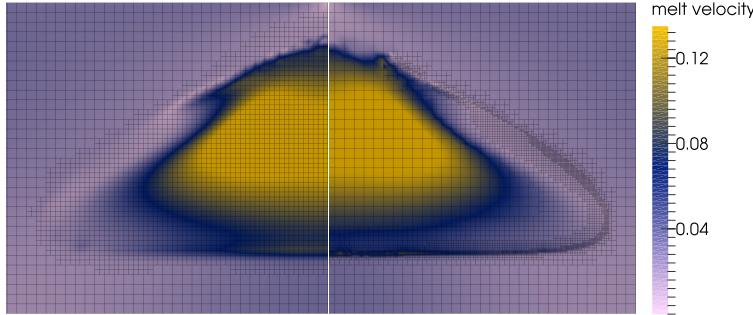


Figure 63: Mesh after a time of 3.6 million years for a model using the composition threshold refinement strategy (left) and the compaction length refinement strategy (right). Background colors indicate the melt velocity. Its sharp gradients at the interface between regions with and without melt can only be resolved using the compaction length refinement strategy.

- Changing the speed of melt migration: The velocity of the melt with respect to the solid velocity is determined by the permeability and the melt viscosity (and the pressure gradients in the melt). Increasing the permeability (by setting a different “Reference permeability” in the melt simple model) can lead to higher melt velocities, melt reaching the depth of freezing faster, and hence lower overall porosity values at steady state.
- Making the viscosity law more realistic: In this simple model, the viscosity only depends on the amount of melt that is present and is otherwise constant. This could be the reason why melt can not flow up all the way up at the ridge axis, but freezes before it reaches the surface. Introducing a temperature-dependent rheology could improve this behavior (and in reality, plastic effects might also play a role).

5.4 Benchmarks

Benchmarks are used to verify that a solver solves the problem correctly, i.e., to *verify* correctness of a code.³⁵ Over the past decades, the geodynamics community has come up with a large number of benchmarks. Depending on the goals of their original inventors, they describe stationary problems in which only the solution of the flow problem is of interest (but the flow may be compressible or incompressible, with constant or variable viscosity, etc), or they may actually model time-dependent processes. Some of them have solutions that are analytically known and can be compared with, while for others, there are only sets of numbers that are approximately known. We have implemented a number of them in ASPECT to convince ourselves (and our users) that ASPECT indeed works as intended and advertised. Some of these benchmarks are discussed below. Numerical results for several of these benchmarks are also presented in [KHB12] in much more detail than shown here.

5.4.1 Running benchmarks that require code

Some of the benchmarks require plugins like custom material models, boundary conditions, or postprocessors. To not pollute ASPECT with all these purpose-built plugins, they are kept separate from the more generic plugins in the normal source tree. Instead, the benchmarks have all the necessary code in .cc files in the benchmark directories. Those are then compiled into a shared library that will be used by ASPECT if it is

³⁵Verification is the first half of the *verification and validation* (V&V) procedure: *verification* intends to ensure that the mathematical model is solved correctly, while *validation* intends to ensure that the mathematical model is correct. Obviously, much of the aim of computational geodynamics is to validate the models that we have.

referenced in a `.prm` file. Let's take the SolCx benchmark as an example (see Section 5.4.4). The directory contains:

- `solcx.cc` – the code file containing a material model “SolCxMaterial” and a postprocessor “SolCx-Postprocessor”,
- `solcx.prm` – the parameter file referencing these plugins,
- `CMakeLists.txt` – a `cmake` configuration that allows you to compile `solcx.cc`.

To run this benchmark you need to follow the general outline of steps discussed in Section 6.2. For the current case, this amounts to the following:

1. Move into the directory of that particular benchmark:

```
$ cd benchmarks/solcx
```

2. Set up the project:

```
$ cmake .
```

By default, `cmake` will look for the ASPECT binary and other information in a number of directories relative to the current one. If it is unable to pick up where ASPECT was built and installed, you can specify this directory explicitly this using `-D Aspect_DIR=<...>` as an additional flag to `cmake`, where `<...>` is the path to the build directory.

3. Build the library:

```
$ make
```

This will generate the file `libsolcx.so`.

Finally, you can run ASPECT with `solcx.prm`:

```
$ ../../aspect solcx.prm
```

where again you may have to use the appropriate path to get to the ASPECT executable. You will need to run ASPECT from the current directory because `solcx.prm` refers to the plugin as `./libsolcx.so`, i.e., in the current directory.

5.4.2 Onset of convection benchmark

This section was contributed by Max Rudolph, based on a course assignment for “Geodynamic Modeling” at Portland State University.

Here we use ASPECT to numerically reproduce the results of a linear stability analysis for the onset of convection in a fluid layer heated from below. This exercise was assigned to students at Portland State University as a first step towards setting up a nominally Earth-like mantle convection model. Hence, representative length scales and transport properties for Earth are used. This cookbook consists of a jupyter notebook (`benchmarks/onset-of-convection/onset-of-convection.ipynb`) that is used to run ASPECT and analyze the results of several calculations. To use this code, you must compile ASPECT and give the path to the executable in the notebook as `aspect_bin`.

The linear stability analysis for the onset of convection appears in Turcotte and Schubert [TS14] (section 6.19). The linear stability analysis assumes the Boussinesq approximation and makes predictions for the growth rate (vertical velocity) of instabilities and the critical Rayleigh number Ra_{cr} above which convection will occur. Ra_{cr} depends only on the dimensionless wavelength of the perturbation, which is assumed to be

equal to the width of the domain. The domain has height b and width λ and the perturbation is described by

$$T'(x, y) = T'_0 \cos\left(\frac{2\pi x}{\lambda}\right) \sin\left(\frac{\pi y}{b}\right),$$

where T'_0 is the amplitude of the perturbation. Note that because we place the bottom boundary of the domain at $y = 0$ and the top at $y = b$, the perturbation vanishes at the top and bottom boundaries. This departs slightly from the setup in [TS14], where the top and bottom boundaries of the domain are at $y = \pm b/2$. The analytic expression for the critical Rayleigh number, Ra_{cr} is given in Turcotte and Schubert [TS14] equation (6.319):

$$Ra_{cr} = \frac{\left(\pi^2 + \frac{4\pi^2 b^2}{\lambda^2}\right)^3}{\frac{4\pi^2 b^2}{\lambda^2}}.$$

The linear stability analysis also makes a prediction for the dimensionless growth rate of the instability α' (Turcotte and Schubert [TS14], equation (6.315)). The maximum vertical velocity is given by

$$v_{y,max} = \frac{2\pi}{\lambda} \phi'_0 e^{\alpha' t},$$

where

$$\phi'_0 = -\frac{2\pi}{\lambda} \frac{\rho_0 g \alpha T'_0}{\mu} \left(\frac{4\pi^2}{\lambda^2} + \frac{\pi^2}{b^2} \right)^{-2},$$

and

$$\alpha' = \frac{\kappa}{b^2} \left[\frac{\rho_0 g \alpha b^3 \Delta T}{\mu \kappa} \left(\frac{\frac{4\pi^2 b^2}{\lambda^2}}{\left(\frac{4\pi^2 b^2}{\lambda^2} + \pi^2 \right)^2} \right) - \left(\pi^2 + \frac{4\pi^2 b^2}{\lambda^2} \right) \right].$$

We use bisection to determine Ra_{cr} for specific choices of the domain geometry, keeping the depth b constant and varying the domain width λ . If the vertical velocity increases from the first to the second timestep, the system is unstable to convection. Otherwise, it is stable and convection will not occur. Each calculation is terminated after the second timestep. Fig. 64 shows the numerically-determined threshold for the onset of convection, which can be compared directly with the theoretical prediction (green curve) and Fig. 6.39 of [TS14]. The relative error between the numerically-determined value of Ra_{cr} and the analytic solution are shown in the right panel of Fig. 64.

5.4.3 The van Keken thermochemical composition benchmark

This section is a co-production of Cedric Thieulot, Juliane Dannberg, Timo Heister and Wolfgang Bangerth with an extension to this benchmark provided by the Virginia Tech Department of Geosciences class “Geodynamics and ASPECT” co-taught by Scott King and D. Sarah Stamps.

One of the most widely used benchmarks for mantle convection codes is the isoviscous Rayleigh-Taylor case (“case 1a”) published by van Keken *et al.* in [vKKS⁹⁷]. The benchmark considers a 2d situation where a lighter fluid underlies a heavier one with a non-horizontal interface between the two of them. This unstable layering causes the lighter fluid to start rising at the point where the interface is highest. Fig. 65 shows a time series of images to illustrate this.

Although van Keken’s paper title suggests that the paper is really about thermochemical convection, the part we look here can equally be considered as thermal or chemical convection: all that is necessary is that we describe the fluid’s density somehow. We can do that by using an inhomogeneous initial temperature field, or an inhomogeneous initial composition field. We will use the input file in `cookbooks/van-keken-discontinuous.prm` as input, the central piece of which is as follows (go to the actual input file to see the remainder of the input parameters):

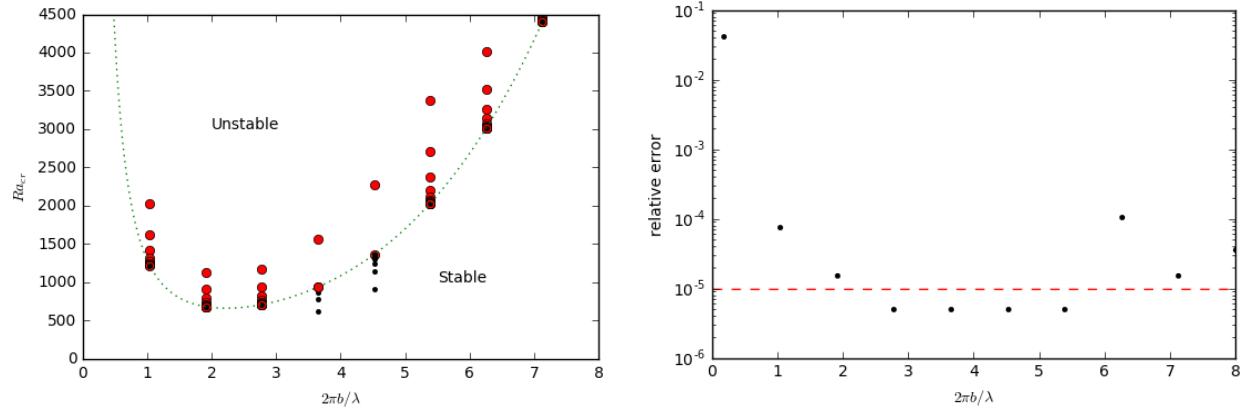


Figure 64: Left: Comparison of numerically-determined and theoretical values for Ra_{cr} . Red circles indicate numerically-simulations unstable to convection, black circles indicate simulations that are stable. The green dashed curve indicates the theoretical prediction. Right: Relative error in determination of Ra_{cr} . The dashed red line indicates the error tolerance used in bisection procedure.

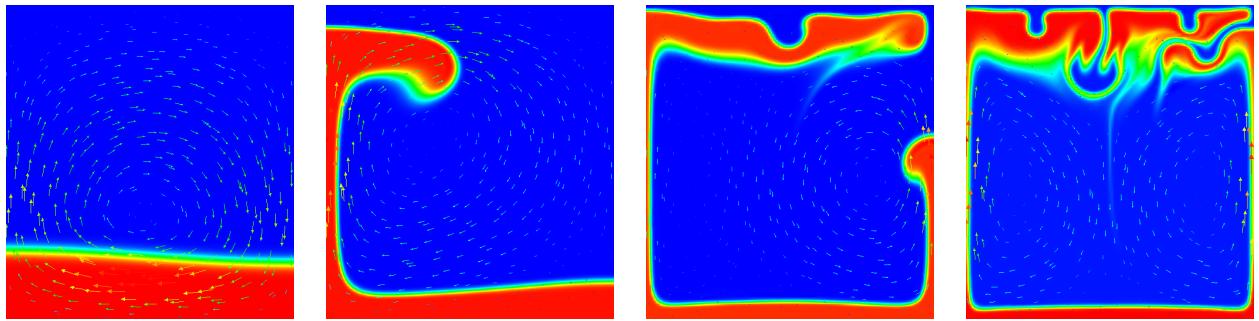


Figure 65: *Van Keken benchmark (using a smoothed out interface, see the main text): Compositional field at times $t = 0, 300, 900, 1800$.*

```

subsection Material model
  set Model name = simple
  subsection Simple model
    set Viscosity           = 1e2
    set Thermal expansion coefficient = 0
    set Density differential for compositional field 1 = -10
  end
end

subsection Initial composition model
  set Model name = function
  subsection Function
    set Variable names      = x,z
    set Function constants  = pi=3.14159
    set Function expression = if( (z>0.2+0.02*cos(pi*x/0.9142)) , 0 , 1 )
  end
end

```

The first part of this selects the `simple` material model and sets the thermal expansion to zero (resulting in a density that does not depend on the temperature, making the temperature a passively advected field) and instead makes the density depend on the first compositional field. The second section prescribes that the first compositional field's initial conditions are 0 above a line describes by a cosine and 1 below it. Because the dependence of the density on the compositional field is negative, this means that a lighter fluid underlies a heavier one.

The dynamics of the resulting flow have already been shown in Fig. 65. The measure commonly considered in papers comparing different methods is the root mean square of the velocity, which we can get using the following block in the input file (the actual input file also enables other postprocessors):

```

subsection Postprocess
  set List of postprocessors = velocity statistics
end

```

Using this, we can plot the evolution of the fluid's average velocity over time, as shown in the left panel of Fig. 66. Looking at this graph, we find that both the timing and the height of the first peak is already well converged on a simple 32×32 mesh (5 global refinements) and is very consistent (to better than 1% accuracy) with the results in the van Keken paper.

That said, it is startling that the second peak does not appear to converge despite the fact that the various codes compared in [vKKS⁹⁷] show good agreement in this comparison. Tracking down the cause for this proved to be a lesson in benchmark design; in hindsight, it may also explain why van Keken *et al.* stated presciently in their abstract that “...good agreement is found for the initial rise of the unstable lower layer; however, the timing and location of the later smaller-scale instabilities may differ between methods.” To understand what is happening here, note that the first peak in these plots corresponds to the plume that rises along the left edge of the domain and whose evolution is primarily determined by the large-scale shape of the initial interface (i.e., the cosine used to describe the initial conditions in the input file). This is a first order deterministic effect, and is obviously resolved already on the coarsest mesh shown used. The second peak corresponds to the plume that rises along the right edge, and its origin along the interface is much harder to trace – its position and the timing when it starts to rise is certainly not obvious from the initial location of the interface. Now recall that we are using a finite element field using continuous shape functions for the composition that determines the density differences that drive the flow. But this interface is neither aligned with the mesh, nor can a discontinuous function be represented by continuous shape functions to begin with. In other words, we may *input* the initial conditions as a discontinuous functions of zero and one in the parameter file, but the initial conditions used in the program are in fact different: they are the *interpolated* values of this discontinuous function on a finite element mesh. This is shown in Fig. 67. It is obvious that these initial conditions agree on the large scale (the determinant of the first plume), but not in

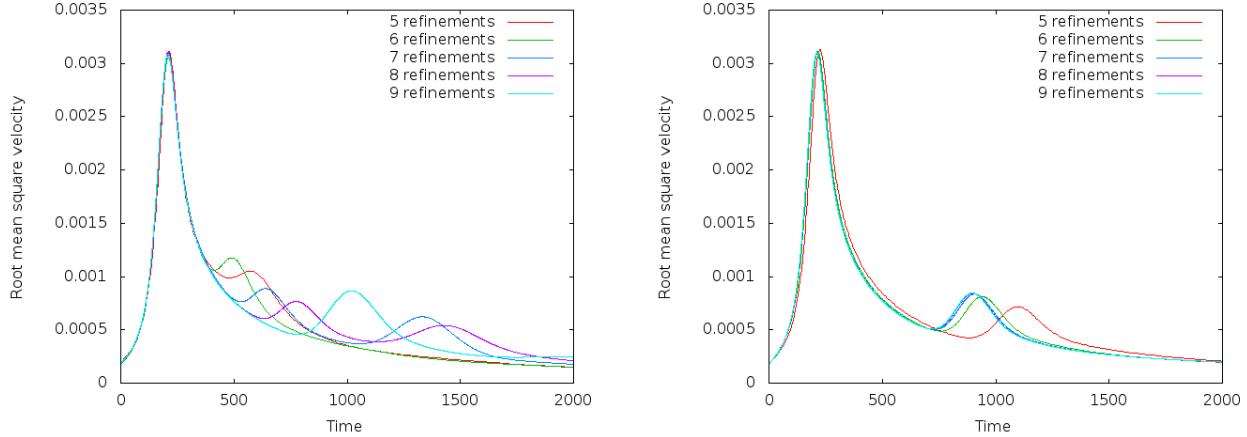


Figure 66: Van Keken benchmark with discontinuous (left) and smoothed, continuous (right) initial conditions for the compositional field: Evolution of the root mean square velocity $\left(\frac{1}{|\Omega|} \int_{\Omega} |\mathbf{u}(\mathbf{x}, t)|^2 d\mathbf{x}\right)^{1/2}$ as a function of time for different numbers of global mesh refinements. 5 global refinements correspond to a 32×32 mesh, 9 refinements to a 512×512 mesh.

the steps that may (and do, in fact) determine when and where the second plume will rise. The evolution of the resulting compositional field is shown in Fig. 68 and it is obvious that the second, smaller plume starts to rise from a completely different location – no wonder the second peak in the root mean square velocity plot is in a different location and with different height!

The conclusion one can draw from this is that if the outcome of a computational experiment depends so critically on very small details like the steps of an initial condition, then it's probably not a particularly good measure to look at in a benchmark. That said, the benchmark is what it is, and so we should try to come up with ways to look at the benchmark in a way that allows us to reproduce what van Keken *et al.* had agreed upon. To this end, note that the codes compared in that paper use all sorts of different methods, and one can certainly agree on the fact that these methods are not identical on small length scales. One approach to make the setup more mesh-independent is to replace the original discontinuous initial condition with a smoothed out version; of course, we can still not represent it exactly on any given mesh, but we can at least get closer to it than for discontinuous variables. Consequently, let us use the following initial conditions instead (see also the file [cookbooks/van-keken-smooth.prm](#)):

```
subsection Initial composition model
  set Model name = function
  subsection Function
    set Variable names      = x,z
    set Function constants  = pi=3.14159
    set Function expression = 0.5*(1+tanh((0.2+0.02*cos(pi*x/0.9142)-z)/0.02))
  end
end
```

This replaces the discontinuous initial conditions with a smoothed out version with a half width of around 0.01. Using this, the root mean square plot now looks as shown in the right panel of Fig. 66. Here, the second peak also converges quickly, as hoped for.

The exact location and height of the two peaks is in good agreement with those given in the paper by van Keken *et al.*, but not exactly where desired (the error is within a couple of per cent for the first peak, and probably better for the second, for both the timing and height of the peaks). This has to do with the fact that they depend on the exact size of the smoothing parameter (the division by 0.02 in the formula for the smoothed initial condition). However, for more exact results, one can choose this half width

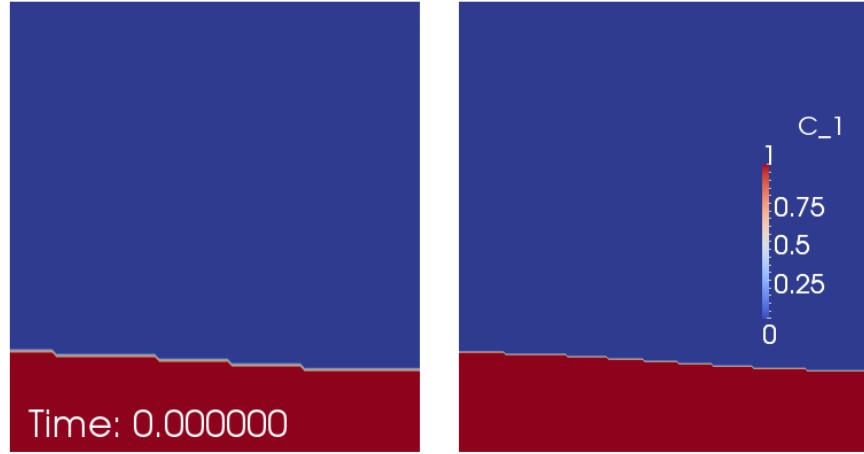


Figure 67: *Van Keken benchmark with discontinuous initial conditions for the compositional field: Initial compositional field interpolated onto a 32×32 (left) and 64×64 finite element mesh (right).*

parameter proportional to the mesh size and thereby get more accurate results. The point of the section was to demonstrate the reason for the lack of convergence.

In this section we extend the van Keken cookbook following up the work previously completed by Cedric Thieulot, Julianne Dannberg, Timo Heister and Wolfgang Bangerth. *This section contributed by Grant Euen, Tahiry Rajaonarison, and Shangxin Liu as part of the Geodynamics and ASPECT class at Virginia Tech.*

As already mentioned above, using a half width parameter proportional to the mesh size allows for more accurate results. We test the effect of the half width size of the smoothed discontinuity by changing the division by 0.02, the smoothing parameter, in the formula for the smoothed initial conditions into values proportional to the mesh size. We use 7 global refinements because the root mean square velocity converges at greater resolution while keeping average runtime around 5 to 25 minutes. These runtimes were produced by the BlueRidge cluster of the Advanced Research Computing (ARC) program at Virginia Tech. BlueRidge is a 408-node Cray CS-300 cluster; each node outfitted with two octa-core Intel Sandy Bridge CPUs and 64 GB of memory. A chart of average runtimes for 5 through 10 global refinements on one node can be seen in Table 5. For 7 global refinements (128×128 mesh size), the size of the mesh is 0.0078 corresponding to a half width parameter of 0.0039. The smooth model allows for much better convergence of the secondary plumes, although they are still more scattered than the primary plumes.

Global Refinements	Number of Processors			
	4	8	12	16
5	28.1 seconds	19.8 seconds	19.6 seconds	17.1 seconds
6	3.07 minutes	1.95 minutes	1.49 minutes	1.21 minutes
7	23.33 minutes	13.92 minutes	9.87 minutes	7.33 minutes
8	3.08 hours	1.83 hours	1.30 hours	56.33 minutes
9	1.03 days	15.39 hours	10.44 hours	7.53 hours
10	More than 6 days	More than 6 days	3.39 days	2.56 days

Table 5: *Average runtimes for the van Keken Benchmark with smoothed initial conditions. These times are for the entire computation, a final time step number of 2000. All of these tests were run using ASPECT version 1.3 in release mode, and used different numbers of processors on one node on the BlueRidge cluster of ARC at Virginia Tech.*

This convergence is due to changing the smoothing parameter, which controls how much of the problem is smoothed over. As the parameter is increased, the smoothed boundary grows and vice versa. As the

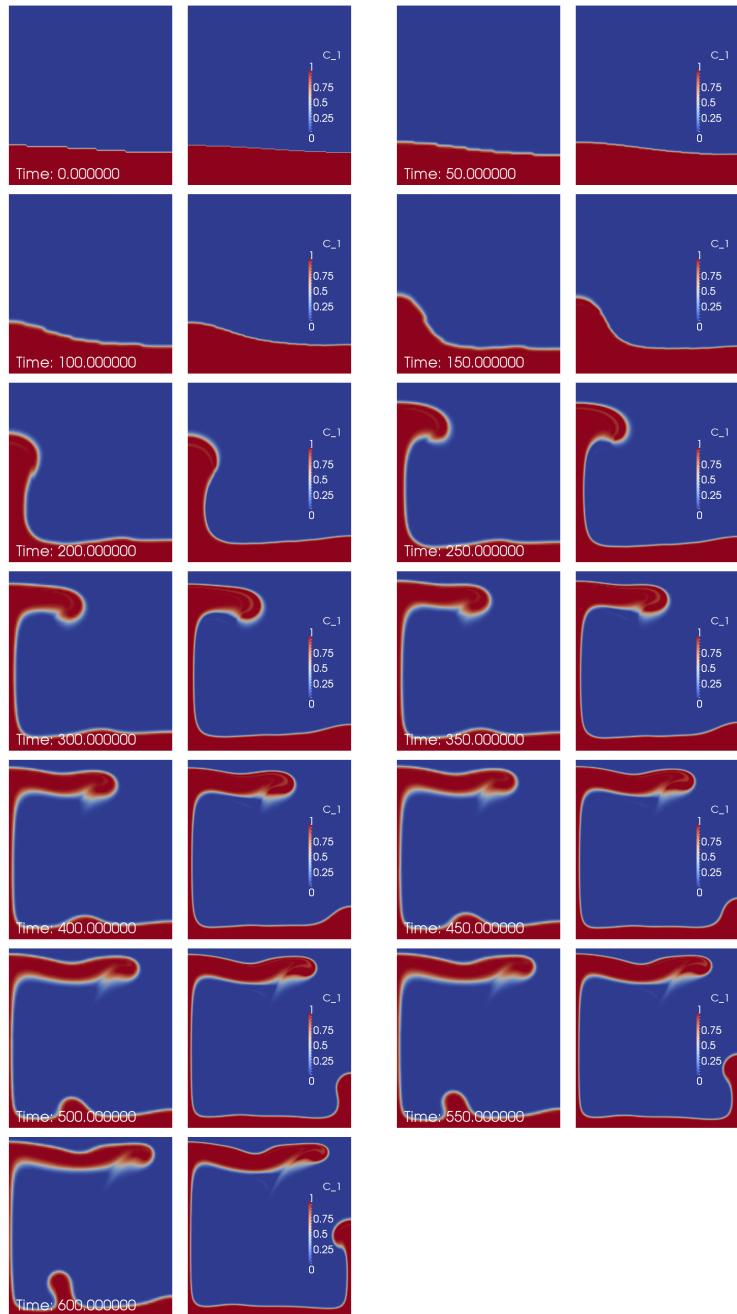


Figure 68: *Van Keken benchmark with discontinuous initial conditions for the compositional field: Evolution of the compositional field over time on a 32×32 (first and third column; left to right and top to bottom) and 64×64 finite element mesh (second and fourth column).*

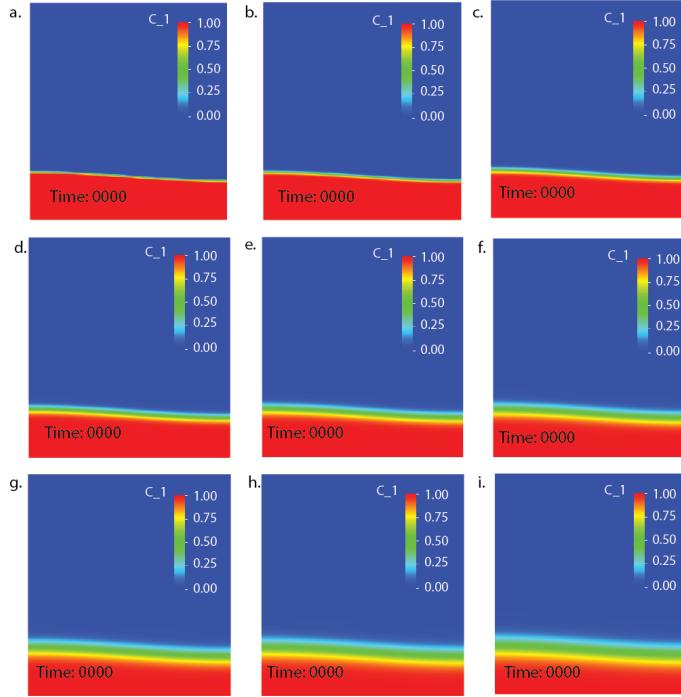


Figure 69: *Van Keken Benchmark using smoothed out interface at 7 global refinements: compositional field at time $t = 0$ using smoothing parameter size: a) 0.0039, b) 0.0078, c) 0.0156, d) 0.0234, e) 0.0312, f) 0.0390, g) 0.0468, h) 0.0546, i) 0.0624.*

smoothed boundary shrinks it becomes sharper until the original discontinuous behavior is revealed. As it grows, the two layers eventually become one large, transitioning layer rather than two distinct layers separated by a boundary. These effects can be seen in Fig. 69. The overall effect is that the secondary rise is at different times based on these conditions. In general, as the smoothing parameter is decreased, the smoothed boundary shrinks and the plumes rise more quickly. As it is increased, the boundary grows and the plumes rise more slowly. This trend can be used to force a more accurate convergence from the secondary plumes.

The evolution in time of the resulting compositional fields (Fig. 70) shows that the first peak converges as the smoothed interface decreases. There is a good agreement for the first peak for all smoothing parameters. As the width of the discontinuity increases, the second peak rises both later and more slowly.

Now let us further add a two-layer viscosity model to the domain. This is done to recreate the two non-isoviscous Rayleigh-Taylor instability cases (“cases 1b and 1c”) published in van Keken *et al.* in [vKKS⁺97]. Let’s assume the viscosity value of the upper heavier layer is η_t and the viscosity value of the lower lighter layer is η_b . Based on the initial constant viscosity value 1×10^2 Pa s, we set the viscosity proportion $\frac{\eta_t}{\eta_b} = 0.1, 0.01$, meaning the viscosity of the upper, heavier layer is still 1×10^2 Pa s, but the viscosity of the lower, lighter layer is now either 10 or 1 Pa s, respectively. The viscosity profiles of the discontinuous and smooth models are shown in Fig. 71.

For both benchmark cases, discontinuous and smooth, and both viscosity proportions, 0.1 and 0.01, the results are shown at the end time step number, 2000, in Fig. 72. This was generated using the original input parameter file, running the cases with 8 global refinement steps, and also adding the two-layer viscosity model.

Compared to the results of the constant viscosity throughout the domain, the plumes rise faster when

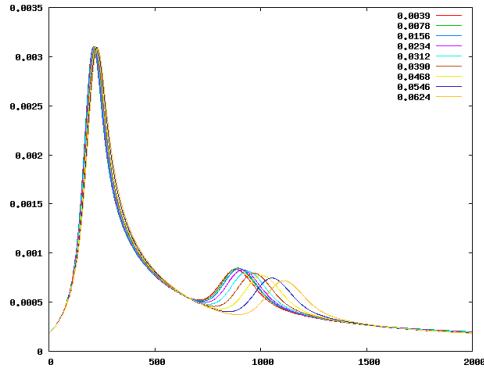


Figure 70: *Van Keken benchmark with smoothed initial conditions for the compositional field using 7 global refinements for different smoothing parameters. Number of the time step is shown on the x-axis, while root mean square velocity is shown on the y-axis.*

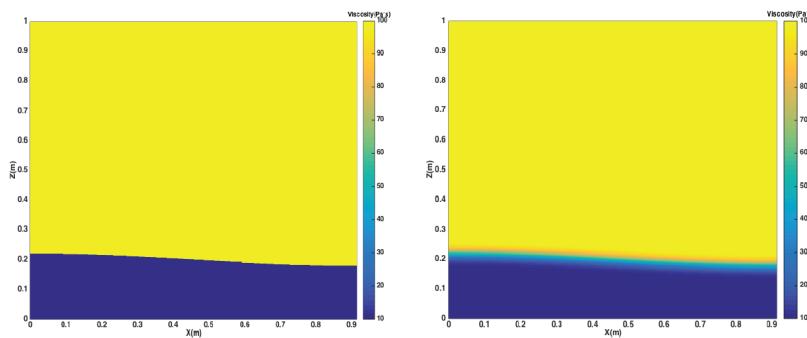


Figure 71: *Van Keken benchmark using layers of different viscosities. The left image is the discontinuous case, while right is the smooth. Both are shown at $t=0$.*

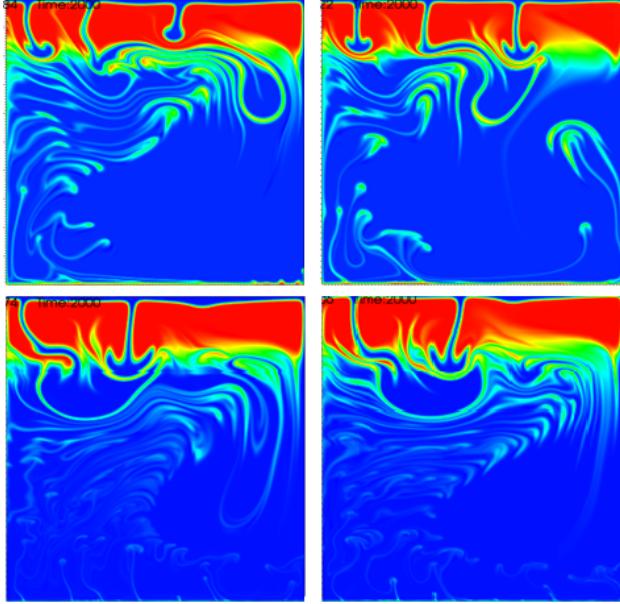


Figure 72: *Van Keken benchmark two-layer viscosity model at final time step number, 2000.* These images show layers of different compositions and viscosities. Discontinuous cases are the left images, smooth cases are the right. The upper images are $\frac{\eta_t}{\eta_b} = 0.1$, and the lower are $\frac{\eta_t}{\eta_b} = 0.01$.

adding the two-layer viscosity model. Also, the larger the viscosity difference is, the earlier the plumes appear and the faster their ascent. To further reveal the effect of the two-layer viscosity model, we also plot the evolution of the fluids' average velocity over time, as shown in Fig. 73.

We can observe that when the two-layer viscosity model is added, there is only one apparent peak for each case. The first peaks of the 0.01 viscosity contrast tests appear earlier and are larger in magnitude than those of 0.1 viscosity contrast tests. There are no secondary plumes and the whole system tends to reach stability after around 500 time steps.

5.4.4 The SolCx Stokes benchmark

The SolCx benchmark is intended to test the accuracy of the solution to a problem that has a large jump in the viscosity along a line through the domain. Such situations are common in geophysics: for example, the viscosity in a cold, subducting slab is much larger than in the surrounding, relatively hot mantle material.

The SolCx benchmark computes the Stokes flow field of a fluid driven by spatial density variations, subject to a spatially variable viscosity. Specifically, the domain is $\Omega = [0, 1]^2$, gravity is $\mathbf{g} = (0, -1)^T$ and the density is given by $\rho(\mathbf{x}) = \sin(\pi x_1) \cos(\pi x_2)$; this can be considered a density perturbation to a constant background density. The viscosity is

$$\eta(\mathbf{x}) = \begin{cases} 1 & \text{for } x_1 \leq 0.5, \\ 10^6 & \text{for } x_1 > 0.5. \end{cases}$$

This strongly discontinuous viscosity field yields an almost stagnant flow in the right half of the domain and consequently a singularity in the pressure along the interface. Boundary conditions are free slip on all of $\partial\Omega$. The temperature plays no role in this benchmark. The prescribed density field and the resulting velocity field are shown in Fig. 74.

The SolCx benchmark was previously used in [DMGT11, Section 4.1.1] (references to earlier uses of the benchmark are available there) and its analytic solution is given in [Zho96]. ASPECT contains an

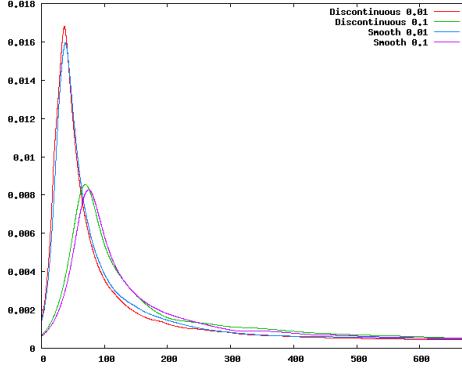


Figure 73: Van Keken benchmark: Evolution of the root mean square velocity as a function of time for different viscosity contrast proportions (0.1/0.01) for both discontinuous and smooth models.

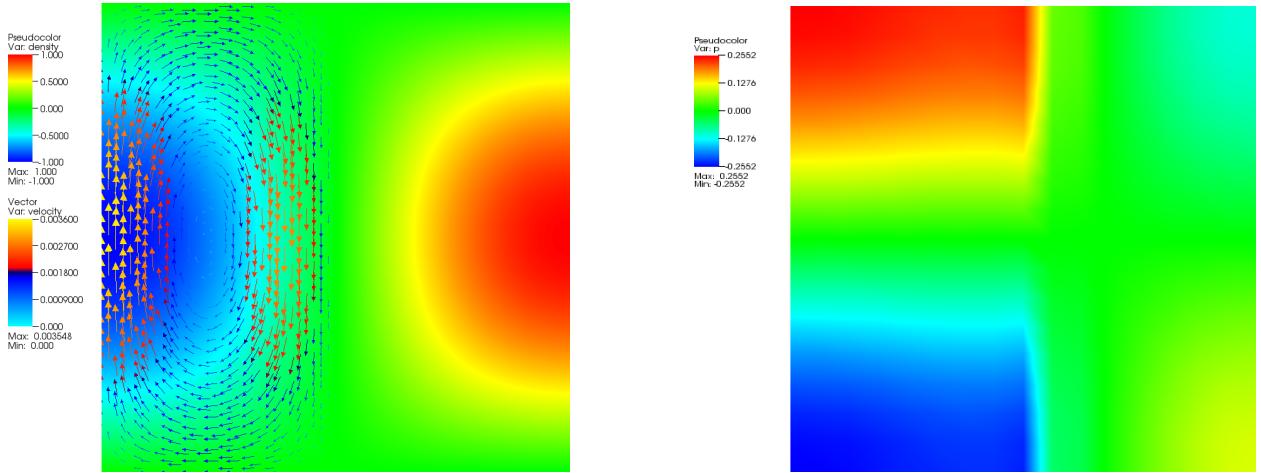


Figure 74: SolCx Stokes benchmark. Left: The density perturbation field and overlaid to it some velocity vectors. The viscosity is very large in the right hand, leading to a stagnant flow in this region. Right: The pressure on a relatively coarse mesh, showing the internal layer along the line where the viscosity jumps.

implementation of this analytic solution taken from the Underworld package (see [MQL⁺07] and <http://www.underworldproject.org/>, and correcting for the mismatch in sign between the implementation and the description in [DMGT11]).

To run this benchmark, the following input file will do (see the files in `benchmarks/solcx/` to rerun the benchmark):

```
set Additional shared libraries          = ./libsolcx.so

##### Global parameters

set Dimension                          = 2
set Start time                         = 0
set End time                           = 0
set Output directory                   = output
```

```

set Pressure normalization          = volume

##### Parameters describing the model

subsection Geometry model
  set Model name = box

  subsection Box
    set X extent = 1
    set Y extent = 1
  end
end

subsection Boundary velocity model
  set Tangential velocity boundary indicators = left, right, bottom, top
end

subsection Material model
  set Model name = SolCxMaterial

  subsection SolCx
    set Viscosity jump = 1e6
  end
end

subsection Gravity model
  set Model name = vertical
end

##### Parameters describing the temperature field

subsection Boundary temperature model
  set List of model names = box
end

subsection Initial temperature model
  set Model name = perturbed box
end

##### Parameters describing the discretization

subsection Discretization
  set Stokes velocity polynomial degree      = 2
  set Use locally conservative discretization = false
end

```

```

subsection Mesh refinement
  set Initial adaptive refinement      = 0
  set Initial global refinement       = 4
end

#####
# Parameters describing what to do with the solution

subsection Postprocess
  set List of postprocessors = SolCxPostprocessor, visualization
end

```

Since this is the first cookbook in the benchmarking section, let us go through the different parts of this file in more detail:

- The material model and the postprocessor
- The first part consists of parameter setting for overall parameters. Specifically, we set the dimension in which this benchmark runs to two and choose an output directory. Since we are not interested in a time dependent solution, we set the end time equal to the start time, which results in only a single time step being computed.
The last parameter of this section, `Pressure normalization`, is set in such a way that the pressure is chosen so that its *domain* average is zero, rather than the pressure along the surface, see Section 2.5.
- The next part of the input file describes the setup of the benchmark. Specifically, we have to say how the geometry should look like (a box of size 1×1) and what the velocity boundary conditions shall be (tangential flow all around – the box geometry defines four boundary indicators for the left, right, bottom and top boundaries, see also Section A.42). This is followed by subsections choosing the material model (where we choose a particular model implemented in ASPECT that describes the spatially variable density and viscosity fields, along with the size of the viscosity jump) and finally the chosen gravity model (a gravity field that is the constant vector $(0, -1)^T$, see Section A.52).
- The part that follows this describes the boundary and initial values for the temperature. While we are not interested in the evolution of the temperature field in this benchmark, we nevertheless need to set something. The values given here are the minimal set of inputs.
- The second-to-last part sets discretization parameters. Specifically, it determines what kind of Stokes element to choose (see Section A.38 and the extensive discussion in [KHB12]). We do not adaptively refine the mesh but only do four global refinement steps at the very beginning. This is obviously a parameter worth playing with.
- The final section on postprocessors determines what to do with the solution once computed. Here, we do two things: we ask ASPECT to compute the error in the solution using the setup described in the Duretz et al. paper [DMGT11], and we request that output files for later visualization are generated and placed in the output directory. The functions that compute the error automatically query which kind of material model had been chosen, i.e., they can know whether we are solving the SolCx benchmark or one of the other benchmarks discussed in the following subsections.

Upon running ASPECT with this input file, you will get output of the following kind (obviously with different timings, and details of the output may also change as development of the code continues):

```

aspect/cookbooks> ../../aspect solcx.prm
Number of active cells: 256 (on 5 levels)
Number of degrees of freedom: 3,556 (2,178+289+1,089)

*** Timestep 0: t=0 years

```

```

Solving temperature system... 0 iterations.
Rebuilding Stokes preconditioner...
Solving Stokes system... 30+3 iterations.

Postprocessing:
Errors u_L1, p_L1, u_L2, p_L2: 1.125997e-06, 2.994143e-03, 1.670009e-06, 9.778441e-03
Writing graphical output:      output/solution/solution-00000

```

Total wallclock time elapsed since start	no. calls	wall time	% of total
		1.51s	
Section	no. calls	wall time	% of total
Assemble Stokes system	1	0.114s	7.6%
Assemble temperature system	1	0.284s	19%
Build Stokes preconditioner	1	0.0935s	6.2%
Build temperature preconditioner	1	0.0043s	0.29%
Solve Stokes system	1	0.0717s	4.8%
Solve temperature system	1	0.000753s	0.05%
Postprocessing	1	0.627s	42%
Setup dof systems	1	0.19s	13%

One can then visualize the solution in a number of different ways (see Section 4.4), yielding pictures like those shown in Fig. 74. One can also analyze the error as shown in various different ways, for example as a function of the mesh refinement level, the element chosen, etc.; we have done so extensively in [KHB12].

5.4.5 The SolKz Stokes benchmark

The SolKz benchmark is another variation on the same theme as the SolCx benchmark above: it solves a Stokes problem with a spatially variable viscosity but this time the viscosity is not a discontinuous function but grows exponentially with the vertical coordinate so that its overall variation is again 10^6 . The forcing is again chosen by imposing a spatially variable density variation. For details, refer again to [DMGT11].

The following input file, only a small variation of the one in the previous section, solves the benchmark (see [benchmarks/solkz/](#)):

```

# A description of the SolKz benchmark for which a known solution
# is available. See the manual for more information.

set Additional shared libraries          = ./libsolkz.so

##### Global parameters

set Dimension                           = 2
set Start time                         = 0
set End time                           = 0
set Output directory                   = output
set Pressure normalization             = volume

##### Parameters describing the model

```

```

subsection Geometry model
  set Model name = box

  subsection Box
    set X extent = 1
    set Y extent = 1
  end
end

subsection Boundary velocity model
  set Tangential velocity boundary indicators = left, right, bottom, top
end

subsection Material model
  set Model name = SolKzMaterial
end

subsection Gravity model
  set Model name = vertical
end

##### Parameters describing the temperature field

subsection Boundary temperature model
  set List of model names = box
end

subsection Initial temperature model
  set Model name = perturbed box
end

##### Parameters describing the discretization

subsection Discretization
  set Stokes velocity polynomial degree      = 2
  set Use locally conservative discretization = false
end

subsection Mesh refinement
  set Initial adaptive refinement           = 0
  set Initial global refinement            = 4
end

##### Parameters describing what to do with the solution

subsection Postprocess

```

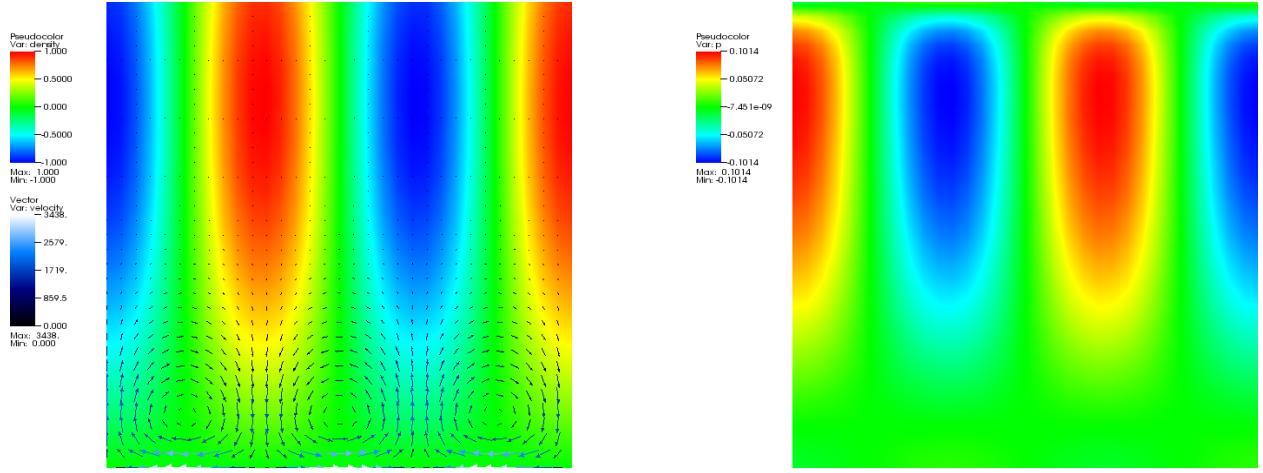


Figure 75: *SolKz Stokes benchmark.* *Left:* The density perturbation field and overlaid to it some velocity vectors. The viscosity grows exponentially in the vertical direction, leading to small velocities at the top despite the large density variations. *Right:* The pressure.

```
set List of postprocessors = SolKzPostprocessor, visualization
end
```

The output when running ASPECT on this parameter file looks similar to the one shown for the SolCx case. The solution when computed with one more level of global refinement is visualized in Fig. 75.

5.4.6 The “inclusion” Stokes benchmark

The “inclusion” benchmark again solves a problem with a discontinuous viscosity, but this time the viscosity is chosen in such a way that the discontinuity is along a circle. This ensures that, unlike in the SolCx benchmark discussed above, the discontinuity in the viscosity never aligns to cell boundaries, leading to much larger difficulties in obtaining an accurate representation of the pressure. Specifically, the almost discontinuous pressure along this interface leads to oscillations in the numerical solution. This can be seen in the visualizations shown in Fig. 76. As before, for details we refer to [DMGT11]. The analytic solution against which we compare is given in [SP03]. An extensive discussion of convergence properties is given in [KHB12].

The benchmark can be run using the parameter files in `benchmarks/inclusion/`. The material model, boundary condition, and postprocessor are defined in `benchmarks/inclusion/inclusion.cc`. Consequently, this code needs to be compiled into a shared lib before you can run the tests.

```
#####
# Global parameters

set Additional shared libraries      = ./libinclusion.so

set Dimension                         = 2
set Start time                        = 0
set End time                          = 0
set Output directory                  = output
set Pressure normalization           = volume
```

Link to a general section on how you can compile libs for the benchmarks. Revisit this once we have the machinery in place to choose nonzero boundary conditions in a more elegant way.

The following prm file isn't annotated yet. How to annotate if we have a .lib?

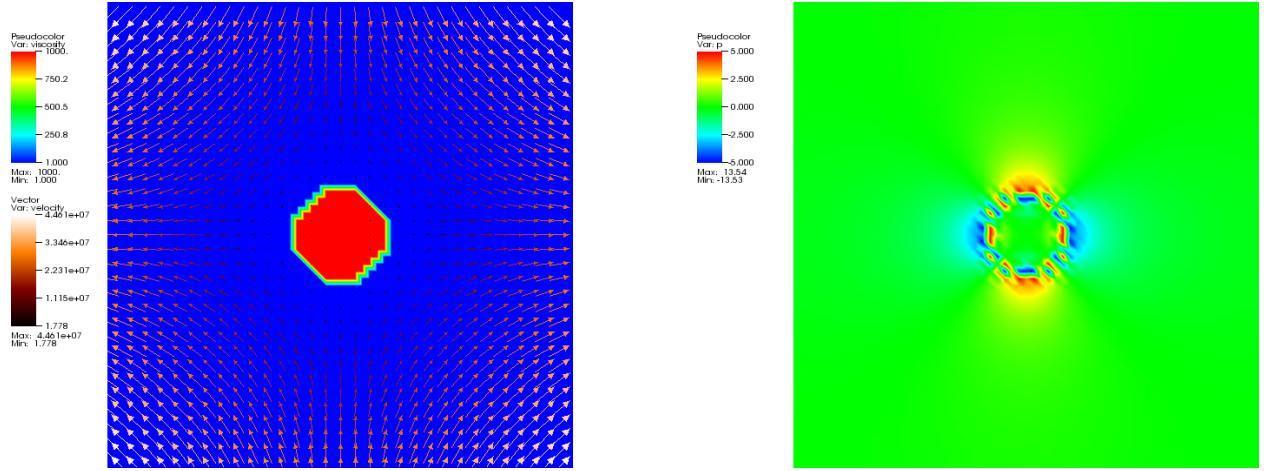


Figure 76: *Inclusion Stokes* benchmark. Left: The viscosity field when interpolated onto the mesh (internally, the “exact” viscosity field – large inside a circle, small outside – is used), and overlaid to it some velocity vectors. Right: The pressure with its oscillations along the interface. The oscillations become more localized as the mesh is refined.

```
#####
##### Parameters describing the model

subsection Geometry model
  set Model name = box

  subsection Box
    set X extent = 2
    set Y extent = 2
  end
end

subsection Boundary velocity model
  set Prescribed velocity boundary indicators = left : InclusionBoundary, \
                                             right : InclusionBoundary, \
                                             bottom: InclusionBoundary, \
                                             top   : InclusionBoundary
end

subsection Material model
  set Model name = InclusionMaterial

  subsection Inclusion
    set Viscosity jump = 1e3
  end
end

subsection Gravity model
```

```

    set Model name = vertical
end

#####
# Parameters describing the temperature field

subsection Boundary temperature model
    set List of model names = box
end

subsection Initial temperature model
    set Model name = perturbed box
end

#####
# Parameters describing the discretization

subsection Discretization
    set Stokes velocity polynomial degree      = 2
    set Use locally conservative discretization = false
end

subsection Mesh refinement
    set Initial adaptive refinement          = 0
    set Initial global refinement           = 6
end

#####
# Parameters describing what to do with the solution

subsection Postprocess
    set List of postprocessors = InclusionPostprocessor, visualization
end

```

5.4.7 The Burstedde variable viscosity benchmark

This section was contributed by Iris van Zelst.

This benchmark is intended to test solvers for variable viscosity Stokes problems. It begins with postulating a smooth exact polynomial solution to the Stokes equation for a unit cube, first proposed by [DB14] and also described by [BSA⁺13]:

$$\mathbf{u} = \begin{pmatrix} x + x^2 + xy + x^3y \\ y + xy + y^2 + x^2y^2 \\ -2z - 3xz - 3yz - 5x^2yz \end{pmatrix} \quad (55)$$

$$p = xyz + x^3y^3z - \frac{5}{32}. \quad (56)$$

It is then trivial to verify that the velocity field is divergence-free. The constant $-\frac{5}{32}$ has been added to the expression of p to ensure that the volume pressure normalization of ASPECT can be used in this benchmark (in other words, to ensure that the exact pressure has mean value zero and, consequently, can easily be compared with the numerically computed pressure). Following [BSA⁺13], the viscosity μ is given

by the smoothly varying function

$$\mu = \exp \{1 - \beta [x(1-x) + y(1-y) + z(1-z)]\}. \quad (57)$$

The maximum of this function is $\mu = e$, for example at $(x, y, z) = (0, 0, 0)$, and the minimum of this function is $\mu = \exp \left(1 - \frac{3\beta}{4}\right)$ at $(x, y, z) = (0.5, 0.5, 0.5)$. The viscosity ratio μ^* is then given by

$$\mu^* = \frac{\exp \left(1 - \frac{3\beta}{4}\right)}{\exp(1)} = \exp \left(\frac{-3\beta}{4}\right). \quad (58)$$

Hence, by varying β between 1 and 20, a difference of up to 7 orders of magnitude viscosity is obtained. β will be one of the parameters that can be selected in the input file that accompanies this benchmark.

The corresponding body force of the Stokes equation can then be computed by inserting this solution into the momentum equation,

$$\nabla p - \nabla \cdot (2\mu \epsilon(\mathbf{u})) = \rho \mathbf{g}. \quad (59)$$

Using equations (55), (56) and (57) in the momentum equation (59), the following expression for the body force $\rho \mathbf{g}$ can be found:

$$\begin{aligned} \rho \mathbf{g} = & \begin{pmatrix} yz + 3x^2y^3z \\ xz + 3x^3y^2z \\ xy + x^3y^3 \end{pmatrix} - \mu \begin{pmatrix} 2 + 6xy \\ 2 + 2x^2 + 2y^2 \\ -10yz \end{pmatrix} \\ & + (1 - 2x)\beta\mu \begin{pmatrix} 2 + 4x + 2y + 6x^2y \\ x + y + 2xy^2 + x^3 \\ -3z - 10xyz \end{pmatrix} + (1 - 2y)\beta\mu \begin{pmatrix} x + y + 2xy^2 + x^3 \\ 2 + 2x + 4y + 4x^2y \\ -3z - 5x^2z \end{pmatrix} \\ & + (1 - 2z)\beta\mu \begin{pmatrix} -3z - 10xyz \\ -3z - 5x^2z \\ -4 - 6x - 6y - 10x^2y \end{pmatrix} \end{aligned} \quad (60)$$

Assuming $\rho = 1$, the above expression translates into an expression for the gravity vector \mathbf{g} . This expression for the gravity (even though it is completely unphysical), has consequently been incorporated into the `BursteddeGravity` gravity model that is described in the `benchmarks/burstedde/burstedde.cc` file that accompanies this benchmark.

We will use the input file `benchmarks/burstedde/burstedde.prm` as input, which is very similar to the input file `benchmarks/inclusion/adaptive.prm` discussed above in Section 5.4.6. The major changes for the 3D polynomial Stokes benchmark are listed below:

```

subsection Solver parameters
  subsection Stokes solver parameters
    set Linear solver tolerance          = 1e-12
  end
end

# Boundary conditions
subsection Boundary velocity model
  set Prescribed velocity boundary indicators = left : BursteddeBoundary, \
                                                right : BursteddeBoundary, \
                                                front : BursteddeBoundary, \
                                                back : BursteddeBoundary, \
                                                bottom: BursteddeBoundary, \
                                                top   : BursteddeBoundary
end

```

```

subsection Material model
  set Model name = BursteddeMaterial
end

subsection Gravity model
  set Model name = BursteddeGravity
end

subsection Burstedde benchmark
  # Viscosity parameter is beta
  set Viscosity parameter          = 20
end

subsection Postprocess
  set List of postprocessors = visualization, velocity statistics, BursteddePostprocessor
end

```

The boundary conditions that are used are simply the velocities from equation (55) prescribed on each boundary. The viscosity parameter in the input file is β . Furthermore, in order to compute the velocity and pressure L_1 and L_2 norm, the postprocessor `BursteddePostprocessor` is used. Please note that the linear solver tolerance is set to a very small value (deviating from the default value), in order to ensure that the solver can solve the system accurately enough to make sure that the iteration error is smaller than the discretization error.

Expected analytical solutions at two locations are summarised in Table 6 and can be deduced from equations (55) and (56). Figure 77 shows that the analytical solution is indeed retrieved by the model.

Table 6: Analytical solutions

Quantity	$\mathbf{r} = (0, 0, 0)$	$\mathbf{r} = (1, 1, 1)$
p	-0.15625	1.84375
\mathbf{u}	(0, 0, 0)	(4, 4, -13)
$ \mathbf{u} $	0	14.177

The convergence of the numerical error of this benchmark has been analysed by playing with the mesh refinement level in the input file, and results can be found in Figure 78. The velocity shows cubic error convergence, while the pressure shows quadratic convergence in the L_1 and L_2 norms, as one would hope for using Q_2 elements for the velocity and Q_1 elements for the pressure.

5.4.8 The hollow sphere benchmark

This benchmark is based on Thieulot [In prep.] in which an analytical solution to the isoviscous incompressible Stokes equations is derived in a spherical shell geometry. The velocity and pressure fields are as follows:

$$v_r(r, \theta) = g(r) \cos \theta, \quad (61)$$

$$v_\theta(r, \theta) = f(r) \sin \theta, \quad (62)$$

$$v_\phi(r, \theta) = f(r) \sin \theta, \quad (63)$$

$$p(r, \theta) = h(r) \cos \theta, \quad (64)$$

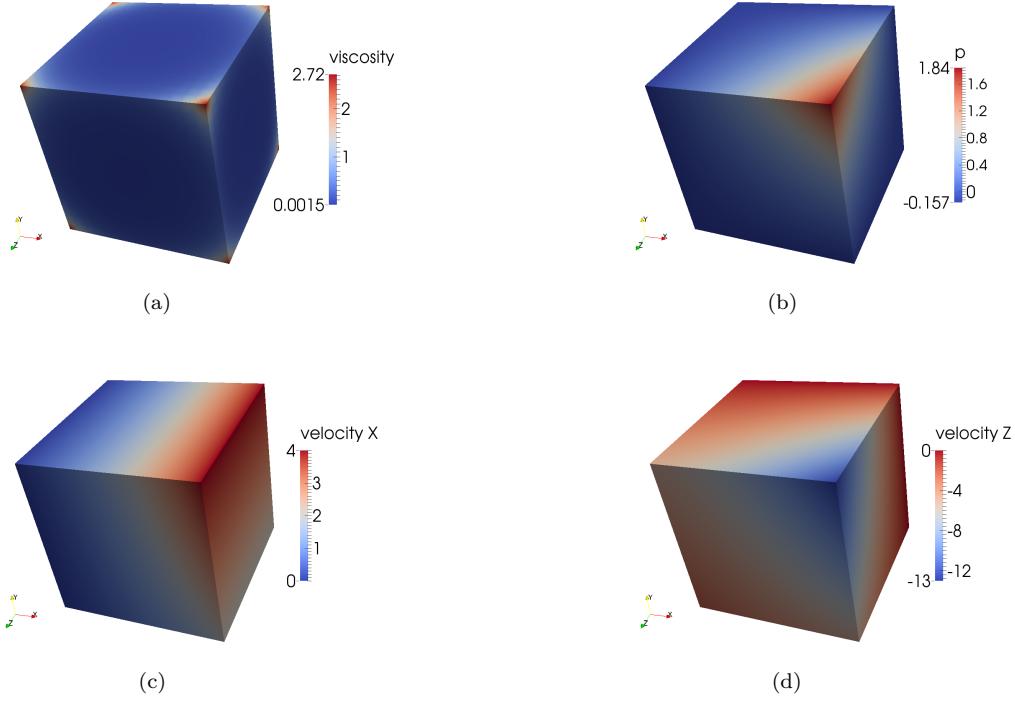


Figure 77: Burstedde benchmark: Results for the 3D polynomial Stokes benchmark, obtained with a resolution of 16×16 elements, with $\beta = 10$.

where

$$f(r) = \frac{\alpha}{r^2} + \beta r, \quad (65)$$

$$g(r) = -\frac{2}{r^2} \left(\alpha \ln r + \frac{\beta}{3} r^3 + \gamma \right), \quad (66)$$

$$h(r) = \frac{2\mu_0}{r} g(r), \quad (67)$$

with

$$\alpha = -\gamma \frac{R_2^3 - R_1^3}{R_2^3 \ln R_1 - R_1^3 \ln R_2}, \quad (68)$$

$$\beta = -3\gamma \frac{\ln R_2 - \ln R_1}{R_1^3 \ln R_2 - R_2^3 \ln R_1}. \quad (69)$$

These two parameters are chosen so that $v_r(R_1) = v_r(R_2) = 0$, i.e. the velocity is tangential to both inner and outer surfaces. The gravity vector is radial and of unit length, while the density is given by:

$$\rho(r, \theta) = \left(\frac{\alpha}{r^4} (8 \ln r - 6) + \frac{8\beta}{3r} + 8 \frac{\gamma}{r^4} \right) \cos \theta. \quad (70)$$

We set $R_1 = 0.5$, $R_2 = 1$ and $\gamma = -1$. The pressure is zero on both surfaces so that the surface pressure normalization is used. The boundary conditions that are used are simply the analytical velocity prescribed on both boundaries. The velocity and pressure fields are shown in Fig. 79.

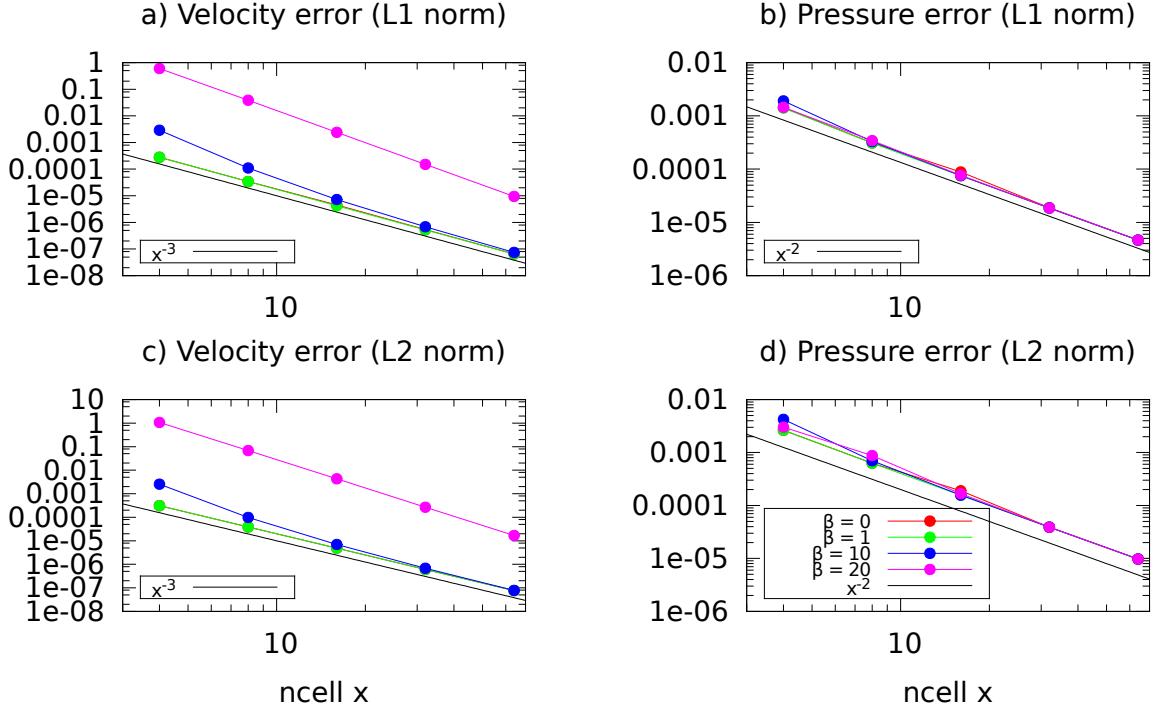


Figure 78: Burstedde benchmark: Error convergence for the 3D polynomial Stokes benchmark.

Fig. 80 shows the velocity and pressure errors in the L_2 -norm as a function of the mesh size h (taken in this case as the radial extent of the elements). As expected we recover a third-order convergence rate for the velocity and a second-order convergence rate for the pressure.

5.4.9 The 2D annulus benchmark

This benchmark is based on Thieulot & Puckett [In prep.] in which an analytical solution to the isoviscous incompressible Stokes equations is derived in an annulus geometry. The velocity and pressure fields are as follows:

$$v_r(r, \theta) = g(r)k \sin(k\theta), \quad (71)$$

$$v_\theta(r, \theta) = f(r) \cos(k\theta), \quad (72)$$

$$p(r, \theta) = kh(r) \sin(k\theta), \quad (73)$$

$$\rho(r, \theta) = \aleph(r)k \sin(k\theta), \quad (74)$$

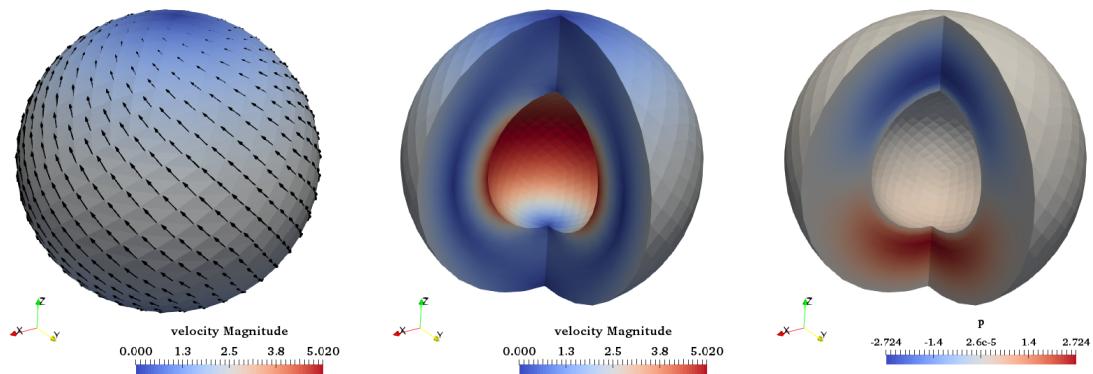


Figure 79: Velocity and pressure fields for the hollow sphere benchmark.

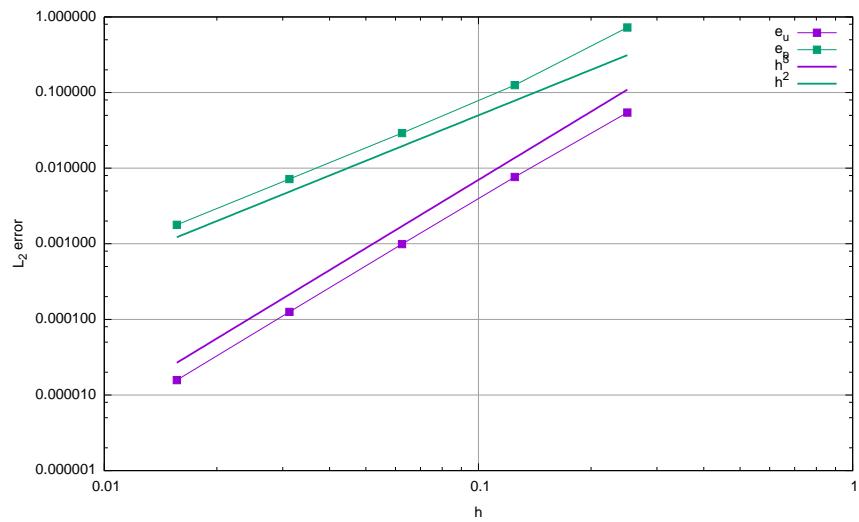


Figure 80: Velocity and pressure errors in the L_2 -norm as a function of the mesh size.

with

$$f(r) = Ar + B/r, \quad (75)$$

$$g(r) = \frac{A}{2}r + \frac{B}{r} \ln r + \frac{C}{r}, \quad (76)$$

$$h(r) = \frac{2g(r) - f(r)}{r}, \quad (77)$$

$$\aleph(r) = g'' - \frac{g'}{r} - \frac{g}{r^2}(k^2 - 1) + \frac{f}{r^2} + \frac{f'}{r}, \quad (78)$$

$$A = -C \frac{2(\ln R_1 - \ln R_2)}{R_2^2 \ln R_1 - R_1^2 \ln R_2}, \quad (79)$$

$$B = -C \frac{R_2^2 - R_1^2}{R_2^2 \ln R_1 - R_1^2 \ln R_2}. \quad (80)$$

The parameters A and B are chosen so that $v_r(R_1) = v_r(R_2) = 0$, i.e. the velocity is tangential to both inner and outer surfaces. The gravity vector is radial and of unit length.

The parameter k controls the number of convection cells present in the domain, as shown in Fig. 81.

In the present case, we set $R_1 = 1$, $R_2 = 2$ and $C = -1$. Fig. 82 shows the velocity and pressure errors in the L_2 -norm as a function of the mesh size h (taken in this case as the radial extent of the elements). As expected we recover a third-order convergence rate for the velocity and a second-order convergence rate for the pressure.

5.4.10 The “Stokes’ law” benchmark

This section was contributed by Juliane Dannberg.

Stokes’ law was derived by George Gabriel Stokes in 1851 and describes the frictional force a sphere with a density different than the surrounding fluid experiences in a laminar flowing viscous medium. A setup for testing this law is a sphere with the radius r falling in a highly viscous fluid with lower density. Due to its higher density the sphere is accelerated by the gravitational force. While the frictional force increases with the velocity of the falling particle, the buoyancy force remains constant. Thus, after some time the forces will be balanced and the settling velocity of the sphere v_s will remain constant:

$$\underbrace{6\pi \eta r v_s}_{\text{frictional force}} = \underbrace{4/3\pi r^3 \Delta\rho g}_{\text{buoyancy force}}, \quad (81)$$

where η is the dynamic viscosity of the fluid, $\Delta\rho$ is the density difference between sphere and fluid and g the gravitational acceleration. The resulting settling velocity is then given by

$$v_s = \frac{2}{9} \frac{\Delta\rho r^2 g}{\eta}. \quad (82)$$

Because we do not take into account inertia in our numerical computation, the falling particle will reach the constant settling velocity right after the first timestep.

For the setup of this benchmark, we chose the following parameters:

$$r = 200 \text{ km}$$

$$\Delta\rho = 100 \text{ kg/m}^3$$

$$\eta = 10^{22} \text{ Pa s}$$

$$g = 9.81 \text{ m/s}^2.$$

With these values, the exact value of sinking velocity is $v_s = 8.72 \cdot 10^{-10} \text{ m/s}$.

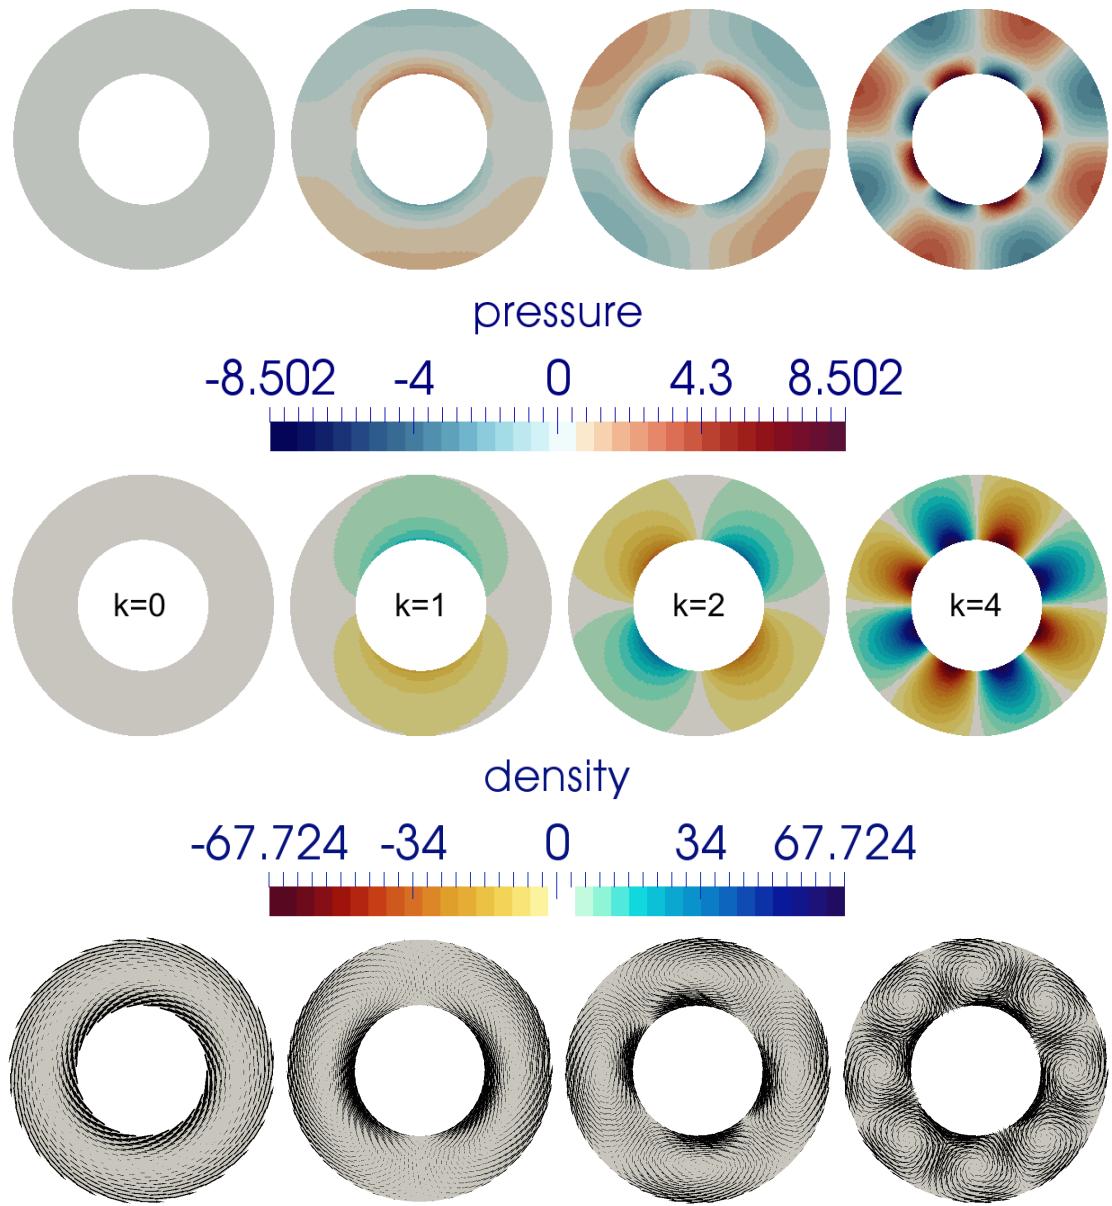


Figure 81: Pressure, density and velocity fields for $k = 0, 1, 2, 3$ for the 2D annulus benchmark.

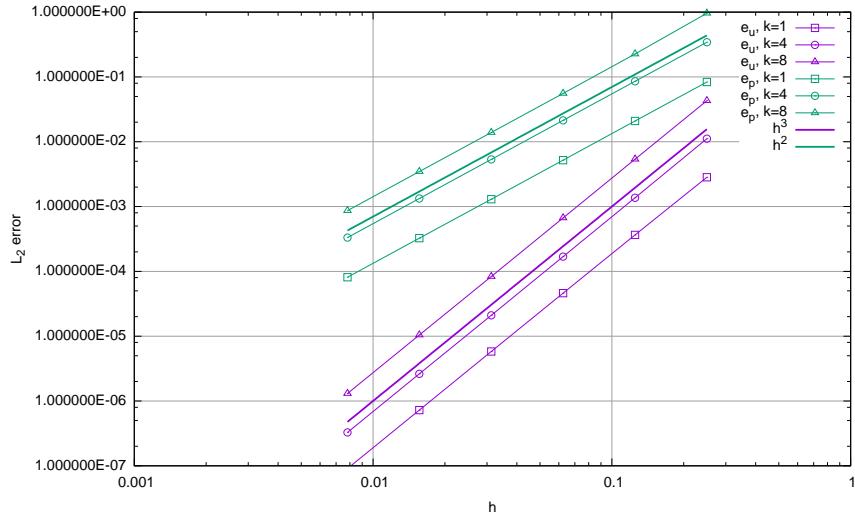


Figure 82: Velocity and pressure errors in the L_2 -norm as a function of the mesh size for the 2D annulus benchmark.

To run this benchmark, we need to set up an input file that describes the situation. In principle, what we need to do is to describe a spherical object with a density that is larger than the surrounding material. There are multiple ways of doing this. For example, we could simply set the initial temperature of the material in the sphere to a lower value, yielding a higher density with any of the common material models. Or, we could use ASPECT's facilities to advect along what are called "compositional fields" and make the density dependent on these fields.

We will go with the second approach and tell ASPECT to advect a single compositional field. The initial conditions for this field will be zero outside the sphere and one inside. We then need to also tell the material model to increase the density by $\Delta\rho = 100\text{kg m}^{-3}$ times the concentration of the compositional field. This can be done, like everything else, from the input file.

All of this setup is then described by the following input file. (You can find the input file to run this cookbook example in [cookbooks/stokes.prm](#). For your first runs you will probably want to reduce the number of mesh refinement steps to make things run more quickly.)

```
#####
# Global parameters
# We use a 3d setup. Since we are only interested
# in a steady state solution, we set the end time
# equal to the start time to force a single time
# step before the program terminates.

set Dimension          = 3
set Start time        = 0
set End time          = 0
set Use years in output instead of seconds = false
set Output directory   = output-stokes

#####
# Parameters describing the model
# The setup is a 3d box with edge length 2890000 in which
# all 6 sides have free slip boundary conditions. Because
# the temperature plays no role in this model we need not
```

```

# bother to describe temperature boundary conditions or
# the material parameters that pertain to the temperature.

subsection Geometry model
  set Model name = box

  subsection Box
    set X extent = 2890000
    set Y extent = 2890000
    set Z extent = 2890000
  end
end

subsection Boundary velocity model
  set Tangential velocity boundary indicators = left, right, front, back, bottom, top
end

subsection Material model
  set Model name = simple

  subsection Simple model
    set Reference density          = 3300
    set Viscosity                  = 1e22
  end
end

subsection Gravity model
  set Model name = vertical

  subsection Vertical
    set Magnitude = 9.81
  end
end

##### Parameters describing the temperature field
# As above, there is no need to set anything for the
# temperature boundary conditions.

subsection Boundary temperature model
  set List of model names = box
end

subsection Initial temperature model
  set Model name = function

  subsection Function
    set Function expression = 0
  end
end

##### Parameters describing the compositional field

```

```

# This, however, is the more important part: We need to describe
# the compositional field and its influence on the density
# function. The following blocks say that we want to
# advect a single compositional field and that we give it an
# initial value that is zero outside a sphere of radius
# r=200000m and centered at the point (p,p,p) with
# p=1445000 (which is half the diameter of the box) and one inside.
# The last block re-opens the material model and sets the
# density differential per unit change in compositional field to
# 100.

subsection Compositional fields
  set Number of fields = 1
end

subsection Initial composition model
  set Model name = function

  subsection Function
    set Variable names      = x,y,z
    set Function constants  = r=200000,p=1445000
    set Function expression = if(sqrt((x-p)*(x-p)+(y-p)*(y-p)+(z-p)*(z-p)) > r, 0, 1)
  end
end

subsection Material model
  subsection Simple model
    set Density differential for compositional field 1 = 100
  end
end

#####
##### Parameters describing the discretization
# The following parameters describe how often we want to refine
# the mesh globally and adaptively, what fraction of cells should
# be refined in each adaptive refinement step, and what refinement
# indicator to use when refining the mesh adaptively.

subsection Mesh refinement
  set Initial adaptive refinement      = 4
  set Initial global refinement       = 4
  set Refinement fraction            = 0.2
  set Strategy                      = velocity
end

#####
##### Parameters describing what to do with the solution
# The final section allows us to choose which postprocessors to
# run at the end of each time step. We select to generate graphical
# output that will consist of the primary variables (velocity, pressure,
# temperature and the compositional fields) as well as the density and
# viscosity. We also select to compute some statistics about the
# velocity field.

```

```

subsection Postprocess
  set List of postprocessors = visualization, velocity statistics

  subsection Visualization
    set List of output variables = density, viscosity
  end
end

```

Using this input file, let us try to evaluate the results of the current computations for the settling velocity of the sphere. You can visualize the output in different ways, one of it being ParaView and shown in Fig. 83 (an alternative is to use Visit as described in Section 4.4; 3d images of this simulation using Visit are shown in Fig. 84). Here, Paraview has the advantage that you can calculate the average velocity of the sphere using the following filters:

1. Threshold (Scalars: C_1, Lower Threshold 0.5, Upper Threshold 1),
2. Integrate Variables,
3. Cell Data to Point Data,
4. Calculator (use the formula $\text{sqrt}(\text{velocity_x}^2 + \text{velocity_y}^2 + \text{velocity_z}^2)/\text{Volume}$).

If you then look at the Calculator object in the Spreadsheet View, you can see the average sinking velocity of the sphere in the column “Result” and compare it to the theoretical value $v_s = 8.72 \cdot 10^{-10}$ m/s. In this case, the numerical result is $8.865 \cdot 10^{-10}$ m/s when you add a few more refinement steps to actually resolve the 3d flow field adequately. The “velocity statistics” postprocessor we have selected above also provides us with a maximal velocity that is on the same order of magnitude. The difference between the analytical and the numerical values can be explained by different at least the following three points: (i) In our case the sphere is viscous and not rigid as assumed in Stokes’ initial model, leading to a velocity field that varies inside the sphere rather than being constant. (ii) Stokes’ law is derived using an infinite domain but we have a finite box instead. (iii) The mesh may not yet fine enough to provide a fully converges solution. Nevertheless, the fact that we get a result that is accurate to less than 2% is a good indication that ASPECT implements the equations correctly.

5.4.11 Latent heat benchmark

This section was contributed by Juliane Dannberg.

The setup of this benchmark is taken from Schubert, Turcotte and Olson [STO01] (part 1, p. 194) and is illustrated in Fig. 85. It tests whether the latent heat production when material crosses a phase transition is calculated correctly according to the laws of thermodynamics. The material model defines two phases in the model domain with the phase transition approximately in the center. The material flows in from the top due to a prescribed downward velocity, and crosses the phase transition before it leaves the model domain at the bottom. As initial condition, the model uses a uniform temperature field, however, upon the phase change, latent heat is released. This leads to a characteristic temperature profile across the phase transition with a higher temperature in the bottom half of the domain. To compute it, we have to solve equation (3) or its reformulation (5). For steady-state one-dimensional downward flow with vertical velocity v_y , it simplifies to the following:

$$\rho C_p v_y \frac{\partial T}{\partial y} = \rho T \Delta S v_y \frac{\partial X}{\partial y} + \rho C_p \kappa \frac{\partial^2 T}{\partial y^2}.$$

Here, $\rho C_p \kappa = k$ with k the thermal conductivity and κ the thermal diffusivity. The first term on the right-hand side of the equation describes the latent heat produced at the phase transition: It is proportional to the temperature T , the entropy change ΔS across the phase transition divided by the specific heat capacity and the derivative of the phase function X . If the velocity is smaller than a critical value, and under the

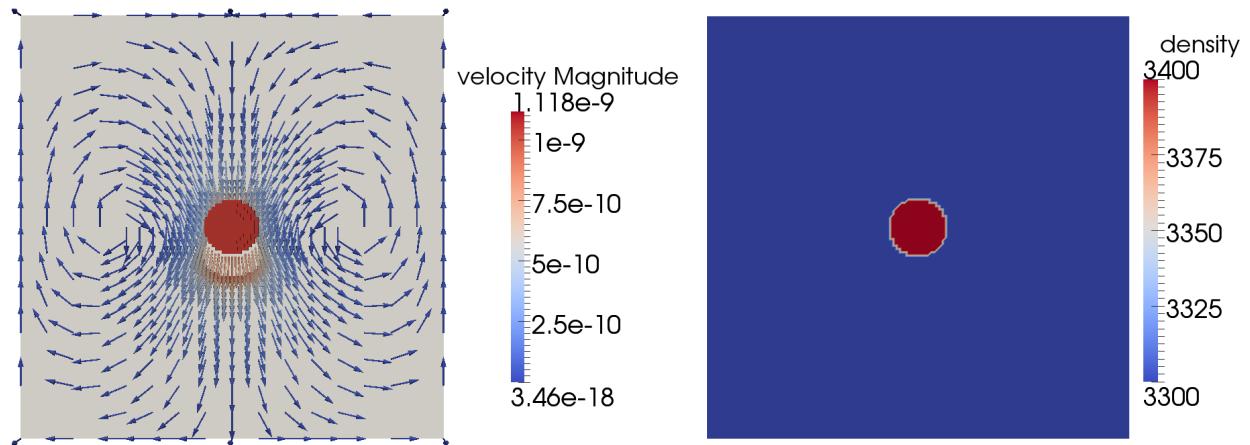


Figure 83: *Stokes benchmark*. Both figures show only a 2D slice of the three-dimensional model. Left: The compositional field and overlaid to it some velocity vectors. The composition is 1 inside a sphere with the radius of 200 km and 0 outside of this sphere. As the velocity vectors show, the sphere sinks in the viscous medium. Right: The density distribution of the model. The compositional density contrast of 100 kg/m^3 leads to a higher density inside of the sphere.

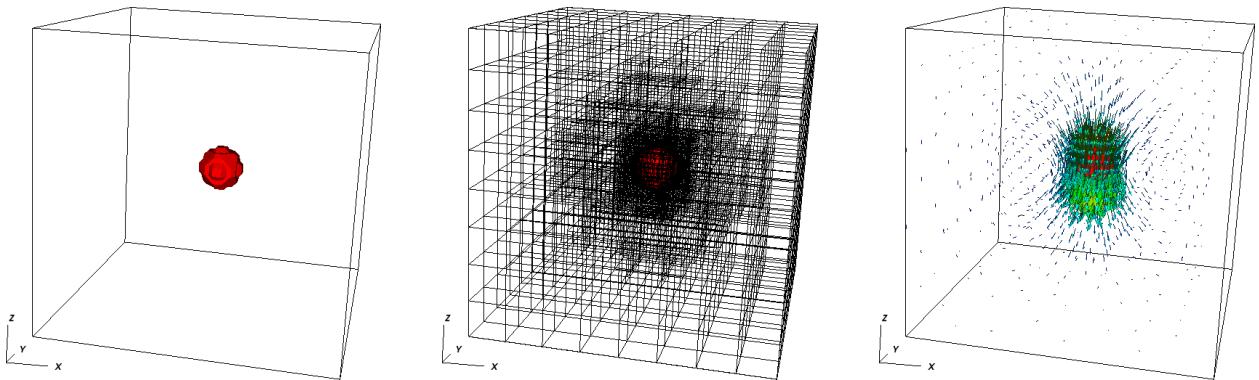


Figure 84: *Stokes benchmark*. Three-dimensional views of the compositional field (left), the adaptively refined mesh (center) and the resulting velocity field (right).

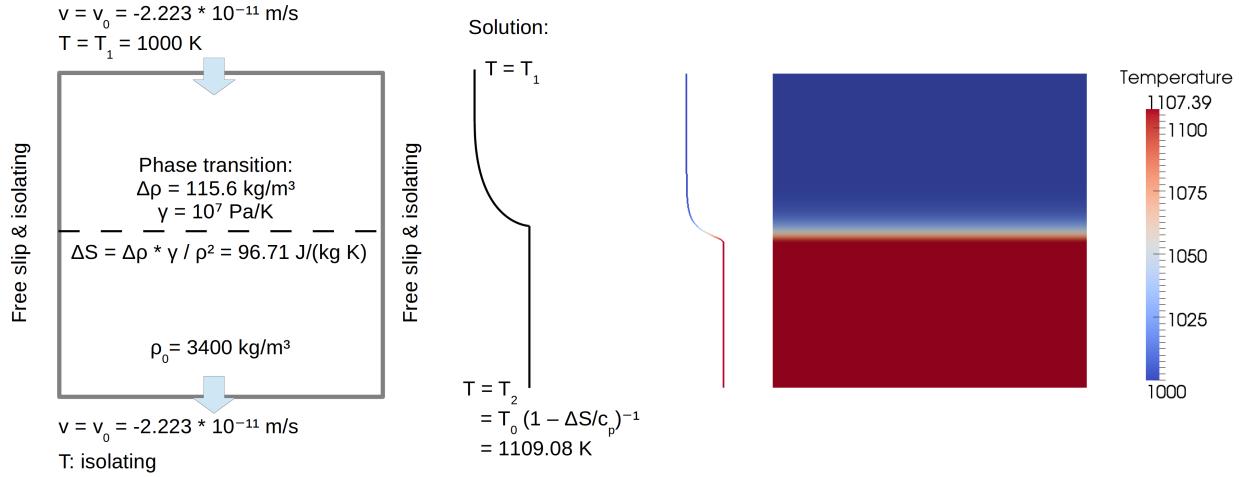


Figure 85: *Latent heat benchmark*. Both figures show the 2D box model domain. Left: Setup of the benchmark together with a sketch of the expected temperature profile across the phase transition. The dashed line marks the phase transition. Material flows in with a prescribed temperature and velocity at the top, crosses the phase transition in the center and flows out at the bottom. The predicted bottom temperature is $T_2 = 1109.08 \text{ K}$. Right: Temperature distribution of the model together with the associated temperature profile across the phase transition. The modelled bottom temperature is $T_2 = 1107.39 \text{ K}$.

assumption of a discontinuous phase transition (i.e. with a step function as phase function), this latent heating term will be zero everywhere except for the one point y_{tr} where the phase transition takes place. This means, we have a region above the phase transition with only phase 1, and below a certain depth a jump to a region with only phase 2. Inside of these one-phase regions, we can solve the equation above (using the boundary conditions $T = T_1$ for $y \rightarrow \infty$ and $T = T_2$ for $y \rightarrow -\infty$) and get

$$T(y) = \begin{cases} T_1 + (T_2 - T_1)e^{\frac{v_y(y-y_{tr})}{\kappa}}, & y > y_{tr} \\ T_2, & y < y_{tr} \end{cases}$$

While it is not entirely obvious while this equation for $T(y)$ should be correct (in particular why it should be asymmetric), it is not difficult to verify that it indeed satisfies the equation stated above for both $y < y_{tr}$ and $y > y_{tr}$. Furthermore, it indeed satisfies the jump condition we get by evaluating the equation at $y = y_{tr}$. Indeed, the jump condition can be reinterpreted as a balance of heat conduction: We know the amount of heat that is produced at the phase boundary, and as we consider only steady-state, the same amount of heat is conducted upwards from the transition:

$$\underbrace{\rho v_y T \Delta S}_{\text{latent heat release}} = \underbrace{\frac{\kappa}{\rho_0 C_p} \frac{\partial T}{\partial y} \Big|_{y=y_{tr}-}}_{\text{heat conduction}} = \frac{v_y}{\rho_0 C_p} (T_2 - T_1)$$

In contrast to [STO01], we also consider the density change $\Delta\rho$ across the phase transition: While the heat conduction takes place above the transition and the density can be assumed as $\rho = \rho_0 = \text{const.}$, the latent heat is released directly at the phase transition. Thus, we assume an average density $\rho = \rho_0 + 0.5\Delta\rho$ for the left side of the equation. Rearranging this equation gives

$$T_2 = \frac{T_1}{1 - (1 + \frac{\Delta\rho}{2\rho_0})\frac{\Delta S}{C_p}}$$

In addition, we have tested the approach exactly as it is described in [STO01] by setting the entropy change to a specific value and in spite of that using a constant density. However, this is physically inconsistent, as the entropy change is proportional to the density change across the phase transition. With this method, we could reproduce the analytic results from [STO01].

The exact values of the parameters used for this benchmark can be found in Fig. 85. They result in a predicted value of $T_2 = 1109.08\text{ K}$ for the temperature in the bottom half of the model, and we will demonstrate below that we can match this value in our numerical computations. However, it is not as simple as suggested above. In actual numerical computations, we can not exactly reproduce the behavior of Dirac delta functions as would result from taking the derivative $\frac{\partial X}{\partial y}$ of a discontinuous function $X(y)$. Rather, we have to model $X(y)$ as a function that has a smooth transition from one value to another, over a depth region of a certain width. In the material model plugin we will use below, this depth is an input parameter and we will play with it in the numerical results shown after the input file.

To run this benchmark, we need to set up an input file that describes the situation. In principle, what we need to do is to describe the position and entropy change of the phase transition in addition to the previously outlined boundary and initial conditions. For this purpose, we use the “latent heat” material model that allows us to set the density change $\Delta\rho$ and Clapeyron slope γ (which together determine the entropy change via $\Delta S = \gamma \frac{\Delta\rho}{\rho^2}$) as well as the depth of the phase transition as input parameters.

All of this setup is then described by the input file `cookbooks/latent-heat.prm` that models flow in a box of 10^6 meters of height and width, and a fixed downward velocity. The following section shows the central part of this file:

```

subsection Material model
  set Model name = latent heat
  subsection Latent heat

    # The change of density across the phase transition. Together with the
    # Clapeyron slope, this is what determines the entropy change.
    set Phase transition density jumps           = 115.6
    set Corresponding phase for density jump     = 0

    # If the temperature is equal to the phase transition temperature, the
    # phase transition will occur at the phase transition depth. However,
    # if the temperature deviates from this value, the Clapeyron slope
    # determines how much the pressure (and depth) of the phase boundary
    # changes. Here, the phase transition will be in the middle of the box
    # for T=T1.
    set Phase transition depths                  = 500000
    set Phase transition temperatures            = 1000
    set Phase transition Clapeyron slopes        = 1e7

    # We set the width of the phase transition to 5 km. You may want to
    # change this parameter to see how latent heating depends on the width
    # of the phase transition.
    set Phase transition widths                 = 5000

    set Reference density                      = 3400
    set Reference specific heat                = 1000
    set Reference temperature                 = 1000
    set Thermal conductivity                  = 2.38

```

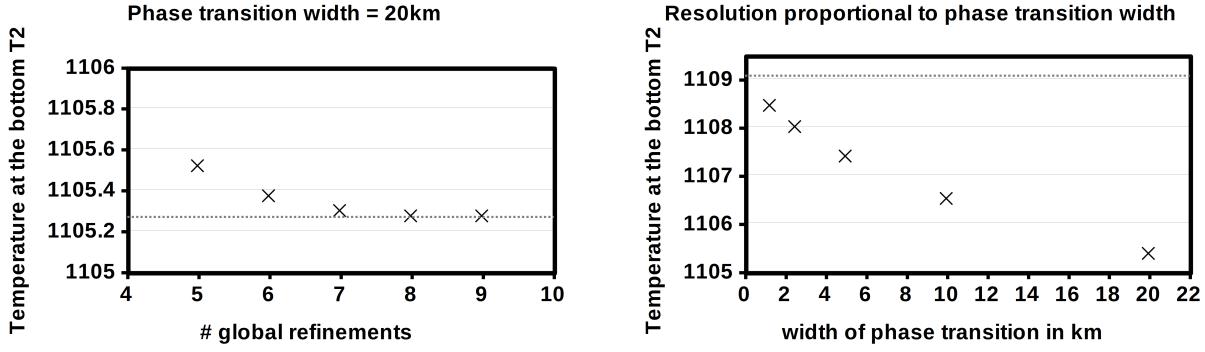


Figure 86: *Results of the latent heat benchmark.* Both figures show the modelled temperature T_2 at the bottom of the model domain. Left: T_2 in dependence of resolution using a constant phase transition width of 20 km. With an increasing number of global refinements of the mesh, the bottom temperature converges against a value of $T_2 = 1105.27$ K. Right: T_2 in dependence of phase transition width. The model resolution is chosen proportional to the phase transition width, starting with 5 global refinements for a width of 20 km. With decreasing phase transition width, T_2 approaches the theoretical value of 1109.08 K

```
# We set the thermal expansion and the compressibility to zero, so that
# all temperature (and density) changes are caused by advection, diffusion
# and latent heating.
set Thermal expansion coefficient           = 0.0
set Compressibility                         = 0.0

# Viscosity is constant.
set Thermal viscosity exponent             = 0.0
set Viscosity                             = 8.44e21
set Viscosity prefactors                  = 1.0, 1.0
set Composition viscosity prefactor       = 1.0
end
end
```

The complete input file referenced above also sets the number of mesh refinement steps. For your first runs you will probably want to reduce the number of mesh refinement steps to make things run more quickly. Later on, you might also want to change the phase transition width to look how this influences the result.

Using this input file, let us try to evaluate the results of the current computations. We note that it takes some time for the model to reach a steady state and only then does the bottom temperature reach the theoretical value. Therefore, we use the last output step to compare predicted and computed values. You can visualize the output in different ways, one of it being ParaView and shown in Fig. 85 on the right side (an alternative is to use Visit as described in Section 4.4). In ParaView, you can plot the temperature profile using the filter “Plot Over Line” (Point1: 500000,0,0; Point2: 500000,1000000,0, then go to the “Display” tab and select “T” as only variable in the “Line series” section) or “Calculator” (as seen in Fig. 85). In Fig. 86 (left) we can see that with increasing resolution, the value for the bottom temperature converges to a value of $T_2 = 1105.27$ K.

However, this is not what the analytic solution predicted. The reason for this difference is the width of the phase transition with which we smooth out the Dirac delta function that results from differentiating the $X(y)$ we would have liked to use in an ideal world. (In reality, however, for the Earth’s mantle we also expect phase transitions that are distributed over a certain depth range and so the smoothed out approach may not be a bad approximation.) Of course, the results shown above result from an analytical approach that is

only correct if the phase transition is discontinuous and constrained to one specific depth $y = y_{tr}$. Instead, we chose a hyperbolic tangent as our phase function. Moreover, Fig. 86 (right) illustrates what happens to the temperature at the bottom when we vary the width of the phase transition: The smaller the width, the closer the temperature gets to the predicted value of $T_2 = 1109.08$ K, demonstrating that we converge to the correct solution.

5.4.12 The 2D cylindrical shell benchmarks by Davies et al.

This section was contributed by William Durkin and Wolfgang Bangerth.

All of the benchmarks presented so far take place in a Cartesian domain. Davies et al. describe a benchmark (in a paper that is currently still being written) for a 2D spherical Earth that is nondimensionalized such that

$$\begin{aligned} r_{\min} &= 1.22 & T|_{r_{\min}} &= 1 \\ r_{\max} &= 2.22 & T|_{r_{\max}} &= 0 \end{aligned}$$

The benchmark is run for a series of approximations (Boussinesq, Extended Boussinesq, Truncated Anelastic Liquid, and Anelastic Liquid), and temperature, velocity, and heat flux calculations are compared with the results of other mantle modeling programs. ASPECT will output all of these values directly except for the Nusselt number, which we must calculate ourselves from the heat fluxes that ASPECT can compute. The Nusselt number of the top and bottom surfaces, Nu_T and Nu_B , respectively, are defined by the authors of the benchmarks as

$$Nu_T = \frac{\ln(f)}{2\pi r_{\max}(1-f)} \int_0^{2\pi} \frac{\partial T}{\partial r} d\theta \quad (83)$$

and

$$Nu_B = \frac{f \ln(f)}{2\pi r_{\min}(1-f)} \int_0^{2\pi} \frac{\partial T}{\partial r} d\theta$$

where f is the ratio $\frac{r_{\min}}{r_{\max}}$.

We can put this in terms of heat flux

$$q_r = -k \frac{\partial T}{\partial r}$$

through the inner and outer surfaces, where q_r is heat flux in the radial direction. Let Q be the total heat that flows through a surface,

$$Q = \int_0^{2\pi} q_r d\theta,$$

then (83) becomes

$$Nu_T = \frac{-Q_T \ln(f)}{2\pi r_{\max}(1-f)k}$$

and similarly

$$Nu_B = \frac{-Q_B f \ln(f)}{2\pi r_{\min}(1-f)k}.$$

Q_T and Q_B are heat fluxes that ASPECT can readily compute through the `heat flux statistics` post-processor (see Section A.118). For further details on the nondimensionalization and equations used for each approximation, refer to Davies et al.

The series of benchmarks is then defined by a number of cases relating to the exact equations chosen to model the fluid. We will discuss these in the following.

Case 1.1: BA_Ra104_Iso_ZS. This case is run with the following settings:

- Boussinesq Approximation
- Boundary Condition: Zero-Slip
- Rayleigh Number = 10^4
- Initial Conditions: $D = 0, O = 4$
- $\eta(T) = 1$

where D and O refer to the degree and order of a spherical harmonic that describes the initial temperature. While the initial conditions matter, what is important here though is that the system evolves to four convective cells since we are only interested in the long term, steady state behavior.

The model is relatively straightforward to set up, basing the input file on that discussed in Section 5.3.1. The full input file can be found at [benchmarks/davies_et_al/case-1.1.prm](#), with the interesting parts excerpted as follows:

```
#####
# Parameters describing the model

subsection Geometry model
  set Model name = spherical shell
  subsection Spherical shell
    set Inner radius   = 1.22
    set Opening angle = 360
    set Outer radius  = 2.22
  end
end

# [...]

subsection Material model
  set Model name = simple
  subsection Simple model
    set Reference density      = 1
    set Reference specific heat = 1.
    set Reference temperature  = 0
    set Thermal conductivity   = 1
    set Thermal expansion coefficient = 1e-6
    set Viscosity              = 1
  end
end

#####
# Parameters describing the temperature field
# Angular mode is set to 4 in order to match the number of
# convective cells reported by Davies et al.

subsection Initial temperature model
  set Model name = spherical hexagonal perturbation
  subsection Spherical hexagonal perturbation
    set Angular mode       = 4
    set Rotation offset    = 0
  end
end
```

```

##### Prescribe the Rayleigh number as g*alpha
# Here, Ra = 10^4 and alpha was chosen as 10^-6 above.
subsection Gravity model
    set Model name = radial constant
    subsection Radial constant
        set Magnitude = 1e10
    end
end

# [...]

```

We use the same trick here as in Section 5.2.1 to produce a model in which the density $\rho(T)$ in the temperature equation (3) is almost constant (namely, by choosing a very small thermal expansion coefficient) as required by the benchmark, and instead prescribe the desired Rayleigh number by choosing a correspondingly large gravity.

Results for this and the other cases are shown below.

Case 2.1: BA_Ra104_Iso_FS. Case 2.1 uses the following setup, differing only in the boundary conditions:

- Boussinesq Approximation
- Boundary Condition: Free-Slip
- Rayleigh Number = 10^4
- Initial Conditions: $D = 0, O = 4$
- $\eta(T) = 1$

As a consequence of the free slip boundary conditions, any solid body rotation of the entire system satisfies the Stokes equations with their boundary conditions. In other words, the solution of the problem is not unique: given a solution, adding a solid body rotation yields another solution. We select arbitrarily the one that has no net rotation (see Section A.117). The section in the input file that is relevant is then as follows (the full input file resides at [benchmarks/davies_et_al/case-2.1.prm](#)):

```

subsection Nullspace removal
    set Remove nullspace                  = net rotation
end

subsection Boundary temperature model
    set Fixed temperature boundary indicators = 0,1
end

subsection Boundary velocity model
    set Tangential velocity boundary indicators = 0,1
end

```

Again, results are shown below.

Case 2.2: BA_Ra105_Iso_FS. Case 2.2 is described as follows:

- Boussinesq Approximation
- Boundary Condition: Free-Slip

- Rayleigh Number = 10^5
- Initial Conditions: Final conditions of case 2.1 (BA_Ra104_Iso_FS)
- $\eta(T) = 1$

In other words, we have an increased Rayleigh number and begin with the final steady state of case 2.1. To start the model where case 2.1 left off, the input file of case 2.1, [benchmarks/davies_et_al/case-2.1.prm](#), instructs ASPECT to checkpoint itself every few time steps (see Section 4.5). If case 2.2 uses the same output directory, we can then resume the computations from this checkpoint with an input file that prescribes a different Rayleigh number and a later input time:

```
#####
# Global parameters
# Case 2.2 begins with the final steady state solution of Case 2.1
# "Resume computation" must be set to true, and "Output directory" must
# point to the folder that contains the results of Case 2.1.

set CFL number          = 10
set End time             = 3
set Output directory     = output
set Resume computation   = true
```

We increase the Rayleigh number to 10^5 by increasing the magnitude of gravity in the input file. The full script for case 2.2 is located in [benchmarks/davies_et_al/case-2.2.prm](#)

Case 2.3: BA_Ra103_vv_FS. Case 2.3 is a variation on the previous one:

- Boussinesq Approximation
- Boundary Condition: Free-Slip
- Rayleigh Number = 10^3
- Initial Conditions: Final conditions of case 2.1 (BA_Ra104_Iso_FS)
- $\eta(T) = 1000^{-T}$

The Rayleigh number is smaller here (and is selected using the gravity parameter in the input file, as before), but the more important change is that the viscosity is now a function of temperature. At the time of writing, there is no material model that would implement such a viscosity, so we create a plugin that does so for us (see Sections 6 and 6.2 in general, and Section 6.3.1 for material models in particular). The code for it is located in [benchmarks/davies_et_al/case-2.3-plugin/VoT.cc](#) (where “VoT” is short for “viscosity as a function of temperature”) and is essentially a copy of the `simpler` material model. The primary change compared to the `simpler` material model is the line about the viscosity in the following function:

```
template <int dim>
void
VoT<dim>::
evaluate(const typename Interface<dim>::MaterialModelInputs &in,
         typename Interface<dim>::MaterialModelOutputs &out) const
{
    for (unsigned int i=0; i<in.position.size(); ++i)
    {
        out.viscosities[i] = eta*std::pow(1000, (-in.temperature[i]));
        out.densities[i] = reference_rho * (1.0 - thermal_alpha * (in.temperature[i] - reference_T));
        out.thermal_expansion_coefficients[i] = thermal_alpha;
```

```

        out.specific_heat[i] = reference_specific_heat;
        out.thermal_conductivities[i] = k_value;
        out.compressibilities[i] = 0.0;
    }
}

```

Using the method described in Sections 5.4.1 and 6.2, and the files in the `benchmarks/davies_et_al/case-2.3-plugin`, we can compile our new material model into a shared library that we can then reference from the input file. The complete input file for case 2.3 is located in `benchmarks/davies_et_al/case-2.3.prm` and contains among others the following parts:

```

set Additional shared libraries          = ./case-2.3-plugin/libVoT.so

subsection Material model
    set Model name = VoT

    subsection VoT model
        set Reference density      = 1
        set Reference specific heat = 1.
        set Reference temperature   = 0
        set Thermal conductivity   = 1
        set Thermal expansion coefficient = 1e-5
        set Viscosity               = 1
    end
end

```

Results. In the following, let us discuss some of the results of the benchmark setups discussed above. First, the final steady state temperature fields are shown in Fig. 87. It is immediately obvious how the different Rayleigh numbers affect the width of the plumes. If one imagines a setup with constant gravity, constant inner and outer temperatures and constant thermal expansion coefficient (this is not how we describe it in the input files, but we could have done so and it is closer to how we intuit about fluids than adjusting the gravity), then the Rayleigh number is inversely proportional to the viscosity – and it is immediately clear that larger Rayleigh numbers (corresponding to lower viscosities) then lead to thinner plumes. This is nicely reflected in the visualizations.

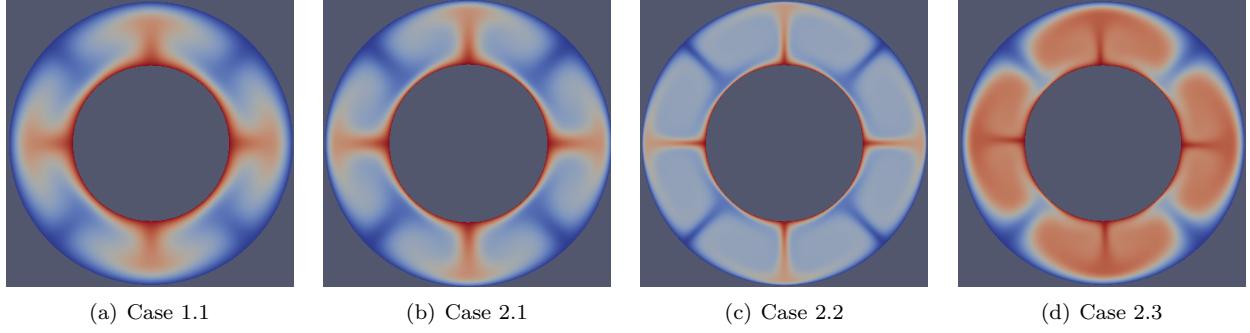


Figure 87: Davies et al. benchmarks: Final steady state temperature fields for the 2D cylindrical benchmark cases.

Secondly, Fig. 88 shows the root mean square velocity as a function of time for the various cases. It is obvious that they all converge to steady state solutions. However, there is an initial transient stage and, in cases 2.2 and 2.3, a sudden jolt to the system at the time where we switch from the model used to compute up to time $t = 2$ to the different models used after that.

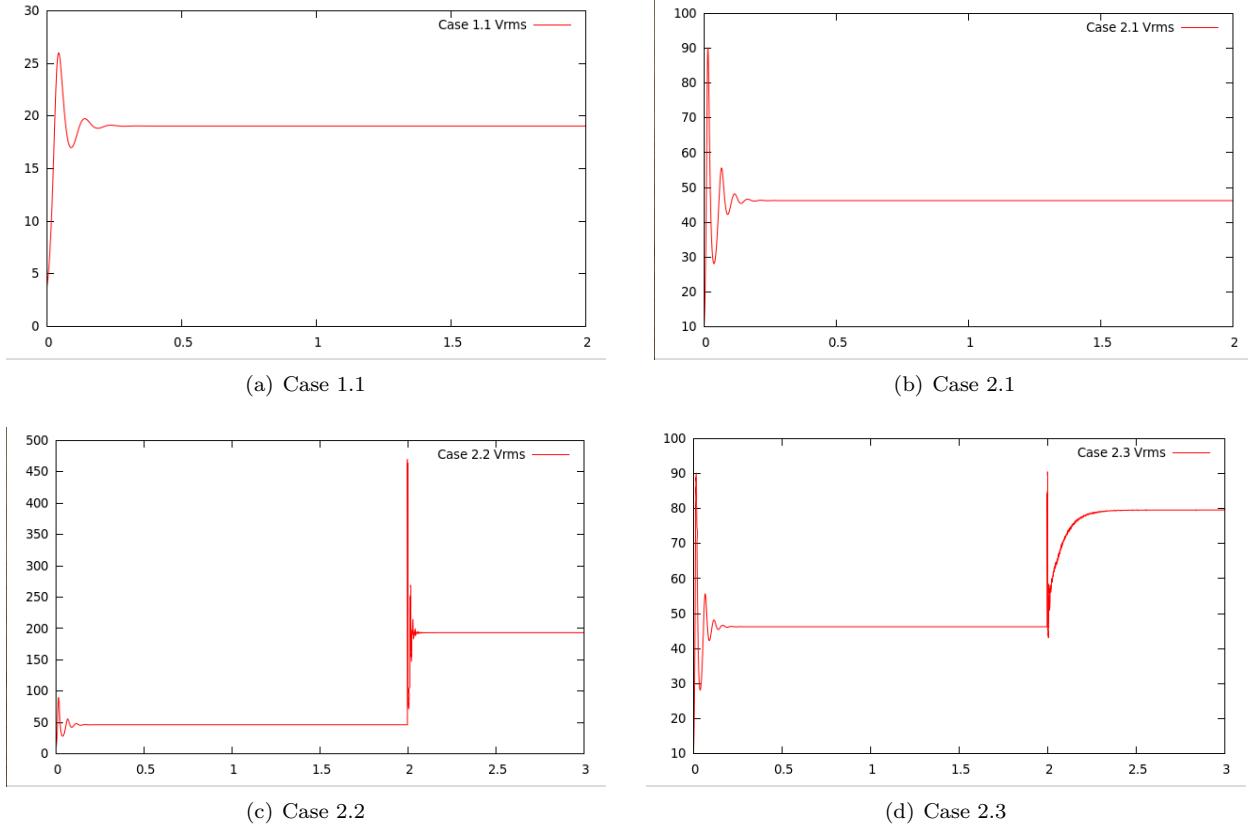


Figure 88: Davies et al. benchmarks: V_{rms} for 2D Cylindrical Cases. Large jumps occur when transitioning from case 2.1 to cases 2.2 and 2.3 due to the instantaneous change of parameter settings.

These runs also produce quantitative data that will be published along with the concise descriptions of the benchmarks and a comparison with other codes. In particular, some of the criteria listed above to judge the accuracy of results are listed in Table 7.³⁶

5.4.13 The Cramer et al. benchmarks

This section was contributed by Ian Rose.

³⁶The input files available in the `benchmarks/davies_et_al` directory use 5 global refinements in order to provide cases that can be run without excessive trouble on a normal computer. However, this is not enough to achieve reasonable accuracy and both the data shown below and the data submitted to the benchmarking effort uses 7 global refinement steps, corresponding to a mesh with 1536 cells in tangential and 128 cells in radial direction. Computing on such meshes is not cheap, as it leads to a problem size of more than 2.5 million unknowns. It is best done using a parallel computation.

Case	$\langle T \rangle$	Nu_T	Nu_B	V_{rms}
1.1	0.403	2.464	2.468	19.053
2.1	0.382	4.7000	4.706	46.244
2.2	0.382	9.548	9.584	193.371
2.3	0.582	5.102	5.121	79.632

Table 7: Davies et al. benchmarks: Numerical results for some of the output quantities required by the benchmarks and the various cases considered.



Figure 89: Setup for the topography relaxation benchmark. The box is 2800 km wide and 700 km high, with a 100 km lid on top. The lid has a viscosity of 10^{23} Pa s, while the mantle has a viscosity of 10^{21} Pa s. The sides are free slip, the bottom is no slip, and the top is a free surface. Both the lid and the mantle have a density of 3300 kg/m^3 , and gravity is 10 m/s^2 . There is a 7 km sinusoidal initial topography on the free surface.

This section follows the two free surface benchmarks described by Crameri et al. [CSG⁺12].

Case 1: Relaxation of topography. The first benchmark involves a high viscosity lid sitting on top of a lower viscosity mantle. There is an initial sinusoidal topography which is then allowed to relax. This benchmark has a semi-analytical solution (which is exact for infinitesimally small topography). Details for the benchmark setup are in Figure 89.

The complete parameter file for this benchmark can be found in [benchmarks/crameri_et_al/case_1/crameri_benchmark_1.prm](#), the most relevant parts of which are excerpted here:

```

set CFL number          = 0.01

set Additional shared libraries = ./libcrameri_benchmark_1.so

subsection Geometry model
    set Model name = rebound box
    subsection Rebound Box
        set Order = 1
        set Amplitude = 7.e3
    end
    subsection Box
        set X extent = 28.e5
        set Y extent = 7.e5
        set X repetitions = 300
        set Y repetitions = 75
    end
end

```

In particular, this benchmark uses a custom geometry model to set the initial geometry. This geometry model, called “ReboundBox”, is based on the Box geometry model. It generates a domain in using the same parameters as Box, but then displaces all the nodes vertically with a sinusoidal perturbation, where the magnitude and order of that perturbation are specified in the ReboundBox subsection.

The characteristic timescales of topography relaxation are significantly smaller than those of mantle convection. Taking timesteps larger than this relaxation timescale tends to cause sloshing instabilities, which are described further in Section 2.12. Some sort of stabilization is required to take large timesteps. In this benchmark, however, we are interested in the relaxation timescale, so we are free to take very small

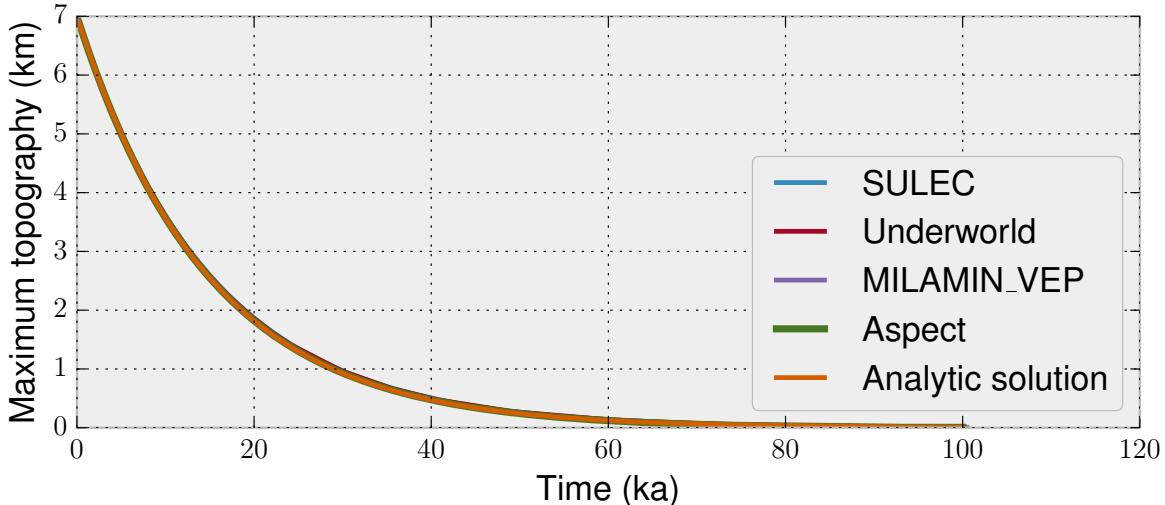


Figure 90: Results for the topography relaxation benchmark, showing maximum topography versus time. Over about 100 ka the topography completely disappears. The results of four free surface codes, as well as the semi-analytic solution, are nearly identical.

timesteps (in this case, 0.01 times the CFL number). As can be seen in Figure 90, the results of all the codes which are included in this comparison are basically indistinguishable.

Case 2: Dynamic topography. Case two is more complicated. Unlike the case one, it occurs over mantle convection timescales. In this benchmark there is the same high viscosity lid over a lower viscosity mantle. However, now there is a blob of buoyant material rising in the center of the domain, causing dynamic topography at the surface. The details for the setup are in the caption of Figure 91.

Case two requires higher resolution and longer time integrations than case one. The benchmark is over 20 million years and builds dynamic topography of ~ 800 meters.

Again, we excerpt the most relevant parts of the parameter file for this benchmark, with the full thing available in [benchmarks/crameri_et_al/case_2/crameri_benchmark_2.prm](#). Here we use the “Multicomponent” material model, which allows us to easily set up a number of compositional fields with different material properties. The first compositional field corresponds to background mantle, the second corresponds to the rising blob, and the third corresponds to the viscous lid.

Furthermore, the results of this benchmark are sensitive to the mesh refinement and timestepping parameters. Here we have nine refinement levels, and refine according to density and the compositional fields.

```

set CFL number          = 0.1

subsection Material model
  set Model name = multicomponent
  subsection Multicomponent
    set Densities = 3300, 3200, 3300
    set Viscosities = 1.e21, 1.e20, 1.e23
    set Viscosity averaging scheme = harmonic
  end
end

subsection Mesh refinement

```

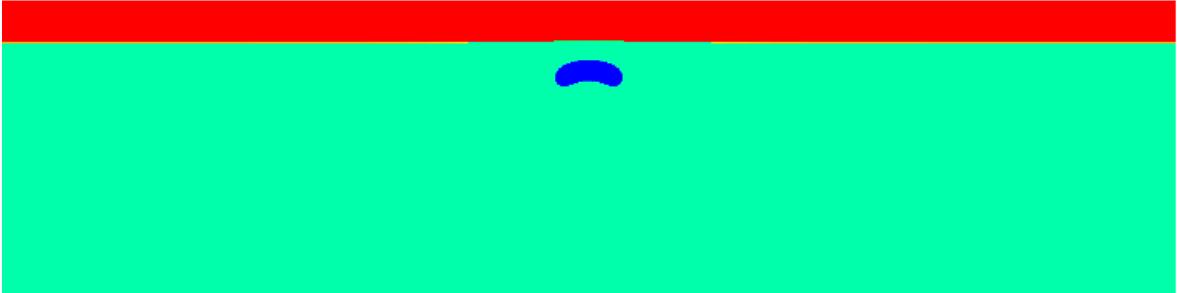


Figure 91: Setup for the dynamic topography benchmark. Again, the domain is 2800 km wide and 700 km high. A 100 km thick lid with viscosity 10^{23} overlies a mantle with viscosity 10^{21} . Both the lid and the mantle have a density of 3300 kg/m^3 . A blob with diameter 100 km lies 300 km from the bottom of the domain. The blob has a density of 3200 kg/m^3 and a viscosity of 10^{20} Pa s .

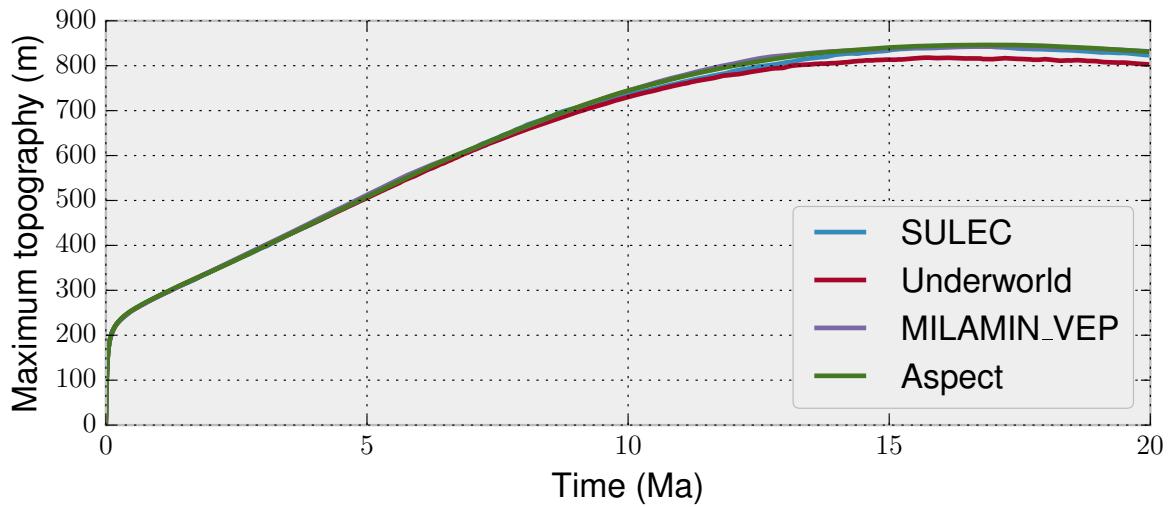


Figure 92: Evolution of topography for the dynamic topography benchmark. The maximum topography is shown as a function of time, for ASPECT as well as for several other codes participating in the benchmark. This benchmark shows considerably more scatter between the codes.

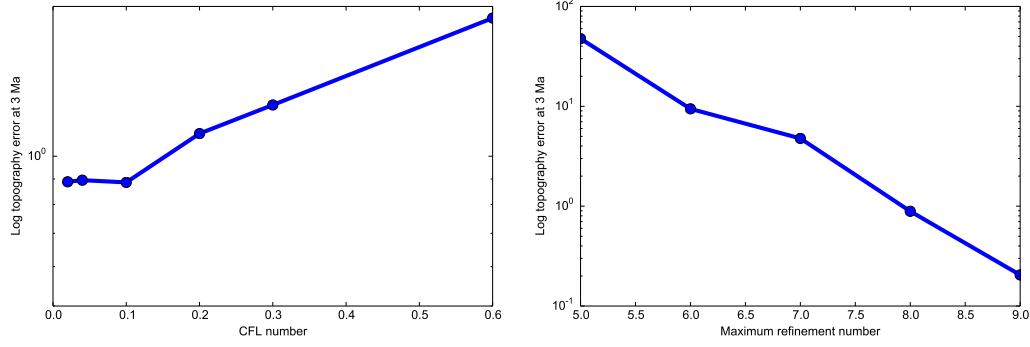


Figure 93: *Convergence for case two. Left: Logarithm of the error with decreasing CFL number. As the CFL number decreases, the error gets smaller. However, once it reaches a value of ~ 0.1 , there stops being much improvement in accuracy. Right: Logarithm of the error with increasing maximum mesh resolution. As the resolution increases, so does the accuracy.*

```

set Additional refinement times      =
set Initial adaptive refinement    = 4
set Initial global refinement      = 5
set Refinement fraction           = 0.3
set Coarsening fraction           = 0.0
set Strategy                      = density,composition
set Refinement criteria merge operation = plus
set Time steps between mesh refinement = 5
end

```

Unlike the first benchmark, for case two there is no (semi) analytical solution to compare against. Furthermore, the time integration for this benchmark is much longer, allowing for errors to accumulate. As such, there is considerably more scatter between the participating codes. ASPECT does, however, fall within the range of the other results, and the curve is somewhat less wiggly. The results for maximum topography versus time are shown in 92

The precise values for topography at a given time are quite dependent on the resolution and timestepping parameters. Following [CSG⁺¹²] we investigate the convergence of the maximum topography at 3 Ma as a function of CFL number and mesh resolution. The results are shown in figure 93.

We find that at 3 Ma ASPECT converges to a maximum topography of ~ 396 meters. This is slightly different from what MILAMIN_VEP reported as its convergent value in [CSG⁺¹²], but still well within the range of variation of the codes. Additionally, we note that ASPECT is able to achieve good results with relatively less mesh resolution due to the ability to adaptively refine in the regions of interest (namely, the blob and the high viscosity lid).

Accuracy improves roughly linearly with decreasing CFL number, though stops improving at CFL ~ 0.1 . Accuracy also improves with increasing mesh resolution, though its convergence order does not seem to be excellent. It is possible that other mesh refinement parameters than we tried in this benchmark could improve the convergence. The primary challenge in accuracy is limiting numerical diffusion of the rising blob. If the blob becomes too diffuse, its ability to lift topography is diminished. It would be instructive to compare the results of this benchmark using particles with the results using compositional fields.

5.4.14 The solitary wave benchmark

This section was contributed by Juliane Dannberg and is based on a section in [DH16] by Juliane Dannberg and Timo Heister.

One of the most widely used benchmarks for codes that model the migration of melt through a compacting and dilating matrix is the propagation of solitary waves (e.g. [SS11, KMK13, Sch00]). The benchmark is intended to test the accuracy of the solution of the two-phase flow equations as described in Section 2.13 and makes use of the fact that there is an analytical solution for the shape of solitary waves that travel through a partially molten rock with a constant background porosity without changing their shape and with a constant wave speed. Here, we follow the setup of the benchmark as it is described in [BR86], which considers one-dimensional solitary waves.

The model features a perturbation of higher porosity with the amplitude $A\phi_0$ in a uniform low-porosity ($\phi = \phi_0$) background. Due to its lower density, melt migrates upwards, dilating the solid matrix at its front and compacting it at its end.

Assuming constant shear and compaction viscosities and using a permeability law of the form

$$k_\phi = k_0\phi^3, \quad \text{implying a Darcy coefficient } K_D(\phi) = \frac{k_0}{\eta_f}\phi^3,$$

and the non-dimensionalization

$$\begin{aligned} x &= \delta x' && \text{with the compaction length } \delta = \sqrt{K_D(\phi_0)(\xi + \frac{4}{3}\eta)}, \\ \phi &= \phi_0\phi' && \text{with the background porosity } \phi_0, \\ (\mathbf{u}_s, \mathbf{u}_f) &= u_0(\mathbf{u}_s, \mathbf{u}_f)' && \text{with the separation flux } \phi_0 u_0 = K_D(\phi_0)\Delta\rho g, \end{aligned}$$

the analytical solution for the shape of the solitary wave can be written in implicit form as:

$$x(\phi) = \pm(A + 0.5) \left[-2\sqrt{A - \phi} + \frac{1}{\sqrt{A - 1}} \ln \frac{\sqrt{A - 1} - \sqrt{A - \phi}}{\sqrt{A - 1} + \sqrt{A - \phi}} \right]$$

and the phase speed c , scaled back to physical units, is $c = u_0(2A + 1)$. This is only valid in the limit of small porosity $\phi_0 \ll 1$. Figure 94 illustrates the model setup.

The parameter file and material model for this setup can be found in `benchmarks/solitary_wave/solitary_wave.prm` and `benchmarks/solitary_wave/solitary_wave.cc`. The most relevant sections are shown in the following paragraph.

```
# Listing of Parameters
#
# Set up the solitary wave benchmark
# (Barcilon & Richter, 1986; Simpson & Spiegelman, 2011;
# Keller et al., 2013; Schmelting, 200

set Additional shared libraries      = ./libsolitary_wave.so

# A non-linear solver has to be used for models with melt migration
set Nonlinear solver scheme        = iterated Advection and Stokes
set Max nonlinear iterations       = 10
set Nonlinear solver tolerance     = 1e-5

# The end time is chosen in such a way that the solitary wave travels
# approximately 5 times its wavelength during the model time.
set End time                      = 6e6

# To model melt migration, there has to be a compositional field with
# the name 'porosity'.
subsection Compositional fields
  set Number of fields = 1
```

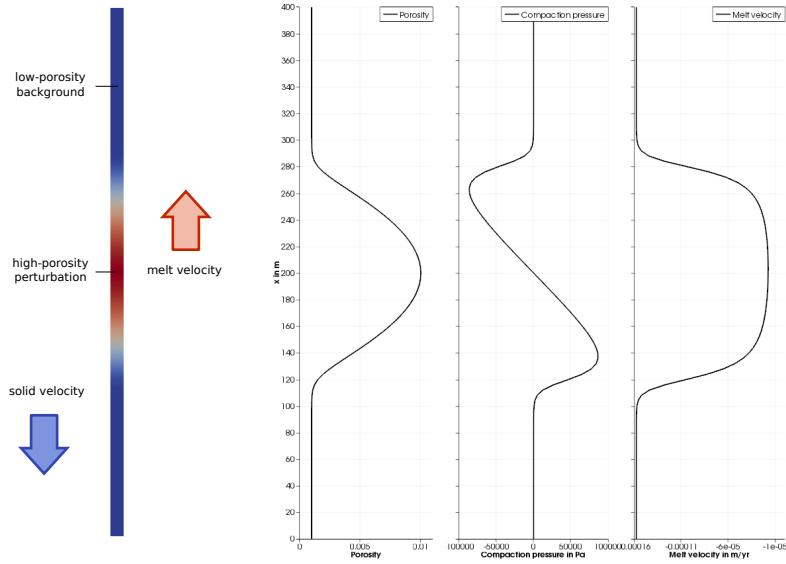


Figure 94: Setup of the solitary wave benchmark. The domain is 400 m high and includes a low porosity ($\phi = 0.001$) background with an initial perturbation ($\phi = 0.1$). The solid density is 3300 kg/m^3 and the melt density is 2500 kg/m^3 . We apply the negative phase speed of the solitary wave $\mathbf{u}_s = -c \mathbf{e}_z$ as velocity boundary condition, so that the wave will stay at its original position while the background is moving.

```

set Names of fields = porosity
end

# Enable modelling of melt migration in addition to the advection of
# solid material.
subsection Melt settings
    set Include melt transport          = true
end

##### Parameters for the porosity field #####
# We use the initial conditions and material model from the
# solitary wave plugin and choose a wave with an amplitude of
# 0.01 and a background porosity of 0.001.
subsection Initial composition model
    set Model name = Solitary wave initial condition
    subsection Solitary wave initial condition
        set Offset = 200
        set Read solution from file = true
        set Amplitude = 0.01
        set Background porosity = 0.001
    end
end

subsection Material model
    set Model name = Solitary Wave
end

```

```

# As material is flowing in, we prescribe the porosity at the
# upper and lower boundary.
subsection Boundary composition model
    set List of model names = initial composition
    set Fixed composition boundary indicators = 2,3
end

# As we know that our solution does not have any steep gradients
# we can use a low stabilization to avoid too much diffusion.
subsection Discretization
    subsection Stabilization parameters
        set beta = 0.001
    end
end

##### Model geometry #####
# Our domain is a pseudo-1D-profile 400 m in height, but only a few elements wide
subsection Geometry model
    set Model name = box
    subsection Box
        set X extent = 10
        set Y extent = 400
        set Y repetitions = 40
    end
end

##### Velocity boundary conditions #####
# We apply the phase speed of the wave here, so that it always stays in the
# same place in our model. The phase speed is c = 5.25e-11 m/s, but we have
# to convert it to m/years using the same conversion that is used internally
# in Aspect: year_in_seconds = 60*60*24*365.2425.
subsection Boundary velocity model
    set Tangential velocity boundary indicators = 0,1
    set Prescribed velocity boundary indicators = 2:function, 3:function
    subsection Function
        set Function expression = 0;-1.65673998e-4
    end
end

# Postprocessor for the error calculation
subsection Postprocess
    set List of postprocessors = solitary wave statistics
end

subsection Solver parameters
    subsection Stokes solver parameters
        set Linear solver tolerance = 1e-10
    end
end

```

The benchmark uses a custom model to generate the initial condition for the porosity field as specified by the analytical solution, and its own material model, which includes the additional material properties needed by models with melt migration, such as the permeability, melt density and compaction viscosity. The solitary wave postprocessor compares the porosity and pressure in the model to the analytical solution, and

computes the errors for the shape of the porosity, shape of the compaction pressure and the phase speed. We apply the negative phase speed of the solitary wave as a boundary condition for the solid velocity. This changes the reference frame, so that the solitary wave stays in the center of the domain, while the solid moves downwards. The temperature evolution does not play a role in this benchmark, so all temperature and heating-related parameters are disabled or set to zero.

And extensive discussion of the results and convergence behavior can be found in [DH16].

5.4.15 Benchmarks for operator splitting

This section was contributed by Juliane Dannberg.

Models of mantle convection and lithosphere dynamics often also contain reactions between materials with different chemical compositions, or processes that can be described as reactions. The most common example is mantle melting: When mantle temperatures exceed the solidus, rocks start to melt. As this is only partial melting, and rocks are a mixture of different minerals, which all contain different chemical components, melting is not only a phase transition, but also leads to reactions between solid and molten rock. Some components are more compatible with the mineral structure, and preferentially stay in the solid rock, other components will mainly move into the mantle melt. This means that the composition of both solid and melt change over time depending on the melt fraction.

Usually, it is assumed that these reactions are much faster than convection in the mantle. In other words, these reactions are so fast that melt is assumed to be always in equilibrium with the surrounding solid rock. In some cases, the formation of new oceanic crust, which is also caused by partial melting, is approximated by a conversion from an average, peridotitic mantle composition to mid-ocean ridge basalt, forming the crust, and harzburgitic lithosphere, once material reaches a given depth. This process can also be considered as a reaction between different compositional fields.

This can cause accuracy problems in geodynamic simulations: The way the equations are formulated (see Equations 1–4), ideally we would need to know reaction rates (q_i) between the different components instead of the equilibrium value (which would then have to be compared with some sort of “old solution” of the compositional fields). Sometimes we also may not know the equilibrium, and would only be able to find it iteratively, starting from the current composition. In addition, the reaction rate for a given compositional field usually depends on the value of the field itself, but can also depend on other compositional fields or the temperature and pressure, and the dependence can be nonlinear.

Hence, ASPECT has the option to decouple the advection from reactions between compositional fields, using operator splitting.

Instead of solving the coupled equation

$$\frac{\partial \mathbf{c}(t)}{\partial t} + \mathbf{u} \cdot \nabla \mathbf{c}(t) = q(\mathbf{c}(t)), \quad (84)$$

and directly obtaining the composition value $\mathbf{c}(t^{n+1})$ for the time step $n + 1$ from the value $\mathbf{c}(t^n)$ from the previous time step n , we do a first-order operator split, first solving the advection problem

$$\frac{\partial \mathbf{c}(t)}{\partial t} + \mathbf{u} \cdot \nabla \mathbf{c}(t) = 0, \quad \text{obtaining } \Delta \mathbf{c}_A(t^{n+1}) \text{ from } \mathbf{c}(t^n), \quad (85)$$

using the advection time step $\Delta t_A = t^{n+1} - t^n$. Then we solve the reactions as a series of coupled ordinary differential equations

$$\frac{\partial \Delta \mathbf{c}_R(t)}{\partial t} = q(\mathbf{c}(t^n)) + \Delta \mathbf{c}_A(t^{n+1}) + \Delta \mathbf{c}_R(t), \quad \text{obtaining } \Delta \mathbf{c}_R(t^{n+1}) \text{ from } \mathbf{c}(t^n) + \Delta \mathbf{c}_A(t^{n+1}). \quad (86)$$

This can be done in several iterations, choosing a different, smaller time step size $\Delta t_R \leq \Delta t_A$ for the time discretization. The updated value of the compositional field after the overall (advection + reaction) time step is then obtained as

$$\mathbf{c}(t^{n+1}) = \mathbf{c}(t^n) + \Delta \mathbf{c}_A(t^{n+1}) + \Delta \mathbf{c}_R(t^{n+1}). \quad (87)$$

This is very useful if the time scales of reactions are different from the time scales of convection. The same scheme can also be used for the temperature: If we want to model latent heat of melting, the temperature evolution is controlled by the melting rate, and hence the temperature changes on the same time scale as the reactions.

We here illustrate the way this operator splitting works using the simple example of exponential decay in a stationary advection field. We will start with a model that has a constant initial temperature and composition and no advection. The reactions for exponential decay

$$c(t) = c_0 e^{\lambda t} \text{ with } \lambda = -\log(2)/t_{1/2}, \quad (88)$$

where c_0 is the initial composition and $t_{1/2}$ is the half life, are implemented in a shared library ([benchmarks/operator_splitting/exponential_decay/exponential_decay.cc](#)). As we split the time-stepping of advection and reactions, there are now two different time steps in the model: We control the advection time step using the ‘Maximum time step’ parameter (as the velocity is essentially 0, we can not use the CFL number), and we set the reaction time step using the ‘Reaction time step’ parameter.

```
set Additional shared libraries = ./libexponential_decay.so

set Dimension = 2
set Start time = 0
set End time = 100

# We use a new solver scheme option that enables the operator split.
set Nonlinear solver scheme = single Advection, single Stokes
set Use operator splitting = true

subsection Solver parameters
    subsection Operator splitting parameters
        set Reaction time step = 0.0005
    end
end
set Maximum time step = 10
```

To illustrate convergence, we will vary both parameters in different model runs.

In our example, we choose $c_0 = 1$, and specify this as initial condition using the `function` plugin for both composition and temperature. We also set $t_{1/2} = 10$, which is implemented as a parameter in the exponential decay material model and the exponential decay heating model.

```
# Both initial temperature and composition are set to 1,
# and will decay starting from this value.
subsection Initial temperature model
    set Model name = function

    subsection Function
        set Variable names = x,z
        set Function expression = 1.0
    end
end

subsection Initial composition model
    set Model name = function

    subsection Function
        set Variable names = x,z
        set Function expression = 1.0
    end
```

```

end

# We choose material and heating models that let temperature
# and composition decay over time, and that is implemented in
# a plugin.
subsection Heating model
    set List of model names = exponential decay heating

    subsection Exponential decay heating
        set Half life = 10
    end
end

subsection Material model
    set Model name = exponential decay

    subsection Exponential decay
        set Half life = 10
    end
end

```

The complete parameter file for this setup can be found in [benchmarks/operator_splitting/exponential_decay/exponential_decay.base.prm](#).

Figure 95 shows the convergence behavior of these models: As there is no advection, the advection time step does not influence the error (blue data points). As we use a first-order operator split, the error is expected to converge linearly with the reaction time step Δt_R , which is indeed the case (red data points). Errors are the same for both composition and temperature, as both fields have identical initial conditions and reactions, and we use the same methods to solve for these variables.

For the second benchmark case, we want to see the effect of advection on convergence. In order to do this, we choose an initial temperature and composition that depends on x (in this case a sine), a decay rate that linearly depends on z, and we apply a constant velocity in x-direction on all boundaries. Our new analytical solution for the evolution of composition is now

$$c(t) = \sin(2\pi(x - tv_0)) c_0 e^{\lambda z t}. \quad (89)$$

v_0 is the constant velocity, which we set to 0.01 m/s. The parameter file for this setup can be found in [benchmarks/operator_splitting/advection_reaction/advection_reaction.base.prm](#). Figure 96 shows the convergence behavior in this second set of models: If we choose the same resolution as in the previous example (left panel), for large advection time steps $\Delta t_A > 0.1$ the error is dominated by advection, and converges with decreasing advection time step size (blue data points). However, for smaller advection time steps, the error stagnates. The data series where the reaction time step varies also shows a stagnating error. The reason for that is probably that our analytical solution is not in the finite element space we chose, and so neither decreasing the advection or the reaction time step will improve the error. If we increase the resolution by a factor of 4 (right panel), we see that that errors converge both with decreasing advection and reaction time steps.

The results shown here can be reproduced using the bash scripts `run.sh` in the corresponding benchmark folders.

6 Extending and contributing to ASPECT

ASPECT is designed to be an extensible code. In particular, it uses both a plugin architecture and a set of signals through which it is trivial to replace or extend certain components of the program. Examples of things that are simple to extend are:

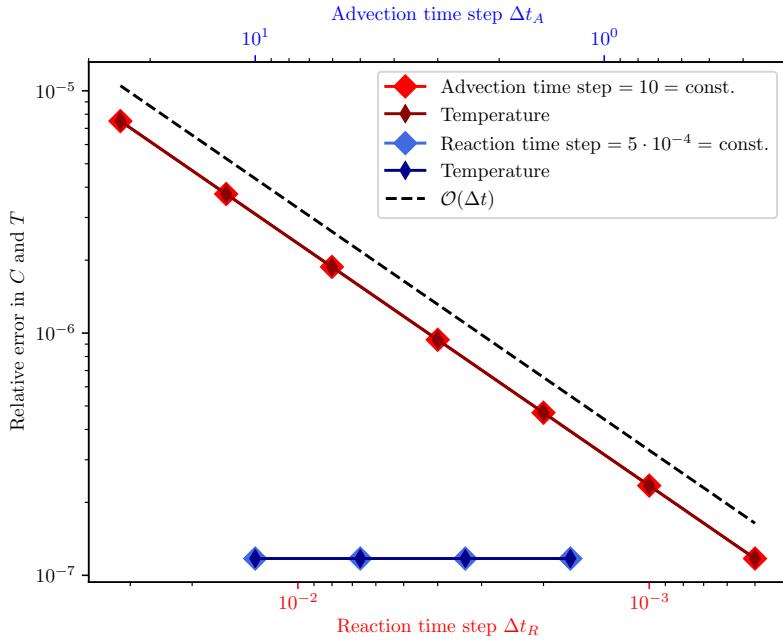


Figure 95: Error for both compositional field and temperature compared to the analytical solution, varying the time steps of advection (blue data points and top/blue x axis) and reactions (red data points and bottom/red x axis), while keeping the other one constant, respectively.

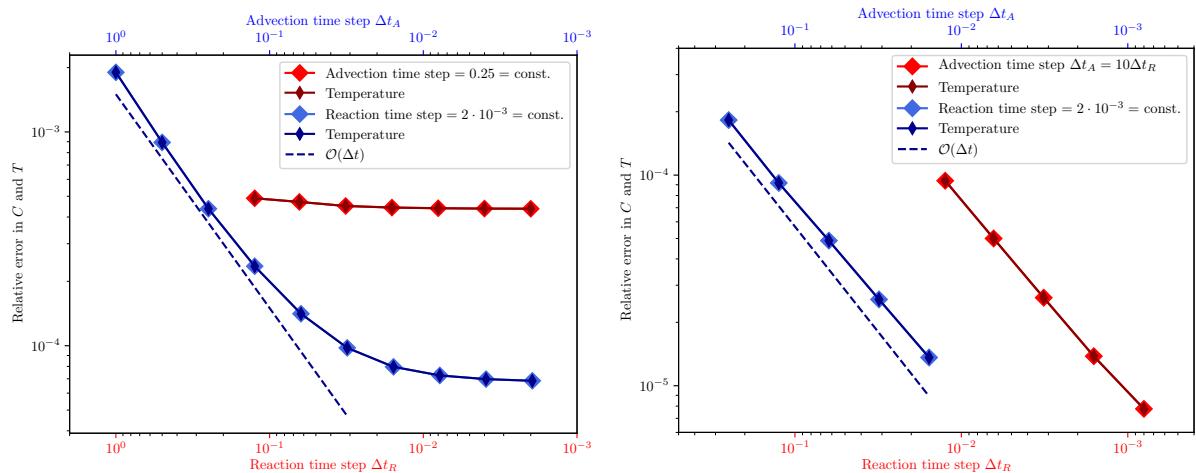


Figure 96: Error for both compositional field and temperature compared to the analytical solution, varying the time steps of advection (blue data points and top/blue x axis) and reactions (red data points and bottom/red x axis), while keeping the other one constant, respectively.

- the material description,
- the geometry,
- the gravity description,
- the initial conditions,
- the boundary conditions,
- the functions that postprocess the solution, i.e., that can compute derived quantities such as heat fluxes over part of the boundary, mean velocities, etc.,
- the functions that generate derived quantities that can be put into graphical output files for visualization such as fields that depict the strength of the friction heating term, spatially dependent actual viscosities, and so on,
- the computation of refinement indicators,
- the determination of how long a computation should run.

This list may also have grown since this section was written. We will discuss the way this is achieved in Sections 6.1 and 6.3. Changing the core functionality, i.e., the basic equations (1)–(3), and how they are solved is arguably more involved. We will discuss this in Section 6.6.

Note: The purpose of coming up with ways to make extensibility simple is that if you want to extend ASPECT for your own purposes, you can do this in a separate set of files that describe your situation, rather than by modifying the ASPECT source files themselves. This is important, because (i) it makes it possible for you to update ASPECT itself to a newer version without losing the functionality you added (because you did not make any changes to the ASPECT files themselves), (ii) because it makes it possible to keep unrelated changes separate in your own set of files, in a place where they are simple to find, and (iii) because it makes it much easier for you to share your modifications and additions with others.

Since ASPECT is written in C++ using the DEAL.II library, you will have to be proficient in C++. You will also likely have to familiarize yourself with this library for which there is an extensive amount of documentation:

- The manual at <https://www.dealii.org/developer/doxygen/deal.II/index.html> that describes in detail what every class, function and variable in DEAL.II does.
- A collection of modules at <https://www.dealii.org/developer/doxygen/deal.II/modules.html> that give an overview of whole groups of classes and functions and how they work together to achieve their goal.
- The DEAL.II tutorial at <https://www.dealii.org/developer/doxygen/tutorial/index.html> that provides a step-by-step introduction to the library using a sequence of several dozen programs that introduce gradually more complex topics. In particular, you will learn DEAL.II’s way of *dimension independent programming* that allows you to write the program once, test it in 2d, and run the exact same code in 3d without having to debug it a second time.
- The step-31 and step-32 tutorial programs at https://www.dealii.org/developer/doxygen/deal.II/step_31.html and https://www.dealii.org/developer/doxygen/deal.II/step_32.html from which ASPECT directly descends.

- An overview of many general approaches to numerical methods, but also a discussion of DEAL.II and tools we use in programming, debugging and visualizing data are given in Wolfgang Bangerth's video lectures. These are linked from the DEAL.II website at <https://www.dealii.org/> and directly available at <http://www.math.tamu.edu/~bangerth/videos.html>.
- The DEAL.II Frequently Asked Questions at <https://github.com/dealii/dealii/wiki/Frequently-Asked-Questions> that also have extensive sections on developing code with DEAL.II as well as on debugging. It also answers a number of questions we frequently get about the use of C++ in DEAL.II.
- Several other parts of the DEAL.II website at <https://www.dealii.org/> also have information that may be relevant if you dive deeper into developing code. If you have questions, the mailing lists at <https://www.dealii.org/mail.html> are also of general help.
- A general overview of DEAL.II is also provided in the paper [BHK07].

As a general note, by default ASPECT utilizes a DEAL.II feature called *debug mode*, see also the introduction to this topic in Section 4.3. If you develop code, you will definitely want this feature to be on, as it will capture the vast majority of bugs you will invariably introduce in your code.

When you write new functionality and run the code for the first time, you will almost invariably first have to deal with a number of these assertions that point out problems in your code. While this may be annoying at first, remember that these are actual bugs in your code that have to be fixed anyway and that are much easier to find if the program aborts than if you have to go by their more indirect results such as wrong answers. The Frequently Asked Questions at <https://github.com/dealii/dealii/wiki/Frequently-Asked-Questions> contain a section on how to debug DEAL.II programs.

The downside of debug mode, as mentioned before, is that it makes the program much slower. Consequently, once you are confident that your program actually does what it is intended to do – **but no earlier!** –, you may want to switch to optimized mode that links ASPECT with a version of the DEAL.II libraries that uses compiler optimizations and that does not contain the `assert` statements discussed above. This switch can be facilitated by editing the top of the ASPECT `Makefile` and recompiling the program.

In addition to these general comments, ASPECT is itself extensively documented. You can find documentation on all classes, functions and namespaces starting from the <doc/doxygen/index.html> page.

6.1 The idea of plugins and the `SimulatorAccess` and `Introspection` classes

The most common modification you will probably want to do to ASPECT are to switch to a different material model (i.e., have different values of functional dependencies for the coefficients η, ρ, C_p, \dots discussed in Section 2.2); change the geometry; change the direction and magnitude of the gravity vector \mathbf{g} ; or change the initial and boundary conditions.

To make this as simple as possible, all of these parts of the program (and some more) have been separated into what we call *plugins* that can be replaced quickly and where it is simple to add a new implementation and make it available to the rest of the program and the input parameter file. There are *a lot* of plugins already, see Fig. 97, that will often be useful starting points and examples if you want to implement plugins yourself.

The way this is achieved is through the following two steps:

- The core of ASPECT really only communicates with material models, geometry descriptions, etc., through a simple and very basic interface. These interfaces are declared in the `include/aspect/material_model/interface.h`, `include/aspect/geometry_model/interface.h`, etc., header files. These classes are always called `Interface`, are located in namespaces that identify their purpose, and their documentation can be found from the general class overview in <doc/doxygen/classes.html>.

To show an example of a rather minimal case, here is the declaration of the `aspect::GravityModel::Interface` class (documentation comments have been removed):

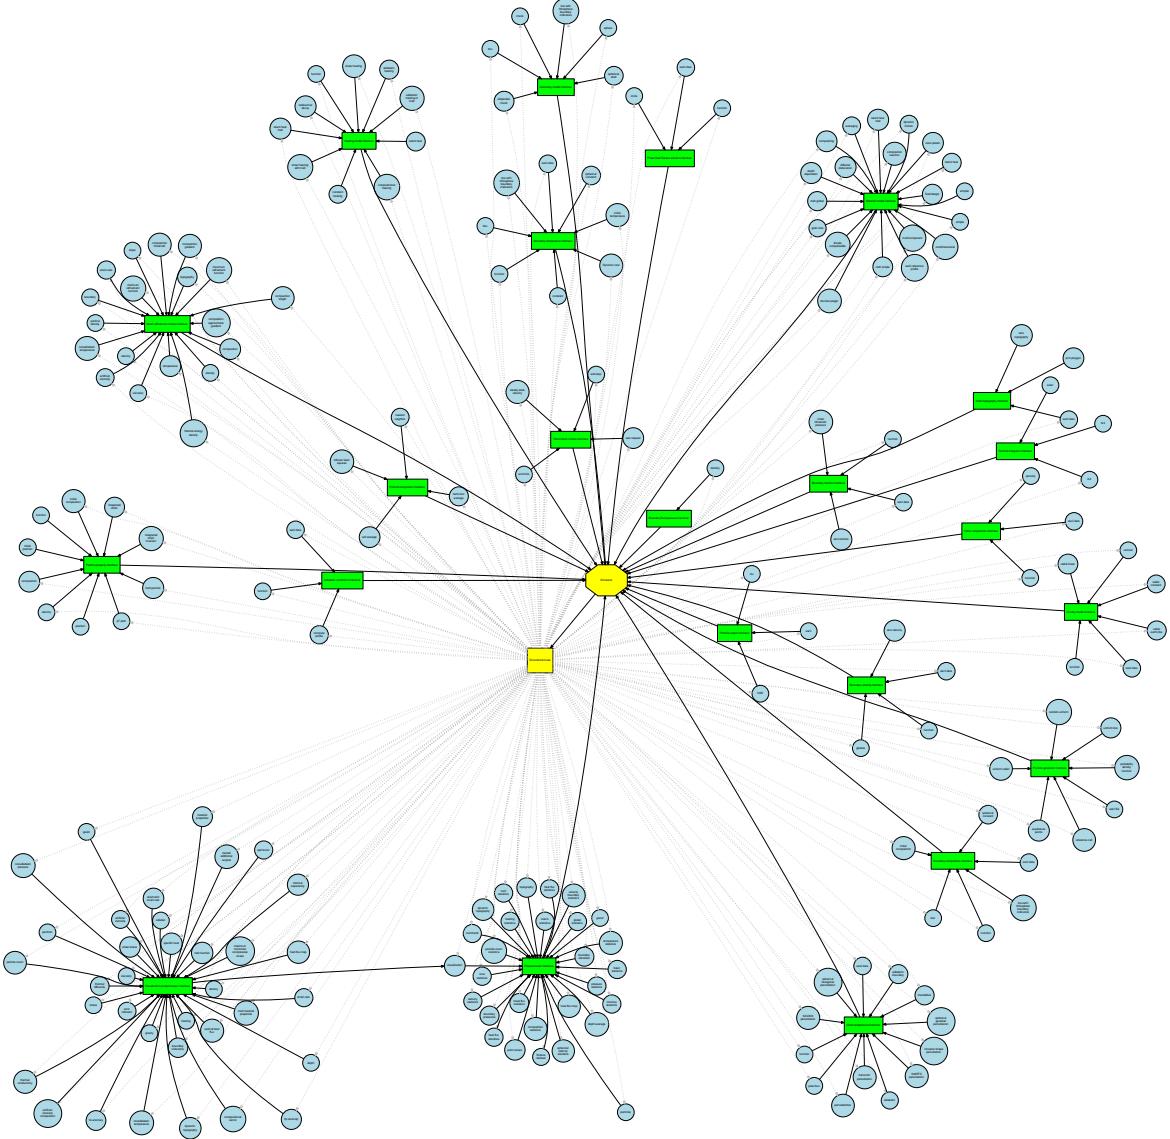


Figure 97: The graph of all current plugins of ASPECT. The yellow octagon and square represent the `Simulator` and `SimulatorAccess` classes. The green boxes are interface classes for everything that can be changed by plugins. Blue circles correspond to plugins that implement particular behavior. The graph is of course too large to allow reading individual plugin names (unless you zoom far into the page), but is intended to illustrate the architecture of ASPECT.

```

class Interface
{
public:
    virtual ~Interface();

    virtual
    Tensor<1,dim>
    gravity_vector (const Point<dim> &position) const = 0;

    static void declare_parameters (ParameterHandler &prm);

    virtual void parse_parameters (ParameterHandler &prm);
};

```

If you want to implement a new model for gravity, you just need to write a class that derives from this base class and implements the `gravity_vector` function. If your model wants to read parameters from the input file, you also need to have functions called `declare_parameters` and `parse_parameters` in your class with the same signatures as the ones above. On the other hand, if the new model does not need any run-time parameters, you do not need to overload these functions.³⁷

Each of the categories above that allow plugins have several implementations of their respective interfaces that you can use to get an idea of how to implement a new model.

- At the end of the file where you implement your new model, you need to have a call to the macro `ASPECT_REGISTER_GRAVITY_MODEL` (or the equivalent for the other kinds of plugins). For example, let us say that you had implemented a gravity model that takes actual gravimetric readings from the GRACE satellites into account, and had put everything that is necessary into a class `aspect::GravityModel::GRACE`. Then you need a statement like this at the bottom of the file:

```

ASPECT_REGISTER_GRAVITY_MODEL
(GRACE,
 "grace",
 "A gravity model derived from GRACE "
 "data. Run-time parameters are read from the parameter "
 "file in subsection 'Radial constant'.");

```

Here, the first argument to the macro is the name of the class. The second is the name by which this model can be selected in the parameter file. And the third one is a documentation string that describes the purpose of the class (see, for example, Section A.52 for an example of how existing models describe themselves).

This little piece of code ensures several things: (i) That the parameters this class declares are known when reading the parameter file. (ii) That you can select this model (by the name “grace”) via the run-time parameter `Gravity model/Model name`. (iii) That ASPECT can create an object of this kind when selected in the parameter file.

Note that you need not announce the existence of this class in any other part of the code: Everything

³⁷At first glance one may think that only the `parse_parameters` function can be overloaded since `declare_parameters` is not virtual. However, while the latter is called by the class that manages plugins through pointers to the interface class, the former function is called essentially at the time of registering a plugin, from code that knows the actual type and name of the class you are implementing. Thus, it can call the function – if it exists in your class, or the default implementation in the base class if it doesn’t – even without it being declared as virtual.

should just work automatically.³⁸ This has the advantage that things are neatly separated: You do not need to understand the core of ASPECT to be able to add a new gravity model that can then be selected in an input file. In fact, this is true for all of the plugins we have: by and large, they just receive some data from the simulator and do something with it (e.g., postprocessors), or they just provide information (e.g., initial meshes, gravity models), but their writing does not require that you have a fundamental understanding of what the core of the program does.

The procedure for the other areas where plugins are supported works essentially the same, with the obvious change in namespace for the interface class and macro name.

In the following, we will discuss the requirements for individual plugins. Before doing so, however, let us discuss ways in which plugins can query other information, in particular about the current state of the simulation. To this end, let us not consider those plugins that by and large just provide information without any context of the simulation, such as gravity models, prescribed boundary velocities, or initial temperatures. Rather, let us consider things like postprocessors that can compute things like boundary heat fluxes. Taking this as an example (see Section 6.3.8), you are required to write a function with the following interface

```
template <int dim>
class MyPostprocessor : public aspect::Postprocess::Interface
{
public:
    virtual
    std::pair<std::string, std::string>
    execute (TableHandler &statistics);

// ... more things ...
}
```

The idea is that in the implementation of the `execute` function you would compute whatever you are interested in (e.g., heat fluxes) and return this information in the `statistics` object that then gets written to a file (see Sections 4.1 and 4.4.2). A postprocessor may also generate other files if it so likes – e.g., graphical output, a file that stores the locations of particles, etc. To do so, obviously you need access to the current solution. This is stored in a vector somewhere in the core of ASPECT. However, this vector is, by itself, not sufficient: you also need to know the finite element space it is associated with, and for that the triangulation it is defined on. Furthermore, you may need to know what the current simulation time is. A variety of other pieces of information enters computations in these kinds of plugins.

All of this information is of course part of the core of ASPECT, as part of the `aspect::Simulator` class. However, this is a rather heavy class: it's got dozens of member variables and functions, and it is the one that does all of the numerical heavy lifting. Furthermore, to access data in this class would require that you need to learn about the internals, the data structures, and the design of this class. It would be poor design if plugins had to access information from this core class directly. Rather, the way this works is that those plugin classes that wish to access information about the state of the simulation inherit from the `aspect::SimulatorAccess` class. This class has an interface that looks like this:

```
template <int dim>
class SimulatorAccess
{
protected:
    double      get_time () const;

    std::string  get_output_directory () const;

    const LinearAlgebra::BlockVector &
```

³⁸The existing implementations of models of the gravity and other interfaces declare the class in a header file and define the member functions in a .cc file. This is done so that these classes show up in our doxygen-generated documentation, but it is not necessary: you can put your entire class declaration and implementation into a single file as long as you call the macro discussed above on it. This single file is all you need to touch to add a new model.

```

get_solution () const;

const DoFHandler<dim> &
get_dof_handler () const;

// ... many more things ...

```

This way, `SimulatorAccess` makes information available to plugins without the need for them to understand details of the core of ASPECT. Rather, if the core changes, the `SimulatorAccess` class can still provide exactly the same interface. Thus, it insulates plugins from having to know the core. Equally importantly, since `SimulatorAccess` only offers its information in a read-only way it insulates the core from plugins since they can not interfere in the workings of the core except through the interface they themselves provide to the core.

Using this class, if a plugin class `MyPostprocess` is then not only derived from the corresponding `Interface` class but *also* from the `SimulatorAccess` class (as indeed most plugins are, see the dashed arrows in Fig. 97), then you can write a member function of the following kind (a nonsensical but instructive example; see Section 6.3.8 for more details on what postprocessors do and how they are implemented):³⁹

```

template <int dim>
std::pair<std::string, std::string>
MyPostprocessor<dim>::execute (TableHandler &statistics)
{
    // compute the mean value of vector component 'dim' of the solution
    // (which here is the pressure block) using a deal.II function:
    const double
        average_pressure = VectorTools::compute_mean_value (this->get_mapping(),
                                                          this->get_dof_handler(),
                                                          QGauss<dim>(2),
                                                          this->get_solution(),
                                                          dim);

    statistics.add_value ("Average pressure", average_pressure);

    // return that there is nothing to print to screen (a useful
    // plugin would produce something more elaborate here):
    return std::pair<std::string, std::string>();
}

```

The second piece of information that plugins can use is called “introspection”. In the code snippet above, we had to use that the pressure variable is at position `dim`. This kind of *implicit knowledge* is usually bad style: it is error prone because one can easily forget where each component is located; and it is an obstacle to the extensibility of a code if this kind of knowledge is scattered all across the code base.

Introspection is a way out of this dilemma. Using the `SimulatorAccess::introspection()` function returns a reference to an object (of type `aspect::Introspection`) that plugins can use to learn about these sort of conventions. For example, `this->introspection().component_mask.pressure` returns a component mask (a deal.II concept that describes a list of booleans for each component in a finite element that are true if a component is part of a variable we would like to select and false otherwise) that describes which component of the finite element corresponds to the pressure. The variable, `dim`, we need above to indicate that we want the pressure component can be accessed as `this->introspection().component_indices.pressure`. While this is certainly not shorter than just writing `dim`, it may in fact be easier to remember. It is most definitely less prone to errors and makes it simpler to extend the code in the future because we don’t litter the sources with “magic constants” like the one above.

³⁹For complicated, technical reasons, in the code below we need to access elements of the `SimulatorAccess` class using the notation `this->get_solution()`, etc. This is due to the fact that both the current class and the base class are templates. A long description of why it is necessary to use `this->` can be found in the DEAL.II Frequently Asked Questions.

This `aspect::Introspection` class has a significant number of variables that can be used in this way, i.e., they provide symbolic names for things one frequently has to do and that would otherwise require implicit knowledge of things such as the order of variables, etc.

6.2 How to write a plugin

Before discussing what each kind of plugin actually has to implement (see the next subsection), let us briefly go over what you actually have to do when implementing a new plugin. Essentially, the following steps are all you need to do:

- Create a file, say `my_plugin.cc` that contains the declaration of the class you want to implement. This class must be derived from one of the `Interface` classes we will discuss below. The file also needs to contain the implementation of all member functions of your class.

As discussed above, it is possible – but not necessary – to split this file into two: a header file, say `my_plugin.h`, and the `my_plugin.cc` file (or, if you prefer, into multiple source files). We do this for all the existing plugins in ASPECT so that the documentation of these plugins shows up in the doxygen-generated documentation. However, for your own plugins, there is typically no need for this split. The only occasion where this would be useful is if some plugin actually makes use of a different plugin (e.g., the implementation of a gravity model of your own may want to query some specifics of a geometry model you also implemented); in that case the *using* plugin needs to be able to see the declaration of the class of the *used* plugin, and for this you will need to put the declaration of the latter into a header file.

- At the bottom of the `my_plugin.cc` file, put a statement that instantiates the plugin, documents it, and makes it available to the parameter file handlers by registering it. This is always done using one of the `ASPECT_REGISTER_*` macros that will be discussed in the next subsections; take a look at how they are used in the existing plugins in the ASPECT source files.
- You need to compile the file. There are two ways by which this can be achieved:
 - Put the `my_plugin.cc` into one of the ASPECT source directories and call `cmake .` followed by `make` to ensure that it actually gets compiled. This approach has the advantage that you do not need to worry much about how the file actually gets compiled. On the other hand, every time you modify the file, calling `make` requires not only compiling this one file, but also link ASPECT. Furthermore, when you upgrade from one version of ASPECT to another, you need to remember to copy the `my_plugin.cc` file.
 - Put the `my_plugin.cc` file into a directory of your choice and compile it into a shared library yourself. This may be as easy as calling

```
g++ -I/path/to/aspect/include -I/path/to/deal.II/include \
     -fPIC -shared my_plugin.cc -o my_plugin.so
```

on Linux, but the command may be different on other systems. Now you only need to tell ASPECT to load this shared library at startup so that the plugin becomes available at run time and can be selected from the input parameter file. This is done using the `Additional shared libraries` parameter in the input file, see Section A.1. This approach has the upside that you can keep all files that define new plugins in your own directories where you also run the simulations, also making it easier to keep around your plugins as you upgrade your ASPECT installation. On the other hand, compiling the file into a shared library is a bit more than you need to do yourself. Nevertheless, this is the preferred approach.

In practice, the compiler line above can become tedious because it includes paths to the ASPECT and DEAL.II header files, but possibly also other things such as Trilinos headers, etc. Having to

remember all of these pieces is a hassle, and a much easier way is in fact to set up a mini-CMake project for this. To this end, simply copy the file [doc/plugin-CMakeLists.txt](#) to the directory where you have your plugin source files and rename it to `CMakeLists.txt`.

You can then just run the commands

```
cmake -DAspect_DIR=/path/to/aspect .
make
```

and it should compile your plugin files into a shared library `my_plugin.so`. A concrete example of this process is discussed in Section [5.4.1](#). Of course, you may want to choose different names for the source files `source_1.cc`, `source_2.cc` or the name of the plugin `my_plugin`.

In essence, what these few lines do is that they find an ASPECT installation (i.e., the directory where you configured and compiled it, which may be the same directory as where you keep your sources, or a different one, as discussed in Section [3](#)) in either the directory explicitly specified in the `Aspect_DIR` variable passed to `cmake`, the shell environment variable `ASPECT_DIR`, or just one directory up. It then sets up compiler paths and similar, and the following lines simply define the name of a plugin, list the source files for it, and define everything that's necessary to compile them into a shared library. Calling `make` on the command line then simply compiles everything.

Note: Complex projects built on ASPECT often require plugins of more than just one kind. For example, they may have plugins for the geometry, the material model, and for postprocessing. In such cases, you can either define multiple shared libraries by repeating the calls to `PROJECT`, `ADD_LIBRARY` and `ASPECT_SETUP_PLUGIN` for each shared library in your `CMakeLists.txt` file above, or you can just compile all of your source files into a single shared library. In the latter case, you only need to list a single library in your input file, but each plugin will still be selectable in the various sections of your input file as long as each of your classes has a corresponding `ASPECT_REGISTER_*` statement somewhere in the file where you have its definition. An even simpler approach is to just put everything into a single file – there is no requirement that different plugins are in separate files, though this is often convenient from a code organization point of view.

Note: If you choose to compile your plugins into a shared library yourself, you will need to recompile them every time you upgrade your ASPECT installation since we do not guarantee that the ASPECT application binary interface (ABI) will remain stable, even if it may not be necessary to actually change anything in the *implementation* of your plugin.

6.3 Materials, geometries, gravitation and other plugin types

6.3.1 Material models

The material model is responsible for describing the various coefficients in the equations that ASPECT solves. To implement a new material model, you need to overload the `aspect::MaterialModel::Interface` class and use the `ASPECT_REGISTER_MATERIAL_MODEL` macro to register your new class. The implementation of the new class should be in namespace `aspect::MaterialModel`. An example of a material model implemented this way is given in Section [5.4.12](#).

Specifically, your new class needs to implement the following interface:

```
template <int dim>
class aspect::MaterialModel::Interface
{
```

```

public:
    // Physical parameters used in the basic equations
    virtual void evaluate(const MaterialModelInputs &in, MaterialModelOutputs &out) const=0;

    virtual bool is_compressible () const = 0;

    // Reference quantities
    virtual double reference_viscosity () const = 0;

    // Functions used in dealing with run-time parameters
    static void
    declare_parameters (ParameterHandler &prm);

    virtual void
    parse_parameters (ParameterHandler &prm);

    // Optional:
    virtual void initialize ();

    virtual void update ();
}

```

The main properties of the material are computed in the function `evaluate()` that takes a struct of type `MaterialModelInputs` and is supposed to fill a `MaterialModelOutputs` structure. For performance reasons this function is handling lookups at an arbitrary number of positions, so for each variable (for example viscosity), a `std::vector` is returned. The following members of `MaterialModelOutputs` need to be filled:

```

struct MaterialModelOutputs
{
    std::vector<double> viscosities;
    std::vector<double> densities;
    std::vector<double> thermal_expansion_coefficients;
    std::vector<double> specific_heat;
    std::vector<double> thermal_conductivities;
    std::vector<double> compressibilities;
}

```

The variables refer to the coefficients η, C_p, k, ρ in equations (1)–(3), each as a function of temperature, pressure, position, compositional fields and, in the case of the viscosity, the strain rate (all handed in by `MaterialModelInputs`). Implementations of `evaluate()` may of course choose to ignore dependencies on any of these arguments.

The remaining functions are used in postprocessing as well as handling run-time parameters. The exact meaning of these member functions is documented in the [aspect::MaterialModel::Interface class documentation](#). Note that some of the functions listed above have a default implementation, as discussed on the documentation page just mentioned.

The function `is_compressible` returns whether we should consider the material as compressible or not, see Section 2.10.3 on the Boussinesq model. As discussed there, incompressibility as described by this function does not necessarily imply that the density is constant; rather, it may still depend on temperature or pressure. In the current context, compressibility simply means whether we should solve the continuity equation as $\nabla \cdot (\rho \mathbf{u}) = 0$ (compressible Stokes) or as $\nabla \cdot \mathbf{u} = 0$ (incompressible Stokes).

The purpose of the parameter handling functions has been discussed in the general overview of plugins above.

The functions `initialize()` and `update()` can be implemented if desired (the default implementation does nothing) and are useful if the material model has internal state. The function `initialize()` is called once during the initialization of ASPECT and can be used to allocate memory, initialize state, or read information from an external file. The function `update()` is called at the beginning of every time step.

Additionally, every material model has a member variable “`model_dependence`”, declared in the Interface class, which can be accessed from the plugin as “`this->model_dependence`”. This structure describes the nonlinear dependence of the various coefficients on pressure, temperature, composition or strain rate. This information will be used in future versions of ASPECT to implement a fully nonlinear solution scheme based on, for example, a Newton iteration. The initialization of this variable is optional, but only plugins that declare correct dependencies can benefit from these solver types. All packaged material models declare their dependencies in the `parse_parameters()` function and can be used as a starting point for implementations of new material models.

Older versions of ASPECT used to have individual functions like `viscosity()` instead of the `evaluate()` function discussed above. This old interface is no longer supported, restructure your plugin to implement `evaluate()` instead (even if this function only calls the old functions).

6.3.2 Heating models

The heating model is responsible for describing the various terms in the energy equation (3), using the coefficients provided by the material model. These can be source terms such as radiogenic heat production or shear heating, they can be terms on the left-hand side of the equation, such as part of the latent heating terms, or they can be heating processes related to reactions. Each of these terms is described by a “heating model”, and a simulation can have none, one, or many heating models that are active throughout a simulation, with each heating model usually only implementing the terms for one specific heating process. One can then decide in the input file which heating processes should be included in the computation by providing a list of heating models in the input file.

When the equations are assembled and solved, the heating terms from all heating models used in the computation are added up.

To implement a new heating model, you need to overload the `aspect::HeatingModel::Interface` class and use the `ASPECT_REGISTER_HEATING_MODEL` macro to register your new class. The implementation of the new class should be in namespace `aspect::HeatingModel`.

Specifically, your new class needs to implement the following basic interface:

```
template <int dim>
class aspect::HeatingModel::Interface
{
public:
    // compute heating terms used in the energy equation
    virtual
    void
    evaluate (const MaterialModel::MaterialModelInputs<dim> &material_model_inputs,
              const MaterialModel::MaterialModelOutputs<dim> &material_model_outputs,
              HeatingModel::HeatingModelOutputs &heating_model_outputs) const;

    // All the following functions are optional:
    virtual
    void
    initialize ();

    virtual
    void
    update ();

    // Functions used in dealing with run-time parameters
```

```

    static
    void
    declare_parameters (ParameterHandler &prm);

    virtual
    void
    parse_parameters (ParameterHandler &prm);

    // Allow the heating model to attach additional material model outputs in case it needs
    // them to compute the heating terms
    virtual
    void
    create_additional_material_model_outputs(MaterialModel::MaterialModelOutputs<dim> &) const;
};


```

The main properties of the material are computed in the function `evaluate()` that takes references to `MaterialModelInputs` and `MaterialModelOutputs` objects and is supposed to fill the `HeatingModelOutputs` structure. As in the material model, this function is handling lookups at an arbitrary number of positions, so for each heating term (for example the heating source terms), a `std::vector` is returned. The following members of `HeatingModelOutputs` need to be filled:

```

struct HeatingModelOutputs
{
    std::vector<double> heating_source_terms;
    std::vector<double> lhs_latent_heat_terms;

    // optional:
    std::vector<double> rates_of_temperature_change;
}

```

Heating source terms are terms on the right-hand side of the equations, such as the adiabatic heating $\alpha T (\mathbf{u} \cdot \nabla p)$ in equation (3). An example for a left-hand side heating term is the temperature-derivative term $\rho T \Delta S \frac{\partial X}{\partial T}$ that is part of latent heat production (see equation (5)).⁴⁰ Rates of temperature change⁴¹ are used when the heating term is related to a reaction process, happening on a faster time scale than the temperature advection. All of these terms can depend on any of the material model inputs or outputs. Implementations of `evaluate()` may of course choose to ignore dependencies on any of these arguments.

The remaining functions are used in postprocessing as well as handling run-time parameters. The exact meaning of these member functions is documented in the [aspect::HeatingModel::Interface class documentation](#). Note that some of the functions listed above have a default implementation, as discussed on the documentation page just mentioned.

Just like for material models, the functions `initialize()` and `update()` can be implemented if desired (the default implementation does nothing) and are useful if the heating model has an internal state. The function `initialize()` is called once during the initialization of ASPECT and can be used to allocate memory for the heating model, initialize state, or read information from an external file. The function `update()` is called at the beginning of every time step.

6.3.3 Geometry models

The geometry model is responsible for describing the domain in which we want to solve the equations. A domain is described in DEAL.II by a coarse mesh and, if necessary, an object that characterizes the boundary.

⁴⁰Whether a term should go on the left or right hand side of the equation is, in some sense, a choice one can make. Putting a term onto the right hand side makes it an explicit term as far as time stepping is concerned, and so may imply a time step restriction if its dynamics are too fast. On the other hand, it does not introduce a nonlinearity if it depends on more than just a multiple of the temperature (such as the term $\alpha T (\mathbf{u} \cdot \nabla p)$). In practice, whether one wants to put a specific term on one side or the other may be a judgment call based on experience with numerical methods.

⁴¹Or, more correctly: Rates of *thermal energy change*.

Together, these two suffice to reconstruct any domain by adaptively refining the coarse mesh and placing new nodes generated by refining cells onto the surface described by the boundary object. The geometry model is also responsible for marking different parts of the boundary with different *boundary indicators* for which one can then, in the input file, select whether these boundaries should be Dirichlet-type (fixed temperature) or Neumann-type (no heat flux) boundaries for the temperature, and what kind of velocity conditions should hold there. In DEAL.II, a boundary indicator is a number of type `types::boundary_id`, but since boundaries are hard to remember and get right in input files, geometry models also have a function that provide a map from symbolic names that can be used to describe pieces of the boundary to the corresponding boundary indicators. For example, the simple `box` geometry model in 2d provides the map `{"left"→0, "right"→1, "bottom"→2, "top"→3}`, and we have consistently used these symbolic names in the input files used in this manual.

To implement a new geometry model, you need to overload the `aspect::GeometryModel::Interface` class and use the `ASPECT_REGISTER_GEOMETRY_MODEL` macro to register your new class. The implementation of the new class should be in namespace `aspect::GeometryModel`.

Specifically, your new class needs to implement the following basic interface:

```
template <int dim>
class aspect::GeometryModel::Interface
{
public:
    virtual
    void
    create_coarse_mesh (parallel::distributed::Triangulation<dim> &coarse_grid) const = 0;

    virtual
    double
    length_scale () const = 0;

    virtual
    double depth(const Point<dim> &position) const = 0;

    virtual
    Point<dim> representative_point(const double depth) const = 0;

    virtual
    double maximal_depth() const = 0;

    virtual
    std::set<types::boundary_id_t>
    get_used_boundary_indicators () const = 0;

    virtual
    std::map<std::string,types::boundary_id>
    get_symbolic_boundary_names_map () const;

    static
    void
    declare_parameters (ParameterHandler &prm);

    virtual
    void
    parse_parameters (ParameterHandler &prm);
};
```

The kind of information these functions need to provide is extensively discussed in the documentation of

this interface class at `aspect::GeometryModel::Interface`. The purpose of the last two functions has been discussed in the general overview of plugins above.

The `create_coarse_mesh` function does not only create the actual mesh (i.e., the locations of the vertices of the coarse mesh and how they connect to cells) but it must also set the boundary indicators for all parts of the boundary of the mesh. The DEAL.II glossary describes the purpose of boundary indicators as follows:

In a `Triangulation` object, every part of the boundary is associated with a unique number (of type `types::boundary_id`) that is used to identify which boundary geometry object is responsible to generate new points when the mesh is refined. By convention, this boundary indicator is also often used to determine what kinds of boundary conditions are to be applied to a particular part of a boundary. The boundary is composed of the faces of the cells and, in 3d, the edges of these faces.

By default, all boundary indicators of a mesh are zero, unless you are reading from a mesh file that specifically sets them to something different, or unless you use one of the mesh generation functions in namespace `GridGenerator` that have a 'colorize' option. A typical piece of code that sets the boundary indicator on part of the boundary to something else would look like this, here setting the boundary indicator to 42 for all faces located at $x = -1$:

```
for (typename Triangulation<dim>::active_cell_iterator
     cell = triangulation.begin_active();
     cell != triangulation.end();
     ++cell)
  for (unsigned int f=0; f<GeometryInfo<dim>::faces_per_cell; ++f)
    if (cell->face(f)->at_boundary())
      if (cell->face(f)->center()[0] == -1)
        cell->face(f)->set_boundary_indicator (42);
```

This calls functions `TriaAccessor::set_boundary_indicator`. In 3d, it may also be appropriate to call `TriaAccessor::set_all_boundary_indicators` instead on each of the selected faces. To query the boundary indicator of a particular face or edge, use `TriaAccessor::boundary_indicator`.

The code above only sets the boundary indicators of a particular part of the boundary, but it does not by itself change the way the `Triangulation` class treats this boundary for the purposes of mesh refinement. For this, you need to call `Triangulation::set_boundary` to associate a boundary object with a particular boundary indicator. This allows the `Triangulation` object to use a different method of finding new points on faces and edges to be refined; the default is to use a `StraightBoundary` object for all faces and edges. The results section of step-49 has a worked example that shows all of this in action.

The second use of boundary indicators is to describe not only which geometry object to use on a particular boundary but to select a part of the boundary for particular boundary conditions.
[...]

Note: Boundary indicators are inherited from mother faces and edges to their children upon mesh refinement. Some more information about boundary indicators is also presented in a section of the documentation of the `Triangulation` class.

Two comments are in order here. First, if a coarse triangulation's faces already accurately represent where you want to pose which boundary condition (for example to set temperature values or determine which are no-flow and which are tangential flow boundary conditions), then it is sufficient to set these boundary indicators only once at the beginning of the program since they will be inherited upon mesh refinement to the child faces. Here, *at the beginning of the program* is equivalent to inside the `create_coarse_mesh()` function of the geometry module shown above that generates the coarse mesh.

Secondly, however, if you can only accurately determine which boundary indicator should hold where on a refined mesh – for example because the coarse mesh is the cube $[0, L]^3$ and you want to have a fixed

velocity boundary describing an extending slab only for those faces for which $z > L - L_{\text{slab}}$ – then you need a way to set the boundary indicator for all boundary faces either to the value representing the slab or the fluid underneath *after every mesh refinement step*. By doing so, child faces can obtain boundary indicators different from that of their parents. DEAL.II triangulations support this kind of operations using a so-called *post-refinement signal*. In essence, what this means is that you can provide a function that will be called by the triangulation immediately after every mesh refinement step.

The way to do this is by writing a function that sets boundary indicators and that will be called by the `Triangulation` class. The triangulation does not provide a pointer to itself to the function being called, nor any other information, so the trick is to get this information into the function. C++ provides a nice mechanism for this that is best explained using an example:

```
#include <deal.II/base/std_cxx1x/bind.h>

template <int dim>
void set_boundary_indicators (parallel::distributed::Triangulation<dim> &triangulation)
{
    ... set boundary indicators on the triangulation object ...
}

template <int dim>
void
MyGeometry<dim>::
create_coarse_mesh (parallel::distributed::Triangulation<dim> &coarse_grid) const
{
    ... create the coarse mesh ...

    coarse_grid.signals.post_refinement.connect
        (std_cxx1x::bind (&set_boundary_indicators<dim>,
                         std_cxx1x::ref(coarse_grid)));
}
```

What the call to `std_cxx1x::bind` does is to produce an object that can be called like a function with no arguments. It does so by taking the address of a function that does, in fact, take an argument but permanently fix this one argument to a reference to the coarse grid triangulation. After each refinement step, the triangulation will then call the object so created which will in turn call `set_boundary_indicators<dim>` with the reference to the coarse grid as argument.

This approach can be generalized. In the example above, we have used a global function that will be called. However, sometimes it is necessary that this function is in fact a member function of the class that generates the mesh, for example because it needs to access run-time parameters. This can be achieved as follows: assuming the `set_boundary_indicators()` function has been declared as a (non-static, but possibly private) member function of the `MyGeometry` class, then the following will work:

```
#include <deal.II/base/std_cxx1x/bind.h>

template <int dim>
void
MyGeometry<dim>::
set_boundary_indicators (parallel::distributed::Triangulation<dim> &triangulation) const
{
    ... set boundary indicators on the triangulation object ...
}

template <int dim>
void
MyGeometry<dim>::
```

```

create_coarse_mesh (parallel::distributed::Triangulation<dim> &coarse_grid) const
{
    ... create the coarse mesh ...

    coarse_grid.signals.post_refinement.connect
        (std_cxx1x::bind (&MyGeometry<dim>::set_boundary_indicators,
                          std_cxx1x::cref(*this),
                          std_cxx1x::ref(coarse_grid)));
}

```

Here, like any other member function, `set_boundary_indicators` implicitly takes a pointer or reference to the object it belongs to as first argument. `std::bind` again creates an object that can be called like a global function with no arguments, and this object in turn calls `set_boundary_indicators` with a pointer to the current object and a reference to the triangulation to work on. Note that because the `create_coarse_mesh` function is declared as `const`, it is necessary that the `set_boundary_indicators` function is also declared `const`.

Note: For reasons that have to do with the way the `parallel::distributed::Triangulation` is implemented, functions that have been attached to the post-refinement signal of the triangulation are called more than once, sometimes several times, every time the triangulation is actually refined.

6.3.4 Gravity models

The gravity model is responsible for describing the magnitude and direction of the gravity vector at each point inside the domain. To implement a new gravity model, you need to overload the `aspect::GravityModel::Interface` class and use the `ASPECT_REGISTER_GRAVITY_MODEL` macro to register your new class. The implementation of the new class should be in namespace `aspect::GravityModel`.

Specifically, your new class needs to implement the following basic interface:

```

template <int dim>
class aspect::GravityModel::Interface
{
public:
    virtual
    Tensor<1,dim>
    gravity_vector (const Point<dim> &position) const = 0;

    virtual
    void
    update ();

    static
    void
    declare_parameters (ParameterHandler &prm);

    virtual
    void
    parse_parameters (ParameterHandler &prm);
};

```

The kind of information these functions need to provide is discussed in the documentation of this interface class at `aspect::GravityModel::Interface`. The first needs to return a gravity vector at a given position, whereas the second is called at the beginning of each time step, for example to allow a model to update itself based on the current time or the solution of the previous time step. The purpose of the last two functions has been discussed in the general overview of plugins above.

6.3.5 Initial conditions

The initial conditions model is responsible for describing the initial temperature distribution throughout the domain. It essentially has to provide a function that for each point can return the initial temperature. Note that the model (1)–(3) does not require initial values for the pressure or velocity. However, if coefficients are nonlinear, one can significantly reduce the number of initial nonlinear iterations if a good guess for them is available; consequently, ASPECT initializes the pressure with the adiabatically computed hydrostatic pressure, and a zero velocity. Neither of these two has to be provided by the objects considered in this section.

To implement a new initial conditions model, you need to overload the `aspect::InitialConditions::Interface` class and use the `ASPECT_REGISTER_INITIAL_CONDITIONS` macro to register your new class. The implementation of the new class should be in namespace `aspect::InitialConditions`.

Specifically, your new class needs to implement the following basic interface:

```
template <int dim>
class aspect::InitialConditions::Interface
{
public:
    void
    initialize (const GeometryModel::Interface<dim>      &geometry_model,
                const BoundaryTemperature::Interface<dim> &boundary_temperature,
                const AdiabaticConditions<dim>           &adiabatic_conditions);

    virtual
    double
    initial_temperature (const Point<dim> &position) const = 0;

    static
    void
    declare_parameters (ParameterHandler &prm);

    virtual
    void
    parse_parameters (ParameterHandler &prm);
};
```

The meaning of the first class should be clear. The purpose of the last two functions has been discussed in the general overview of plugins above.

6.3.6 Prescribed velocity boundary conditions

Most of the time, one chooses relatively simple boundary values for the velocity: either a zero boundary velocity, a tangential flow model in which the tangential velocity is unspecified but the normal velocity is zero at the boundary, or one in which all components of the velocity are unspecified (i.e., for example, an outflow or inflow condition where the total stress in the fluid is assumed to be zero). However, sometimes we want to choose a velocity model in which the velocity on the boundary equals some prescribed value. A typical example is one in which plate velocities are known, for example their current values or historical reconstructions. In that case, one needs a model in which one needs to be able to evaluate the velocity at individual points at the boundary. This can be implemented via plugins.

To implement a new boundary velocity model, you need to overload the `aspect::VelocityBoundaryConditions::Interface` class and use the `ASPECT_REGISTER_VELOCITY_BOUNDARY_CONDITIONS` macro to register your new class. The implementation of the new class should be in namespace `aspect::VelocityBoundaryConditions`.

Specifically, your new class needs to implement the following basic interface:

```
template <int dim>
```

```

class aspect::VelocityBoundaryConditions::Interface
{
public:
    virtual
    Tensor<1,dim>
    boundary_velocity (const Point<dim> &position) const = 0;

    virtual
    void
    initialize (const GeometryModel::Interface<dim> &geometry_model);

    virtual
    void
    update ();

    static
    void
    declare_parameters (ParameterHandler &prm);

    virtual
    void
    parse_parameters (ParameterHandler &prm);
};


```

The first of these functions needs to provide the velocity at the given point. The next two are other member functions that can (but need not) be overloaded if a model wants to do initialization steps at the beginning of the program or at the beginning of each time step. Examples are models that need to call an external program to obtain plate velocities for the current time, or from historical records, in which case it is far cheaper to do so only once at the beginning of the time step than for every boundary point separately. See, for example, the [aspect::VelocityBoundaryConditions::GPPlates](#) class.

The remaining functions are obvious, and are also discussed in the documentation of this interface class at [aspect::VelocityBoundaryConditions::Interface](#). The purpose of the last two functions has been discussed in the general overview of plugins above.

6.3.7 Temperature boundary conditions

The boundary conditions are responsible for describing the temperature values at those parts of the boundary at which the temperature is fixed (see Section 6.3.3 for how it is determined which parts of the boundary this applies to).

To implement a new boundary conditions model, you need to overload the [aspect::BoundaryTemperature::Interface](#) class and use the `ASPECT_REGISTER_BOUNDARY_TEMPERATURE_MODEL` macro to register your new class. The implementation of the new class should be in namespace `aspect::BoundaryTemperature`.

Specifically, your new class needs to implement the following basic interface:

```

template <int dim>
class aspect::BoundaryTemperature::Interface
{
public:
    virtual
    double
    temperature (const GeometryModel::Interface<dim> &geometry_model,
                 const unsigned int                  boundary_indicator,
                 const Point<dim>                  &location) const = 0;

    virtual

```

```

    double minimal_temperature () const = 0;

    virtual
    double maximal_temperature () const = 0;

    static
    void
    declare_parameters (ParameterHandler &prm);

    virtual
    void
    parse_parameters (ParameterHandler &prm);
};


```

The first of these functions needs to provide the fixed temperature at the given point. The geometry model and the boundary indicator of the particular piece of boundary on which the point is located is also given as a hint in determining where this point may be located; this may, for example, be used to determine if a point is on the inner or outer boundary of a spherical shell. The remaining functions are obvious, and are also discussed in the documentation of this interface class at [aspect::BoundaryTemperature::Interface](#). The purpose of the last two functions has been discussed in the general overview of plugins above.

6.3.8 Postprocessors: Evaluating the solution after each time step

Postprocessors are arguably the most complex and powerful of the plugins available in ASPECT since they do not only passively provide any information but can actually compute quantities derived from the solution. They are executed once at the end of each time step and, unlike all the other plugins discussed above, there can be an arbitrary number of active postprocessors in the same program (for the plugins discussed in previous sections it was clear that there is always exactly one material model, geometry model, etc.).

Motivation. The original motivation for postprocessors is that the goal of a simulation is of course not the simulation itself, but that we want to do something with the solution. Examples for already existing postprocessors are:

- Generating output in file formats that are understood by visualization programs. This is facilitated by the [aspect::Postprocess::Visualization](#) class and a separate class of visualization postprocessors, see Section 6.3.9.
- Computing statistics about the velocity field (e.g., computing minimal, maximal, and average velocities), temperature field (minimal, maximal, and average temperatures), or about the heat fluxes across boundaries of the domain. This is provided by the [aspect::Postprocess::VelocityStatistics](#), [aspect::Postprocess::TemperatureStatistics](#), [aspect::Postprocess::HeatFluxStatistics](#) classes, respectively.

Since writing this text, there may have been other additions as well.

However, postprocessors can be more powerful than this. For example, while the ones listed above are by and large stateless, i.e., they do not carry information from one invocation at one timestep to the next invocation,⁴² there is nothing that prohibits postprocessors from doing so. For example, the following ideas would fit nicely into the postprocessor framework:

- *Passive particles:* If one would like to follow the trajectory of material as it is advected along with the flow field, one technique is to use particles. To implement this, one would start with an initial population of particles distributed in a certain way, for example close to the core-mantle boundary. At the end of each time step, one would then need to move them forward with the flow field by one

⁴²This is not entirely true. The visualization plugin keeps track of how many output files it has already generated, so that they can be numbered consecutively.

time increment. As long as these particles do not affect the flow field (i.e., they do not carry any information that feeds into material properties; in other words, they are *passive*), their location could well be stored in a postprocessor object and then be output in periodic intervals for visualization. In fact, such a passive particle postprocessor is already available.

- *Surface or crustal processes:* Another possibility would be to keep track of surface or crustal processes induced by mantle flow. An example would be to keep track of the thermal history of a piece of crust by updating it every time step with the heat flux from the mantle below. One could also imagine integrating changes in the surface topography by considering the surface divergence of the surface velocity computed in the previous time step: if the surface divergence is positive, the topography is lowered, eventually forming a trench; if the divergence is negative, a mountain belt eventually forms.

In all of these cases, the essential limitation is that postprocessors are *passive*, i.e., that they do not affect the simulation but only observe it.

The statistics file. Postprocessors fall into two categories: ones that produce lots of output every time they run (e.g., the visualization postprocessor), and ones that only produce one, two, or in any case a small and fixed number of often numerical results (e.g., the postprocessors computing velocity, temperature, or heat flux statistics). While the former are on their own in implementing how they want to store their data to disk, there is a mechanism in place that allows the latter class of postprocessors to store their data into a central file that is updated at the end of each time step, after all postprocessors are run.

To this end, the function that executes each of the postprocessors is given a reference to a `dealii::TableHandler` object that allows to store data in named columns, with one row for each time step. This table is then stored in the `statistics` file in the directory designated for output in the input parameter file. It allows for easy visualization of trends over all time steps. To see how to put data into this statistics object, take a look at the existing postprocessor objects.

Note that the data deposited into the statistics object need not be numeric in type, though it often is. An example of text-based entries in this table is the visualization class that stores the name of the graphical output file written in a particular time step.

Implementing a postprocessor. Ultimately, implementing a new postprocessor is no different than any of the other plugins. Specifically, you'll have to write a class that overloads the `aspect::Postprocess::Interface` base class and use the `ASPECT_REGISTER_POSTPROCESSOR` macro to register your new class. The implementation of the new class should be in namespace `aspect::Postprocess`.

In reality, however, implementing new postprocessors is often more difficult. Primarily, this difficulty results from two facts:

- Postprocessors are not self-contained (only providing information) but in fact need to access the solution of the model at each time step. That is, of course, the purpose of postprocessors, but it requires that the writer of a plugin has a certain amount of knowledge of how the solution is computed by the main `Simulator` class, and how it is represented in data structures. To alleviate this somewhat, and to insulate the two worlds from each other, postprocessors do not directly access the data structures of the simulator class. Rather, in addition to deriving from the `aspect::Postprocess::Interface` base class, postprocessors also derive from the `SimulatorAccess` class that has a number of member functions postprocessors can call to obtain read-only access to some of the information stored in the main class of ASPECT. See [the documentation of this class](#) to see what kind of information is available to postprocessors. See also Section 6.1 for more information about the `SimulatorAccess` class.
- Writing a new postprocessor typically requires a fair amount of knowledge how to leverage the DEAL.II library to extract information from the solution. The existing postprocessors are certainly good examples to start from in trying to understand how to do this.

Given these comments, the interface a postprocessor class has to implement is rather basic:

```

template <int dim>
class aspect::Postprocess::Interface
{
public:
    virtual std::pair<std::string, std::string>
    execute (TableHandler &statistics) = 0;

    virtual void
    save (std::map<std::string, std::string> &status_strings) const;

    virtual void
    load (const std::map<std::string, std::string> &status_strings);

    static void
    declare_parameters (ParameterHandler &prm);

    virtual void
    parse_parameters (ParameterHandler &prm);
};

```

The purpose of these functions is described in detail in the documentation of the [aspect::Postprocess::Interface](#) class. While the first one is responsible for evaluating the solution at the end of a time step, the `save`/`load` functions are used in checkpointing the program and restarting it at a previously saved point during the simulation. The first of these functions therefore needs to store the status of the object as a string under a unique key in the database described by the argument, while the latter function restores the same state as before by looking up the status string under the same key. The default implementation of these functions is to do nothing; postprocessors that do have non-static member variables that contain a state need to overload these functions.

There are numerous postprocessors already implemented. If you want to implement a new one, it would be helpful to look at the existing ones to see how they implement their functionality.

Postprocessors and checkpoint/restart. Postprocessors have `save()` and `load()` functions that are used to write the data a postprocessor has into a checkpoint file, and to load it again upon restart. This is important since many postprocessors store some state – say, a temporal average over all the time steps seen so far, or the number of the last graphical output file generated so that we know how the next one needs to be numbered.

The typical case is that this state is the same across all processors of a parallel computation. Consequently, what ASPECT writes into the checkpoint file is only the state obtained from the postprocessors on processor 0 of a parallel computation. On restart, all processors read from the same file and the postprocessors on *all* processors will be initialized by what the same postprocessor on processor 0 wrote.

There are situations where postprocessors do in fact store complementary information on different processors. At the time of writing this, one example is the postprocessor that supports advecting passive particles along the velocity field: on every processor, it handles only those particles that lie inside the part of the domain that is owned by this MPI rank. The serialization approach outlined above can not work in this case, for obvious reasons. In cases like this, one needs to implement the `save()` and `load()` differently than usual: one needs to put all variables that are common across processors into the maps of string as usual, but one then also needs to save all state that is different across processors, from all processors. There are two ways: If the amount of data is small, you can use MPI communications to send the state of all processors

to processor zero, and have processor zero store it in the result so that it gets written into the checkpoint file; in the `load()` function, you will then have to identify which part of the text written by processor 0 is relevant to the current processor. Or, if your postprocessor stores a large amount of data, you may want to open a restart file specifically for this postprocessor, use MPI I/O or other ways to write into it, and do the reverse operation in `load()`.

Note that this approach requires that ASPECT actually calls the `save()` function on all processors. This in fact happens – though ASPECT also discards the result on all but processor zero.

6.3.9 Visualization postprocessors

As mentioned in the previous section, one of the postprocessors that are already implemented in ASPECT is the `aspect::Postprocess::Visualization` class that takes the solution and outputs it as a collection of files that can then be visualized graphically, see Section 4.4. The question is which variables to output: the solution of the basic equations we solve here is characterized by the velocity, pressure and temperature; on the other hand, we are frequently interested in derived, spatially and temporally variable quantities such as the viscosity for the actual pressure, temperature and strain rate at a given location, or seismic wave speeds.

ASPECT already implements a good number of such derived quantities that one may want to visualize. On the other hand, always outputting *all* of them would yield very large output files, and would furthermore not scale very well as the list continues to grow. Consequently, as with the postprocessors described in the previous section, what *can* be computed is implemented in a number of plugins and what *is* computed is selected in the input parameter file (see Section A.138).

Defining visualization postprocessors works in much the same way as for the other plugins discussed in this section. Specifically, an implementation of such a plugin needs to be a class that derives from interface classes, should by convention be in namespace `aspect::Postprocess::VisualizationPostprocessors`, and is registered using a macro, here called `ASPECT_REGISTER_VISUALIZATION_POSTPROCESSOR`. Like the postprocessor plugins, visualization postprocessors can derive from class `aspect::Postprocess::SimulatorAccess` if they need to know specifics of the simulation such as access to the material models and to get access to the introspection facility outlined in Section 6.1. A typical example is the plugin that produces the viscosity as a spatially variable field by evaluating the viscosity function of the material model using the pressure, temperature and location of each visualization point (implemented in the `aspect::Postprocess::VisualizationPostprocessors::Viscosity` class). On the other hand, a hypothetical plugin that simply outputs the norm of the strain rate $\sqrt{\varepsilon(\mathbf{u}) : \varepsilon(\mathbf{u})}$ would not need access to anything but the solution vector (which the plugin’s main function is given as an argument) and consequently is not derived from the `aspect::Postprocess::SimulatorAccess` class.⁴³

Visualization plugins can come in two flavors:

- *Plugins that compute things from the solution in a point-wise way:* The classes in this group are derived not only from the respective interface class (and possibly the `SimulatorAccess` class) but also from the deal.II class `DataPostprocessor` or any of the classes like `DataPostprocessorScalar` or `DataPostprocessorVector`. These classes can be thought of as filters: DataOut will call a function in them for every cell and this function will transform the values or gradients of the solution and other information such as the location of quadrature points into the desired quantity to output. A typical case would be if the quantity $g(x)$ you want to output can be written as a function $g(x) = G(u(x), \nabla u(x), x, \dots)$ in a point-wise sense where $u(x)$ is the value of the solution vector (i.e., the velocities, pressure, temperature, etc) at an evaluation point. In the context of this program an example would be to output the density of the medium as a spatially variable function since this is a quantity that for realistic media depends point-wise on the values of the solution.

To sum this, slightly confusing multiple inheritance up, visualization postprocessors do the following:

⁴³The actual plugin `aspect::Postprocess::VisualizationPostprocessors::StrainRate` only computes $\sqrt{\varepsilon(\mathbf{u}) : \varepsilon(\mathbf{u})}$ in the incompressible case. In the compressible case, it computes $\sqrt{[\varepsilon(\mathbf{u}) - \frac{1}{3}(\text{tr } \varepsilon(\mathbf{u}))\mathbf{I}] : [\varepsilon(\mathbf{u}) - \frac{1}{3}(\text{tr } \varepsilon(\mathbf{u}))\mathbf{I}]}$ instead. To test whether the model is compressible or not, the plugin needs access to the material model object, which the class gains by deriving from `aspect::Postprocess::SimulatorAccess` and then calling `this->get_material_model().is_compressible()`.

- If necessary, they derive from `aspect::Postprocess::SimulatorAccess`.
- They derive from `aspect::Postprocess::VisualizationPostprocessors::Interface`. The functions of this interface class are all already implemented as doing nothing in the base class but can be overridden in a plugin. Specifically, the following functions exist:

```
class Interface
{
public:
    static
    void
    declare_parameters (ParameterHandler &prm);

    virtual
    void
    parse_parameters (ParameterHandler &prm);

    virtual
    void save (std::map<std::string, std::string> &status_strings) const;

    virtual
    void load (const std::map<std::string, std::string> &status_strings);
};
```

- They derive from either the `dealii::DataPostprocessor` class, or the simpler to use `dealii::DataPostprocessor` or `dealii::DataPostprocessorVector` classes. For example, to derive from the second of these classes, the following interface functions has to be implemented:

```
class dealii::DataPostprocessorScalar
{
public:
    virtual
    void
    compute_derived_quantities_vector
        (const std::vector<Vector<double> > &uh,
         const std::vector<std::vector<Tensor<1,dim> > > &duh,
         const std::vector<std::vector<Tensor<2,dim> > > &dduh,
         const std::vector<Point<dim> > &normals,
         const std::vector<Point<dim> > &evaluation_points,
         std::vector<Vector<double> > &computed_quantities) const;
};
```

What this function does is described in detail in the deal.II documentation. In addition, one has to write a suitable constructor to call `dealii::DataPostprocessorScalar::DataPostprocessorScalar`.

- *Plugins that compute things from the solution in a cell-wise way:* The second possibility is for a class to not derive from `dealii::DataPostprocessor` but instead from the `aspect::Postprocess::VisualizationPostprocessors::Cell` class. In this case, a visualization postprocessor would generate and return a vector that consists of one element per cell. The intent of this option is to output quantities that are not point-wise functions of the solution but instead can only be computed as integrals or other functionals on a per-cell basis. A typical case would be error estimators that do depend on the solution but not in a point-wise sense; rather, they yield one value per cell of the mesh. See the documentation of the `CellDataVectorCreator` class for more information.

If all of this sounds confusing, we recommend consulting the implementation of the various visualization plugins that already exist in the ASPECT sources, and using them as a template.

6.3.10 Mesh refinement criteria

Despite research since the mid-1980s, it isn't completely clear how to refine meshes for complex situations like the ones modeled by ASPECT. The basic problem is that mesh refinement criteria either can refine based on some variable such as the temperature, the pressure, the velocity, or a compositional field, but that oftentimes this by itself is not quite what one wants. For example, we know that Earth has discontinuities, e.g., at 440km and 610km depth. In these places, densities and other material properties suddenly change. Their resolution in computation models is important as we know that they affect convection patterns. At the same time, there is only a small effect on the primary variables in a computation – maybe a jump in the second or third derivative, for example, but not a discontinuity that would be clear to see. As a consequence, automatic refinement criteria do not always refine these interfaces as well as necessary.

To alleviate this, ASPECT has plugins for mesh refinement. Through the parameters in Section A.107, one can select when to refine but also which refinement criteria should be used and how they should be combined if multiple refinement criteria are selected. Furthermore, through the usual plugin mechanism, one can extend the list of available mesh refinement criteria (see the parameter “Strategy” in Section A.107). Each such plugin is responsible for producing a vector of values (one per active cell on the current processor, though only those values for cells that the current processor owns are used) with an indicator of how badly this cell needs to be refined: large values mean that the cell should be refined, small values that the cell may be coarsened away.

To implement a new mesh refinement criterion, you need to overload the `aspect::MeshRefinement::Interface` class and use the `ASPECT_REGISTER_MESH_REFINEMENT_CRITERION` macro to register your new class. The implementation of the new class should be in namespace `aspect::MeshRefinement`.

Specifically, your new class needs to implement the following basic interface:

```
template <int dim>
class aspect::MeshRefinement::Interface
{
public:
    virtual
    void
    execute (Vector<float> &error_indicators) const = 0;

    static
    void
    declare_parameters (ParameterHandler &prm);

    virtual
    void
    parse_parameters (ParameterHandler &prm);
};
```

The first of these functions computes the set of refinement criteria (one per cell) and returns it in the given argument. Typical examples can be found in the existing implementations in the `source/mesh_refinement` directory. As usual, your termination criterion implementation will likely need to be derived from the `SimulatorAccess` to get access to the current state of the simulation.

The remaining functions are obvious, and are also discussed in the documentation of this interface class at `aspect::MeshRefinement::Interface`. The purpose of the last two functions has been discussed in the general overview of plugins above.

6.3.11 Criteria for terminating a simulation

ASPECT allows for different ways of terminating a simulation. For example, the simulation may have reached a final time specified in the input file. However, it also allows for ways to terminate a simulation when it has reached a steady state (or, rather, some criterion determines that it is close enough to steady

state), or by an external action such as placing a specially named file in the output directory. The criteria determining termination of a simulation are all implemented in plugins. The parameters describing these criteria are listed in Section A.158.

To implement a termination criterion, you need to overload the `aspect::TerminationCriteria::Interface` class and use the `ASPECT_REGISTER_TERMINATION_CRITERION` macro to register your new class. The implementation of the new class should be in namespace `aspect::TerminationCriteria`.

Specifically, your new class needs to implement the following basic interface:

```
template <int dim>
class aspect::TerminationCriteria::Interface
{
public:
    virtual
    bool
    execute () const = 0;

    static
    void
    declare_parameters (ParameterHandler &prm);

    virtual
    void
    parse_parameters (ParameterHandler &prm);
};
```

The first of these functions returns a value that indicates whether the simulation should be terminated. Typical examples can be found in the existing implementations in the `source/termination_criteria` directory. As usual, your termination criterion implementation will likely need to be derived from the `SimulatorAccess` to get access to the current state of the simulation.

The remaining functions are obvious, and are also discussed in the documentation of this interface class at `aspect::TerminationCriteria::Interface`. The purpose of the last two functions has been discussed in the general overview of plugins above.

6.4 Compatibility of plugins with newer ASPECT versions

We strive to maintain compatibility for user written plugins with new versions of ASPECT for as long as possible. However, occasionally we have to restructure interface classes to improve ASPECT further. This is in particular true for new major versions. In order to allow running old plugins with newer ASPECT versions we provide scripts that can automatically update existing plugins to the new syntax. Executing `doc/update_source_files.sh` with one or more plugin files as arguments will create a backup of the old file (named `old_filename.bak`), and replace the existing file with a version that should work with the current ASPECT version. Using this script would look like this:

```
bash doc/update_source_files.sh cookbooks/finite_strain/finite_strain.cc
```

Note: Not all text replacements are unique, and the structure of plugin files allows for constructs the script can not properly parse. Thus, it is important that you check your updated plugin file for errors. That being said, all plugin files in the main ASPECT repository are updated successfully using this script.

6.5 Extending ASPECT through the signals mechanism

Not all things you may want to do fit neatly into the list of plugins of the previous sections. Rather, there are cases where you may want to change things that are more of the one-off kind and that require code

that is at a lower level and requires more knowledge about ASPECT's internal workings. For such changes, we still want to stick with the general principle outlined at the beginning of Section 6: You should be able to make all of your changes and extensions in your own files, without having to modify ASPECT's own sources.

To support this, ASPECT uses a “signals” mechanism. Signals are, in essence, objects that represent *events*, for example the fact that the solver has finished a time step. The core of ASPECT defines a number of such signals, and *triggers* them at the appropriate points. The idea of signals is now that you can *connect* to them: you can tell the signal that it should call a particular function every time the signal is triggered. The functions that are connected to a signal are called “slots” in common diction. One, several, or no slots may be connected to each signal.

There are two kinds of signals that ASPECT provides:

- Signals that are triggered at startup of the program: These are, in essence, signals that live in some kind of global scope. Examples are signals that declare additional parameters for use in input files, or that read the values of these parameters from a `ParameterHandler` object. These signals are static member variables of the structure that contains them and consequently exist only once for the entire program.
- Signals that reference specific events that happen inside a simulator object. These are regular member variables of the structure that contains them, and because each simulator object has such a structure, the signals exist once per simulator object. (Which in practice is only once per program, of course.)

For both of these kinds, a user-written plugin file can (but does not need) to register functions that connect functions in this file (i.e., slots) to their respective signals.

In the first case, code that registers slots with global signals would look like this:

```
// A function that will be called at the time when parameters are declared.
// It receives the dimension in which ASPECT will be run as the first argument,
// and the ParameterHandler object that holds the runtime parameter
// declarations as second argument.
void declare_parameters(const unsigned int dim,
                      ParameterHandler &prm)
{
    prm.declare_entry("My parameter", ...);
}

// The same for parsing parameters. 'my_parameter' is a parameter
// that stores something we want to read from the input file
// and use in other functions in this file (which we don't show here).
// For simplicity, we assume that it is an integer.
//
// The function also receives a first argument that contains all
// of the other (already parsed) arguments of the simulation, in
// case what you want to do here wants to refer to other parameters.
int my_parameter;

template <int dim>
void parse_parameters(const Parameters<dim> &parameters,
                      ParameterHandler &prm)
{
    my_parameter = prm.get_integer ("My parameter");
}

// Now have a function that connects slots (i.e., the two functions
```

```

// above) to the static signals. Do this for both the 2d and 3d
// case for generality.
void parameter_connector ()
{
    SimulatorSignals<2>::declare_additional_parameters.connect (&declare_parameters);
    SimulatorSignals<3>::declare_additional_parameters.connect (&declare_parameters);

    SimulatorSignals<2>::parse_additional_parameters.connect (&parse_parameters<2>);
    SimulatorSignals<3>::parse_additional_parameters.connect (&parse_parameters<3>);
}

// Finally register the connector function above to make sure it gets run
// whenever we load a user plugin that is mentioned among the additional
// shared libraries in the input file:
ASPECT_REGISTER_SIGNALS_PARAMETER_CONNECTOR(parameter_connector)

```

The second kind of signal can be connected to once a simulator object has been created. As above, one needs to define the slots, define a connector function, and register the connector function. The following gives an example:

```

// A function that is called at the end of creating the current constraints
// on degrees of freedom (i.e., the constraints that describe, for example,
// hanging nodes, boundary conditions, etc).
template <int dim>
void post_constraints_creation (const SimulatorAccess<dim> &simulator_access,
                                ConstraintMatrix &current_constraints)
{
    ...; // do whatever you want to do here
}

// A function that is called from the simulator object and that can connect
// a slot (such as the function above) to any of the signals declared in the
// structure passed as argument:
template <int dim>
void signal_connector (SimulatorSignals<dim> &signals)
{
    signals.post_constraints_creation.connect (&post_constraints_creation<dim>);
}

// Finally register the connector function so that it is called whenever
// a simulator object has been set up. For technical reasons, we need to
// register both 2d and 3d versions of this function:
ASPECT_REGISTER_SIGNALS_CONNECTOR(signal_connector<2>,
                                  signal_connector<3>)

```

As mentioned above, each signal may be connected to zero, one, or many slots. Consequently, you could have multiple plugins each of which connect to the same slot, or the connector function above may just connect multiple slots (i.e., functions in your program) to the same signal.

So what could one do in a place like this? One option would be to just monitor what is going on, e.g., in code like this that simply outputs into the statistics file (see Section 4.4.2):

```

template <int dim>
void post_constraints_creation (const SimulatorAccess<dim> &simulator_access,

```

```

        ConstraintMatrix &current_constraints)
{
    simulator_access.get_statistics_object()
        .add_value ("number of constraints",
                    current_constraints.n_constraints());
}

```

This will produce, for every time step (because this is how often the signal is called) an entry in a new column in the statistics file that records the number of constraints. On the other hand, it is equally possible to also modify the constraints object at this point. An application would be if you wanted to run a simulation where you prescribe the velocity in a part of the domain, e.g., for a subducting slab (see Section 5.2.9).

Signals exist for various waypoints in a simulation and you can consequently monitor and change what is happening inside a simulation by connecting your own functions to these signals. It would be pointless to list here what signals actually exist – simply refer to the documentation of the [SimulatorSignals class](#) for a complete list of signals you can connect to.

As a final note, it is generally true that writing functions that can connect to signals require significantly more internal knowledge of the workings of ASPECT than writing plugins through the mechanisms outlined above. It also allows to affect the course of a simulation by working on the internal data structures of ASPECT in ways that are not available to the largely passive and reactive plugins discussed in previous sections. With this obviously also comes the potential for trouble. On the other hand, it also allows to do things with ASPECT that were not initially intended by the authors, and that would be hard or impossible to implement through plugins. An example would be to couple different codes by exchanging details of the internal data structures, or even update the solution vectors using information received from another code.

Note: Chances are that if you think about using the signal mechanism, there is not yet a signal that is triggered at exactly the point where you need it. Consequently, you will be tempted to just put your code into the place where it fits inside ASPECT where it fits best. This is poor practice: it prevents you from upgrading to a newer version of ASPECT at a later time because this would overwrite the code you inserted.

Rather, a more productive approach would be to either define a new signal that is triggered where you need it, and connect a function (slot) in your own plugin file to this signal using the mechanisms outlined above. Then send the code that defines and triggers the signal to the developers of ASPECT to make sure that it is also included in the next release. Alternatively, you can also simply ask on the mailing lists for someone to add such a signal in the place where you want it. Either way, adding signals is something that is easy to do, and we would much rather add signals than have people who modify the ASPECT source files for their own needs and are then stuck on a particular version.

6.6 Extending the basic solver

The core functionality of the code, i.e., that part of the code that implements the time stepping, assembles matrices, solves linear and nonlinear systems, etc., is in the `aspect::Simulator` class (see the [doxygen documentation of this class](#)). Since the implementation of this class has more than 3,000 lines of code, it is split into several files that are all located in the `source/simulator` directory. Specifically, functionality is split into the following files:

- `source/simulator/core.cc`: This file contains the functions that drive the overall algorithm (in particular `Simulator::run`) through the main time stepping loop and the functions immediately called by `Simulator::run`.
- `source/simulator/assembly.cc`: This is where all the functions are located that are related to assembling linear systems.

- `source/simulator/solver.cc`: This file provides everything that has to do with solving and preconditioning the linear systems.
- `source/simulator/initial_conditions.cc`: The functions in this file deal with setting initial conditions for all variables.
- `source/simulator/checkpoint_restart.cc`: The location of functionality related to saving the current state of the program to a set of files and restoring it from these files again.
- `source/simulator/helper_functions.cc`: This file contains a set of functions that do the odd thing in support of the rest of the simulator class.
- `source/simulator/parameters.cc`: This is where we define and read run-time parameters that pertain to the top-level functionality of the program.

Obviously, if you want to extend this core functionality, it is useful to first understand the numerical methods this class implements. To this end, take a look at the paper that describes these methods, see [KHB12]. Further, there are two predecessor programs whose extensive documentation is at a much higher level than the one typically found inside ASPECT itself, since they are meant to teach the basic components of convection simulators as part of the DEAL.II tutorial:

- The step-31 program at https://www.dealii.org/developer/doxygen/deal.II/step_31.html: This program is the first version of a convection solver. It does not run in parallel, but it introduces many of the concepts relating to the time discretization, the linear solvers, etc.
- The step-32 program at https://www.dealii.org/developer/doxygen/deal.II/step_32.html: This is a parallel version of the step-31 program that already solves on a spherical shell geometry. The focus of the documentation in this program is on the techniques necessary to make the program run in parallel, as well as some of the consequences of making things run with realistic geometries, material models, etc.

Neither of these two programs is nearly as modular as ASPECT, but that was also not the goal in creating them. They will, however, serve as good introductions to the general approach for solving thermal convection problems.

Note: Neither this manual, nor the documentation in ASPECT makes much of an attempt at teaching how to use the DEAL.II library upon which ASPECT is built. Nevertheless, you will likely have to know at least the basics of DEAL.II to successfully work on the ASPECT code. We refer to the resources listed at the beginning of this section as well as references [BHK07, BHK12].

6.7 Testing ASPECT

ASPECT makes use of a large suite of tests to ensure correct behavior. The test suite is run automatically for each change to the Github repository, and it is good practice to add new tests for any new functionality.

6.7.1 Running tests

In order to run the tests, it is necessary to have either Diff or Numdiff to compare the results to the known good case. Diff is installed by default on most Linux systems, and Numdiff is usually available as a package so this is not a severe limitation. While it is possible to use Diff, Numdiff is preferred due to being able to more accurately identify whether a variation in numerical output is significant. The test suite is run using the `ctest` program that comes with `cmake`, and should therefore be available on all systems that have compiled ASPECT.

After running `cmake` and then compiling ASPECT, you can run the testsuite by saying `ctest`. By default, this will only run a small subset of all tests given that both setting up all tests (several hundred) and running them takes a non-trivial amount of time. To set up the full test suite, you can run

```
make setup_tests
```

in the build directory. To run the entire set of tests, then execute

```
ctest
```

Unless you have a very fast machine with lots of processors, running the entire testsuite will take hours, though it can be made substantially faster if you use

```
ctest -j <N>
```

where $<N>$ is the number of tests you want `ctest` to run in parallel; you may want to choose $<N>$ equal to or slightly smaller than the number of processors you have. Alternatively, you can run only a subset of all tests by saying

```
ctest -R <regex>
```

where $<\text{regex}>$ is a regular expression and the only tests that will be run are those whose names match the expression.

When `ctest` runs a test, it will ultimately output results of the form

```
build> ctest -R additional_outputs
Test project /home/fac/f/bangerth/p/deal.II/1/projects/build
  Start 1: additional_outputs
  1/3 Test #1: additional_outputs ..... Passed    2.03 sec
  Start 2: additional_outputs_02
  2/3 Test #2: additional_outputs_02 ..... Passed    1.84 sec
  Start 3: additional_outputs_03
  3/3 Test #3: additional_outputs_03 ..... Passed    1.91 sec

100% tests passed, 0 tests failed out of 3

Total Test time (real) = 5.88 sec
```

While the small default subset of tests should work on almost all platforms, you will find that some of the tests fail on your machine when you run the entire testsuite. This is because success or failure of a test is determined by looking at whether its output matches the one saved at the time when the test was written to the last digit, both as far as numerical output in floating point precision is concerned (e.g., for heat fluxes or other things we compute via postprocessors) as well as for integers such as the number of iterations that is printed in the screen output.⁴⁴ Unfortunately, systems almost always differ by compiler version, processor type and version, system libraries, etc, that can all lead to small changes in output – generally (and hopefully!) not large enough to produce *qualitatively* different results, but *quantitatively* large enough to change the number of iterations necessary to reach a specific tolerance, or to change the computed heat flux by one part in a million. This leads to `ctest` reporting that a test failed, when in reality it produced output that is qualitatively correct.

Given that some tests are expected to fail on any given system raises the question why it makes sense to have tests at all? The answer is that there is *one* system on which all tests are supposed to succeed: This system is a machine that churns through all tests every time someone proposes a change to the ASPECT code base via the ASPECT GitHub page.⁴⁵ Upon completion of the test suite, both the general summary

⁴⁴This is not actually completely true. Rather, if `cmake` finds a program called `numdiff` on your system, it uses `numdiff` to compare the output of a test run against the saved output, and calls two files the same if all numbers differ by no more than some tolerance.

⁴⁵This is again not completely true: The test machine will only go to work for pull requests by a set of trusted maintainers since the test machine will execute code that is proposed as part of the pull request – posing a security risk if anyone's patches were allowed to run on that system. For pull requests by others, one of the trusted maintainers has to specifically request a test run, and this will usually happen as soon as the patch developer indicates the patch is ready for review.

(pass/fail) and a full verbose log will available from the GitHub page. Because the official test setup is set up in a Docker container, it is simple to replicate the results on a local machine. To this end, follow the instructions in Section 3.1 to set up Docker, and then run the following command in any terminal (replace ASPECT_SOURCE_DIR with the path to your ASPECT directory):

```
docker run -v ASPECT_SOURCE_DIR:/home/dealii/aspect \
--name=aspect-tester --rm -it \
dealii/dealii:v8.5.0-gcc-mpi-fulldepscandi-debugrelease \
bash /home/dealii/aspect/cmake/compile_and_update_tests.sh
```

This command executes the shell script `cmake/compile_and_update_tests.sh` *inside* the docker container that contains the official ASPECT test system. Note that by mounting your ASPECT folder into the container you are actually updating the reference test results on the host system (i.e. your computer).

6.7.2 Writing tests

To write a test for a new feature, take one of the existing parameter files in the `tests/` directory, modify it to use the new feature, check that the new feature does what it is supposed to do, and then just add the parameter file to the tests directory. You will then need to add another folder to that directory that is named like the parameter file, and add the model output files that prove that the feature is working (rename `log.txt` to `screen-output` for historical reasons). The test and output files should be as small and quick to run as possible.

When you make a new test part of a pull request on GitHub, then as explained above that will lead to a run of all tests – including your new one – on a “reference machine”. The reference machine that runs the tests may of course produce slightly different results than the machine on which a pull request was developed and from which the output was taken. If this has been confirmed to be the source of a failed test run, a file that contains the differences between the patch content and the tester output will be available from the GitHub page and can be applied to the patch to make it pass the tester. On the other hand, if a change leads to even a single *existing* test failing on that system, then we know that some more investigation as to the causes is necessary.

6.8 Contributing to ASPECT’s development

To end this section, let us repeat something already stated in the introduction:

Note: ASPECT is a community project. As such, we strongly encourage contributions from the community to improve this code over time. Obvious candidates for such contributions are implementations of new plugins as discussed in Section 6.3, since they are typically self-contained and do not require much knowledge of the details of the remaining code. Other much appreciated contributions are new test models or benchmarks, extended documentation (every paragraph helps), and in particular fixing typos or updating outdated documentation. Obviously, however, we also encourage contributions to the core functionality in any form!

Let us assume you found something in ASPECT to improve, something you did not understand, or something that is simply wrong. Do something about it! No matter whether you are a C++ expert or first-time user, there are no such things as too-unimportant contributions, and if you struggled with something, it is most likely somebody else will as well. The process of contributing to a new project can be daunting, but we appreciate every contribution and are happy to work with you on improving ASPECT. To get you started we have collected a set of guidelines and advice on how to get involved in the community. To avoid duplication we store these guidelines in a separate file `CONTRIBUTING.md` in the main folder of the repository, and you can also access them online at <https://github.com/geodynamics/aspect/blob/master/CONTRIBUTING.md>. Even if something in that file is not clear, this is an opportunity for you to ask your question on the mailing list (see Section 8, and let us know that file needs improvement.

7 Future plans for ASPECT

We have a number of near-term plans for ASPECT that we hope to implement soon:

- *Better ways to deal with the nonlinearity:* The simplest algorithm to deal with the nonlinearity of the equations is to use the velocity, pressure and temperature of the previous time step to evaluate the coefficients that appear in the flow equations (1)–(2); and the velocity and pressure of the current time step as well as the previous time step’s temperature to evaluate the coefficients in the temperature equation (3). A slight improvement of this scheme – and the one that is the default in ASPECT – is to use velocities, pressures, and temperatures that are extrapolated from previous time steps to the current time. This is an appropriate strategy if the model is not too nonlinear; however, it introduces inaccuracies and limits the size of the time step if coefficients strongly depend on the solution variables.

To avoid this, one can iterate out the equations using either a fixed point or Newton scheme. Both approaches ensure that at the end of a time step, the values of coefficients and solution variables are consistent. On the other hand, one may have to solve the linear systems that describe a time step more than once, increasing the computational effort.

ASPECT currently already implements methods that iterate out the nonlinearity based on various fixed point strategies. These are based on some earlier experiments by Jennifer Worthen [Wor12]. What is missing, though, is a proper Newton method.

- *Faster 3d computations:* Whichever way you look at it, 3d computations are expensive. In parallel computations, the Stokes solve currently takes upward of 90% of the overall wallclock time, suggesting an obvious target for improvements based on better algorithms as well as from profiling the code to find hot spots. In particular, playing with better solver and/or preconditioner options would seem to be a useful goal.
- *More, and more realistic material models:* The number of material models available in ASPECT is sizeable, but it does not include all models that have in the past been used in the literature. On the other hand, how realistic a simulation is depends on how realistic the used material model is. We hope to obtain more descriptions of realistic material descriptions over time, either given analytically or based on table-lookup of material properties.

8 Finding answers to more questions

If you have questions that go beyond this manual, there are a number of resources:

- For questions on the source code of ASPECT, portability, installation, etc., use the ASPECT development mailing list at <http://lists.geodynamics.org/cgi-bin/mailman/listinfo/aspect-devel>. This mailing list is where the ASPECT developers all hang out.
- ASPECT is primarily based on the deal.II library (the dependency on Trilinos and p4est is primarily through deal.II, and not directly visible in the ASPECT source code). If you have particular questions about deal.II, contact the mailing lists described at <https://www.dealii.org/mail.html>.
- In case of more general questions about mantle convection, you can contact the CIG mantle convection mailing lists at <http://lists.geodynamics.org/cgi-bin/mailman/listinfo/cig-MC>.
- If you have specific questions about ASPECT that are not suitable for public and archived mailing lists, you can contact the primary developers:
 - Wolfgang Bangerth: bangerth@colostate.edu.
 - Juliane Dannberg: judannberg@gmail.com
 - Rene Gassmöller: rene.gassmoeller@mailbox.org
 - Timo Heister: heister@clemson.edu.

A Run-time input parameters

A.1 Global parameters

- *Parameter name:* Additional shared libraries

Value:

Default:

Description: A list of names of additional shared libraries that should be loaded upon starting up the program. The names of these files can contain absolute or relative paths (relative to the directory in which you call ASPECT). In fact, file names that do not contain any directory information (i.e., only the name of a file such as <myplugin.so> will not be found if they are not located in one of the directories listed in the LD_LIBRARY_PATH environment variable. In order to load a library in the current directory, use <./myplugin.so> instead.

The typical use of this parameter is so that you can implement additional plugins in your own directories, rather than in the ASPECT source directories. You can then simply compile these plugins into a shared library without having to re-compile all of ASPECT. See the section of the manual discussing writing extensions for more information on how to compile additional files into a shared library.

Possible values: A list of 0 to 4294967295 elements where each element is [an input filename]

- *Parameter name:* Adiabatic surface temperature

Value: 0

Default: 0

Description: In order to make the problem in the first time step easier to solve, we need a reasonable guess for the temperature and pressure. To obtain it, we use an adiabatic pressure and temperature field. This parameter describes what the ‘adiabatic’ temperature would be at the surface of the domain (i.e. at depth zero). Note that this value need not coincide with the boundary condition posed at this point. Rather, the boundary condition may differ significantly from the adiabatic value, and then typically induce a thermal boundary layer.

For more information, see the section in the manual that discusses the general mathematical model.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* CFL number

Value: 1.0

Default: 1.0

Description: In computations, the time step k is chosen according to $k = c \min_K \frac{h_K}{\|u\|_{\infty, K p_T}}$ where h_K is the diameter of cell K , and the denominator is the maximal magnitude of the velocity on cell K times the polynomial degree p_T of the temperature discretization. The dimensionless constant c is called the CFL number in this program. For time discretizations that have explicit components, c must be less than a constant that depends on the details of the time discretization and that is no larger than one. On the other hand, for implicit discretizations such as the one chosen here, one can choose the time step as large as one wants (in particular, one can choose $c > 1$) though a CFL number significantly larger than one will yield rather diffusive solutions. Units: None.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Dimension

Value: 2

Default: 2

Description: The number of space dimensions you want to run this program in. ASPECT can run in 2 and 3 space dimensions.

Possible values: An integer n such that $2 \leq n \leq 4$

- *Parameter name:* `End_time`

Value: 0.0

Default: 5.69e+300

Description: The end time of the simulation. The default value is a number so that when converted from years to seconds it is approximately equal to the largest number representable in floating point arithmetic. For all practical purposes, this equals infinity. Units: Years if the 'Use years in output instead of seconds' parameter is set; seconds otherwise.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `Max_nonlinear_iterations`

Value: 10

Default: 10

Description: The maximal number of nonlinear iterations to be performed.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* `Max_nonlinear_iterations_in_pre-refinement`

Value: 2147483647

Default: 2147483647

Description: The maximal number of nonlinear iterations to be performed in the pre-refinement steps. This does not include the last refinement step before moving to timestep 1. When this parameter has a larger value than max nonlinear iterations, the latter is used.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* `Maximum_relative_increase_in_time_step`

Value: 2147483647

Default: 2147483647

Description: Set a percentage with which the time step is limited to increase. Generally the time step based on the CFL number should be sufficient, but for complicated models which may suddenly drastically change behavior, it may be useful to limit the increase in the time step, without limiting the time step size of the whole simulation to a particular number. For example, if this parameter is set to 50, then that means that the time step can at most increase by 50% from one time step to the next, or by a factor of 1.5. Units: %

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `Maximum_time_step`

Value: 5.69e+300

Default: 5.69e+300

Description: Set a maximum time step size for the solver to use. Generally the time step based on the CFL number should be sufficient, but for complicated models or benchmarking it may be useful to limit the time step to some value. The default value is a value so that when converted from years into seconds it equals the largest number representable by a floating point number, implying an unlimited time step. Units: Years or seconds, depending on the "Use years in output instead of seconds" parameter.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Nonlinear solver scheme

Value: single Advection, single Stokes

Default: single Advection, single Stokes

Description: The kind of scheme used to resolve the nonlinearity in the system. ‘single Advection, single Stokes’ means that no nonlinear iterations are done, and the temperature, compositional fields and Stokes equations are solved exactly once per time step, one after the other. The ‘iterated Advection and Stokes’ scheme iterates this decoupled approach by alternating the solution of the temperature, composition and Stokes systems. The ‘single Advection, iterated Stokes’ scheme solves the temperature and composition equation once at the beginning of each time step and then iterates out the solution of the Stokes equation. The ‘no Advection, iterated Stokes’ scheme only solves the Stokes system, iterating out the solution, and ignores compositions and the temperature equation (careful, the material model must not depend on the temperature or composition; this is mostly useful for Stokes benchmarks). The ‘single Advection, no Stokes’ scheme only solves the temperature and other advection systems once, and instead of solving for the Stokes system, a prescribed velocity and pressure is used. The ‘iterated Advection and Newton Stokes’ scheme iterates by alternating the solution of the temperature, composition and Stokes equations, using Picard iterations for the temperature and composition, and Newton iterations for the Stokes system. The ‘IMPES’ scheme is deprecated and only allowed for reasons of backwards compatibility. It is the same as ‘single Advection, single Stokes’. The ‘iterated IMPES’ scheme is deprecated and only allowed for reasons of backwards compatibility. It is the same as ‘iterated Advection and Stokes’. The ‘iterated Stokes’ scheme is deprecated and only allowed for reasons of backwards compatibility. It is the same as ‘single Advection, iterated Stokes’. The ‘Stokes only’ scheme is deprecated and only allowed for reasons of backwards compatibility. It is the same as ‘no Advection, iterated Stokes’. The ‘Advection only’ scheme is deprecated and only allowed for reasons of backwards compatibility. It is the same as ‘single Advection, no Stokes’. The ‘Newton Stokes’ scheme is deprecated and only allowed for reasons of backwards compatibility. It is the same as ‘iterated Advection and Newton Stokes’.

Possible values: Any one of single Advection, single Stokes, iterated Advection and Stokes, single Advection, iterated Stokes, no Advection, iterated Stokes, iterated Advection and Newton Stokes, single Advection, no Stokes, IMPES, iterated IMPES, iterated Stokes, Newton Stokes, Stokes only, Advection only

- *Parameter name:* Nonlinear solver tolerance

Value: 1e-5

Default: 1e-5

Description: A relative tolerance up to which the nonlinear solver will iterate. This parameter is only relevant if the ‘Nonlinear solver scheme’ does nonlinear iterations, in other words, if it is set to something other than ‘single Advection, single Stokes’ or ‘single Advection, no Stokes’.

Possible values: A floating point number v such that $0 \leq v \leq 1$

- *Parameter name:* Output directory

Value: output

Default: output

Description: The name of the directory into which all output files should be placed. This may be an absolute or a relative path.

Possible values: A directory name

- *Parameter name:* Pressure normalization

Value: surface

Default: surface

Description: If and how to normalize the pressure after the solution step. This is necessary because depending on boundary conditions, in many cases the pressure is only determined by the model up to a constant. On the other hand, we often would like to have a well-determined pressure, for example for table lookups of material properties in models or for comparing solutions. If the given value is ‘surface’, then normalization at the end of each time steps adds a constant value to the pressure in such a way that the average pressure at the surface of the domain is what is set in the ‘Surface pressure’ parameter; the surface of the domain is determined by asking the geometry model whether a particular face of the geometry has a zero or small ‘depth’. If the value of this parameter is ‘volume’ then the pressure is normalized so that the domain average is zero. If ‘no’ is given, the no pressure normalization is performed.

Possible values: Any one of surface, volume, no

- *Parameter name:* Resume computation

Value: false

Default: false

Description: A flag indicating whether the computation should be resumed from a previously saved state (if true) or start from scratch (if false). If auto is selected, models will be resumed if there is an existing checkpoint file, otherwise started from scratch.

Possible values: Any one of true, false, auto

- *Parameter name:* Start time

Value: 0

Default: 0

Description: The start time of the simulation. Units: Years if the ‘Use years in output instead of seconds’ parameter is set; seconds otherwise.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Surface pressure

Value: 0

Default: 0

Description: The value the pressure is normalized to in each time step when ‘Pressure normalization’ is set to ‘surface’ with default value 0. This setting is ignored in all other cases.

The mathematical equations that describe thermal convection only determine the pressure up to an arbitrary constant. On the other hand, for comparison and for looking up material parameters it is important that the pressure be normalized somehow. We do this by enforcing a particular average pressure value at the surface of the domain, where the geometry model determines where the surface is. This parameter describes what this average surface pressure value is supposed to be. By default, it is set to zero, but one may want to choose a different value for example for simulating only the volume of the mantle below the lithosphere, in which case the surface pressure should be the lithostatic pressure at the bottom of the lithosphere.

For more information, see the section in the manual that discusses the general mathematical model.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Timing output frequency

Value: 100

Default: 100

Description: How frequently in timesteps to output timing information. This is generally adjusted only for debugging and timing purposes. If the value is set to zero it will also output timing information at the initiation timesteps.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* Use conduction timestep

Value: false

Default: false

Description: Mantle convection simulations are often focused on convection dominated systems. However, these codes can also be used to investigate systems where heat conduction plays a dominant role. This parameter indicates whether the simulator should also use heat conduction in determining the length of each time step.

Possible values: A boolean value (true or false)

- *Parameter name:* Use operator splitting

Value: false

Default: false

Description: If set to true, the advection and reactions of compositional fields and temperature are solved separately, and can use different time steps. Note that this will only work if the material/heating model fills the reaction_rates/heating_reaction_rates structures. Operator splitting can be used with any existing solver schemes that solve the temperature/composition equations.

Possible values: A boolean value (true or false)

- *Parameter name:* Use years in output instead of seconds

Value: true

Default: true

Description: When computing results for mantle convection simulations, it is often difficult to judge the order of magnitude of results when they are stated in MKS units involving seconds. Rather, some kinds of results such as velocities are often stated in terms of meters per year (or, sometimes, centimeters per year). On the other hand, for non-dimensional computations, one wants results in their natural unit system as used inside the code. If this flag is set to ‘true’ conversion to years happens; if it is ‘false’, no such conversion happens. Note that when ‘true’, some input such as prescribed velocities should also use years instead of seconds.

Possible values: A boolean value (true or false)

A.2 Parameters in section Adiabatic conditions model

- *Parameter name:* Model name

Value: compute profile

Default: compute profile

Description: Select one of the following models:

‘ascii data’: A model in which the adiabatic profile is read from a file that describes the reference state. Note the required format of the input data: The first lines may contain any number of comments if they begin with '#', but one of these lines needs to contain the number of points in the reference state as for example '# POINTS: 3'. Following the comment lines there has to be a single line containing the names of all data columns, separated by arbitrarily many spaces. Column names are not allowed to contain spaces. The file can contain unnecessary columns, but for this plugin it needs to at least

provide columns named ‘temperature’, ‘pressure’, and ‘density’. Note that the data lines in the file need to be sorted in order of increasing depth from 0 to the maximal depth in the model domain. Points in the model that are outside of the provided depth range will be assigned the maximum or minimum depth values, respectively. Points do not need to be equidistant, but the computation of properties is optimized in speed if they are.

‘compute profile’: A model in which the adiabatic profile is calculated by solving the hydrostatic equations for pressure and temperature in depth. The gravity is assumed to be in depth direction and the composition is either given by the initial composition at reference points or computed as a reference depth-function. All material parameters are computed by the material model plugin. The surface conditions are either constant or changing over time as prescribed by a user-provided function.

‘function’: A model in which the adiabatic profile is specified by a user defined function. The supplied function has to contain temperature, pressure, and density as a function of depth in this order.

Possible values: Any one of ascii data, compute profile, function

A.3 Parameters in section Adiabatic conditions model/Ascii data model

- *Parameter name:* Data directory

Value: \$ASPECT_SOURCE_DIR/tests/adiabatic-conditions/ascii-data/test/

Default: \$ASPECT_SOURCE_DIR/tests/adiabatic-conditions/ascii-data/test/

Description: The name of a directory that contains the model data. This path may either be absolute (if starting with a '/') or relative to the current directory. The path may also include the special text '\$ASPECT_SOURCE_DIR' which will be interpreted as the path in which the ASPECT source files were located when ASPECT was compiled. This interpretation allows, for example, to reference files located in the ‘data/’ subdirectory of ASPECT.

Possible values: A directory name

- *Parameter name:* Data file name

Value:

Default:

Description: The file name of the material data. Provide file in format: (Velocity file name).%s%d where %s is a string specifying the boundary of the model according to the names of the boundary indicators (of a box or a spherical shell).%d is any sprintf integer qualifier, specifying the format of the current file number.

Possible values: Any string

- *Parameter name:* Scale factor

Value: 1

Default: 1

Description: Scalar factor, which is applied to the boundary velocity. You might want to use this to scale the velocities to a reference model (e.g. with free-slip boundary) or another plate reconstruction. Another way to use this factor is to convert units of the input files. The unit is assumed to bem/s or m/yr depending on the ‘Use years in output instead of seconds’ flag. If you provide velocities in cm/yr set this factor to 0.01.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.4 Parameters in section Adiabatic conditions model/Compute profile

- **Parameter name:** Composition reference profile

Value: initial composition

Default: initial composition

Description: Select how the reference profile for composition is computed. This profile is used to evaluate the material model, when computing the pressure and temperature profile.

Possible values: Any one of initial composition, function

- **Parameter name:** Function constants

Value:

Default:

Description: Sometimes it is convenient to use symbolic constants in the expression that describes the function, rather than having to use its numeric value everywhere the constant appears. These values can be defined using this parameter, in the form ‘var1=value1, var2=value2, ...’

A typical example would be to set this runtime parameter to ‘pi=3.1415926536’ and then use ‘pi’ in the expression of the actual formula. (That said, for convenience this class actually defines both ‘pi’ and ‘Pi’ by default, but you get the idea.)

Possible values: Any string

- **Parameter name:** Function expression

Value: 0

Default: 0

Description: The formula that denotes the function you want to evaluate for particular values of the independent variables. This expression may contain any of the usual operations such as addition or multiplication, as well as all of the common functions such as ‘sin’ or ‘cos’. In addition, it may contain expressions like ‘if(x>0, 1, -1)’ where the expression evaluates to the second argument if the first argument is true, and to the third argument otherwise. For a full overview of possible expressions accepted see the documentation of the muparser library at <http://muparser.beltofion.de/>.

If the function you are describing represents a vector-valued function with multiple components, then separate the expressions for individual components by a semicolon.

Possible values: Any string

- **Parameter name:** Number of points

Value: 1000

Default: 1000

Description: The number of points we use to compute the adiabatic profile. The higher the number of points, the more accurate the downward integration from the adiabatic surface temperature will be.

Possible values: An integer n such that $5 \leq n \leq 2147483647$

- **Parameter name:** Use surface condition function

Value: false

Default: false

Description: Whether to use the ‘Surface condition function’ to determine surface conditions, or the ‘Adiabatic surface temperature’ and ‘Surface pressure’ parameters. If this is set to true the reference profile is updated every timestep. The function expression of the function should be independent of

space, but can depend on time 't'. The function must return two components, the first one being reference surface pressure, the second one being reference surface temperature.

Possible values: A boolean value (true or false)

- *Parameter name:* Variable names

Value: x,t

Default: x,t

Description: The name of the variables as they will be used in the function, separated by commas. By default, the names of variables at which the function will be evaluated is 'x' (in 1d), 'x,y' (in 2d) or 'x,y,z' (in 3d) for spatial coordinates and 't' for time. You can then use these variable names in your function expression and they will be replaced by the values of these variables at which the function is currently evaluated. However, you can also choose a different set of names for the independent variables at which to evaluate your function expression. For example, if you work in spherical coordinates, you may wish to set this input parameter to 'r,phi,theta,t' and then use these variable names in your function expression.

Possible values: Any string

A.5 Parameters in section Adiabatic conditions model/Compute profile/Surface condition function

- *Parameter name:* Function constants

Value:

Default:

Description: Sometimes it is convenient to use symbolic constants in the expression that describes the function, rather than having to use its numeric value everywhere the constant appears. These values can be defined using this parameter, in the form 'var1=value1, var2=value2, ...'.

A typical example would be to set this runtime parameter to 'pi=3.1415926536' and then use 'pi' in the expression of the actual formula. (That said, for convenience this class actually defines both 'pi' and 'Pi' by default, but you get the idea.)

Possible values: Any string

- *Parameter name:* Function expression

Value: 0; 0

Default: 0; 0

Description: The formula that denotes the function you want to evaluate for particular values of the independent variables. This expression may contain any of the usual operations such as addition or multiplication, as well as all of the common functions such as 'sin' or 'cos'. In addition, it may contain expressions like 'if(x>0, 1, -1)' where the expression evaluates to the second argument if the first argument is true, and to the third argument otherwise. For a full overview of possible expressions accepted see the documentation of the muparser library at <http://muparser.beltoforion.de/>.

If the function you are describing represents a vector-valued function with multiple components, then separate the expressions for individual components by a semicolon.

Possible values: Any string

- *Parameter name:* Variable names

Value: x,t

Default: x,t

Description: The name of the variables as they will be used in the function, separated by commas. By default, the names of variables at which the function will be evaluated is ‘x’ (in 1d), ‘x,y’ (in 2d) or ‘x,y,z’ (in 3d) for spatial coordinates and ‘t’ for time. You can then use these variable names in your function expression and they will be replaced by the values of these variables at which the function is currently evaluated. However, you can also choose a different set of names for the independent variables at which to evaluate your function expression. For example, if you work in spherical coordinates, you may wish to set this input parameter to ‘r,phi,theta,t’ and then use these variable names in your function expression.

Possible values: Any string

A.6 Parameters in section Adiabatic conditions model/Function

- *Parameter name:* Function constants

Value:

Default:

Description: Sometimes it is convenient to use symbolic constants in the expression that describes the function, rather than having to use its numeric value everywhere the constant appears. These values can be defined using this parameter, in the form ‘var1=value1, var2=value2, ...’.

A typical example would be to set this runtime parameter to ‘pi=3.1415926536’ and then use ‘pi’ in the expression of the actual formula. (That said, for convenience this class actually defines both ‘pi’ and ‘Pi’ by default, but you get the idea.)

Possible values: Any string

- *Parameter name:* Function expression

Value: 0.0; 0.0; 1.0

Default: 0.0; 0.0; 1.0

Description: Expression for the adiabatic temperature, pressure, and density separated by semicolons as a function of ‘depth’.

Possible values: Any string

- *Parameter name:* Variable names

Value: depth

Default: depth

Possible values: Any string

A.7 Parameters in section Boundary composition model

- *Parameter name:* Fixed composition boundary indicators

Value:

Default:

Description: A comma separated list of names denoting those boundaries on which the composition is fixed and described by the boundary composition object selected in its own section of this input file. All boundary indicators used by the geometry but not explicitly listed here will end up with no-flux (insulating) boundary conditions.

The names of the boundaries listed here can either be numbers (in which case they correspond to the numerical boundary indicators assigned by the geometry object), or they can correspond to any of the

symbolic names the geometry object may have provided for each part of the boundary. You may want to compare this with the documentation of the geometry model you use in your model.

This parameter only describes which boundaries have a fixed composition, but not what composition should hold on these boundaries. The latter piece of information needs to be implemented in a plugin in the BoundaryComposition group, unless an existing implementation in this group already provides what you want.

Possible values: A list of 0 to 4294967295 elements where each element is [Any string]

- *Parameter name:* List of model names

Value:

Default:

Description: A comma-separated list of boundary composition models that will be used to initialize the composition. These plugins are loaded in the order given, and modify the existing composition field via the operators listed in 'List of model operators'.

The following boundary composition models are available:

‘ascii data’: Implementation of a model in which the boundary composition is derived from files containing data in ascii format. Note the required format of the input data: The first lines may contain any number of comments if they begin with '#', but one of these lines needs to contain the number of grid points in each dimension as for example '# POINTS: 3 3'. The order of the data columns has to be ‘x’, ‘composition1’, ‘composition2’, etc. in a 2d model and ‘x’, ‘y’, ‘composition1’, ‘composition2’, etc., in a 3d model, according to the number of compositional fields, which means that there has to be a single column for every composition in the model. Note that the data in the input files need to be sorted in a specific order: the first coordinate needs to ascend first, followed by the second in order to assign the correct data to the prescribed coordinates. If you use a spherical model, then the data will still be handled as Cartesian, however the assumed grid changes. ‘x’ will be replaced by the radial distance of the point to the bottom of the model, ‘y’ by the azimuth angle and ‘z’ by the polar angle measured positive from the north pole. The grid will be assumed to be a latitude-longitude grid. Note that the order of spherical coordinates is ‘r’, ‘phi’, ‘theta’ and not ‘r’, ‘theta’, ‘phi’, since this allows for dimension independent expressions.

‘box’: A model in which the composition is chosen constant on the sides of a box which are selected by the parameters Left/Right/Top/Bottom/Front/Back composition

‘box with lithosphere boundary indicators’: A model in which the composition is chosen constant on all the sides of a box. Additional boundary indicators are added to the lithospheric parts of the vertical boundaries. This model is to be used with the ‘Two Merged Boxes’ Geometry Model.

‘function’: Implementation of a model in which the boundary composition is given in terms of an explicit formula that is elaborated in the parameters in section “Boundary composition model|Function”.

Since the symbol t indicating time may appear in the formulas for the prescribed composition, it is interpreted as having units seconds unless the global input parameter “Use years in output instead of seconds” is set, in which case we interpret the formula expressions as having units year.

The format of these functions follows the syntax understood by the muparser library, see Section 4.7.3.

‘initial composition’: A model in which the composition at the boundary is chosen to be the same as given in the initial conditions.

Because this class simply takes what the initial composition had described, this class can not know certain pieces of information such as the minimal and maximal composition on the boundary. For operations that require this, for example in post-processing, this boundary composition model must therefore be told what the minimal and maximal values on the boundary are. This is done using parameters set in section “Boundary composition model/Initial composition”.

‘spherical constant’: A model in which the composition is chosen constant on the inner and outer boundaries of a surface, spherical shell, chunk or ellipsoidal chunk. Parameters are read from subsection ‘Spherical constant’.

Possible values: A comma-separated list of any of ascii data, box, box with lithosphere boundary indicators, function, initial composition, spherical constant

- *Parameter name:* List of model operators

Value: add

Default: add

Description: A comma-separated list of operators that will be used to append the listed composition models onto the previous models. If only one operator is given, the same operator is applied to all models.

Possible values: A comma-separated list of any of add, subtract, minimum, maximum

- *Parameter name:* Model name

Value: unspecified

Default: unspecified

Description: Select one of the following models:

‘ascii data’: Implementation of a model in which the boundary composition is derived from files containing data in ascii format. Note the required format of the input data: The first lines may contain any number of comments if they begin with '#', but one of these lines needs to contain the number of grid points in each dimension as for example '# POINTS: 3 3'. The order of the data columns has to be ‘x’, ‘composition1’, ‘composition2’, etc. in a 2d model and ‘x’, ‘y’, ‘composition1’, ‘composition2’, etc., in a 3d model, according to the number of compositional fields, which means that there has to be a single column for every composition in the model. Note that the data in the input files need to be sorted in a specific order: the first coordinate needs to ascend first, followed by the second in order to assign the correct data to the prescribed coordinates. If you use a spherical model, then the data will still be handled as Cartesian, however the assumed grid changes. ‘x’ will be replaced by the radial distance of the point to the bottom of the model, ‘y’ by the azimuth angle and ‘z’ by the polar angle measured positive from the north pole. The grid will be assumed to be a latitude-longitude grid. Note that the order of spherical coordinates is ‘r’, ‘phi’, ‘theta’ and not ‘r’, ‘theta’, ‘phi’, since this allows for dimension independent expressions.

‘box’: A model in which the composition is chosen constant on the sides of a box which are selected by the parameters Left/Right/Top/Bottom/Front/Back composition

‘box with lithosphere boundary indicators’: A model in which the composition is chosen constant on all the sides of a box. Additional boundary indicators are added to the lithospheric parts of the vertical boundaries. This model is to be used with the ‘Two Merged Boxes’ Geometry Model.

‘function’: Implementation of a model in which the boundary composition is given in terms of an explicit formula that is elaborated in the parameters in section “Boundary composition model|Function”.

Since the symbol t indicating time may appear in the formulas for the prescribed composition, it is interpreted as having units seconds unless the global input parameter “Use years in output instead of seconds” is set, in which case we interpret the formula expressions as having units year.

The format of these functions follows the syntax understood by the muparser library, see Section 4.7.3.

‘initial composition’: A model in which the composition at the boundary is chosen to be the same as given in the initial conditions.

Because this class simply takes what the initial composition had described, this class can not know certain pieces of information such as the minimal and maximal composition on the boundary. For

operations that require this, for example in post-processing, this boundary composition model must therefore be told what the minimal and maximal values on the boundary are. This is done using parameters set in section “Boundary composition model/Initial composition”.

‘spherical constant’: A model in which the composition is chosen constant on the inner and outer boundaries of a surface, spherical shell, chunk or ellipsoidal chunk. Parameters are read from subsection ‘Spherical constant’.

Warning: This parameter provides an old and deprecated way of specifying boundary composition models and shouldn’t be used. Please use ‘List of model names’ instead.

Possible values: Any one of ascii data, box, box with lithosphere boundary indicators, function, initial composition, spherical constant, unspecified

A.8 Parameters in section Boundary composition model/Ascii data model

- *Parameter name:* Data directory

Value: \$ASPECT_SOURCE_DIR/data/boundary-composition/ascii-data/test/

Default: \$ASPECT_SOURCE_DIR/data/boundary-composition/ascii-data/test/

Description: The name of a directory that contains the model data. This path may either be absolute (if starting with a '/') or relative to the current directory. The path may also include the special text '\$ASPECT_SOURCE_DIR' which will be interpreted as the path in which the ASPECT source files were located when ASPECT was compiled. This interpretation allows, for example, to reference files located in the ‘data/’ subdirectory of ASPECT.

Possible values: A directory name

- *Parameter name:* Data file name

Value: box_2d_%s.%d.txt

Default: box_2d_%s.%d.txt

Description: The file name of the material data. Provide file in format: (Velocity file name).%s%d where %s is a string specifying the boundary of the model according to the names of the boundary indicators (of a box or a spherical shell).%d is any sprintf integer qualifier, specifying the format of the current file number.

Possible values: Any string

- *Parameter name:* Data file time step

Value: 1e6

Default: 1e6

Description: Time step between following velocity files. Depending on the setting of the global ‘Use years in output instead of seconds’ flag in the input file, this number is either interpreted as seconds or as years. The default is one million, i.e., either one million seconds or one million years.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Decreasing file order

Value: false

Default: false

Description: In some cases the boundary files are not numbered in increasing but in decreasing order (e.g. ‘Ma BP’). If this flag is set to ‘True’ the plugin will first load the file with the number ‘First velocity file number’ and decrease the file number during the model run.

Possible values: A boolean value (true or false)

- *Parameter name:* First data file model time

Value: 0

Default: 0

Description: Time from which on the velocity file with number 'First velocity file number' is used as boundary condition. Previous to this time, a no-slip boundary condition is assumed. Depending on the setting of the global 'Use years in output instead of seconds' flag in the input file, this number is either interpreted as seconds or as years.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* First data file number

Value: 0

Default: 0

Description: Number of the first velocity file to be loaded when the model time is larger than 'First velocity file model time'.

Possible values: An integer n such that $-2147483648 \leq n \leq 2147483647$

- *Parameter name:* Scale factor

Value: 1

Default: 1

Description: Scalar factor, which is applied to the boundary velocity. You might want to use this to scale the velocities to a reference model (e.g. with free-slip boundary) or another plate reconstruction. Another way to use this factor is to convert units of the input files. The unit is assumed to be m/s or m/yr depending on the 'Use years in output instead of seconds' flag. If you provide velocities in cm/yr set this factor to 0.01.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.9 Parameters in section Boundary composition model/Box

- *Parameter name:* Bottom composition

Value:

Default:

Description: A comma separated list of composition boundary values at the bottom boundary (at minimal y-value in 2d, or minimal z-value in 3d). This list must have as many entries as there are compositional fields. Units: none.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Left composition

Value:

Default:

Description: A comma separated list of composition boundary values at the left boundary (at minimal x-value). This list must have as many entries as there are compositional fields. Units: none.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Right composition

Value:

Default:

Description: A comma separated list of composition boundary values at the right boundary (at maximal x-value). This list must have as many entries as there are compositional fields. Units: none.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Top composition

Value:

Default:

Description: A comma separated list of composition boundary values at the top boundary (at maximal y-value in 2d, or maximal z-value in 3d). This list must have as many entries as there are compositional fields. Units: none.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

A.10 Parameters in section Boundary composition model/Box with lithosphere boundary indicators

- *Parameter name:* Bottom composition

Value:

Default:

Description: A comma separated list of composition boundary values at the bottom boundary (at minimal y-value in 2d, or minimal z-value in 3d). This list must have as many entries as there are compositional fields. Units: none.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Left composition

Value:

Default:

Description: A comma separated list of composition boundary values at the left boundary (at minimal x-value). This list must have as many entries as there are compositional fields. Units: none.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Left composition lithosphere

Value:

Default:

Description: A comma separated list of composition boundary values at the left boundary (at minimal x-value). This list must have as many entries as there are compositional fields. Units: none.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Right composition

Value:

Default:

Description: A comma separated list of composition boundary values at the right boundary (at maximal x-value). This list must have as many entries as there are compositional fields. Units: none.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Right composition lithosphere

Value:

Default:

Description: A comma separated list of composition boundary values at the right boundary (at maximal x-value). This list must have as many entries as there are compositional fields. Units: none.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Top composition

Value:

Default:

Description: A comma separated list of composition boundary values at the top boundary (at maximal y-value in 2d, or maximal z-value in 3d). This list must have as many entries as there are compositional fields. Units: none.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

A.11 Parameters in section Boundary composition model/Function

- *Parameter name:* Coordinate system

Value: cartesian

Default: cartesian

Description: A selection that determines the assumed coordinate system for the function variables. Allowed values are 'cartesian', 'spherical', and 'depth'. 'spherical' coordinates are interpreted as r,phi or r,phi,theta in 2D/3D respectively with theta being the polar angle. 'depth' will create a function, in which only the first parameter is non-zero, which is interpreted to be the depth of the point.

Possible values: Any one of cartesian, spherical, depth

- *Parameter name:* Function constants

Value:

Default:

Description: Sometimes it is convenient to use symbolic constants in the expression that describes the function, rather than having to use its numeric value everywhere the constant appears. These values can be defined using this parameter, in the form 'var1=value1, var2=value2, ...'

A typical example would be to set this runtime parameter to 'pi=3.1415926536' and then use 'pi' in the expression of the actual formula. (That said, for convenience this class actually defines both 'pi' and 'Pi' by default, but you get the idea.)

Possible values: Any string

- *Parameter name:* Function expression

Value: 0

Default: 0

Description: The formula that denotes the function you want to evaluate for particular values of the independent variables. This expression may contain any of the usual operations such as addition or multiplication, as well as all of the common functions such as ‘sin’ or ‘cos’. In addition, it may contain expressions like ‘if($x > 0$, 1, -1)’ where the expression evaluates to the second argument if the first argument is true, and to the third argument otherwise. For a full overview of possible expressions accepted see the documentation of the muparser library at <http://muparser.beltoforion.de/>.

If the function you are describing represents a vector-valued function with multiple components, then separate the expressions for individual components by a semicolon.

Possible values: Any string

- *Parameter name:* Variable names

Value: x,y,t

Default: x,y,t

Description: The name of the variables as they will be used in the function, separated by commas. By default, the names of variables at which the function will be evaluated is ‘x’ (in 1d), ‘x,y’ (in 2d) or ‘x,y,z’ (in 3d) for spatial coordinates and ‘t’ for time. You can then use these variable names in your function expression and they will be replaced by the values of these variables at which the function is currently evaluated. However, you can also choose a different set of names for the independent variables at which to evaluate your function expression. For example, if you work in spherical coordinates, you may wish to set this input parameter to ‘r,phi,theta,t’ and then use these variable names in your function expression.

Possible values: Any string

A.12 Parameters in section Boundary composition model/Initial composition

- *Parameter name:* Maximal composition

Value: 1

Default: 1

Description: Maximal composition. Units: none.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Minimal composition

Value: 0

Default: 0

Description: Minimal composition. Units: none.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

A.13 Parameters in section Boundary composition model/Spherical constant

- *Parameter name:* Inner composition

Value: 1

Default: 1

Description: Composition at the inner boundary (core mantle boundary). Units: none.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Outer composition

Value: 0

Default: 0

Description: Composition at the outer boundary (lithosphere water/air). For a spherical geometry model, this is the only boundary. Units: none.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

A.14 Parameters in section Boundary fluid pressure model

- *Parameter name:* Plugin name

Value: density

Default: density

Description: Select one of the following plugins:

'density': A plugin that prescribes the fluid pressure gradient at the boundary based on fluid/solid density from the material model.

Possible values: Any one of density

A.15 Parameters in section Boundary fluid pressure model/Density

- *Parameter name:* Density formulation

Value: solid density

Default: solid density

Description: The density formulation used to compute the fluid pressure gradient at the model boundary.

'solid density' prescribes the gradient of the fluid pressure as solid density times gravity (which is the lithostatic pressure) and leads to approximately the same pressure in the melt as in the solid, so that fluid is only flowing in or out due to differences in dynamic pressure.

'fluid density' prescribes the gradient of the fluid pressure as fluid density times gravity and causes melt to flow in with the same velocity as inflowing solid material, or no melt flowing in or out if the solid velocity normal to the boundary is zero.

'average density' prescribes the gradient of the fluid pressure as the averaged fluid and solid density times gravity (which is a better approximation for the lithostatic pressure than just the solid density) and leads to approximately the same pressure in the melt as in the solid, so that fluid is only flowing in or out due to differences in dynamic pressure.

Possible values: Any one of solid density, fluid density, average density

A.16 Parameters in section Boundary temperature model

- *Parameter name:* Fixed temperature boundary indicators

Value:

Default:

Description: A comma separated list of names denoting those boundaries on which the temperature is fixed and described by the boundary temperature object selected in its own section of this input file. All boundary indicators used by the geometry but not explicitly listed here will end up with no-flux (insulating) boundary conditions.

The names of the boundaries listed here can either be numbers (in which case they correspond to the numerical boundary indicators assigned by the geometry object), or they can correspond to any of the symbolic names the geometry object may have provided for each part of the boundary. You may want to compare this with the documentation of the geometry model you use in your model.

This parameter only describes which boundaries have a fixed temperature, but not what temperature should hold on these boundaries. The latter piece of information needs to be implemented in a plugin in the BoundaryTemperature group, unless an existing implementation in this group already provides what you want.

Possible values: A list of 0 to 4294967295 elements where each element is [Any string]

- *Parameter name:* List of model names

Value: box

Default:

Description: A comma-separated list of boundary temperature models that will be used to initialize the temperature. These plugins are loaded in the order given, and modify the existing temperature field via the operators listed in 'List of model operators'.

The following boundary temperature models are available:

‘Dynamic core’: This is a boundary temperature model working only with spherical shell geometry and core statistics postprocessor. The temperature at the top is constant, and the core mantle boundary temperature is dynamically evolving through time by calculating the heat flux into the core and solving the core energy balance. The formulation is mainly following Nimmo et al. [2004], and the plugin is used in Zhang et al. [2016]. The energy of core cooling and freeing of the inner core is included in the plugin. However, current plugin can not deal with the energy balance if the core is in the ‘snowing core’ regime (i.e. core solidifies from top instead of bottom)

‘ascii data’: Implementation of a model in which the boundary data is derived from files containing data in ascii format. Note the required format of the input data: The first lines may contain any number of comments if they begin with '#', but one of these lines needs to contain the number of grid points in each dimension as for example '# POINTS: 3 3'. The order of the data columns has to be ‘x’, ‘Temperature [K]’ in a 2d model and ‘x’, ‘y’, ‘Temperature [K]’ in a 3d model, which means that there has to be a single column containing the temperature. Note that the data in the input files need to be sorted in a specific order: the first coordinate needs to ascend first, followed by the second in order to assign the correct data to the prescribed coordinates. If you use a spherical model, then the data will still be handled as Cartesian, however the assumed grid changes. ‘x’ will be replaced by the radial distance of the point to the bottom of the model, ‘y’ by the azimuth angle and ‘z’ by the polar angle measured positive from the north pole. The grid will be assumed to be a latitude-longitude grid. Note that the order of spherical coordinates is ‘r’, ‘phi’, ‘theta’ and not ‘r’, ‘theta’, ‘phi’, since this allows for dimension independent expressions.

‘box’: A model in which the temperature is chosen constant on the sides of a box which are selected by the parameters Left/Right/Top/Bottom/Front/Back temperature

‘box with lithosphere boundary indicators’: A model in which the temperature is chosen constant on all the sides of a box. Additional boundary indicators are added to the lithospheric parts of the vertical boundaries. This model is to be used with the ‘Two Merged Boxes’ Geometry Model.

‘constant’: A model in which the temperature is chosen constant on a given boundary indicator. Parameters are read from the subsection ‘Constant’.

‘function’: Implementation of a model in which the boundary temperature is given in terms of an explicit formula that is elaborated in the parameters in section “Boundary temperature modelFunction”.

Since the symbol t indicating time may appear in the formulas for the prescribed temperatures, it is interpreted as having units seconds unless the global input parameter “Use years in output instead of seconds” is set, in which case we interpret the formula expressions as having units year.

Because this class simply takes what the function calculates, this class can not know certain pieces of information such as the minimal and maximal temperature on the boundary. For operations that require this, for example in post-processing, this boundary temperature model must therefore be told what the minimal and maximal values on the boundary are. This is done using parameters set in section “Boundary temperature model/Initial temperature”.

The format of these functions follows the syntax understood by the muparser library, see Section 4.7.3.

‘initial temperature’: A model in which the temperature at the boundary is chosen to be the same as given in the initial conditions.

Because this class simply takes what the initial temperature had described, this class can not know certain pieces of information such as the minimal and maximal temperature on the boundary. For operations that require this, for example in post-processing, this boundary temperature model must therefore be told what the minimal and maximal values on the boundary are. This is done using parameters set in section “Boundary temperature model/Initial temperature”.

‘spherical constant’: A model in which the temperature is chosen constant on the inner and outer boundaries of a spherical shell, ellipsoidal chunk or chunk. Parameters are read from subsection ‘Spherical constant’.

Possible values: A comma-separated list of any of Dynamic core, ascii data, box, box with lithosphere boundary indicators, constant, function, initial temperature, spherical constant

- *Parameter name:* List of model operators

Value: add

Default: add

Description: A comma-separated list of operators that will be used to append the listed temperature models onto the previous models. If only one operator is given, the same operator is applied to all models.

Possible values: A comma-separated list of any of add, subtract, minimum, maximum

- *Parameter name:* Model name

Value: unspecified

Default: unspecified

Description: Select one of the following models:

‘Dynamic core’: This is a boundary temperature model working only with spherical shell geometry and core statistics postprocessor. The temperature at the top is constant, and the core mantle boundary temperature is dynamically evolving through time by calculating the heat flux into the core and solving the core energy balance. The formulation is mainly following Nimmo et al. [2004], and the plugin is used in Zhang et al. [2016]. The energy of core cooling and freeing of the inner core is included in the plugin. However, current plugin can not deal with the energy balance if the core is in the ‘snowing core’ regime (i.e. core solidifies from top instead of bottom)

‘ascii data’: Implementation of a model in which the boundary data is derived from files containing data in ascii format. Note the required format of the input data: The first lines may contain any number of comments if they begin with '#', but one of these lines needs to contain the number of grid points in each dimension as for example '# POINTS: 3 3'. The order of the data columns has to be ‘x’, ‘Temperature [K]’ in a 2d model and ‘x’, ‘y’, ‘Temperature [K]’ in a 3d model, which means that there has to be a single column containing the temperature. Note that the data in the input files need to be sorted in a specific order: the first coordinate needs to ascend first, followed by the second in order to assign the correct data to the prescribed coordinates. If you use a spherical model, then the data will still be handled as Cartesian, however the assumed grid changes. ‘x’ will be replaced by the radial distance of the point to the bottom of the model, ‘y’ by the azimuth angle and ‘z’ by the polar angle

measured positive from the north pole. The grid will be assumed to be a latitude-longitude grid. Note that the order of spherical coordinates is ‘r’, ‘phi’, ‘theta’ and not ‘r’, ‘theta’, ‘phi’, since this allows for dimension independent expressions.

‘box’: A model in which the temperature is chosen constant on the sides of a box which are selected by the parameters Left/Right/Top/Bottom/Front/Back temperature

‘box with lithosphere boundary indicators’: A model in which the temperature is chosen constant on all the sides of a box. Additional boundary indicators are added to the lithospheric parts of the vertical boundaries. This model is to be used with the ‘Two Merged Boxes’ Geometry Model.

‘constant’: A model in which the temperature is chosen constant on a given boundary indicator. Parameters are read from the subsection ‘Constant’.

‘function’: Implementation of a model in which the boundary temperature is given in terms of an explicit formula that is elaborated in the parameters in section “Boundary temperature model|Function”.

Since the symbol t indicating time may appear in the formulas for the prescribed temperatures, it is interpreted as having units seconds unless the global input parameter “Use years in output instead of seconds” is set, in which case we interpret the formula expressions as having units year.

Because this class simply takes what the function calculates, this class can not know certain pieces of information such as the minimal and maximal temperature on the boundary. For operations that require this, for example in post-processing, this boundary temperature model must therefore be told what the minimal and maximal values on the boundary are. This is done using parameters set in section “Boundary temperature model/Initial temperature”.

The format of these functions follows the syntax understood by the muparser library, see Section 4.7.3.

‘initial temperature’: A model in which the temperature at the boundary is chosen to be the same as given in the initial conditions.

Because this class simply takes what the initial temperature had described, this class can not know certain pieces of information such as the minimal and maximal temperature on the boundary. For operations that require this, for example in post-processing, this boundary temperature model must therefore be told what the minimal and maximal values on the boundary are. This is done using parameters set in section “Boundary temperature model/Initial temperature”.

‘spherical constant’: A model in which the temperature is chosen constant on the inner and outer boundaries of a spherical shell, ellipsoidal chunk or chunk. Parameters are read from subsection ‘Spherical constant’.

Warning: This parameter provides an old and deprecated way of specifying boundary temperature models and shouldn’t be used. Please use ‘List of model names’ instead.

Possible values: Any one of Dynamic core, ascii data, box, box with lithosphere boundary indicators, constant, function, initial temperature, spherical constant, unspecified

A.17 Parameters in section Boundary temperature model/Ascii data model

- *Parameter name:* Data directory

Value: \$ASPECT_SOURCE_DIR/data/boundary-temperature/ascii-data/test/

Default: \$ASPECT_SOURCE_DIR/data/boundary-temperature/ascii-data/test/

Description: The name of a directory that contains the model data. This path may either be absolute (if starting with a '/') or relative to the current directory. The path may also include the special text '\$ASPECT_SOURCE_DIR' which will be interpreted as the path in which the ASPECT source files were located when ASPECT was compiled. This interpretation allows, for example, to reference files located in the ‘data/’ subdirectory of ASPECT.

Possible values: A directory name

- **Parameter name:** `Data file name`

Value: `box_2d_%s.%d.txt`

Default: `box_2d_%s.%d.txt`

Description: The file name of the material data. Provide file in format: (Velocity file name).%s%d where %s is a string specifying the boundary of the model according to the names of the boundary indicators (of a box or a spherical shell).%d is any sprintf integer qualifier, specifying the format of the current file number.

Possible values: Any string

- **Parameter name:** `Data file time step`

Value: `1e6`

Default: `1e6`

Description: Time step between following velocity files. Depending on the setting of the global 'Use years in output instead of seconds' flag in the input file, this number is either interpreted as seconds or as years. The default is one million, i.e., either one million seconds or one million years.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** `Decreasing file order`

Value: `false`

Default: `false`

Description: In some cases the boundary files are not numbered in increasing but in decreasing order (e.g. 'Ma BP'). If this flag is set to 'True' the plugin will first load the file with the number 'First velocity file number' and decrease the file number during the model run.

Possible values: A boolean value (true or false)

- **Parameter name:** `First data file model time`

Value: `0`

Default: `0`

Description: Time from which on the velocity file with number 'First velocity file number' is used as boundary condition. Previous to this time, a no-slip boundary condition is assumed. Depending on the setting of the global 'Use years in output instead of seconds' flag in the input file, this number is either interpreted as seconds or as years.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** `First data file number`

Value: `0`

Default: `0`

Description: Number of the first velocity file to be loaded when the model time is larger than 'First velocity file model time'.

Possible values: An integer n such that $-2147483648 \leq n \leq 2147483647$

- **Parameter name:** `Scale factor`

Value: `1`

Default: `1`

Description: Scalar factor, which is applied to the boundary velocity. You might want to use this to scale the velocities to a reference model (e.g. with free-slip boundary) or another plate reconstruction.

Another way to use this factor is to convert units of the input files. The unit is assumed to be m/s or m/yr depending on the 'Use years in output instead of seconds' flag. If you provide velocities in cm/yr set this factor to 0.01.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.18 Parameters in section Boundary temperature model/Box

- *Parameter name:* Bottom temperature

Value: 0

Default: 0

Description: Temperature at the bottom boundary (at minimal z-value). Units: K.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Left temperature

Value: 1

Default: 1

Description: Temperature at the left boundary (at minimal x-value). Units: K.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Right temperature

Value: 0

Default: 0

Description: Temperature at the right boundary (at maximal x-value). Units: K.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Top temperature

Value: 0

Default: 0

Description: Temperature at the top boundary (at maximal x-value). Units: K.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

A.19 Parameters in section Boundary temperature model/Box with lithosphere boundary indicators

- *Parameter name:* Bottom temperature

Value: 0

Default: 0

Description: Temperature at the bottom boundary (at minimal z-value). Units: K.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Left temperature

Value: 1

Default: 1

Description: Temperature at the left boundary (at minimal x-value). Units: K.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Left temperature lithosphere
Value: 0
Default: 0
Description: Temperature at the additional left lithosphere boundary (specified by user in Geometry Model). Units: K.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Right temperature
Value: 0
Default: 0
Description: Temperature at the right boundary (at maximal x-value). Units: K.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Right temperature lithosphere
Value: 0
Default: 0
Description: Temperature at the additional right lithosphere boundary (specified by user in Geometry Model). Units: K.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Top temperature
Value: 0
Default: 0
Description: Temperature at the top boundary (at maximal x-value). Units: K.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

A.20 Parameters in section Boundary temperature model/Constant

- *Parameter name:* Boundary indicator to temperature mappings
Value:
Default:
Description: A comma separated list of mappings between boundary indicators and the temperature associated with the boundary indicators. The format for this list is “indicator1 : value1, indicator2 : value2, ...”, where each indicator is a valid boundary indicator (either a number or the symbolic name of a boundary as provided by the geometry model) and each value is the temperature of that boundary.
Possible values: A key-value map of 0 to 4294967295 elements where each key is [Any string] and each value is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

A.21 Parameters in section Boundary temperature model/Dynamic core

- *Parameter name:* Alpha
Value: 1.35e-5
Default: 1.35e-5
Description: Core thermal expansivity. Unit: 1/K
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Beta_composition

Value: 1.1

Default: 1.1

Description: Compositional expansion coefficient Beta_c . Referring to Nimmo et al. (2004) for more details.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* CMB_pressure

Value: 0.14e12

Default: 0.14e12

Description: Pressure at CMB. Units: Pa.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Core_conductivity

Value: 60

Default: 60

Description: Core heat conductivity k_c . Unit: W/m/K

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Core_density

Value: 12.5e3

Default: 12.5e3

Description: Density of the core. Units: kg/m^3

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Core_heat_capacity

Value: 840

Default: 840

Description: Heat capacity of the core. Units: $\text{J}/\text{kg}/\text{K}$

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Delta

Value: 0.5

Default: 0.5

Description: Partition coefficient of the light element.

Possible values: A floating point number v such that $0 \leq v \leq 1$

- *Parameter name:* Gravity_acceleration

Value: 9.8

Default: 9.8

Description: Gravitation acceleration at CMB. Units: m/s^2 .

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Initial light composition
Value: 0.01
Default: 0.01
Description: Initial light composition (eg. S,O) concentration in weight fraction.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Inner temperature
Value: 6000
Default: 6000
Description: Temperature at the inner boundary (core mantle boundary) at the beginning. Units: K.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* K0
Value: 4.111e11
Default: 4.111e11
Description: Core compressibility at zero pressure. Referring to Nimmo et al. (2004) for more details.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Lh
Value: 750e3
Default: 750e3
Description: The latent heat of core freeze. Unit: J/kg
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Max iteration
Value: 30000
Default: 30000
Description: The max iterations for nonlinear core energy solver.
Possible values: An integer n such that $0 \leq n \leq 2147483647$
- *Parameter name:* Outer temperature
Value: 0
Default: 0
Description: Temperature at the outer boundary (lithosphere water/air). Units: K.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Rh
Value: -27.7e6
Default: -27.7e6
Description: The heat of reaction. Unit: J/kg
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Rho0
Value: 7.019e3
Default: 7.019e3
Description: Core density at zero pressure. Units: kg/m^3 . Referring to Nimmo et al. (2004) for more details.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* dR over dt
Value: 0
Default: 0
Description: Initial inner core radius changing rate. Units: km/year
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* dT over dt
Value: 0
Default: 0
Description: Initial CMB temperature changing rate. Units: K/year
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* dX over dt
Value: 0
Default: 0
Description: Initial light composition changing rate. Units: 1/year
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

A.22 Parameters in section Boundary temperature model/Dynamic core/Geotherm parameters

- *Parameter name:* Composition dependency
Value: true
Default: true
Description: If melting curve dependent on composition.
Possible values: A boolean value (true or false)
- *Parameter name:* Theta
Value: 0.11
Default: 0.11
Description: Melting curve (Nimmo et al. [2004] eq. (40)) parameter Theta.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Tm0
Value: 1695
Default: 1695
Description: Melting curve (Nimmo et al. [2004] eq. (40)) parameter Tm0. Unit: K
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Tm1

Value: 10.9

Default: 10.9

Description: Melting curve (Nimmo et al. [2004] eq. (40)) parameter Tm1. Unit: $1/Tpa$

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Tm2

Value: -8.0

Default: -8.0

Description: Melting curve (Nimmo et al. [2004] eq. (40)) parameter Tm2. Unit: $1/TPa^2$

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Use_BW11

Value: false

Default: false

Description: If using the Fe-FeS system solidus from Buono & Walker (2011) instead.

Possible values: A boolean value (true or false)

A.23 Parameters in section Boundary temperature model/Dynamic core/Other energy source

- *Parameter name:* File_name

Value:

Default:

Description: Data file name for other energy source into the core. The 'other energy source' is used for external core energy source. For example if someone want to test the early lunar core powered by precession (Dwyer, C. A., et al. (2011). A long-lived lunar dynamo driven by continuous mechanical stirring. Nature 479(7372): 212-214.) Format [Time(Gyr) Energy rate(W)]

Possible values: Any string

A.24 Parameters in section Boundary temperature model/Dynamic core/Radioactive heat source

- *Parameter name:* Half_life_times

Value:

Default:

Description: Half decay times of different elements (Ga)

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Heating_rates

Value:

Default:

Description: Heating rates of different elements (W/kg)

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Initial concentrations

Value:

Default:

Description: Initial concentrations of different elements (ppm)

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Number of radioactive heating elements

Value: 0

Default: 0

Description: Number of different radioactive heating elements in core

Possible values: An integer n such that $0 \leq n \leq 2147483647$

A.25 Parameters in section Boundary temperature model/Function

- *Parameter name:* Coordinate system

Value: cartesian

Default: cartesian

Description: A selection that determines the assumed coordinate system for the function variables. Allowed values are ‘cartesian’, ‘spherical’, and ‘depth’. ‘spherical’ coordinates are interpreted as r,phi or r,phi,theta in 2D/3D respectively with theta being the polar angle. ‘depth’ will create a function, in which only the first parameter is non-zero, which is interpreted to be the depth of the point.

Possible values: Any one of cartesian, spherical, depth

- *Parameter name:* Function constants

Value:

Default:

Description: Sometimes it is convenient to use symbolic constants in the expression that describes the function, rather than having to use its numeric value everywhere the constant appears. These values can be defined using this parameter, in the form ‘var1=value1, var2=value2, ...’

A typical example would be to set this runtime parameter to ‘pi=3.1415926536’ and then use ‘pi’ in the expression of the actual formula. (That said, for convenience this class actually defines both ‘pi’ and ‘Pi’ by default, but you get the idea.)

Possible values: Any string

- *Parameter name:* Function expression

Value: 0

Default: 0

Description: The formula that denotes the function you want to evaluate for particular values of the independent variables. This expression may contain any of the usual operations such as addition or multiplication, as well as all of the common functions such as ‘sin’ or ‘cos’. In addition, it may contain expressions like ‘if(x>0, 1, -1)’ where the expression evaluates to the second argument if the first argument is true, and to the third argument otherwise. For a full overview of possible expressions accepted see the documentation of the muparser library at <http://muparser.beltoforion.de/>.

If the function you are describing represents a vector-valued function with multiple components, then separate the expressions for individual components by a semicolon.

Possible values: Any string

- *Parameter name:* Maximal temperature
Value: 3773
Default: 3773
Description: Maximal temperature. Units: K.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Minimal temperature
Value: 273
Default: 273
Description: Minimal temperature. Units: K.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Variable names
Value: x,y,t
Default: x,y,t
Description: The name of the variables as they will be used in the function, separated by commas. By default, the names of variables at which the function will be evaluated is ‘x’ (in 1d), ‘x,y’ (in 2d) or ‘x,y,z’ (in 3d) for spatial coordinates and ‘t’ for time. You can then use these variable names in your function expression and they will be replaced by the values of these variables at which the function is currently evaluated. However, you can also choose a different set of names for the independent variables at which to evaluate your function expression. For example, if you work in spherical coordinates, you may wish to set this input parameter to ‘r,phi,theta,t’ and then use these variable names in your function expression.
Possible values: Any string

A.26 Parameters in section Boundary temperature model/Initial temperature

- *Parameter name:* Maximal temperature
Value: 3773
Default: 3773
Description: Maximal temperature. Units: K.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Minimal temperature
Value: 0
Default: 0
Description: Minimal temperature. Units: K.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

A.27 Parameters in section Boundary temperature model/Spherical constant

- *Parameter name:* Inner temperature
Value: 6000
Default: 6000
Description: Temperature at the inner boundary (core mantle boundary). Units: K.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Outer temperature

Value: 0

Default: 0

Description: Temperature at the outer boundary (lithosphere water/air). Units: K.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

A.28 Parameters in section Boundary traction model

- *Parameter name:* Prescribed traction boundary indicators

Value:

Default:

Description: A comma separated list denoting those boundaries on which a traction force is prescribed, i.e., where known external forces act, resulting in an unknown velocity. This is often used to model “open” boundaries where we only know the pressure. This pressure then produces a force that is normal to the boundary and proportional to the pressure.

The format of valid entries for this parameter is that of a map given as “key1 [selector]: value1, key2 [selector]: value2, key3: value3, ...” where each key must be a valid boundary indicator (which is either an integer or the symbolic name the geometry model in use may have provided for this part of the boundary) and each value must be one of the currently implemented boundary traction models. “selector” is an optional string given as a subset of the letters ‘xyz’ that allows you to apply the boundary conditions only to the components listed. As an example, ‘1 y: function’ applies the type ‘function’ to the y component on boundary 1. Without a selector it will affect all components of the traction.

Possible values: A key-value map of 0 to 4294967295 elements where each key is [Any string] and each value is [Any one of ascii data, function, initial lithostatic pressure, zero traction]

A.29 Parameters in section Boundary traction model/Ascii data model

- *Parameter name:* Data directory

Value: \$ASPECT_SOURCE_DIR/data/boundary-traction/ascii-data/test/

Default: \$ASPECT_SOURCE_DIR/data/boundary-traction/ascii-data/test/

Description: The name of a directory that contains the model data. This path may either be absolute (if starting with a '/') or relative to the current directory. The path may also include the special text '\$ASPECT_SOURCE_DIR' which will be interpreted as the path in which the ASPECT source files were located when ASPECT was compiled. This interpretation allows, for example, to reference files located in the ‘data/’ subdirectory of ASPECT.

Possible values: A directory name

- *Parameter name:* Data file name

Value: box_2d_%s.%d.txt

Default: box_2d_%s.%d.txt

Description: The file name of the material data. Provide file in format: (Velocity file name).%s%d where %s is a string specifying the boundary of the model according to the names of the boundary indicators (of a box or a spherical shell).%d is any sprintf integer qualifier, specifying the format of the current file number.

Possible values: Any string

- *Parameter name:* `Data file time step`

Value: 1e6

Default: 1e6

Description: Time step between following velocity files. Depending on the setting of the global 'Use years in output instead of seconds' flag in the input file, this number is either interpreted as seconds or as years. The default is one million, i.e., either one million seconds or one million years.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `Decreasing file order`

Value: false

Default: false

Description: In some cases the boundary files are not numbered in increasing but in decreasing order (e.g. 'Ma BP'). If this flag is set to 'True' the plugin will first load the file with the number 'First velocity file number' and decrease the file number during the model run.

Possible values: A boolean value (true or false)

- *Parameter name:* `First data file model time`

Value: 0

Default: 0

Description: Time from which on the velocity file with number 'First velocity file number' is used as boundary condition. Previous to this time, a no-slip boundary condition is assumed. Depending on the setting of the global 'Use years in output instead of seconds' flag in the input file, this number is either interpreted as seconds or as years.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `First data file number`

Value: 0

Default: 0

Description: Number of the first velocity file to be loaded when the model time is larger than 'First velocity file model time'.

Possible values: An integer n such that $-2147483648 \leq n \leq 2147483647$

- *Parameter name:* `Scale factor`

Value: 1

Default: 1

Description: Scalar factor, which is applied to the boundary velocity. You might want to use this to scale the velocities to a reference model (e.g. with free-slip boundary) or another plate reconstruction. Another way to use this factor is to convert units of the input files. The unit is assumed to bem/s or m/yr depending on the 'Use years in output instead of seconds' flag. If you provide velocities in cm/yr set this factor to 0.01.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.30 Parameters in section Boundary traction model/Function

- **Parameter name:** Coordinate system

Value: cartesian

Default: cartesian

Description: A selection that determines the assumed coordinate system for the function variables. Allowed values are ‘cartesian’, ‘spherical’, and ‘depth’. ‘spherical’ coordinates are interpreted as r,phi or r,phi,theta in 2D/3D respectively with theta being the polar angle. ‘depth’ will create a function, in which only the first parameter is non-zero, which is interpreted to be the depth of the point.

Possible values: Any one of cartesian, spherical, depth

- **Parameter name:** Function constants

Value:

Default:

Description: Sometimes it is convenient to use symbolic constants in the expression that describes the function, rather than having to use its numeric value everywhere the constant appears. These values can be defined using this parameter, in the form ‘var1=value1, var2=value2, ...’

A typical example would be to set this runtime parameter to ‘pi=3.1415926536’ and then use ‘pi’ in the expression of the actual formula. (That said, for convenience this class actually defines both ‘pi’ and ‘Pi’ by default, but you get the idea.)

Possible values: Any string

- **Parameter name:** Function expression

Value: 0; 0

Default: 0; 0

Description: The formula that denotes the function you want to evaluate for particular values of the independent variables. This expression may contain any of the usual operations such as addition or multiplication, as well as all of the common functions such as ‘sin’ or ‘cos’. In addition, it may contain expressions like ‘if(x>0, 1, -1)’ where the expression evaluates to the second argument if the first argument is true, and to the third argument otherwise. For a full overview of possible expressions accepted see the documentation of the muparser library at <http://muparser.beltoforion.de/>.

If the function you are describing represents a vector-valued function with multiple components, then separate the expressions for individual components by a semicolon.

Possible values: Any string

- **Parameter name:** Variable names

Value: x,y,t

Default: x,y,t

Description: The name of the variables as they will be used in the function, separated by commas. By default, the names of variables at which the function will be evaluated is ‘x’ (in 1d), ‘x,y’ (in 2d) or ‘x,y,z’ (in 3d) for spatial coordinates and ‘t’ for time. You can then use these variable names in your function expression and they will be replaced by the values of these variables at which the function is currently evaluated. However, you can also choose a different set of names for the independent variables at which to evaluate your function expression. For example, if you work in spherical coordinates, you may wish to set this input parameter to ‘r,phi,theta,t’ and then use these variable names in your function expression.

Possible values: Any string

A.31 Parameters in section Boundary traction model/Initial lithostatic pressure

- *Parameter name:* Number of integration points

Value: 1000

Default: 1000

Description: The number of integration points over which we integrate the lithostatic pressure downwards.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* Representative point

Value:

Default:

Description: The point where the pressure profile will be calculated. Cartesian coordinates when geometry is a box, otherwise enter radius, longitude, and in 3D latitude. Units: m or degrees.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

A.32 Parameters in section Boundary velocity model

- *Parameter name:* Prescribed velocity boundary indicators

Value:

Default:

Description: A comma separated list denoting those boundaries on which the velocity is prescribed, i.e., where unknown external forces act to prescribe a particular velocity. This is often used to prescribe a velocity that equals that of overlying plates.

The format of valid entries for this parameter is that of a map given as “key1 [selector]: value1, key2 [selector]: value2, key3: value3, ...” where each key must be a valid boundary indicator (which is either an integer or the symbolic name the geometry model in use may have provided for this part of the boundary) and each value must be one of the currently implemented boundary velocity models. “selector” is an optional string given as a subset of the letters ‘xyz’ that allows you to apply the boundary conditions only to the components listed. As an example, ‘1 y: function’ applies the type ‘function’ to the y component on boundary 1. Without a selector it will affect all components of the velocity.

Note that the no-slip boundary condition is a special case of the current one where the prescribed velocity happens to be zero. It can thus be implemented by indicating that a particular boundary is part of the ones selected using the current parameter and using “zero velocity” as the boundary values. Alternatively, you can simply list the part of the boundary on which the velocity is to be zero with the parameter “Zero velocity boundary indicator” in the current parameter section.

Note that when “Use years in output instead of seconds” is set to true, velocity should be given in m/yr. The following boundary velocity models are available:

‘ascii data’: Implementation of a model in which the boundary velocity is derived from files containing data in ascii format. Note the required format of the input data: The first lines may contain any number of comments if they begin with '#', but one of these lines needs to contain the number of grid points in each dimension as for example '# POINTS: 3 3'. The order of the data columns has to be ‘x’, ‘velocity_x’, ‘velocity_y’ in a 2d model or ‘x’, ‘y’, ‘velocity_x’, ‘velocity_y’, ‘velocity_z’ in a 3d model. Note that the data in the input files need to be sorted in a specific order: the first coordinate needs to ascend first, followed by the second in order to assign the correct data to the prescribed coordinates. If you use a spherical model, then the velocities will still be handled as Cartesian, however the assumed

grid changes. ‘x’ will be replaced by the radial distance of the point to the bottom of the model, ‘y’ by the azimuth angle and ‘z’ by the polar angle measured positive from the north pole. The grid will be assumed to be a latitude-longitude grid. Note that the order of spherical coordinates is ‘r’, ‘phi’, ‘theta’ and not ‘r’, ‘theta’, ‘phi’, since this allows for dimension independent expressions. No matter which geometry model is chosen, the unit of the velocities is assumed to be m/s or m/yr depending on the ‘Use years in output instead of seconds’ flag.

‘function’: Implementation of a model in which the boundary velocity is given in terms of an explicit formula that is elaborated in the parameters in section “Boundary velocity model|Function”. The format of these functions follows the syntax understood by the muparser library, see Section 4.7.3.

The formula you describe in the mentioned section is a semicolon separated list of velocities for each of the d components of the velocity vector. These d formulas are interpreted as having units m/s, unless the global input parameter “Use years in output instead of seconds” is set, in which case we interpret the formula expressions as having units m/year.

Likewise, since the symbol t indicating time may appear in the formulas for the prescribed velocities, it is interpreted as having units seconds unless the global parameter above has been set.

‘gplates’: Implementation of a model in which the boundary velocity is derived from files that are generated by the GPlates program.

‘zero velocity’: Implementation of a model in which the boundary velocity is zero. This is commonly referred to as a “stick boundary condition”, indicating that the material “sticks” to the material on the other side of the boundary.

Possible values: A key-value map of 0 to 4294967295 elements where each key is [Any string] and each value is [Any one of ascii data, function, gplates, zero velocity]

- **Parameter name:** Tangential velocity boundary indicators

Value:

Default:

Description: A comma separated list of names denoting those boundaries on which the velocity is tangential and unrestrained, i.e., free-slip where no external forces act to prescribe a particular tangential velocity (although there is a force that requires the flow to be tangential).

The names of the boundaries listed here can either be numbers (in which case they correspond to the numerical boundary indicators assigned by the geometry object), or they can correspond to any of the symbolic names the geometry object may have provided for each part of the boundary. You may want to compare this with the documentation of the geometry model you use in your model.

Possible values: A list of 0 to 4294967295 elements where each element is [Any string]

- **Parameter name:** Zero velocity boundary indicators

Value:

Default:

Description: A comma separated list of names denoting those boundaries on which the velocity is zero.

The names of the boundaries listed here can either be numbers (in which case they correspond to the numerical boundary indicators assigned by the geometry object), or they can correspond to any of the symbolic names the geometry object may have provided for each part of the boundary. You may want to compare this with the documentation of the geometry model you use in your model.

Possible values: A list of 0 to 4294967295 elements where each element is [Any string]

A.33 Parameters in section Boundary velocity model/Ascii data model

- *Parameter name:* Data directory

Value: \$ASPECT_SOURCE_DIR/data/boundary-velocity/ascii-data/test/

Default: \$ASPECT_SOURCE_DIR/data/boundary-velocity/ascii-data/test/

Description: The name of a directory that contains the model data. This path may either be absolute (if starting with a '/') or relative to the current directory. The path may also include the special text '\$ASPECT_SOURCE_DIR' which will be interpreted as the path in which the ASPECT source files were located when ASPECT was compiled. This interpretation allows, for example, to reference files located in the 'data/' subdirectory of ASPECT.

Possible values: A directory name

- *Parameter name:* Data file name

Value: box_2d_%s.%d.txt

Default: box_2d_%s.%d.txt

Description: The file name of the material data. Provide file in format: (Velocity file name).%s%d where %s is a string specifying the boundary of the model according to the names of the boundary indicators (of a box or a spherical shell).%d is any sprintf integer qualifier, specifying the format of the current file number.

Possible values: Any string

- *Parameter name:* Data file time step

Value: 1e6

Default: 1e6

Description: Time step between following velocity files. Depending on the setting of the global 'Use years in output instead of seconds' flag in the input file, this number is either interpreted as seconds or as years. The default is one million, i.e., either one million seconds or one million years.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Decreasing file order

Value: false

Default: false

Description: In some cases the boundary files are not numbered in increasing but in decreasing order (e.g. 'Ma BP'). If this flag is set to 'True' the plugin will first load the file with the number 'First velocity file number' and decrease the file number during the model run.

Possible values: A boolean value (true or false)

- *Parameter name:* First data file model time

Value: 0

Default: 0

Description: Time from which on the velocity file with number 'First velocity file number' is used as boundary condition. Previous to this time, a no-slip boundary condition is assumed. Depending on the setting of the global 'Use years in output instead of seconds' flag in the input file, this number is either interpreted as seconds or as years.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* First data file number

Value: 0

Default: 0

Description: Number of the first velocity file to be loaded when the model time is larger than 'First velocity file model time'.

Possible values: An integer n such that $-2147483648 \leq n \leq 2147483647$

- *Parameter name:* Scale factor

Value: 1

Default: 1

Description: Scalar factor, which is applied to the boundary velocity. You might want to use this to scale the velocities to a reference model (e.g. with free-slip boundary) or another plate reconstruction. Another way to use this factor is to convert units of the input files. The unit is assumed to be m/s or m/yr depending on the 'Use years in output instead of seconds' flag. If you provide velocities in cm/yr set this factor to 0.01.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.34 Parameters in section Boundary velocity model/Function

- *Parameter name:* Coordinate system

Value: cartesian

Default: cartesian

Description: A selection that determines the assumed coordinate system for the function variables. Allowed values are 'cartesian', 'spherical', and 'depth'. 'spherical' coordinates are interpreted as r,phi or r,phi,theta in 2D/3D respectively with theta being the polar angle. 'depth' will create a function, in which only the first parameter is non-zero, which is interpreted to be the depth of the point.

Possible values: Any one of cartesian, spherical, depth

- *Parameter name:* Function constants

Value:

Default:

Description: Sometimes it is convenient to use symbolic constants in the expression that describes the function, rather than having to use its numeric value everywhere the constant appears. These values can be defined using this parameter, in the form 'var1=value1, var2=value2, ...'.

A typical example would be to set this runtime parameter to 'pi=3.1415926536' and then use 'pi' in the expression of the actual formula. (That said, for convenience this class actually defines both 'pi' and 'Pi' by default, but you get the idea.)

Possible values: Any string

- *Parameter name:* Function expression

Value: 0; 0

Default: 0; 0

Description: The formula that denotes the function you want to evaluate for particular values of the independent variables. This expression may contain any of the usual operations such as addition or multiplication, as well as all of the common functions such as 'sin' or 'cos'. In addition, it may contain expressions like 'if(x>0, 1, -1)' where the expression evaluates to the second argument if the first

argument is true, and to the third argument otherwise. For a full overview of possible expressions accepted see the documentation of the muparser library at <http://muparser.beltoforion.de/>.

If the function you are describing represents a vector-valued function with multiple components, then separate the expressions for individual components by a semicolon.

Possible values: Any string

- *Parameter name:* **Variable names**

Value: x,y,t

Default: x,y,t

Description: The name of the variables as they will be used in the function, separated by commas. By default, the names of variables at which the function will be evaluated is ‘x’ (in 1d), ‘x,y’ (in 2d) or ‘x,y,z’ (in 3d) for spatial coordinates and ‘t’ for time. You can then use these variable names in your function expression and they will be replaced by the values of these variables at which the function is currently evaluated. However, you can also choose a different set of names for the independent variables at which to evaluate your function expression. For example, if you work in spherical coordinates, you may wish to set this input parameter to ‘r,phi,theta,t’ and then use these variable names in your function expression.

Possible values: Any string

A.35 Parameters in section Boundary velocity model/GPlates model

- *Parameter name:* **Data directory**

Value: \$ASPECT_SOURCE_DIR/data/boundary-velocity/gplates/

Default: \$ASPECT_SOURCE_DIR/data/boundary-velocity/gplates/

Description: The name of a directory that contains the model data. This path may either be absolute (if starting with a '/') or relative to the current directory. The path may also include the special text '\$ASPECT_SOURCE_DIR' which will be interpreted as the path in which the ASPECT source files were located when ASPECT was compiled. This interpretation allows, for example, to reference files located in the ‘data/’ subdirectory of ASPECT.

Possible values: A directory name

- *Parameter name:* **Data file time step**

Value: 1e6

Default: 1e6

Description: Time step between following velocity files. Depending on the setting of the global ‘Use years in output instead of seconds’ flag in the input file, this number is either interpreted as seconds or as years. The default is one million, i.e., either one million seconds or one million years.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* **Decreasing file order**

Value: false

Default: false

Description: In some cases the boundary files are not numbered in increasing but in decreasing order (e.g. ‘Ma BP’). If this flag is set to ‘True’ the plugin will first load the file with the number ‘First velocity file number’ and decrease the file number during the model run.

Possible values: A boolean value (true or false)

- *Parameter name:* `First data file model time`

Value: 0

Default: 0

Description: Time from which on the velocity file with number 'First velocity file number' is used as boundary condition. Previous to this time, a no-slip boundary condition is assumed. Depending on the setting of the global 'Use years in output instead of seconds' flag in the input file, this number is either interpreted as seconds or as years.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `First data file number`

Value: 0

Default: 0

Description: Number of the first velocity file to be loaded when the model time is larger than 'First velocity file model time'.

Possible values: An integer n such that $-2147483648 \leq n \leq 2147483647$

- *Parameter name:* `Lithosphere thickness`

Value: 100000

Default: 100000

Description: Determines the depth of the lithosphere, so that the GPlates velocities can be applied at the sides of the model as well as at the surface.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `Point one`

Value: 1.570796,0.0

Default: 1.570796,0.0

Description: Point that determines the plane in which a 2D model lies in. Has to be in the format 'a,b' where a and b are theta (polar angle) and phi in radians.

Possible values: Any string

- *Parameter name:* `Point two`

Value: 1.570796,1.570796

Default: 1.570796,1.570796

Description: Point that determines the plane in which a 2D model lies in. Has to be in the format 'a,b' where a and b are theta (polar angle) and phi in radians.

Possible values: Any string

- *Parameter name:* `Scale factor`

Value: 1

Default: 1

Description: Scalar factor, which is applied to the boundary velocity. You might want to use this to scale the velocities to a reference model (e.g. with free-slip boundary) or another plate reconstruction.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Velocity file name

Value: phi.%d

Default: phi.%d

Description: The file name of the material data. Provide file in format: (Velocity file name).%d.gpml where %d is any sprintf integer qualifier, specifying the format of the current file number.

Possible values: Any string

A.36 Parameters in section Checkpointing

- *Parameter name:* Steps between checkpoint

Value: 0

Default: 0

Description: The number of timesteps between performing checkpoints. If 0 and time between checkpoint is not specified, checkpointing will not be performed. Units: None.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* Time between checkpoint

Value: 0

Default: 0

Description: The wall time between performing checkpoints. If 0, will use the checkpoint step frequency instead. Units: Seconds.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

A.37 Parameters in section Compositional fields

- *Parameter name:* Compositional field methods

Value:

Default:

Description: A comma separated list denoting the solution method of each compositional field. Each entry of the list must be one of the currently implemented field types: “field”, “particles”, or “static”.

These choices correspond to the following methods by which compositional fields gain their values:

- “field”: If a compositional field is marked with this method, then its values are computed in each time step by advecting along the values of the previous time step using the velocity field, and applying reaction rates to it. In other words, this corresponds to the usual notion of a composition field as mentioned in Section 2.7.
- “particles”: If a compositional field is marked with this method, then its values are obtained in each time step by interpolating the corresponding properties from the particles located on each cell. The time evolution therefore happens because particles move along with the velocity field, and particle properties can react with each other as well. See Section 2.15 for more information about how particles behave.
- “static”: If a compositional field is marked this way, then it does not evolve at all. Its values are simply set to the initial conditions, and will then never change.

Possible values: A list of 0 to 4294967295 elements where each element is [Any one of field, particles, static]

- *Parameter name:* `List of normalized fields`

Value:

Default:

Description: A list of integers smaller than or equal to the number of compositional fields. All compositional fields in this list will be normalized before the first timestep. The normalization is implemented in the following way: First, the sum of the fields to be normalized is calculated at every point and the global maximum is determined. Second, the compositional fields to be normalized are divided by this maximum.

Possible values: A list of 0 to 4294967295 elements where each element is [An integer n such that $0 \leq n \leq 2147483647$]

- *Parameter name:* `Mapped particle properties`

Value:

Default:

Description: A comma separated list denoting the particle properties that will be projected to those compositional fields that are of the “particles” field type.

The format of valid entries for this parameter is that of a map given as “key1: value1, key2: value2 [component2], key3: value3 [component4], ...” where each key must be a valid field name of the “particles” type, and each value must be one of the currently selected particle properties. Component is a component index of the particle property that is 0 by default, but can be set up to n-1, where n is the number of vector components of this particle property. The component indicator only needs to be set if not the first component of the particle property should be mapped (e.g. the y -component of the velocity at the particle positions).

Possible values: A key-value map of 0 to 4294967295 elements where each key is [Any string] and each value is [Any string]

- *Parameter name:* `Names of fields`

Value:

Default:

Description: A user-defined name for each of the compositional fields requested.

Possible values: A list of 0 to 4294967295 elements where each element is [Any string]

- *Parameter name:* `Number of fields`

Value: 0

Default: 0

Description: The number of fields that will be advected along with the flow field, excluding velocity, pressure and temperature.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

A.38 Parameters in section Discretization

- *Parameter name:* `Composition polynomial degree`

Value: 2

Default: 2

Description: The polynomial degree to use for the composition variable(s). As an example, a value of 2 for this parameter will yield either the element Q_2 or DGQ_2 for the compositional field(s), depending on whether we use continuous or discontinuous field(s). Units: None.

Possible values: An integer n such that $1 \leq n \leq 2147483647$

- *Parameter name:* Stokes velocity polynomial degree

Value: 2

Default: 2

Description: The polynomial degree to use for the velocity variables in the Stokes system. The polynomial degree for the pressure variable will then be one less in order to make the velocity/pressure pair conform with the usual LBB (Babuska-Brezzi) condition. In other words, we are using a Taylor-Hood element for the Stokes equations and this parameter indicates the polynomial degree of it. As an example, a value of 2 for this parameter will yield the element $Q_2^d \times Q_1$ for the d velocity components and the pressure, respectively (unless the ‘Use locally conservative discretization’ parameter is set, which modifies the pressure element). Units: None.

Possible values: An integer n such that $1 \leq n \leq 2147483647$

- *Parameter name:* Temperature polynomial degree

Value: 2

Default: 2

Description: The polynomial degree to use for the temperature variable. As an example, a value of 2 for this parameter will yield either the element Q_2 or DGQ_2 for the temperature field, depending on whether we use a continuous or discontinuous field. Units: None.

Possible values: An integer n such that $1 \leq n \leq 2147483647$

- *Parameter name:* Use discontinuous composition discretization

Value: false

Default: false

Description: Whether to use a composition discretization that is discontinuous as opposed to continuous. This then requires the assembly of face terms between cells, and weak imposition of boundary terms for the composition field via the discontinuous Galerkin method.

Possible values: A boolean value (true or false)

- *Parameter name:* Use discontinuous temperature discretization

Value: false

Default: false

Description: Whether to use a temperature discretization that is discontinuous as opposed to continuous. This then requires the assembly of face terms between cells, and weak imposition of boundary terms for the temperature field via the interior-penalty discontinuous Galerkin method.

Possible values: A boolean value (true or false)

- *Parameter name:* Use locally conservative discretization

Value: false

Default: false

Description: Whether to use a Stokes discretization that is locally conservative at the expense of a larger number of degrees of freedom (true), or to go with a cheaper discretization that does not locally conserve mass, although it is globally conservative (false).

When using a locally conservative discretization, the finite element space for the pressure is discontinuous between cells and is the polynomial space P_{-q} of polynomials of degree q in each variable separately. Here, q is one less than the value given in the parameter “Stokes velocity polynomial degree”. As a

consequence of choosing this element, it can be shown if the medium is considered incompressible that the computed discrete velocity field \mathbf{u}_h satisfies the property $\int_{\partial K} \mathbf{u}_h \cdot \mathbf{n} = 0$ for every cell K , i.e., for each cell inflow and outflow exactly balance each other as one would expect for an incompressible medium. In other words, the velocity field is locally conservative.

On the other hand, if this parameter is set to “false”, then the finite element space is chosen as Q_q . This choice does not yield the local conservation property but has the advantage of requiring fewer degrees of freedom. Furthermore, the error is generally smaller with this choice.

For an in-depth discussion of these issues and a quantitative evaluation of the different choices, see [KHB12].

Possible values: A boolean value (true or false)

A.39 Parameters in section Discretization/Stabilization parameters

- *Parameter name:* Discontinuous penalty

Value: 10

Default: 10

Description: The value used to penalize discontinuities in the discontinuous Galerkin method. This is used only for the temperature field, and not for the composition field, as pure advection does not use the interior penalty method. This is largely empirically decided – it must be large enough to ensure the bilinear form is coercive, but not so large as to penalize discontinuity at all costs.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Global composition maximum

Value: 1.7976931348623157e+308

Default: 1.7976931348623157e+308

Description: The maximum global composition values that will be used in the bound preserving limiter for the discontinuous solutions from composition advection fields. The number of the input ‘Global composition maximum’ values separated by ‘,’ has to be the same as the number of the compositional fields

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Global composition minimum

Value: -1.7976931348623157e+308

Default: -1.7976931348623157e+308

Description: The minimum global composition value that will be used in the bound preserving limiter for the discontinuous solutions from composition advection fields. The number of the input ‘Global composition minimum’ values separated by ‘,’ has to be the same as the number of the compositional fields

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Global temperature maximum

Value: 1.7976931348623157e+308

Default: 1.7976931348623157e+308

Description: The maximum global temperature value that will be used in the bound preserving limiter for the discontinuous solutions from temperature advection fields.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Global temperature minimum

Value: -1.7976931348623157e+308

Default: -1.7976931348623157e+308

Description: The minimum global temperature value that will be used in the bound preserving limiter for the discontinuous solutions from temperature advection fields.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Use artificial viscosity smoothing

Value: false

Default: false

Description: If set to false, the artificial viscosity of a cell is computed and is computed on every cell separately as discussed in [KHB12]. If set to true, the maximum of the artificial viscosity in the cell as well as the neighbors of the cell is computed and used instead.

Possible values: A boolean value (true or false)

- *Parameter name:* Use limiter for discontinuous composition solution

Value: false

Default: false

Description: Whether to apply the bound preserving limiter as a correction after having the discontinuous composition solution. Currently we apply this only to the compositional solution if the 'Global composition maximum' and 'Global composition minimum' are already defined in the .prm file. This limiter keeps the discontinuous solution in the range given by 'Global composition maximum' and 'Global composition minimum'.

Possible values: A boolean value (true or false)

- *Parameter name:* Use limiter for discontinuous temperature solution

Value: false

Default: false

Description: Whether to apply the bound preserving limiter as a correction after computing the discontinuous temperature solution. Currently we apply this only to the temperature solution if the 'Global temperature maximum' and 'Global temperature minimum' are already defined in the .prm file. This limiter keeps the discontinuous solution in the range given by 'Global temperature maximum' and 'Global temperature minimum'.

Possible values: A boolean value (true or false)

- *Parameter name:* alpha

Value: 2

Default: 2

Description: The exponent α in the entropy viscosity stabilization. Valid options are 1 or 2. The recommended setting is 2. (This parameter does not correspond to any variable in the 2012 paper by Kronbichler, Heister and Bangerth that describes ASPECT, see [KHB12]. Rather, the paper always uses 2 as the exponent in the definition of the entropy, following equation (15) of the paper. The full approach is discussed in [GPP11].) Note that this is not the thermal expansion coefficient, also commonly referred to as α .*Units:* None.

Possible values: An integer n such that $1 \leq n \leq 2$

- Parameter name: **beta**

Value: 0.078

Default: 0.078

Description: The β factor in the artificial viscosity stabilization. This parameter controls the maximum dissipation of the entropy viscosity, which is the part that only scales with the cell diameter and the maximum velocity in the cell, but does not depend on the solution field itself or its residual. An appropriate value for 2d is 0.078 and 0.117 for 3d. (For historical reasons, the name used here is different from the one used in the 2012 paper by Kronbichler, Heister and Bangerth that describes ASPECT, see [KHB12]. This parameter corresponds to the factor α_{\max} in the formulas following equation (15) of the paper. After further experiments, we have also chosen to use a different value than described there: It can be chosen as stated there for uniformly refined meshes, but it needs to be chosen larger if the mesh has cells that are not squares or cubes.) Units: None.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: **cR**

Value: 0.33

Default: 0.33

Description: The c_R factor in the entropy viscosity stabilization. This parameter controls the part of the entropy viscosity that depends on the solution field itself and its residual in addition to the cell diameter and the maximum velocity in the cell. (For historical reasons, the name used here is different from the one used in the 2012 paper by Kronbichler, Heister and Bangerth that describes ASPECT, see [KHB12]. This parameter corresponds to the factor α_E in the formulas following equation (15) of the paper. After further experiments, we have also chosen to use a different value than described there.) Units: None.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: **gamma**

Value: 0.0

Default: 0.0

Description: The strain rate scaling factor in the artificial viscosity stabilization. This parameter determines how much the strain rate (in addition to the velocity) should influence the stabilization. (This parameter does not correspond to any variable in the 2012 paper by Kronbichler, Heister and Bangerth that describes ASPECT, see [KHB12]. Rather, the paper always uses 0, i.e. they specify the maximum dissipation ν_h^{\max} as $\nu_h^{\max}|_K = \alpha_{\max} h_K \|\mathbf{u}\|_{\infty,K}$. Here, we use $\||\mathbf{u}| + \gamma h_K |\boldsymbol{\varepsilon}(\mathbf{u})|\|_{\infty,K}$ instead of $\|\mathbf{u}\|_{\infty,K}$. Units: None.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.40 Parameters in section Formulation

- Parameter name: **Enable additional Stokes RHS**

Value: false

Default: false

Description: Whether to ask the material model for additional terms for the right-hand side of the Stokes equation. This feature is likely only used when implementing force vectors for manufactured solution problems and requires filling additional outputs of type AdditionalMaterialOutputsStokesRHS.

Possible values: A boolean value (true or false)

- *Parameter name:* Formulation

Value: custom

Default: custom

Description: Select a formulation for the basic equations. Different published formulations are available in ASPECT (see the list of possible values for this parameter in the manual for available options). Two ASPECT specific options are

1. ‘isothermal compression’: ASPECT’s original formulation, using the explicit compressible mass equation, and the full density for the temperature equation.
2. ‘custom’: A custom selection of ‘Mass conservation’ and ‘Temperature equation’.

Note: Warning: The ‘custom’ option is implemented for advanced users that want full control over the equations solved. It is possible to choose inconsistent formulations and no error checking is performed on the consistency of the resulting equations.

Note: The ‘anelastic liquid approximation’ option here can also be used to set up the ‘truncated anelastic liquid approximation’ as long as this option is chosen together with a material model that defines a density that depends on temperature and depth and not on the pressure.

Possible values: Any one of isothermal compression, custom, anelastic liquid approximation, Boussinesq approximation

- *Parameter name:* Mass conservation

Value: ask material model

Default: ask material model

Description: Possible approximations for the density derivatives in the mass conservation equation. Note that this parameter is only evaluated if ‘Formulation’ is set to ‘custom’. Other formulations ignore the value of this parameter.

Possible values: Any one of incompressible, isothermal compression, hydrostatic compression, reference density profile, implicit reference density profile, ask material model

- *Parameter name:* Temperature equation

Value: real density

Default: real density

Description: Possible approximations for the density in the temperature equation. Possible approximations are ‘real density’ and ‘reference density profile’. Note that this parameter is only evaluated if ‘Formulation’ is set to ‘custom’. Other formulations ignore the value of this parameter.

Possible values: Any one of real density, reference density profile

A.41 Parameters in section Free surface

- *Parameter name:* Additional tangential mesh velocity boundary indicators

Value:

Default:

Description: A comma separated list of names denoting those boundaries where the mesh is allowed to move tangential to the boundary. All tangential mesh movements along those boundaries that have tangential material velocity boundary conditions are allowed by default, this parameter allows to generate mesh movements along other boundaries that are open, or have prescribed material velocities or tractions.

The names of the boundaries listed here can either be numbers (in which case they correspond to the numerical boundary indicators assigned by the geometry object), or they can correspond to any of the symbolic names the geometry object may have provided for each part of the boundary. You may want to compare this with the documentation of the geometry model you use in your model.

Possible values: A list of 0 to 4294967295 elements where each element is [Any string]

- *Parameter name:* Free surface boundary indicators

Value:

Default:

Description: A comma separated list of names denoting those boundaries where there is a free surface. Set to nothing to disable all free surface computations.

The names of the boundaries listed here can either be numbers (in which case they correspond to the numerical boundary indicators assigned by the geometry object), or they can correspond to any of the symbolic names the geometry object may have provided for each part of the boundary. You may want to compare this with the documentation of the geometry model you use in your model.

Possible values: A list of 0 to 4294967295 elements where each element is [Any string]

- *Parameter name:* Free surface stabilization theta

Value: 0.5

Default: 0.5

Description: Theta parameter described in Kaus et. al. 2010. An unstabilized free surface can overshoot its equilibrium position quite easily and generate unphysical results. One solution is to use a quasi-implicit correction term to the forces near the free surface. This parameter describes how much the free surface is stabilized with this term, where zero is no stabilization, and one is fully implicit.

Possible values: A floating point number v such that $0 \leq v \leq 1$

- *Parameter name:* Surface velocity projection

Value: normal

Default: normal

Description: After each time step the free surface must be advected in the direction of the velocity field. Mass conservation requires that the mesh velocity is in the normal direction of the surface. However, for steep topography or large curvature, advection in the normal direction can become ill-conditioned, and instabilities in the mesh can form. Projection of the mesh velocity onto the local vertical direction can preserve the mesh quality better, but at the cost of slightly poorer mass conservation of the domain.

Possible values: Any one of normal, vertical

A.42 Parameters in section Geometry model

- *Parameter name:* Model name

Value: box

Default: unspecified

Description: Select one of the following models:

‘box’: A box geometry parallel to the coordinate directions. The extent of the box in each coordinate direction is set in the parameter file. The box geometry labels its $2^*\dim$ sides as follows: in 2d, boundary indicators 0 through 3 denote the left, right, bottom and top boundaries; in 3d, boundary indicators 0 through 5 indicate left, right, front, back, bottom and top boundaries (see also the documentation of the deal.II class “GeometryInfo”). You can also use symbolic names “left”, “right”, etc., to refer to these boundaries in input files. It is also possible to add initial topography to the box model. Note however that this is done after the last initial adaptive refinement cycle. Also, initial topography is supposed to be small, as it is not taken into account when depth or a representative point is computed.

‘box with lithosphere boundary indicators’: A box geometry parallel to the coordinate directions. The extent of the box in each coordinate direction is set in the parameter file. This geometry model labels its sides with $2^*\dim+2^*(\dim-1)$ boundary indicators: in 2d, boundary indicators 0 through 3 denote the left, right, bottom and top boundaries, while indicators 4 and 5 denote the upper part of the left and right vertical boundary, respectively. In 3d, boundary indicators 0 through 5 indicate left, right, front, back, bottom and top boundaries (see also the documentation of the deal.II class “GeometryInfo”), while indicators 6, 7, 8 and 9 denote the left, right, front and back upper parts of the vertical boundaries, respectively. You can also use symbolic names “left”, “right”, “left lithosphere”, etc., to refer to these boundaries in input files.

Note that for a given “Global refinement level” and no user-specified “Repetitions”, the lithosphere part of the mesh will be more refined.

The additional boundary indicators for the lithosphere allow for selecting boundary conditions for the lithosphere different from those for the underlying mantle. An example application of this geometry is to prescribe a velocity on the lithospheric plates, but use open boundary conditions underneath.

‘chunk’: A geometry which can be described as a chunk of a spherical shell, bounded by lines of longitude, latitude and radius. The minimum and maximum longitude, (latitude) and depth of the chunk is set in the parameter file. The chunk geometry labels its $2^*\dim$ sides as follows: “west” and “east”: minimum and maximum longitude, “south” and “north”: minimum and maximum latitude, “inner” and “outer”: minimum and maximum radii. Names in the parameter files are as follows: Chunk (minimum || maximum) (longitude || latitude): edges of geographical quadrangle (in degrees) Chunk (inner || outer) radius: Radii at bottom and top of chunk(Longitude || Latitude || Radius) repetitions: number of cells in each coordinate direction.

‘ellipsoidal chunk’: A 3D chunk geometry that accounts for Earth’s ellipticity (default assuming the WGS84 ellipsoid definition) which can be defined in non-coordinate directions. In the description of the ellipsoidal chunk, two of the ellipsoidal axes have the same length so that there is only a semi-major axis and a semi-minor axis. The user has two options for creating an ellipsoidal chunk geometry: 1) by defining two opposing points (SW and NE or NW and SE) a coordinate parallel ellipsoidal chunk geometry will be created. 2) by defining three points a non-coordinate parallel ellipsoidal chunk will be created. The points are defined in the input file by longitude:latitude. It is also possible to define additional subdivisions of the mesh in each direction. Faces of the model are defined as 0, west; 1,east; 2, south; 3, north; 4, inner; 5, outer.

This geometry model supports initial topography for deforming the initial mesh.

‘sphere’: Geometry model for sphere with a user specified radius. This geometry has only a single boundary, so the only valid boundary indicator to specify in the input file is “0”. It can also be referenced by the symbolic name “surface” in input files.

‘spherical shell’: A geometry representing a spherical shell or a piece of it. Inner and outer radii are read from the parameter file in subsection ‘Spherical shell’.

The model assigns boundary indicators as follows: In 2d, inner and outer boundaries get boundary indicators zero and one, and if the opening angle set in the input file is less than 360, then left and right boundaries are assigned indicators two and three. These boundaries can also be referenced using the symbolic names ‘inner’, ‘outer’ and (if applicable) ‘left’, ‘right’.

In 3d, inner and outer indicators are treated as in 2d. If the opening angle is chosen as 90 degrees, i.e., the domain is the intersection of a spherical shell and the first octant, then indicator 2 is at the face $x = 0$, 3 at $y = 0$, and 4 at $z = 0$. These last three boundaries can then also be referred to as ‘east’, ‘west’ and ‘south’ symbolically in input files.

Possible values: Any one of box, box with lithosphere boundary indicators, chunk, ellipsoidal chunk, sphere, spherical shell, unspecified

A.43 Parameters in section Geometry model/Box

- *Parameter name:* Box origin X coordinate

Value: 0

Default: 0

Description: X coordinate of box origin. Units: m.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Box origin Y coordinate

Value: 0

Default: 0

Description: Y coordinate of box origin. Units: m.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Box origin Z coordinate

Value: 0

Default: 0

Description: Z coordinate of box origin. This value is ignored if the simulation is in 2d. Units: m.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* X extent

Value: 1

Default: 1

Description: Extent of the box in x-direction. Units: m.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* X periodic

Value: false

Default: false

Description: Whether the box should be periodic in X direction

Possible values: A boolean value (true or false)

- *Parameter name:* X repetitions

Value: 1

Default: 1

Description: Number of cells in X direction.

Possible values: An integer n such that $1 \leq n \leq 2147483647$

- *Parameter name:* Y extent
Value: 1
Default: 1
Description: Extent of the box in y-direction. Units: m.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Y periodic
Value: false
Default: false
Description: Whether the box should be periodic in Y direction
Possible values: A boolean value (true or false)
- *Parameter name:* Y repetitions
Value: 1
Default: 1
Description: Number of cells in Y direction.
Possible values: An integer n such that $1 \leq n \leq 2147483647$
- *Parameter name:* Z extent
Value: 1
Default: 1
Description: Extent of the box in z-direction. This value is ignored if the simulation is in 2d. Units: m.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Z periodic
Value: false
Default: false
Description: Whether the box should be periodic in Z direction
Possible values: A boolean value (true or false)
- *Parameter name:* Z repetitions
Value: 1
Default: 1
Description: Number of cells in Z direction.
Possible values: An integer n such that $1 \leq n \leq 2147483647$

A.44 Parameters in section Geometry model/Box with lithosphere boundary indicators

- *Parameter name:* Box origin X coordinate
Value: 0
Default: 0
Description: X coordinate of box origin. Units: m.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Box origin Y coordinate
Value: 0
Default: 0
Description: Y coordinate of box origin. Units: m.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$
- **Parameter name:** Box origin Z coordinate
Value: 0
Default: 0
Description: Z coordinate of box origin. This value is ignored if the simulation is in 2d. Units: m.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$
- **Parameter name:** Lithospheric thickness
Value: 0.2
Default: 0.2
Description: The thickness of the lithosphere used to create additional boundary indicators to set specific boundary conditions for the lithosphere.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- **Parameter name:** X extent
Value: 1
Default: 1
Description: Extent of the box in x-direction. Units: m.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- **Parameter name:** X periodic
Value: false
Default: false
Description: Whether the box should be periodic in X direction.
Possible values: A boolean value (true or false)
- **Parameter name:** X periodic lithosphere
Value: false
Default: false
Description: Whether the box should be periodic in X direction in the lithosphere.
Possible values: A boolean value (true or false)
- **Parameter name:** X repetitions
Value: 1
Default: 1
Description: Number of cells in X direction of the lower box. The same number of repetitions will be used in the upper box.
Possible values: An integer n such that $1 \leq n \leq 2147483647$

- *Parameter name:* Y extent

Value: 1

Default: 1

Description: Extent of the box in y-direction. Units: m.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Y periodic

Value: false

Default: false

Description: Whether the box should be periodic in Y direction.

Possible values: A boolean value (true or false)

- *Parameter name:* Y periodic lithosphere

Value: false

Default: false

Description: Whether the box should be periodic in Y direction in the lithosphere. This value is ignored if the simulation is in 2d.

Possible values: A boolean value (true or false)

- *Parameter name:* Y repetitions

Value: 1

Default: 1

Description: Number of cells in Y direction of the lower box. If the simulation is in 3d, the same number of repetitions will be used in the upper box.

Possible values: An integer n such that $1 \leq n \leq 2147483647$

- *Parameter name:* Y repetitions lithosphere

Value: 1

Default: 1

Description: Number of cells in Y direction in the lithosphere. This value is ignored if the simulation is in 3d.

Possible values: An integer n such that $1 \leq n \leq 2147483647$

- *Parameter name:* Z extent

Value: 1

Default: 1

Description: Extent of the box in z-direction. This value is ignored if the simulation is in 2d. Units: m.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Z periodic

Value: false

Default: false

Description: Whether the box should be periodic in Z direction. This value is ignored if the simulation is in 2d.

Possible values: A boolean value (true or false)

- *Parameter name:* Z repetitions

Value: 1

Default: 1

Description: Number of cells in Z direction of the lower box. This value is ignored if the simulation is in 2d.

Possible values: An integer n such that $1 \leq n \leq 2147483647$

- *Parameter name:* Z repetitions lithosphere

Value: 1

Default: 1

Description: Number of cells in Z direction in the lithosphere. This value is ignored if the simulation is in 2d.

Possible values: An integer n such that $1 \leq n \leq 2147483647$

A.45 Parameters in section Geometry model/Chunk

- *Parameter name:* Chunk inner radius

Value: 0

Default: 0

Description: Radius at the bottom surface of the chunk. Units: m.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Chunk maximum latitude

Value: 1

Default: 1

Description: Maximum latitude of the chunk. This value is ignored if the simulation is in 2d. Units: degrees.

Possible values: A floating point number v such that $-90 \leq v \leq 90$

- *Parameter name:* Chunk maximum longitude

Value: 1

Default: 1

Description: Maximum longitude of the chunk. Units: degrees.

Possible values: A floating point number v such that $-180 \leq v \leq 360$

- *Parameter name:* Chunk minimum latitude

Value: 0

Default: 0

Description: Minimum latitude of the chunk. This value is ignored if the simulation is in 2d. Units: degrees.

Possible values: A floating point number v such that $-90 \leq v \leq 90$

- *Parameter name:* `Chunk minimum longitude`
Value: 0
Default: 0
Description: Minimum longitude of the chunk. Units: degrees.
Possible values: A floating point number v such that $-180 \leq v \leq 360$
- *Parameter name:* `Chunk outer radius`
Value: 1
Default: 1
Description: Radius at the top surface of the chunk. Units: m.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* `Latitude repetitions`
Value: 1
Default: 1
Description: Number of cells in latitude. This value is ignored if the simulation is in 2d
Possible values: An integer n such that $1 \leq n \leq 2147483647$
- *Parameter name:* `Longitude repetitions`
Value: 1
Default: 1
Description: Number of cells in longitude.
Possible values: An integer n such that $1 \leq n \leq 2147483647$
- *Parameter name:* `Radius repetitions`
Value: 1
Default: 1
Description: Number of cells in radius.
Possible values: An integer n such that $1 \leq n \leq 2147483647$

A.46 Parameters in section Geometry model/Ellipsoidal chunk

- *Parameter name:* `Depth`
Value: 500000.0
Default: 500000.0
Description: Bottom depth of model region.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* `Depth subdivisions`
Value: 1
Default: 1
Description: The number of subdivisions of the coarse (initial) mesh in depth.
Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* **East-West subdivisions**

Value: 1

Default: 1

Description: The number of subdivisions of the coarse (initial) mesh in the East-West direction.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* **Eccentricity**

Value: 8.1819190842622e-2

Default: 8.1819190842622e-2

Description: Eccentricity of the ellipsoid. Zero is a perfect sphere, default (8.1819190842622e-2) is WGS84.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* **NE corner**

Value:

Default:

Description: Longitude:latitude in degrees of the North-East corner point of model region. The North-East direction is positive. If one of the three corners is not provided the missing corner value will be calculated so all faces are parallel.

Possible values: Any string

- *Parameter name:* **NW corner**

Value:

Default:

Description: Longitude:latitude in degrees of the North-West corner point of model region. The North-East direction is positive. If one of the three corners is not provided the missing corner value will be calculated so all faces are parallel.

Possible values: Any string

- *Parameter name:* **North-South subdivisions**

Value: 1

Default: 1

Description: The number of subdivisions of the coarse (initial) mesh in the North-South direction.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* **SE corner**

Value:

Default:

Description: Longitude:latitude in degrees of the South-East corner point of model region. The North-East direction is positive. If one of the three corners is not provided the missing corner value will be calculated so all faces are parallel.

Possible values: Any string

- *Parameter name:* SW_corner

Value:

Default:

Description: Longitude:latitude in degrees of the South-West corner point of model region. The North-East direction is positive. If one of the three corners is not provided the missing corner value will be calculated so all faces are parallel.

Possible values: Any string

- *Parameter name:* Semi-major_axis

Value: 6378137.0

Default: 6378137.0

Description: The semi-major axis (a) of an ellipsoid. This is the radius for a sphere (eccentricity=0). Default WGS84 semi-major axis.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.47 Parameters in section Geometry model/Initial topography model

- *Parameter name:* Model_name

Value: zero topography

Default: zero topography

Description: Select one of the following models:

‘ascii data’: Implementation of a model in which the surface topography is derived from a file containing data in ascii format. The following geometry models are currently supported: box, chunk, spherical shell. Note the required format of the input data: The first lines may contain any number of comments if they begin with ‘#’, but one of these lines needs to contain the number of grid points in each dimension as for example ‘# POINTS: 3 3’. The order of the data columns has to be ‘x’, ‘Topography [m]’ in a 2d model and ‘x’, ‘y’, ‘Topography [m]’ in a 3d model, which means that there has to be a single column containing the topography. Note that the data in the input file needs to be sorted in a specific order: the first coordinate needs to ascend first, followed by the second in order to assign the correct data to the prescribed coordinates. If you use a spherical model, then the data will still be handled as Cartesian, however the assumed grid changes. ‘x’ will be replaced by the azimuth angle in radians and ‘y’ by the polar angle in radians measured positive from the north pole. The grid will be assumed to be a longitude-colatitude grid. Note that the order of spherical coordinates is ‘phi’, ‘theta’ and not ‘theta’, ‘phi’, since this allows for dimension independent expressions.

‘prm polygon’: An initial topography model that defines the initial topography as constant inside each of a set of polygonal parts of the surface. The polygons, and their associated surface elevation, are defined in the ‘Geometry model/Initial topography/Prm polygon’ section.

‘zero topography’: Implementation of a model in which the initial topography is zero.

Possible values: Any one of ascii data, prm polygon, zero topography

A.48 Parameters in section Geometry model/Initial topography model/Ascii data model

- *Parameter name:* Data_directory

Value: \$ASPECT_SOURCE_DIR/data/geometry-model/initial-topography-model/ascii-data/test/

Default: \$ASPECT_SOURCE_DIR/data/geometry-model/initial-topography-model/ascii-data/test/

Description: The name of a directory that contains the model data. This path may either be absolute (if starting with a '/') or relative to the current directory. The path may also include the special text '\$ASPECT_SOURCE_DIR' which will be interpreted as the path in which the ASPECT source files were located when ASPECT was compiled. This interpretation allows, for example, to reference files located in the 'data/' subdirectory of ASPECT.

Possible values: A directory name

- *Parameter name:* `Data file name`

Value: `box_2d_%s.0.txt`

Default: `box_2d_%s.0.txt`

Description: The file name of the material data. Provide file in format: (Velocity file name).%s%d where %s is a string specifying the boundary of the model according to the names of the boundary indicators (of a box or a spherical shell).%d is any sprintf integer qualifier, specifying the format of the current file number.

Possible values: Any string

- *Parameter name:* `Scale factor`

Value: 1

Default: 1

Description: Scalar factor, which is applied to the boundary velocity. You might want to use this to scale the velocities to a reference model (e.g. with free-slip boundary) or another plate reconstruction. Another way to use this factor is to convert units of the input files. The unit is assumed to be m/s or m/yr depending on the 'Use years in output instead of seconds' flag. If you provide velocities in cm/yr set this factor to 0.01.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.49 Parameters in section Geometry model/Initial topography model/Prm polygon

- *Parameter name:* `Topography parameters`

Value:

Default:

Description: Set the topography height and the polygon which should be set to that height. The format is : "The topography height extgreater The point list describing a polygon & The next topography height extgreater the next point list describing a polygon." The format for the point list describing the polygon is "x1,y1;x2,y2". For example for two triangular areas of 100 and -100 meters high set: '100 extgreater 0,0;5,5;0,10 & -100 extgreater 10,10;10,15;20,15'. Units of the height are always in meters. The units of the coordinates are dependent on the geometry model. In the box model they are in meters, in the chunks they are in degrees, etc. Please refer to the manual of the individual geometry model to see how the topography is implemented.

Possible values: Any string

A.50 Parameters in section Geometry model/Sphere

- *Parameter name:* `Radius`

Value: 6371000

Default: 6371000

Description: Radius of the sphere. Units: m.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.51 Parameters in section Geometry model/Spherical shell

- *Parameter name:* Cells along circumference

Value: 0

Default: 0

Description: The number of cells in circumferential direction that are created in the coarse mesh in 2d. If zero, this number is chosen automatically in a way that produces meshes in which cells have a reasonable aspect ratio for models in which the depth of the mantle is roughly that of the Earth. For planets with much shallower mantles and larger cores, you may want to choose a larger number to avoid cells that are elongated in tangential and compressed in radial direction.

In 3d, the number of cells is computed differently and does not have an easy interpretation. Valid values for this parameter in 3d are 0 (let this class choose), 6, 12 and 96. Other possible values may be discussed in the documentation of the deal.II function GridGenerator::hyper_shell. The parameter is best left at its default in 3d.

In either case, this parameter is ignored unless the opening angle of the domain is 360 degrees.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* Inner radius

Value: 3481000

Default: 3481000

Description: Inner radius of the spherical shell. Units: m.

Note: The default value of 3,481,000 m equals the radius of a sphere with equal volume as Earth (i.e., 6371 km) minus the average depth of the core-mantle boundary (i.e., 2890 km).

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Opening angle

Value: 360

Default: 360

Description: Opening angle in degrees of the section of the shell that we want to build. The only opening angles that are allowed for this geometry are 90, 180, and 360 in 2d; and 90 and 360 in 3d. Units: degrees.

Possible values: A floating point number v such that $0 \leq v \leq 360$

- *Parameter name:* Outer radius

Value: 6336000

Default: 6336000

Description: Outer radius of the spherical shell. Units: m.

Note: The default value of 6,336,000 m equals the radius of a sphere with equal volume as Earth (i.e., 6371 km) minus the average depth of the mantle-crust interface (i.e., 35 km).

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.52 Parameters in section Gravity model

- *Parameter name:* Model name

Value: vertical

Default: unspecified

Description: Select one of the following models:

‘ascii data’: Gravity is read from a file that describes the reference state. The default profile follows the preliminary reference Earth model (PREM, Dziewonski and Anderson, 1981). Note the required format of the input data: The first lines may contain any number of comments if they begin with ‘#’, but one of these lines needs to contain the number of points in the reference state as for example ‘# POINTS: 3’. Following the comment lines there has to be a single line containing the names of all data columns, separated by arbitrarily many spaces. Column names are not allowed to contain spaces. The file can contain unnecessary columns, but for this plugin it needs to at least provide a column named ‘gravity’. Note that the data lines in the file need to be sorted in order of increasing depth from 0 to the maximal depth in the model domain. Points in the model that are outside of the provided depth range will be assigned the maximum or minimum depth values, respectively. Points do not need to be equidistant, but the computation of properties is optimized in speed if they are.

‘function’: Gravity is given in terms of an explicit formula that is elaborated in the parameters in section “Gravity model|Function”. The format of these functions follows the syntax understood by the muparser library, see Section 4.7.3.

‘radial constant’: A gravity model in which the gravity has a constant magnitude and the direction is radial (pointing inward if the value is positive). The magnitude is read from the parameter file in subsection ‘Radial constant’.

‘radial earth-like’: This plugin has been removed due to its misleading name. The included profile was hard-coded and was less earth-like than the ‘ascii data’ plugin, which uses the profile of the Preliminary Reference Earth Model (PREM). Use ‘ascii data’ instead of ‘radial earth-like’.

‘radial linear’: A gravity model which is radial (pointing inward if the gravity is positive) and the magnitude changes linearly with depth. The magnitude of gravity at the surface and bottom is read from the input file in a section “Gravity model/Radial linear”.

‘vertical’: A gravity model in which the gravity direction is vertical (pointing downward for positive values) and at a constant magnitude by default equal to one.

Possible values: Any one of ascii data, function, radial constant, radial earth-like, radial linear, vertical, unspecified

A.53 Parameters in section Gravity model/Ascii data model

- *Parameter name:* Data directory

Value: \$ASPECT_SOURCE_DIR/data/gravity-model/

Default: \$ASPECT_SOURCE_DIR/data/gravity-model/

Description: The name of a directory that contains the model data. This path may either be absolute (if starting with a '/') or relative to the current directory. The path may also include the special text ‘\$ASPECT_SOURCE_DIR’ which will be interpreted as the path in which the ASPECT source files were located when ASPECT was compiled. This interpretation allows, for example, to reference files located in the ‘data/’ subdirectory of ASPECT.

Possible values: A directory name

- *Parameter name:* Data file name

Value: prem.txt

Default: prem.txt

Description: The file name of the material data. Provide file in format: (Velocity file name).%s%d where %s is a string specifying the boundary of the model according to the names of the boundary indicators (of a box or a spherical shell).%d is any sprintf integer qualifier, specifying the format of the current file number.

Possible values: Any string

- *Parameter name:* Scale factor

Value: 1

Default: 1

Description: Scalar factor, which is applied to the boundary velocity. You might want to use this to scale the velocities to a reference model (e.g. with free-slip boundary) or another plate reconstruction. Another way to use this factor is to convert units of the input files. The unit is assumed to be m/s or m/yr depending on the 'Use years in output instead of seconds' flag. If you provide velocities in cm/yr set this factor to 0.01.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.54 Parameters in section Gravity model/Function

- *Parameter name:* Function constants

Value:

Default:

Description: Sometimes it is convenient to use symbolic constants in the expression that describes the function, rather than having to use its numeric value everywhere the constant appears. These values can be defined using this parameter, in the form 'var1=value1, var2=value2, ...'.

A typical example would be to set this runtime parameter to 'pi=3.1415926536' and then use 'pi' in the expression of the actual formula. (That said, for convenience this class actually defines both 'pi' and 'Pi' by default, but you get the idea.)

Possible values: Any string

- *Parameter name:* Function expression

Value: 0; 0

Default: 0; 0

Description: The formula that denotes the function you want to evaluate for particular values of the independent variables. This expression may contain any of the usual operations such as addition or multiplication, as well as all of the common functions such as 'sin' or 'cos'. In addition, it may contain expressions like 'if(x>0, 1, -1)' where the expression evaluates to the second argument if the first argument is true, and to the third argument otherwise. For a full overview of possible expressions accepted see the documentation of the muparser library at <http://muparser.beltoforion.de/>.

If the function you are describing represents a vector-valued function with multiple components, then separate the expressions for individual components by a semicolon.

Possible values: Any string

- *Parameter name:* Variable names

Value: x,y,t

Default: x,y,t

Description: The name of the variables as they will be used in the function, separated by commas. By default, the names of variables at which the function will be evaluated is ‘x’ (in 1d), ‘x,y’ (in 2d) or ‘x,y,z’ (in 3d) for spatial coordinates and ‘t’ for time. You can then use these variable names in your function expression and they will be replaced by the values of these variables at which the function is currently evaluated. However, you can also choose a different set of names for the independent variables at which to evaluate your function expression. For example, if you work in spherical coordinates, you may wish to set this input parameter to ‘r,phi,theta,t’ and then use these variable names in your function expression.

Possible values: Any string

A.55 Parameters in section Gravity model/Radial constant

- *Parameter name:* Magnitude

Value: 9.81

Default: 9.81

Description: Magnitude of the gravity vector in m/s^2 . For positive values the direction is radially inward towards the center of the earth.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

A.56 Parameters in section Gravity model/Radial linear

- *Parameter name:* Magnitude at bottom

Value: 10.7

Default: 10.7

Description: Magnitude of the radial gravity vector at the bottom of the domain. ‘Bottom’ means the maximum depth in the chosen geometry, and for example represents the core-mantle boundary in the case of the ‘spherical shell’ geometry model, and the center in the case of the ‘sphere’ geometry model. Units: m/s^2

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Magnitude at surface

Value: 9.8

Default: 9.8

Description: Magnitude of the radial gravity vector at the surface of the domain. Units: m/s^2

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

A.57 Parameters in section Gravity model/Vertical

- *Parameter name:* Magnitude

Value: 1

Default: 1

Description: Value of the gravity vector in m/s^2 directed along negative y (2D) or z (3D) axis (if the magnitude is positive).

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

A.58 Parameters in section Heating model

- *Parameter name:* List of model names

Value:

Default:

Description: A comma separated list of heating models that will be used to calculate the heating terms in the energy equation. The results of each of these criteria, i.e., the heating source terms and the latent heat terms for the left hand side will be added.

The following heating models are available:

‘adiabatic heating’: Implementation of a standard and a simplified model of adiabatic heating.

‘adiabatic heating of melt’: Implementation of a standard and a simplified model of adiabatic heating of melt. The full model implements the heating term $\alpha T(-\phi \mathbf{u}_s \cdot \nabla p) + \alpha T(\phi \mathbf{u}_f \cdot ablap)$. For full adiabatic heating, this has to be used in combination with the heating model ‘adiabatic heating’ to also include adiabatic heating for the solid part, and the full heating term is then $\alpha T((1-\phi)\mathbf{u}_s \cdot \nabla p) + \alpha T(\phi \mathbf{u}_f \cdot \nabla p)$.

‘compositional heating’: Implementation of a model in which magnitude of internal heat production is determined from fixed values assigned to each compositional field. These values are interpreted as having units W/m^3 .

‘constant heating’: Implementation of a model in which the heating rate is constant.

‘function’: Implementation of a model in which the heating rate is given in terms of an explicit formula that is elaborated in the parameters in section “Heating model|Function”. The format of these functions follows the syntax understood by the muparser library, see Section 4.7.3.

The formula is interpreted as having units W/kg .

Since the symbol t indicating time may appear in the formulas for the heating rate, it is interpreted as having units seconds unless the global parameter “Use years in output instead of seconds” is set.

‘latent heat’: Implementation of a standard model for latent heat.

‘latent heat melt’: Implementation of a standard model for latent heat of melting. This assumes that there is a compositional field called porosity, and it uses the reaction term of this field (the fraction of material that melted in the current time step) multiplied by a constant entropy change for melting all of the material as source term of the heating model. If there is no field called porosity, the heating terms are 0.

‘radioactive decay’: Implementation of a model in which the internal heating rate is radioactive decaying in the following rule:

$$(\text{initial concentration}) \cdot 0.5^{\text{time}/(\text{half life})}$$

The crust and mantle can have different concentrations, and the crust can be defined either by depth or by a certain compositional field. The formula is interpreted as having units W/kg .

‘shear heating’: Implementation of a standard model for shear heating. Adds the term: $2\eta (\varepsilon - \frac{1}{3}\text{tr}\varepsilon\mathbf{1}) : (\varepsilon - \frac{1}{3}\text{tr}\varepsilon\mathbf{1})$ to the right-hand side of the temperature equation.

‘shear heating with melt’: Implementation of a standard model for shear heating of migrating melt, including bulk (compression) heating $\xi (\nabla \cdot \mathbf{u}_s)^2$ and heating due to melt segregation $\frac{\eta_f \phi^2}{k} (\mathbf{u}_f - \mathbf{u}_s)^2$. For full shear heating, this has to be used in combination with the heating model shear heating to also include shear heating for the solid part.

Possible values: A comma-separated list of any of adiabatic heating, adiabatic heating of melt, compositional heating, constant heating, function, latent heat, latent heat melt, radioactive decay, shear heating, shear heating with melt

A.59 Parameters in section Heating model/Adiabatic heating

- *Parameter name:* Use simplified adiabatic heating

Value: false

Default: false

Description: A flag indicating whether the adiabatic heating should be simplified from $\alpha T(\mathbf{u} \cdot \nabla p)$ to $\alpha \rho T(\mathbf{u} \cdot \mathbf{g})$.

Possible values: A boolean value (true or false)

A.60 Parameters in section Heating model/Adiabatic heating of melt

- *Parameter name:* Use simplified adiabatic heating

Value: false

Default: false

Description: A flag indicating whether the adiabatic heating should be simplified from $\alpha T(\mathbf{u} \cdot \nabla p)$ to $\alpha \rho T(\mathbf{u} \cdot \mathbf{g})$.

Possible values: A boolean value (true or false)

A.61 Parameters in section Heating model/Compositional heating

- *Parameter name:* Compositional heating values

Value: 0

Default: 0

Description: List of heat production per unit volume values for background and compositional fields, for a total of N+1 values, where the first value corresponds to the background material, and N is the number of compositional fields. Units: W/m^3 .

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Use compositional field for heat production averaging

Value: 1

Default: 1

Description: A list of integers with as many entries as compositional fields plus one. The first entry corresponds to the background material, each following entry corresponds to a particular compositional field. If the entry for a field is '1' this field is considered during the computation of volume fractions, if it is '0' the field is ignored. This is useful if some compositional fields are used to track properties like finite strain that should not contribute to heat production. The first entry determines whether the background field contributes to heat production or not (essentially similar to setting its 'Compositional heating values' to zero, but included for consistency in the length of the input lists).

Possible values: A list of 0 to 4294967295 elements where each element is [An integer n such that $0 \leq n \leq 1$]

A.62 Parameters in section Heating model/Constant heating

- *Parameter name:* Radiogenic heating rate

Value: 0e0

Default: 0e0

Description: The specific rate of heating due to radioactive decay (or other bulk sources you may want to describe). This parameter corresponds to the variable H in the temperature equation stated in the manual, and the heating term is $h_0 H$. Units: W/kg.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.63 Parameters in section Heating model/Function

- *Parameter name:* Function constants

Value:

Default:

Description: Sometimes it is convenient to use symbolic constants in the expression that describes the function, rather than having to use its numeric value everywhere the constant appears. These values can be defined using this parameter, in the form ‘var1=value1, var2=value2, ...’.

A typical example would be to set this runtime parameter to ‘pi=3.1415926536’ and then use ‘pi’ in the expression of the actual formula. (That said, for convenience this class actually defines both ‘pi’ and ‘Pi’ by default, but you get the idea.)

Possible values: Any string

- *Parameter name:* Function expression

Value: 0

Default: 0

Description: The formula that denotes the function you want to evaluate for particular values of the independent variables. This expression may contain any of the usual operations such as addition or multiplication, as well as all of the common functions such as ‘sin’ or ‘cos’. In addition, it may contain expressions like ‘if($x > 0$, 1, -1)’ where the expression evaluates to the second argument if the first argument is true, and to the third argument otherwise. For a full overview of possible expressions accepted see the documentation of the muparser library at <http://muparser.beltofornion.de/>.

If the function you are describing represents a vector-valued function with multiple components, then separate the expressions for individual components by a semicolon.

Possible values: Any string

- *Parameter name:* Variable names

Value: x,y,t

Default: x,y,t

Description: The name of the variables as they will be used in the function, separated by commas. By default, the names of variables at which the function will be evaluated is ‘x’ (in 1d), ‘x,y’ (in 2d) or ‘x,y,z’ (in 3d) for spatial coordinates and ‘t’ for time. You can then use these variable names in your function expression and they will be replaced by the values of these variables at which the function is currently evaluated. However, you can also choose a different set of names for the independent variables at which to evaluate your function expression. For example, if you work in spherical coordinates, you may wish to set this input parameter to ‘r,phi,theta,t’ and then use these variable names in your function expression.

Possible values: Any string

A.64 Parameters in section Heating model/Latent heat melt

- *Parameter name:* Melting entropy change

Value: -300

Default: -300

Description: The entropy change for the phase transition from solid to melt. Units: $J/(kgK)$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

A.65 Parameters in section Heating model/Radioactive decay

- *Parameter name:* Crust composition number

Value: 0

Default: 0

Description: Which composition field should be treated as crust

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* Crust defined by composition

Value: false

Default: false

Description: Whether crust defined by composition or depth

Possible values: A boolean value (true or false)

- *Parameter name:* Crust depth

Value: 0

Default: 0

Description: Depth of the crust when crust if defined by depth. Units: m

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Half decay times

Value:

Default:

Description: Half decay times. Units: (Seconds), or (Years) if set ‘use years instead of seconds’.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Heating rates

Value:

Default:

Description: Heating rates of different elements (W/kg)

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Initial concentrations crust

Value:

Default:

Description: Initial concentrations of different elements (ppm)

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Initial concentrations mantle

Value:

Default:

Description: Initial concentrations of different elements (ppm)

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Number of elements

Value: 0

Default: 0

Description: Number of radioactive elements

Possible values: An integer n such that $0 \leq n \leq 2147483647$

A.66 Parameters in section Initial composition model

- *Parameter name:* List of model names

Value:

Default:

Description: A comma-separated list of initial composition models that together describe the initial composition field. These plugins are loaded in the order given, and modify the existing composition field via the operators listed in 'List of model operators'.

The following composition models are available:

'ascii data': Implementation of a model in which the initial composition is derived from files containing data in ascii format. Note the required format of the input data: The first lines may contain any number of comments if they begin with '#', but one of these lines needs to contain the number of grid points in each dimension as for example '# POINTS: 3 3'. The order of the data columns has to be 'x', 'y', 'composition1', 'composition2', etc. in a 2d model and 'x', 'y', 'z', 'composition1', 'composition2', etc. in a 3d model, according to the number of compositional fields, which means that there has to be a single column for every composition in the model. Note that the data in the input files need to be sorted in a specific order: the first coordinate needs to ascend first, followed by the second and the third at last in order to assign the correct data to the prescribed coordinates. If you use a spherical model, then the data will still be handled as Cartesian, however the assumed grid changes. 'x' will be replaced by the radial distance of the point to the bottom of the model, 'y' by the azimuth angle and 'z' by the polar angle measured positive from the north pole. The grid will be assumed to be a latitude-longitude grid. Note that the order of spherical coordinates is 'r', 'phi', 'theta' and not 'r', 'theta', 'phi', since this allows for dimension independent expressions.

'function': Specify the composition in terms of an explicit formula. The format of these functions follows the syntax understood by the muparser library, see Section 4.7.3.

'porosity': A class that implements initial conditions for the porosity field by computing the equilibrium melt fraction for the given initial condition and reference pressure profile. Note that this plugin only

works if there is a compositional field called ‘porosity’, and the used material model implements the ‘MeltFractionModel’ interface. For all compositional fields except porosity this plugin returns 0.0, and they are therefore not changed as long as the default ‘add’ operator is selected for this plugin.

Possible values: A comma-separated list of any of ascii data, function, porosity

- *Parameter name:* List of model operators

Value: add

Default: add

Description: A comma-separated list of operators that will be used to append the listed composition models onto the previous models. If only one operator is given, the same operator is applied to all models.

Possible values: A comma-separated list of any of add, subtract, minimum, maximum

- *Parameter name:* Model name

Value: unspecified

Default: unspecified

Description: Select one of the following models:

‘ascii data’: Implementation of a model in which the initial composition is derived from files containing data in ascii format. Note the required format of the input data: The first lines may contain any number of comments if they begin with '#', but one of these lines needs to contain the number of grid points in each dimension as for example '# POINTS: 3 3'. The order of the data columns has to be ‘x’, ‘y’, ‘composition1’, ‘composition2’, etc. in a 2d model and ‘x’, ‘y’, ‘z’, ‘composition1’, ‘composition2’, etc. in a 3d model, according to the number of compositional fields, which means that there has to be a single column for every composition in the model. Note that the data in the input files need to be sorted in a specific order: the first coordinate needs to ascend first, followed by the second and the third at last in order to assign the correct data to the prescribed coordinates. If you use a spherical model, then the data will still be handled as Cartesian, however the assumed grid changes. ‘x’ will be replaced by the radial distance of the point to the bottom of the model, ‘y’ by the azimuth angle and ‘z’ by the polar angle measured positive from the north pole. The grid will be assumed to be a latitude-longitude grid. Note that the order of spherical coordinates is ‘r’, ‘phi’, ‘theta’ and not ‘r’, ‘theta’, ‘phi’, since this allows for dimension independent expressions.

‘function’: Specify the composition in terms of an explicit formula. The format of these functions follows the syntax understood by the muparser library, see Section 4.7.3.

‘porosity’: A class that implements initial conditions for the porosity field by computing the equilibrium melt fraction for the given initial condition and reference pressure profile. Note that this plugin only works if there is a compositional field called ‘porosity’, and the used material model implements the ‘MeltFractionModel’ interface. For all compositional fields except porosity this plugin returns 0.0, and they are therefore not changed as long as the default ‘add’ operator is selected for this plugin.

Warning: This parameter provides an old and deprecated way of specifying initial composition models and shouldn’t be used. Please use ‘List of model names’ instead.

Possible values: Any one of ascii data, function, porosity, unspecified

A.67 Parameters in section Initial composition model/Ascii data model

- *Parameter name:* Data directory

Value: \$ASPECT_SOURCE_DIR/data/initial-composition/ascii-data/test/

Default: \$ASPECT_SOURCE_DIR/data/initial-composition/ascii-data/test/

Description: The name of a directory that contains the model data. This path may either be absolute (if starting with a '/') or relative to the current directory. The path may also include the special text '\$ASPECT_SOURCE_DIR' which will be interpreted as the path in which the ASPECT source files were located when ASPECT was compiled. This interpretation allows, for example, to reference files located in the 'data/' subdirectory of ASPECT.

Possible values: A directory name

- *Parameter name:* **Data file name**

Value: box_2d.txt

Default: box_2d.txt

Description: The file name of the material data. Provide file in format: (Velocity file name).%s%d where %s is a string specifying the boundary of the model according to the names of the boundary indicators (of a box or a spherical shell).%d is any sprintf integer qualifier, specifying the format of the current file number.

Possible values: Any string

- *Parameter name:* **Scale factor**

Value: 1

Default: 1

Description: Scalar factor, which is applied to the boundary velocity. You might want to use this to scale the velocities to a reference model (e.g. with free-slip boundary) or another plate reconstruction. Another way to use this factor is to convert units of the input files. The unit is assumed to be m/s or m/yr depending on the 'Use years in output instead of seconds' flag. If you provide velocities in cm/yr set this factor to 0.01.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.68 Parameters in section Initial composition model/Function

- *Parameter name:* **Coordinate system**

Value: cartesian

Default: cartesian

Description: A selection that determines the assumed coordinate system for the function variables. Allowed values are 'cartesian', 'spherical', and 'depth'. 'spherical' coordinates are interpreted as r,phi or r,phi,theta in 2D/3D respectively with theta being the polar angle. 'depth' will create a function, in which only the first parameter is non-zero, which is interpreted to be the depth of the point.

Possible values: Any one of cartesian, spherical, depth

- *Parameter name:* **Function constants**

Value:

Default:

Description: Sometimes it is convenient to use symbolic constants in the expression that describes the function, rather than having to use its numeric value everywhere the constant appears. These values can be defined using this parameter, in the form 'var1=value1, var2=value2, ...'.

A typical example would be to set this runtime parameter to 'pi=3.1415926536' and then use 'pi' in the expression of the actual formula. (That said, for convenience this class actually defines both 'pi' and 'Pi' by default, but you get the idea.)

Possible values: Any string

- *Parameter name:* Function expression

Value: 0

Default: 0

Description: The formula that denotes the function you want to evaluate for particular values of the independent variables. This expression may contain any of the usual operations such as addition or multiplication, as well as all of the common functions such as ‘sin’ or ‘cos’. In addition, it may contain expressions like ‘if($x > 0$, 1, -1)’ where the expression evaluates to the second argument if the first argument is true, and to the third argument otherwise. For a full overview of possible expressions accepted see the documentation of the muparser library at <http://muparser.beltoforion.de/>.

If the function you are describing represents a vector-valued function with multiple components, then separate the expressions for individual components by a semicolon.

Possible values: Any string

- *Parameter name:* Variable names

Value: x,y,t

Default: x,y,t

Description: The name of the variables as they will be used in the function, separated by commas. By default, the names of variables at which the function will be evaluated is ‘x’ (in 1d), ‘x,y’ (in 2d) or ‘x,y,z’ (in 3d) for spatial coordinates and ‘t’ for time. You can then use these variable names in your function expression and they will be replaced by the values of these variables at which the function is currently evaluated. However, you can also choose a different set of names for the independent variables at which to evaluate your function expression. For example, if you work in spherical coordinates, you may wish to set this input parameter to ‘r,phi,theta,t’ and then use these variable names in your function expression.

Possible values: Any string

A.69 Parameters in section Initial temperature model

- *Parameter name:* List of model names

Value:

Default:

Description: A comma-separated list of initial temperature models that will be used to initialize the temperature. These plugins are loaded in the order given, and modify the existing temperature field via the operators listed in ‘List of model operators’.

The following initial temperature models are available:

‘S40RTS perturbation’: An initial temperature field in which the temperature is perturbed following the S20RTS or S40RTS shear wave velocity model by Ritsema and others, which can be downloaded here <http://www.earth.lsa.umich.edu/~jritsema/research.html>. Information on the vs model can be found in Ritsema, J., Deuss, A., van Heijst, H.J. & Woodhouse, J.H., 2011. S40RTS: a degree-40 shear-velocity model for the mantle from new Rayleigh wave dispersion, teleseismic traveltimes and normal-mode splitting function measurements, Geophys. J. Int. 184, 1223-1236. The scaling between the shear wave perturbation and the temperature perturbation can be set by the user with the ‘Vs to density scaling’ parameter and the ‘Thermal expansion coefficient in initial temperature scaling’ parameter. The scaling is as follows: $\delta \ln \rho(r, \theta, \phi) = \xi \cdot \delta \ln v_s(r, \theta, \phi)$ and $\delta T(r, \theta, \phi) = -\frac{1}{\alpha} \delta \ln \rho(r, \theta, \phi)$. ξ is the ‘vs to density scaling’ parameter and α is the ‘Thermal expansion coefficient in initial temperature scaling’ parameter. The temperature perturbation is added to an otherwise constant temperature (incompressible model) or adiabatic reference profile (compressible model). If a depth is specified in

'Remove temperature heterogeneity down to specified depth', there is no temperature perturbation prescribed down to that depth.

'SAVANI perturbation': An initial temperature field in which the temperature is perturbed following the SAVANI shear wave velocity model by Auer and others, which can be downloaded here <http://n.ethz.ch/~auerl/savani.tar.bz2>. Information on the vs model can be found in Auer, L., Boschi, L., Becker, T.W., Nissen-Meyer, T. & Giardini, D., 2014. Savani: A variable resolution whole-mantle model of anisotropic shear velocity variations based on multiple data sets. *Journal of Geophysical Research: Solid Earth* 119.4 (2014): 3006-3034. The scaling between the shear wave perturbation and the temperature perturbation can be set by the user with the 'Vs to density scaling' parameter and the 'Thermal expansion coefficient in initial temperature scaling' parameter. The scaling is as follows: $\delta \ln \rho(r, \theta, \phi) = \xi \cdot \delta \ln v_s(r, \theta, \phi)$ and $\delta T(r, \theta, \phi) = -\frac{1}{\alpha} \delta \ln \rho(r, \theta, \phi)$. ξ is the 'vs to density scaling' parameter and α is the 'Thermal expansion coefficient in initial temperature scaling' parameter. The temperature perturbation is added to an otherwise constant temperature (incompressible model) or adiabatic reference profile (compressible model).

'adiabatic': Temperature is prescribed as an adiabatic profile with upper and lower thermal boundary layers, whose ages are given as input parameters.

'adiabatic boundary': An initial temperature condition that allows for discretizing the model domain into two layers separated by a user-defined isothermal boundary using a table look-up approach. The user includes an input ascii data file that is formatted as 3 columns of 'latitude', 'longitude', and 'depth', where 'depth' is in kilometers and represents the depth of an initial temperature of 1673.15 K (by default). The temperature is defined from the surface (273.15 K) to the isotherm as a linear gradient. Below the isotherm the temperature increases approximately adiabatically (0.0005 K per meter). This initial temperature condition is designed specifically for the ellipsoidal chunk geometry model.

'ascii data': Implementation of a model in which the initial temperature is derived from files containing data in ascii format. Note the required format of the input data: The first lines may contain any number of comments if they begin with '#', but one of these lines needs to contain the number of grid points in each dimension as for example '# POINTS: 3 3'. The order of the data columns has to be 'x', 'y', 'Temperature [K]' in a 2d model and 'x', 'y', 'z', 'Temperature [K]' in a 3d model, which means that there has to be a single column containing the temperature. Note that the data in the input files need to be sorted in a specific order: the first coordinate needs to ascend first, followed by the second and the third at last in order to assign the correct data to the prescribed coordinates. If you use a spherical model, then the data will still be handled as Cartesian, however the assumed grid changes. 'x' will be replaced by the radial distance of the point to the bottom of the model, 'y' by the azimuth angle and 'z' by the polar angle measured positive from the north pole. The grid will be assumed to be a latitude-longitude grid. Note that the order of spherical coordinates is 'r', 'phi', 'theta' and not 'r', 'theta', 'phi', since this allows for dimension independent expressions.

'function': Specify the initial temperature in terms of an explicit formula. The format of these functions follows the syntax understood by the muparser library, see Section 4.7.3.

'harmonic perturbation': An initial temperature field in which the temperature is perturbed following a harmonic function (spherical harmonic or sine depending on geometry and dimension) in lateral and radial direction from an otherwise constant temperature (incompressible model) or adiabatic reference profile (compressible model).

'inclusion shape perturbation': An initial temperature field in which there is an inclusion in a constant-temperature box field. The size, shape, gradient, position, and temperature of the inclusion are defined by parameters.

'mandelbox': Fractal-shaped temperature field.

'perturbed box': An initial temperature field in which the temperature is perturbed slightly from an otherwise constant value equal to one. The perturbation is chosen in such a way that the initial

temperature is constant to one along the entire boundary.

‘polar box’: An initial temperature field in which the temperature is perturbed slightly from an otherwise constant value equal to one. The perturbation is such that there are two poles on opposing corners of the box.

‘spherical gaussian perturbation’: An initial temperature field in which the temperature is perturbed by a single Gaussian added to an otherwise spherically symmetric state. Additional parameters are read from the parameter file in subsection ‘Spherical gaussian perturbation’.

‘spherical hexagonal perturbation’: An initial temperature field in which the temperature is perturbed following an N -fold pattern in a specified direction from an otherwise spherically symmetric state. The class’s name comes from previous versions when the only option was $N = 6$.

Possible values: A comma-separated list of any of S40RTS perturbation, SAVANI perturbation, adiabatic, adiabatic boundary, ascii data, function, harmonic perturbation, inclusion shape perturbation, mandelbox, perturbed box, polar box, spherical gaussian perturbation, spherical hexagonal perturbation

- *Parameter name:* List of model operators

Value: add

Default: add

Description: A comma-separated list of operators that will be used to append the listed temperature models onto the previous models. If only one operator is given, the same operator is applied to all models.

Possible values: A comma-separated list of any of add, subtract, minimum, maximum

- *Parameter name:* Model name

Value: adiabatic

Default: unspecified

Description: Select one of the following models:

‘S40RTS perturbation’: An initial temperature field in which the temperature is perturbed following the S20RTS or S40RTS shear wave velocity model by Ritsema and others, which can be downloaded here <http://www.earth.lsa.umich.edu/~jritsema/research.html>. Information on the vs model can be found in Ritsema, J., Deuss, A., van Heijst, H.J. & Woodhouse, J.H., 2011. S40RTS: a degree-40 shear-velocity model for the mantle from new Rayleigh wave dispersion, teleseismic traveltimes and normal-mode splitting function measurements, Geophys. J. Int. 184, 1223-1236. The scaling between the shear wave perturbation and the temperature perturbation can be set by the user with the ‘Vs to density scaling’ parameter and the ‘Thermal expansion coefficient in initial temperature scaling’ parameter. The scaling is as follows: $\delta ln\rho(r, \theta, \phi) = \xi \cdot \delta ln v_s(r, \theta, \phi)$ and $\delta T(r, \theta, \phi) = -\frac{1}{\alpha} \delta ln\rho(r, \theta, \phi)$. ξ is the ‘vs to density scaling’ parameter and α is the ‘Thermal expansion coefficient in initial temperature scaling’ parameter. The temperature perturbation is added to an otherwise constant temperature (incompressible model) or adiabatic reference profile (compressible model). If a depth is specified in ‘Remove temperature heterogeneity down to specified depth’, there is no temperature perturbation prescribed down to that depth.

‘SAVANI perturbation’: An initial temperature field in which the temperature is perturbed following the SAVANI shear wave velocity model by Auer and others, which can be downloaded here <http://n.ethz.ch/~auer1/savani.tar.bz2>. Information on the vs model can be found in Auer, L., Boschi, L., Becker, T.W., Nissen-Meyer, T. & Giardini, D., 2014. Savani: A variable resolution whole-mantle model of anisotropic shear velocity variations based on multiple data sets. Journal of Geophysical Research: Solid Earth 119.4 (2014): 3006-3034. The scaling between the shear wave perturbation and the temperature perturbation can be set by the user with the ‘Vs to density scaling’ parameter

and the 'Thermal expansion coefficient in initial temperature scaling' parameter. The scaling is as follows: $\delta ln\rho(r, \theta, \phi) = \xi \cdot \delta ln v_s(r, \theta, \phi)$ and $\delta T(r, \theta, \phi) = -\frac{1}{\alpha} \delta ln\rho(r, \theta, \phi)$. ξ is the 'vs to density scaling' parameter and α is the 'Thermal expansion coefficient in initial temperature scaling' parameter. The temperature perturbation is added to an otherwise constant temperature (incompressible model) or adiabatic reference profile (compressible model).

'adiabatic': Temperature is prescribed as an adiabatic profile with upper and lower thermal boundary layers, whose ages are given as input parameters.

'adiabatic boundary': An initial temperature condition that allows for discretizing the model domain into two layers separated by a user-defined isothermal boundary using a table look-up approach. The user includes an input ascii data file that is formatted as 3 columns of 'latitude', 'longitude', and 'depth', where 'depth' is in kilometers and represents the depth of an initial temperature of 1673.15 K (by default). The temperature is defined from the surface (273.15 K) to the isotherm as a linear gradient. Below the isotherm the temperature increases approximately adiabatically (0.0005 K per meter). This initial temperature condition is designed specifically for the ellipsoidal chunk geometry model.

'ascii data': Implementation of a model in which the initial temperature is derived from files containing data in ascii format. Note the required format of the input data: The first lines may contain any number of comments if they begin with '#', but one of these lines needs to contain the number of grid points in each dimension as for example '# POINTS: 3 3'. The order of the data columns has to be 'x', 'y', 'Temperature [K]' in a 2d model and 'x', 'y', 'z', 'Temperature [K]' in a 3d model, which means that there has to be a single column containing the temperature. Note that the data in the input files need to be sorted in a specific order: the first coordinate needs to ascend first, followed by the second and the third at last in order to assign the correct data to the prescribed coordinates. If you use a spherical model, then the data will still be handled as Cartesian, however the assumed grid changes. 'x' will be replaced by the radial distance of the point to the bottom of the model, 'y' by the azimuth angle and 'z' by the polar angle measured positive from the north pole. The grid will be assumed to be a latitude-longitude grid. Note that the order of spherical coordinates is 'r', 'phi', 'theta' and not 'r', 'theta', 'phi', since this allows for dimension independent expressions.

'function': Specify the initial temperature in terms of an explicit formula. The format of these functions follows the syntax understood by the muparser library, see Section 4.7.3.

'harmonic perturbation': An initial temperature field in which the temperature is perturbed following a harmonic function (spherical harmonic or sine depending on geometry and dimension) in lateral and radial direction from an otherwise constant temperature (incompressible model) or adiabatic reference profile (compressible model).

'inclusion shape perturbation': An initial temperature field in which there is an inclusion in a constant-temperature box field. The size, shape, gradient, position, and temperature of the inclusion are defined by parameters.

'mandelbox': Fractal-shaped temperature field.

'perturbed box': An initial temperature field in which the temperature is perturbed slightly from an otherwise constant value equal to one. The perturbation is chosen in such a way that the initial temperature is constant to one along the entire boundary.

'polar box': An initial temperature field in which the temperature is perturbed slightly from an otherwise constant value equal to one. The perturbation is such that there are two poles on opposing corners of the box.

'spherical gaussian perturbation': An initial temperature field in which the temperature is perturbed by a single Gaussian added to an otherwise spherically symmetric state. Additional parameters are read from the parameter file in subsection 'Spherical gaussian perturbation'.

‘spherical hexagonal perturbation’: An initial temperature field in which the temperature is perturbed following an N -fold pattern in a specified direction from an otherwise spherically symmetric state. The class’s name comes from previous versions when the only option was $N = 6$.

Warning: This parameter provides an old and deprecated way of specifying initial temperature models and shouldn’t be used. Please use ‘List of model names’ instead.

Possible values: Any one of S40RTS perturbation, SAVANI perturbation, adiabatic, adiabatic boundary, ascii data, function, harmonic perturbation, inclusion shape perturbation, mandelbox, perturbed box, polar box, spherical gaussian perturbation, spherical hexagonal perturbation, unspecified

A.70 Parameters in section Initial temperature model/Adiabatic

- *Parameter name:* Age bottom boundary layer

Value: 0e0

Default: 0e0

Description: The age of the lower thermal boundary layer, used for the calculation of the half-space cooling model temperature. Units: years if the ‘Use years in output instead of seconds’ parameter is set; seconds otherwise.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Age top boundary layer

Value: 0e0

Default: 0e0

Description: The age of the upper thermal boundary layer, used for the calculation of the half-space cooling model temperature. Units: years if the ‘Use years in output instead of seconds’ parameter is set; seconds otherwise.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Amplitude

Value: 0e0

Default: 0e0

Description: The amplitude (in K) of the initial spherical temperature perturbation at the bottom of the model domain. This perturbation will be added to the adiabatic temperature profile, but not to the bottom thermal boundary layer. Instead, the maximum of the perturbation and the bottom boundary layer temperature will be used.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Position

Value: center

Default: center

Description: Where the initial temperature perturbation should be placed. If ‘center’ is given, then the perturbation will be centered along a ‘midpoint’ of some sort of the bottom boundary. For example, in the case of a box geometry, this is the center of the bottom face; in the case of a spherical shell geometry, it is along the inner surface halfway between the bounding radial lines.

Possible values: Any one of center

- *Parameter name:* Radius

Value: 0e0

Default: 0e0

Description: The Radius (in m) of the initial spherical temperature perturbation at the bottom of the model domain.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Subadiabaticity

Value: 0e0

Default: 0e0

Description: If this value is larger than 0, the initial temperature profile will not be adiabatic, but subadiabatic. This value gives the maximal deviation from adiabaticity. Set to 0 for an adiabatic temperature profile. Units: K.

The function object in the Function subsection represents the compositional fields that will be used as a reference profile for calculating the thermal diffusivity. This function is one-dimensional and depends only on depth. The format of this functions follows the syntax understood by the muparser library, see Section 4.7.3.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.71 Parameters in section Initial temperature model/Adiabatic/Function

- *Parameter name:* Function constants

Value:

Default:

Description: Sometimes it is convenient to use symbolic constants in the expression that describes the function, rather than having to use its numeric value everywhere the constant appears. These values can be defined using this parameter, in the form ‘var1=value1, var2=value2, ...’.

A typical example would be to set this runtime parameter to ‘pi=3.1415926536’ and then use ‘pi’ in the expression of the actual formula. (That said, for convenience this class actually defines both ‘pi’ and ‘Pi’ by default, but you get the idea.)

Possible values: Any string

- *Parameter name:* Function expression

Value: 0

Default: 0

Description: The formula that denotes the function you want to evaluate for particular values of the independent variables. This expression may contain any of the usual operations such as addition or multiplication, as well as all of the common functions such as ‘sin’ or ‘cos’. In addition, it may contain expressions like ‘if(x>0, 1, -1)’ where the expression evaluates to the second argument if the first argument is true, and to the third argument otherwise. For a full overview of possible expressions accepted see the documentation of the muparser library at <http://muparser.beltoforion.de/>.

If the function you are describing represents a vector-valued function with multiple components, then separate the expressions for individual components by a semicolon.

Possible values: Any string

- *Parameter name:* **Variable names**

Value: x,t

Default: x,t

Description: The name of the variables as they will be used in the function, separated by commas. By default, the names of variables at which the function will be evaluated is ‘x’ (in 1d), ‘x,y’ (in 2d) or ‘x,y,z’ (in 3d) for spatial coordinates and ‘t’ for time. You can then use these variable names in your function expression and they will be replaced by the values of these variables at which the function is currently evaluated. However, you can also choose a different set of names for the independent variables at which to evaluate your function expression. For example, if you work in spherical coordinates, you may wish to set this input parameter to ‘r,phi,theta,t’ and then use these variable names in your function expression.

Possible values: Any string

A.72 Parameters in section Initial temperature model/Adiabatic boundary

- *Parameter name:* **Adiabatic temperature gradient**

Value: 0.0005

Default: 0.0005

Description: The value of the adiabatic temperature gradient. Units: Km^{-1} .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* **Data directory**

Value: \$ASPECT_SOURCE_DIR/data/initial-temperature/adiabatic-boundary/

Default: \$ASPECT_SOURCE_DIR/data/initial-temperature/adiabatic-boundary/

Description: The path to the isotherm depth data file

Possible values: A directory name

- *Parameter name:* **Isotherm depth filename**

Value: adiabatic_boundary.txt

Default: adiabatic_boundary.txt

Description: File from which the isotherm depth data is read.

Possible values: an input filename

- *Parameter name:* **Isotherm temperature**

Value: 1673.15

Default: 1673.15

Description: The value of the isothermal boundary temperature. Units: Kelvin.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* **Surface temperature**

Value: 273.15

Default: 273.15

Description: The value of the surface temperature. Units: Kelvin.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.73 Parameters in section Initial temperature model/Ascii data model

- *Parameter name:* Data directory

Value: \$ASPECT_SOURCE_DIR/data/initial-temperature/ascii-data/test/

Default: \$ASPECT_SOURCE_DIR/data/initial-temperature/ascii-data/test/

Description: The name of a directory that contains the model data. This path may either be absolute (if starting with a '/') or relative to the current directory. The path may also include the special text '\$ASPECT_SOURCE_DIR' which will be interpreted as the path in which the ASPECT source files were located when ASPECT was compiled. This interpretation allows, for example, to reference files located in the 'data/' subdirectory of ASPECT.

Possible values: A directory name

- *Parameter name:* Data file name

Value: box_2d.txt

Default: box_2d.txt

Description: The file name of the material data. Provide file in format: (Velocity file name).%s%d where %s is a string specifying the boundary of the model according to the names of the boundary indicators (of a box or a spherical shell).%d is any sprintf integer qualifier, specifying the format of the current file number.

Possible values: Any string

- *Parameter name:* Scale factor

Value: 1

Default: 1

Description: Scalar factor, which is applied to the boundary velocity. You might want to use this to scale the velocities to a reference model (e.g. with free-slip boundary) or another plate reconstruction. Another way to use this factor is to convert units of the input files. The unit is assumed to be m/s or m/yr depending on the 'Use years in output instead of seconds' flag. If you provide velocities in cm/yr set this factor to 0.01.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.74 Parameters in section Initial temperature model/Function

- *Parameter name:* Coordinate system

Value: cartesian

Default: cartesian

Description: A selection that determines the assumed coordinate system for the function variables. Allowed values are 'cartesian', 'spherical', and 'depth'. 'spherical' coordinates are interpreted as r,phi or r,phi,theta in 2D/3D respectively with theta being the polar angle. 'depth' will create a function, in which only the first parameter is non-zero, which is interpreted to be the depth of the point.

Possible values: Any one of cartesian, spherical, depth

- *Parameter name:* Function constants

Value:

Default:

Description: Sometimes it is convenient to use symbolic constants in the expression that describes the function, rather than having to use its numeric value everywhere the constant appears. These values can be defined using this parameter, in the form 'var1=value1, var2=value2, ...'.

A typical example would be to set this runtime parameter to ‘pi=3.1415926536’ and then use ‘pi’ in the expression of the actual formula. (That said, for convenience this class actually defines both ‘pi’ and ‘Pi’ by default, but you get the idea.)

Possible values: Any string

- *Parameter name:* Function expression

Value: 0

Default: 0

Description: The formula that denotes the function you want to evaluate for particular values of the independent variables. This expression may contain any of the usual operations such as addition or multiplication, as well as all of the common functions such as ‘sin’ or ‘cos’. In addition, it may contain expressions like ‘if($x > 0$, 1, -1)’ where the expression evaluates to the second argument if the first argument is true, and to the third argument otherwise. For a full overview of possible expressions accepted see the documentation of the muparser library at <http://muparser.beltoforion.de/>.

If the function you are describing represents a vector-valued function with multiple components, then separate the expressions for individual components by a semicolon.

Possible values: Any string

- *Parameter name:* Variable names

Value: x,y,t

Default: x,y,t

Description: The name of the variables as they will be used in the function, separated by commas. By default, the names of variables at which the function will be evaluated is ‘x’ (in 1d), ‘x,y’ (in 2d) or ‘x,y,z’ (in 3d) for spatial coordinates and ‘t’ for time. You can then use these variable names in your function expression and they will be replaced by the values of these variables at which the function is currently evaluated. However, you can also choose a different set of names for the independent variables at which to evaluate your function expression. For example, if you work in spherical coordinates, you may wish to set this input parameter to ‘r,phi,theta,t’ and then use these variable names in your function expression.

Possible values: Any string

A.75 Parameters in section Initial temperature model/Harmonic perturbation

- *Parameter name:* Lateral wave number one

Value: 3

Default: 3

Description: Doubled first lateral wave number of the harmonic perturbation. Equals the spherical harmonic degree in 3D spherical shells. In all other cases one equals half of a sine period over the model domain. This allows for single up-/downswings. Negative numbers reverse the sign of the perturbation but are not allowed for the spherical harmonic case.

Possible values: An integer n such that $-2147483648 \leq n \leq 2147483647$

- *Parameter name:* Lateral wave number two

Value: 2

Default: 2

Description: Doubled second lateral wave number of the harmonic perturbation. Equals the spherical harmonic order in 3D spherical shells. In all other cases one equals half of a sine period over the model domain. This allows for single up-/downswings. Negative numbers reverse the sign of the perturbation.

Possible values: An integer n such that $-2147483648 \leq n \leq 2147483647$

- *Parameter name:* `Magnitude`

Value: 1.0

Default: 1.0

Description: The magnitude of the Harmonic perturbation.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `Reference temperature`

Value: 1600.0

Default: 1600.0

Description: The reference temperature that is perturbed by the harmonic function. Only used in incompressible models.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `Vertical wave number`

Value: 1

Default: 1

Description: Doubled radial wave number of the harmonic perturbation. One equals half of a sine period over the model domain. This allows for single up-/downswings. Negative numbers reverse the sign of the perturbation.

Possible values: An integer n such that $-2147483648 \leq n \leq 2147483647$

A.76 Parameters in section Initial temperature model/Inclusion shape perturbation

- *Parameter name:* `Ambient temperature`

Value: 1.0

Default: 1.0

Description: The background temperature for the temperature field.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `Center X`

Value: 0.5

Default: 0.5

Description: The X coordinate for the center of the shape.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `Center Y`

Value: 0.5

Default: 0.5

Description: The Y coordinate for the center of the shape.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `Center Z`

Value: 0.5

Default: 0.5

Description: The Z coordinate for the center of the shape. This is only necessary for three-dimensional fields.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `Inclusion gradient`

Value: constant

Default: constant

Description: The gradient of the inclusion to be generated.

Possible values: Any one of gaussian, linear, constant

- *Parameter name:* `Inclusion shape`

Value: circle

Default: circle

Description: The shape of the inclusion to be generated.

Possible values: Any one of square, circle

- *Parameter name:* `Inclusion temperature`

Value: 0.0

Default: 0.0

Description: The temperature of the inclusion shape. This is only the true temperature in the case of the constant gradient. In all other cases, it gives one endpoint of the temperature gradient for the shape.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `Shape radius`

Value: 1.0

Default: 1.0

Description: The radius of the inclusion to be generated. For shapes with no radius (e.g. square), this will be the width, and for shapes with no width, this gives a general guideline for the size of the shape.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.77 Parameters in section Initial temperature model/S4ORTS perturbation

- *Parameter name:* `Data directory`

Value: \$ASPECT_SOURCE_DIR/data/initial-temperature/S4ORTS/

Default: \$ASPECT_SOURCE_DIR/data/initial-temperature/S4ORTS/

Description: The path to the model data.

Possible values: A directory name

- *Parameter name:* Initial condition file name

Value: S40RTS.sph

Default: S40RTS.sph

Description: The file name of the spherical harmonics coefficients from Ritsema et al.

Possible values: Any string

- *Parameter name:* Maximum order

Value: 20

Default: 20

Description: The maximum order the users specify when reading the data file of spherical harmonic coefficients, which must be smaller than the maximum order the data file stored. This parameter will be used only if 'Specify a lower maximum order' is set to true.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* Reference temperature

Value: 1600.0

Default: 1600.0

Description: The reference temperature that is perturbed by the spherical harmonic functions. Only used in incompressible models.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Remove degree 0 from perturbation

Value: true

Default: true

Description: Option to remove the degree zero component from the perturbation, which will ensure that the laterally averaged temperature for a fixed depth is equal to the background temperature.

Possible values: A boolean value (true or false)

- *Parameter name:* Remove temperature heterogeneity down to specified depth

Value: -1.7976931348623157e+308

Default: -1.7976931348623157e+308

Description: This will set the heterogeneity prescribed by S20RTS or S40RTS to zero down to the specified depth (in meters). Note that your resolution has to be adequate to capture this cutoff. For example if you specify a depth of 660km, but your closest spherical depth layers are only at 500km and 750km (due to a coarse resolution) it will only zero out heterogeneities down to 500km. Similar caution has to be taken when using adaptive meshing.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Specify a lower maximum order

Value: false

Default: false

Description: Option to use a lower maximum order when reading the data file of spherical harmonic coefficients. This is probably used for the faster tests or when the users only want to see the spherical harmonic pattern up to a certain order.

Possible values: A boolean value (true or false)

- *Parameter name:* Spline knots depth file name
Value: Spline_knots.txt
Default: Spline_knots.txt
Description: The file name of the spline knot locations from Ritsema et al.
Possible values: Any string
- *Parameter name:* Thermal expansion coefficient in initial temperature scaling
Value: 2e-5
Default: 2e-5
Description: The value of the thermal expansion coefficient β . Units: $1/K$.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Vs to density scaling
Value: 0.25
Default: 0.25
Description: This parameter specifies how the perturbation in shear wave velocity as prescribed by S20RTS or S40RTS is scaled into a density perturbation. See the general description of this model for more detailed information.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.78 Parameters in section Initial temperature model/SAVANI perturbation

- *Parameter name:* Data directory
Value: \$ASPECT_SOURCE_DIR/data/initial-temperature/SAVANI/
Default: \$ASPECT_SOURCE_DIR/data/initial-temperature/SAVANI/
Description: The path to the model data.
Possible values: A directory name
- *Parameter name:* Initial condition file name
Value: savani.dlnvs.60.m.ab
Default: savani.dlnvs.60.m.ab
Description: The file name of the spherical harmonics coefficients from Auer et al.
Possible values: Any string
- *Parameter name:* Maximum order
Value: 20
Default: 20
Description: The maximum order the users specify when reading the data file of spherical harmonic coefficients, which must be smaller than the maximum order the data file stored. This parameter will be used only if 'Specify a lower maximum order' is set to true.
Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* Reference temperature

Value: 1600.0

Default: 1600.0

Description: The reference temperature that is perturbed by the spherical harmonic functions. Only used in incompressible models.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Remove degree 0 from perturbation

Value: true

Default: true

Description: Option to remove the degree zero component from the perturbation, which will ensure that the laterally averaged temperature for a fixed depth is equal to the background temperature.

Possible values: A boolean value (true or false)

- *Parameter name:* Remove temperature heterogeneity down to specified depth

Value: -1.7976931348623157e+308

Default: -1.7976931348623157e+308

Description: This will set the heterogeneity prescribed by SAVANI to zero down to the specified depth (in meters). Note that your resolution has to be adequate to capture this cutoff. For example if you specify a depth of 660km, but your closest spherical depth layers are only at 500km and 750km (due to a coarse resolution) it will only zero out heterogeneities down to 500km. Similar caution has to be taken when using adaptive meshing.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Specify a lower maximum order

Value: false

Default: false

Description: Option to use a lower maximum order when reading the data file of spherical harmonic coefficients. This is probably used for the faster tests or when the users only want to see the spherical harmonic pattern up to a certain order.

Possible values: A boolean value (true or false)

- *Parameter name:* Spline knots depth file name

Value: Spline_knots.txt

Default: Spline_knots.txt

Description: The file name of the spline knots taken from the 28 spherical layers of SAVANI tomography model.

Possible values: Any string

- *Parameter name:* Thermal expansion coefficient in initial temperature scaling

Value: 2e-5

Default: 2e-5

Description: The value of the thermal expansion coefficient β . Units: $1/K$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: `Vs to density scaling`

Value: 0.25

Default: 0.25

Description: This parameter specifies how the perturbation in shear wave velocity as prescribed by SAVANI is scaled into a density perturbation. See the general description of this model for more detailed information.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.79 Parameters in section Initial temperature model/Spherical gaussian perturbation

- Parameter name: `Amplitude`

Value: 0.01

Default: 0.01

Description: The amplitude of the perturbation.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: `Angle`

Value: 0e0

Default: 0e0

Description: The angle where the center of the perturbation is placed.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: `Filename for initial geotherm table`

Value: initial-geotherm-table

Default: initial-geotherm-table

Description: The file from which the initial geotherm table is to be read. The format of the file is defined by what is read in source/initial_conditions/spherical_shell.cc.

Possible values: an input filename

- Parameter name: `Non-dimensional depth`

Value: 0.7

Default: 0.7

Description: The non-dimensional radial distance where the center of the perturbation is placed.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: `Sigma`

Value: 0.2

Default: 0.2

Description: The standard deviation of the Gaussian perturbation.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: `Sign`

Value: 1

Default: 1

Description: The sign of the perturbation.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

A.80 Parameters in section Initial temperature model/Spherical hexagonal perturbation

- *Parameter name:* Angular mode

Value: 6

Default: 6

Description: The number of convection cells with which to perturb the system.

Possible values: An integer n such that $-2147483648 \leq n \leq 2147483647$

- *Parameter name:* Rotation offset

Value: -45

Default: -45

Description: Amount of clockwise rotation in degrees to apply to the perturbations. Default is set to -45 in order to provide backwards compatibility.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

A.81 Parameters in section Material model

- *Parameter name:* Material averaging

Value: none

Default: none

Description: Whether or not (and in the first case, how) to do any averaging of material model output data when constructing the linear systems for velocity/pressure, temperature, and compositions in each time step, as well as their corresponding preconditioners.

Possible choices: none|arithmetic average|harmonic average|geometric average|pick largest|project to Q1|log average

The process of averaging, and where it may be used, is discussed in more detail in Section 5.2.8.

More averaging schemes are available in the averaging material model. This material model is a “compositing material model” which can be used in combination with other material models.

Possible values: Any one of none, arithmetic average, harmonic average, geometric average, pick largest, project to Q1, log average

- *Parameter name:* Model name

Value: simple

Default: unspecified

Description: The name of the material model to be used in this simulation. There are many material models you can choose from, as listed below. They generally fall into two category: (i) models that implement a particular case of material behavior, (ii) models that modify other models in some way. We sometimes call the latter “compositing models”. An example of a compositing model is the “depth dependent” model below in that it takes another, freely choosable model as its base and then modifies that model’s output in some way.

You can select one of the following models:

‘Steinberger’: This material model looks up the viscosity from the tables that correspond to the paper of Steinberger and Calderwood 2006 (“Models of large-scale viscous flow in the Earth’s mantle with constraints from mineral physics and surface observations”, Geophys. J. Int., 167, 1461-1481, <http://dx.doi.org/10.1111/j.1365-246X.2006.03131.x>) and material data from a database generated by the thermodynamics code PerpleX, see <http://www.perplex.ethz.ch/>. The default example data

builds upon the thermodynamic database by Stixrude 2011 and assumes a pyrolytic composition by Ringwood 1988 but is easily replaceable by other data files.

‘ascii reference profile’: A material model that reads in a reference state for density, thermal expansivity, compressibility and specific heat from a text file. Note the required format of the input data: The first lines may contain any number of comments if they begin with ‘#’, but one of these lines needs to contain the number of points in the reference state as for example ‘# POINTS: 3’. Following the comment lines there has to be a single line containing the names of all data columns, separated by arbitrarily many spaces. Column names are not allowed to contain spaces. The file can contain unnecessary columns, but for this plugin it needs to at least provide the columns named ‘density’, ‘thermal_expansivity’, ‘specific_heat’, and ‘compressibility’. Note that the data lines in the file need to be sorted in order of increasing depth from 0 to the maximal depth in the model domain. Points in the model that are outside of the provided depth range will be assigned the maximum or minimum depth values, respectively. Points do not need to be equidistant, but the computation of properties is optimized in speed if they are.

The viscosity η is computed as

$$\eta(z, T) = \eta_r(z)\eta_0 \exp\left(-A \frac{T - T_{\text{adi}}}{T_{\text{adi}}}\right), \quad (90)$$

where $\eta_r(z)$ is the depth-dependence, which is a piecewise constant function computed according to the list of “Viscosity prefactors” and “Transition depths”, η_0 is the reference viscosity specified by the parameter “Viscosity” and A describes the dependence on temperature and corresponds to the parameter “Thermal viscosity exponent”.

‘averaging’: The “averaging” Material model applies an averaging of the quadrature points within a cell. The values to average are supplied by any of the other available material models. In other words, it is a “compositing material model”. Parameters related to the average model are read from a subsection “Material model/Averaging”.

The user must specify a “Base model” from which material properties are derived. Furthermore an averaging operation must be selected, where the Choice should be from the list none|arithmetic average|harmonic average|geometric average|pick largest|log average|NWD arithmetic average|NWD harmonic average|NWD geometric average.

NWD stands for Normalized Weighed Distance. The models with this in front of their name work with a weighed average, which means each quadrature point requires an individual weight. The weight is determined by the distance, where the exact relation is determined by a bell shaped curve. A bell shaped curve is a continuous function which is one at its maximum and exactly zero at and beyond its limit. This bell shaped curve is spanned around each quadrature point to determine the weighting map for each quadrature point. The used bell shape comes from Lucy (1977). The distance is normalized so the largest distance becomes one. This means that if variable “Bell shape limit” is exactly one, the farthest quadrature point is just on the limit and its weight will be exactly zero. In this plugin it is not implemented as larger and equal than the limit, but larger than, to ensure the quadrature point at distance zero is always included.

‘compositing’: The “compositing” Material model selects material model properties from a given set of other material models, and is intended to make mixing different material models easier.

Specifically, this material model works by allowing to specify the name of another material model for each coefficient that material models are asked for (such as the viscosity, density, etc.). Whenever the material model is asked for the values of coefficients, it then evaluates all of the “base models” that were listed for the various coefficients, and copies the values returned by these base models into the output structure.

The implementation of this material model is somewhat expensive because it has to evaluate all material coefficients of all underlying material models. Consequently, if performance of assembly and

postprocessing is important, then implementing a separate material model is a better choice than using this material model.

‘composition reaction’: A material model that behaves in the same way as the simple material model, but includes two compositional fields and a reaction between them. Above a depth given in the input file, the first fields gets converted to the second field.

‘depth dependent’: The “depth dependent” Material model applies a depth-dependent scaling to any of the other available material models. In other words, it is a “compositing material model”.

Parameters related to the depth dependent model are read from a subsection “Material model/Depth dependent model”. The user must specify a “Base model” from which material properties are derived. Currently the depth dependent model only allows depth dependence of viscosity - other material properties are taken from the “Base model”. Viscosity η at depth z is calculated according to:

$$\eta(z, p, T, X, \dots) = \eta(z)\eta_b(p, T, X, \dots)/\eta_{rb} \quad (91)$$

where $\eta(z)$ is the depth-dependence specified by the depth dependent model, $\eta_b(p, T, X, \dots)$ is the viscosity calculated from the base model, and η_{rb} is the reference viscosity of the “Base model”. In addition to the specification of the “Base model”, the user must specify the method to be used to calculate the depth-dependent viscosity $\eta(z)$ as “Material model/Depth dependent model/Depth dependence method”, which can be chosen among “None|Function|File|List”. Each method and the associated parameters are as follows:

“Function”: read a user-specified parsed function from the input file in a subsection “Material model/Depth dependent model/Viscosity depth function”. By default, this function is uniformly equal to 1.0e21. Specifying a function that returns a value less than or equal to 0.0 anywhere in the model domain will produce an error.

“File”: read a user-specified file containing viscosity values at specified depths. The file containing depth-dependent viscosities is read from a directory specified by the user as “Material model/Depth dependent model/Data directory”, from a file with name specified as “Material model/Depth dependent model/Viscosity depth file”. The format of this file is ascii text and contains two columns with one header line:

example Viscosity depth file:

```
Depth (m) Viscosity (Pa-s)
0.0000000e+00 1.0000000e+21
6.7000000e+05 1.0000000e+22
```

Viscosity is interpolated from this file using linear interpolation. “None”: no depth-dependence. Viscosity is taken directly from “Base model”

“List”: read a comma-separated list of depth values corresponding to the maximum depths of layers having constant depth-dependence $\eta(z)$. The layers must be specified in order of increasing depth, and the last layer in the list must have a depth greater than or equal to the maximal depth of the model. The list of layer depths is specified as “Material model/Depth dependent model/Depth list” and the corresponding list of layer viscosities is specified as “Material model/Depth dependent model/Viscosity list”

‘diffusion dislocation’: An implementation of a viscous rheology including diffusion and dislocation creep. Compositional fields can each be assigned individual activation energies, reference densities, thermal expansivities, and stress exponents. The effective viscosity is defined as

$$\eta_{\text{eff}} = \left(\frac{1}{\eta_{\text{eff}}^{\text{diff}}} + \frac{1}{\eta_{\text{eff}}^{\text{dis}}} \right)^{-1}$$

where

$$\eta_i = 0.5 A_i^{-\frac{1}{n_i}} d^{\frac{m_i}{n_i}} \dot{\varepsilon}_i^{\frac{1-n_i}{n_i}} \exp\left(\frac{E_i^* + PV_i^*}{n_i R T}\right)$$

where d is grain size, i corresponds to diffusion or dislocation creep, $\dot{\varepsilon}$ is the square root of the second invariant of the strain rate tensor, R is the gas constant, T is temperature, and P is pressure. A_i are prefactors, n_i and m_i are stress and grain size exponents E_i are the activation energies and V_i are the activation volumes.

This form of the viscosity equation is commonly used in geodynamic simulations. See, for example, Billen and Hirth (2007), G3, 8, Q08012. Significantly, other studies may use slightly different forms of the viscosity equation leading to variations in how specific terms are defined or combined. For example, the grain size exponent should always be positive in the diffusion viscosity equation used here, while other studies place the grain size term in the denominator and invert the sign of the grain size exponent. When examining previous work, one should carefully check how the viscous prefactor and grain size terms are defined.

The ratio of diffusion to dislocation strain rate is found by Newton's method, iterating to find the stress which satisfies the above equations. The value for the components of this formula and additional parameters are read from the parameter file in subsection 'Material model/DiffusionDislocation'.

'drucker prager': A material model that has constant values for all coefficients but the density and viscosity. The defaults for all coefficients are chosen to be similar to what is believed to be correct for Earth's mantle. All of the values that define this model are read from a section "Material model/-Drucker Prager" in the input file, see Section A.90. Note that the model does not take into account any dependencies of material properties on compositional fields.

The viscosity is computed according to the Drucker Prager frictional plasticity criterion (non-associative) based on a user-defined internal friction angle ϕ and cohesion C . In 3D: $\sigma_y = \frac{6C \cos(\phi)}{\sqrt{(3)(3+\sin(\phi))}} + \frac{2P \sin(\phi)}{\sqrt{(3)(3+\sin(\phi))}}$, where P is the pressure. See for example Zienkiewicz, O. C., Humphezon, C. and Lewis, R. W. (1975), Géotechnique 25, No. 4, 671-689. With this formulation we circumscribe instead of inscribe the Mohr Coulomb yield surface. In 2D the Drucker Prager yield surface is the same as the Mohr Coulomb surface: $\sigma_y = P \sin(\phi) + C \cos(\phi)$. Note that in 2D for $\phi = 0$, these criteria revert to the von Mises criterion (no pressure dependence). See for example Thieulot, C. (2011), PEPI 188, 47-68.

Note that we enforce the pressure to be positive to prevent negative yield strengths and viscosities.

We then use the computed yield strength to scale back the viscosity on to the yield surface using the Viscosity Rescaling Method described in Kachanov, L. M. (2004), Fundamentals of the Theory of Plasticity, Dover Publications, Inc. (Not Radial Return.) A similar implementation can be found in GALE (<https://geodynamics.org/cig/software/gale/gale-manual.pdf>).

To avoid numerically unfavourably large (or even negative) viscosity ranges, we cut off the viscosity with a user-defined minimum and maximum viscosity: $\eta_{eff} = \frac{1}{\eta_m^{1/n} + \eta_a^{1/\alpha}}$.

Note that this model uses the formulation that assumes an incompressible medium despite the fact that the density follows the law $\rho(T) = \rho_0(1 - \beta(T - T_{ref}))$.

'dynamic friction': This model is for use with an arbitrary number of compositional fields, where each field represents a rock type which can have completely different properties from the others. Each rock type itself has constant material properties, with the exception of viscosity which is modified according to a Drucker-Prager yield criterion. Unlike the drucker prager or visco plastic material models, the angle of internal friction is a function of velocity. This relationship is similar to rate-and-state friction constitutive relationships, which are applicable to the strength of rocks during earthquakes. The formulation used here is derived from van Dinther et al. 2013, JGR. Each compositional field is interpreted as a volume fraction. If the sum of the fields is greater than one, they are renormalized. If

it is less than one, material properties for “background material” make up the rest. When more than one field is present, the material properties are averaged arithmetically. An exception is the viscosity, where the averaging should make more of a difference. For this, the user selects between arithmetic, harmonic, geometric, or maximum composition averaging.

‘grain size’: A material model that relies on compositional fields that stand for average grain sizes of a mineral phase and source terms for them that determine the grain size evolution in dependence of the strain rate, temperature, phase transitions, and the creep regime. This material model only works if a compositional field named ‘grain_size’ is present. In the diffusion creep regime, the viscosity depends on this grain size field. We use the grain size evolution laws described in Behn et al., 2009. Implications of grain size evolution on the seismic structure of the oceanic upper mantle, *Earth Planet. Sci. Letters*, 282, 178–189. Other material parameters are either prescribed similar to the ‘simple’ material model, or read from data files that were generated by the PerpleX or Hefesto software. The material model is described in more detail in Dannberg, J., Z. Eilon, U. Faul, R. Gassmoeller, P. Moulik, and R. Myhill (2017), The importance of grain size to mantle dynamics and seismological observations, *Geochem. Geophys. Geosyst.*, 18, 3034–3061, doi:10.1002/2017GC006944.

‘latent heat’: A material model that includes phase transitions and the possibility that latent heat is released or absorbed when material crosses one of the phase transitions of up to two different materials (compositional fields). This model implements a standard approximation of the latent heat terms following Christensen & Yuen, 1985. The change of entropy is calculated as $\Delta S = \gamma \frac{\Delta \rho}{\rho^2}$ with the Clapeyron slope γ and the density change $\Delta \rho$ of the phase transition being input parameters. The model employs an analytic phase function in the form $X = 0.5 \left(1 + \tanh \left(\frac{\Delta p}{\Delta p_0} \right) \right)$ with $\Delta p = p - p_{\text{transition}} - \gamma(T - T_{\text{transition}})$ and Δp_0 being the pressure difference over the width of the phase transition (specified as input parameter).

‘latent heat melt’: A material model that includes the latent heat of melting for two materials: peridotite and pyroxenite. The melting model for peridotite is taken from Katz et al., 2003 (A new parameterization of hydrous mantle melting) and the one for pyroxenite from Sobolev et al., 2011 (Linking mantle plumes, large igneous provinces and environmental catastrophes). The model assumes a constant entropy change for melting 100% of the material, which can be specified in the input file. The partial derivatives of entropy with respect to temperature and pressure required for calculating the latent heat consumption are then calculated as product of this constant entropy change, and the respective derivative of the function describes the melt fraction. This is linearly averaged with respect to the fractions of the two materials present. If no compositional fields are specified in the input file, the model assumes that the material is peridotite. If compositional fields are specified, the model assumes that the first compositional field is the fraction of pyroxenite and the rest of the material is peridotite.

Otherwise, this material model has a temperature- and pressure-dependent density and viscosity and the density and thermal expansivity depend on the melt fraction present. It is possible to extent this model to include a melt fraction dependence of all the material parameters by calling the function melt_fraction in the calculation of the respective parameter. However, melt and solid move with the same velocity and melt extraction is not taken into account (batch melting).

‘melt global’: A material model that implements a simple formulation of the material parameters required for the modelling of melt transport, including a source term for the porosity according to a simplified linear melting model similar to [Sch06]: $\phi_{\text{equilibrium}} = \frac{T - T_{\text{sol}}}{T_{\text{liq}} - T_{\text{sol}}}$ with $T_{\text{sol}} = T_{\text{sol},0} + \Delta T_p p + \Delta T_c C T_{\text{liq}} = T_{\text{sol}} + \Delta T_{\text{sol-liq}}$.

‘melt simple’: A material model that implements a simple formulation of the material parameters required for the modelling of melt transport, including a source term for the porosity according to the melting model for dry peridotite of [KSL03]. This also includes a computation of the latent heat of melting (if the ‘latent heat’ heating model is active).

Most of the material properties are constant, except for the shear, viscosity η , the compaction viscosity

ξ , and the permeability k , which depend on the porosity; and the solid and melt densities, which depend on temperature and pressure: $\eta(\phi, T) = \eta_0 e^{\alpha(\phi - \phi_0)} e^{-\beta(T - T_0)/T_0}$, $\xi(\phi, T) = \xi_0 \frac{\phi_0}{\phi} e^{-\beta(T - T_0)/T_0}$, $k = k_0 \phi^n (1 - \phi)^m$, $\rho = \rho_0 (1 - \alpha(T - T_{\text{adi}})) e^{\kappa p}$.

The model is compressible only if this is specified in the input file, and contains compressibility for both solid and melt.

‘multicomponent’: This model is for use with an arbitrary number of compositional fields, where each field represents a rock type which can have completely different properties from the others. However, each rock type itself has constant material properties. The value of the compositional field is interpreted as a volume fraction. If the sum of the fields is greater than one, they are renormalized. If it is less than one, material properties for “background mantle” make up the rest. When more than one field is present, the material properties are averaged arithmetically. An exception is the viscosity, where the averaging should make more of a difference. For this, the user selects between arithmetic, harmonic, geometric, or maximum composition averaging.

‘nondimensional’: A material model for nondimensionalized computations for compressible or incompressible computations defined through Rayleigh number extRa and Dissipation number Di. This model is made to be used with the Boussinesq, ALA, or TALA formulation.

The viscosity is defined as

$$\eta = \text{Di/Ra} \cdot \exp(-bT' + cz)$$

where T' is the temperature variation from the adiabatic temperature, z is the depth, b is given by “Viscosity temperature prefactor”, and c by “Viscosity depth prefactor”. If Di is zero, it will be replaced by 1.0 in η .

The density is defined as

$$\rho = \exp(\text{Di}/\gamma \cdot z)(1.0 - \alpha T' + \text{Di}\gamma p'),$$

where $\alpha = \text{Di}$ is the thermal expansion coefficient, γ is the Grueneisen parameter, and p' is the pressure variation from the adiabatic pressure. The pressure dependent term is not present if “TALA” is enabled.

‘simple’: A material model that has constant values for all coefficients but the density and viscosity. The defaults for all coefficients are chosen to be similar to what is believed to be correct for Earth’s mantle. All of the values that define this model are read from a section “Material model/Simple model” in the input file, see Section A.102.

This model uses the following set of equations for the two coefficients that are non-constant:

$$\eta(p, T, \mathbf{c}) = \tau(T)\zeta(\mathbf{c})\eta_0, \quad (92)$$

$$\rho(p, T, \mathbf{c}) = (1 - \alpha(T - T_0))\rho_0 + \Delta\rho c_0, \quad (93)$$

where c_0 is the first component of the compositional vector \mathbf{c} if the model uses compositional fields, or zero otherwise.

The temperature pre-factor for the viscosity formula above is defined as

$$\tau(T) = H \left(e^{-\beta(T - T_0)/T_0} \right), \quad (94)$$

with

$$H(x) = \begin{cases} \tau_{\min} & \text{if } x < \tau_{\min}, \\ x & \text{if } 10^{-2} \leq x \leq 10^2, \\ \tau_{\max} & \text{if } x > \tau_{\max}, \end{cases} \quad (95)$$

where $x = e^{-\beta(T-T_0)/T_0}$, β corresponds to the input parameter “Thermal viscosity exponent”, and T_0 to the parameter “Reference temperature”. If you set $T_0 = 0$ in the input file, the thermal pre-factor $\tau(T) = 1$. The parameters τ_{\min} and τ_{\max} set the minimum and maximum values of the temperature pre-factor and are set using “Maximum thermal prefactor” and “Minimum thermal prefactor”. Specifying a value of 0.0 for the minimum or maximum values will disable pre-factor limiting.

The compositional pre-factor for the viscosity is defined as

$$\zeta(\mathfrak{c}) = \xi^{co} \quad (96)$$

if the model has compositional fields and equals one otherwise. ξ corresponds to the parameter “Composition viscosity prefactor” in the input file.

Finally, in the formula for the density, α corresponds to the “Thermal expansion coefficient” and $\Delta\rho$ corresponds to the parameter “Density differential for compositional field 1”.

Note that this model uses the formulation that assumes an incompressible medium despite the fact that the density follows the law $\rho(T) = \rho_0(1 - \alpha(T - T_{ref}))$.

Note: Despite its name, this material model is not exactly “simple”, as indicated by the formulas above. While it was originally intended to be simple, it has over time acquired all sorts of temperature and compositional dependencies that weren’t initially intended. Consequently, there is now a “simpler” material model that now fills the role the current model was originally intended to fill.

‘simple compressible’: A material model that has constant values for all coefficients but the density. The defaults for all coefficients are chosen to be similar to what is believed to be correct for Earth’s mantle. All of the values that define this model are read from a section “Material model/Simple compressible model” in the input file, see Section [A.101](#).

This model uses the following equations for the density:

$$\rho(p, T) = \rho_0 (1 - \alpha(T - T_a)) \exp \beta(P - P_0) \quad (97)$$

‘simpler’: A material model that has constant values except for density, which depends linearly on temperature:

$$\rho(p, T) = (1 - \alpha(T - T_0)) \rho_0. \quad (98)$$

Note: This material model fills the role the “simple” material model was originally intended to fill, before the latter acquired all sorts of complicated temperature and compositional dependencies.

‘visco plastic’: An implementation of a visco-plastic rheology with options for selecting dislocation creep, diffusion creep or composite viscous flow laws. Plasticity limits viscous stresses through a Drucker Prager yield criterion. The model is incompressible. Note that this material model is based heavily on the DiffusionDislocation (Bob Myhill) and DruckerPrager (Anne Glerum) material models.

The viscosity for dislocation or diffusion creep is defined as

$$v = \frac{1}{2} A^{-\frac{1}{n}} d^{\frac{m}{n}} \dot{\varepsilon}_{ii}^{\frac{1-n}{n}} \exp \left(\frac{E + PV}{nRT} \right)$$

where A is the prefactor, n is the stress exponent, $\dot{\varepsilon}_{ii}$ is the square root of the deviatoric strain rate tensor second invariant, d is grain size, m is the grain size exponent, E is activation energy, V is

activation volume, P is pressure, R is the gas exponent and T is temperature. This form of the viscosity equation is commonly used in geodynamic simulations. See, for example, Billen and Hirth (2007), G3, 8, Q08012. Significantly, other studies may use slightly different forms of the viscosity equation leading to variations in how specific terms are defined or combined. For example, the grain size exponent should always be positive in the diffusion viscosity equation used here, while other studies place the grain size term in the denominator and invert the sign of the grain size exponent. When examining previous work, one should carefully check how the viscous prefactor and grain size terms are defined.

One may select to use the diffusion (v_{diff} ; $n = 1$, $m! = 0$), dislocation (v_{disl} , $n > 1$, $m = 0$) or composite $\frac{v_{\text{diff}} * v_{\text{disl}}}{v_{\text{diff}} + v_{\text{disl}}}$ equation form.

Viscosity is limited through one of two different ‘yielding’ mechanisms.

Plasticity limits viscous stress through a Drucker Prager yield criterion, where the yield stress in 3D is $\sigma_y = \frac{6*C*\cos(\phi)+2*P*\sin(\phi)}{\sqrt{(3)*(3+\sin(\phi))}}$ and $\sigma_y = C \cos(\phi) + P \sin(\phi)$ in 2D. Above, C is cohesion and ϕ is the angle of internal friction. Note that the 2D form is equivalent to the Mohr Coulomb yield surface. If ϕ is 0, the yield stress is fixed and equal to the cohesion (Von Mises yield criterion). When the viscous stress ($2v\varepsilon_{ii}$) exceeds the yield stress, the viscosity is rescaled back to the yield surface: $v_y = \sigma_y/(2\varepsilon_{ii})$. This form of plasticity is commonly used in geodynamic models. See, for example, Thieulot, C. (2011), PEPI 188, pp. 47-68.

The user has the option to linearly reduce the cohesion and internal friction angle as a function of the finite strain magnitude. The finite strain invariant or full strain tensor is calculated through compositional fields within the material model. This implementation is identical to the compositional field finite strain plugin and cookbook described in the manual (author: Gassmoeller, Dannberg). If the user selects to track the finite strain invariant (e_{ii}), a single compositional field tracks the value derived from $e_{ii}^t = (e_{ii})^{(t-1)} + \dot{e}_{ii} dt$, where t and $t - 1$ are the current and prior time steps, \dot{e}_{ii} is the second invariant of the strain rate tensor and dt is the time step size. In the case of the full strain tensor F , the finite strain magnitude is derived from the second invariant of the symmetric stretching tensor L , where $L = F[F]^T$. The user must specify a single compositional field for the finite strain invariant or multiple fields (4 in 2D, 9 in 3D) for the finite strain tensor. These field(s) must be the first listed compositional fields in the parameter file. Note that one or more of the finite strain tensor components must be assigned a non-zero value initially. This value can be quite small (e.g., 1.e-8), but still non-zero. While the option to track and use the full finite strain tensor exists, tracking the associated compositional fields is computationally expensive in 3D. Similarly, the finite strain magnitudes may in fact decrease if the orientation of the deformation field switches through time. Consequently, the ideal solution is track the finite strain invariant (single compositional) field within the material and track the full finite strain tensor through particles.

Viscous stress may also be limited by a non-linear stress limiter that has a form similar to the Peierls creep mechanism. This stress limiter assigns an effective viscosity $\sigma_{\text{eff}} = \frac{\tau_y}{2\varepsilon_y} \frac{\varepsilon_{ii}}{\varepsilon_y}^{\frac{1}{n_y}-1}$. Above τ_y is a yield stress, ε_y is the reference strain rate, ε_{ii} is the strain rate and n_y is the stress limiter exponent. The yield stress, τ_y , is defined through the Drucker Prager yield criterion formulation. This method of limiting viscous stress has been used in various forms within the geodynamic literature, including Christensen (1992), JGR, 97(B2), pp. 2015-2036; Cizkova and Bina (2013), EPSL, 379, pp. 95-103; Cizkova and Bina (2015), EPSL, 430, pp. 408-415. When n_y is 1, it essentially becomes a linear viscosity model, and in the limit $n_y \rightarrow \infty$ it converges to the standard viscosity rescaling method (concretely, values $n_y > 20$ are large enough).

Compositional fields can each be assigned individual values of thermal diffusivity, heat capacity, density, thermal expansivity and rheological parameters.

If more than one compositional field is present at a given point, viscosities are averaged with an arithmetic, geometric harmonic (default) or maximum composition scheme.

The value for the components of this formula and additional parameters are read from the parameter file in subsection 'Material model/Visco Plastic'.

Possible values: Any one of Steinberger, ascii reference profile, averaging, compositing, composition reaction, depth dependent, diffusion dislocation, drucker prager, dynamic friction, grain size, latent heat, latent heat melt, melt global, melt simple, multicomponent, nondimensional, simple, simple compressible, simpler, visco plastic, unspecified

A.82 Parameters in section Material model/Ascii reference profile

- *Parameter name:* Thermal conductivity

Value: 4.0

Default: 4.0

Description: Reference conductivity

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Thermal viscosity exponent

Value: 0.0

Default: 0.0

Description: The temperature dependence of viscosity. Dimensionless exponent.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Transition depths

Value: 1.5e5, 4.1e5, 6.6e5

Default: 1.5e5, 4.1e5, 6.6e5

Description: A list of depths where the viscosity changes. Values must monotonically increase. Units: m.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Use TALA

Value: false

Default: false

Description: Whether to use the TALA instead of the ALA approximation.

Possible values: A boolean value (true or false)

- *Parameter name:* Viscosity

Value: 1e21

Default: 1e21

Description: Viscosity

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Viscosity prefactors

Value: 10, 0.1, 1, 10

Default: 10, 0.1, 1, 10

Description: A list of prefactors for the viscosity that determine the viscosity profile. Each prefactor is applied in a depth range specified by the list of 'Transition depths', i.e. the first prefactor is applied

above the first transition depth, the second one between the first and second transition depth, and so on. To compute the viscosity profile, this prefactor is multiplied by the reference viscosity specified through the parameter ‘Viscosity’. List must have one more entry than Transition depths. Units: non-dimensional.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

A.83 Parameters in section Material model/Ascii reference profile/Ascii data model

- *Parameter name:* Data directory

Value: \$ASPECT_SOURCE_DIR/data/adiabatic-conditions/ascii-data/

Default: \$ASPECT_SOURCE_DIR/data/adiabatic-conditions/ascii-data/

Description: The name of a directory that contains the model data. This path may either be absolute (if starting with a '/') or relative to the current directory. The path may also include the special text '\$ASPECT_SOURCE_DIR' which will be interpreted as the path in which the ASPECT source files were located when ASPECT was compiled. This interpretation allows, for example, to reference files located in the ‘data/’ subdirectory of ASPECT.

Possible values: A directory name

- *Parameter name:* Data file name

Value:

Default:

Description: The file name of the material data. Provide file in format: (Velocity file name).%s%d where %s is a string specifying the boundary of the model according to the names of the boundary indicators (of a box or a spherical shell).%d is any sprintf integer qualifier, specifying the format of the current file number.

Possible values: Any string

- *Parameter name:* Scale factor

Value: 1

Default: 1

Description: Scalar factor, which is applied to the boundary velocity. You might want to use this to scale the velocities to a reference model (e.g. with free-slip boundary) or another plate reconstruction. Another way to use this factor is to convert units of the input files. The unit is assumed to be m/s or m/yr depending on the ‘Use years in output instead of seconds’ flag. If you provide velocities in cm/yr set this factor to 0.01.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.84 Parameters in section Material model/Averaging

- *Parameter name:* Averaging operation

Value: none

Default: none

Description: Choose the averaging operation to use.

Possible values: Any one of none, arithmetic average, harmonic average, geometric average, pick largest, log average, nwd arithmetic average, nwd harmonic average, nwd geometric average

- *Parameter name:* Base model

Value: simple

Default: simple

Description: The name of a material model that will be modified by an averaging operation. Valid values for this parameter are the names of models that are also valid for the “Material models/Model name” parameter. See the documentation for that for more information.

Possible values: Any one of Steinberger, ascii reference profile, averaging, compositing, composition reaction, depth dependent, diffusion dislocation, drucker prager, dynamic friction, grain size, latent heat, latent heat melt, melt global, melt simple, multicomponent, nondimensional, simple, simple compressible, simpler, visco plastic

- *Parameter name:* Bell shape limit

Value: 1

Default: 1

Description: The limit normalized distance between 0 and 1 where the bell shape becomes zero. See the manual for a more information.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.85 Parameters in section Material model/Compositing

- *Parameter name:* Compressibility

Value: unspecified

Default: unspecified

Description: Material model to use for Compressibility. Valid values for this parameter are the names of models that are also valid for the “Material models/Model name” parameter. See the documentation for that for more information.

Possible values: Any one of Steinberger, ascii reference profile, averaging, compositing, composition reaction, depth dependent, diffusion dislocation, drucker prager, dynamic friction, grain size, latent heat, latent heat melt, melt global, melt simple, multicomponent, nondimensional, simple, simple compressible, simpler, visco plastic, unspecified

- *Parameter name:* Density

Value: unspecified

Default: unspecified

Description: Material model to use for Density. Valid values for this parameter are the names of models that are also valid for the “Material models/Model name” parameter. See the documentation for that for more information.

Possible values: Any one of Steinberger, ascii reference profile, averaging, compositing, composition reaction, depth dependent, diffusion dislocation, drucker prager, dynamic friction, grain size, latent heat, latent heat melt, melt global, melt simple, multicomponent, nondimensional, simple, simple compressible, simpler, visco plastic, unspecified

- *Parameter name:* Entropy derivative pressure

Value: unspecified

Default: unspecified

Description: Material model to use for Entropy derivative pressure. Valid values for this parameter are the names of models that are also valid for the “Material models/Model name” parameter. See the documentation for that for more information.

Possible values: Any one of Steinberger, ascii reference profile, averaging, compositing, composition reaction, depth dependent, diffusion dislocation, drucker prager, dynamic friction, grain size, latent heat, latent heat melt, melt global, melt simple, multicomponent, nondimensional, simple, simple compressible, simpler, visco plastic, unspecified

- *Parameter name:* Entropy derivative temperature

Value: unspecified

Default: unspecified

Description: Material model to use for Entropy derivative temperature. Valid values for this parameter are the names of models that are also valid for the “Material models/Model name” parameter. See the documentation for that for more information.

Possible values: Any one of Steinberger, ascii reference profile, averaging, compositing, composition reaction, depth dependent, diffusion dislocation, drucker prager, dynamic friction, grain size, latent heat, latent heat melt, melt global, melt simple, multicomponent, nondimensional, simple, simple compressible, simpler, visco plastic, unspecified

- *Parameter name:* Reaction terms

Value: unspecified

Default: unspecified

Description: Material model to use for Reaction terms. Valid values for this parameter are the names of models that are also valid for the “Material models/Model name” parameter. See the documentation for that for more information.

Possible values: Any one of Steinberger, ascii reference profile, averaging, compositing, composition reaction, depth dependent, diffusion dislocation, drucker prager, dynamic friction, grain size, latent heat, latent heat melt, melt global, melt simple, multicomponent, nondimensional, simple, simple compressible, simpler, visco plastic, unspecified

- *Parameter name:* Specific heat

Value: unspecified

Default: unspecified

Description: Material model to use for Specific heat. Valid values for this parameter are the names of models that are also valid for the “Material models/Model name” parameter. See the documentation for that for more information.

Possible values: Any one of Steinberger, ascii reference profile, averaging, compositing, composition reaction, depth dependent, diffusion dislocation, drucker prager, dynamic friction, grain size, latent heat, latent heat melt, melt global, melt simple, multicomponent, nondimensional, simple, simple compressible, simpler, visco plastic, unspecified

- *Parameter name:* Thermal conductivity

Value: unspecified

Default: unspecified

Description: Material model to use for Thermal conductivity. Valid values for this parameter are the names of models that are also valid for the “Material models/Model name” parameter. See the documentation for that for more information.

Possible values: Any one of Steinberger, ascii reference profile, averaging, compositing, composition reaction, depth dependent, diffusion dislocation, drucker prager, dynamic friction, grain size, latent heat, latent heat melt, melt global, melt simple, multicomponent, nondimensional, simple, simple compressible, simpler, visco plastic, unspecified

- *Parameter name:* Thermal expansion coefficient

Value: unspecified

Default: unspecified

Description: Material model to use for Thermal expansion coefficient. Valid values for this parameter are the names of models that are also valid for the “Material models/Model name” parameter. See the documentation for that for more information.

Possible values: Any one of Steinberger, ascii reference profile, averaging, compositing, composition reaction, depth dependent, diffusion dislocation, drucker prager, dynamic friction, grain size, latent heat, latent heat melt, melt global, melt simple, multicomponent, nondimensional, simple, simple compressible, simpler, visco plastic, unspecified

- *Parameter name:* Viscosity

Value: unspecified

Default: unspecified

Description: Material model to use for Viscosity. Valid values for this parameter are the names of models that are also valid for the “Material models/Model name” parameter. See the documentation for that for more information.

Possible values: Any one of Steinberger, ascii reference profile, averaging, compositing, composition reaction, depth dependent, diffusion dislocation, drucker prager, dynamic friction, grain size, latent heat, latent heat melt, melt global, melt simple, multicomponent, nondimensional, simple, simple compressible, simpler, visco plastic, unspecified

A.86 Parameters in section Material model/Composition reaction model

- *Parameter name:* Composition viscosity prefactor 1

Value: 1.0

Default: 1.0

Description: A linear dependency of viscosity on the first compositional field. Dimensionless prefactor. With a value of 1.0 (the default) the viscosity does not depend on the composition.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Composition viscosity prefactor 2

Value: 1.0

Default: 1.0

Description: A linear dependency of viscosity on the second compositional field. Dimensionless prefactor. With a value of 1.0 (the default) the viscosity does not depend on the composition.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Density differential for compositional field 1

Value: 0

Default: 0

Description: If compositional fields are used, then one would frequently want to make the density depend on these fields. In this simple material model, we make the following assumptions: if no compositional fields are used in the current simulation, then the density is simply the usual one with its linear dependence on the temperature. If there are compositional fields, then the density only depends on the first and the second one in such a way that the density has an additional term of the kind $+\Delta\rho c_1(\mathbf{x})$. This parameter describes the value of $\Delta\rho$ for the first field. Units: kg/m^3 /unit change in composition.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Density differential for compositional field 2

Value: 0

Default: 0

Description: If compositional fields are used, then one would frequently want to make the density depend on these fields. In this simple material model, we make the following assumptions: if no compositional fields are used in the current simulation, then the density is simply the usual one with its linear dependence on the temperature. If there are compositional fields, then the density only depends on the first and the second one in such a way that the density has an additional term of the kind $+\Delta\rho c_1(\mathbf{x})$. This parameter describes the value of $\Delta\rho$ for the second field. Units: kg/m^3 /unit change in composition.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reaction depth

Value: 0

Default: 0

Description: Above this depth the compositional fields react: The first field gets converted to the second field. Units: m .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference density

Value: 3300

Default: 3300

Description: Reference density ρ_0 . Units: kg/m^3 .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference specific heat

Value: 1250

Default: 1250

Description: The value of the specific heat C_p . Units: $J/kg/K$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference temperature

Value: 293

Default: 293

Description: The reference temperature T_0 . Units: K .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Thermal conductivity
Value: 4.7
Default: 4.7
Description: The value of the thermal conductivity k . Units: $W/m/K$.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Thermal expansion coefficient
Value: 2e-5
Default: 2e-5
Description: The value of the thermal expansion coefficient β . Units: $1/K$.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Thermal viscosity exponent
Value: 0.0
Default: 0.0
Description: The temperature dependence of viscosity. Dimensionless exponent.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Viscosity
Value: 5e24
Default: 5e24
Description: The value of the constant viscosity. Units: $kg/m/s$.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.87 Parameters in section Material model/Depth dependent model

- *Parameter name:* Base model
Value: simple
Default: simple
Description: The name of a material model that will be modified by a depth dependent viscosity. Valid values for this parameter are the names of models that are also valid for the “Material models/Model name” parameter. See the documentation for that for more information.
Possible values: Any one of Steinberger, ascii reference profile, averaging, compositing, composition reaction, depth dependent, diffusion dislocation, drucker prager, dynamic friction, grain size, latent heat, latent heat melt, melt global, melt simple, multicomponent, nondimensional, simple, simple compressible, simpler, visco plastic
- *Parameter name:* Data directory
Value: ./
Default: ./
Description: The path to the model data. The path may also include the special text ‘\$ASPECT_SOURCE_DIR’ which will be interpreted as the path in which the ASPECT source files were located when ASPECT was compiled. This interpretation allows, for example, to reference files located in the ‘data/’ subdirectory of ASPECT.
Possible values: A directory name

- *Parameter name:* Depth dependence method

Value: None

Default: None

Description: Method that is used to specify how the viscosity should vary with depth.

Possible values: Any one of Function, File, List, None

- *Parameter name:* Depth list

Value:

Default:

Description: A comma-separated list of depth values for use with the “List” “Depth dependence method”. The list must be provided in order of increasing depth, and the last value must be greater than or equal to the maximal depth of the model. The depth list is interpreted as a layered viscosity structure and the depth values specify the maximum depths of each layer.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Viscosity depth file

Value: visc-depth.txt

Default: visc-depth.txt

Description: The name of the file containing depth-dependent viscosity data.

Possible values: Any string

- *Parameter name:* Viscosity list

Value:

Default:

Description: A comma-separated list of viscosity values, corresponding to the depth values provided in “Depth list”. The number of viscosity values specified here must be the same as the number of depths provided in “Depth list”

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

A.88 Parameters in section Material model/Depth dependent model/Viscosity depth function

- *Parameter name:* Function constants

Value:

Default:

Description: Sometimes it is convenient to use symbolic constants in the expression that describes the function, rather than having to use its numeric value everywhere the constant appears. These values can be defined using this parameter, in the form ‘var1=value1, var2=value2, ...’.

A typical example would be to set this runtime parameter to ‘pi=3.1415926536’ and then use ‘pi’ in the expression of the actual formula. (That said, for convenience this class actually defines both ‘pi’ and ‘Pi’ by default, but you get the idea.)

Possible values: Any string

- **Parameter name:** Function expression

Value: 1.0e21

Default: 1.0e21

Possible values: Any string

- **Parameter name:** Variable names

Value: x,t

Default: x,t

Description: The name of the variables as they will be used in the function, separated by commas. By default, the names of variables at which the function will be evaluated is ‘x’ (in 1d), ‘x,y’ (in 2d) or ‘x,y,z’ (in 3d) for spatial coordinates and ‘t’ for time. You can then use these variable names in your function expression and they will be replaced by the values of these variables at which the function is currently evaluated. However, you can also choose a different set of names for the independent variables at which to evaluate your function expression. For example, if you work in spherical coordinates, you may wish to set this input parameter to ‘r,phi,theta,t’ and then use these variable names in your function expression.

Possible values: Any string

A.89 Parameters in section Material model/Diffusion dislocation

- **Parameter name:** Activation energies for diffusion creep

Value: 375e3

Default: 375e3

Description: List of activation energies, E_a , for background mantle and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: J/mol

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- **Parameter name:** Activation energies for dislocation creep

Value: 530e3

Default: 530e3

Description: List of activation energies, E_a , for background mantle and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: J/mol

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- **Parameter name:** Activation volumes for diffusion creep

Value: 6e-6

Default: 6e-6

Description: List of activation volumes, V_a , for background mantle and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: m^3/mol

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Activation volumes for dislocation creep

Value: 1.4e-5

Default: 1.4e-5

Description: List of activation volumes, V_a , for background mantle and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: m^3/mol

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Densities

Value: 3300.

Default: 3300.

Description: List of densities, ρ , for background mantle and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: kg/m^3

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Effective viscosity coefficient

Value: 1.0

Default: 1.0

Description: Scaling coefficient for effective viscosity.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Grain size

Value: 1e-3

Default: 1e-3

Description: Units: m

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Grain size exponents for diffusion creep

Value: 3

Default: 3

Description: List of grain size exponents, $m_{\text{diffusion}}$, for background mantle and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: None

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Heat capacity

Value: 1.25e3

Default: 1.25e3

Description: The value of the specific heat C_p . Units: $J/kg/K$

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Maximum strain rate ratio iterations
Value: 40
Default: 40
Description: Maximum number of iterations to find the correct diffusion/dislocation strain rate ratio.
Possible values: An integer n such that $0 \leq n \leq 2147483647$
- **Parameter name:** Maximum viscosity
Value: 1e28
Default: 1e28
Description: Upper cutoff for effective viscosity. Units: Pas
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- **Parameter name:** Minimum strain rate
Value: 1.4e-20
Default: 1.4e-20
Description: Stabilizes strain dependent viscosity. Units: 1/s
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- **Parameter name:** Minimum viscosity
Value: 1e17
Default: 1e17
Description: Lower cutoff for effective viscosity. Units: Pas
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- **Parameter name:** Prefactors for diffusion creep
Value: 1.5e-15
Default: 1.5e-15
Description: List of viscosity prefactors, A , for background mantle and compositional fields, for a total of $N+1$ values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: $\text{Pa}^{-1} \text{m}^{m_{\text{diffusion}}} \text{s}^{-1}$
Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]
- **Parameter name:** Prefactors for dislocation creep
Value: 1.1e-16
Default: 1.1e-16
Description: List of viscosity prefactors, A , for background mantle and compositional fields, for a total of $N+1$ values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: $\text{Pa}^{-n_{\text{dislocation}}} \text{s}^{-1}$
Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]
- **Parameter name:** Reference temperature
Value: 293
Default: 293
Description: For calculating density by thermal expansivity. Units: K
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference viscosity

Value: 1e22

Default: 1e22

Description: Reference viscosity for nondimensionalization. Units Pas

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Strain rate residual tolerance

Value: 1e-22

Default: 1e-22

Description: Tolerance for correct diffusion/dislocation strain rate ratio.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Stress exponents for diffusion creep

Value: 1

Default: 1

Description: List of stress exponents, $n_{\text{diffusion}}$, for background mantle and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: None

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Stress exponents for dislocation creep

Value: 3.5

Default: 3.5

Description: List of stress exponents, $n_{\text{dislocation}}$, for background mantle and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: None

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Thermal diffusivity

Value: 0.8e-6

Default: 0.8e-6

Description: Units: m^2/s

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Thermal expansivities

Value: 3.5e-5

Default: 3.5e-5

Description: List of thermal expansivities for background mantle and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: $1/K$

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Viscosity averaging scheme

Value: harmonic

Default: harmonic

Description: When more than one compositional field is present at a point with different viscosities, we need to come up with an average viscosity at that point. Select a weighted harmonic, arithmetic, geometric, or maximum composition.

Possible values: Any one of arithmetic, harmonic, geometric, maximum composition

A.90 Parameters in section Material model/Drucker Prager

- *Parameter name:* Reference density

Value: 3300

Default: 3300

Description: The reference density ρ_0 . Units: kg/m^3 .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference specific heat

Value: 1250

Default: 1250

Description: The value of the specific heat C_p . Units: $J/kg/K$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference temperature

Value: 293

Default: 293

Description: The reference temperature T_0 . The reference temperature is used in the density calculation. Units: K .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference viscosity

Value: 1e22

Default: 1e22

Description: The value of the reference viscosity η_0 . Units: $kg/m/s$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Thermal conductivity

Value: 4.7

Default: 4.7

Description: The value of the thermal conductivity k . Units: $W/m/K$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Thermal expansion coefficient

Value: 2e-5

Default: 2e-5

Description: The value of the thermal expansion coefficient β . Units: $1/K$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.91 Parameters in section Material model/Drucker Prager/Viscosity

- Parameter name: Angle of internal friction

Value: 0

Default: 0

Description: The value of the angle of internal friction ϕ . For a value of zero, in 2D the von Mises criterion is retrieved. Angles higher than 30 degrees are harder to solve numerically. Units: degrees.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: Cohesion

Value: 2e7

Default: 2e7

Description: The value of the cohesion C . Units: Pa.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: Maximum viscosity

Value: 1e24

Default: 1e24

Description: The value of the maximum viscosity cutoff η_{max} . Units: Pa s.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: Minimum viscosity

Value: 1e19

Default: 1e19

Description: The value of the minimum viscosity cutoff η_{min} . Units: Pa s.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: Reference strain rate

Value: 1e-15

Default: 1e-15

Description: The value of the initial strain rate prescribed during the first nonlinear iteration $\dot{\epsilon}_{ref}$. Units: 1/s.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.92 Parameters in section Material model/Dynamic Friction

- Parameter name: Densities

Value: 3300.

Default: 3300.

Description: List of densities for background mantle and compositional fields,for a total of N+1 values, where N is the number of compositional fields.If only one value is given, then all use the same value. Units: kg/m³

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Reference temperature

Value: 293

Default: 293

Description: The reference temperature T_0 . Units: K .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Specific heats

Value: 1250.

Default: 1250.

Description: List of specific heats C_p for background mantle and compositional fields,for a total of $N+1$ values, where N is the number of compositional fields.If only one value is given, then all use the same value. Units: $J/kg/K$

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Thermal conductivities

Value: 4.7

Default: 4.7

Description: List of thermal conductivities for background mantle and compositional fields,for a total of $N+1$ values, where N is the number of compositional fields.If only one value is given, then all use the same value. Units: $W/m/K$

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Thermal expansivities

Value: 4.e-5

Default: 4.e-5

Description: List of thermal expansivities for background mantle and compositional fields,for a total of $N+1$ values, where N is the number of compositional fields.If only one value is given, then all use the same value. Units: $1/K$

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Viscosity averaging scheme

Value: harmonic

Default: harmonic

Description: When more than one compositional field is present at a point with different viscosities, we need to come up with an average viscosity at that point. Select a weighted harmonic, arithmetic, geometric, or maximum composition.

Possible values: Any one of arithmetic, harmonic, geometric, maximum composition

A.93 Parameters in section Material model/Dynamic Friction/Viscosities

- Parameter name: Background Viscosities

Value: 1.e20

Default: 1.e20

Description: List of background viscosities for mantle and compositional fields,for a total of N+1 values, where N is the number of compositional fields.If only one value is given, then all use the same value. Units: Pas

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- Parameter name: Coefficients of dynamic friction

Value: 0.4

Default: 0.4

Description: List of coefficients of dynamic friction for background mantle and compositional fields,for a total of N+1 values, where N is the number of compositional fields.If only one value is given, then all use the same value. Units: dimensionless

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- Parameter name: Coefficients of static friction

Value: 0.5

Default: 0.5

Description: List of coefficients of static friction for background mantle and compositional fields,for a total of N+1 values, where N is the number of compositional fields.If only one value is given, then all use the same value. Units: dimensionless

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- Parameter name: Cohesions

Value: 4.e6

Default: 4.e6

Description: List of cohesions for background mantle and compositional fields,for a total of N+1 values, where N is the number of compositional fields.If only one value is given, then all use the same value. Units: Pa

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- Parameter name: Maximum viscosity

Value: 1e24

Default: 1e24

Description: The value of the maximum viscosity cutoff η_{max} . Units: Pa s.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: Minimum viscosity

Value: 1e19

Default: 1e19

Description: The value of the minimum viscosity cutoff η_{min} . Units: Pa s.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference strain rate

Value: 1e-15

Default: 1e-15

Description: The value of the initial strain rate prescribed during the first nonlinear iteration $\dot{\epsilon}_{ref}$. Units: 1/s.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.94 Parameters in section Material model/Grain size model

- *Parameter name:* Advect logarithm of grain size

Value: false

Default: false

Description: Whether to advect the logarithm of the grain size or the grain size. The equation and the physics are the same, but for problems with high grain size gradients it might be preferable to advect the logarithm.

Possible values: A boolean value (true or false)

- *Parameter name:* Average specific grain boundary energy

Value: 1.0

Default: 1.0

Description: The average specific grain boundary energy γ . Units: J/m².

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Bilinear interpolation

Value: true

Default: true

Description: Whether to use bilinear interpolation to compute material properties (slower but more accurate).

Possible values: A boolean value (true or false)

- *Parameter name:* Data directory

Value: \$ASPECT_SOURCE_DIR/data/material-model/steinberger/

Default: \$ASPECT_SOURCE_DIR/data/material-model/steinberger/

Description: The path to the model data. The path may also include the special text '\$ASPECT_SOURCE_DIR' which will be interpreted as the path in which the ASPECT source files were located when ASPECT was compiled. This interpretation allows, for example, to reference files located in the 'data/' subdirectory of ASPECT.

Possible values: A directory name

- *Parameter name:* Derivatives file names

Value:

Default:

Description: The file names of the enthalpy derivatives data. List with as many components as active compositional fields (material data is assumed to be in order with the ordering of the fields).

Possible values: A list of 0 to 4294967295 elements where each element is [Any string]

- *Parameter name:* Diffusion activation energy

Value: 3.35e5

Default: 3.35e5

Description: The activation energy for diffusion creep E_{diff} . Units: J/mol .

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Diffusion activation volume

Value: 4e-6

Default: 4e-6

Description: The activation volume for diffusion creep V_{diff} . Units: m^3/mol .

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Diffusion creep exponent

Value: 1

Default: 1

Description: Power-law exponent n_{diff} for diffusion creep. Units: none.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Diffusion creep grain size exponent

Value: 3

Default: 3

Description: Diffusion creep grain size exponent p_{diff} that determines the dependence of viscosity on grain size. Units: none.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Diffusion creep prefactor

Value: 7.4e-15

Default: 7.4e-15

Description: Prefactor for the diffusion creep law A_{diff} . Units: $m^{p_{diff}} Pa^{-n_{diff}}/s$.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- **Parameter name:** Dislocation activation energy

Value: 4.8e5

Default: 4.8e5

Description: The activation energy for dislocation creep E_{dis} . Units: J/mol .

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- **Parameter name:** Dislocation activation volume

Value: 1.1e-5

Default: 1.1e-5

Description: The activation volume for dislocation creep V_{dis} . Units: m^3/mol .

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- **Parameter name:** Dislocation creep exponent

Value: 3.5

Default: 3.5

Description: Power-law exponent n_{dis} for dislocation creep. Units: none.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- **Parameter name:** Dislocation creep prefactor

Value: 4.5e-15

Default: 4.5e-15

Description: Prefactor for the dislocation creep law A_{dis} . Units: $Pa^{-n_{dis}}/s$.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- **Parameter name:** Dislocation viscosity iteration number

Value: 100

Default: 100

Description: We need to perform an iteration inside the computation of the dislocation viscosity, because it depends on the dislocation strain rate, which depends on the dislocation viscosity itself. This number determines the maximum number of iterations that are performed.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- **Parameter name:** Dislocation viscosity iteration threshold

Value: 1e-3

Default: 1e-3

Description: We need to perform an iteration inside the computation of the dislocation viscosity, because it depends on the dislocation strain rate, which depends on the dislocation viscosity itself. This number determines the termination accuracy, i.e. if the dislocation viscosity changes by less than this factor we terminate the iteration.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Geometric constant

Value: 3

Default: 3

Description: Geometric constant c used in the paleowattmeter grain size reduction law. Units: none.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Grain growth activation energy

Value: 3.5e5

Default: 3.5e5

Description: The activation energy for grain growth E_g . Units: J/mol .

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Grain growth activation volume

Value: 8e-6

Default: 8e-6

Description: The activation volume for grain growth V_g . Units: m^3/mol .

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Grain growth exponent

Value: 3

Default: 3

Description: Exponent of the grain growth law p_g . This is an experimentally determined grain growth constant. Units: none.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Grain growth rate constant

Value: 1.5e-5

Default: 1.5e-5

Description: Prefactor of the Ostwald ripening grain growth law G_0 . This is dependent on water content, which is assumed to be 50 H/10⁶ Si for the default value. Units: m^{p_g}/s .

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Lower mantle grain size scaling

Value: 1.0

Default: 1.0

Description: A scaling factor for the grain size in the lower mantle. In models where the high grain size contrast between the upper and lower mantle causes numerical problems, the grain size in the lower mantle can be scaled to a larger value, simultaneously scaling the viscosity prefactors and grain growth parameters to keep the same physical behavior. Differences to the original formulation only occur when material with a smaller grain size than the recrystallization grain size cross the upper-lower

mantle boundary. The real grain size can be obtained by dividing the model grain size by this value.
Units: none.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `Material file format`

Value: perplex

Default: perplex

Description: The material file format to be read in the property tables.

Possible values: Any one of perplex, hefesto

- *Parameter name:* `Material file names`

Value: pyr-ringwood88.txt

Default: pyr-ringwood88.txt

Description: The file names of the material data. List with as many components as active compositional fields (material data is assumed to be in order with the ordering of the fields).

Possible values: A list of 0 to 4294967295 elements where each element is [Any string]

- *Parameter name:* `Maximum latent heat substeps`

Value: 1

Default: 1

Description: The maximum number of substeps over the temperature pressure range to calculate the averaged enthalpy gradient over a cell.

Possible values: An integer n such that $1 \leq n \leq 2147483647$

- *Parameter name:* `Maximum specific heat`

Value: 6000

Default: 6000

Description: The maximum specific heat that is allowed in the whole model domain. Units: J/kg/K.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `Maximum temperature dependence of viscosity`

Value: 100

Default: 100

Description: The factor by which viscosity at adiabatic temperature and ambient temperature are allowed to differ (a value of x means that the viscosity can be x times higher or x times lower compared to the value at adiabatic temperature. This parameter is introduced to limit local viscosity contrasts, but still allow for a widely varying viscosity over the whole mantle range. Units: none.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `Maximum thermal expansivity`

Value: 1e-3

Default: 1e-3

Description: The maximum thermal expansivity that is allowed in the whole model domain. Units: 1/K.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Maximum viscosity

Value: 1e26

Default: 1e26

Description: The maximum viscosity that is allowed in the whole model domain. Units: Pa s.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Minimum grain size

Value: 1e-5

Default: 1e-5

Description: The minimum grain size that is used for the material model. This parameter is introduced to limit local viscosity contrasts, but still allows for a widely varying viscosity over the whole mantle range. Units: m.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Minimum specific heat

Value: 500

Default: 500

Description: The minimum specific heat that is allowed in the whole model domain. Units: J/kg/K.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Minimum thermal expansivity

Value: 1e-5

Default: 1e-5

Description: The minimum thermal expansivity that is allowed in the whole model domain. Units: 1/K.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Minimum viscosity

Value: 1e18

Default: 1e18

Description: The minimum viscosity that is allowed in the whole model domain. Units: Pa s.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Phase transition Clapeyron slopes

Value:

Default:

Description: A list of Clapeyron slopes for each phase transition. A positive Clapeyron slope indicates that the phase transition will occur in a greater depth, if the temperature is higher than the one given in Phase transition temperatures and in a smaller depth, if the temperature is smaller than the one given in Phase transition temperatures. For negative slopes the other way round. List must have the same number of entries as Phase transition depths. Units: Pa/K.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Phase transition depths

Value:

Default:

Description: A list of depths where phase transitions occur. Values must monotonically increase.
Units: m.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Phase transition temperatures

Value:

Default:

Description: A list of temperatures where phase transitions occur. Higher or lower temperatures lead to phase transition occurring in smaller or greater depths than given in Phase transition depths, depending on the Clapeyron slope given in Phase transition Clapeyron slopes. List must have the same number of entries as Phase transition depths. Units: K.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Phase transition widths

Value:

Default:

Description: A list of widths for each phase transition. This is only use to specify the region where the recrystallized grain size is assigned after material has crossed a phase transition and should accordingly be chosen similar to the maximum cell width expected at the phase transition. List must have the same number of entries as Phase transition depths. Units: m.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Reciprocal required strain

Value: 10

Default: 10

Description: This parameters λ gives an estimate of the strain necessary to achieve a new grain size.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Recrystallized grain size

Value:

Default:

Description: The grain size d_{ph} to that a phase will be reduced to when crossing a phase transition. When set to zero, grain size will not be reduced. Units: m.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Reference compressibility

Value: 4e-12

Default: 4e-12

Description: The value of the reference compressibility. Units: $1/Pa$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference density

Value: 3300

Default: 3300

Description: Reference density ρ_0 . Units: kg/m^3 .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference specific heat

Value: 1250

Default: 1250

Description: The value of the specific heat cp . Units: $J/kg/K$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference temperature

Value: 293

Default: 293

Description: The reference temperature T_0 . Units: K .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Thermal conductivity

Value: 4.7

Default: 4.7

Description: The value of the thermal conductivity k . Units: $W/m/K$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Thermal expansion coefficient

Value: 2e-5

Default: 2e-5

Description: The value of the thermal expansion coefficient α . Units: $1/K$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Use enthalpy for material properties

Value: true

Default: true

Description: Whether to use the enthalpy to calculate thermal expansivity and specific heat (if true) or use the thermal expansivity and specific heat values from the material properties table directly (if false).

Possible values: A boolean value (true or false)

- *Parameter name:* Use paleowattmeter

Value: true

Default: true

Description: A flag indicating whether the computation should use the paleowattmeter approach of Austin and Evans (2007) for grain size reduction in the dislocation creep regime (if true) or the paleopiezometer approach from Hall and Parmetier (2003) (if false).

Possible values: A boolean value (true or false)

- *Parameter name:* Use_table_properties

Value: false

Default: false

Description: Whether to use the table properties also for density, thermal expansivity and specific heat. If false the properties are generated as in the simple compressible plugin.

Possible values: A boolean value (true or false)

- *Parameter name:* Viscosity

Value: 5e24

Default: 5e24

Description: The value of the constant viscosity. Units: $kg/m/s$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Work_fraction_for_boundary_area_change

Value: 0.1

Default: 0.1

Description: The fraction χ of work done by dislocation creep to change the grain boundary area. Units: J/m^2 .

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

A.95 Parameters in section Material model/Latent heat

- *Parameter name:* Composition_viscosity_prefactor

Value: 1.0

Default: 1.0

Description: A linear dependency of viscosity on composition. Dimensionless prefactor.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Compressibility

Value: 5.124e-12

Default: 5.124e-12

Description: The value of the compressibility κ . Units: $1/Pa$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Corresponding_phase_for_density_jump

Value:

Default:

Description: A list of phases, which correspond to the Phase transition density jumps. The density jumps occur only in the phase that is given by this phase value. 0 stands for the 1st compositional

fields, 1 for the second compositional field and -1 for none of them. List must have the same number of entries as Phase transition depths. Units: Pa/K .

Possible values: A list of 0 to 4294967295 elements where each element is [An integer n such that $0 \leq n \leq 2147483647$]

- *Parameter name:* Define transition by depth instead of pressure

Value: true

Default: true

Description: Whether to list phase transitions by depth or pressure. If this parameter is true, then the input file will use Phase transitions depths and Phase transition widths to define the phase transition. If it is false, the parameter file will read in phase transition data from Phase transition pressures and Phase transition pressure widths.

Possible values: A boolean value (true or false)

- *Parameter name:* Density differential for compositional field 1

Value: 0

Default: 0

Description: If compositional fields are used, then one would frequently want to make the density depend on these fields. In this simple material model, we make the following assumptions: if no compositional fields are used in the current simulation, then the density is simply the usual one with its linear dependence on the temperature. If there are compositional fields, then the density only depends on the first one in such a way that the density has an additional term of the kind $+\Delta\rho c_1(\mathbf{x})$. This parameter describes the value of $\Delta\rho$. Units: $\text{kg}/\text{m}^3/\text{unit change in composition}$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Maximum viscosity

Value: 1e24

Default: 1e24

Description: Limit for the maximum viscosity in the model. Units: Pa s.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Minimum viscosity

Value: 1e19

Default: 1e19

Description: Limit for the minimum viscosity in the model. Units: Pa s.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Phase transition Clapeyron slopes

Value:

Default:

Description: A list of Clapeyron slopes for each phase transition. A positive Clapeyron slope indicates that the phase transition will occur in a greater depth, if the temperature is higher than the one given in Phase transition temperatures and in a smaller depth, if the temperature is smaller than the one given in Phase transition temperatures. For negative slopes the other way round. List must have the same number of entries as Phase transition depths. Units: Pa/K .

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Phase transition density jumps

Value:

Default:

Description: A list of density jumps at each phase transition. A positive value means that the density increases with depth. The corresponding entry in Corresponding phase for density jump determines if the density jump occurs in peridotite, eclogite or none of them. List must have the same number of entries as Phase transition depths. Units: kg/m^3 .

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Phase transition depths

Value:

Default:

Description: A list of depths where phase transitions occur. Values must monotonically increase. Units: m .

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Phase transition pressure widths

Value:

Default:

Description: A list of widths for each phase transition, in terms of pressure. The phase functions are scaled with these values, leading to a jump between phases for a value of zero and a gradual transition for larger values. List must have the same number of entries as Phase transition pressures. Define transition by depth instead of pressure must be set to false to use this parameter. Units: Pa .

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Phase transition pressures

Value:

Default:

Description: A list of pressures where phase transitions occur. Values must monotonically increase. Define transition by depth instead of pressure must be set to false to use this parameter. Units: Pa .

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Phase transition temperatures

Value:

Default:

Description: A list of temperatures where phase transitions occur. Higher or lower temperatures lead to phase transition occurring in smaller or greater depths than given in Phase transition depths, depending on the Clapeyron slope given in Phase transition Clapeyron slopes. List must have the same number of entries as Phase transition depths. Units: K .

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Phase transition widths

Value:

Default:

Description: A list of widths for each phase transition, in terms of depth. The phase functions are scaled with these values, leading to a jump between phases for a value of zero and a gradual transition for larger values. List must have the same number of entries as Phase transition depths. Units: m.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Reference density

Value: 3300

Default: 3300

Description: Reference density ρ_0 . Units: kg/m³.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference specific heat

Value: 1250

Default: 1250

Description: The value of the specific heat C_p . Units: J/kg/K.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference temperature

Value: 293

Default: 293

Description: The reference temperature T_0 . Units: K.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Thermal conductivity

Value: 2.38

Default: 2.38

Description: The value of the thermal conductivity k . Units: W/m/K.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Thermal expansion coefficient

Value: 4e-5

Default: 4e-5

Description: The value of the thermal expansion coefficient β . Units: 1/K.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Thermal viscosity exponent

Value: 0.0

Default: 0.0

Description: The temperature dependence of viscosity. Dimensionless exponent.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* **Viscosity**

Value: 5e24

Default: 5e24

Description: The value of the constant viscosity. Units: $kg/m/s$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* **Viscosity prefactors**

Value:

Default:

Description: A list of prefactors for the viscosity for each phase. The reference viscosity will be multiplied by this factor to get the corresponding viscosity for each phase. List must have one more entry than Phase transition depths. Units: non-dimensional.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

A.96 Parameters in section Material model/Latent heat melt

- *Parameter name:* **A1**

Value: 1085.7

Default: 1085.7

Description: Constant parameter in the quadratic function that approximates the solidus of peridotite. Units: K .

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* **A2**

Value: 1.329e-7

Default: 1.329e-7

Description: Prefactor of the linear pressure term in the quadratic function that approximates the solidus of peridotite. Units: K/Pa .

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* **A3**

Value: -5.1e-18

Default: -5.1e-18

Description: Prefactor of the quadratic pressure term in the quadratic function that approximates the solidus of peridotite. Units: $\text{K}/(\text{Pa}^2)$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* **B1**

Value: 1475.0

Default: 1475.0

Description: Constant parameter in the quadratic function that approximates the lherzolite liquidus used for calculating the fraction of peridotite-derived melt. Units: K .

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* B2

Value: 8.0e-8

Default: 8.0e-8

Description: Prefactor of the linear pressure term in the quadratic function that approximates the lherzolite liquidus used for calculating the fraction of peridotite-derived melt. Units: $\text{Å}^2\text{C}/\text{Pa}$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* B3

Value: -3.2e-18

Default: -3.2e-18

Description: Prefactor of the quadratic pressure term in the quadratic function that approximates the lherzolite liquidus used for calculating the fraction of peridotite-derived melt. Units: $\text{Å}^2\text{C}/(\text{Pa}^2)$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* C1

Value: 1780.0

Default: 1780.0

Description: Constant parameter in the quadratic function that approximates the liquidus of peridotite. Units: $\text{Å}^2\text{C}$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* C2

Value: 4.50e-8

Default: 4.50e-8

Description: Prefactor of the linear pressure term in the quadratic function that approximates the liquidus of peridotite. Units: $\text{Å}^2\text{C}/\text{Pa}$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* C3

Value: -2.0e-18

Default: -2.0e-18

Description: Prefactor of the quadratic pressure term in the quadratic function that approximates the liquidus of peridotite. Units: $\text{Å}^2\text{C}/(\text{Pa}^2)$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Composition viscosity prefactor

Value: 1.0

Default: 1.0

Description: A linear dependency of viscosity on composition. Dimensionless prefactor.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Compressibility

Value: 5.124e-12

Default: 5.124e-12

Description: The value of the compressibility κ . Units: $1/\text{Pa}$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: D1

Value: 976.0

Default: 976.0

Description: Constant parameter in the quadratic function that approximates the solidus of pyroxenite. Units: K .

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: D2

Value: 1.329e-7

Default: 1.329e-7

Description: Prefactor of the linear pressure term in the quadratic function that approximates the solidus of pyroxenite. Note that this factor is different from the value given in Sobolev, 2011, because they use the potential temperature whereas we use the absolute temperature. Units: K/Pa .

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: D3

Value: -5.1e-18

Default: -5.1e-18

Description: Prefactor of the quadratic pressure term in the quadratic function that approximates the solidus of pyroxenite. Units: $\text{K}/(\text{Pa}^2)$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: Density differential for compositional field 1

Value: 0

Default: 0

Description: If compositional fields are used, then one would frequently want to make the density depend on these fields. In this simple material model, we make the following assumptions: if no compositional fields are used in the current simulation, then the density is simply the usual one with its linear dependence on the temperature. If there are compositional fields, then the density only depends on the first one in such a way that the density has an additional term of the kind $+\Delta\rho c_1(\mathbf{x})$. This parameter describes the value of $\Delta\rho$. Units: $\text{kg}/\text{m}^3/\text{unit change in composition}$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: E1

Value: 663.8

Default: 663.8

Description: Prefactor of the linear depletion term in the quadratic function that approximates the melt fraction of pyroxenite. Units: K/Pa .

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: E2

Value: -611.4

Default: -611.4

Description: Prefactor of the quadratic depletion term in the quadratic function that approximates the melt fraction of pyroxenite. Units: $\text{K}/(\text{Pa}^2)$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Mass fraction cpx

Value: 0.15

Default: 0.15

Description: Mass fraction of clinopyroxene in the peridotite to be molten. Units: non-dimensional.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Maximum pyroxenite melt fraction

Value: 0.5429

Default: 0.5429

Description: Maximum melt fraction of pyroxenite in this parameterization. At higher temperatures peridotite begins to melt.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Peridotite melting entropy change

Value: -300

Default: -300

Description: The entropy change for the phase transition from solid to melt of peridotite. Units: $J/(kgK)$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Pyroxenite melting entropy change

Value: -400

Default: -400

Description: The entropy change for the phase transition from solid to melt of pyroxenite. Units: $J/(kgK)$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Reference density

Value: 3300

Default: 3300

Description: Reference density ρ_0 . Units: kg/m^3 .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Reference specific heat

Value: 1250

Default: 1250

Description: The value of the specific heat C_p . Units: $J/kg/K$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Reference temperature

Value: 293

Default: 293

Description: The reference temperature T_0 . Units: K .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Relative density of melt
Value: 0.9
Default: 0.9
Description: The relative density of melt compared to the solid material. This means, the density change upon melting is this parameter times the density of solid material. Units: non-dimensional.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$
- **Parameter name:** Thermal conductivity
Value: 2.38
Default: 2.38
Description: The value of the thermal conductivity k . Units: $W/m/K$.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- **Parameter name:** Thermal expansion coefficient
Value: 4e-5
Default: 4e-5
Description: The value of the thermal expansion coefficient α_s . Units: $1/K$.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- **Parameter name:** Thermal expansion coefficient of melt
Value: 6.8e-5
Default: 6.8e-5
Description: The value of the thermal expansion coefficient α_f . Units: $1/K$.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- **Parameter name:** Thermal viscosity exponent
Value: 0.0
Default: 0.0
Description: The temperature dependence of viscosity. Dimensionless exponent.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- **Parameter name:** Viscosity
Value: 5e24
Default: 5e24
Description: The value of the constant viscosity. Units: $kg/m/s$.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- **Parameter name:** beta
Value: 1.5
Default: 1.5
Description: Exponent of the melting temperature in the melt fraction calculation. Units: non-dimensional.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `r1`

Value: 0.5

Default: 0.5

Description: Constant in the linear function that approximates the clinopyroxene reaction coefficient.
Units: non-dimensional.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `r2`

Value: 8e-11

Default: 8e-11

Description: Prefactor of the linear pressure term in the linear function that approximates the clinopyroxene reaction coefficient. Units: $1/\text{Pa}$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

A.97 Parameters in section Material model/Melt global

- *Parameter name:* Depletion density change

Value: 0.0

Default: 0.0

Description: The density contrast between material with a depletion of 1 and a depletion of zero. Negative values indicate lower densities of depleted material. Depletion is indicated by the compositional field with the name peridotite. Not used if this field does not exist in the model. Units: kg/m^3 .

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Depletion solidus change

Value: 200.0

Default: 200.0

Description: The solidus temperature change for a depletion of 100%. For positive values, the solidus gets increased for a positive peridotite field (depletion) and lowered for a negative peridotite field (enrichment). Scaling with depletion is linear. Only active when fractional melting is used. Units: K .

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Exponential melt weakening factor

Value: 27

Default: 27

Description: The porosity dependence of the viscosity. Units: dimensionless.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Include melting and freezing

Value: true

Default: true

Description: Whether to include melting and freezing (according to a simplified linear melting approximation in the model (if true), or not (if false)).

Possible values: A boolean value (true or false)

- **Parameter name:** Melt bulk modulus derivative

Value: 0.0

Default: 0.0

Description: The value of the pressure derivative of the melt bulk modulus. Units: None.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Melt compressibility

Value: 0.0

Default: 0.0

Description: The value of the compressibility of the melt. Units: $1/\text{Pa}$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Melting time scale for operator splitting

Value: 1e3

Default: 1e3

Description: In case the operator splitting scheme is used, the porosity field can not be set to a new equilibrium melt fraction instantly, but the model has to provide a melting time scale instead. This time scale defines how fast melting happens, or more specifically, the parameter defines the time after which the deviation of the porosity from the equilibrium melt fraction will be reduced to a fraction of $1/e$. So if the melting time scale is small compared to the time step size, the reaction will be so fast that the porosity is very close to the equilibrium melt fraction after reactions are computed. Conversely, if the melting time scale is large compared to the time step size, almost no melting and freezing will occur.

Also note that the melting time scale has to be larger than or equal to the reaction time step used in the operator splitting scheme, otherwise reactions can not be computed. If the model does not use operator splitting, this parameter is not used. Units: yr or s, depending on the “Use years in output instead of seconds” parameter.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Pressure solidus change

Value: 6e-8

Default: 6e-8

Description: The linear solidus temperature change with pressure. For positive values, the solidus gets increased for positive pressures. Units: $1/\text{Pa}$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Reference bulk viscosity

Value: 1e22

Default: 1e22

Description: The value of the constant bulk viscosity ξ_0 of the solid matrix. This viscosity may be modified by both temperature and porosity dependencies. Units: Pas.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Reference melt density

Value: 2500

Default: 2500

Description: Reference density of the melt/fluid $\rho_{f,0}$. Units: kg/m^3 .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference melt viscosity

Value: 10

Default: 10

Description: The value of the constant melt viscosity η_f . Units: *Pas.*

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference permeability

Value: 1e-8

Default: 1e-8

Description: Reference permeability of the solid host rock. Units: m^2 .

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference shear viscosity

Value: 5e20

Default: 5e20

Description: The value of the constant viscosity η_0 of the solid matrix. This viscosity may be modified by both temperature and porosity dependencies. Units: *Pas.*

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference solid density

Value: 3000

Default: 3000

Description: Reference density of the solid $\rho_{s,0}$. Units: kg/m^3 .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference specific heat

Value: 1250

Default: 1250

Description: The value of the specific heat C_p . Units: $J/kg/K$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference temperature

Value: 293

Default: 293

Description: The reference temperature T_0 . The reference temperature is used in both the density and viscosity formulas. Units: K .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Solid compressibility

Value: 0.0

Default: 0.0

Description: The value of the compressibility of the solid matrix. Units: $1/Pa$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Surface solidus
Value: 1300
Default: 1300
Description: Solidus for a pressure of zero. Units: K.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Thermal bulk viscosity exponent
Value: 0.0
Default: 0.0
Description: The temperature dependence of the bulk viscosity. Dimensionless exponent. See the general documentation of this model for a formula that states the dependence of the viscosity on this factor, which is called β there.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Thermal conductivity
Value: 4.7
Default: 4.7
Description: The value of the thermal conductivity k . Units: W/m/K.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Thermal expansion coefficient
Value: 2e-5
Default: 2e-5
Description: The value of the thermal expansion coefficient β . Units: 1/K.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Thermal viscosity exponent
Value: 0.0
Default: 0.0
Description: The temperature dependence of the shear viscosity. Dimensionless exponent. See the general documentation of this model for a formula that states the dependence of the viscosity on this factor, which is called β there.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.98 Parameters in section Material model/Melt simple

- *Parameter name:* A1
Value: 1085.7
Default: 1085.7
Description: Constant parameter in the quadratic function that approximates the solidus of peridotite. Units: °C.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* A2

Value: 1.329e-7

Default: 1.329e-7

Description: Prefactor of the linear pressure term in the quadratic function that approximates the solidus of peridotite. Units: $^{\circ}\text{C}/\text{Pa}$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* A3

Value: -5.1e-18

Default: -5.1e-18

Description: Prefactor of the quadratic pressure term in the quadratic function that approximates the solidus of peridotite. Units: $^{\circ}\text{C}/(\text{Pa}^2)$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* B1

Value: 1475.0

Default: 1475.0

Description: Constant parameter in the quadratic function that approximates the lherzolite liquidus used for calculating the fraction of peridotite-derived melt. Units: $^{\circ}\text{C}$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* B2

Value: 8.0e-8

Default: 8.0e-8

Description: Prefactor of the linear pressure term in the quadratic function that approximates the lherzolite liquidus used for calculating the fraction of peridotite-derived melt. Units: $^{\circ}\text{C}/\text{Pa}$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* B3

Value: -3.2e-18

Default: -3.2e-18

Description: Prefactor of the quadratic pressure term in the quadratic function that approximates the lherzolite liquidus used for calculating the fraction of peridotite-derived melt. Units: $^{\circ}\text{C}/(\text{Pa}^2)$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* C1

Value: 1780.0

Default: 1780.0

Description: Constant parameter in the quadratic function that approximates the liquidus of peridotite. Units: $^{\circ}\text{C}$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* C2

Value: 4.50e-8

Default: 4.50e-8

Description: Prefactor of the linear pressure term in the quadratic function that approximates the liquidus of peridotite. Units: $^{\circ}C/Pa$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* C3

Value: -2.0e-18

Default: -2.0e-18

Description: Prefactor of the quadratic pressure term in the quadratic function that approximates the liquidus of peridotite. Units: $^{\circ}C/(Pa^2)$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Depletion density change

Value: 0.0

Default: 0.0

Description: The density contrast between material with a depletion of 1 and a depletion of zero. Negative values indicate lower densities of depleted material. Depletion is indicated by the compositional field with the name peridotite. Not used if this field does not exist in the model. Units: kg/m^3 .

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Depletion solidus change

Value: 200.0

Default: 200.0

Description: The solidus temperature change for a depletion of 100%. For positive values, the solidus gets increased for a positive peridotite field (depletion) and lowered for a negative peridotite field (enrichment). Scaling with depletion is linear. Only active when fractional melting is used. Units: K .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Exponential melt weakening factor

Value: 27

Default: 27

Description: The porosity dependence of the viscosity. Units: dimensionless.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Freezing rate

Value: 0.0

Default: 0.0

Description: Freezing rate of melt when in subsolidus regions. If this parameter is set to a number larger than 0.0, it specifies the fraction of melt that will freeze per year (or per second, depending on the “Use years in output instead of seconds” parameter), as soon as the porosity exceeds the equilibrium melt fraction, and the equilibrium melt fraction falls below the depletion. In this case, melt will freeze according to the given rate until one of those conditions is not fulfilled anymore. The reasoning behind this is that there should not be more melt present than the equilibrium melt fraction, as melt production decreases with increasing depletion, but the freezing process of melt also reduces the depletion by the same amount, and as soon as the depletion falls below the equilibrium melt fraction, we expect that material should melt again (no matter how much melt is present). This is quite a simplification and not a realistic freezing parameterization, but without tracking the melt composition, there is no way to compute freezing rates accurately. If this parameter is set to zero, no freezing will occur. Note

that freezing can never be faster than determined by the “Melting time scale for operator splitting”. The product of the “Freezing rate” and the “Melting time scale for operator splitting” defines how fast freezing occurs with respect to melting (if the product is 0.5, melting will occur twice as fast as freezing). Units: 1/yr or 1/s, depending on the “Use years in output instead of seconds” parameter.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `Mass fraction cpx`

Value: 0.15

Default: 0.15

Description: Mass fraction of clinopyroxene in the peridotite to be molten. Units: non-dimensional.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `Melt bulk modulus derivative`

Value: 0.0

Default: 0.0

Description: The value of the pressure derivative of the melt bulk modulus. Units: None.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `Melt compressibility`

Value: 0.0

Default: 0.0

Description: The value of the compressibility of the melt. Units: $1/\text{Pa}$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `Melt extraction depth`

Value: 1000.0

Default: 1000.0

Description: Depth above that melt will be extracted from the model, which is done by a negative reaction term proportional to the porosity field. Units: m .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `Melting time scale for operator splitting`

Value: 1e3

Default: 1e3

Description: Because the operator splitting scheme is used, the porosity field can not be set to a new equilibrium melt fraction instantly, but the model has to provide a melting time scale instead. This time scale defines how fast melting happens, or more specifically, the parameter defines the time after which the deviation of the porosity from the equilibrium melt fraction will be reduced to a fraction of $1/e$. So if the melting time scale is small compared to the time step size, the reaction will be so fast that the porosity is very close to the equilibrium melt fraction after reactions are computed. Conversely, if the melting time scale is large compared to the time step size, almost no melting and freezing will occur.

Also note that the melting time scale has to be larger than or equal to the reaction time step used in the operator splitting scheme, otherwise reactions can not be computed. Units: yr or s, depending on the “Use years in output instead of seconds” parameter.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Peridotite melting entropy change
Value: -300
Default: -300
Description: The entropy change for the phase transition from solid to melt of peridotite. Units: $J/(kgK)$.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$
- **Parameter name:** Reference bulk viscosity
Value: 1e22
Default: 1e22
Description: The value of the constant bulk viscosity ξ_0 of the solid matrix. This viscosity may be modified by both temperature and porosity dependencies. Units: *Pas*.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- **Parameter name:** Reference melt density
Value: 2500
Default: 2500
Description: Reference density of the melt/fluid $\rho_{f,0}$. Units: kg/m^3 .
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- **Parameter name:** Reference melt viscosity
Value: 10
Default: 10
Description: The value of the constant melt viscosity η_f . Units: *Pas*.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- **Parameter name:** Reference permeability
Value: 1e-8
Default: 1e-8
Description: Reference permeability of the solid host rock. Units: m^2 .
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$
- **Parameter name:** Reference shear viscosity
Value: 5e20
Default: 5e20
Description: The value of the constant viscosity η_0 of the solid matrix. This viscosity may be modified by both temperature and porosity dependencies. Units: *Pas*.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- **Parameter name:** Reference solid density
Value: 3000
Default: 3000
Description: Reference density of the solid $\rho_{s,0}$. Units: kg/m^3 .
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference specific heat

Value: 1250

Default: 1250

Description: The value of the specific heat C_p . Units: $J/kg/K$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference temperature

Value: 293

Default: 293

Description: The reference temperature T_0 . The reference temperature is used in both the density and viscosity formulas. Units: K .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Solid compressibility

Value: 0.0

Default: 0.0

Description: The value of the compressibility of the solid matrix. Units: $1/Pa$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Thermal bulk viscosity exponent

Value: 0.0

Default: 0.0

Description: The temperature dependence of the bulk viscosity. Dimensionless exponent. See the general documentation of this model for a formula that states the dependence of the viscosity on this factor, which is called β there.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Thermal conductivity

Value: 4.7

Default: 4.7

Description: The value of the thermal conductivity k . Units: $W/m/K$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Thermal expansion coefficient

Value: 2e-5

Default: 2e-5

Description: The value of the thermal expansion coefficient β . Units: $1/K$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Thermal viscosity exponent

Value: 0.0

Default: 0.0

Description: The temperature dependence of the shear viscosity. Dimensionless exponent. See the general documentation of this model for a formula that states the dependence of the viscosity on this factor, which is called β there.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `Use fractional melting`

Value: false

Default: false

Description: If fractional melting should be used (if true), including a solidus change based on depletion (in this case, the amount of melt that has migrated away from its origin), and freezing of melt when it has moved to a region with temperatures lower than the solidus; or if batch melting should be used (if false), assuming that the melt fraction only depends on temperature and pressure, and how much melt has already been generated at a given point, but not considering movement of melt in the melting parameterization.

Note that melt does not freeze unless the 'Freezing rate' parameter is set to a value larger than 0.

Possible values: A boolean value (true or false)

- *Parameter name:* `Use full compressibility`

Value: false

Default: false

Description: If the compressibility should be used everywhere in the code (if true), changing the volume of material when the density changes, or only in the momentum conservation and advection equations (if false).

Possible values: A boolean value (true or false)

- *Parameter name:* `beta`

Value: 1.5

Default: 1.5

Description: Exponent of the melting temperature in the melt fraction calculation. Units: non-dimensional.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `r1`

Value: 0.5

Default: 0.5

Description: Constant in the linear function that approximates the clinopyroxene reaction coefficient. Units: non-dimensional.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `r2`

Value: 8e-11

Default: 8e-11

Description: Prefactor of the linear pressure term in the linear function that approximates the clinopyroxene reaction coefficient. Units: $1/\text{Pa}$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

A.99 Parameters in section Material model/Multicomponent

- Parameter name: **Densities**

Value: 3300.

Default: 3300.

Description: List of densities for background mantle and compositional fields,for a total of N+1 values, where N is the number of compositional fields.If only one value is given, then all use the same value. Units: kg/m^3

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- Parameter name: **Reference temperature**

Value: 293

Default: 293

Description: The reference temperature T_0 . Units: K .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: **Specific heats**

Value: 1250.

Default: 1250.

Description: List of specific heats C_p for background mantle and compositional fields,for a total of N+1 values, where N is the number of compositional fields.If only one value is given, then all use the same value. Units: $J/kg/K$

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- Parameter name: **Thermal conductivities**

Value: 4.7

Default: 4.7

Description: List of thermal conductivities for background mantle and compositional fields,for a total of N+1 values, where N is the number of compositional fields.If only one value is given, then all use the same value. Units: $W/m/K$

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- Parameter name: **Thermal expansivities**

Value: 4.e-5

Default: 4.e-5

Description: List of thermal expansivities for background mantle and compositional fields,for a total of N+1 values, where N is the number of compositional fields.If only one value is given, then all use the same value. Units: $1/K$

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- Parameter name: **Viscosities**

Value: 1.e21

Default: 1.e21

Description: List of viscosities for background mantle and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: Pas

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Viscosity averaging scheme

Value: harmonic

Default: harmonic

Description: When more than one compositional field is present at a point with different viscosities, we need to come up with an average viscosity at that point. Select a weighted harmonic, arithmetic, geometric, or maximum composition.

Possible values: Any one of arithmetic, harmonic, geometric, maximum composition

A.100 Parameters in section Material model/Nondimensional model

- *Parameter name:* Di

Value: 0.0

Default: 0.0

Description: Dissipation number. Pick 0.0 for incompressible computations.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Ra

Value: 1e4

Default: 1e4

Description: Rayleigh number Ra

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference density

Value: 1.0

Default: 1.0

Description: Reference density ρ_0 . Units: kg/m^3 .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference specific heat

Value: 1.0

Default: 1.0

Description: The value of the specific heat C_p . Units: $\text{J}/\text{kg}/\text{K}$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Use TALA

Value: false

Default: false

Description: Whether to use the TALA instead of the ALA approximation.

Possible values: A boolean value (true or false)

- *Parameter name:* Viscosity depth prefactor
Value: 0.0
Default: 0.0
Description: Exponential depth prefactor for viscosity.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Viscosity temperature prefactor
Value: 0.0
Default: 0.0
Description: Exponential temperature prefactor for viscosity.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* gamma
Value: 1.0
Default: 1.0
Description: Grueneisen parameter
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.101 Parameters in section Material model/Simple compressible model

- *Parameter name:* Reference compressibility
Value: 4e-12
Default: 4e-12
Description: The value of the reference compressibility. Units: $1/\text{Pa}$.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Reference density
Value: 3300
Default: 3300
Description: Reference density ρ_0 . Units: kg/m^3 .
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Reference specific heat
Value: 1250
Default: 1250
Description: The value of the specific heat C_p . Units: $\text{J}/\text{kg}/\text{K}$.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Thermal conductivity
Value: 4.7
Default: 4.7
Description: The value of the thermal conductivity k . Units: $\text{W}/\text{m}/\text{K}$.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Thermal expansion coefficient
Value: 2e-5
Default: 2e-5
Description: The value of the thermal expansion coefficient α . Units: $1/K$.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Viscosity
Value: 1e21
Default: 1e21
Description: The value of the constant viscosity η_0 . Units: $kg/m/s$.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.102 Parameters in section Material model/Simple model

- *Parameter name:* Composition viscosity prefactor
Value: 1.0
Default: 1.0
Description: A linear dependency of viscosity on the first compositional field. Dimensionless prefactor. With a value of 1.0 (the default) the viscosity does not depend on the composition. See the general documentation of this model for a formula that states the dependence of the viscosity on this factor, which is called ξ there.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Density differential for compositional field 1
Value: 0
Default: 0
Description: If compositional fields are used, then one would frequently want to make the density depend on these fields. In this simple material model, we make the following assumptions: if no compositional fields are used in the current simulation, then the density is simply the usual one with its linear dependence on the temperature. If there are compositional fields, then the density only depends on the first one in such a way that the density has an additional term of the kind $+\Delta\rho c_1(\mathbf{x})$. This parameter describes the value of $\Delta\rho$. Units: $kg/m^3/\text{unit change in composition}$.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Maximum thermal prefactor
Value: 1.0e2
Default: 1.0e2
Description: The maximum value of the viscosity prefactor associated with temperature dependence.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Minimum thermal prefactor
Value: 1.0e-2
Default: 1.0e-2
Description: The minimum value of the viscosity prefactor associated with temperature dependence.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Reference density

Value: 3300

Default: 3300

Description: Reference density ρ_0 . Units: kg/m^3 .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Reference specific heat

Value: 1250

Default: 1250

Description: The value of the specific heat C_p . Units: $J/kg/K$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Reference temperature

Value: 293

Default: 293

Description: The reference temperature T_0 . The reference temperature is used in both the density and viscosity formulas. Units: K .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Thermal conductivity

Value: 4.7

Default: 4.7

Description: The value of the thermal conductivity k . Units: $W/m/K$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Thermal expansion coefficient

Value: 2e-5

Default: 2e-5

Description: The value of the thermal expansion coefficient α . Units: $1/K$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Thermal viscosity exponent

Value: 0.0

Default: 0.0

Description: The temperature dependence of viscosity. Dimensionless exponent. See the general documentation of this model for a formula that states the dependence of the viscosity on this factor, which is called β there.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Viscosity

Value: 5e24

Default: 5e24

Description: The value of the constant viscosity η_0 . This viscosity may be modified by both temperature and compositional dependencies. Units: $kg/m/s$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.103 Parameters in section Material model/Simpler model

- *Parameter name:* Reference density

Value: 3300

Default: 3300

Description: Reference density ρ_0 . Units: kg/m^3 .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference specific heat

Value: 1250

Default: 1250

Description: The value of the specific heat C_p . Units: $J/kg/K$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reference temperature

Value: 293

Default: 293

Description: The reference temperature T_0 . The reference temperature is used in the density formula. Units: K .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Thermal conductivity

Value: 4.7

Default: 4.7

Description: The value of the thermal conductivity k . Units: $W/m/K$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Thermal expansion coefficient

Value: 2e-5

Default: 2e-5

Description: The value of the thermal expansion coefficient β . Units: $1/K$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Viscosity

Value: 5e24

Default: 5e24

Description: The value of the viscosity η . Units: $kg/m/s$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.104 Parameters in section Material model/Steinberger model

- *Parameter name:* Bilinear interpolation

Value: true

Default: true

Description: Whether to use bilinear interpolation to compute material properties (slower but more accurate).

Possible values: A boolean value (true or false)

- *Parameter name:* Data directory

Value: \$ASPECT_SOURCE_DIR/data/material-model/steinberger/

Default: \$ASPECT_SOURCE_DIR/data/material-model/steinberger/

Description: The path to the model data. The path may also include the special text '\$ASPECT_SOURCE_DIR' which will be interpreted as the path in which the ASPECT source files were located when ASPECT was compiled. This interpretation allows, for example, to reference files located in the 'data/' subdirectory of ASPECT.

Possible values: A directory name

- *Parameter name:* Latent heat

Value: false

Default: false

Description: Whether to include latent heat effects in the calculation of thermal expansivity and specific heat. Following the approach of Nakagawa et al. 2009.

Possible values: A boolean value (true or false)

- *Parameter name:* Lateral viscosity file name

Value: temp-viscosity-prefactor.txt

Default: temp-viscosity-prefactor.txt

Description: The file name of the lateral viscosity data.

Possible values: Any string

- *Parameter name:* Material file names

Value: pyr-ringwood88.txt

Default: pyr-ringwood88.txt

Description: The file names of the material data (material data is assumed to be in order with the ordering of the compositional fields). Note that there are three options on how many files need to be listed here: 1. If only one file is provided, it is used for the whole model domain, and compositional fields are ignored. 2. If there is one more file name than the number of compositional fields, then the first file is assumed to define a 'background composition' that is modified by the compositional fields. If there are exactly as many files as compositional fields, the fields are assumed to represent the fractions of different materials and the average property is computed as a sum of the value of the compositional field times the material property of that field.

Possible values: A list of 0 to 4294967295 elements where each element is [Any string]

- *Parameter name:* Maximum lateral viscosity variation

Value: 1e2

Default: 1e2

Description: The relative cutoff value for lateral viscosity variations caused by temperature deviations. The viscosity may vary laterally by this factor squared.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Maximum viscosity

Value: 1e23

Default: 1e23

Description: The maximum viscosity that is allowed in the viscosity calculation. Larger values will be cut off.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Minimum viscosity

Value: 1e19

Default: 1e19

Description: The minimum viscosity that is allowed in the viscosity calculation. Smaller values will be cut off.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Number lateral average bands

Value: 10

Default: 10

Description: Number of bands to compute laterally averaged temperature within.

Possible values: An integer n such that $1 \leq n \leq 2147483647$

- *Parameter name:* Radial viscosity file name

Value: radial-visc.txt

Default: radial-visc.txt

Description: The file name of the radial viscosity data.

Possible values: Any string

- *Parameter name:* Reference viscosity

Value: 1e23

Default: 1e23

Description: The reference viscosity that is used for pressure scaling.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Thermal conductivity

Value: 4.7

Default: 4.7

Description: The value of the thermal conductivity k . Units: $W/m/K$.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Use lateral average temperature for viscosity

Value: true

Default: true

Description: Whether to use to use the laterally averaged temperature instead of the adiabatic temperature as reference for the viscosity calculation. This ensures that the laterally averaged viscosities remain more or less constant over the model runtime. This behaviour might or might not be desired.

Possible values: A boolean value (true or false)

A.105 Parameters in section Material model/Visco Plastic

- *Parameter name:* Activation energies for diffusion creep

Value: 375e3

Default: 375e3

Description: List of activation energies, E_a , for background material and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: J/mol

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Activation energies for dislocation creep

Value: 530e3

Default: 530e3

Description: List of activation energies, E_a , for background material and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: J/mol

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Activation volumes for diffusion creep

Value: 6e-6

Default: 6e-6

Description: List of activation volumes, V_a , for background material and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: m³/mol

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Activation volumes for dislocation creep

Value: 1.4e-5

Default: 1.4e-5

Description: List of activation volumes, V_a , for background material and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: m³/mol

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Angles of internal friction

Value: 0

Default: 0

Description: List of angles of internal friction, ϕ , for background material and compositional fields, for a total of N+1 values, where N is the number of compositional fields. For a value of zero, in 2D the von Mises criterion is retrieved. Angles higher than 30 degrees are harder to solve numerically. Units: degrees.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Cohesion strain weakening factors

Value: 1.

Default: 1.

Description: List of cohesion strain weakening factors for background material and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: None

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Cohesions

Value: 1e20

Default: 1e20

Description: List of cohesions, C , for background material and compositional fields, for a total of N+1 values, where N is the number of compositional fields. The extremely large default cohesion value (1e20 Pa) prevents the viscous stress from exceeding the yield stress. Units: Pa.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Densities

Value: 3300.

Default: 3300.

Description: List of densities, ρ , for background material and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: kg/m³

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* End strain weakening intervals

Value: 1.

Default: 1.

Description: List of strain weakening interval final strains for background material and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: None

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- **Parameter name:** Friction strain weakening factors

Value: 1.

Default: 1.

Description: List of friction strain weakening factors for background material and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: None

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- **Parameter name:** Grain size

Value: 1e-3

Default: 1e-3

Description: Units: m

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Grain size exponents for diffusion creep

Value: 3

Default: 3

Description: List of grain size exponents, $m_{\text{diffusion}}$, for background material and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: None

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- **Parameter name:** Heat capacities

Value: 1.25e3

Default: 1.25e3

Description: List of heat capacities C_p , for background material and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: J/kg/K

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- **Parameter name:** Maximum viscosity

Value: 1e28

Default: 1e28

Description: Upper cutoff for effective viscosity. Units: Pas

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Minimum strain rate

Value: 1.0e-20

Default: 1.0e-20

Description: Stabilizes strain dependent viscosity. Units: 1/s

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Minimum viscosity

Value: 1e17

Default: 1e17

Description: Lower cutoff for effective viscosity. Units: Pas

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Prefactors for diffusion creep

Value: 1.5e-15

Default: 1.5e-15

Description: List of viscosity prefactors, A , for background material and compositional fields, for a total of $N+1$ values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: $\text{Pa}^{-1}\text{m}^m\text{diffusion s}^{-1}$

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- **Parameter name:** Prefactors for dislocation creep

Value: 1.1e-16

Default: 1.1e-16

Description: List of viscosity prefactors, A , for background material and compositional fields, for a total of $N+1$ values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: $\text{Pa}^{-n_{\text{dislocation}}}\text{s}^{-1}$

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- **Parameter name:** Reference strain rate

Value: 1.0e-15

Default: 1.0e-15

Description: Reference strain rate for first time step. Units: 1/s

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Reference temperature

Value: 293

Default: 293

Description: For calculating density by thermal expansivity. Units: K

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Reference viscosity

Value: 1e22

Default: 1e22

Description: Reference viscosity for nondimensionalization. Units Pas

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Start strain weakening intervals

Value: 0.

Default: 0.

Description: List of strain weakening interval initial strains for background material and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: None

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Stress exponents for diffusion creep

Value: 1

Default: 1

Description: List of stress exponents, $n_{\text{diffusion}}$, for background material and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: None

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Stress exponents for dislocation creep

Value: 3.5

Default: 3.5

Description: List of stress exponents, $n_{\text{dislocation}}$, for background material and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: None

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Stress limiter exponents

Value: 1.0

Default: 1.0

Description: List of stress limiter exponents, n_{lim} , for background material and compositional fields, for a total of N+1 values, where N is the number of compositional fields. Units: none.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Thermal diffusivities

Value: 0.8e-6

Default: 0.8e-6

Description: List of thermal diffusivities, for background material and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: m^2/s

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Thermal expansivities

Value: 3.5e-5

Default: 3.5e-5

Description: List of thermal expansivities for background material and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: $1/K$

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* `Use finite strain tensor`

Value: false

Default: false

Description: Track and use the full finite strain tensor for strain weakening. Units: None

Possible values: A boolean value (true or false)

- *Parameter name:* `Use strain weakening`

Value: false

Default: false

Description: Apply strain weakening to viscosity, cohesion and internal angle of friction based on accumulated finite strain. Units: None

Possible values: A boolean value (true or false)

- *Parameter name:* `Viscosity averaging scheme`

Value: harmonic

Default: harmonic

Description: When more than one compositional field is present at a point with different viscosities, we need to come up with an average viscosity at that point. Select a weighted harmonic, arithmetic, geometric, or maximum composition.

Possible values: Any one of arithmetic, harmonic, geometric, maximum composition

- *Parameter name:* `Viscous flow law`

Value: composite

Default: composite

Description: Select what type of viscosity law to use between diffusion, dislocation and composite options. Soon there will be an option to select a specific flow law for each assigned composition

Possible values: Any one of diffusion, dislocation, composite

- *Parameter name:* `Viscous strain weakening factors`

Value: 1.

Default: 1.

Description: List of viscous strain weakening factors for background material and compositional fields, for a total of N+1 values, where N is the number of compositional fields. If only one value is given, then all use the same value. Units: None

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* `Yield mechanism`

Value: drucker

Default: drucker

Description: Select what type of yield mechanism to use between Drucker Prager and stress limiter options.

Possible values: Any one of drucker, limiter

A.106 Parameters in section Melt settings

- *Parameter name:* Average melt velocity

Value: true

Default: true

Description: Whether to cell-wise average the material properties that are used to compute the melt velocity or not. The melt velocity is computed as the sum of the solid velocity and the phase separation flux $-K_D/\phi(\nabla p_f - \rho_f \mathbf{g})$. If this parameter is set to true, K_D and ϕ will be averaged cell-wise in the computation of the phase separation flux. This is useful because in some models the melt velocity can have spikes close to the interface between regions of melt and no melt, as both K_D and ϕ go to zero for vanishing melt fraction. As the melt velocity is used for computing the time step size, and in models that use heat transport by melt or shear heating of melt, setting this parameter to true can speed up the model and make it more stable. In computations where accuracy and convergence behavior of the melt velocity is important (like in benchmark cases with an analytical solution), this parameter should probably be set to 'false'.

Possible values: A boolean value (true or false)

- *Parameter name:* Heat advection by melt

Value: false

Default: false

Description: Whether to use a porosity weighted average of the melt and solid velocity to advect heat in the temperature equation or not. If this is set to true, additional terms are assembled on the left-hand side of the temperature advection equation. Only used if Include melt transport is true. If this is set to false, only the solid velocity is used (as in models without melt migration).

Possible values: A boolean value (true or false)

- *Parameter name:* Include melt transport

Value: false

Default: false

Description: Whether to include the transport of melt into the model or not. If this is set to true, two additional pressures (the fluid pressure and the compaction pressure) will be added to the finite element. Including melt transport in the simulation also requires that there is one compositional field that has the name 'porosity'. This field will be used for computing the additional pressures and the melt velocity, and has a different advection equation than other compositional fields, as it is effectively advected with the melt velocity.

Possible values: A boolean value (true or false)

- *Parameter name:* Melt scaling factor threshold

Value: 1e-3

Default: 1e-3

Description: The factor by how much the Darcy coefficient K_D in a cell can be smaller than the reference Darcy coefficient for this cell still to be considered a melt cell (for which the melt transport equations are solved). For smaller Darcy coefficients, the Stokes equations (without melt) are solved instead. Only used if "Include melt transport" is true.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Use discontinuous compaction pressure

Value: true

Default: true

Description: Whether to use a discontinuous element for the compaction pressure or not. From our preliminary tests, continuous elements seem to work better in models where the porosity is > 0 everywhere in the domain, and discontinuous elements work better in models where in parts of the domain the porosity = 0.

Possible values: A boolean value (true or false)

A.107 Parameters in section Mesh refinement

- *Parameter name:* Adapt by fraction of cells

Value: false

Default: false

Description: Use fraction of the total number of cells instead of fraction of the total error as the limit for refinement and coarsening.

Possible values: A boolean value (true or false)

- *Parameter name:* Additional refinement times

Value:

Default:

Description: A list of times so that if the end time of a time step is beyond this time, an additional round of mesh refinement is triggered. This is mostly useful to make sure we can get through the initial transient phase of a simulation on a relatively coarse mesh, and then refine again when we are in a time range that we are interested in and where we would like to use a finer mesh. Units: Each element of the list has units years if the 'Use years in output instead of seconds' parameter is set; seconds otherwise.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Coarsening fraction

Value: 0.05

Default: 0.05

Description: The fraction of cells with the smallest error that should be flagged for coarsening.

Possible values: A floating point number v such that $0 \leq v \leq 1$

- *Parameter name:* Initial adaptive refinement

Value: 0

Default: 0

Description: The number of adaptive refinement steps performed after initial global refinement but while still within the first time step.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* Initial global refinement

Value: 2

Default: 2

Description: The number of global refinement steps performed on the initial coarse mesh, before the problem is first solved there.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* Minimum refinement level

Value: 0

Default: 0

Description: The minimum refinement level each cell should have, and that can not be exceeded by coarsening. Should not be higher than the ‘Initial global refinement’ parameter.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* Normalize individual refinement criteria

Value: true

Default: true

Description: If multiple refinement criteria are specified in the “Strategy” parameter, then they need to be combined somehow to form the final refinement indicators. This is done using the method described by the “Refinement criteria merge operation” parameter which can either operate on the raw refinement indicators returned by each strategy (i.e., dimensional quantities) or using normalized values where the indicators of each strategy are first normalized to the interval $[0, 1]$ (which also makes them non-dimensional). This parameter determines whether this normalization will happen.

Possible values: A boolean value (true or false)

- *Parameter name:* Refinement criteria merge operation

Value: max

Default: max

Description: If multiple mesh refinement criteria are computed for each cell (by passing a list of more than element to the **Strategy** parameter in this section of the input file) then one will have to decide which one should win when deciding which cell to refine. The operation that selects from these competing criteria is the one that is selected here. The options are:

- plus: Add the various error indicators together and refine those cells on which the sum of indicators is largest.
- max: Take the maximum of the various error indicators and refine those cells on which the maximal indicators is largest.

The refinement indicators computed by each strategy are modified by the “Normalize individual refinement criteria” and “Refinement criteria scale factors” parameters.

Possible values: Any one of plus, max

- *Parameter name:* Refinement criteria scaling factors

Value:

Default:

Description: A list of scaling factors by which every individual refinement criterion will be multiplied by. If only a single refinement criterion is selected (using the “Strategy” parameter, then this parameter has no particular meaning. On the other hand, if multiple criteria are chosen, then these factors are used to weigh the various indicators relative to each other.

If “Normalize individual refinement criteria” is set to true, then the criteria will first be normalized to the interval $[0, 1]$ and then multiplied by the factors specified here. You will likely want to choose the

factors to be not too far from 1 in that case, say between 1 and 10, to avoid essentially disabling those criteria with small weights. On the other hand, if the criteria are not normalized to [0, 1] using the parameter mentioned above, then the factors you specify here need to take into account the relative numerical size of refinement indicators (which in that case carry physical units).

You can experimentally play with these scaling factors by choosing to output the refinement indicators into the graphical output of a run.

If the list of indicators given in this parameter is empty, then this indicates that they should all be chosen equal to one. If the list is not empty then it needs to have as many entries as there are indicators chosen in the “Strategy” parameter.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Refinement fraction

Value: 0.3

Default: 0.3

Description: The fraction of cells with the largest error that should be flagged for refinement.

Possible values: A floating point number v such that $0 \leq v \leq 1$

- *Parameter name:* Run postprocessors on initial refinement

Value: false

Default: false

Description: Whether or not the postprocessors should be executed after each of the initial adaptive refinement cycles that are run at the start of the simulation.

Possible values: A boolean value (true or false)

- *Parameter name:* Strategy

Value: thermal energy density

Default: thermal energy density

Description: A comma separated list of mesh refinement criteria that will be run whenever mesh refinement is required. The results of each of these criteria, i.e., the refinement indicators they produce for all the cells of the mesh will then be normalized to a range between zero and one and the results of different criteria will then be merged through the operation selected in this section.

The following criteria are available:

‘artificial viscosity’: A mesh refinement criterion that computes refinement indicators from the artificial viscosity of the temperature or compositional fields based on user specified weights.

‘boundary’: A class that implements a mesh refinement criterion which always flags all cells on specified boundaries for refinement. This is useful to provide high accuracy for processes at or close to the edge of the model domain.

To use this refinement criterion, you may want to combine it with other refinement criteria, setting the ‘Normalize individual refinement criteria’ flag and using the ‘max’ setting for ‘Refinement criteria merge operation’.

‘compaction length’: A mesh refinement criterion for models with melt transport that computes refinement indicators based on the compaction length, defined as $\delta = \sqrt{\frac{(\xi + 4\eta/3)k}{\eta_f}}$. ξ is the bulk viscosity, η is the shear viscosity, k is the permeability and η_f is the melt viscosity. If the cell width or height exceeds a multiple (which is specified as an input parameter) of this compaction length, the cell is marked for refinement.

‘composition’: A mesh refinement criterion that computes refinement indicators from the compositional fields. If there is more than one compositional field, then it simply takes the sum of the indicators computed from each of the compositional field.

The way these indicators are computed is by evaluating the ‘Kelly error indicator’ on each compositional field. This error indicator takes the finite element approximation of the compositional field and uses it to compute an approximation of the second derivatives of the composition for each cell. This approximation is then multiplied by an appropriate power of the cell’s diameter to yield an indicator for how large the error is likely going to be on this cell. This construction rests on the observation that for many partial differential equations, the error on each cell is proportional to some power of the cell’s diameter times the second derivatives of the solution on that cell.

For complex equations such as those we solve here, this observation may not be strictly true in the mathematical sense, but it often yields meshes that are surprisingly good.

‘composition approximate gradient’: A mesh refinement criterion that computes refinement indicators from the gradients of compositional fields. If there is more than one compositional field, then it simply takes the sum of the indicators times a user-specified weight for each field.

In contrast to the ‘composition gradient’ refinement criterion, the current criterion does not compute the gradient at quadrature points on each cell, but by a finite difference approximation between the centers of cells. Consequently, it also works if the compositional fields are computed using discontinuous finite elements.

‘composition gradient’: A mesh refinement criterion that computes refinement indicators from the gradients of compositional fields. If there is more than one compositional field, then it simply takes the sum of the indicators times a user-specified weight for each field.

This refinement criterion computes the gradient of the compositional field at quadrature points on each cell, and then averages them in some way to obtain a refinement indicator for each cell. This will give a reasonable approximation of the true gradient of the compositional field if you are using a continuous finite element.

On the other hand, for discontinuous finite elements (see the ‘Use discontinuous composition discretization’ parameter in the ‘Discretization’ section), the gradient at quadrature points does not include the contribution of jumps in the compositional field between cells, and consequently will not be an accurate approximation of the true gradient. As an extreme example, consider the case of using piecewise constant finite elements for compositional fields; in that case, the gradient of the solution at quadrature points inside each cell will always be exactly zero, even if the finite element solution is different from each cell to the next. Consequently, the current refinement criterion will likely not be useful in this situation. That said, the ‘composition approximate gradient’ refinement criterion exists for exactly this purpose.

‘composition threshold’: A mesh refinement criterion that computes refinement indicators from the compositional fields. If any field exceeds the threshold given in the input file, the cell is marked for refinement.

‘density’: A mesh refinement criterion that computes refinement indicators from a field that describes the spatial variability of the density, ρ . Because this quantity may not be a continuous function (ρ and C_p may be discontinuous functions along discontinuities in the medium, for example due to phase changes), we approximate the gradient of this quantity to refine the mesh. The error indicator defined here takes the magnitude of the approximate gradient and scales it by $h_K^{1+d/2}$ where h_K is the diameter of each cell and d is the dimension. This scaling ensures that the error indicators converge to zero as $h_K \rightarrow 0$ even if the energy density is discontinuous, since the gradient of a discontinuous function grows like $1/h_K$.

‘maximum refinement function’: A mesh refinement criterion that ensures a maximum refinement level described by an explicit formula with the depth or position as argument. Which coordinate

representation is used is determined by an input parameter. Whatever the coordinate system chosen, the function you provide in the input file will by default depend on variables ‘x’, ‘y’ and ‘z’ (if in 3d). However, the meaning of these symbols depends on the coordinate system. In the Cartesian coordinate system, they simply refer to their natural meaning. If you have selected ‘depth’ for the coordinate system, then ‘x’ refers to the depth variable and ‘y’ and ‘z’ will simply always be zero. If you have selected a spherical coordinate system, then ‘x’ will refer to the radial distance of the point to the origin, ‘y’ to the azimuth angle and ‘z’ to the polar angle measured positive from the north pole. Note that the order of spherical coordinates is r,phi,theta and not r,theta,phi, since this allows for dimension independent expressions. Each coordinate system also includes a final ‘t’ variable which represents the model time, evaluated in years if the ‘Use years in output instead of seconds’ parameter is set, otherwise evaluated in seconds. After evaluating the function, its values are rounded to the nearest integer.

The format of these functions follows the syntax understood by the muparser library, see Section 4.7.3.

‘minimum refinement function’: A mesh refinement criterion that ensures a minimum refinement level described by an explicit formula with the depth or position as argument. Which coordinate representation is used is determined by an input parameter. Whatever the coordinate system chosen, the function you provide in the input file will by default depend on variables ‘x’, ‘y’ and ‘z’ (if in 3d). However, the meaning of these symbols depends on the coordinate system. In the Cartesian coordinate system, they simply refer to their natural meaning. If you have selected ‘depth’ for the coordinate system, then ‘x’ refers to the depth variable and ‘y’ and ‘z’ will simply always be zero. If you have selected a spherical coordinate system, then ‘x’ will refer to the radial distance of the point to the origin, ‘y’ to the azimuth angle and ‘z’ to the polar angle measured positive from the north pole. Note that the order of spherical coordinates is r,phi,theta and not r,theta,phi, since this allows for dimension independent expressions. Each coordinate system also includes a final ‘t’ variable which represents the model time, evaluated in years if the ‘Use years in output instead of seconds’ parameter is set, otherwise evaluated in seconds. After evaluating the function, its values are rounded to the nearest integer.

The format of these functions follows the syntax understood by the muparser library, see Section 4.7.3.

‘nonadiabatic temperature’: A mesh refinement criterion that computes refinement indicators from the excess temperature(difference between temperature and adiabatic temperature).

‘particle density’: A mesh refinement criterion that computes refinement indicators based on the density of particles. In practice this plugin equilibrates the number of particles per cell, leading to fine cells in high particle density regions and coarse cells in low particle density regions. This plugin is mostly useful for models with inhomogeneous particle density, e.g. when tracking an initial interface with a high particle density, or when the spatial particle density denotes the region of interest. Additionally, this plugin tends to balance the computational load between processes in parallel computations, because the particle and mesh density is more aligned.

‘slope’: A class that implements a mesh refinement criterion intended for use with a free surface. It calculates a local slope based on the angle between the surface normal and the local gravity vector. Cells with larger angles are marked for refinement.

To use this refinement criterion, you may want to combine it with other refinement criteria, setting the ‘Normalize individual refinement criteria’ flag and using the ‘max’ setting for ‘Refinement criteria merge operation’.

‘strain rate’: A mesh refinement criterion that computes the refinement indicators equal to the strain rate norm computed at the center of the elements.

‘temperature’: A mesh refinement criterion that computes refinement indicators from the temperature field.

The way these indicators are computed is by evaluating the ‘Kelly error indicator’ on the temperature field. This error indicator takes the finite element approximation of the temperature field and uses

it to compute an approximation of the second derivatives of the temperature for each cell. This approximation is then multiplied by an appropriate power of the cell's diameter to yield an indicator for how large the error is likely going to be on this cell. This construction rests on the observation that for many partial differential equations, the error on each cell is proportional to some power of the cell's diameter times the second derivatives of the solution on that cell.

For complex equations such as those we solve here, this observation may not be strictly true in the mathematical sense, but it often yields meshes that are surprisingly good.

'thermal energy density': A mesh refinement criterion that computes refinement indicators from a field that describes the spatial variability of the thermal energy density, $\rho C_p T$. Because this quantity may not be a continuous function (ρ and C_p may be discontinuous functions along discontinuities in the medium, for example due to phase changes), we approximate the gradient of this quantity to refine the mesh. The error indicator defined here takes the magnitude of the approximate gradient and scales it by $h_K^{1.5}$ where h_K is the diameter of each cell. This scaling ensures that the error indicators converge to zero as $h_K \rightarrow 0$ even if the energy density is discontinuous, since the gradient of a discontinuous function grows like $1/h_K$.

'topography': A class that implements a mesh refinement criterion, which always flags all cells in the uppermost layer for refinement. This is useful to provide high accuracy for processes at or close to the surface.

To use this refinement criterion, you may want to combine it with other refinement criteria, setting the 'Normalize individual refinement criteria' flag and using the 'max' setting for 'Refinement criteria merge operation'.

'velocity': A mesh refinement criterion that computes refinement indicators from the velocity field.

The way these indicators are computed is by evaluating the 'Kelly error indicator' on the velocity field. This error indicator takes the finite element approximation of the velocity field and uses it to compute an approximation of the second derivatives of the velocity for each cell. This approximation is then multiplied by an appropriate power of the cell's diameter to yield an indicator for how large the error is likely going to be on this cell. This construction rests on the observation that for many partial differential equations, the error on each cell is proportional to some power of the cell's diameter times the second derivatives of the solution on that cell.

For complex equations such as those we solve here, this observation may not be strictly true in the mathematical sense, but it often yields meshes that are surprisingly good.

'viscosity': A mesh refinement criterion that computes refinement indicators from a field that describes the spatial variability of the logarithm of the viscosity, $\log \eta$. (We choose the logarithm of the viscosity because it can vary by orders of magnitude.) Because this quantity may not be a continuous function (η may be a discontinuous function along discontinuities in the medium, for example due to phase changes), we approximate the gradient of this quantity to refine the mesh. The error indicator defined here takes the magnitude of the approximate gradient and scales it by $h_K^{1+d/2}$ where h_K is the diameter of each cell and d is the dimension. This scaling ensures that the error indicators converge to zero as $h_K \rightarrow 0$ even if the energy density is discontinuous, since the gradient of a discontinuous function grows like $1/h_K$.

Possible values: A comma-separated list of any of artificial viscosity, boundary, compaction length, composition, composition approximate gradient, composition gradient, composition threshold, density, maximum refinement function, minimum refinement function, nonadiabatic temperature, particle density, slope, strain rate, temperature, thermal energy density, topography, velocity, viscosity

- **Parameter name:** Time steps between mesh refinement

Value: 10

Default: 10

Description: The number of time steps after which the mesh is to be adapted again based on computed error indicators. If 0 then the mesh will never be changed.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

A.108 Parameters in section Mesh refinement/Artificial viscosity

- *Parameter name:* Compositional field scaling factors

Value:

Default:

Description: A list of scaling factors by which every individual compositional field will be multiplied. These factors are used to weigh the various indicators relative to each other and to the temperature.

If the list of scaling factors given in this parameter is empty, then this indicates that they should all be chosen equal to 0. If the list is not empty then it needs to have as many entries as there are compositional fields.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Temperature scaling factor

Value: 0.0

Default: 0.0

Description: A scaling factor for the artificial viscosity of the temperature equation. Use 0.0 to disable.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.109 Parameters in section Mesh refinement/Boundary

- *Parameter name:* Boundary refinement indicators

Value:

Default:

Description: A comma separated list of names denoting those boundaries where there should be mesh refinement.

The names of the boundaries listed here can either be numbers (in which case they correspond to the numerical boundary indicators assigned by the geometry object), or they can correspond to any of the symbolic names the geometry object may have provided for each part of the boundary. You may want to compare this with the documentation of the geometry model you use in your model.

Possible values: A list of 0 to 4294967295 elements where each element is [Any string]

A.110 Parameters in section Mesh refinement/Compaction length

- *Parameter name:* Mesh cells per compaction length

Value: 1.0

Default: 1.0

Description: The desired ratio between compaction length and size of the mesh cells, or, in other words, how many cells the mesh should (at least) have per compaction length. Every cell where this ratio is smaller than the value specified by this parameter (in places with fewer mesh cells per compaction length) is marked for refinement.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.111 Parameters in section Mesh refinement/Composition

- Parameter name: Compositional field scaling factors

Value:

Default:

Description: A list of scaling factors by which every individual compositional field will be multiplied by. If only a single compositional field exists, then this parameter has no particular meaning. On the other hand, if multiple criteria are chosen, then these factors are used to weigh the various indicators relative to each other.

If the list of scaling factors given in this parameter is empty, then this indicates that they should all be chosen equal to one. If the list is not empty then it needs to have as many entries as there are compositional fields.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

A.112 Parameters in section Mesh refinement/Composition approximate gradient

- Parameter name: Compositional field scaling factors

Value:

Default:

Description: A list of scaling factors by which every individual compositional field gradient will be multiplied. If only a single compositional field exists, then this parameter has no particular meaning. On the other hand, if multiple criteria are chosen, then these factors are used to weigh the various indicators relative to each other.

If the list of scaling factors given in this parameter is empty, then this indicates that they should all be chosen equal to one. If the list is not empty then it needs to have as many entries as there are compositional fields.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

A.113 Parameters in section Mesh refinement/Composition gradient

- Parameter name: Compositional field scaling factors

Value:

Default:

Description: A list of scaling factors by which every individual compositional field gradient will be multiplied. If only a single compositional field exists, then this parameter has no particular meaning. On the other hand, if multiple criteria are chosen, then these factors are used to weigh the various indicators relative to each other.

If the list of scaling factors given in this parameter is empty, then this indicates that they should all be chosen equal to one. If the list is not empty then it needs to have as many entries as there are compositional fields.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

A.114 Parameters in section Mesh refinement/Composition threshold

- *Parameter name:* Compositional field thresholds

Value:

Default:

Description: A list of thresholds that every individual compositional field will be evaluated against.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

A.115 Parameters in section Mesh refinement/Maximum refinement function

- *Parameter name:* Coordinate system

Value: depth

Default: depth

Description: A selection that determines the assumed coordinate system for the function variables. Allowed values are ‘depth’, ‘cartesian’ and ‘spherical’. ‘depth’ will create a function, in which only the first variable is non-zero, which is interpreted to be the depth of the point. ‘spherical’ coordinates are interpreted as r,phi or r,phi,theta in 2D/3D respectively with theta being the polar angle.

Possible values: Any one of depth, cartesian, spherical

- *Parameter name:* Function constants

Value:

Default:

Description: Sometimes it is convenient to use symbolic constants in the expression that describes the function, rather than having to use its numeric value everywhere the constant appears. These values can be defined using this parameter, in the form ‘var1=value1, var2=value2, ...’

A typical example would be to set this runtime parameter to ‘pi=3.1415926536’ and then use ‘pi’ in the expression of the actual formula. (That said, for convenience this class actually defines both ‘pi’ and ‘Pi’ by default, but you get the idea.)

Possible values: Any string

- *Parameter name:* Function expression

Value: 0

Default: 0

Description: The formula that denotes the function you want to evaluate for particular values of the independent variables. This expression may contain any of the usual operations such as addition or multiplication, as well as all of the common functions such as ‘sin’ or ‘cos’. In addition, it may contain expressions like ‘if($x > 0$, 1, -1)’ where the expression evaluates to the second argument if the first argument is true, and to the third argument otherwise. For a full overview of possible expressions accepted see the documentation of the muparser library at <http://muparser.beltoforion.de/>.

If the function you are describing represents a vector-valued function with multiple components, then separate the expressions for individual components by a semicolon.

Possible values: Any string

- *Parameter name:* Variable names

Value: x,y,t

Default: x,y,t

Description: The name of the variables as they will be used in the function, separated by commas. By default, the names of variables at which the function will be evaluated is ‘x’ (in 1d), ‘x,y’ (in 2d) or ‘x,y,z’ (in 3d) for spatial coordinates and ‘t’ for time. You can then use these variable names in your function expression and they will be replaced by the values of these variables at which the function is currently evaluated. However, you can also choose a different set of names for the independent variables at which to evaluate your function expression. For example, if you work in spherical coordinates, you may wish to set this input parameter to ‘r,phi,theta,t’ and then use these variable names in your function expression.

Possible values: Any string

A.116 Parameters in section Mesh refinement/Minimum refinement function

- *Parameter name:* Coordinate system

Value: depth

Default: depth

Description: A selection that determines the assumed coordinate system for the function variables. Allowed values are ‘depth’, ‘cartesian’ and ‘spherical’. ‘depth’ will create a function, in which only the first variable is non-zero, which is interpreted to be the depth of the point. ‘spherical’ coordinates are interpreted as r,phi or r,phi,theta in 2D/3D respectively with theta being the polar angle.

Possible values: Any one of depth, cartesian, spherical

- *Parameter name:* Function constants

Value:

Default:

Description: Sometimes it is convenient to use symbolic constants in the expression that describes the function, rather than having to use its numeric value everywhere the constant appears. These values can be defined using this parameter, in the form ‘var1=value1, var2=value2, ...’.

A typical example would be to set this runtime parameter to ‘pi=3.1415926536’ and then use ‘pi’ in the expression of the actual formula. (That said, for convenience this class actually defines both ‘pi’ and ‘Pi’ by default, but you get the idea.)

Possible values: Any string

- *Parameter name:* Function expression

Value: 0

Default: 0

Description: The formula that denotes the function you want to evaluate for particular values of the independent variables. This expression may contain any of the usual operations such as addition or multiplication, as well as all of the common functions such as ‘sin’ or ‘cos’. In addition, it may contain expressions like ‘if(x>0, 1, -1)’ where the expression evaluates to the second argument if the first argument is true, and to the third argument otherwise. For a full overview of possible expressions accepted see the documentation of the muparser library at <http://muparser.beltoforion.de/>.

If the function you are describing represents a vector-valued function with multiple components, then separate the expressions for individual components by a semicolon.

Possible values: Any string

- *Parameter name:* Variable names

Value: x,y,t

Default: x,y,t

Description: The name of the variables as they will be used in the function, separated by commas. By default, the names of variables at which the function will be evaluated is ‘x’ (in 1d), ‘x,y’ (in 2d) or ‘x,y,z’ (in 3d) for spatial coordinates and ‘t’ for time. You can then use these variable names in your function expression and they will be replaced by the values of these variables at which the function is currently evaluated. However, you can also choose a different set of names for the independent variables at which to evaluate your function expression. For example, if you work in spherical coordinates, you may wish to set this input parameter to ‘r,phi,theta,t’ and then use these variable names in your function expression.

Possible values: Any string

A.117 Parameters in section Nullspace removal

- *Parameter name:* Remove nullspace

Value:

Default:

Description: Choose none, one or several from

- net rotation
- angular momentum
- net translation
- linear momentum
- net x translation
- net y translation
- net z translation
- linear x momentum
- linear y momentum
- linear z momentum

These are a selection of operations to remove certain parts of the nullspace from the velocity after solving. For some geometries and certain boundary conditions the velocity field is not uniquely determined but contains free translations and/or rotations. Depending on what you specify here, these non-determined modes will be removed from the velocity field at the end of the Stokes solve step.

The “angular momentum” option removes a rotation such that the net angular momentum is zero. The “linear * momentum” options remove translations such that the net momentum in the relevant direction is zero. The “net rotation” option removes the net rotation of the domain, and the “net * translation” options remove the net translations in the relevant directions. For most problems there should not be a significant difference between the momentum and rotation/translation versions of nullspace removal, although the momentum versions are more physically motivated. They are equivalent for constant density simulations, and approximately equivalent when the density variations are small.

Note that while more than one operation can be selected it only makes sense to pick one rotational and one translational operation.

Possible values: A comma-separated list of any of net rotation, angular momentum, net translation, linear momentum, net x translation, net y translation, net z translation, linear x momentum, linear y momentum, linear z momentum

A.118 Parameters in section Postprocess

- *Parameter name:* List of postprocessors

Value:

Default:

Description: A comma separated list of postprocessor objects that should be run at the end of each time step. Some of these postprocessors will declare their own parameters which may, for example, include that they will actually do something only every so many time steps or years. Alternatively, the text ‘all’ indicates that all available postprocessors should be run after each time step.

The following postprocessors are available:

‘Stokes residual’: A postprocessor that outputs the Stokes residuals during the iterative solver algorithm into a file `stokes_residuals.txt` in the output directory.

‘basic statistics’: A postprocessor that outputs some simplified statistics like the Rayleigh number and other quantities that only make sense in certain model setups. The output is written after completing initial adaptive refinement steps. The postprocessor assumes a point at the surface at the adiabatic surface temperature and pressure is a reasonable reference condition for computing these properties. Furthermore, the Rayleigh number is computed using the model depth (i.e. not the radius of the Earth), as we need a definition that is geometry independent. Take care when comparing these values to published studies and make sure they use the same definitions.

‘boundary densities’: A postprocessor that computes the laterally averaged density at the top and bottom of the domain.

‘boundary pressures’: A postprocessor that computes the laterally averaged pressure at the top and bottom of the domain.

‘command’: A postprocessor that executes a command line process.

‘composition statistics’: A postprocessor that computes some statistics about the compositional fields, if present in this simulation. In particular, it computes maximal and minimal values of each field, as well as the total mass contained in this field as defined by the integral $m_i(t) = \int_{\Omega} c_i(\mathbf{x}, t) dx$.

‘core statistics’: A postprocessor that computes some statistics about the core evolution. (Working only with dynamic core boundary temperature plugin)

‘depth average’: A postprocessor that computes depth averaged quantities and writes them into a file `<depth_average.ext>` in the output directory, where the extension of the file is determined by the output format you select. In addition to the output format, a number of other parameters also influence this postprocessor, and they can be set in the section `Postprocess/Depth average` in the input file.

In the output files, the x -value of each data point corresponds to the depth, whereas the y -value corresponds to the simulation time. The time is provided in seconds or, if the global “Use years in output instead of seconds” parameter is set, in years.

‘dynamic topography’: A postprocessor that computes a measure of dynamic topography based on the stress at the surface and bottom. The data is written into text files named ‘`dynamic_topography.NNNNN`’ in the output directory, where NNNNN is the number of the time step.

The exact approach works as follows: At the centers of all cells that sit along the top surface, we evaluate the stress and evaluate the component of it in the direction in which gravity acts. In other words, we compute $\sigma_{rr} = \hat{g}^T(2\eta\varepsilon(\mathbf{u}) - \frac{1}{3}(\operatorname{div} \mathbf{u})I)\hat{g} - p_d$ where $\hat{g} = \mathbf{g}/\|\mathbf{g}\|$ is the direction of the gravity vector \mathbf{g} and $p_d = p - p_a$ is the dynamic pressure computed by subtracting the adiabatic pressure p_a from the total pressure p computed as part of the Stokes solve. From this, the dynamic topography is computed using the formula $h = \frac{\sigma_{rr}}{(\mathbf{g} \cdot \mathbf{n})\rho}$ where ρ is the density at the cell center. For the bottom surface we chose the convection that positive values are up (out) and negative values are in (down),

analogous to the deformation of the upper surface. Note that this implementation takes the direction of gravity into account, which means that reversing the flow in backward advection calculations will not reverse the instantaneous topography because the reverse flow will be divided by the reverse surface gravity. The file format then consists of lines with Euclidean coordinates followed by the corresponding topography value.

(As a side note, the postprocessor chooses the cell center instead of the center of the cell face at the surface, where we really are interested in the quantity, since this often gives better accuracy. The results should in essence be the same, though.)

‘geoid’: A postprocessor that computes a representation of the geoid based on the density structure in the mantle, as well as the dynamic topography at the surface and core mantle boundary (CMB). The geoid is computed from a spherical harmonic expansion, so the geometry of the domain must be a 3D spherical shell.

‘global statistics’: A postprocessor that outputs all the global statistics information, e.g. the time of the simulation, the timestep number, number of degrees of freedom and solver iterations for each timestep. The postprocessor can output different formats, the first printing one line in the statistics file per nonlinear solver iteration (if a nonlinear solver scheme is selected). The second prints one line per timestep, summing the information about all nonlinear iterations in this line. Note that this postprocessor is always active independent on whether or not it is selected in the parameter file.

‘heat flux densities’: A postprocessor that computes some statistics about the (conductive) heat flux density for each boundary id. The heat flux density is computed in outward direction, i.e., from the domain to the outside, using the formula $\frac{1}{|\Gamma_i|} \int_{\Gamma_i} -k\nabla T \cdot \mathbf{n}$ where Γ_i is the part of the boundary with indicator i , k is the thermal conductivity as reported by the material model, T is the temperature, and \mathbf{n} is the outward normal. Note that the quantity so computed does not include any energy transported across the boundary by material transport in cases where $\mathbf{u} \cdot \mathbf{n} \neq 0$.

Note that the “heat flux” postprocessor computes the same quantity as the one here, but not divided by the area of the surface. In other words, it computes the *total* heat flux through each boundary.

‘heat flux map’: A postprocessor that computes the (conductive) heat flux density across each boundary. The heat density flux is computed in outward direction, i.e., from the domain to the outside, using the formula $-k\nabla T \cdot \mathbf{n}$, where k is the thermal conductivity as reported by the material model, T is the temperature, and \mathbf{n} is the outward normal. Note that the quantity so computed does not include any energy transported across the boundary by material transport in cases where $\mathbf{u} \cdot \mathbf{n} \neq 0$. The integrated heat flux for each boundary can be obtained from the heat flux statistics postprocessor.

‘heat flux statistics’: A postprocessor that computes some statistics about the (conductive) heat flux across boundaries. For each boundary indicator (see your geometry description for which boundary indicators are used), the heat flux is computed in outward direction, i.e., from the domain to the outside, using the formula $\int_{\Gamma_i} -k\nabla T \cdot \mathbf{n}$ where Γ_i is the part of the boundary with indicator i , k is the thermal conductivity as reported by the material model, T is the temperature, and \mathbf{n} is the outward normal. Note that the quantity so computed does not include any energy transported across the boundary by material transport in cases where $\mathbf{u} \cdot \mathbf{n} \neq 0$. The point-wise heat flux can be obtained from the heat flux map postprocessor, which outputs the heat flux to a file, or the heat flux map visualization postprocessor, which outputs the heat flux for visualization.

As stated, this postprocessor computes the *outbound* heat flux. If you are interested in the opposite direction, for example from the core into the mantle when the domain describes the mantle, then you need to multiply the result by -1.

Note: In geodynamics, the term “heat flux” is often understood to be the quantity $-k\nabla T$, which is really a heat flux *density*, i.e., a vector-valued field. In contrast to this, the current postprocessor only computes the integrated flux over each part of the boundary. Consequently, the units of the quantity computed here are $W = \frac{J}{s}$.

The “heat flux densities” postprocessor computes the same quantity as the one here, but divided by the area of the surface.

‘heating statistics’: A postprocessor that computes some statistics about heating, averaged by volume.

‘mass flux statistics’: A postprocessor that computes some statistics about the mass flux across boundaries. For each boundary indicator (see your geometry description for which boundary indicators are used), the mass flux is computed in outward direction, i.e., from the domain to the outside, using the formula $\int_{\Gamma_i} \rho \mathbf{v} \cdot \mathbf{n}$ where Γ_i is the part of the boundary with indicator i , ρ is the density as reported by the material model, \mathbf{v} is the velocity, and \mathbf{n} is the outward normal.

As stated, this postprocessor computes the *outbound* mass flux. If you are interested in the opposite direction, for example from the core into the mantle when the domain describes the mantle, then you need to multiply the result by -1.

Note: In geodynamics, the term “mass flux” is often understood to be the quantity $\rho \mathbf{v}$, which is really a mass flux *density*, i.e., a vector-valued field. In contrast to this, the current postprocessor only computes the integrated flux over each part of the boundary. Consequently, the units of the quantity computed here are $\frac{\text{kg}}{\text{s}}$.

‘matrix statistics’: A postprocessor that computes some statistics about the matrices. In particular, it outputs total memory consumption, total non-zero elements, and non-zero elements per block, for system matrix and system preconditioner matrix.

‘melt statistics’: A postprocessor that computes some statistics about the melt (volume) fraction. If the material model does not implement a melt fraction function, the output is set to zero.

‘memory statistics’: A postprocessor that computes some statistics about the memory consumption. In particular, it computes the memory usage of the system matrix, triangulation, p4est, DoFHandler, current constraints, and solution vector, all in MB. It also outputs the memory usage of the system matrix to the screen.

‘particle count statistics’: A postprocessor that computes some statistics about the particle distribution, if present in this simulation. In particular, it computes minimal, average and maximal values of particles per cell in the global domain.

‘particles’: A Postprocessor that creates particles that follow the velocity field of the simulation. The particles can be generated and propagated in various ways and they can carry a number of constant or time-varying properties. The postprocessor can write output positions and properties of all particles at chosen intervals, although this is not mandatory. It also allows other parts of the code to query the particles for information.

‘point values’: A postprocessor that evaluates the solution (i.e., velocity, pressure, temperature, and compositional fields along with other fields that are treated as primary variables) at the end of every time step at a given set of points and then writes this data into the file <point_values.txt> in the output directory. The points at which the solution should be evaluated are specified in the section **Postprocess/Point values** in the input file.

In the output file, data is organized as (i) time, (ii) the 2 or 3 coordinates of the evaluation points, and (iii) followed by the values of the solution vector at this point. The time is provided in seconds or, if the global “Use years in output instead of seconds” parameter is set, in years. In the latter case, the velocity is also converted to meters/year, instead of meters/second.

Note: Evaluating the solution of a finite element field at arbitrarily chosen points is an expensive process. Using this postprocessor will only be efficient if the number of evaluation points is relatively small. If you need a very large number of evaluation points, you should consider extracting this information from the visualization program you use to display the output of the ‘visualization’ postprocessor.

‘pressure statistics’: A postprocessor that computes some statistics about the pressure field.
 ‘spherical velocity statistics’: A postprocessor that computes radial, tangential and total RMS velocity.
 ‘temperature statistics’: A postprocessor that computes some statistics about the temperature field.
 ‘topography’: A postprocessor intended for use with a free surface. After every step it loops over all the vertices on the top surface and determines the maximum and minimum topography relative to a reference datum (initial box height for a box geometry model or initial radius for a sphere/spherical shell geometry model). If ‘Topography.Output to file’ is set to true, also outputs topography into text files named ‘topography.NNNNN’ in the output directory, where NNNNN is the number of the time step. The file format then consists of lines with Euclidean coordinates followed by the corresponding topography value. Topography is printed/written in meters.
 ‘velocity boundary statistics’: A postprocessor that computes some statistics about the velocity along the boundaries. For each boundary indicator (see your geometry description for which boundary indicators are used), the min and max velocity magnitude is computed.
 ‘velocity statistics’: A postprocessor that computes some statistics about the velocity field.
 ‘visualization’: A postprocessor that takes the solution and writes it into files that can be read by a graphical visualization program. Additional run time parameters are read from the parameter subsection ‘Visualization’.

Possible values: A comma-separated list of any of Stokes residual, basic statistics, boundary densities, boundary pressures, command, composition statistics, core statistics, depth average, dynamic topography, geoid, global statistics, heat flux densities, heat flux map, heat flux statistics, heating statistics, mass flux statistics, matrix statistics, melt statistics, memory statistics, particle count statistics, particles, point values, pressure statistics, spherical velocity statistics, temperature statistics, topography, velocity boundary statistics, velocity statistics, visualization

- *Parameter name:* Run postprocessors on nonlinear iterations

Value: false

Default: false

Description: Whether or not the postprocessors should be executed after each of the nonlinear iterations done within one time step. As this is mainly an option for the purposes of debugging, it is not supported when the ‘Time between graphical output’ is larger than zero, or when the postprocessor is not intended to be run more than once per timestep.

Possible values: A boolean value (true or false)

A.119 Parameters in section Postprocess/Command

- *Parameter name:* Command

Value:

Default:

Description: Command to execute.

Possible values: Any string

- *Parameter name:* Run on all processes

Value: false

Default: false

Description: Whether to run command from all processes (true), or only on process 0 (false).

Possible values: A boolean value (true or false)

- *Parameter name:* Terminate on failure

Value: false

Default: false

Description: Select whether ASPECT should terminate if the command returns a non-zero exit status.

Possible values: A boolean value (true or false)

A.120 Parameters in section Postprocess/Depth average

- *Parameter name:* List of output variables

Value: all

Default: all

Description: A comma separated list which specifies which quantities to average in each depth slice. It defaults to averaging all available quantities, but this can be an expensive operation, so you may want to select only a few.

List of options: all|temperature|composition|adiabatic temperature|adiabatic pressure|adiabatic density|adiabatic density derivative|velocity magnitude|sinking velocity|Vs|Vp|viscosity|vertical heat flux

Possible values: A comma-separated list of any of all, temperature, composition, adiabatic temperature, adiabatic pressure, adiabatic density, adiabatic density derivative, velocity magnitude, sinking velocity, Vs, Vp, viscosity, vertical heat flux

- *Parameter name:* Number of zones

Value: 10

Default: 10

Description: The number of zones in depth direction within which we are to compute averages. By default, we subdivide the entire domain into 10 depth zones and compute temperature and other averages within each of these zones. However, if you have a very coarse mesh, it may not make much sense to subdivide the domain into so many zones and you may wish to choose less than this default. It may also make computations slightly faster. On the other hand, if you have an extremely highly resolved mesh, choosing more zones might also make sense.

Possible values: An integer n such that $1 \leq n \leq 2147483647$

- *Parameter name:* Output format

Value: gnuplot

Default: gnuplot

Description: The format in which the output shall be produced. The format in which the output is generated also determines the extension of the file into which data is written.

Possible values: Any one of none, dx, ucd, gnuplot, povray, eps, gmv, tecplot, tecplot_binary, vtk, vtu, hdf5, svg, deal.II intermediate, txt

- *Parameter name:* Time between graphical output

Value: 1e8

Default: 1e8

Description: The time interval between each generation of graphical output files. A value of zero indicates that output should be generated in each time step. Units: years if the 'Use years in output instead of seconds' parameter is set; seconds otherwise.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.121 Parameters in section Postprocess/Dynamic core statistics

- *Parameter name:* Excess entropy only

Value: false

Default: false

Description: Output the excess entropy only instead the each entropy terms.

Possible values: A boolean value (true or false)

A.122 Parameters in section Postprocess/Dynamic topography

- *Parameter name:* Density above

Value: 0

Default: 0

Description: Dynamic topography is calculated as the excess or lack of mass that is supported by mantle flow. This value depends on the density of material that is moved up or down, i.e. crustal rock, and the density of the material that is displaced (generally water or air). While the density of crustal rock is part of the material model, this parameter 'Density above' allows the user to specify the density value of material that is displaced above the solid surface. By default this material is assumed to be air, with a density of 0. Units: kg/m^3 .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Density below

Value: 9900

Default: 9900

Description: Dynamic topography is calculated as the excess or lack of mass that is supported by mantle flow. This value depends on the density of material that is moved up or down, i.e. mantle above CMB, and the density of the material that is displaced (generally outer core material). While the density of mantle rock is part of the material model, this parameter 'Density below' allows the user to specify the density value of material that is displaced below the solid surface. By default this material is assumed to be outer core material with a density of 9900. Units: kg/m^3 .

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Output bottom

Value: true

Default: true

Description: Whether to output a file containing the bottom (i.e., CMB) dynamic topography.

Possible values: A boolean value (true or false)

- *Parameter name:* Output surface
Value: true
Default: true
Description: Whether to output a file containing the surface dynamic topography.
Possible values: A boolean value (true or false)

A.123 Parameters in section Postprocess/Geoid

- *Parameter name:* Also output the spherical harmonic coefficients of CMB dynamic topography contribution
Value: false
Default: false
Description: Option to also output the spherical harmonic coefficients of the CMB dynamic topography contribution to the maximum degree. The default is false.
Possible values: A boolean value (true or false)
- *Parameter name:* Also output the spherical harmonic coefficients of density anomaly contribution
Value: false
Default: false
Description: Option to also output the spherical harmonic coefficients of the density anomaly contribution to the maximum degree. The default is false.
Possible values: A boolean value (true or false)
- *Parameter name:* Also output the spherical harmonic coefficients of geoid anomaly
Value: false
Default: false
Description: Option to also output the spherical harmonic coefficients of the geoid anomaly up to the maximum degree. The default is false, so postprocess will only output the geoid anomaly in grid format.
Possible values: A boolean value (true or false)
- *Parameter name:* Also output the spherical harmonic coefficients of surface dynamic topography contribution
Value: false
Default: false
Description: Option to also output the spherical harmonic coefficients of the surface dynamic topography contribution to the maximum degree. The default is false.
Possible values: A boolean value (true or false)
- *Parameter name:* Density above
Value: 0
Default: 0
Description: The density value above the surface boundary.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Density below
Value: 9900
Default: 9900
Description: The density value below the CMB boundary.
Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Maximum degree
Value: 20
Default: 20
Description: This parameter can be a random positive integer. However, the value normally should not exceed the maximum degree of the initial perturbed temperature field. For example, if the initial temperature uses S40RTS, the maximum degree should not be larger than 40.
Possible values: An integer n such that $0 \leq n \leq 2147483647$
- *Parameter name:* Minimum degree
Value: 2
Default: 2
Description: This parameter normally is set to 2 since the perturbed gravitational potential at degree 1 always vanishes in a reference frame with the planetary center of mass same as the center of figure.
Possible values: An integer n such that $0 \leq n \leq 2147483647$
- *Parameter name:* Output data in geographical coordinates
Value: false
Default: false
Description: Option to output the geoid anomaly in geographical coordinates (latitude and longitude). The default is false, so postprocess will output the data in geocentric coordinates (x,y,z) as normally.
Possible values: A boolean value (true or false)

A.124 Parameters in section Postprocess/Global statistics

- *Parameter name:* Write statistics for each nonlinear iteration
Value: false
Default: false
Description: Whether to put every nonlinear iteration into a separate line in the statistics file (if true), or to output only one line per time step that contains the total number of iterations of the Stokes and advection linear system solver.
Possible values: A boolean value (true or false)

A.125 Parameters in section Postprocess/Particles

- *Parameter name:* Data output format
Value: vtu
Default: vtu
Description: File format to output raw particle data in. If you select ‘none’ no output will be written. Select one of the following models:

‘ascii’: This particle output plugin writes particle positions and properties into space separated ascii files.

‘hdf5’: This particle output plugin writes particle positions and properties into hdf5 files.

‘vtu’: This particle output plugin writes particle positions and properties into vtu files.

Possible values: Any one of ascii, hdf5, vtu, none

- *Parameter name:* Integration scheme

Value: rk2

Default: rk2

Description: This parameter is used to decide which method to use to solve the equation that describes the position of particles, i.e., $\frac{d}{dt} \mathbf{x}_k(t) = \mathbf{u}(\mathbf{x}_k(t), t)$, where k is an index that runs over all particles, and $\mathbf{u}(\mathbf{x}, t)$ is the velocity field that results from the Stokes equations.

In practice, the exact velocity $\mathbf{u}(\mathbf{x}, t)$ is of course not available, but only a numerical approximation $\mathbf{u}_h(\mathbf{x}, t)$. Furthermore, this approximation is only available at discrete time steps, $\mathbf{u}^n(\mathbf{x}) = \mathbf{u}(\mathbf{x}, t^n)$, and these need to be interpolated between time steps if the integrator for the equation above requires an evaluation at time points between the discrete time steps. If we denote this interpolation in time by $\tilde{\mathbf{u}}_h(\mathbf{x}, t)$ where $\tilde{\mathbf{u}}_h(\mathbf{x}, t^n) = \mathbf{u}^n(\mathbf{x})$, then the equation the differential equation solver really tries to solve is $\frac{d}{dt} \tilde{\mathbf{x}}_k(t) = \tilde{\mathbf{u}}_h(\mathbf{x}_k(t), t)$.

As a consequence of these considerations, if you try to assess convergence properties of an ODE integrator – for example to verify that the RK4 integrator converges with fourth order –, it is important to recall that the integrator may not solve the equation you think it solves. If, for example, we call the numerical solution of the ODE $\tilde{\mathbf{x}}_{k,h}(t)$, then the error will typically satisfy a relationship like

$$\|\tilde{\mathbf{x}}_k(T) - \tilde{\mathbf{x}}_{k,h}(T)\| \leq C(T)\Delta t^p$$

where Δt is the time step and p the convergence order of the method, and $C(T)$ is a (generally unknown) constant that depends on the end time T at which one compares the solutions. On the other hand, an analytically computed trajectory would likely use the *exact* velocity, and one may be tempted to compute $\|\mathbf{x}_k(T) - \tilde{\mathbf{x}}_{k,h}(T)\|$, but this quantity will, in the best case, only satisfy an estimate of the form

$$\|\mathbf{x}_k(T) - \tilde{\mathbf{x}}_{k,h}(T)\| \leq C_1(T)\Delta t^p + C_2(T)\|\mathbf{u} - \mathbf{u}_h\| + C_3(T)\|\mathbf{u}_h - \tilde{\mathbf{u}}_h\|$$

with appropriately chosen norms for the second and third term. These second and third terms typically converge to zero at relatively low rates (compared to the order p of the integrator, which can often be chosen relatively high) in the mesh size h and the time step size

Deltat, limiting the overall accuracy of the ODE integrator.

Select one of the following models:

‘euler’: Explicit Euler scheme integrator, where $y_{n+1} = y_n + dt * v(y_n)$. This requires only one integration substep per timestep.

‘rk2’: Second Order Runge Kutta integrator $y_{n+1} = y_n + dt * v(t_{n+1/2}, y_n + 0.5 * k_1)$ where $k_1 = y_n + 0.5 * dt * v(t_n, y_n)$

‘rk4’: Runge Kutta fourth order integrator, where $y_{n+1} = y_n + (1/6)*k_1 + (1/3)*k_2 + (1/3)*k_3 + (1/6)*k_4$ and k_1, k_2, k_3, k_4 are defined as usual.

Possible values: Any one of euler, rk2, rk4

- *Parameter name:* Interpolation scheme

Value: cell average

Default: cell average

Description: Select one of the following models:

‘bilinear least squares’: Interpolates particle properties onto a vector of points using a bilinear least squares method. Currently only 2D models are supported. Note that deal.II must be configured with BLAS/LAPACK.

‘cell average’: Return the average of all particle properties in the given cell.

‘harmonic average’: Return the harmonic average of all particle properties in the given cell. If the cell contains no particles, return the harmonic average of the properties in the neighboring cells.

‘nearest neighbor’: Return the properties of the nearest neighboring particle in the current cell, or nearest particle in nearest neighboring cell if current cell is empty.

Possible values: Any one of bilinear least squares, cell average, harmonic average, nearest neighbor

- *Parameter name:* List of particle properties

Value:

Default:

Description: A comma separated list of particle properties that should be tracked. By default none is selected, which means only position, velocity and id of the particles are output.

The following properties are available:

‘composition’: Implementation of a plugin in which the particle property is defined by the compositional fields in the model. This can be used to track solid compositionevolution over time.

‘function’: Implementation of a model in which the particle property is set by evaluating an explicit function at the initial position of each particle. The function is defined in the parameters in section “Particles|Function”. The format of these functions follows the syntax understood by the muparser library, see Section 4.7.3.

‘initial composition’: Implementation of a plugin in which the particle property is given as the initial composition at the particle’s initial position. The particle gets as many properties as there are compositional fields.

‘initial position’: Implementation of a plugin in which the particle property is given as the initial position of the particle. This property is vector-valued with as many components as there are space dimensions. In practice, it is often most useful to only visualize one of the components of this vector, or the magnitude of the vector. For example, in a spherical mantle simulation, the magnitude of this property equals the starting radius of a particle, and is thereby indicative of which part of the mantle a particle comes from.

‘integrated strain’: A plugin in which the particle property tensor is defined as the deformation gradient tensor \mathbf{F} this particle has experienced. \mathbf{F} can be polar-decomposed into the left stretching tensor \mathbf{L} (the finite strain we are interested in), and the rotation tensor \mathbf{Q} . See the corresponding cookbook in the manual for more detailed information.

‘integrated strain invariant’: A plugin in which the particle property is defined as the finite strain invariant (ε_{ii}). This property is calculated with the timestep (dt) and the second invariant of the deviatoric strain rate tensor ($\dot{\varepsilon}_{ii}$), where the value at time step n is $\varepsilon_{ii}^n = \varepsilon_{ii}^{n-1} + dt\dot{\varepsilon}_{ii}$.

‘melt particle’: Implementation of a plugin in which the particle property is defined as presence of melt above a threshold, which can be set as an input parameter. This property is set to 0 if melt is not present and set to 1 if melt is present.

‘pT path’: Implementation of a plugin in which the particle property is defined as the current pressure and temperature at this position. This can be used to generate pressure-temperature paths of material points over time.

‘position’: Implementation of a plugin in which the particle property is defined as the current position.

'velocity': Implementation of a plugin in which the particle property is defined as the recent velocity at this position.

Possible values: A comma-separated list of any of composition, function, initial composition, initial position, integrated strain, integrated strain invariant, melt particle, pT path, position, velocity

- **Parameter name:** Load balancing strategy

Value: repartition

Default: repartition

Description: Strategy that is used to balance the computational load across processors for adaptive meshes.

Possible values: A comma-separated list of any of none, remove particles, add particles, remove and add particles, repartition

- **Parameter name:** Maximum particles per cell

Value: 100

Default: 100

Description: Upper limit for particle number per cell. This limit is useful for adaptive meshes to prevent coarse cells from slowing down the whole model. It will be checked and enforced after mesh refinement, after MPI transfer of particles and after particle movement. If there are `n_number_of_particles > max_particles_per_cell` particles in one cell then `n_number_of_particles - max_particles_per_cell` particles in this cell are randomly chosen and destroyed.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- **Parameter name:** Minimum particles per cell

Value: 0

Default: 0

Description: Lower limit for particle number per cell. This limit is useful for adaptive meshes to prevent fine cells from being empty of particles. It will be checked and enforced after mesh refinement and after particle movement. If there are `n_number_of_particles < min_particles_per_cell` particles in one cell then `min_particles_per_cell - n_number_of_particles` particles are generated and randomly placed in this cell. If the particles carry properties the individual property plugins control how the properties of the new particles are initialized.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- **Parameter name:** Number of particles

Value: 1000

Default: 1000

Description: Total number of particles to create (not per processor or per element). The number is parsed as a floating point number (so that one can specify, for example, '1e4' particles) but it is interpreted as an integer, of course.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- **Parameter name:** Particle generator name

Value: random uniform

Default: random uniform

Description: Select one of the following models:

‘ascii file’: Generates a distribution of particles from coordinates specified in an Ascii data file. The file format is a simple text file, with as many columns as spatial dimensions and as many lines as particles to be generated. Initial comment lines starting with ‘#’ will be discarded. All of the values that define this generator are read from a section “Postprocess/Particles/Generator/Ascii file” in the input file, see Section A.128.

‘probability density function’: Generate a random distribution of particles over the entire simulation domain. The probability density is prescribed in the form of a user-prescribed function. The format of this function follows the syntax understood by the muparser library, see Section 4.7.3. The return value of the function is always checked to be a non-negative probability density but it can be zero in parts of the domain.

‘quadrature points’: Generates particles at the quadrature points of each active cell of the triangulation. Here, Gauss quadrature of degree (velocity_degree + 1), is used similarly to the assembly of Stokes matrix.

‘random uniform’: Generates a random uniform distribution of particles over the entire simulation domain.

‘reference cell’: Generate a uniform distribution of particles per cell and spatial direction in the unit cell and transforms each of the particles back to real region in the model domain. Uniform here means the particles will be generated with an equal spacing in each spatial dimension

‘uniform box’: Generate a uniform distribution of particles over a rectangular domain in 2D or 3D. Uniform here means the particles will be generated with an equal spacing in each spatial dimension. Note that in order to produce a regular distribution the number of generated particles might not exactly match the one specified in the input file.

‘uniform radial’: Generate a uniform distribution of particles over a spherical domain in 2D or 3D. Uniform here means the particles will be generated with an equal spacing in each spherical spatial dimension, i.e., the particles are created at positions that increase linearly with equal spacing in radius, colatitude and longitude around a certain center point. Note that in order to produce a regular distribution the number of generated particles might not exactly match the one specified in the input file.

Possible values: Any one of ascii file, probability density function, quadrature points, random uniform, reference cell, uniform box, uniform radial

- *Parameter name:* Particle weight

Value: 10

Default: 10

Description: Weight that is associated with the computational load of a single particle. The sum of particle weights will be added to the sum of cell weights to determine the partitioning of the mesh if the ‘repartition’ particle load balancing strategy is selected. The optimal weight depends on the used integrator and particle properties. In general for a more expensive integrator and more expensive properties a larger particle weight is recommended. Before adding the weights of particles, each cell already carries a weight of 1000 to account for the cost of field-based computations.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* Time between data output

Value: 1e8

Default: 1e8

Description: The time interval between each generation of output files. A value of zero indicates that output should be generated every time step.

Units: years if the 'Use years in output instead of seconds' parameter is set; seconds otherwise.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Update ghost particles

Value: false

Default: false

Description: Some particle interpolation algorithms require knowledge about particles in neighboring cells. To allow this, particles in ghost cells need to be exchanged between the processes neighboring this cell. This parameter determines whether this transport is happening.

Possible values: A boolean value (true or false)

A.126 Parameters in section Postprocess/Particles/Function

- *Parameter name:* Function constants

Value:

Default:

Description: Sometimes it is convenient to use symbolic constants in the expression that describes the function, rather than having to use its numeric value everywhere the constant appears. These values can be defined using this parameter, in the form 'var1=value1, var2=value2, ...'.

A typical example would be to set this runtime parameter to 'pi=3.1415926536' and then use 'pi' in the expression of the actual formula. (That said, for convenience this class actually defines both 'pi' and 'Pi' by default, but you get the idea.)

Possible values: Any string

- *Parameter name:* Function expression

Value: 0

Default: 0

Description: The formula that denotes the function you want to evaluate for particular values of the independent variables. This expression may contain any of the usual operations such as addition or multiplication, as well as all of the common functions such as 'sin' or 'cos'. In addition, it may contain expressions like 'if(x>0, 1, -1)' where the expression evaluates to the second argument if the first argument is true, and to the third argument otherwise. For a full overview of possible expressions accepted see the documentation of the muparser library at <http://muparser.beltoforion.de/>.

If the function you are describing represents a vector-valued function with multiple components, then separate the expressions for individual components by a semicolon.

Possible values: Any string

- *Parameter name:* Number of components

Value: 1

Default: 1

Description: The number of function components where each component is described by a function expression delimited by a ';'.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* Variable names

Value: x,y,t

Default: x,y,t

Description: The name of the variables as they will be used in the function, separated by commas. By default, the names of variables at which the function will be evaluated is ‘x’ (in 1d), ‘x,y’ (in 2d) or ‘x,y,z’ (in 3d) for spatial coordinates and ‘t’ for time. You can then use these variable names in your function expression and they will be replaced by the values of these variables at which the function is currently evaluated. However, you can also choose a different set of names for the independent variables at which to evaluate your function expression. For example, if you work in spherical coordinates, you may wish to set this input parameter to ‘r,phi,theta,t’ and then use these variable names in your function expression.

Possible values: Any string

A.127 Parameters in section Postprocess/Particles/Generator

A.128 Parameters in section Postprocess/Particles/Generator/Ascii file

- *Parameter name:* Data directory

Value: \$ASPECT_SOURCE_DIR/data/particle/generator/ascii/

Default: \$ASPECT_SOURCE_DIR/data/particle/generator/ascii/

Description: The name of a directory that contains the particle data. This path may either be absolute (if starting with a '/') or relative to the current directory. The path may also include the special text '\$ASPECT_SOURCE_DIR' which will be interpreted as the path in which the ASPECT source files were located when ASPECT was compiled. This interpretation allows, for example, to reference files located in the ‘data/’ subdirectory of ASPECT.

Possible values: A directory name

- *Parameter name:* Data file name

Value: particle.dat

Default: particle.dat

Description: The name of the particle file.

Possible values: Any string

A.129 Parameters in section Postprocess/Particles/Generator/Probability density function

- *Parameter name:* Function constants

Value:

Default:

Description: Sometimes it is convenient to use symbolic constants in the expression that describes the function, rather than having to use its numeric value everywhere the constant appears. These values can be defined using this parameter, in the form ‘var1=value1, var2=value2, ...’.

A typical example would be to set this runtime parameter to ‘pi=3.1415926536’ and then use ‘pi’ in the expression of the actual formula. (That said, for convenience this class actually defines both ‘pi’ and ‘Pi’ by default, but you get the idea.)

Possible values: Any string

- **Parameter name:** Function expression

Value: 0

Default: 0

Description: The formula that denotes the function you want to evaluate for particular values of the independent variables. This expression may contain any of the usual operations such as addition or multiplication, as well as all of the common functions such as ‘sin’ or ‘cos’. In addition, it may contain expressions like ‘if($x > 0$, 1, -1)’ where the expression evaluates to the second argument if the first argument is true, and to the third argument otherwise. For a full overview of possible expressions accepted see the documentation of the muparser library at <http://muparser.beltoforion.de/>.

If the function you are describing represents a vector-valued function with multiple components, then separate the expressions for individual components by a semicolon.

Possible values: Any string

- **Parameter name:** Random cell selection

Value: true

Default: true

Description: If true, particle numbers per cell are calculated randomly according to their respective probability density. This means particle numbers per cell can deviate statistically from the integral of the probability density. If false, first determine how many particles each cell should have based on the integral of the density over each of the cells, and then once we know how many particles we want on each cell, choose their locations randomly within each cell.

Possible values: A boolean value (true or false)

- **Parameter name:** Random number seed

Value: 5432

Default: 5432

Description: The seed for the random number generator that controls the particle generation. Keep constant to generate identical particle distributions in subsequent model runs. Change to get a different distribution. In parallel computations the seed is further modified on each process to ensure different particle patterns on different processes.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- **Parameter name:** Variable names

Value: x,y,t

Default: x,y,t

Description: The name of the variables as they will be used in the function, separated by commas. By default, the names of variables at which the function will be evaluated is ‘x’ (in 1d), ‘x,y’ (in 2d) or ‘x,y,z’ (in 3d) for spatial coordinates and ‘t’ for time. You can then use these variable names in your function expression and they will be replaced by the values of these variables at which the function is currently evaluated. However, you can also choose a different set of names for the independent variables at which to evaluate your function expression. For example, if you work in spherical coordinates, you may wish to set this input parameter to ‘r,phi,theta,t’ and then use these variable names in your function expression.

Possible values: Any string

A.130 Parameters in section Postprocess/Particles/Generator/Reference cell

- *Parameter name:* Number of particles per cell per direction

Value: 2

Default: 2

Description: List of number of particles to create per cell and spatial dimension. The size of the list is the number of spatial dimensions. If only one value is given, then each spatial dimension is set to the same value. The list of numbers are parsed as a floating point number (so that one can specify, for example, '1e4' particles) but it is interpreted as an integer, of course.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$]

A.131 Parameters in section Postprocess/Particles/Generator/Uniform box

- *Parameter name:* Maximum x

Value: 1

Default: 1

Description: Maximum x coordinate for the region of particles.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Maximum y

Value: 1

Default: 1

Description: Maximum y coordinate for the region of particles.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Maximum z

Value: 1

Default: 1

Description: Maximum z coordinate for the region of particles.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Minimum x

Value: 0

Default: 0

Description: Minimum x coordinate for the region of particles.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Minimum y

Value: 0

Default: 0

Description: Minimum y coordinate for the region of particles.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Minimum z
Value: 0
Default: 0
Description: Minimum z coordinate for the region of particles.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

A.132 Parameters in section Postprocess/Particles/Generator/Uniform radial

- *Parameter name:* Center x
Value: 0
Default: 0
Description: x coordinate for the center of the spherical region, where particles are generated.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Center y
Value: 0
Default: 0
Description: y coordinate for the center of the spherical region, where particles are generated.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Center z
Value: 0
Default: 0
Description: z coordinate for the center of the spherical region, where particles are generated.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$
- *Parameter name:* Maximum latitude
Value: 3.1415
Default: 3.1415
Description: Maximum latitude coordinate for the region of particles in degrees. Measured from the center position.
Possible values: A floating point number v such that $0 \leq v \leq 180$
- *Parameter name:* Maximum longitude
Value: 3.1415
Default: 3.1415
Description: Maximum longitude coordinate for the region of particles in degrees. Measured from the center position.
Possible values: A floating point number v such that $0 \leq v \leq 360$
- *Parameter name:* Maximum radius
Value: 1
Default: 1
Description: Maximum radial coordinate for the region of particles. Measured from the center position.
Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Minimum latitude

Value: 0

Default: 0

Description: Minimum latitude coordinate for the region of particles in degrees. Measured from the center position.

Possible values: A floating point number v such that $0 \leq v \leq 180$

- *Parameter name:* Minimum longitude

Value: 0

Default: 0

Description: Minimum longitude coordinate for the region of particles in degrees. Measured from the center position.

Possible values: A floating point number v such that $0 \leq v \leq 360$

- *Parameter name:* Minimum radius

Value: 0

Default: 0

Description: Minimum radial coordinate for the region of particles. Measured from the center position.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Radial layers

Value: 1

Default: 1

Description: The number of radial shells of particles that will be generated around the central point.

Possible values: An integer n such that $1 \leq n \leq 2147483647$

A.133 Parameters in section Postprocess/Particles/Interpolator

A.134 Parameters in section Postprocess/Particles/Interpolator/Bilinear least squares

- *Parameter name:* Global particle property maximum

Value: 1.7976931348623157e+308

Default: 1.7976931348623157e+308

Description: The maximum global particle property values that will be used as a limiter for the bilinear least squares interpolation. The number of the input 'Global particle property maximum' values separated by ',' has to be the same as the number of particle properties.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Global particle property minimum

Value: -1.7976931348623157e+308

Default: -1.7976931348623157e+308

Description: The minimum global particle property that will be used as a limiter for the bilinear least squares interpolation. The number of the input 'Global particle property minimum' values separated by ',' has to be the same as the number of particle properties.

Possible values: A list of 0 to 4294967295 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]

- *Parameter name:* Use limiter

Value: false

Default: false

Description: Whether to apply a global particle property limiting scheme to the interpolated particle properties.

Possible values: A boolean value (true or false)

A.135 Parameters in section Postprocess/Particles/Melt particle

- *Parameter name:* Threshold for melt presence

Value: 1e-3

Default: 1e-3

Description: The minimum porosity that has to be present at the position of a particle for it to be considered a melt particle (in the sense that the melt presence property is set to 1).

Possible values: A floating point number v such that $0 \leq v \leq 1$

A.136 Parameters in section Postprocess/Point values

- *Parameter name:* Evaluation points

Value:

Default:

Description: The list of points at which the solution should be evaluated. Points need to be separated by semicolons, and coordinates of each point need to be separated by commas.

Possible values: A list of 0 to 4294967295 elements separated by <;> where each element is [A list of 2 to 2 elements where each element is [A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$]]

A.137 Parameters in section Postprocess/Topography

- *Parameter name:* Output to file

Value: false

Default: false

Description: Whether or not to write topography to a text file named 'topography.NNNNN' in the output directory

Possible values: A list of 0 to 4294967295 elements where each element is [A boolean value (true or false)]

- *Parameter name:* Time between text output

Value: 0.

Default: 0.

Description: The time interval between each generation of text output files. A value of zero indicates that output should be generated in each time step. Units: years if the 'Use years in output instead of seconds' parameter is set; seconds otherwise.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.138 Parameters in section Postprocess/Visualization

- Parameter name: `Filter output`

Value: false

Default: false

Description: deal.II offers the possibility to filter duplicate vertices for HDF5 output files. This merges the vertices of adjacent cells and therefore saves disk space, but misrepresents discontinuous output properties. Activating this function reduces the disk space by about a factor of 2^{dim} for HDF5 output, and currently has no effect on other output formats.

Note: Warning: Setting this flag to true will result in visualization output that does not accurately represent discontinuous fields. This may be because you are using a discontinuous finite element for the pressure, temperature, or compositional variables, or because you use a visualization postprocessor that outputs quantities as discontinuous fields (e.g., the strain rate, viscosity, etc.). These will then all be visualized as *continuous* quantities even though, internally, ASPECT considers them as discontinuous fields.

Possible values: A boolean value (true or false)

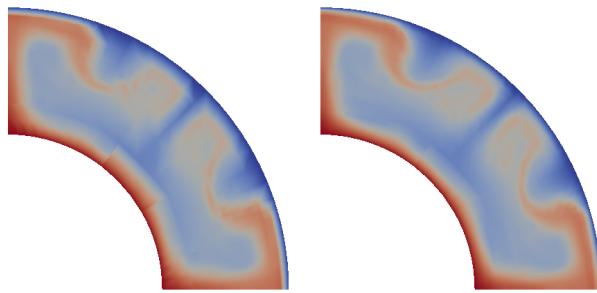
- Parameter name: `Interpolate output`

Value: false

Default: false

Description: deal.II offers the possibility to linearly interpolate output fields of higher order elements to a finer resolution. This somewhat compensates the fact that most visualization software only offers linear interpolation between grid points and therefore the output file is a very coarse representation of the actual solution field. Activating this option increases the spatial resolution in each dimension by a factor equal to the polynomial degree used for the velocity finite element (usually 2). In other words, instead of showing one quadrilateral or hexahedron in the visualization per cell on which ASPECT computes, it shows multiple (for quadratic elements, it will describe each cell of the mesh on which we compute as 2×2 or $2 \times 2 \times 2$ cells in 2d and 3d, respectively; correspondingly more subdivisions are used if you use cubic, quartic, or even higher order elements for the velocity).

The effect of using this option can be seen in the following picture showing a variation of the output produced with the input files from Section 5.3.1:



Here, the left picture shows one visualization cell per computational cell (i.e., the option is switch off, as is the default), and the right picture shows the same simulation with the option switched on. The

images show the same data, demonstrating that interpolating the solution onto bilinear shape functions as is commonly done in visualizing data loses information.

Of course, activating this option also greatly increases the amount of data ASPECT will write to disk: approximately by a factor of 4 in 2d, and a factor of 8 in 3d, when using quadratic elements for the velocity, and correspondingly more for even higher order elements.

Possible values: A boolean value (true or false)

- *Parameter name:* **List of output variables**

Value:

Default:

Description: A comma separated list of visualization objects that should be run whenever writing graphical output. By default, the graphical output files will always contain the primary variables velocity, pressure, and temperature. However, one frequently wants to also visualize derived quantities, such as the thermodynamic phase that corresponds to a given temperature-pressure value, or the corresponding seismic wave speeds. The visualization objects do exactly this: they compute such derived quantities and place them into the output file. The current parameter is the place where you decide which of these additional output variables you want to have in your output file.

The following postprocessors are available:

‘Vp anomaly’: A visualization output object that generates output showing the percentage anomaly in the seismic compressional wave speed V_p as a spatially variable function with one value per cell. This anomaly is either shown as a percentage anomaly relative to the reference profile given by adiabatic conditions (with the compositions given by the current composition, such that the reference could potentially change through time), or as a percentage change relative to the laterally averaged velocity at the depth of the cell. This velocity is calculated by linear interpolation between average values calculated within equally thick depth slices. The number of depth slices in the domain is user-defined. Typically, the best results will be obtained if the number of depth slices is balanced between being large enough to capture step changes in velocities, but small enough to maintain a reasonable number of evaluation points per slice. Bear in mind that lateral averaging subsamples the finite element mesh. Note that this plugin requires a material model that provides seismic velocities.

‘Vs anomaly’: A visualization output object that generates output showing the percentage anomaly in the seismic shear wave speed V_s as a spatially variable function with one value per cell. This anomaly is either shown as a percentage anomaly relative to the reference profile given by adiabatic conditions (with the compositions given by the current composition, such that the reference could potentially change through time), or as a percentage change relative to the laterally averaged velocity at the depth of the cell. This velocity is calculated by linear interpolation between average values calculated within equally thick depth slices. The number of depth slices in the domain is user-defined. Typically, the best results will be obtained if the number of depth slices is balanced between being large enough to capture step changes in velocities, but small enough to maintain a reasonable number of evaluation points per slice. Bear in mind that lateral averaging subsamples the finite element mesh. Note that this plugin requires a material model that provides seismic velocities.

‘adiabat’: A visualization output object that generates adiabatic temperature, pressure, density, and density derivative as produced by AdiabaticConditions.

‘artificial viscosity’: A visualization output object that generates output showing the value of the artificial viscosity on each cell.

‘artificial viscosity composition’: A visualization output object that generates output showing the value of the artificial viscosity for a compositional field on each cell.

‘boundary indicators’: A visualization output object that generates output about the used boundary indicators. In a loop over the active cells, if a cell lies at a domain boundary, the boundary indicator

of the face along the boundary is requested. In case the cell does not lie along any domain boundary, the cell is assigned the value of the largest used boundary indicator plus one. When a cell is situated in one of the corners of the domain, multiple faces will have a boundary indicator. This postprocessor returns the value of the first face along a boundary that is encountered in a loop over all the faces.

‘compositional vector’: A visualization output object that outputs vectors whose components are derived from compositional fields. Input parameters for this postprocessor are defined in section Postprocess/Visualization/Compositional fields as vectors

‘density’: A visualization output object that generates output for the density.

‘depth’: A visualization output postprocessor that outputs the depth for all points inside the domain, as determined by the geometry model.

‘dynamic topography’: A visualization output object that generates output for the dynamic topography at the top and bottom of the model space. The approach to determine the dynamic topography requires us to compute the stress tensor and evaluate the component of it in the direction in which gravity acts. In other words, we compute $\sigma_{rr} = \hat{g}^T(2\eta\varepsilon(\mathbf{u}) - \frac{1}{3}(\operatorname{div} \mathbf{u})I)\hat{g} - p_d$ where $\hat{g} = \mathbf{g}/\|\mathbf{g}\|$ is the direction of the gravity vector \mathbf{g} and $p_d = p - p_a$ is the dynamic pressure computed by subtracting the adiabatic pressure p_a from the total pressure p computed as part of the Stokes solve. From this, the dynamic topography is computed using the formula $h = \frac{\sigma_{rr}}{(\mathbf{g} \cdot \mathbf{n})\rho}$ where ρ is the density at the cell center. For the bottom surface we chose the convection that positive values are up (out) and negative values are in (down), analogous to the deformation of the upper surface. Note that this implementation takes the direction of gravity into account, which means that reversing the flow in backward advection calculations will not reverse the instantaneous topography because the reverse flow will be divided by the reverse surface gravity.

Strictly speaking, the dynamic topography is of course a quantity that is only of interest at the surface. However, we compute it everywhere to make things fit into the framework within which we produce data for visualization. You probably only want to visualize whatever data this postprocessor generates at the surface of your domain and simply ignore the rest of the data generated.

‘error indicator’: A visualization output object that generates output showing the estimated error or other mesh refinement indicator as a spatially variable function with one value per cell.

‘geoid’: Visualization for the geoid solution. The geoid is given by the equivalent water column height due to a gravity perturbation. (Units: m)

‘gravity’: A visualization output object that outputs the gravity vector.

‘heat flux map’: A visualization output object that generates output for the heat flux density across each boundary. The heat flux density is computed in outward direction, i.e., from the domain to the outside, using the formula $-k\nabla T \cdot \mathbf{n}$, where k is the thermal conductivity as reported by the material model, T is the temperature, and \mathbf{n} is the outward normal. Note that the quantity so computed does not include any energy transported across the boundary by material transport in cases where $\mathbf{u} \cdot \mathbf{n} \neq 0$. At the edge of the domain (e.g. top left corner) the heat flux density is calculated as the sum of the heat flux across both boundaries of the cell (e.g. top and left boundary) divided by the sum of both face areas. The integrated heatflux for each boundary can be obtained from the heat flux statistics postprocessor.

‘heating’: A visualization output object that generates output for all the heating terms used in the energy equation.

‘material properties’: A visualization output object that generates output for the material properties given by the material model. There are a number of other visualization postprocessors that offer to write individual material properties. However, they all individually have to evaluate the material model. This is inefficient if one wants to output more than just one or two of the fields provided by the material model. The current postprocessor allows to output a (potentially large) subset of all of

the information provided by material models at once, with just a single material model evaluation per output point.

‘maximum horizontal compressive stress’: A plugin that computes the direction and magnitude of the maximum horizontal component of the compressive stress as a vector field. The direction of this vector can often be used to visualize the principal mode of deformation (e.g., at normal faults or extensional margins) and can be correlated with seismic anisotropy. Recall that the *compressive* stress is simply the negative stress, $\sigma_c = -\sigma = -[2\eta(\varepsilon(\mathbf{u}) - \frac{1}{3}(\nabla \cdot \mathbf{u})I) + pI]$.

Following [LT07], we define the maximum horizontal stress direction as that *horizontal* direction \mathbf{n} that maximizes $\mathbf{n}^T \sigma_c \mathbf{n}$. We call a vector *horizontal* if it is perpendicular to the gravity vector \mathbf{g} .

In two space dimensions, \mathbf{n} is simply a vector that is horizontal (we choose one of the two possible choices). This direction is then scaled by the size of the horizontal stress in this direction, i.e., the plugin outputs the vector $\mathbf{w} = (\mathbf{n}^T \sigma_c \mathbf{n}) \mathbf{n}$.

In three space dimensions, given two horizontal, perpendicular, unit length, but otherwise arbitrarily chosen vectors \mathbf{u}, \mathbf{v} , we can express $\mathbf{n} = (\cos \alpha)\mathbf{u} + (\sin \alpha)\mathbf{v}$ where α maximizes the expression

$$f(\alpha) = \mathbf{n}^T \sigma_c \mathbf{n} = (\mathbf{u}^T \sigma_c \mathbf{u})(\cos \alpha)^2 + 2(\mathbf{u}^T \sigma_c \mathbf{v})(\cos \alpha)(\sin \alpha) + (\mathbf{v}^T \sigma_c \mathbf{v})(\sin \alpha)^2.$$

The maximum of $f(\alpha)$ is attained where $f'(\alpha) = 0$. Evaluating the derivative and using trigonometric identities, one finds that α has to satisfy the equation

$$\tan(2\alpha) = \frac{\mathbf{u}^T \sigma_c \mathbf{v}}{\mathbf{u}^T \sigma_c \mathbf{u} - \mathbf{v}^T \sigma_c \mathbf{v}}.$$

Since the transform $\alpha \mapsto \alpha + \pi$ flips the direction of \mathbf{n} , we only need to seek a solution to this equation in the interval $\alpha \in [0, \pi)$. These are given by $\alpha_1 = \frac{1}{2} \arctan \frac{\mathbf{u}^T \sigma_c \mathbf{v}}{\mathbf{u}^T \sigma_c \mathbf{u} - \mathbf{v}^T \sigma_c \mathbf{v}}$ and $\alpha_2 = \alpha_1 + \frac{\pi}{2}$, one of which will correspond to a minimum and the other to a maximum of $f(\alpha)$. One checks the sign of $f''(\alpha) = -2(\mathbf{u}^T \sigma_c \mathbf{u} - \mathbf{v}^T \sigma_c \mathbf{v}) \cos(2\alpha) - 2(\mathbf{u}^T \sigma_c \mathbf{v}) \sin(2\alpha)$ for each of these to determine the α that maximizes $f(\alpha)$, and from this immediately arrives at the correct form for the maximum horizontal stress \mathbf{n} .

The description above computes a 3d *direction* vector \mathbf{n} . If one were to scale this vector the same way as done in 2d, i.e., with the magnitude of the stress in this direction, one will typically get vectors whose length is principally determined by the hydrostatic pressure at a given location simply because the hydrostatic pressure is the largest component of the overall stress. On the other hand, the hydrostatic pressure does not determine any principle direction because it is an isotropic, anti-compressive force. As a consequence, there are often points in simulations (e.g., at the center of convection rolls) where the stress has no dominant horizontal direction, and the algorithm above will then in essence choose a random direction because the stress is approximately equal in all horizontal directions. If one scaled the output by the magnitude of the stress in this direction (i.e., approximately equal to the hydrostatic pressure at this point), one would get randomly oriented vectors at these locations with significant lengths.

To avoid this problem, we scale the maximal horizontal compressive stress direction \mathbf{n} by the *difference* between the stress in the maximal and minimal horizontal stress directions. In other words, let $\mathbf{n}_\perp = (\sin \alpha)\mathbf{u} - (\cos \alpha)\mathbf{v}$ be the horizontal direction perpendicular to \mathbf{n} , then this plugin outputs the vector quantity $\mathbf{w} = (\mathbf{n}^T \sigma_c \mathbf{n} - \mathbf{n}_\perp^T \sigma_c \mathbf{n}_\perp) \mathbf{n}$. In other words, the length of the vector produced indicates *how dominant* the direction of maximal horizontal compressive strength is.

Fig. 98 shows a simple example for this kind of visualization in 3d.

‘melt fraction’: A visualization output object that generates output for the melt fraction at the temperature and pressure of the current point. If the material model computes a melt fraction, this is the quantity that will be visualized. Otherwise, a specific parametrization for batch melting (as described in the following) will be used. It does not take into account latent heat. If there are no compositional

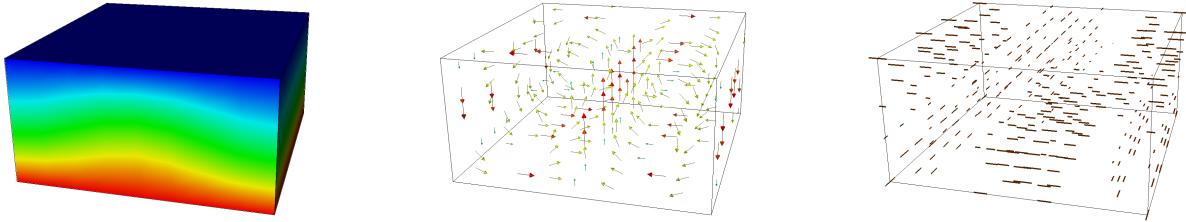


Figure 98: Illustration of the ‘maximum horizontal compressive stress’ visualization plugin. The left figure shows a ridge-like temperature anomaly. Together with no-slip boundary along all six boundaries, this results in two convection rolls (center). The maximal horizontal compressive strength at the bottom center of the domain is perpendicular to the ridge because the flow comes together there from the left and right, yielding a compressive force in left-right direction. At the top of the model, the flow separates outward, leading to a negative compressive stress in left-right direction; because there is no flow in front-back direction, the compressive strength in front-back direction is zero, making the along-ridge direction the dominant one. At the center of the convection rolls, both horizontal directions yield the same stress; the plugin therefore chooses an essentially arbitrary horizontal vector, but then uses a zero magnitude given that the difference between the maximal and minimal horizontal stress is zero at these points.

fields, this postprocessor will visualize the melt fraction of peridotite (calculated using the anhydrous model of Katz, 2003). If there is at least one compositional field, the postprocessor assumes that the first compositional field is the content of pyroxenite, and will visualize the melt fraction for a mixture of peridotite and pyroxenite (using the melting model of Sobolev, 2011 for pyroxenite). All the parameters that were used in these calculations can be changed in the input file, the most relevant maybe being the mass fraction of Cpx in peridotite in the Katz melting model (Mass fraction cpx), which right now has a default of 15%. The corresponding p-T-diagrams can be generated by running the tests `melt_postprocessor_peridotite` and `melt_postprocessor_pyroxenite`.

‘melt material properties’: A visualization output object that generates output for melt related properties of the material model. Note that this postprocessor always outputs the compaction pressure, but can output a large range of additional properties, as selected in the “List of properties” parameter.

‘named additional outputs’: Some material models can compute quantities other than those that typically appear in the equations that ASPECT solves (such as the viscosity, density, etc). Examples of quantities material models may be able to compute are seismic velocities, or other quantities that can be derived from the state variables and the material coefficients such as the stress or stress anisotropies. These quantities are generically referred to as ‘named outputs’ because they are given an explicit name different from the usual outputs of material models.

This visualization postprocessor outputs whatever quantities the material model can compute. What quantities these are is specific to the material model in use for a simulation, and for many models in fact does not contain any named outputs at all.

‘nonadiabatic pressure’: A visualization output object that generates output for the non-adiabatic component of the pressure.

‘nonadiabatic temperature’: A visualization output object that generates output for the non-adiabatic component of the temperature.

‘particle count’: A visualization output object that generates output about the number of particles per cell.

‘partition’: A visualization output object that generates output for the parallel partition that every cell of the mesh is associated with.

‘shear stress’: A visualization output object that generates output for the 3 (in 2d) or 6 (in 3d) components of the shear stress tensor, i.e., for the components of the tensor $2\eta\varepsilon(\mathbf{u})$ in the incompressible case and $2\eta [\varepsilon(\mathbf{u}) - \frac{1}{3}(\text{tr } \varepsilon(\mathbf{u}))\mathbf{I}]$ in the compressible case. The shear stress differs from the full stress tensor by the absence of the pressure.

‘spd factor’: A visualization output object that generates output for the spd factor. The spd factor is a factor which scales a part of the Jacobian used for the Newton solver to make sure that the Jacobian remains positive definite.

‘specific heat’: A visualization output object that generates output for the specific heat C_p .

‘strain rate’: A visualization output object that generates output for the norm of the strain rate, i.e., for the quantity $\sqrt{\varepsilon(\mathbf{u}) : \varepsilon(\mathbf{u})}$ in the incompressible case and $\sqrt{[\varepsilon(\mathbf{u}) - \frac{1}{3}(\text{tr } \varepsilon(\mathbf{u}))\mathbf{I}] : [\varepsilon(\mathbf{u}) - \frac{1}{3}(\text{tr } \varepsilon(\mathbf{u}))\mathbf{I}]}$ in the compressible case.

‘stress’: A visualization output object that generates output for the 3 (in 2d) or 6 (in 3d) components of the stress tensor, i.e., for the components of the tensor $2\eta\varepsilon(\mathbf{u}) + pI$ in the incompressible case and $2\eta [\varepsilon(\mathbf{u}) - \frac{1}{3}(\text{tr } \varepsilon(\mathbf{u}))\mathbf{I}] + pI$ in the compressible case.

‘thermal conductivity’: A visualization output object that generates output for the thermal conductivity k .

‘thermal diffusivity’: A visualization output object that generates output for the thermal diffusivity $\kappa = \frac{k}{\rho C_p}$, with k the thermal conductivity.

‘thermal expansivity’: A visualization output object that generates output for the thermal expansivity.

‘vertical heat flux’: A visualization output object that generates output for the heat flux in the vertical direction, which is the sum of the advective and the conductive heat flux, with the sign convention of positive flux upwards.

‘viscosity’: A visualization output object that generates output for the viscosity.

‘volumetric strain rate’: A visualization output object that generates output for the volumetric strain rate, i.e., for the quantity $\nabla \cdot \mathbf{u} = \text{div } \mathbf{u} = \text{trace } \varepsilon(\mathbf{u})$. This should be zero (in some average sense) in incompressible convection models, but can be non-zero in compressible models and models with melt transport.

Possible values: A comma-separated list of any of Vp anomaly, Vs anomaly, adiabat, artificial viscosity, artificial viscosity composition, boundary indicators, compositional vector, density, depth, dynamic topography, error indicator, geoid, gravity, heat flux map, heating, material properties, maximum horizontal compressive stress, melt fraction, melt material properties, named additional outputs, nonadiabatic pressure, nonadiabatic temperature, particle count, partition, shear stress, spd factor, specific heat, strain rate, stress, thermal conductivity, thermal diffusivity, thermal expansivity, vertical heat flux, viscosity, volumetric strain rate

- *Parameter name:* Number of grouped files

Value: 16

Default: 16

Description: VTU file output supports grouping files from several CPUs into a given number of files using MPI I/O when writing on a parallel filesystem. Select 0 for no grouping. This will disable parallel file output and instead write one file per processor. A value of 1 will generate one big file containing the whole solution, while a larger value will create that many files (at most as many as there are MPI ranks).

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* `Output format`

Value: vtu

Default: vtu

Description: The file format to be used for graphical output.

Possible values: Any one of none, dx, ucd, gnuplot, povray, eps, gmv, tecplot, tecplot_binary, vtk, vtu, hdf5, svg, deal.II intermediate

- *Parameter name:* `Output mesh velocity`

Value: false

Default: false

Description: For free surface computations Aspect uses an Arbitrary-Lagrangian-Eulerian formulation to handle deforming the domain, so the mesh has its own velocity field. This may be written as an output field by setting this parameter to true.

Possible values: A boolean value (true or false)

- *Parameter name:* `Temporary output location`

Value:

Default:

Description: On large clusters it can be advantageous to first write the output to a temporary file on a local file system and later move this file to a network file system. If this variable is set to a non-empty string it will be interpreted as a temporary storage location.

Possible values: Any string

- *Parameter name:* `Time between graphical output`

Value: 1e8

Default: 1e8

Description: The time interval between each generation of graphical output files. A value of zero indicates that output should be generated in each time step. Units: years if the 'Use years in output instead of seconds' parameter is set; seconds otherwise.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* `Time steps between graphical output`

Value: 2147483647

Default: 2147483647

Description: The maximum number of time steps between each generation of graphical output files.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* `Write in background thread`

Value: false

Default: false

Description: File operations can potentially take a long time, blocking the progress of the rest of the model run. Setting this variable to 'true' moves this process into a background thread, while the rest of the model continues.

Possible values: A boolean value (true or false)

A.139 Parameters in section Postprocess/Visualization/Artificial viscosity composition

- *Parameter name:* Name of compositional field

Value:

Default:

Description: The name of the compositional field whose output should be visualized.

Possible values: Any string

A.140 Parameters in section Postprocess/Visualization/Compositional fields as vectors

- *Parameter name:* Names of fields

Value:

Default:

Description: A list of sets of compositional fields which should be output as vectors. Sets are separated from each other by semicolons and vector components within each set are separated by commas (e.g. $vec1_x, vec1_y ; vec2_x, vec2_y$) where each name must be a defined named compositional field. If only one name is given in a set, it is interpreted as the first in a sequence of dim consecutive compositional fields.

Possible values: Any string

- *Parameter name:* Names of vectors

Value:

Default:

Description: Names of vectors as they will appear in the output.

Possible values: A list of 0 to 4294967295 elements where each element is [Any string]

A.141 Parameters in section Postprocess/Visualization/Material properties

- *Parameter name:* List of material properties

Value: density,thermal expansivity,specific heat,viscosity

Default: density,thermal expansivity,specific heat,viscosity

Description: A comma separated list of material properties that should be written whenever writing graphical output. By default, the material properties will always contain the density, thermal expansivity, specific heat and viscosity. The following material properties are available:

viscosity|density|thermal expansivity|specific heat|thermal conductivity|thermal diffusivity|compressibility|entropy derivative temperature|entropy derivative pressure|reaction terms|melt fraction

Possible values: A comma-separated list of any of viscosity, density, thermal expansivity, specific heat, thermal conductivity, thermal diffusivity, compressibility, entropy derivative temperature, entropy derivative pressure, reaction terms, melt fraction

A.142 Parameters in section Postprocess/Visualization/Melt fraction

- Parameter name: A1

Value: 1085.7

Default: 1085.7

Description: Constant parameter in the quadratic function that approximates the solidus of peridotite. Units: $\text{Å}^{\circ}\text{C}$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: A2

Value: 1.329e-7

Default: 1.329e-7

Description: Prefactor of the linear pressure term in the quadratic function that approximates the solidus of peridotite. Units: $\text{Å}^{\circ}\text{C}/\text{Pa}$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: A3

Value: -5.1e-18

Default: -5.1e-18

Description: Prefactor of the quadratic pressure term in the quadratic function that approximates the solidus of peridotite. Units: $\text{Å}^{\circ}\text{C}/(\text{Pa}^2)$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: B1

Value: 1475.0

Default: 1475.0

Description: Constant parameter in the quadratic function that approximates the lherzolite liquidus used for calculating the fraction of peridotite-derived melt. Units: $\text{Å}^{\circ}\text{C}$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: B2

Value: 8.0e-8

Default: 8.0e-8

Description: Prefactor of the linear pressure term in the quadratic function that approximates the lherzolite liquidus used for calculating the fraction of peridotite-derived melt. Units: $\text{Å}^{\circ}\text{C}/\text{Pa}$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- Parameter name: B3

Value: -3.2e-18

Default: -3.2e-18

Description: Prefactor of the quadratic pressure term in the quadratic function that approximates the lherzolite liquidus used for calculating the fraction of peridotite-derived melt. Units: $\text{Å}^{\circ}\text{C}/(\text{Pa}^2)$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* C1

Value: 1780.0

Default: 1780.0

Description: Constant parameter in the quadratic function that approximates the liquidus of peridotite. Units: K .

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* C2

Value: 4.50e-8

Default: 4.50e-8

Description: Prefactor of the linear pressure term in the quadratic function that approximates the liquidus of peridotite. Units: K/Pa .

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* C3

Value: -2.0e-18

Default: -2.0e-18

Description: Prefactor of the quadratic pressure term in the quadratic function that approximates the liquidus of peridotite. Units: $\text{K}/(\text{Pa}^2)$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* D1

Value: 976.0

Default: 976.0

Description: Constant parameter in the quadratic function that approximates the solidus of pyroxenite. Units: K .

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* D2

Value: 1.329e-7

Default: 1.329e-7

Description: Prefactor of the linear pressure term in the quadratic function that approximates the solidus of pyroxenite. Note that this factor is different from the value given in Sobolev, 2011, because they use the potential temperature whereas we use the absolute temperature. Units: K/Pa .

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* D3

Value: -5.1e-18

Default: -5.1e-18

Description: Prefactor of the quadratic pressure term in the quadratic function that approximates the solidus of pyroxenite. Units: $\text{K}/(\text{Pa}^2)$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* E1

Value: 663.8

Default: 663.8

Description: Prefactor of the linear depletion term in the quadratic function that approximates the melt fraction of pyroxenite. Units: $\text{Å}^3\text{C}/\text{Pa}$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* E2

Value: -611.4

Default: -611.4

Description: Prefactor of the quadratic depletion term in the quadratic function that approximates the melt fraction of pyroxenite. Units: $\text{Å}^3\text{C}/(\text{Pa}^2)$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Mass fraction cpx

Value: 0.15

Default: 0.15

Description: Mass fraction of clinopyroxene in the peridotite to be molten. Units: non-dimensional.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* beta

Value: 1.5

Default: 1.5

Description: Exponent of the melting temperature in the melt fraction calculation. Units: non-dimensional.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* r1

Value: 0.5

Default: 0.5

Description: Constant in the linear function that approximates the clinopyroxene reaction coefficient. Units: non-dimensional.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* r2

Value: 8e-11

Default: 8e-11

Description: Prefactor of the linear pressure term in the linear function that approximates the clinopyroxene reaction coefficient. Units: $1/\text{Pa}$.

Possible values: A floating point number v such that $-\text{MAX_DOUBLE} \leq v \leq \text{MAX_DOUBLE}$

A.143 Parameters in section Postprocess/Visualization/Melt material properties

- *Parameter name:* List of properties

Value: compaction viscosity,permeability

Default: compaction viscosity,permeability

Description: A comma separated list of melt properties that should be written whenever writing graphical output. The following material properties are available:

compaction viscosity|fluid viscosity|permeability|fluid density|fluid density gradient|is melt cell|darcy coefficient|darcy coefficient no cutoff

Possible values: A comma-separated list of any of compaction viscosity, fluid viscosity, permeability, fluid density, fluid density gradient, is melt cell, darcy coefficient, darcy coefficient no cutoff

A.144 Parameters in section Postprocess/Visualization/Vp anomaly

- *Parameter name:* Average velocity scheme

Value: reference profile

Default: reference profile

Description: Scheme to compute the average velocity-depth profile. The reference profile option evaluates the conditions along the reference adiabat according to the material model. The lateral average option instead calculates a lateral average from subdivision of the mesh. The lateral average option may produce spurious results where there are sharp velocity changes.

Possible values: Any one of reference profile, lateral average

- *Parameter name:* Number of depth slices

Value: 50

Default: 50

Description: Number of depth slices used to define average seismic compressional wave velocities from which anomalies are calculated. Units: non-dimensional.

Possible values: An integer n such that $1 \leq n \leq 2147483647$

A.145 Parameters in section Postprocess/Visualization/Vs anomaly

- *Parameter name:* Average velocity scheme

Value: reference profile

Default: reference profile

Description: Scheme to compute the average velocity-depth profile. The reference profile option evaluates the conditions along the reference adiabat according to the material model. The lateral average option instead calculates a lateral average from subdivision of the mesh. The lateral average option may produce spurious results where there are sharp velocity changes.

Possible values: Any one of reference profile, lateral average

- *Parameter name:* Number of depth slices

Value: 50

Default: 50

Description: Number of depth slices used to define average seismic shear wave velocities from which anomalies are calculated. Units: non-dimensional.

Possible values: An integer n such that $1 \leq n \leq 2147483647$

A.146 Parameters in section Prescribed Stokes solution

- *Parameter name:* Model name

Value: unspecified

Default: unspecified

Description: Select one of the following models:

‘ascii data’: Implementation of a model in which the velocity is derived from files containing data in ascii format. Note the required format of the input data: The first lines may contain any number of comments if they begin with ‘#’, but one of these lines needs to contain the number of grid points in each dimension as for example ‘# POINTS: 3 3’. The order of the data columns has to be ‘x’, ‘y’, ‘v_x’ , ‘v_y’ in a 2d model and ‘x’, ‘y’, ‘z’, ‘v_x’ , ‘v_y’ , ‘v_z’ in a 3d model. Note that the data in the input files need to be sorted in a specific order: the first coordinate needs to ascend first, followed by the second and the third at last in order to assign the correct data to the prescribed coordinates. If you use a spherical model, then the data will still be handled as Cartesian, however the assumed grid changes. ‘x’ will be replaced by the radial distance of the point to the bottom of the model, ‘y’ by the azimuth angle and ‘z’ by the polar angle measured positive from the north pole. The grid will be assumed to be a latitude-longitude grid. Note that the order of spherical coordinates is ‘r’, ‘phi’, ‘theta’ and not ‘r’, ‘theta’, ‘phi’, since this allows for dimension independent expressions.

‘circle’: This value describes a vector field that rotates around the z-axis with constant angular velocity (i.e., with a velocity that increases with distance from the axis). The pressure is set to zero.

‘function’: This plugin allows to prescribe the Stokes solution for the velocity and pressure field in terms of an explicit formula. The format of these functions follows the syntax understood by the muparser library, see Section 4.7.3.

Possible values: Any one of ascii data, circle, function, unspecified

A.147 Parameters in section Prescribed Stokes solution/Ascii data model

- *Parameter name:* Data directory

Value: \$ASPECT_SOURCE_DIR/data/prescribed-stokes-solution/

Default: \$ASPECT_SOURCE_DIR/data/prescribed-stokes-solution/

Description: The name of a directory that contains the model data. This path may either be absolute (if starting with a '/') or relative to the current directory. The path may also include the special text ‘\$ASPECT_SOURCE_DIR’ which will be interpreted as the path in which the ASPECT source files were located when ASPECT was compiled. This interpretation allows, for example, to reference files located in the ‘data/’ subdirectory of ASPECT.

Possible values: A directory name

- *Parameter name:* Data file name

Value: box_2d.txt

Default: box_2d.txt

Description: The file name of the material data. Provide file in format: (Velocity file name).%s%d where %s is a string specifying the boundary of the model according to the names of the boundary indicators (of a box or a spherical shell).%d is any sprintf integer qualifier, specifying the format of the current file number.

Possible values: Any string

- *Parameter name:* Scale factor

Value: 1

Default: 1

Description: Scalar factor, which is applied to the boundary velocity. You might want to use this to scale the velocities to a reference model (e.g. with free-slip boundary) or another plate reconstruction. Another way to use this factor is to convert units of the input files. The unit is assumed to be m/s or m/yr depending on the 'Use years in output instead of seconds' flag. If you provide velocities in cm/yr set this factor to 0.01.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.148 Parameters in section Prescribed Stokes solution/Compaction pressure function

- *Parameter name:* Function constants

Value:

Default:

Description: Sometimes it is convenient to use symbolic constants in the expression that describes the function, rather than having to use its numeric value everywhere the constant appears. These values can be defined using this parameter, in the form 'var1=value1, var2=value2, ...'.

A typical example would be to set this runtime parameter to 'pi=3.1415926536' and then use 'pi' in the expression of the actual formula. (That said, for convenience this class actually defines both 'pi' and 'Pi' by default, but you get the idea.)

Possible values: Any string

- *Parameter name:* Function expression

Value: 0

Default: 0

Description: The formula that denotes the function you want to evaluate for particular values of the independent variables. This expression may contain any of the usual operations such as addition or multiplication, as well as all of the common functions such as 'sin' or 'cos'. In addition, it may contain expressions like 'if(x>0, 1, -1)' where the expression evaluates to the second argument if the first argument is true, and to the third argument otherwise. For a full overview of possible expressions accepted see the documentation of the muparser library at <http://muparser.beltoforion.de/>.

If the function you are describing represents a vector-valued function with multiple components, then separate the expressions for individual components by a semicolon.

Possible values: Any string

- *Parameter name:* Variable names

Value: x,y,t

Default: x,y,t

Description: The name of the variables as they will be used in the function, separated by commas. By default, the names of variables at which the function will be evaluated is 'x' (in 1d), 'x,y' (in 2d) or 'x,y,z' (in 3d) for spatial coordinates and 't' for time. You can then use these variable names in your function expression and they will be replaced by the values of these variables at which the function is currently evaluated. However, you can also choose a different set of names for the independent variables at which to evaluate your function expression. For example, if you work in spherical coordinates, you

may wish to set this input parameter to ‘r,phi,theta,t’ and then use these variable names in your function expression.

Possible values: Any string

A.149 Parameters in section Prescribed Stokes solution/Fluid pressure function

- *Parameter name:* Function constants

Value:

Default:

Description: Sometimes it is convenient to use symbolic constants in the expression that describes the function, rather than having to use its numeric value everywhere the constant appears. These values can be defined using this parameter, in the form ‘var1=value1, var2=value2, ...’.

A typical example would be to set this runtime parameter to ‘pi=3.1415926536’ and then use ‘pi’ in the expression of the actual formula. (That said, for convenience this class actually defines both ‘pi’ and ‘Pi’ by default, but you get the idea.)

Possible values: Any string

- *Parameter name:* Function expression

Value: 0

Default: 0

Description: The formula that denotes the function you want to evaluate for particular values of the independent variables. This expression may contain any of the usual operations such as addition or multiplication, as well as all of the common functions such as ‘sin’ or ‘cos’. In addition, it may contain expressions like ‘if(x>0, 1, -1)’ where the expression evaluates to the second argument if the first argument is true, and to the third argument otherwise. For a full overview of possible expressions accepted see the documentation of the muparser library at <http://muparser.beltoforion.de/>.

If the function you are describing represents a vector-valued function with multiple components, then separate the expressions for individual components by a semicolon.

Possible values: Any string

- *Parameter name:* Variable names

Value: x,y,t

Default: x,y,t

Description: The name of the variables as they will be used in the function, separated by commas. By default, the names of variables at which the function will be evaluated is ‘x’ (in 1d), ‘x,y’ (in 2d) or ‘x,y,z’ (in 3d) for spatial coordinates and ‘t’ for time. You can then use these variable names in your function expression and they will be replaced by the values of these variables at which the function is currently evaluated. However, you can also choose a different set of names for the independent variables at which to evaluate your function expression. For example, if you work in spherical coordinates, you may wish to set this input parameter to ‘r,phi,theta,t’ and then use these variable names in your function expression.

Possible values: Any string

A.150 Parameters in section Prescribed Stokes solution/Fluid velocity function

- *Parameter name:* Function constants

Value:

Default:

Description: Sometimes it is convenient to use symbolic constants in the expression that describes the function, rather than having to use its numeric value everywhere the constant appears. These values can be defined using this parameter, in the form ‘var1=value1, var2=value2, ...’.

A typical example would be to set this runtime parameter to ‘pi=3.1415926536’ and then use ‘pi’ in the expression of the actual formula. (That said, for convenience this class actually defines both ‘pi’ and ‘Pi’ by default, but you get the idea.)

Possible values: Any string

- *Parameter name:* Function expression

Value: 0; 0

Default: 0; 0

Description: The formula that denotes the function you want to evaluate for particular values of the independent variables. This expression may contain any of the usual operations such as addition or multiplication, as well as all of the common functions such as ‘sin’ or ‘cos’. In addition, it may contain expressions like ‘if(x>0, 1, -1)’ where the expression evaluates to the second argument if the first argument is true, and to the third argument otherwise. For a full overview of possible expressions accepted see the documentation of the muparser library at <http://muparser.beltoforion.de/>.

If the function you are describing represents a vector-valued function with multiple components, then separate the expressions for individual components by a semicolon.

Possible values: Any string

- *Parameter name:* Variable names

Value: x,y,t

Default: x,y,t

Description: The name of the variables as they will be used in the function, separated by commas. By default, the names of variables at which the function will be evaluated is ‘x’ (in 1d), ‘x,y’ (in 2d) or ‘x,y,z’ (in 3d) for spatial coordinates and ‘t’ for time. You can then use these variable names in your function expression and they will be replaced by the values of these variables at which the function is currently evaluated. However, you can also choose a different set of names for the independent variables at which to evaluate your function expression. For example, if you work in spherical coordinates, you may wish to set this input parameter to ‘r,phi,theta,t’ and then use these variable names in your function expression.

Possible values: Any string

A.151 Parameters in section Prescribed Stokes solution/Pressure function

- *Parameter name:* Function constants

Value:

Default:

Description: Sometimes it is convenient to use symbolic constants in the expression that describes the function, rather than having to use its numeric value everywhere the constant appears. These values can be defined using this parameter, in the form ‘var1=value1, var2=value2, ...’.

A typical example would be to set this runtime parameter to ‘pi=3.1415926536’ and then use ‘pi’ in the expression of the actual formula. (That said, for convenience this class actually defines both ‘pi’ and ‘Pi’ by default, but you get the idea.)

Possible values: Any string

- *Parameter name:* Function expression

Value: 0

Default: 0

Description: The formula that denotes the function you want to evaluate for particular values of the independent variables. This expression may contain any of the usual operations such as addition or multiplication, as well as all of the common functions such as ‘sin’ or ‘cos’. In addition, it may contain expressions like ‘if($x > 0$, 1, -1)’ where the expression evaluates to the second argument if the first argument is true, and to the third argument otherwise. For a full overview of possible expressions accepted see the documentation of the muparser library at <http://muparser.beltoforion.de/>.

If the function you are describing represents a vector-valued function with multiple components, then separate the expressions for individual components by a semicolon.

Possible values: Any string

- *Parameter name:* Variable names

Value: x,y,t

Default: x,y,t

Description: The name of the variables as they will be used in the function, separated by commas. By default, the names of variables at which the function will be evaluated is ‘x’ (in 1d), ‘x,y’ (in 2d) or ‘x,y,z’ (in 3d) for spatial coordinates and ‘t’ for time. You can then use these variable names in your function expression and they will be replaced by the values of these variables at which the function is currently evaluated. However, you can also choose a different set of names for the independent variables at which to evaluate your function expression. For example, if you work in spherical coordinates, you may wish to set this input parameter to ‘r,phi,theta,t’ and then use these variable names in your function expression.

Possible values: Any string

A.152 Parameters in section Prescribed Stokes solution/Velocity function

- *Parameter name:* Function constants

Value:

Default:

Description: Sometimes it is convenient to use symbolic constants in the expression that describes the function, rather than having to use its numeric value everywhere the constant appears. These values can be defined using this parameter, in the form ‘var1=value1, var2=value2, ...’.

A typical example would be to set this runtime parameter to ‘pi=3.1415926536’ and then use ‘pi’ in the expression of the actual formula. (That said, for convenience this class actually defines both ‘pi’ and ‘Pi’ by default, but you get the idea.)

Possible values: Any string

- *Parameter name:* Function expression

Value: 0; 0

Default: 0; 0

Description: The formula that denotes the function you want to evaluate for particular values of the independent variables. This expression may contain any of the usual operations such as addition or multiplication, as well as all of the common functions such as ‘sin’ or ‘cos’. In addition, it may contain expressions like ‘if($x > 0$, 1, -1)’ where the expression evaluates to the second argument if the first

argument is true, and to the third argument otherwise. For a full overview of possible expressions accepted see the documentation of the muparser library at <http://muparser.beltoforion.de/>.

If the function you are describing represents a vector-valued function with multiple components, then separate the expressions for individual components by a semicolon.

Possible values: Any string

- *Parameter name:* Variable names

Value: x,y,t

Default: x,y,t

Description: The name of the variables as they will be used in the function, separated by commas. By default, the names of variables at which the function will be evaluated is ‘x’ (in 1d), ‘x,y’ (in 2d) or ‘x,y,z’ (in 3d) for spatial coordinates and ‘t’ for time. You can then use these variable names in your function expression and they will be replaced by the values of these variables at which the function is currently evaluated. However, you can also choose a different set of names for the independent variables at which to evaluate your function expression. For example, if you work in spherical coordinates, you may wish to set this input parameter to ‘r,phi,theta,t’ and then use these variable names in your function expression.

Possible values: Any string

A.153 Parameters in section Solver parameters

- *Parameter name:* Composition solver tolerance

Value: 1e-12

Default: 1e-12

Description: The relative tolerance up to which the linear system for the composition system gets solved. See ‘Stokes solver parameters/Linear solver tolerance’ for more details.

Possible values: A floating point number v such that $0 \leq v \leq 1$

- *Parameter name:* Temperature solver tolerance

Value: 1e-12

Default: 1e-12

Description: The relative tolerance up to which the linear system for the temperature system gets solved. See ‘Stokes solver parameters/Linear solver tolerance’ for more details.

Possible values: A floating point number v such that $0 \leq v \leq 1$

A.154 Parameters in section Solver parameters/AMG parameters

- *Parameter name:* AMG aggregation threshold

Value: 0.001

Default: 0.001

Description: This threshold tells the AMG setup how the coarsening should be performed. In the AMG used by ML, all points that strongly couple with the tentative coarse-level point form one aggregate. The term strong coupling is controlled by the variable aggregation_threshold, meaning that all elements that are not smaller than aggregation_threshold times the diagonal element do couple strongly. The default is strongly recommended. There are indications that for the Newton solver a different value might be better. For extensive benchmarking of various settings of the AMG parameters in this section for the Stokes problem and others, see <https://github.com/geodynamics/aspect/pull/234>.

Possible values: A floating point number v such that $0 \leq v \leq 1$

- *Parameter name:* AMG output details

Value: false

Default: false

Description: Turns on extra information on the AMG solver. Note that this will generate much more output.

Possible values: A boolean value (true or false)

- *Parameter name:* AMG smoother sweeps

Value: 2

Default: 2

Description: Determines how many sweeps of the smoother should be performed. When the flag elliptic is set to true, (which is true for ASPECT), the polynomial degree of the Chebyshev smoother is set to this value. The term sweeps refers to the number of matrix-vector products performed in the Chebyshev case. In the non-elliptic case, this parameter sets the number of SSOR relaxation sweeps for post-smoothing to be performed. The default is strongly recommended. There are indications that for the Newton solver a different value might be better. For extensive benchmarking of various settings of the AMG parameters in this section for the Stokes problem and others, see <https://github.com/geodynamics/aspect/pull/234>.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* AMG smoother type

Value: Chebyshev

Default: Chebyshev

Description: This parameter sets the type of smoother for the AMG. The default is strongly recommended for any normal runs with ASPECT. There are some indications that the symmetric Gauss-Seidel might be better and more stable for the Newton solver. For extensive benchmarking of various settings of the AMG parameters in this section for the Stokes problem and others, see <https://github.com/geodynamics/aspect/pull/234>

Possible values: Any one of Chebyshev, symmetric Gauss-Seidel

A.155 Parameters in section Solver parameters/Newton solver parameters

- *Parameter name:* Max Newton line search iterations

Value: 5

Default: 5

Description: The maximum number of line search iterations allowed. If the criterion is not reached after this number of iteration, we apply the scaled increment even though it does not satisfy the necessary criteria and simply continue with the next Newton iteration.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* Max pre-Newton nonlinear iterations

Value: 10

Default: 10

Description: If the 'Nonlinear Newton solver switch tolerance' is reached before the maximal number of Picard iteration, then the solver switches to Newton solves anyway.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* Maximum linear Stokes solver tolerance

Value: 0.9

Default: 0.9

Description: The linear Stokes solver tolerance is dynamically chosen for the Newton solver, based on the Eisenstat walker 1994 paper (<https://doi.org/10.1137/0917003>), equation 2.2. Because this value can become larger than one, we limit this value by this parameter.

Possible values: A floating point number v such that $0 \leq v \leq 1$

- *Parameter name:* Nonlinear Newton solver switch tolerance

Value: 1e-5

Default: 1e-5

Description: A relative tolerance with respect to the residual of the first iteration, up to which the nonlinear Picard solver will iterate, before changing to the Newton solver.

Possible values: A floating point number v such that $0 \leq v \leq 1$

- *Parameter name:* SPD safety factor

Value: 0.9

Default: 0.9

Description: When stabilizing the Newton matrix, we can encounter situations where the coefficient inside the elliptic (top-left) block becomes negative or zero. This coefficient has the form $1 + x$ where x can sometimes be smaller than -1 . In this case, the top-left block of the matrix is no longer positive definite, and both preconditioners and iterative solvers may fail. To prevent this, the stabilization computes an α so that $1 + \alpha x$ is never negative. This α is chosen as 1 if $x \geq -1$, and $\alpha = -\frac{1}{x}$ otherwise. (Note that this always leads to $0 \leq \alpha \leq 1$.) On the other hand, we also want to stay away from $1 + \alpha x = 0$, and so modify the choice of α to be 1 if $x \geq -c$, and $\alpha = -\frac{c}{x}$ with a c between zero and one. This way, if $c < 1$, we are assured that $1 - \alpha x > c$, i.e., bounded away from zero.

Possible values: A floating point number v such that $0 \leq v \leq 1$

- *Parameter name:* Stabilization preconditioner

Value: SPD

Default: SPD

Description: This parameters allows for the stabilization of the preconditioner. If one derives the Newton method without any modifications, the matrix created for the preconditioning is not necessarily Symmetric Positive Definite. This is problematic (see Fraters et al., in prep). When ‘none’ is chosen, the preconditioner is not stabilized. The ‘symmetric’ parameters symmetrizes the matrix, and ‘PD’ makes the matrix Positive Definite. ‘SPD’ is the full stabilization, where the matrix is guaranteed Symmetric Positive Definite.

Possible values: Any one of SPD, PD, symmetric, none

- *Parameter name:* Stabilization velocity block

Value: SPD

Default: SPD

Description: This parameters allows for the stabilization of the velocity block. If one derives the Newton method without any modifications, the matrix created for the velocity block is not necessarily Symmetric Positive Definite. This is problematic (see Fraters et al., in prep). When ‘none’ is chosen, the velocity block is not stabilized. The ‘symmetric’ parameters symmetrizes the matrix, and ‘PD’

makes the matrix Positive Definite. ‘SPD’ is the full stabilization, where the matrix is guaranteed Symmetric Positive Definite.

Possible values: Any one of SPD, PD, symmetric, none

- *Parameter name:* Use Eisenstat Walker method for Picard iterations

Value: false

Default: false

Description: If set to true, the Picard iteration uses the Eisenstat Walker method to determine how accurately linear systems need to be solved. The Picard iteration is used, for example, in the first few iterations of the Newton method before the matrix is built including derivatives of the model, since the Picard iteration generally converges even from points where Newton’s method does not.

Once derivatives are used in a Newton method, ASPECT always uses the Eisenstat Walker method.

Possible values: A boolean value (true or false)

- *Parameter name:* Use Newton failsafe

Value: false

Default: false

Description: When this parameter is true and the linear solver fails, we try again, but now with SPD stabilization for both the preconditioner and the velocity block. The SPD stabilization will remain active until the next timestep, when the default values are restored.

Possible values: A boolean value (true or false)

- *Parameter name:* Use Newton residual scaling method

Value: false

Default: false

Description: This method allows to slowly introduce the derivatives based on the improvement of the residual. If set to false, the scaling factor for the Newton derivatives is set to one immediately when switching on the Newton solver. When this is set to true, the derivatives are slowly introduced by the following equation: $\max(0.0, (1.0 - (\text{residual}/\text{switch_initial_residual})))$, where `switch_initial_residual` is the residual at the time when the Newton solver is switched on.

Possible values: A boolean value (true or false)

A.156 Parameters in section Solver parameters/Operator splitting parameters

- *Parameter name:* Reaction time step

Value: 1000.0

Default: 1000.0

Description: Set a time step size for computing reactions of compositional fields and the temperature field in case operator splitting is used. This is only used when the nonlinear solver scheme “operator splitting” is selected. The reaction time step must be greater than 0. If you want to prescribe the reaction time step only as a relative value compared to the advection time step as opposed to as an absolute value, you should use the parameter “Reaction time steps per advection step” and set this parameter to the same (or larger) value as the “Maximum time step” (which is 5.69e+300 by default). Units: Years or seconds, depending on the “Use years in output instead of seconds” parameter.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Reaction time steps per advection step

Value: 0

Default: 0

Description: The number of reaction time steps done within one advection time step in case operator splitting is used. This is only used if the nonlinear solver scheme “operator splitting” is selected. If set to zero, this parameter is ignored. Otherwise, the reaction time step size is chosen according to this criterion and the “Reaction time step”, whichever yields the smaller time step. Units: none.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

A.157 Parameters in section Solver parameters/Stokes solver parameters

- *Parameter name:* Linear solver A block tolerance

Value: 1e-2

Default: 1e-2

Description: A relative tolerance up to which the approximate inverse of the A block of the Stokes system is computed. This approximate A is used in the preconditioning used in the GMRES solver. The exact definition of this block preconditioner for the Stokes equation can be found in [KHB12].

Possible values: A floating point number v such that $0 \leq v \leq 1$

- *Parameter name:* Linear solver S block tolerance

Value: 1e-6

Default: 1e-6

Description: A relative tolerance up to which the approximate inverse of the S block (i.e., the Schur complement matrix $S = BA^{-1}B^T$) of the Stokes system is computed. This approximate inverse of the S block is used in the preconditioning used in the GMRES solver. The exact definition of this block preconditioner for the Stokes equation can be found in [KHB12].

Possible values: A floating point number v such that $0 \leq v \leq 1$

- *Parameter name:* Linear solver tolerance

Value: 1e-7

Default: 1e-7

Description: A relative tolerance up to which the linear Stokes systems in each time or nonlinear step should be solved. The absolute tolerance will then be $\|Mx_0 - F\| \cdot \text{tol}$, where $x_0 = (0, p_0)$ is the initial guess of the pressure, M is the system matrix, F is the right-hand side, and tol is the parameter specified here. We include the initial guess of the pressure to remove the dependency of the tolerance on the static pressure. A given tolerance value of 1 would mean that a zero solution vector is an acceptable solution since in that case the norm of the residual of the linear system equals the norm of the right hand side. A given tolerance of 0 would mean that the linear system has to be solved exactly, since this is the only way to obtain a zero residual.

In practice, you should choose the value of this parameter to be so that if you make it smaller the results of your simulation do not change any more (qualitatively) whereas if you make it larger, they do. For most cases, the default value should be sufficient. In fact, a tolerance of 1e-4 might be accurate enough.

Possible values: A floating point number v such that $0 \leq v \leq 1$

- *Parameter name:* Maximum number of expensive Stokes solver steps

Value: 1000

Default: 1000

Description: This sets the maximum number of iterations used in the expensive Stokes solver. If this value is set too low for the size of the problem, the Stokes solver will not converge and return an error message pointing out that the user didn't allow a sufficiently large number of iterations for the iterative solver to converge.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* Number of cheap Stokes solver steps

Value: 200

Default: 200

Description: As explained in the paper that describes ASPECT (Kronbichler, Heister, and Bangerth, 2012, see [KHB12]) we first try to solve the Stokes system in every time step using a GMRES iteration with a poor but cheap preconditioner. By default, we try whether we can converge the GMRES solver in 200 such iterations before deciding that we need a better preconditioner. This is sufficient for simple problems with variable viscosity and we never need the second phase with the more expensive preconditioner. On the other hand, for more complex problems, and in particular for problems with strongly nonlinear viscosity, the 200 cheap iterations don't actually do very much good and one might skip this part right away. In that case, this parameter can be set to zero, i.e., we immediately start with the better but more expensive preconditioner.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* Use direct solver for Stokes system

Value: true

Default: false

Description: If set to true the linear system for the Stokes equation will be solved using Trilinos klu, otherwise an iterative Schur complement solver is used. The direct solver is only efficient for small problems.

Possible values: A boolean value (true or false)

A.158 Parameters in section Termination criteria

- *Parameter name:* Checkpoint on termination

Value: false

Default: false

Description: Whether to checkpoint the simulation right before termination.

Possible values: A boolean value (true or false)

- *Parameter name:* End step

Value: 100

Default: 100

Description: Terminate the simulation once the specified timestep has been reached.

Possible values: An integer n such that $0 \leq n \leq 2147483647$

- *Parameter name:* Termination criteria

Value: end time

Default: end time

Description: A comma separated list of termination criteria that will determine when the simulation should end. Whether explicitly stated or not, the “end time” termination criterion will always be used. The following termination criteria are available:

‘end step’: Terminate the simulation once the specified timestep has been reached.

‘end time’: Terminate the simulation once the end time specified in the input file has been reached. Unlike all other termination criteria, this criterion is *always* active, whether it has been explicitly selected or not in the input file (this is done to preserve historical behavior of ASPECT, but it also likely does not inconvenience anyone since it is what would be selected in most cases anyway).

‘steady state velocity’: A criterion that terminates the simulation when the RMS of the velocity field stays within a certain range for a specified period of time.

‘user request’: Terminate the simulation gracefully when a file with a specified name appears in the output directory. This allows the user to gracefully exit the simulation at any time by simply creating such a file using, for example, `touch output/terminate`. The file’s location is chosen to be in the output directory, rather than in a generic location such as the Aspect directory, so that one can run multiple simulations at the same time (which presumably write to different output directories) and can selectively terminate a particular one.

Possible values: A comma-separated list of any of end step, end time, steady state velocity, user request

A.159 Parameters in section Termination criteria/Steady state velocity

- *Parameter name:* Maximum relative deviation

Value: 0.05

Default: 0.05

Description: The maximum relative deviation of the RMS in recent simulation time for the system to be considered in steady state. If the actual deviation is smaller than this number, then the simulation will be terminated.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

- *Parameter name:* Time in steady state

Value: 1e7

Default: 1e7

Description: The minimum length of simulation time that the system should be in steady state before termination. Units: years if the ‘Use years in output instead of seconds’ parameter is set; seconds otherwise.

Possible values: A floating point number v such that $0 \leq v \leq \text{MAX_DOUBLE}$

A.160 Parameters in section Termination criteria/User request

- *Parameter name:* File name

Value: terminate-aspect

Default: terminate-aspect

Description: The name of a file that, if it exists in the output directory (whose name is also specified in the input file) will lead to termination of the simulation. The file’s location is chosen to be in the

output directory, rather than in a generic location such as the Aspect directory, so that one can run multiple simulations at the same time (which presumably write to different output directories) and can selectively terminate a particular one.

Possible values: an input filename

References

- [AHFT13] V. Allken, R.S. Huismans, H. Fossen, and C. Thieulot. 3D numerical modelling of graben interaction and linkage: a case study of the Canyonlands grabens, Utah. *Basin Research*, 25:1–14, 2013.
- [AHT11] V. Allken, R. Huismans, and C. Thieulot. Three dimensional numerical modelling of upper crustal extensional systems. *J. Geophys. Res.*, 116:B10409, 2011.
- [AHT12] V. Allken, R. Huismans, and C. Thieulot. Factors controlling the mode of rift interaction in brittle-ductile coupled systems: a 3d numerical study. *Geochem. Geophys. Geosyst.*, 13(5):Q05010, 2012.
- [BBC⁺89] B. Blankenbach, F. Busse, U. Christensen, L. Cserepes, D. Gunkel, U. Hansen, H. Harder, G. Jarvis, M. Koch, G. Marquart, D. Moore, P. Olson, H. Schmeling, and T. Schnaubelt. A benchmark comparison for mantle convection codes. *Geophys. J. Int.*, 98:23–38, 1989.
- [BHK07] W. Bangerth, R. Hartmann, and G. Kanschat. deal.II – a general purpose object oriented finite element library. *ACM Trans. Math. Softw.*, 33(4):24, 2007.
- [BHK12] W. Bangerth, T. Heister, and G. Kanschat. *deal.II Differential Equations Analysis Library, Technical Reference*, 2012. <http://www.dealii.org/>.
- [BHPeGeS14] S. Brune, C. Heine, M Pérez Gussinyé, and S. V. Sobolev. Rift migration explains continental margin asymmetry and crustal hyperextension. *Nat. Comm.*, 5(4014), 2014.
- [BKEO03] Thorsten W. Becker, James B. Kellogg, Göran Ekström, and Richard J. O’Connell. Comparison of azimuthal seismic anisotropy from surface waves and finite strain from global mantle-circulation models. *Geophysical Journal International*, 155(2):696–714, nov 2003.
- [BR86] V. Barcilon and F. M. Richter. Nonlinear waves in compacting media. *Journal of Fluid mechanics*, 164:429–448, 1986.
- [BRV⁺04] J. Badro, J.-P. Rueff, G. Vankó, G. Monaco, G. Fiquet, and F. Guyot. Electronic transitions in perovskite: Possible nonconvecting layers in the lower mantle. *Science*, 305:383–386, 2004.
- [BSA⁺13] C. Burstedde, G. Stadler, L. Alisic, L. C. Wilcox, E. Tan, M. Gurnis, and O. Ghattas. Large-scale adaptive mantle convection simulation. *Geophysical Journal International*, 192.3:889–906, 2013.
- [Bui12] S. J. H. Buiter. A review of brittle compressional wedge models. *Tectonophysics*, 530:1–17, 2012.
- [BWG11] C. Burstedde, L. C. Wilcox, and O. Ghattas. **p4est**: Scalable algorithms for parallel adaptive mesh refinement on forests of octrees. *SIAM J. Sci. Comput.*, 33(3):1103–1133, 2011.
- [Cha86] D. S. Chapman. Thermal gradients in the continental crust. *Geol. Soc. London Spec. Publ.*, 24:63–70, 1986.
- [CSG⁺12] F. Crameri, H. Schmeling, G. J. Golabek, T. Duretz, R. Orendt, S. J. H. Buiter, D. A. May, B. J. P. Kaus, T. V. Gerya, and P. J. Tackley. A comparison of numerical surface topography calculations in geodynamic modelling: An evaluation of the ‘sticky air’ method. *Geophysical Journal International*, 189(1):38–54, 2012.
- [DAC13] R. Deguen, T. Alboussière, and P. Cardin. Thermal convection in Earth’s inner core with phase change at its boundary. *Geophysical Journal International*, 194(3):1310–1334, 2013.

- [DB14] C. R. Dohrmann and P. B. Bochev. A stabilized finite element method for the stokes problem based on polynomial pressure projections. *International Journal for Numerical Methods in Fluids*, 46:183–201, 2004.
- [DH16] J. Dannberg and T. Heister. Compressible magma/mantle dynamics: 3D, adaptive simulations in ASPECT. *Geophysical Journal International*, 207(3):1343–1366, 2016.
- [DHPRF04] J. Donea, A. Huerta, J.-Ph. Ponthot, and A. Rodríguez-Ferran. *Arbitrary Lagrangian-Eulerian Methods*. John Wiley & Sons, Ltd, 2004.
- [DK08] Y. Deubelbeiss and B. J. P. Kaus. Comparison of eulerian and lagrangian numerical techniques for the stokes equations in the presence of strongly varying viscosity. *Physics of the Earth and Planetary Interiors*, 171:92–111, 2008.
- [DMGT11] T. Duretz, D. A. May, T. V. Garya, and P. J. Tackley. Discretization errors and free surface stabilization in the finite difference and marker-in-cell method for applied geodynamics: A numerical study. *Geoch. Geoph. Geosystems*, 12:Q07004/1–26, 2011.
- [DT98] FA Dahlen and Jeroen Tromp. *Theoretical global seismology*. Princeton University Press, 1998.
- [GHPB17] R. Gassmöller, E. Heien, E. G. Puckett, and W. Bangerth. Flexible and scalable particle-in-cell methods for massively parallel computations. *submitted*, 2017.
- [GPP11] J.-L. Guermond, R. Pasquetti, and B. Popov. Entropy viscosity method for nonlinear conservation laws. *J. Comput. Phys.*, 230:4248–4267, 2011.
- [GTZ⁺12] M. Gurnis, M. Turner, S. Zahirovic, L. DiCaprio, S. Spasojevic, R. D. Müller, J. Boyden, M. Seton, V. C. Manea, and D. J. Bower. Plate tectonic reconstructions with continuously closing plates. *Computers & Geosciences*, 38:35–42, 2012.
- [H⁺11] M. A. Heroux et al. Trilinos web page, 2011. <http://trilinos.sandia.gov>.
- [HB11] R. Huismans and C. Beaumont. Depth-dependent extension, two-stage breakup and cratonic underplating at rifted margins. *Nature*, 473(71–75), 2011.
- [HBH⁺05] M. A. Heroux, R. A. Bartlett, V. E. Howle, R. J. Hoekstra, J. J. Hu, T. G. Kolda, R. B. Lehoucq, K. R. Long, R. P. Pawlowski, E. T. Phipps, A. G. Salinger, H. K. Thornquist, R. S. Tuminaro, J. M. Willenbring, A. Williams, and K. S. Stanley. An overview of the Trilinos project. *ACM Trans. Math. Softw.*, 31:397–423, 2005.
- [HK04] G. Hirth and D Kohlstedt. Rheology of the upper mantle and the mantle wedge:a view from the experimentalists. In J. M. Eiler, editor, *Inside the Subduction Factory*, Geophys. Monogr. Ser. 138, pages 83–105. American Geophysical Union, Washington, DC, 2004.
- [Jea11] Jean-Luc Guermond and Richard Pasquetti and Bojan Popov. Entropy viscosity method for nonlinear conservation laws. *Journal of Computational Physics*, 230:4248–4267, 2011.
- [Kau10] B.J.P. Kaus. Factors that control the angle of shear bands in geodynamic numerical models of brittle deformation. *Tectonophysics*, 484:36–47, 2010.
- [KHB12] M. Kronbichler, T. Heister, and W. Bangerth. High accuracy mantle convection simulation through modern numerical methods. *Geophysical Journal International*, 191:12–29, 2012.
- [KLVK⁺10] Scott D King, Changyeol Lee, Peter E Van Keken, Wei Leng, Shijie Zhong, Eh Tan, Nicola Tosi, and Masanori C Kameyama. A community benchmark for 2-D Cartesian compressible convection in the Earth’s mantle. *Geophysical Journal International*, 180(1):73–87, 2010.

- [KMK13] Tobias Keller, Dave A. May, and Boris J. P. Kaus. Numerical modelling of magma dynamics coupled to tectonic deformation of lithosphere and crust. *Geophysical Journal International*, 195(3):1406–1442, 2013.
- [KMM10] B. J. P. Kaus, H. Mühlhaus, and D. A. May. A stabilization algorithm for geodynamic numerical simulations with a free surface. *Physics of the Earth and Planetary Interiors*, 181(1):12–20, 2010.
- [KSL03] Richard F. Katz, Marc Spiegelman, and Charles H. Langmuir. A new parameterization of hydrous mantle melting. *Geochemistry, Geophysics, Geosystems*, 4(9):n/a–n/a, 2003.
- [LT07] B. Lund and J. Townend. Calculating horizontal stress orientations with full or partial knowledge of the tectonic stress tensor. *Geophys. J. Int.*, 170:1328–1335, 2007.
- [MD04] C. Morency and M.-P. Doin. Numerical simulations of the mantle lithosphere delamination. *Journal of Geophysical Research: Solid Earth (1978–2012)*, 109(B3), 2004.
- [MJ83] Dan McKenzie and James Jackson. The relationship between strain rates, crustal thickening, palaeomagnetism, finite strain and fault movements within a deforming zone. *Earth and Planetary Science Letters*, 65(1):182–202, 1983.
- [MQL⁺07] L. Moresi, S. Quenette, V. Lemiale, C. Meriaux, B. Appelbe, and H. B. Mühlhaus. Computational approaches to studying non-linear dynamics of the crust and mantle. *Phys. Earth Planet. Interiors*, 163:69–82, 2007.
- [NB15] J. Naliboff and S. Buiter. Rift reactivation and migration during multiphase extension. *Earth Planet. Sci. Lett.*, 421(58–67), 2015.
- [RB04] E. H. Rutter and K. H. Brodie. Experimental grain size-sensitive flow of hot-pressed brazilian quartz aggregates. *J. Struct. Geol.*, 26:2011–2023, 2004.
- [RDvHW11] J. Ritsema, A. Deuss, H. J. van Heijst, and J. H. Woodhouse. S40rts: a degree-40 shear-velocity model for the mantle from new rayleigh wave dispersion, teleseismic traveltimes and normal-mode splitting function measurements. *Geophysical Journal International*, 184:1223–1236, 2011.
- [RGWD06] E. Rybacki, M. Gottschalk, R. Wirth, and G. Dresen. Influence of water fugacity and activation volume on the flow properties of fine-grained anorthite aggregates. *J. Geophys. Res.*, 111(B3), 2006.
- [Rib92] Neil M. Ribe. On the relation between seismic anisotropy and finite strain. *Journal of Geophysical Research*, 97(B6):8737, 1992.
- [RvH00] J. Ritsema and H. J. van Heijst. Seismic imaging of structural heterogeneity in earth’s mantle: Evidence for large-scale mantle flow. *Sci. Progr.*, 83:243–259, 2000.
- [SBE⁺08] H. Schmeling, A. Y. Babeyko, A. Enns, C. Faccenna, F. Funiciello, T. Gerya, G. J. Golabek, S. Grigull, B. J. P. Kaus, G. Morra, S. M. Schmalholz, and J. van Hunen. A benchmark comparison of spontaneous subduction models—towards a free surface. *Physics of the Earth and Planetary Interiors*, 171:198–223, 2008.
- [Sch00] Harro Schmeling. Partial melting and melt segregation in a convecting mantle. In *Physics and Chemistry of Partially Molten Rocks*, pages 141–178. Springer, 2000.
- [Sch06] H Schmeling. A model of episodic melt extraction for plumes. *Journal of Geophysical Research: Solid Earth*, 111(B3), 2006.

- [SP03] D. W. Schmid and Y. Y. Podladchikov. Analytical solutions for deformable elliptical inclusions in general shear. *Geophysical Journal International*, 155(1):269–288, 2003.
- [SS11] Gideon Simpson and Marc Spiegelman. Solitary wave benchmarks in magma dynamics. *Journal of Scientific Computing*, 49(3):268–290, 2011.
- [STO01] G. Schubert, D. L. Turcotte, and P. Olson. *Mantle Convection in the Earth and Planets, Part 1*. Cambridge, 2001.
- [Thi15] C. Thieulet. ELEFANT: a user-friendly multipurpose geodynamics code. Technical report, Utrecht University, 2015.
- [TMK14] M. Thielmann, D. A. May, and B. J. P. Kaus. Discretization errors in the hybrid finite element particle-in-cell method. *Pure and Applied Geophysics*, 171:2165–2184, 2014.
- [TS14] D. L. Turcotte and G. Schubert. *Geodynamics*. Cambridge, 3rd edition, 2014.
- [vKKS⁺97] P. E. van Keken, S. D. King, H. Schmeling, U. R. Christensen, D. Neumeister, and M.-P. Doin. A comparison of methods for the modeling of thermochemical convection. *J. Geophys. Res.*, 102:22477–22495, 1997.
- [Wil99] S. D. Willett. Rheological dependence of extension in wedge models of convergent orogens. *Tectonophysics*, 305:419–435, 1999.
- [Wor12] J. Worthen. *Inverse Problems in Mantle Convection: Models, Algorithms, and Applications*. PhD thesis, University of Texas at Austin, in preparation, 2012.
- [ZGH93] Shijie Zhong, Michael Gurnis, and Gregory Hulbert. Accurate determination of surface normal stress in viscous flow from a consistent boundary flux method. *Physics of the Earth and Planetary Interiors*, 78(1-2):1–8, jun 1993.
- [Zho96] S. Zhong. Analytic solution for Stokes' flow with lateral variations in viscosity. *Geophys. J. Int.*, 124:18–28, 1996.

Index of run-time parameter entries

The following is a listing of all run-time parameters that can be set in the input parameter file. They are all described in Section A and the listed page numbers are where their detailed documentation can be found. A listing of all parameters sorted by the section name in which they are declared is given in the index on page 470 below.

- A1, 374, 382, 442
- A2, 374, 383, 442
- A3, 374, 383, 442
- Activation energies, 143
 - Activation energies for diffusion creep, 354, 397
 - Activation energies for dislocation creep, 354, 397
- Activation volume, 143
 - Activation volumes for diffusion creep, 354, 397
 - Activation volumes for dislocation creep, 355, 397
- Adapt by fraction of cells, 404
- Additional refinement times, 71, 75, 404
- Additional shared libraries, 95, 107, 111, 185, 188, 190, 212, 222, 231, 256
- Additional tangential mesh velocity boundary indicators, 300
- Adiabatic surface temperature, 19, 256
- Adiabatic temperature gradient, 329
- Advect logarithm of grain size, 362
- Age bottom boundary layer, 327
- Age top boundary layer, 327
- Alpha, 278
- alpha, 298
 - Also output the spherical harmonic coefficients of CMB dynamic topography contribution, 421
 - Also output the spherical harmonic coefficients of density anomaly contribution, 421
 - Also output the spherical harmonic coefficients of geoid anomaly, 421
 - Also output the spherical harmonic coefficients of surface dynamic topography contribution, 421
- Ambient temperature, 332
- AMG aggregation threshold, 451
- AMG output details, 452
- AMG smoother sweeps, 452
- AMG smoother type, 452
- Amplitude, 327, 337
- Angle, 337
 - Angle of internal friction, 359
 - Angles of internal friction, 398
- Angular mode, 209, 338
- Average melt velocity, 403
- Average specific grain boundary energy, 362
 - Average velocity scheme, 445
 - Averaging operation, 347
- B1, 374, 383, 442
- B2, 375, 383, 442
- B3, 375, 383, 442
- Background Viscosities, 361
- Base model, 348, 352
- Bell shape limit, 348
- beta, 299, 378, 388, 444
- Beta composition, 279
- Bilinear interpolation, 362, 395
- Bottom composition, 268, 269
- Bottom temperature, 64, 74, 81, 277
- Boundary indicator to temperature mappings, 278
- Boundary refinement indicators, 410
- Box origin X coordinate, 303, 304
- Box origin Y coordinate, 303, 305
- Box origin Z coordinate, 303, 305
- C1, 375, 383, 443
- C2, 375, 383, 443
- C3, 375, 384, 443
- Cells along circumference, 312
- Center X, 332
- Center x, 431
- Center Y, 332
- Center y, 431
- Center Z, 333
- Center z, 431
- CFL number, 56, 94, 211, 256
- Checkpoint on termination, 456
- Chunk inner radius, 307
- Chunk maximum latitude, 307
- Chunk maximum longitude, 307
- Chunk minimum latitude, 307
- Chunk minimum longitude, 308
- Chunk outer radius, 308
- CMB pressure, 279
- Coarsening fraction, 71, 404
- Coefficient of yield stress increase with depth, 143
- Coefficients of dynamic friction, 361
- Coefficients of static friction, 361

Cohesion, 359
 Cohesion strain weakening factors, 398
 Cohesions, 361, 398
 Cohesive strength of rocks at the surface, 143
 Command, 418
 Composition dependency, 281
 Composition polynomial degree, 54, 295
 Composition reference profile, 262
 Composition solver tolerance, 54, 451
 Composition viscosity prefactor, 98, 207, 370, 375, 392
 Composition viscosity prefactor 1, 350
 Composition viscosity prefactor 2, 350
 Compositional field methods, 92, 294
 Compositional field scaling factors, 410, 411
 Compositional field thresholds, 412
 Compositional heating values, 317
 Compressibility, 207, 348, 370, 375
 Coordinate system, 270, 283, 287, 291, 322, 330, 412, 413
 Core conductivity, 279
 Core density, 279
 Core heat capacity, 279
 Corresponding phase for density jump, 206, 370
 cR, 299
 Crust composition number, 319
 Crust defined by composition, 319
 Crust depth, 319
 D1, 376, 443
 D2, 376, 443
 D3, 376, 443
 Data directory, 130, 135, 261, 267, 275, 285, 290, 292, 310, 313, 321, 329, 330, 333, 335, 347, 352, 362, 395, 428, 446
 Data file name, 261, 267, 276, 285, 290, 311, 313, 322, 330, 347, 428, 446
 Data file time step, 267, 276, 286, 290, 292
 Data output format, 89, 92, 422
 Decreasing file order, 267, 276, 286, 290, 292
 Define transition by depth instead of pressure, 371
 Delta, 279
 Densities, 143, 355, 359, 389, 398
 Density, 348
 Density above, 420, 421
 Density below, 420, 422
 Density differential for compositional field 1, 86, 87, 98, 178, 202, 350, 371, 376, 392
 Density differential for compositional field 2, 87, 351
 Density formulation, 272
 Depletion density change, 379, 384
 Depletion solidus change, 379, 384
 Depth, 308
 Depth dependence method, 353
 Depth list, 353
 Depth subdivisions, 308
 Derivatives file names, 363
 Di, 390
 Diffusion activation energy, 363
 Diffusion activation volume, 363
 Diffusion creep exponent, 363
 Diffusion creep grain size exponent, 363
 Diffusion creep prefactor, 363
 Dimension, 12, 42, 44, 45, 56, 61, 63, 74, 80, 97, 117, 123, 185, 188, 190, 200, 222, 256
 Discontinuous penalty, 297
 Dislocation activation energy, 364
 Dislocation activation volume, 364
 Dislocation creep exponent, 364
 Dislocation creep prefactor, 364
 Dislocation viscosity iteration number, 364
 Dislocation viscosity iteration threshold, 364
 dR over dt, 281
 dT over dt, 281
 dX over dt, 281
 E1, 376, 444
 E2, 376, 444
 East-West subdivisions, 309
 Eccentricity, 309
 Effective viscosity coefficient, 355
 Enable additional Stokes RHS, 299
 End step, 456
 End strain weakening intervals, 398
 End time, 42, 56, 63, 80, 97, 117, 123, 185, 188, 190, 200, 211, 222, 257
 Entropy derivative pressure, 348
 Entropy derivative temperature, 349
 Evaluation points, 433
 Excess entropy only, 420
 Exponential melt weakening factor, 379, 384
 File name, 282, 457
 Filename for initial geotherm table, 337
 Filter output, 434
 First data file model time, 268, 276, 286, 290, 293
 First data file number, 268, 276, 286, 291, 293
 Fixed composition boundary indicators, 264
 Fixed temperature boundary indicators, 61, 64, 74, 81, 94, 118, 125, 135, 210, 241, 272
 Formulation, 65, 300
 Free surface boundary indicators, 94, 301

Free surface stabilization theta, 94, 301
 Freezing rate, 384
 Friction strain weakening factors, 399
 Function constants, 58, 64, 74, 79, 81, 143, 178, 179, 202, 262–264, 270, 283, 287, 291, 314, 318, 322, 328, 330, 353, 412, 413, 427, 428, 447–450
 Function expression, 58, 64, 74, 79, 81, 83, 87, 90, 94, 98, 99, 107–109, 126, 143, 178, 179, 201, 202, 222, 262–264, 271, 283, 287, 291, 314, 318, 323, 328, 331, 354, 412, 413, 427, 429, 447–450
 gamma, 299, 391
 Geometric constant, 365
 Global composition maximum, 297
 Global composition minimum, 297
 Global particle property maximum, 432
 Global particle property minimum, 432
 Global temperature maximum, 297
 Global temperature minimum, 298
 Grain growth activation energy, 365
 Grain growth activation volume, 365
 Grain growth exponent, 365
 Grain growth rate constant, 365
 Grain size, 355, 399
 Grain size exponents for diffusion creep, 355, 399
 Gravity acceleration, 279
 Half decay times, 319
 Half life, 223
 Half life times, 282
 Heat advection by melt, 403
 Heat capacities, 399
 Heat capacity, 143, 355
 Heating rates, 282, 319
 Include melt transport, 403
 Include melting and freezing, 379
 Inclusion gradient, 333
 Inclusion shape, 333
 Inclusion temperature, 333
 Initial adaptive refinement, 42, 56, 65, 71, 75, 82, 99, 119, 126, 187, 189, 192, 202, 404
 Initial concentrations, 283
 Initial concentrations crust, 320
 Initial concentrations mantle, 320
 Initial condition file name, 130, 334, 335
 Initial global refinement, 42, 56, 65, 71, 75, 82, 99, 119, 126, 187, 189, 192, 202, 404
 Initial light composition, 280
 Inner composition, 271
 Inner radius, 118, 124, 209, 312
 Inner temperature, 118, 125, 280, 284
 Integration scheme, 423
 Interpolate output, 434
 Interpolation scheme, 92, 423
 Interpolation width, 135
 Isotherm depth filename, 329
 Isotherm temperature, 329
 Jump height, 96
 K0, 280
 Latent heat, 395
 Lateral viscosity file name, 395
 Lateral wave number one, 331
 Lateral wave number two, 331
 Latitude repetitions, 308
 Left composition, 268, 269
 Left composition lithosphere, 269
 Left temperature, 64, 277
 Left temperature lithosphere, 278
 Lh, 280
 Linear solver A block tolerance, 54, 455
 Linear solver S block tolerance, 54, 455
 Linear solver tolerance, 54, 193, 455
 List of material properties, 441
 List of model names, 64, 74, 81, 98, 118, 122, 123, 125, 186, 189, 192, 201, 223, 265, 273, 316, 320, 323
 List of model operators, 266, 274, 321, 325
 List of normalized fields, 295
 List of output variables, 86, 99, 128, 203, 245, 419, 435
 List of particle properties, 90, 92, 424
 List of postprocessors, 42, 54, 65, 75, 82, 85, 86, 88, 94, 99, 119, 126, 128, 178, 187, 190, 192, 194, 203, 242, 415
 List of properties, 445
 Lithosphere thickness, 293
 Lithospheric thickness, 305
 Load balancing strategy, 425
 Longitude repetitions, 308
 Lower mantle grain size scaling, 365
 Lower viscosity, 96
 Magnitude, 64, 71, 75, 98, 201, 210, 315, 332
 Magnitude at bottom, 315
 Magnitude at surface, 315
 Mapped particle properties, 92, 295
 Mass conservation, 300
 Mass fraction cpx, 377, 385, 444

Material averaging, 55, 98, 99, 338
 Material file format, 366
 Material file names, 366, 395
 Max iteration, 280
 Max Newton line search iterations, 452
 Max nonlinear iterations, 257
 Max nonlinear iterations in pre-refinement, 257
 Max pre-Newton nonlinear iterations, 452
 Maximal composition, 271
 Maximal temperature, 284
 Maximum degree, 422
 Maximum latent heat substeps, 366
 Maximum lateral viscosity variation, 395
 Maximum latitude, 431
 Maximum linear Stokes solver tolerance, 453
 Maximum longitude, 431
 Maximum number of expensive Stokes solver steps, 456
 Maximum order, 334, 335
 Maximum particles per cell, 425
 Maximum pyroxenite melt fraction, 377
 Maximum radius, 431
 Maximum relative deviation, 457
 Maximum relative increase in time step, 257
 Maximum specific heat, 366
 Maximum strain rate ratio iterations, 356
 Maximum temperature dependence of viscosity, 366
 Maximum thermal expansivity, 366
 Maximum thermal prefactor, 392
 Maximum time step, 222, 257
 Maximum viscosity, 356, 359, 361, 367, 371, 396, 399
 Maximum x, 430
 Maximum y, 430
 Maximum z, 430
 Melt bulk modulus derivative, 380, 385
 Melt compressibility, 380, 385
 Melt extraction depth, 385
 Melt scaling factor threshold, 403
 Melting entropy change, 319
 Melting time scale for operator splitting, 380, 385
 Mesh cells per compaction length, 410
 Mesh refinement, 247
 Minimal composition, 271
 Minimal temperature, 284
 Minimum degree, 422
 Minimum grain size, 367
 Minimum latitude, 432
 Minimum longitude, 432
 Minimum particles per cell, 425
 Minimum radius, 432
 Minimum refinement level, 405
 Minimum specific heat, 367
 Minimum strain rate, 143, 356, 399
 Minimum thermal expansivity, 367
 Minimum thermal prefactor, 392
 Minimum viscosity, 356, 359, 361, 367, 371, 396, 400
 Minimum x, 430
 Minimum y, 430
 Minimum z, 431
 Model name, 42, 56, 58, 61, 63–65, 74, 75, 80, 81, 83, 85, 87, 94, 96, 98, 99, 109, 111, 117–119, 124, 126, 130, 142, 143, 178, 179, 186, 187, 189, 191, 192, 194, 201, 202, 206, 209, 210, 212, 222, 223, 232, 234, 235, 239, 240, 260, 266, 274, 301, 310, 313, 321, 325, 338, 446
 Name of compositional field, 441
 Names of fields, 92, 142, 295, 441
 Names of vectors, 441
 NE corner, 309
 Non-dimensional depth, 337
 Nonlinear Newton solver switch tolerance, 453
 Nonlinear solver scheme, 222, 258
 Nonlinear solver tolerance, 258
 Normalize individual refinement criteria, 405
 North-South subdivisions, 309
 Number lateral average bands, 396
 Number of cheap Stokes solver steps, 456
 Number of components, 427
 Number of depth slices, 445
 Number of elements, 320
 Number of fields, 83, 92, 98, 109, 142, 202, 295
 Number of grouped files, 52, 119, 126, 128, 439
 Number of integration points, 288
 Number of particles, 88–90, 92, 425
 Number of particles per cell per direction, 430
 Number of points, 262
 Number of radioactive heating elements, 283
 Number of zones, 419
 NW corner, 309
 Opening angle, 118, 209, 312
 Outer composition, 272
 Outer radius, 118, 124, 209, 312
 Outer temperature, 118, 125, 280, 285
 Output bottom, 420
 Output data in geographical coordinates, 422
 Output directory, 42, 52, 56, 63, 80, 89, 97, 117, 123, 185, 188, 190, 200, 211, 258

Output format, 46, 99, 119, 126, 128, 419, 440
 Output mesh velocity, 440
 Output surface, 421
 Output to file, 433

Particle generator name, 92, 425
 Particle weight, 426
 Peridotite melting entropy change, 377, 386
 Phase transition Clapeyron slopes, 206, 367, 371
 Phase transition density jumps, 206, 372
 Phase transition depths, 206, 368, 372
 Phase transition pressure widths, 372
 Phase transition pressures, 372
 Phase transition temperatures, 206, 368, 372
 Phase transition widths, 206, 368, 373
 Plugin name, 272
 Point one, 135, 136, 293
 Point two, 135, 136, 293
 Position, 327
 Preexponential constant for viscous rheology law, 143
 Prefactors for diffusion creep, 356, 400
 Prefactors for dislocation creep, 356, 400
 Prescribe internal velocities, 107
 Prescribed traction boundary indicators, 285
 Prescribed velocity boundary indicators, 79, 81, 135, 191, 193, 240, 288
 Pressure normalization, 18, 55, 63, 97, 186–188, 190, 258
 Pressure solidus change, 380
 Pyroxenite melting entropy change, 377

r1, 379, 388, 444
 r2, 379, 388, 444
 Ra, 390
 Radial layers, 432
 Radial viscosity file name, 396
 Radiogenic heating rate, 317
 Radius, 311, 328
 Radius repetitions, 308
 Random cell selection, 429
 Random number seed, 429
 Reaction depth, 87, 351
 Reaction terms, 349
 Reaction time step, 222, 454
 Reaction time steps per advection step, 455
 Reciprocal required strain, 368
 Recrystallized grain size, 368
 Reference bulk viscosity, 380, 386
 Reference compressibility, 368, 391
 Reference density, 56, 65, 86, 96, 98, 111, 201, 206, 209, 212, 351, 358, 369, 373, 377, 390, 391, 393, 394
 Reference melt density, 380, 386
 Reference melt viscosity, 381, 386
 Reference permeability, 381, 386
 Reference shear viscosity, 381, 386
 Reference solid density, 381, 386
 Reference specific heat, 65, 96, 206, 209, 212, 351, 358, 369, 373, 377, 381, 387, 390, 391, 393, 394
 Reference strain rate, 143, 359, 362, 400
 Reference temperature, 56, 65, 86, 96, 111, 130, 143, 206, 209, 212, 332, 334, 336, 351, 356, 358, 360, 369, 373, 377, 381, 387, 389, 393, 394, 400
 Reference viscosity, 357, 358, 396, 400
 Refinement criteria merge operation, 405
 Refinement criteria scaling factors, 405
 Refinement fraction, 71, 202, 406
 Relative density of melt, 378
 Remove degree 0 from perturbation, 130, 334, 336
 Remove nullspace, 210, 414
 Remove temperature heterogeneity down to specified depth, 334, 336
 Representative point, 288
 Resume computation, 52, 56, 75, 211, 259
 Rh, 280
 Rho0, 281
 Right composition, 269, 270
 Right composition lithosphere, 270
 Right temperature, 64, 277, 278
 Right temperature lithosphere, 278
 Rotation offset, 209, 338
 Run on all processes, 419
 Run postprocessors on initial refinement, 406
 Run postprocessors on nonlinear iterations, 418
 Scale factor, 261, 268, 276, 286, 291, 293, 311, 314, 322, 330, 347, 447
 SE corner, 309
 Semi-major axis, 310
 Shape radius, 333
 Sigma, 337
 Sign, 337
 Solid compressibility, 381, 387
 SPD safety factor, 453
 Specific heat, 349
 Specific heats, 360, 389
 Specify a lower maximum order, 334, 336
 Spline knots depth file name, 130, 335, 336
 Stabilization preconditioner, 453

- Stabilization velocity block, 453
 Start strain weakening intervals, 400
 Start time, 80, 97, 185, 188, 190, 200, 222, 259
 Steps between checkpoint, 75, 126, 294
 Stokes velocity polynomial degree, 73, 186, 187, 189, 192, 296
 Strain rate residual tolerance, 357
 Strategy, 119, 126, 202, 247, 406
 Stress exponents for diffusion creep, 357, 401
 Stress exponents for dislocation creep, 357, 401
 Stress exponents for plastic rheology, 143
 Stress exponents for viscous rheology, 143
 Stress limiter exponents, 401
 Subadiabaticity, 328
 Surface pressure, 18, 63, 259
 Surface solidus, 382
 Surface temperature, 329
 Surface velocity projection, 301
 SW corner, 310
- Tangential velocity boundary indicators, 61, 64, 74, 79, 81, 94, 118, 124, 135, 186, 189, 201, 210, 289
 Temperature equation, 300
 Temperature polynomial degree, 54, 73, 109, 296
 Temperature scaling factor, 410
 Temperature solver tolerance, 54, 65, 451
 Temporary output location, 440
 Terminate on failure, 419
 Termination criteria, 247, 457
 Thermal bulk viscosity exponent, 382, 387
 Thermal conductivities, 360, 389
 Thermal conductivity, 65, 81, 85, 87, 96, 111, 206, 209, 212, 346, 349, 352, 358, 369, 373, 378, 382, 387, 391, 393, 394, 396
 Thermal diffusivities, 401
 Thermal diffusivity, 143, 357
 Thermal expansion coefficient, 65, 81, 85, 87, 96, 98, 111, 117, 124, 178, 207, 209, 212, 350, 352, 358, 369, 373, 378, 382, 387, 392–394
 Thermal expansion coefficient in initial temperature scaling, 130, 335, 336
 Thermal expansion coefficient of melt, 378
 Thermal expansivities, 143, 357, 360, 389, 401
 Thermal viscosity exponent, 111, 207, 346, 352, 373, 378, 382, 387, 393
 Theta, 281
 Threshold for melt presence, 433
 Time between checkpoint, 294
 Time between data output, 89, 92, 426
 Time between graphical output, 65, 75, 82, 86, 88, 94, 99, 119, 126, 420, 440
 Time between text output, 433
 Time in steady state, 457
 Time step, 135
 Time steps between graphical output, 440
 Time steps between mesh refinement, 65, 71, 75, 82, 119, 126, 409
 Timing output frequency, 259
 Tm0, 281
 Tm1, 282
 Tm2, 282
 Top composition, 269, 270
 Top temperature, 64, 74, 81, 277, 278
 Topography parameters, 311
 Transition depths, 346
- Update ghost particles, 427
 Upper viscosity, 96
 Use artificial viscosity smoothing, 109, 298
 Use BW11, 282
 Use compositional field for heat production averaging, 317
 Use conduction timestep, 260
 Use direct solver for Stokes system, 456
 Use discontinuous compaction pressure, 404
 Use discontinuous composition discretization, 296
 Use discontinuous temperature discretization, 296
 Use Eisenstat Walker method for Picard iterations, 454
 Use enthalpy for material properties, 369
 Use finite strain tensor, 402
 Use fractional melting, 388
 Use full compressibility, 388
 Use lateral average temperature for viscosity, 397
 Use limiter, 433
 Use limiter for discontinuous composition solution, 298
 Use limiter for discontinuous temperature solution, 298
 Use locally conservative discretization, 100, 186, 189, 192, 296
 Use Newton failsafe, 454
 Use Newton residual scaling method, 454
 Use operator splitting, 222, 260
 Use paleowattmeter, 369
 Use simplified adiabatic heating, 11, 317
 Use strain weakening, 402
 Use surface condition function, 262
 Use table properties, 370
 Use TALA, 346, 390

Use years in output instead of seconds, 15, 63, 80, 117, 123, 200, 260

Variable names, 58, 64, 74, 79, 81, 83, 87, 90, 94, 99, 107–109, 143, 178, 179, 202, 222, 263, 264, 271, 284, 287, 292, 314, 318, 323, 329, 331, 354, 412, 413, 428, 429, 447–451

Velocity file name, 135, 294

Vertical wave number, 332

Viscosities, 389

Viscosity, 56, 65, 81, 86, 87, 98, 111, 117, 124, 178, 201, 207, 209, 212, 346, 350, 352, 370, 374, 378, 392–394

Viscosity averaging scheme, 358, 360, 390, 402

Viscosity depth file, 353

Viscosity depth prefactor, 391

Viscosity jump, 186, 191

Viscosity list, 353

Viscosity parameter, 194

Viscosity prefactors, 207, 346, 374

Viscosity temperature prefactor, 391

Viscous flow law, 402

Viscous strain weakening factors, 402

Vs to density scaling, 130, 335, 337

Work fraction for boundary area change, 370

Write in background thread, 440

Write statistics for each nonlinear iteration, 422

X extent, 61, 63, 74, 80, 98, 142, 186, 189, 191, 201, 303, 305

X periodic, 303, 305

X periodic lithosphere, 305

X repetitions, 142, 303, 305

Y extent, 61, 63, 74, 80, 98, 142, 186, 189, 191, 201, 304, 306

Y periodic, 304, 306

Y periodic lithosphere, 306

Y repetitions, 304, 306

Y repetitions lithosphere, 306

Yield mechanism, 402

Z extent, 74, 201, 304, 306

Z periodic, 304, 306

Z repetitions, 304, 307

Z repetitions lithosphere, 307

Zero velocity boundary indicators, 98, 118, 124, 289

Index of run-time parameters with section names

The following is a listing of all run-time parameters, sorted by the section in which they appear. To find entries sorted by their name, rather than their section, see the index on page [463](#) above.

- Additional shared libraries, [95](#), [107](#), [111](#), [185](#), [188](#), [190](#), [212](#), [222](#), [231](#), [256](#)
 - Adiabatic conditions model
 - Ascii data model
 - Data directory, [261](#)
 - Data file name, [261](#)
 - Scale factor, [261](#)
 - Compute profile
 - Composition reference profile, [262](#)
 - Function constants, [262](#)
 - Function expression, [262](#)
 - Number of points, [262](#)
 - Surface condition function/Function constants, [263](#)
 - Surface condition function/Function expression, [263](#)
 - Surface condition function/Variable names, [263](#)
 - Use surface condition function, [262](#)
 - Variable names, [263](#)
 - Function
 - Function constants, [264](#)
 - Function expression, [264](#)
 - Variable names, [264](#)
 - Model name, [260](#)
 - Adiabatic surface temperature, [19](#), [256](#)
 - Boundary composition model
 - Ascii data model
 - Data directory, [267](#)
 - Data file name, [267](#)
 - Data file time step, [267](#)
 - Decreasing file order, [267](#)
 - First data file model time, [268](#)
 - First data file number, [268](#)
 - Scale factor, [268](#)
 - Box
 - Bottom composition, [268](#)
 - Left composition, [268](#)
 - Right composition, [269](#)
 - Top composition, [269](#)
 - Box with lithosphere boundary indicators
 - Bottom composition, [269](#)
 - Left composition, [269](#)
 - Left composition lithosphere, [269](#)
 - Right composition, [270](#)
 - Right composition lithosphere, [270](#)
 - Top composition, [270](#)
- Fixed composition boundary indicators, [264](#)
- Function
 - Coordinate system, [270](#)
 - Function constants, [270](#)
 - Function expression, [271](#)
 - Variable names, [271](#)
- Initial composition
 - Maximal composition, [271](#)
 - Minimal composition, [271](#)
- List of model names, [265](#)
- List of model operators, [266](#)
- Model name, [266](#)
- Spherical constant
 - Inner composition, [271](#)
 - Outer composition, [272](#)
- Boundary fluid pressure model
 - Density
 - Density formulation, [272](#)
 - Plugin name, [272](#)
- Boundary temperature model
 - Ascii data model
 - Data directory, [275](#)
 - Data file name, [276](#)
 - Data file time step, [276](#)
 - Decreasing file order, [276](#)
 - First data file model time, [276](#)
 - First data file number, [276](#)
 - Scale factor, [276](#)
 - Box
 - Bottom temperature, [64](#), [74](#), [81](#), [277](#)
 - Left temperature, [64](#), [277](#)
 - Right temperature, [64](#), [277](#)
 - Top temperature, [64](#), [74](#), [81](#), [277](#)
 - Box with lithosphere boundary indicators
 - Bottom temperature, [277](#)
 - Left temperature, [277](#)
 - Left temperature lithosphere, [278](#)
 - Right temperature, [278](#)
 - Right temperature lithosphere, [278](#)
 - Top temperature, [278](#)
- Constant
 - Boundary indicator to temperature mappings, [278](#)
- Dynamic core

Alpha, 278
 Beta composition, 279
 CMB pressure, 279
 Core conductivity, 279
 Core density, 279
 Core heat capacity, 279
 Delta, 279
 dR over dt , 281
 dT over dt , 281
 dX over dt , 281
 Geotherm parameters/Composition dependency, 281
 Geotherm parameters/Theta, 281
 Geotherm parameters/Tm0, 281
 Geotherm parameters/Tm1, 282
 Geotherm parameters/Tm2, 282
 Geotherm parameters/Use BW11, 282
 Gravity acceleration, 279
 Initial light composition, 280
 Inner temperature, 280
 K0, 280
 Lh, 280
 Max iteration, 280
 Other energy source/File name, 282
 Outer temperature, 280
 Radioactive heat source/Half life times, 282
 Radioactive heat source/Heating rates, 282
 Radioactive heat source/Initial concentrations, 283
 Radioactive heat source/Number of radioactive heating elements, 283
 Rh, 280
 Rho0, 281
 Fixed temperature boundary indicators, 61, 64, 74, 81, 94, 118, 125, 135, 210, 241, 272
 Function
 Coordinate system, 283
 Function constants, 283
 Function expression, 283
 Maximal temperature, 284
 Minimal temperature, 284
 Variable names, 284
 Initial temperature
 Maximal temperature, 284
 Minimal temperature, 284
 List of model names, 64, 74, 81, 98, 118, 125, 186, 189, 192, 201, 273
 List of model operators, 274
 Model name, 274
 Spherical constant
 Inner temperature, 118, 125, 284
 Outer temperature, 118, 125, 285
 Boundary traction model
 Ascii data model
 Data directory, 285
 Data file name, 285
 Data file time step, 286
 Decreasing file order, 286
 First data file model time, 286
 First data file number, 286
 Scale factor, 286
 Function
 Coordinate system, 287
 Function constants, 287
 Function expression, 287
 Variable names, 287
 Initial lithostatic pressure
 Number of integration points, 288
 Representative point, 288
 Prescribed traction boundary indicators, 285
 Boundary velocity model
 Ascii data model
 Data directory, 290
 Data file name, 290
 Data file time step, 290
 Decreasing file order, 290
 First data file model time, 290
 First data file number, 291
 Scale factor, 291
 Function
 Coordinate system, 291
 Function constants, 79, 81, 291
 Function expression, 79, 81, 291
 Variable names, 79, 81, 292
 GPlates model
 Data directory, 135, 292
 Data file time step, 292
 Decreasing file order, 292
 First data file model time, 293
 First data file number, 293
 Interpolation width, 135
 Lithosphere thickness, 293
 Point one, 135, 293
 Point two, 135, 293
 Scale factor, 293
 Time step, 135
 Velocity file name, 135, 294
 Prescribed velocity boundary indicators, 79, 81, 135, 191, 193, 240, 288

Tangential velocity boundary indicators, 61, 64, 74, 79, 81, 94, 118, 124, 135, 186, 189, 201, 210, 289
 Zero velocity boundary indicators, 98, 118, 124, 289
 Burstedde benchmark
 Viscosity parameter, 194
 CFL number, 56, 94, 211, 256
 Checkpointing
 Steps between checkpoint, 75, 126, 294
 Time between checkpoint, 294
 Composition solver tolerance, 54
 Compositional fields
 Compositional field methods, 92, 294
 List of normalized fields, 295
 Mapped particle properties, 92, 295
 Names of fields, 92, 142, 295
 Number of fields, 83, 92, 98, 109, 142, 202, 295
 Dimension, 12, 42, 44, 45, 56, 61, 63, 74, 80, 97, 117, 123, 185, 188, 190, 200, 222, 256
 Discretization
 Composition polynomial degree, 54, 295
 Stabilization parameters
 alpha, 298
 beta, 299
 cR, 299
 Discontinuous penalty, 297
 gamma, 299
 Global composition maximum, 297
 Global composition minimum, 297
 Global temperature maximum, 297
 Global temperature minimum, 298
 Use artificial viscosity smoothing, 109, 298
 Use limiter for discontinuous composition solution, 298
 Use limiter for discontinuous temperature solution, 298
 Stokes velocity polynomial degree, 73, 186, 187, 189, 192, 296
 Temperature polynomial degree, 54, 73, 109, 296
 Use discontinuous composition discretization, 296
 Use discontinuous temperature discretization, 296
 Use locally conservative discretization, 100, 186, 189, 192, 296
 End time, 42, 56, 63, 80, 97, 117, 123, 185, 188, 190, 200, 211, 222, 257
 Formulation
 Enable additional Stokes RHS, 299
 Formulation, 65, 300
 Mass conservation, 300
 Temperature equation, 300
 Free surface
 Additional tangential mesh velocity boundary indicators, 300
 Free surface boundary indicators, 94, 301
 Free surface stabilization theta, 94, 301
 Surface velocity projection, 301
 Geometry model
 Box
 Box origin X coordinate, 303
 Box origin Y coordinate, 303
 Box origin Z coordinate, 303
 X extent, 61, 63, 74, 80, 98, 142, 186, 189, 191, 201, 303
 X periodic, 303
 X repetitions, 142, 303
 Y extent, 61, 63, 74, 80, 98, 142, 186, 189, 191, 201, 304
 Y periodic, 304
 Y repetitions, 304
 Z extent, 74, 201, 304
 Z periodic, 304
 Z repetitions, 304
 Box with lithosphere boundary indicators
 Box origin X coordinate, 304
 Box origin Y coordinate, 305
 Box origin Z coordinate, 305
 Lithospheric thickness, 305
 X extent, 305
 X periodic, 305
 X periodic lithosphere, 305
 X repetitions, 305
 Y extent, 306
 Y periodic, 306
 Y periodic lithosphere, 306
 Y repetitions, 306
 Y repetitions lithosphere, 306
 Z extent, 306
 Z periodic, 306
 Z repetitions, 307
 Z repetitions lithosphere, 307
 Chunk
 Chunk inner radius, 307
 Chunk maximum latitude, 307
 Chunk maximum longitude, 307
 Chunk minimum latitude, 307
 Chunk minimum longitude, 308

Chunk outer radius, 308
 Latitude repetitions, 308
 Longitude repetitions, 308
 Radius repetitions, 308
 Ellipsoidal chunk
 Depth, 308
 Depth subdivisions, 308
 East-West subdivisions, 309
 Eccentricity, 309
 NE corner, 309
 North-South subdivisions, 309
 NW corner, 309
 SE corner, 309
 Semi-major axis, 310
 SW corner, 310
 Initial topography model
 Ascii data model/Data directory, 310
 Ascii data model/Data file name, 311
 Ascii data model/Scale factor, 311
 Model name, 310
 Prm polygon/Topography parameters, 311
 Model name, 42, 61, 63, 74, 80, 98, 118, 124, 142, 186, 187, 189, 191, 201, 209, 235, 301
 Sphere
 Radius, 311
 Spherical shell
 Cells along circumference, 312
 Inner radius, 118, 124, 209, 312
 Opening angle, 118, 209, 312
 Outer radius, 118, 124, 209, 312
 Gravity model
 Ascii data model
 Data directory, 313
 Data file name, 313
 Scale factor, 314
 Function
 Function constants, 314
 Function expression, 314
 Variable names, 314
 Model name, 64, 75, 81, 98, 119, 126, 186, 187, 189, 192, 194, 201, 210, 239, 313
 Radial constant
 Magnitude, 210, 315
 Radial linear
 Magnitude at bottom, 315
 Magnitude at surface, 315
 Vertical
 Magnitude, 64, 71, 75, 98, 201, 315
 Heating model
 Adiabatic heating

Use simplified adiabatic heating, 11, 317
 Adiabatic heating of melt
 Use simplified adiabatic heating, 317
 Compositional heating
 Compositional heating values, 317
 Use compositional field for heat production averaging, 317
 Constant heating
 Radiogenic heating rate, 317
 Exponential decay heating
 Half life, 223
 Function
 Function constants, 318
 Function expression, 318
 Variable names, 318
 Latent heat melt
 Melting entropy change, 319
 List of model names, 118, 122, 123, 223, 316
 Model name, 234
 Radioactive decay
 Crust composition number, 319
 Crust defined by composition, 319
 Crust depth, 319
 Half decay times, 319
 Heating rates, 319
 Initial concentrations crust, 320
 Initial concentrations mantle, 320
 Number of elements, 320
 Initial composition model
 Ascii data model
 Data directory, 321
 Data file name, 322
 Scale factor, 322
 Function
 Coordinate system, 322
 Function constants, 178, 179, 202, 322
 Function expression, 83, 87, 99, 109, 143, 178, 179, 202, 222, 323
 Variable names, 83, 87, 99, 109, 143, 178, 179, 202, 222, 323
 List of model names, 320
 List of model operators, 321
 Model name, 83, 87, 99, 109, 142, 178, 179, 202, 222, 321
 Initial conditions
 Model name, 240
 Initial temperature model
 Adiabatic
 Age bottom boundary layer, 327
 Age top boundary layer, 327
 Amplitude, 327

Function/Function constants, 328
 Function/Function expression, 328
 Function/Variable names, 329
 Position, 327
 Radius, 328
 Subadiabaticity, 328
 Adiabatic boundary
 Adiabatic temperature gradient, 329
 Data directory, 329
 Isotherm depth filename, 329
 Isotherm temperature, 329
 Surface temperature, 329
 Ascii data model
 Data directory, 330
 Data file name, 330
 Scale factor, 330
 Function
 Coordinate system, 330
 Function constants, 58, 64, 74, 143, 330
 Function expression, 58, 64, 74, 81, 94, 98, 126, 143, 201, 222, 331
 Variable names, 58, 64, 74, 81, 94, 143, 222, 331
 Harmonic perturbation
 Lateral wave number one, 331
 Lateral wave number two, 331
 Magnitude, 332
 Reference temperature, 332
 Vertical wave number, 332
 Inclusion shape perturbation
 Ambient temperature, 332
 Center X, 332
 Center Y, 332
 Center Z, 333
 Inclusion gradient, 333
 Inclusion shape, 333
 Inclusion temperature, 333
 Shape radius, 333
 List of model names, 323
 List of model operators, 325
 Model name, 58, 63, 74, 81, 94, 98, 119, 126, 130, 143, 186, 189, 192, 201, 209, 222, 325
 S4ORTS perturbation
 Data directory, 130, 333
 Initial condition file name, 130, 334
 Maximum order, 334
 Reference temperature, 130, 334
 Remove degree 0 from perturbation, 130, 334
 Remove temperature heterogeneity down to specified depth, 334
 Specify a lower maximum order, 334
 Spline knots depth file name, 130, 335
 Thermal expansion coefficient in initial temperature scaling, 130, 335
 Vs to density scaling, 130, 335
 SAVANI perturbation
 Data directory, 335
 Initial condition file name, 335
 Maximum order, 335
 Reference temperature, 336
 Remove degree 0 from perturbation, 336
 Remove temperature heterogeneity down to specified depth, 336
 Specify a lower maximum order, 336
 Spline knots depth file name, 336
 Thermal expansion coefficient in initial temperature scaling, 336
 Vs to density scaling, 337
 Spherical gaussian perturbation
 Amplitude, 337
 Angle, 337
 Filename for initial geotherm table, 337
 Non-dimensional depth, 337
 Sigma, 337
 Sign, 337
 Spherical hexagonal perturbation
 Angular mode, 209, 338
 Rotation offset, 209, 338
 Linear solver A block tolerance, 54
 Linear solver S block tolerance, 54
 Linear solver tolerance, 54
 Material model
 Ascii reference profile
 Ascii data model/Data directory, 347
 Ascii data model/Data file name, 347
 Ascii data model/Scale factor, 347
 Thermal conductivity, 346
 Thermal viscosity exponent, 346
 Transition depths, 346
 Use TALA, 346
 Viscosity, 346
 Viscosity prefactors, 346
 Averaging
 Averaging operation, 347
 Base model, 348
 Bell shape limit, 348
 Compositing
 Compressibility, 348
 Density, 348

- Entropy derivative pressure, 348
- Entropy derivative temperature, 349
- Reaction terms, 349
- Specific heat, 349
- Thermal conductivity, 349
- Thermal expansion coefficient, 350
- Viscosity, 350
- Composition reaction model
 - Composition viscosity prefactor 1, 350
 - Composition viscosity prefactor 2, 350
 - Density differential for compositional field 1, 87, 350
 - Density differential for compositional field 2, 87, 351
 - Reaction depth, 87, 351
 - Reference density, 351
 - Reference specific heat, 351
 - Reference temperature, 351
 - Thermal conductivity, 87, 352
 - Thermal expansion coefficient, 87, 352
 - Thermal viscosity exponent, 352
 - Viscosity, 87, 352
- Depth dependent model
 - Base model, 352
 - Data directory, 352
 - Depth dependence method, 353
 - Depth list, 353
 - Viscosity depth file, 353
 - Viscosity depth function/Function constants, 353
 - Viscosity depth function/Function expression, 354
 - Viscosity depth function/Variable names, 354
 - Viscosity list, 353
- Diffusion dislocation
 - Activation energies for diffusion creep, 354
 - Activation energies for dislocation creep, 354
 - Activation volumes for diffusion creep, 354
 - Activation volumes for dislocation creep, 355
 - Densities, 355
 - Effective viscosity coefficient, 355
 - Grain size, 355
 - Grain size exponents for diffusion creep, 355
 - Heat capacity, 355
 - Maximum strain rate ratio iterations, 356
 - Maximum viscosity, 356
 - Minimum strain rate, 356
- Minimum viscosity, 356
- Prefactors for diffusion creep, 356
- Prefactors for dislocation creep, 356
- Reference temperature, 356
- Reference viscosity, 357
- Strain rate residual tolerance, 357
- Stress exponents for diffusion creep, 357
- Stress exponents for dislocation creep, 357
- Thermal diffusivity, 357
- Thermal expansivities, 357
- Viscosity averaging scheme, 358
- Drucker Prager
 - Reference density, 358
 - Reference specific heat, 358
 - Reference temperature, 358
 - Reference viscosity, 358
 - Thermal conductivity, 358
 - Thermal expansion coefficient, 358
 - Viscosity/Angle of internal friction, 359
 - Viscosity/Cohesion, 359
 - Viscosity/Maximum viscosity, 359
 - Viscosity/Minimum viscosity, 359
 - Viscosity/Reference strain rate, 359
- Dynamic Friction
 - Densities, 359
 - Reference temperature, 360
 - Specific heats, 360
 - Thermal conductivities, 360
 - Thermal expansivities, 360
 - Viscosities/Background Viscosities, 361
 - Viscosities/Coefficients of dynamic friction, 361
 - Viscosities/Coefficients of static friction, 361
 - Viscosities/Cohesions, 361
 - Viscosities/Maximum viscosity, 361
 - Viscosities/Minimum viscosity, 361
 - Viscosities/Reference strain rate, 362
 - Viscosity averaging scheme, 360
- Exponential decay
 - Half life, 223
- Grain size model
 - Advect logarithm of grain size, 362
 - Average specific grain boundary energy, 362
 - Bilinear interpolation, 362
 - Data directory, 362
 - Derivatives file names, 363
 - Diffusion activation energy, 363
 - Diffusion activation volume, 363
 - Diffusion creep exponent, 363

Diffusion creep grain size exponent, 363
 Diffusion creep prefactor, 363
 Dislocation activation energy, 364
 Dislocation activation volume, 364
 Dislocation creep exponent, 364
 Dislocation creep prefactor, 364
 Dislocation viscosity iteration number, 364
 Dislocation viscosity iteration threshold, 364
 Geometric constant, 365
 Grain growth activation energy, 365
 Grain growth activation volume, 365
 Grain growth exponent, 365
 Grain growth rate constant, 365
 Lower mantle grain size scaling, 365
 Material file format, 366
 Material file names, 366
 Maximum latent heat substeps, 366
 Maximum specific heat, 366
 Maximum temperature dependence of viscosity, 366
 Maximum thermal expansivity, 366
 Maximum viscosity, 367
 Minimum grain size, 367
 Minimum specific heat, 367
 Minimum thermal expansivity, 367
 Minimum viscosity, 367
 Phase transition Clapeyron slopes, 367
 Phase transition depths, 368
 Phase transition temperatures, 368
 Phase transition widths, 368
 Reciprocal required strain, 368
 Recrystallized grain size, 368
 Reference compressibility, 368
 Reference density, 369
 Reference specific heat, 369
 Reference temperature, 369
 Thermal conductivity, 369
 Thermal expansion coefficient, 369
 Use enthalpy for material properties, 369
 Use paleowattmeter, 369
 Use table properties, 370
 Viscosity, 370
 Work fraction for boundary area change, 370
 Inclusion
 Viscosity jump, 191
 Latent heat
 Composition viscosity prefactor, 207, 370
 Compressibility, 207, 370
 Corresponding phase for density jump, 206, 370
 Define transition by depth instead of pressure, 371
 Density differential for compositional field 1, 371
 Maximum viscosity, 371
 Minimum viscosity, 371
 Phase transition Clapeyron slopes, 206, 371
 Phase transition density jumps, 206, 372
 Phase transition depths, 206, 372
 Phase transition pressure widths, 372
 Phase transition pressures, 372
 Phase transition temperatures, 206, 372
 Phase transition widths, 206, 373
 Reference density, 206, 373
 Reference specific heat, 206, 373
 Reference temperature, 206, 373
 Thermal conductivity, 206, 373
 Thermal expansion coefficient, 207, 373
 Thermal viscosity exponent, 207, 373
 Viscosity, 207, 374
 Viscosity prefactors, 207, 374
 Latent heat melt
 A1, 374
 A2, 374
 A3, 374
 B1, 374
 B2, 375
 B3, 375
 beta, 378
 C1, 375
 C2, 375
 C3, 375
 Composition viscosity prefactor, 375
 Compressibility, 375
 D1, 376
 D2, 376
 D3, 376
 Density differential for compositional field 1, 376
 E1, 376
 E2, 376
 Mass fraction cpx, 377
 Maximum pyroxenite melt fraction, 377
 Peridotite melting entropy change, 377
 Pyroxenite melting entropy change, 377
 r1, 379
 r2, 379
 Reference density, 377
 Reference specific heat, 377

- Reference temperature, 377
 Relative density of melt, 378
 Thermal conductivity, 378
 Thermal expansion coefficient, 378
 Thermal expansion coefficient of melt, 378
 Thermal viscosity exponent, 378
 Viscosity, 378
 Material averaging, 55, 98, 99, 338
 Melt global
 Depletion density change, 379
 Depletion solidus change, 379
 Exponential melt weakening factor, 379
 Include melting and freezing, 379
 Melt bulk modulus derivative, 380
 Melt compressibility, 380
 Melting time scale for operator splitting, 380
 Pressure solidus change, 380
 Reference bulk viscosity, 380
 Reference melt density, 380
 Reference melt viscosity, 381
 Reference permeability, 381
 Reference shear viscosity, 381
 Reference solid density, 381
 Reference specific heat, 381
 Reference temperature, 381
 Solid compressibility, 381
 Surface solidus, 382
 Thermal bulk viscosity exponent, 382
 Thermal conductivity, 382
 Thermal expansion coefficient, 382
 Thermal viscosity exponent, 382
 Melt simple
 A1, 382
 A2, 383
 A3, 383
 B1, 383
 B2, 383
 B3, 383
 beta, 388
 C1, 383
 C2, 383
 C3, 384
 Depletion density change, 384
 Depletion solidus change, 384
 Exponential melt weakening factor, 384
 Freezing rate, 384
 Mass fraction cpx, 385
 Melt bulk modulus derivative, 385
 Melt compressibility, 385
 Melt extraction depth, 385
 Melting time scale for operator splitting, 385
 Peridotite melting entropy change, 386
 r1, 388
 r2, 388
 Reference bulk viscosity, 386
 Reference melt density, 386
 Reference melt viscosity, 386
 Reference permeability, 386
 Reference shear viscosity, 386
 Reference solid density, 386
 Reference specific heat, 387
 Reference temperature, 387
 Solid compressibility, 387
 Thermal bulk viscosity exponent, 387
 Thermal conductivity, 387
 Thermal expansion coefficient, 387
 Thermal viscosity exponent, 387
 Use fractional melting, 388
 Use full compressibility, 388
 Model name, 56, 65, 81, 85, 87, 96, 98, 111, 117, 124, 143, 178, 186, 189, 191, 194, 201, 206, 209, 212, 223, 232, 338
 Morency and Doin
 Activation energies, 143
 Activation volume, 143
 Coefficient of yield stress increase with depth, 143
 Cohesive strength of rocks at the surface, 143
 Densities, 143
 Heat capacity, 143
 Minimum strain rate, 143
 Preexponential constant for viscous rheology law, 143
 Reference strain rate, 143
 Reference temperature, 143
 Stress exponents for plastic rheology, 143
 Stress exponents for viscous rheology, 143
 Thermal diffusivity, 143
 Thermal expansivities, 143
 Multicomponent
 Densities, 389
 Reference temperature, 389
 Specific heats, 389
 Thermal conductivities, 389
 Thermal expansivities, 389
 Viscosities, 389
 Viscosity averaging scheme, 390
 Nondimensional model
 Di, 390

gamma, 391
 Ra, 390
 Reference density, 390
 Reference specific heat, 390
 Use TALA, 390
 Viscosity depth prefactor, 391
 Viscosity temperature prefactor, 391
 Simple compressible model
 Reference compressibility, 391
 Reference density, 391
 Reference specific heat, 391
 Thermal conductivity, 391
 Thermal expansion coefficient, 392
 Viscosity, 392
 Simple model
 Composition viscosity prefactor, 98, 392
 Density differential for compositional field
 1, 86, 98, 178, 202, 392
 Maximum thermal prefactor, 392
 Minimum thermal prefactor, 392
 Reference density, 56, 65, 86, 98, 111, 201,
 209, 393
 Reference specific heat, 65, 209, 393
 Reference temperature, 56, 65, 86, 111,
 209, 393
 Thermal conductivity, 65, 81, 85, 111, 209,
 393
 Thermal expansion coefficient, 65, 81, 85,
 98, 111, 117, 124, 178, 209, 393
 Thermal viscosity exponent, 111, 393
 Viscosity, 56, 65, 81, 86, 98, 111, 117, 124,
 178, 201, 209, 393
 Simpler model
 Reference density, 394
 Reference specific heat, 394
 Reference temperature, 394
 Thermal conductivity, 394
 Thermal expansion coefficient, 394
 Viscosity, 394
 Simpler with crust model
 Jump height, 96
 Lower viscosity, 96
 Reference density, 96
 Reference specific heat, 96
 Reference temperature, 96
 Thermal conductivity, 96
 Thermal expansion coefficient, 96
 Upper viscosity, 96
 SolCx
 Viscosity jump, 186
 Steinberger model

Bilinear interpolation, 395
 Data directory, 395
 Latent heat, 395
 Lateral viscosity file name, 395
 Material file names, 395
 Maximum lateral viscosity variation, 395
 Maximum viscosity, 396
 Minimum viscosity, 396
 Number lateral average bands, 396
 Radial viscosity file name, 396
 Reference viscosity, 396
 Thermal conductivity, 396
 Use lateral average temperature for
 viscosity, 397
 Visco Plastic
 Activation energies for diffusion creep, 397
 Activation energies for dislocation creep,
 397
 Activation volumes for diffusion creep, 397
 Activation volumes for dislocation creep,
 397
 Angles of internal friction, 398
 Cohesion strain weakening factors, 398
 Cohesions, 398
 Densities, 398
 End strain weakening intervals, 398
 Friction strain weakening factors, 399
 Grain size, 399
 Grain size exponents for diffusion creep,
 399
 Heat capacities, 399
 Maximum viscosity, 399
 Minimum strain rate, 399
 Minimum viscosity, 400
 Prefactors for diffusion creep, 400
 Prefactors for dislocation creep, 400
 Reference strain rate, 400
 Reference temperature, 400
 Reference viscosity, 400
 Start strain weakening intervals, 400
 Stress exponents for diffusion creep, 401
 Stress exponents for dislocation creep, 401
 Stress limiter exponents, 401
 Thermal diffusivities, 401
 Thermal expansivities, 401
 Use finite strain tensor, 402
 Use strain weakening, 402
 Viscosity averaging scheme, 402
 Viscous flow law, 402
 Viscous strain weakening factors, 402
 Yield mechanism, 402

VoT model
 Reference density, 212
 Reference specific heat, 212
 Reference temperature, 212
 Thermal conductivity, 212
 Thermal expansion coefficient, 212
 Viscosity, 212
 Max nonlinear iterations, 257
 Max nonlinear iterations in pre-refinement, 257
 Maximum relative increase in time step, 257
 Maximum time step, 222, 257
 Melt settings
 Average melt velocity, 403
 Heat advection by melt, 403
 Include melt transport, 403
 Melt scaling factor threshold, 403
 Use discontinuous compaction pressure, 404
 Mesh refinement, 247
 Adapt by fraction of cells, 404
 Additional refinement times, 71, 75, 404
 Artificial viscosity
 Compositional field scaling factors, 410
 Temperature scaling factor, 410
 Boundary
 Boundary refinement indicators, 410
 Coarsening fraction, 71, 404
 Compaction length
 Mesh cells per compaction length, 410
 Composition
 Compositional field scaling factors, 411
 Composition approximate gradient
 Compositional field scaling factors, 411
 Composition gradient
 Compositional field scaling factors, 411
 Composition threshold
 Compositional field thresholds, 412
 Initial adaptive refinement, 42, 56, 65, 71, 75, 82, 99, 119, 126, 187, 189, 192, 202, 404
 Initial global refinement, 42, 56, 65, 71, 75, 82, 99, 119, 126, 187, 189, 192, 202, 404
 Maximum refinement function
 Coordinate system, 412
 Function constants, 412
 Function expression, 412
 Variable names, 412
 Minimum refinement function
 Coordinate system, 413
 Function constants, 413
 Function expression, 413
 Variable names, 413
 Minimum refinement level, 405
 Normalize individual refinement criteria, 405
 Refinement criteria merge operation, 405
 Refinement criteria scaling factors, 405
 Refinement fraction, 71, 202, 406
 Run postprocessors on initial refinement, 406
 Strategy, 119, 126, 202, 247, 406
 Time steps between mesh refinement, 65, 71, 75, 82, 119, 126, 409
 Nonlinear solver scheme, 222, 258
 Nonlinear solver tolerance, 258
 Nullspace removal
 Remove nullspace, 210, 414
 Output directory, 42, 52, 56, 63, 80, 89, 97, 117, 123, 185, 188, 190, 200, 211, 258
 Point one, 135, 136
 Point two, 136
 Postprocess
 Command
 Command, 418
 Run on all processes, 419
 Terminate on failure, 419
 Depth average
 List of output variables, 419
 Number of zones, 419
 Output format, 126, 419
 Time between graphical output, 119, 126, 420
 Dynamic core statistics
 Excess entropy only, 420
 Dynamic topography
 Density above, 420
 Density below, 420
 Output bottom, 420
 Output surface, 421
 Geoid
 Also output the spherical harmonic coefficients of CMB dynamic topography contribution, 421
 Also output the spherical harmonic coefficients of density anomaly contribution, 421
 Also output the spherical harmonic coefficients of geoid anomaly, 421
 Also output the spherical harmonic coefficients of surface dynamic topography contribution, 421
 Density above, 421
 Density below, 422
 Maximum degree, 422

Minimum degree, 422
 Output data in geographical coordinates, 422
Global statistics
 Write statistics for each nonlinear iteration, 422
 List of postprocessors, 42, 54, 65, 75, 82, 85, 86, 88, 94, 99, 119, 126, 128, 178, 187, 190, 192, 194, 203, 242, 415
Particles
 Data output format, 89, 92, 422
 Function/Function constants, 427
 Function/Function expression, 90, 427
 Function/Number of components, 427
 Function/Variable names, 90, 428
 Generator/Ascii file/Data directory, 428
 Generator/Ascii file/Data file name, 428
 Generator/Probability density function/Function constants, 428
 Generator/Probability density function/Function expression, 429
 Generator/Probability density function/Random cell selection, 429
 Generator/Probability density function/Random number seed, 429
 Generator/Probability density function/Variable names, 429
 Generator/Reference cell/Number of particles per cell per direction, 430
 Generator/Uniform box/Maximum x, 430
 Generator/Uniform box/Maximum y, 430
 Generator/Uniform box/Maximum z, 430
 Generator/Uniform box/Minimum x, 430
 Generator/Uniform box/Minimum y, 430
 Generator/Uniform box/Minimum z, 431
 Generator/Uniform radial/Center x, 431
 Generator/Uniform radial/Center y, 431
 Generator/Uniform radial/Center z, 431
 Generator/Uniform radial/Maximum latitude, 431
 Generator/Uniform radial/Maximum longitude, 431
 Generator/Uniform radial/Maximum radius, 431
 Generator/Uniform radial/Minimum latitude, 432
 Generator/Uniform radial/Minimum longitude, 432
 Generator/Uniform radial/Minimum radius, 432
 Generator/Uniform radial/Radial layers, 432
 Integration scheme, 423
 Interpolation scheme, 92, 423
 Interpolator/Bilinear least squares/Global particle property maximum, 432
 Interpolator/Bilinear least squares/Global particle property minimum, 432
 Interpolator/Bilinear least squares/Use limiter, 433
 List of particle properties, 90, 92, 424
 Load balancing strategy, 425
 Maximum particles per cell, 425
 Melt particle/Threshold for melt presence, 433
 Minimum particles per cell, 425
 Number of particles, 88–90, 92, 425
 Particle generator name, 92, 425
 Particle weight, 426
 Time between data output, 89, 92, 426
 Update ghost particles, 427
Point values
 Evaluation points, 433
 Run postprocessors on nonlinear iterations, 418
Topography
 Output to file, 433
 Time between text output, 433
Visualization
 Artificial viscosity composition/Name of compositional field, 441
 Compositional fields as vectors/Names of fields, 441
 Compositional fields as vectors/Names of vectors, 441
 Filter output, 434
 Interpolate output, 434
 List of output variables, 86, 99, 128, 203, 245, 435
 Material properties/List of material properties, 441
 Melt fraction/A1, 442
 Melt fraction/A2, 442
 Melt fraction/A3, 442
 Melt fraction/B1, 442
 Melt fraction/B2, 442
 Melt fraction/B3, 442
 Melt fraction/beta, 444
 Melt fraction/C1, 443
 Melt fraction/C2, 443
 Melt fraction/C3, 443
 Melt fraction/D1, 443

Melt fraction/D2, 443
 Melt fraction/D3, 443
 Melt fraction/E1, 444
 Melt fraction/E2, 444
 Melt fraction/Mass fraction cpx, 444
 Melt fraction/r1, 444
 Melt fraction/r2, 444
 Melt material properties/List of properties, 445
 Number of grouped files, 52, 119, 126, 128, 439
 Output format, 46, 99, 119, 126, 128, 440
 Output mesh velocity, 440
 Temporary output location, 440
 Time between graphical output, 65, 75, 82, 86, 88, 94, 99, 119, 126, 440
 Time steps between graphical output, 440
 Vp anomaly/Average velocity scheme, 445
 Vp anomaly/Number of depth slices, 445
 Vs anomaly/Average velocity scheme, 445
 Vs anomaly/Number of depth slices, 445
 Write in background thread, 440
 Prescribe internal velocities, 107
 Prescribed Stokes solution
 Ascii data model
 Data directory, 446
 Data file name, 446
 Scale factor, 447
 Compaction pressure function
 Function constants, 447
 Function expression, 447
 Variable names, 447
 Fluid pressure function
 Function constants, 448
 Function expression, 448
 Variable names, 448
 Fluid velocity function
 Function constants, 448
 Function expression, 449
 Variable names, 449
 Model name, 446
 Pressure function
 Function constants, 449
 Function expression, 450
 Variable names, 450
 Velocity function
 Function constants, 450
 Function expression, 450
 Variable names, 451
 Prescribed velocities
 Indicator function

Function expression, 107
 Variable names, 107
 Velocity function
 Function expression, 108
 Variable names, 108
 Pressure normalization, 18, 55, 63, 97, 186–188, 190, 258
 Resume computation, 52, 56, 75, 211, 259
 Solver parameters
 AMG parameters
 AMG aggregation threshold, 451
 AMG output details, 452
 AMG smoother sweeps, 452
 AMG smoother type, 452
 Composition solver tolerance, 451
 Newton solver parameters
 Max Newton line search iterations, 452
 Max pre-Newton nonlinear iterations, 452
 Maximum linear Stokes solver tolerance, 453
 Nonlinear Newton solver switch tolerance, 453
 SPD safety factor, 453
 Stabilization preconditioner, 453
 Stabilization velocity block, 453
 Use Eisenstat Walker method for Picard iterations, 454
 Use Newton failsafe, 454
 Use Newton residual scaling method, 454
 Operator splitting parameters
 Reaction time step, 222, 454
 Reaction time steps per advection step, 455
 Stokes solver parameters
 Linear solver A block tolerance, 455
 Linear solver S block tolerance, 455
 Linear solver tolerance, 193, 455
 Maximum number of expensive Stokes solver steps, 456
 Number of cheap Stokes solver steps, 456
 Use direct solver for Stokes system, 456
 Temperature solver tolerance, 65, 451
 Start time, 80, 97, 185, 188, 190, 200, 222, 259
 Surface pressure, 18, 63, 259
 Temperature solver tolerance, 54
 Termination criteria, 247
 Checkpoint on termination, 456
 End step, 456
 Steady state velocity

Maximum relative deviation, 457
Time in steady state, 457
Termination criteria, 457
User request
File name, 457

Timing output frequency, 259
Use conduction timestep, 260
Use operator splitting, 222, 260
Use years in output instead of seconds, 15, 63,
80, 117, 123, 200, 260